

Apache Kafka

and the Rise of Stream Processing

Guozhang Wang

QCon, April 17, 2017

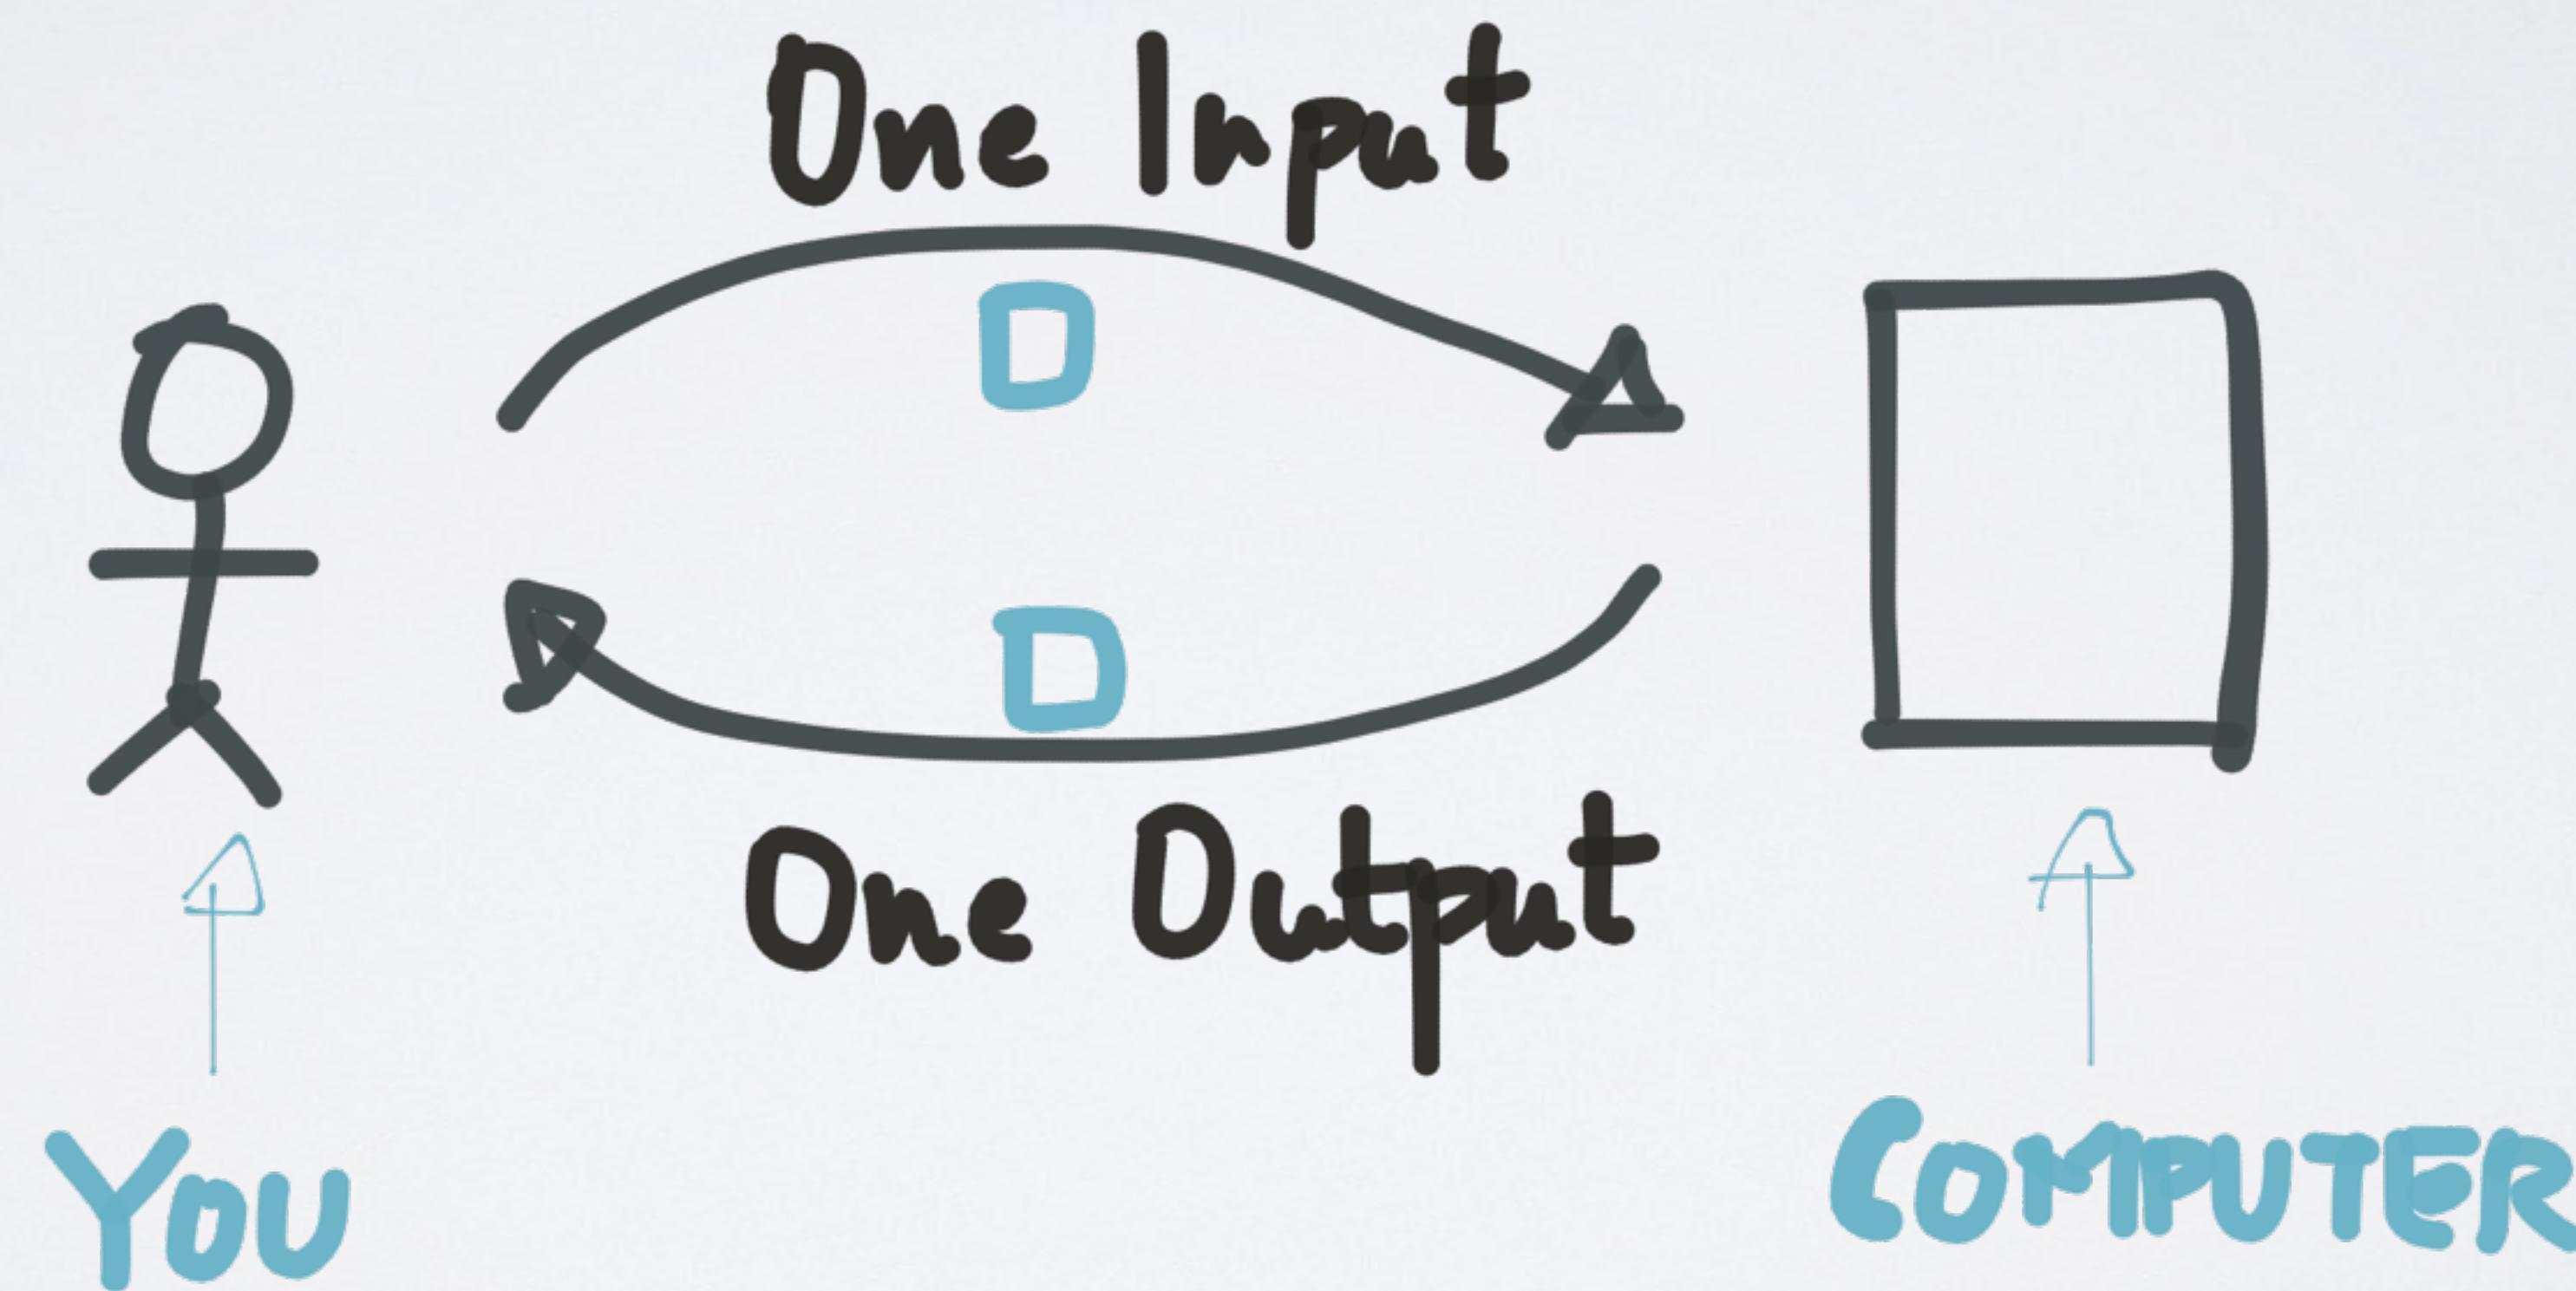


What is NOT Stream Processing?

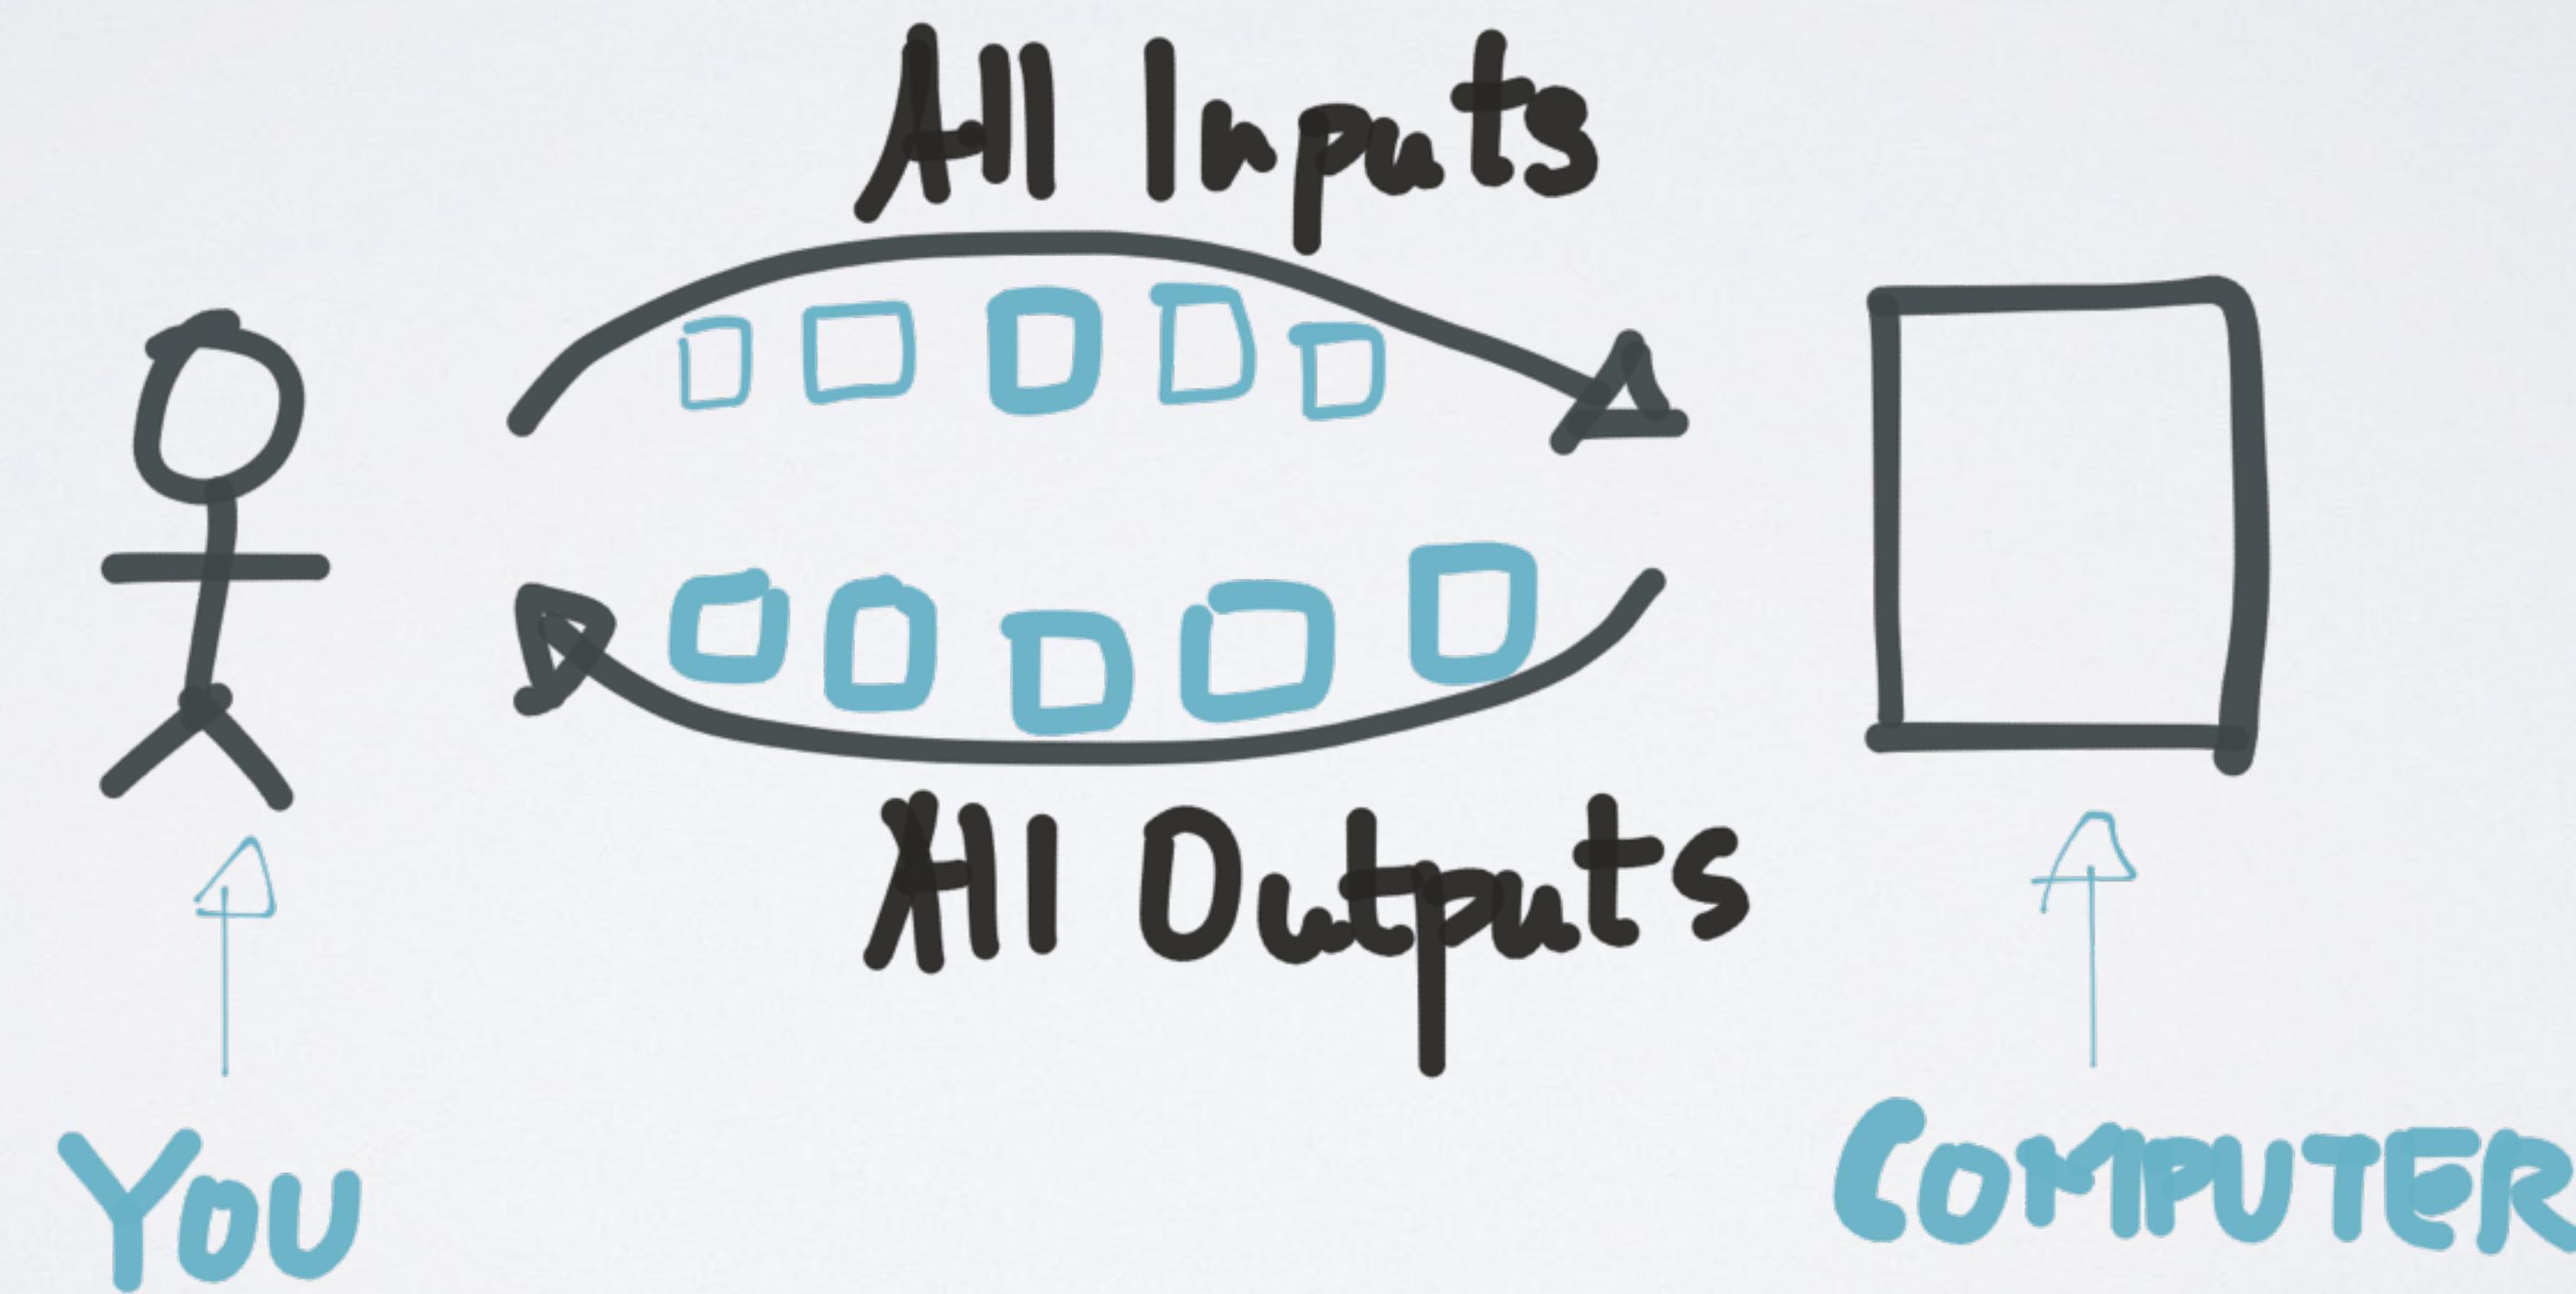
Stream Processing isn't (necessarily)

- *Transient, approximate, lossy...*
- *.. that you must have batch processing as safety net*

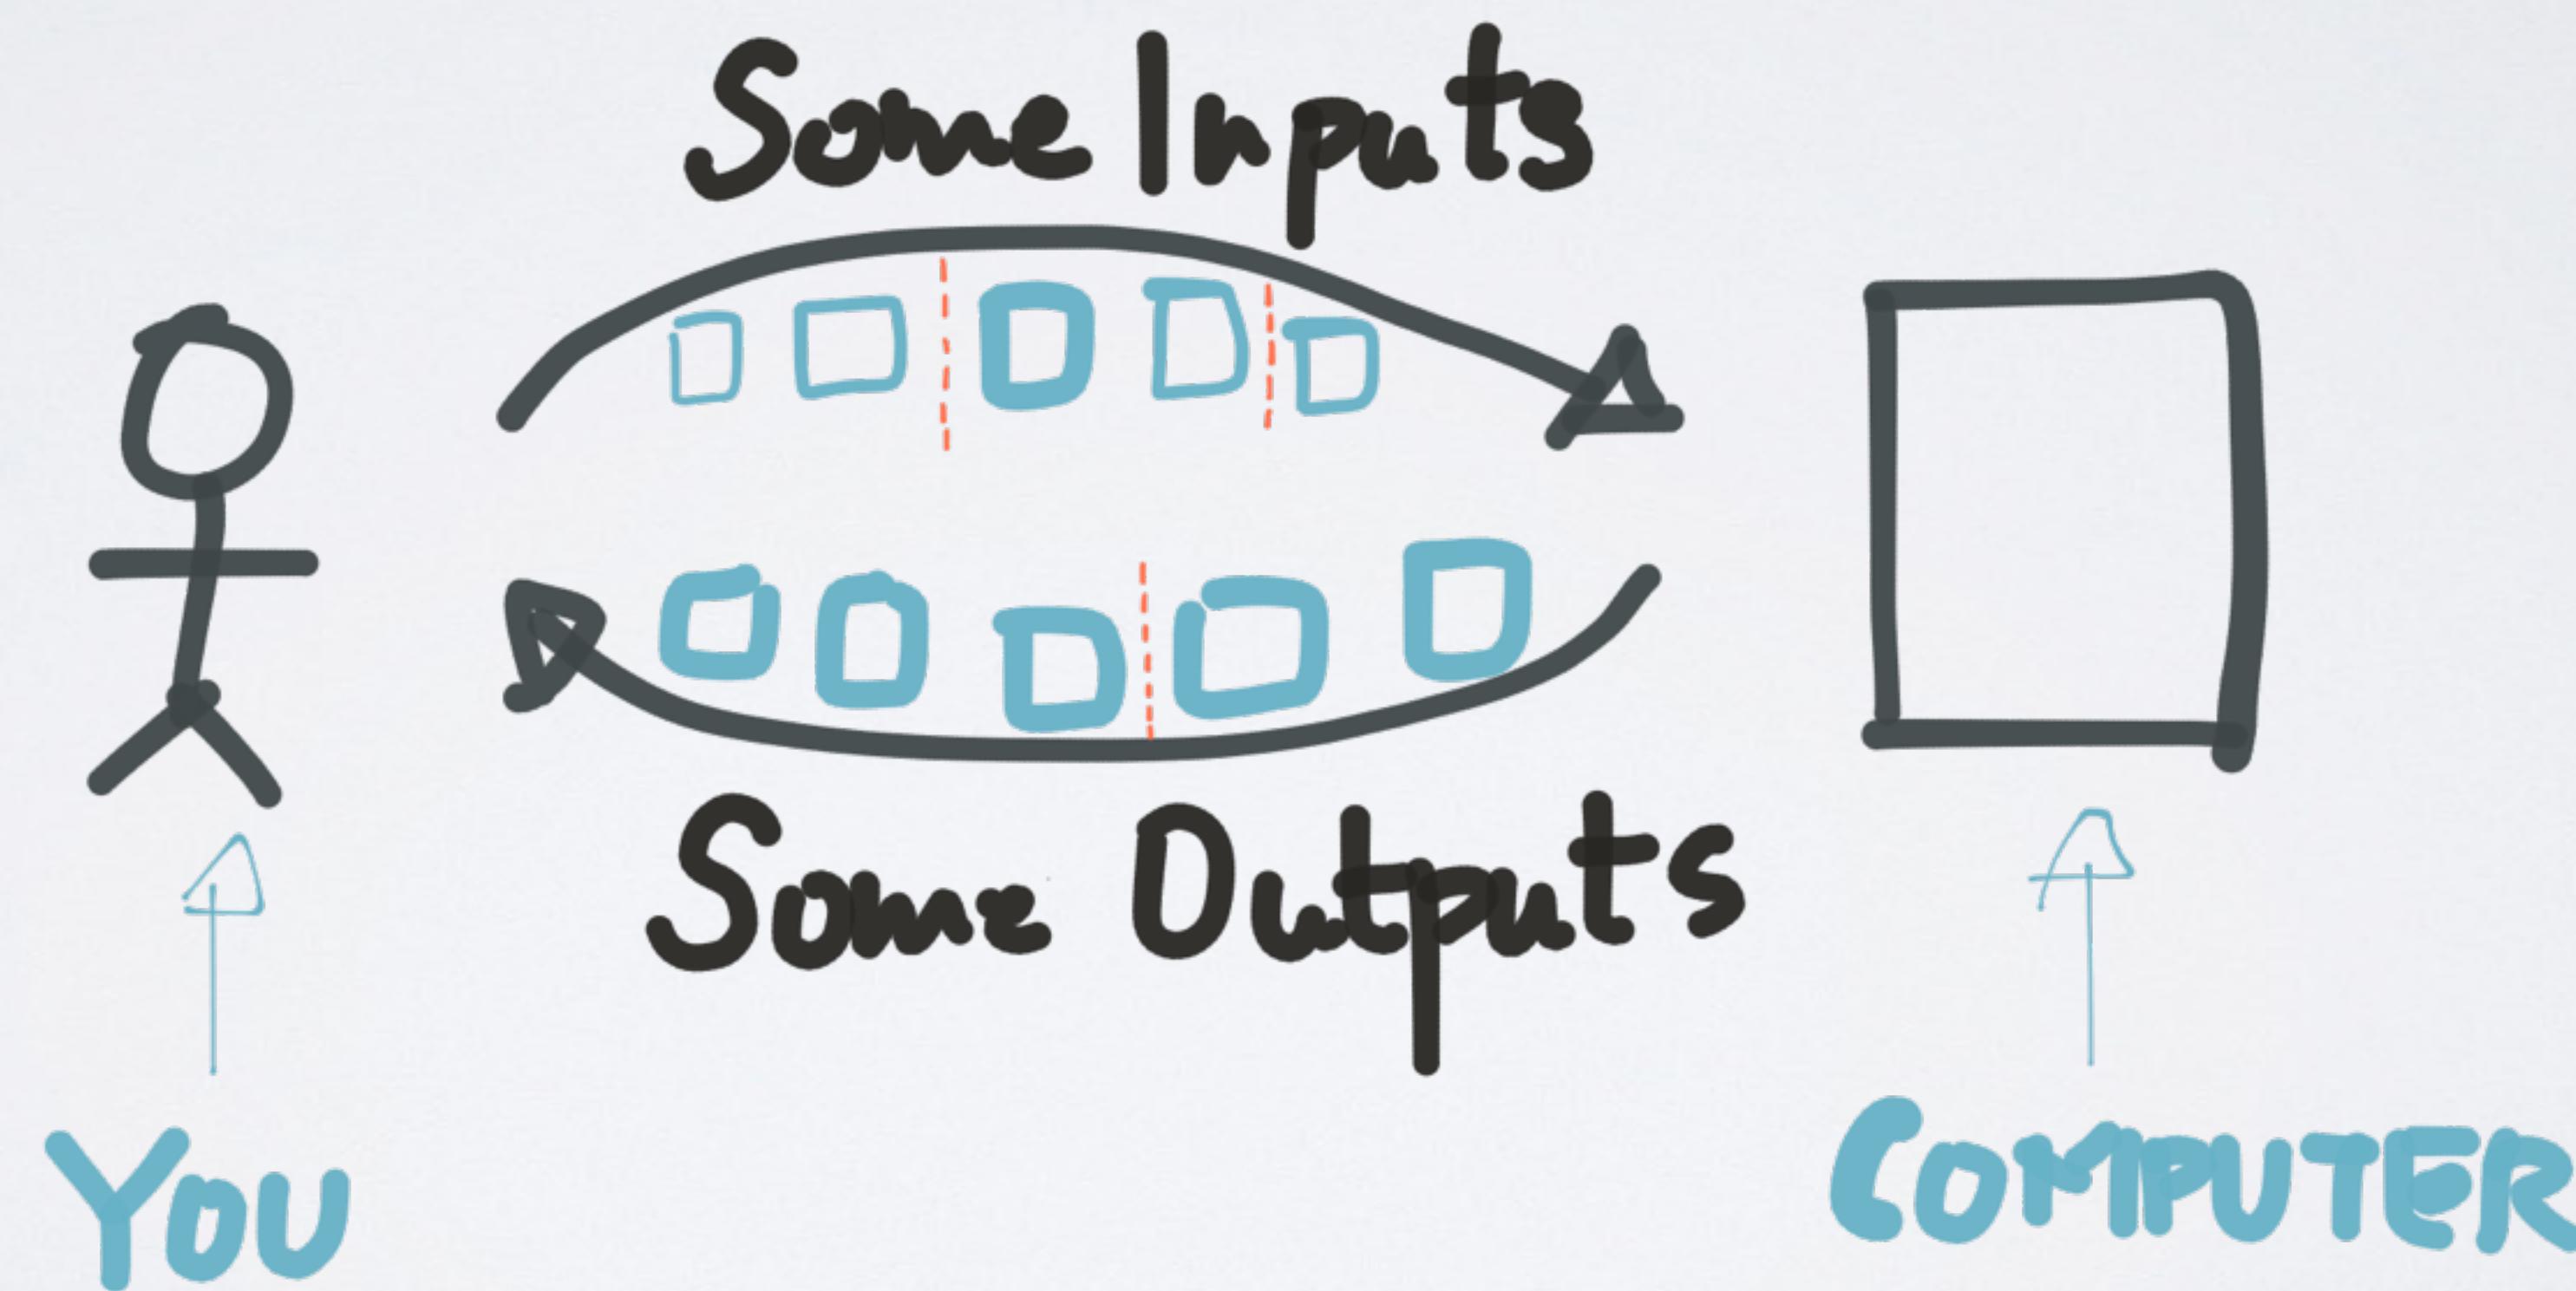
REQUEST / RESPONSE



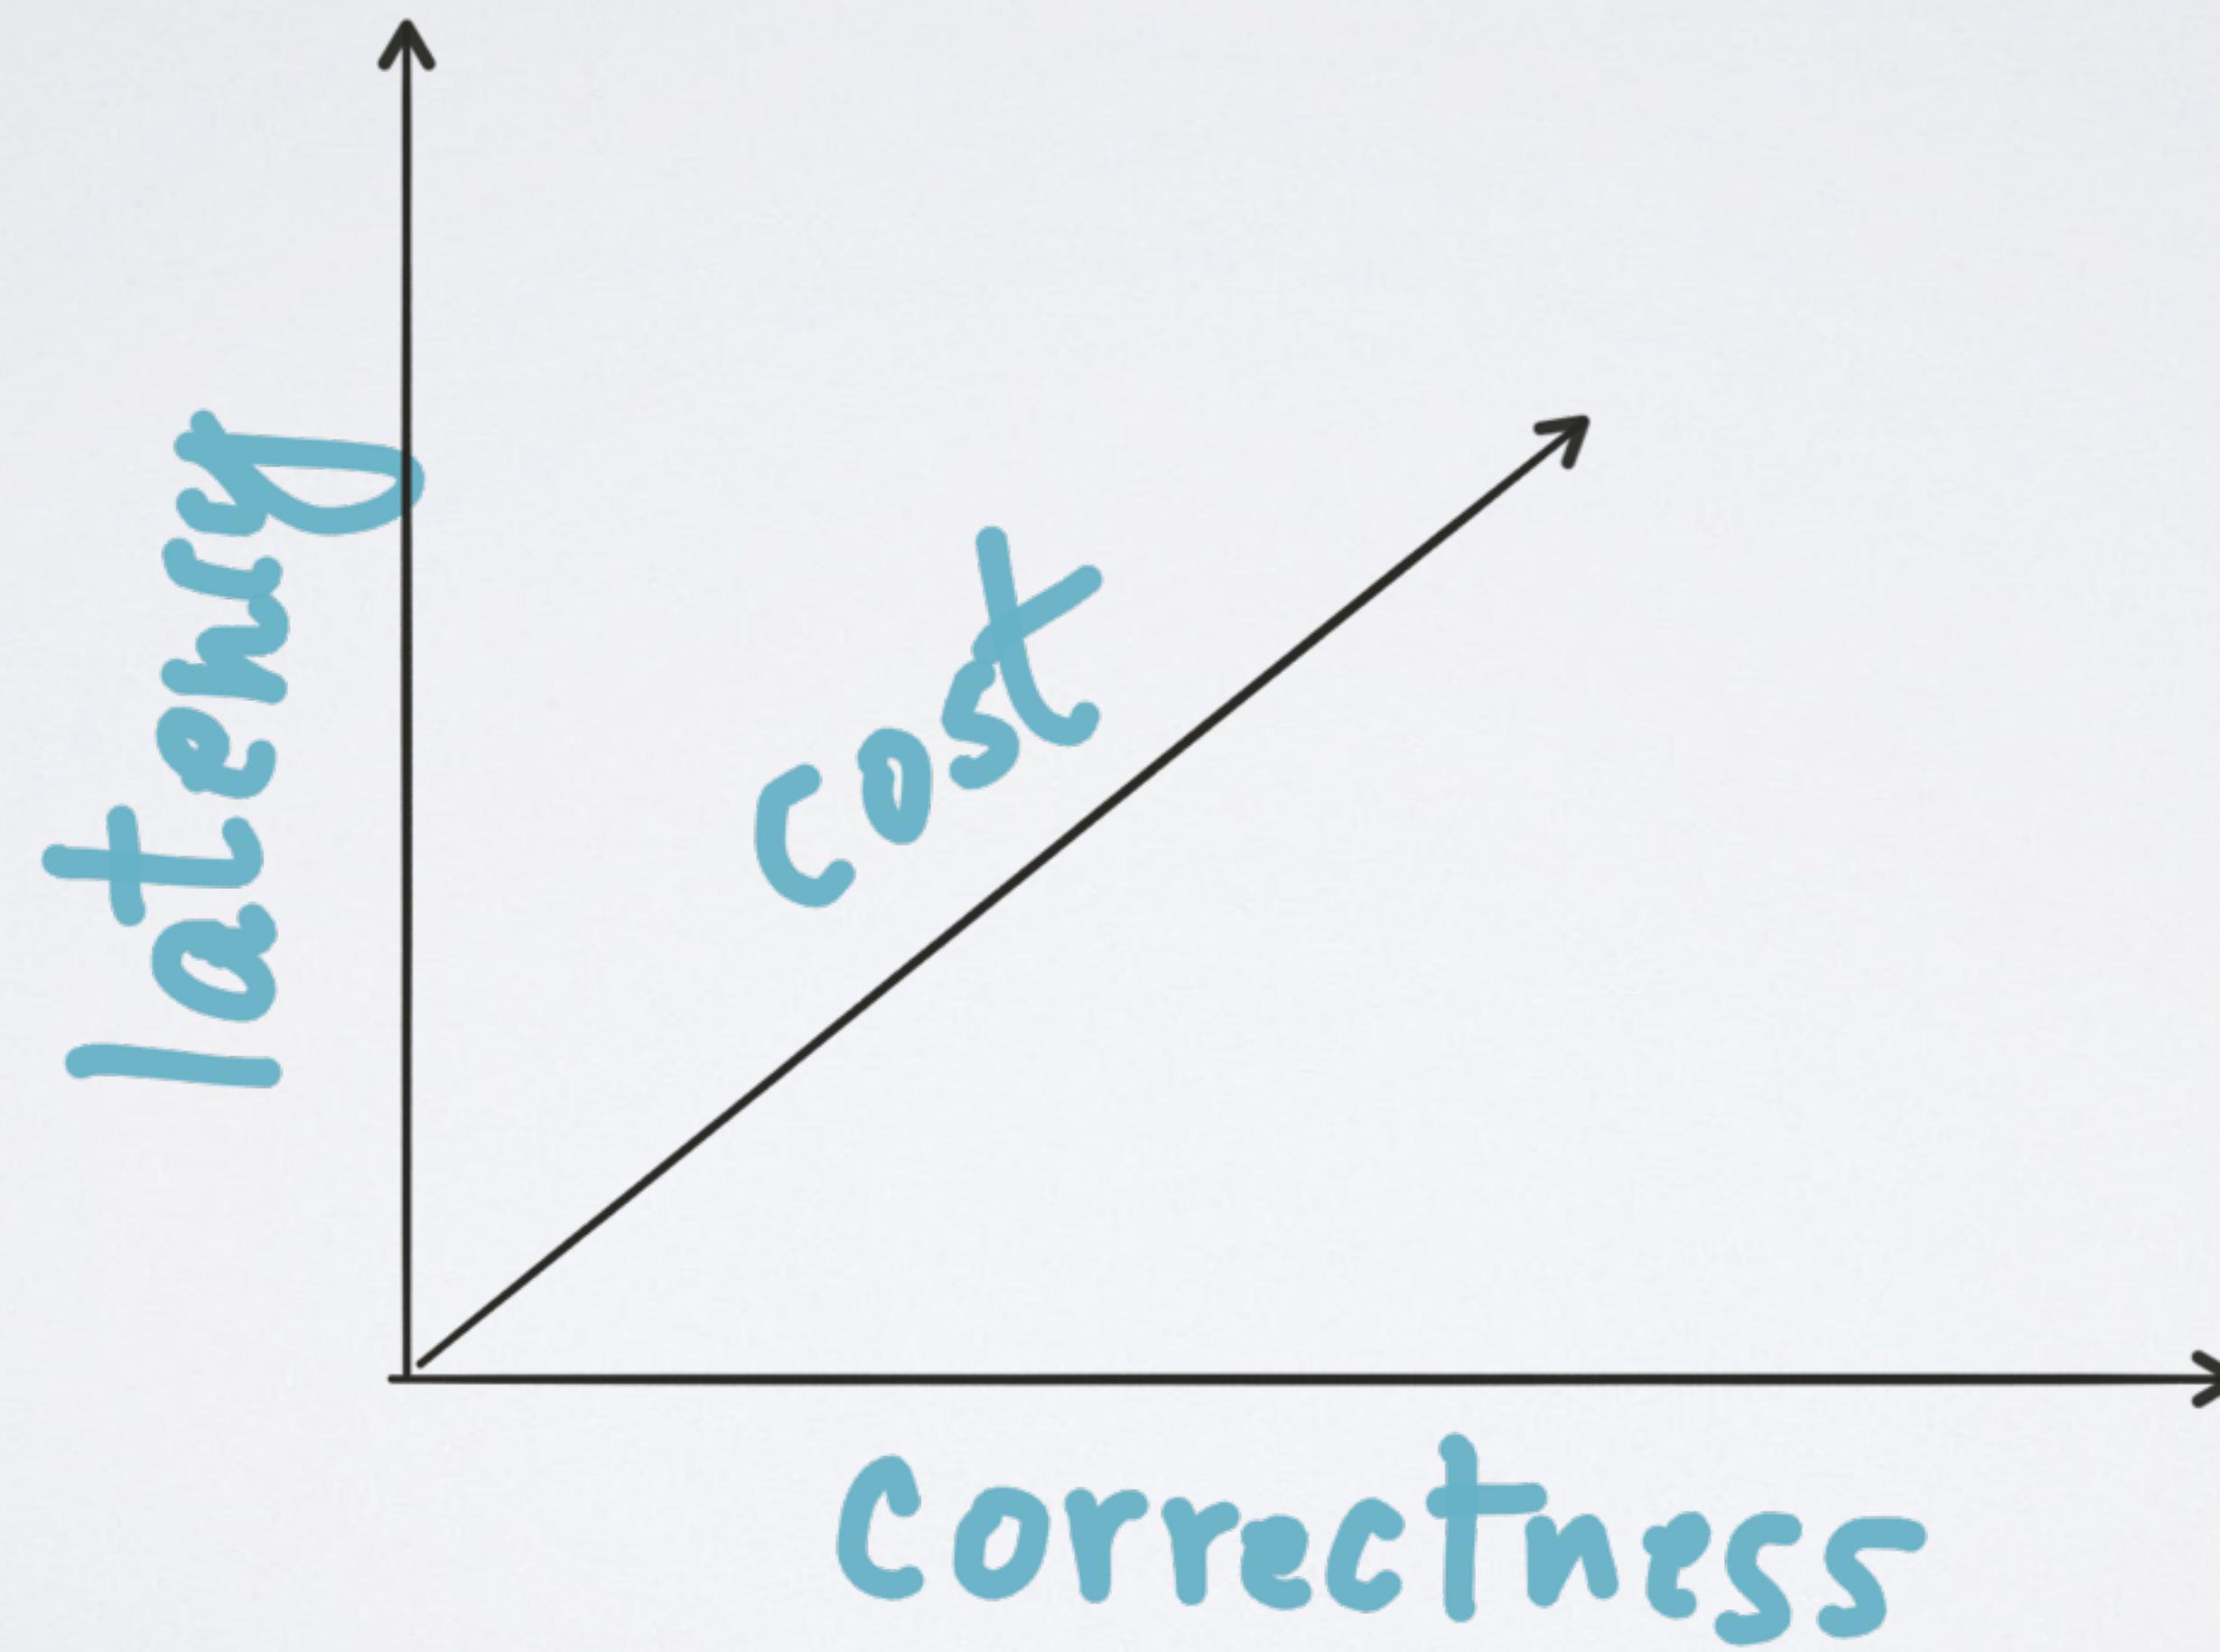
BATCH



STREAM PROCESSING



TRADE OFFS



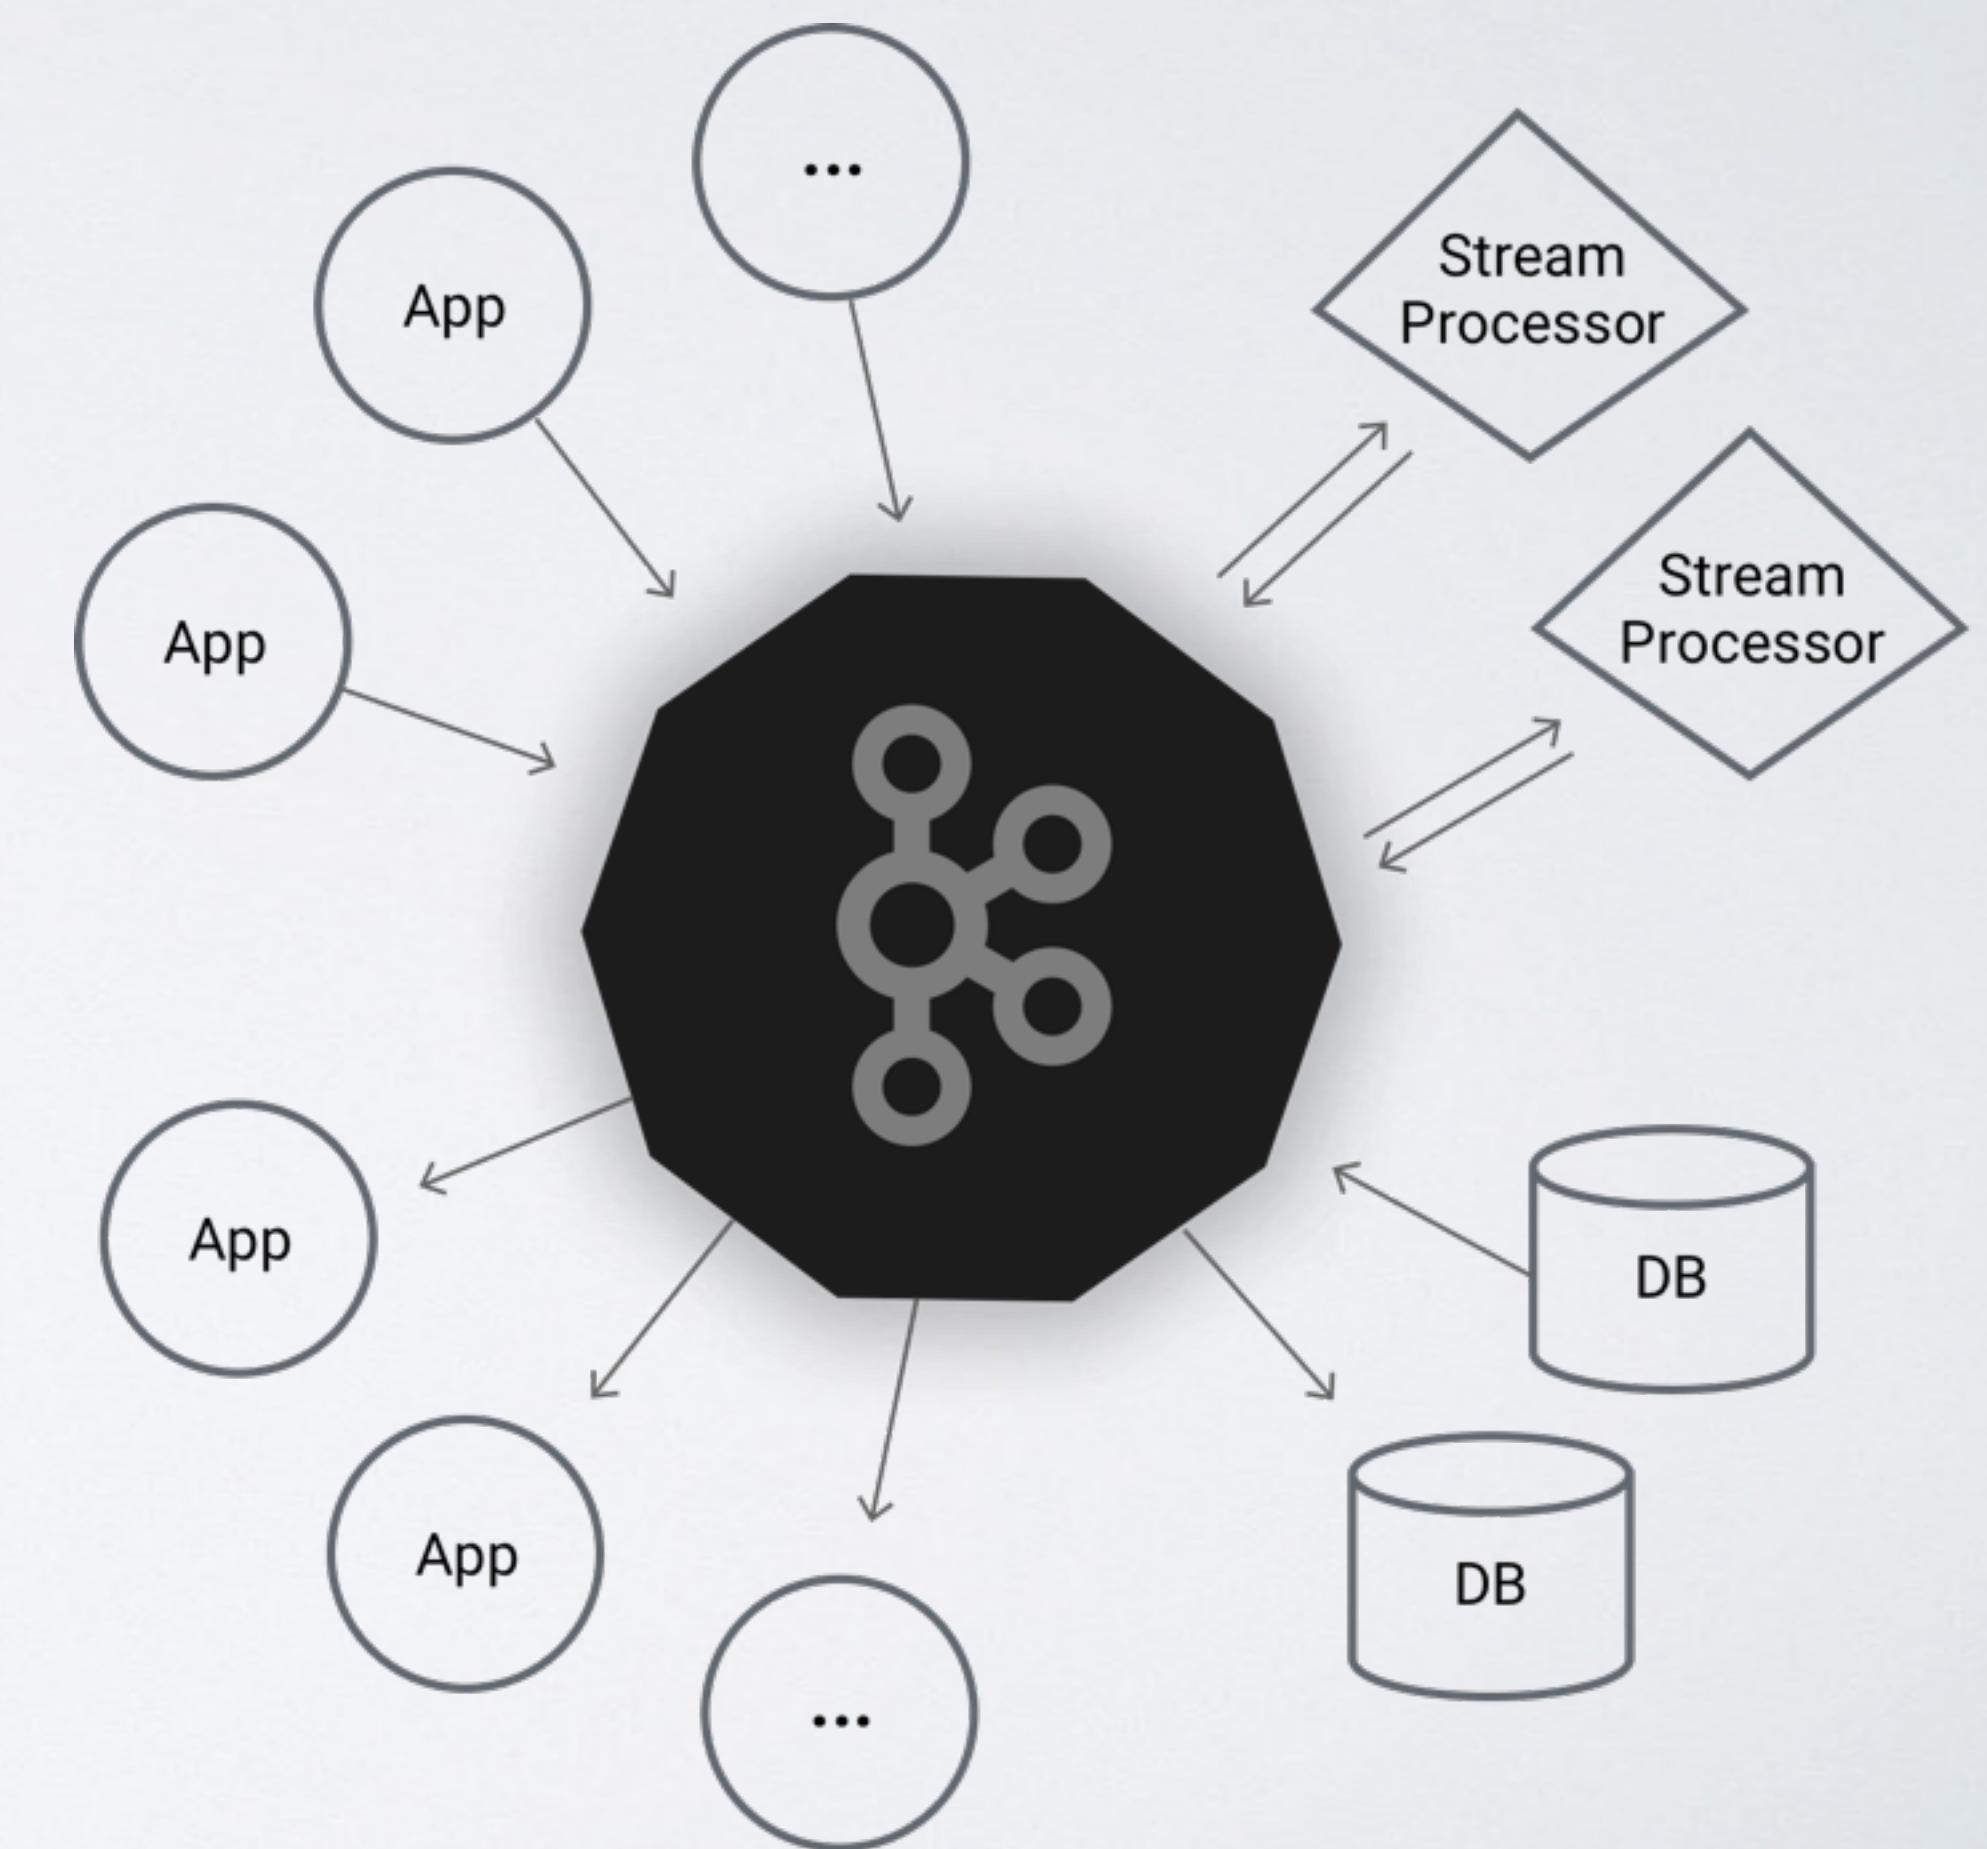
Stream Processing

- A *different programming paradigm*
- .. *that brings computation to **unbounded** data*
- .. *with tradeoffs between **latency** / **cost** / **correctness***

Why Kafka in Stream Processing?

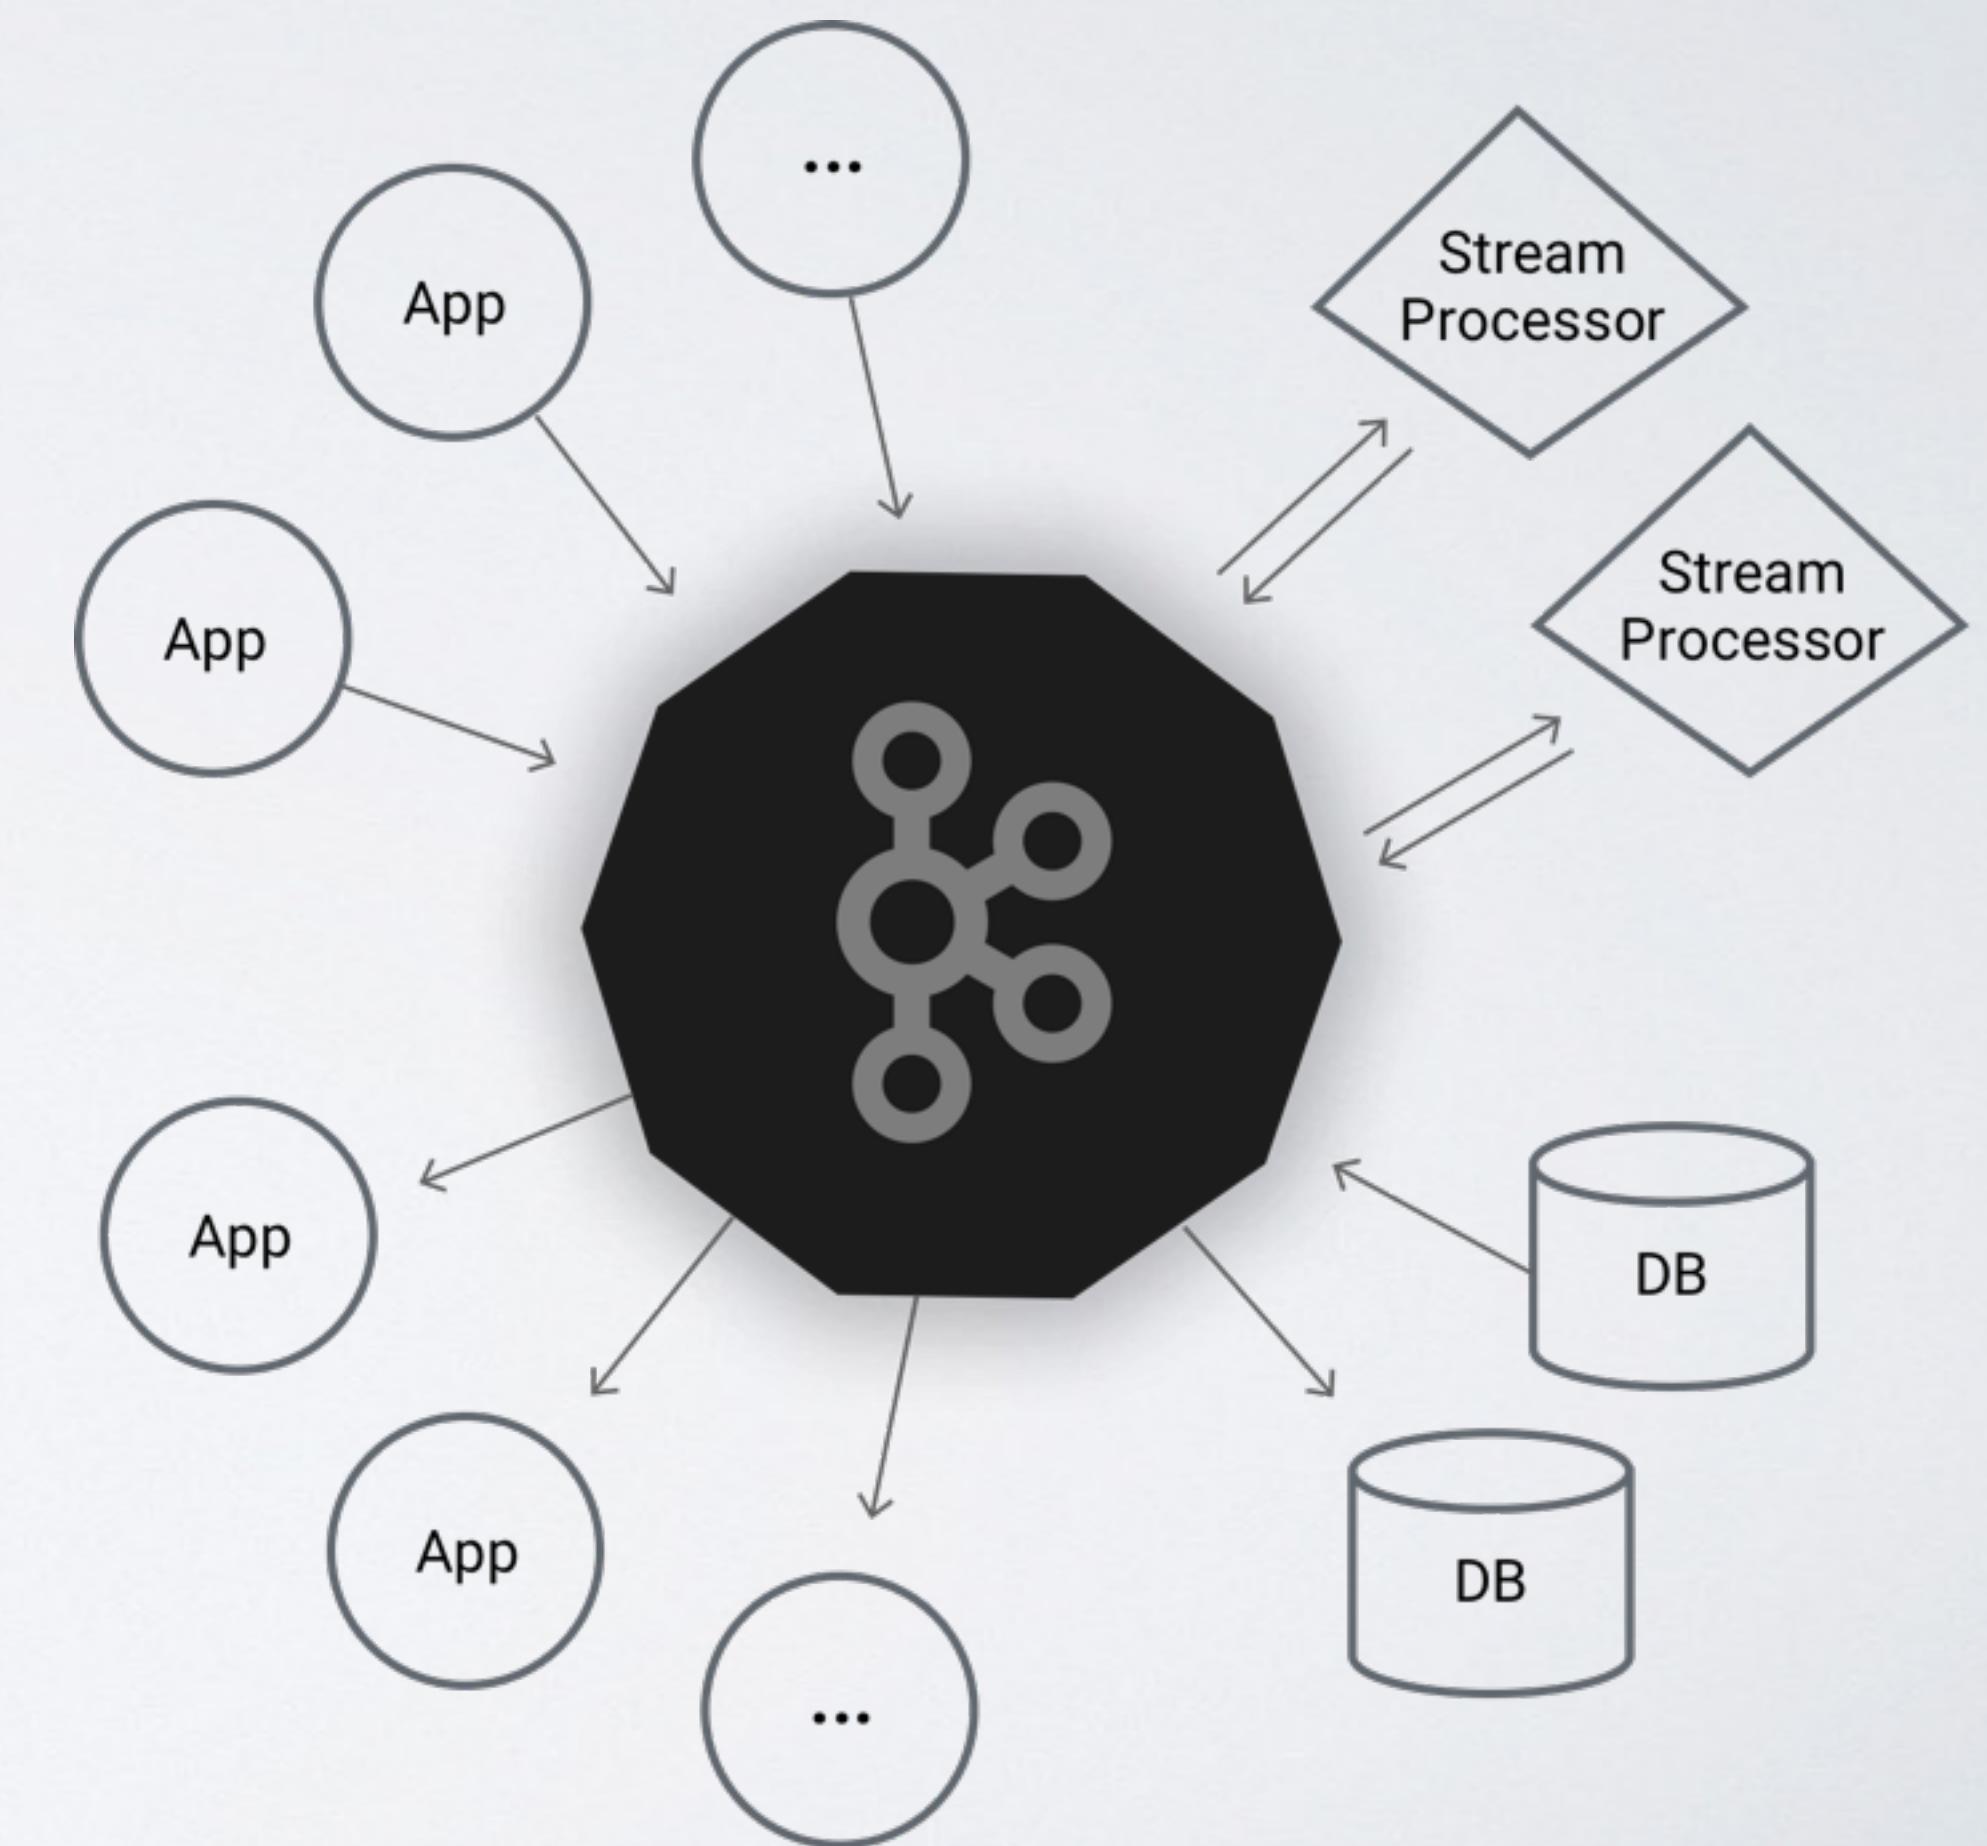
Kafka: Streaming Platform

- *Persistent Buffering*
- *Logical Ordering*
- *Scalable “source-of-truth”*



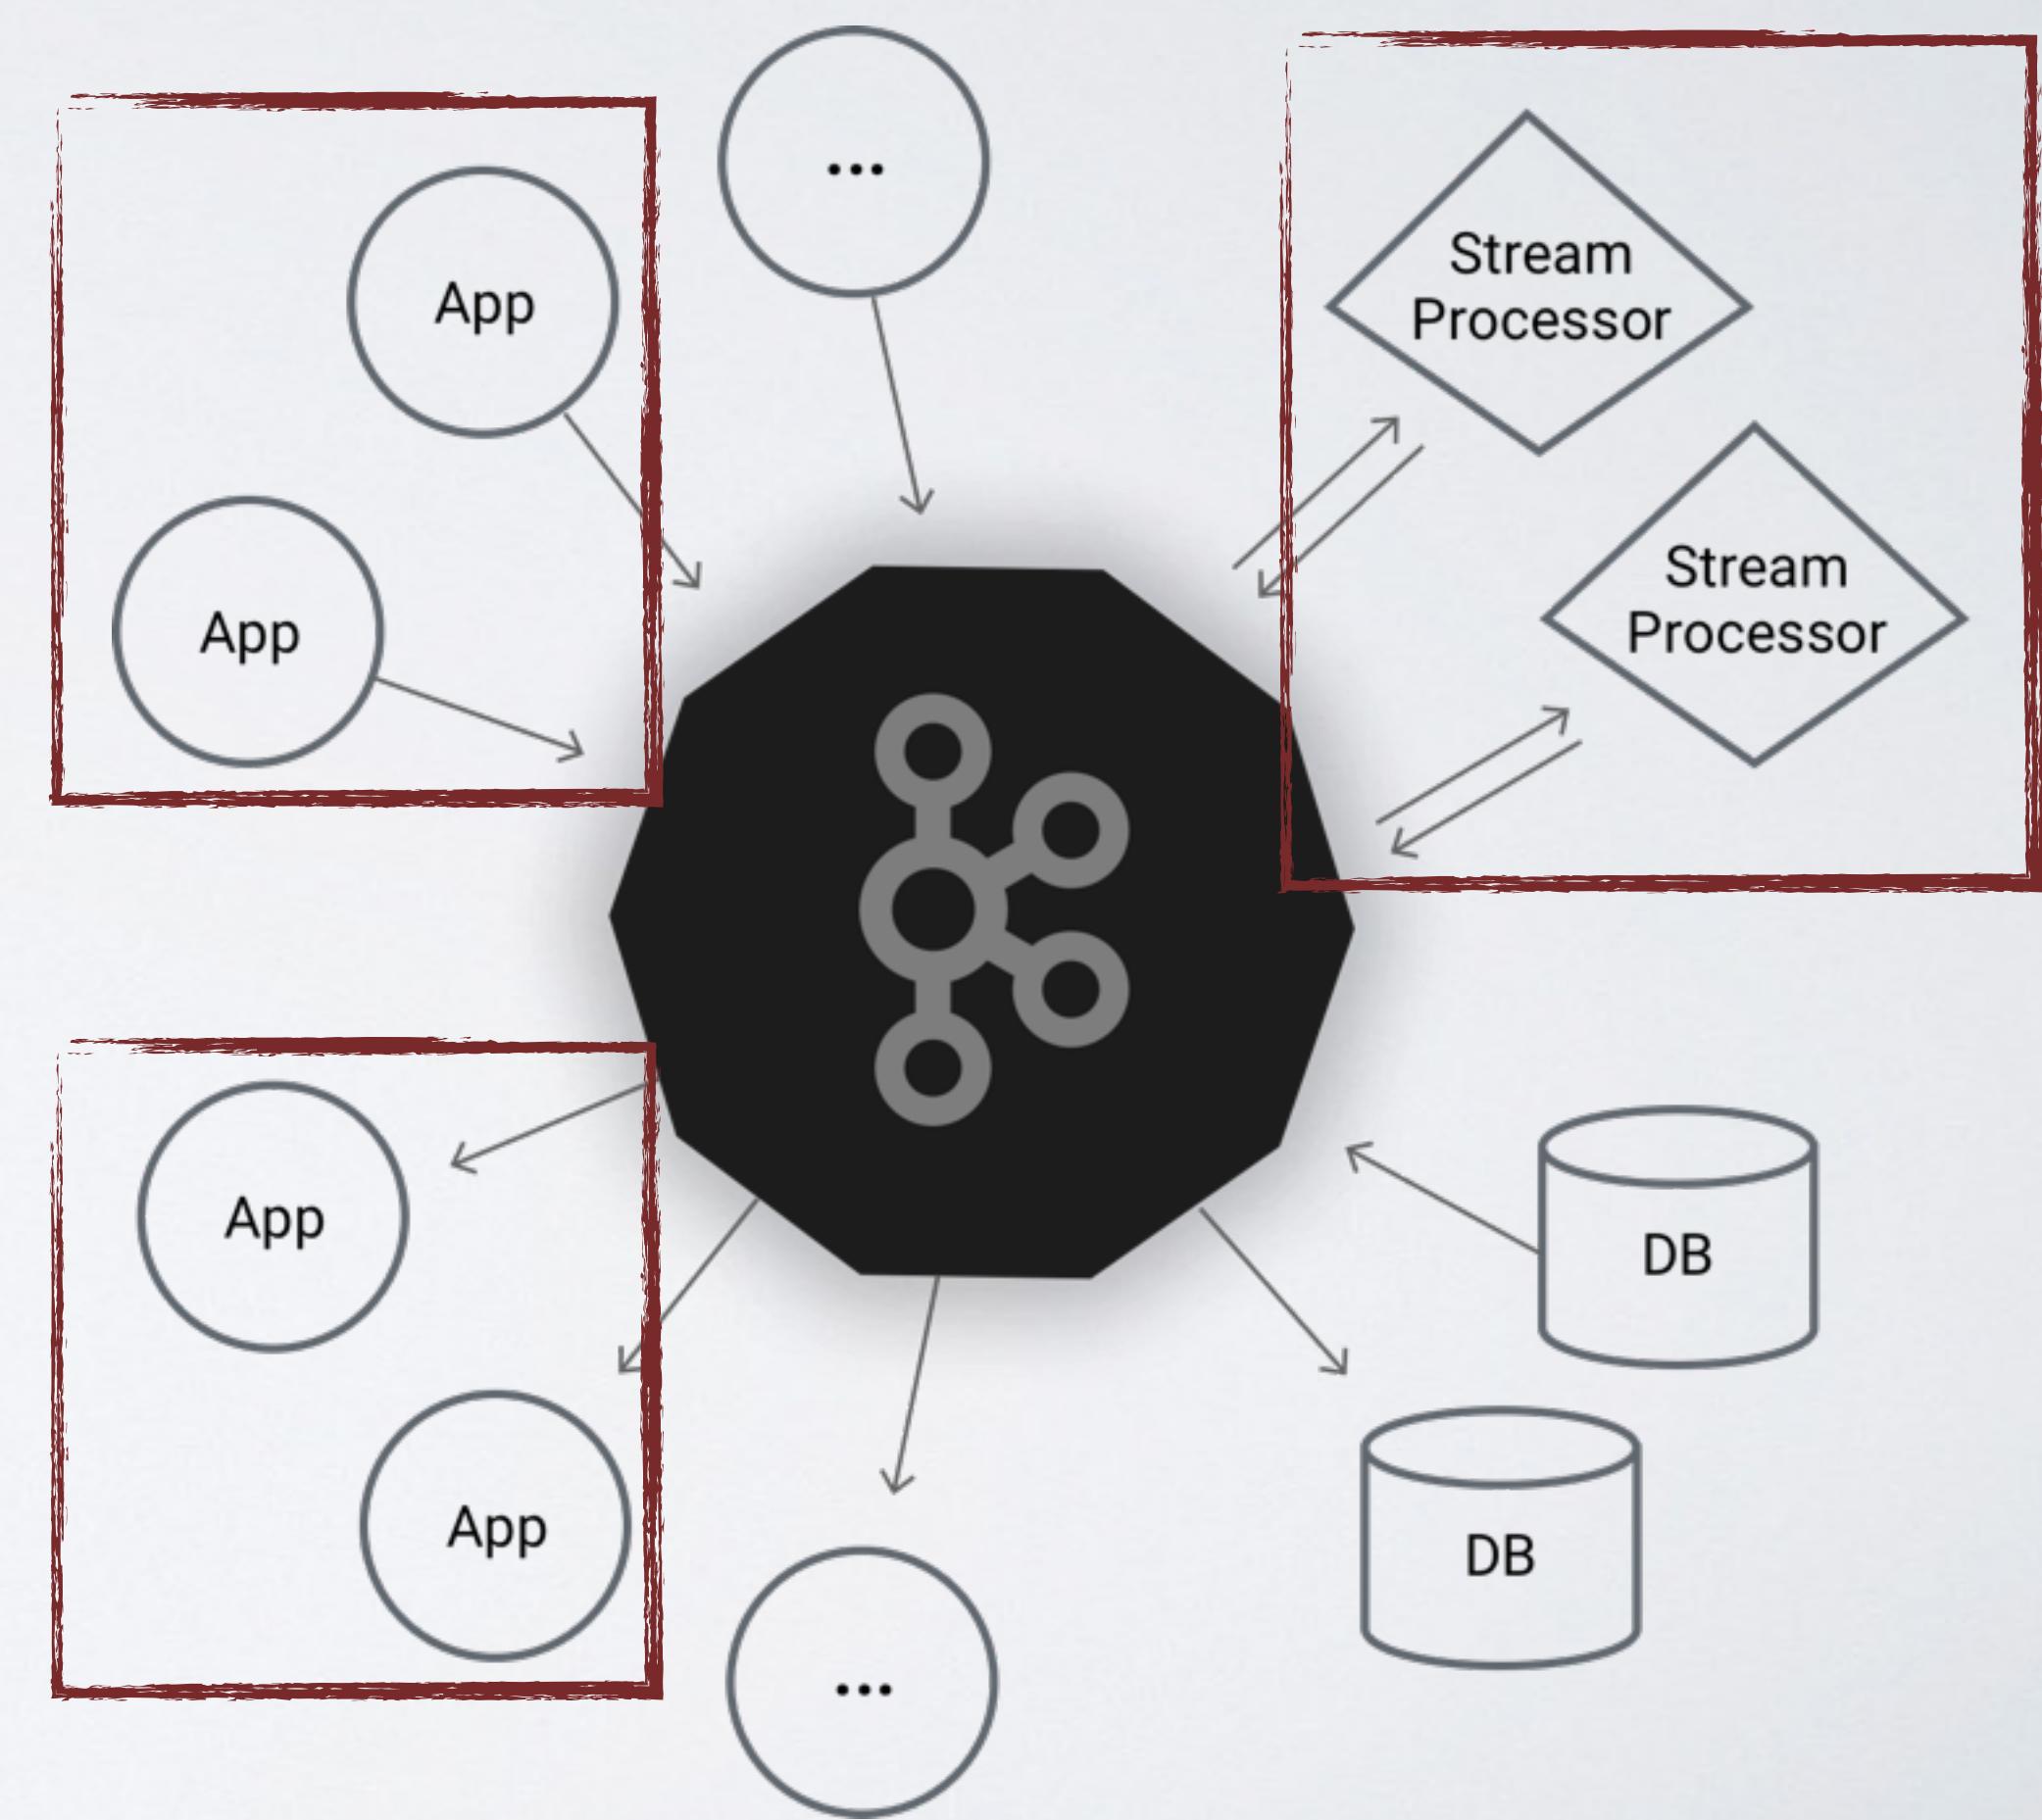
Kafka: Streaming Platform

- ~~Publish/Subscribe~~
- Move data around as *online streams*
- **Logical Ordering**
- **Store Scalable “source-of-truth”**
“Source-of-truth” continuous data
- **Process**
- React / process data in real-time



Kafka: Streaming Platform

- **Publish / Subscribe**
 - Move *data around as online streams*
- **Store**
 - “*Source-of-truth*” continuous data
- **Process**
 - *React / process data in real-time*



Stream Processing with Kafka

Stream Processing with Kafka

- *Option 1: Do It Yourself!*

Stream Processing with Kafka

- *Option 1: Do It Yourself!*

```
while (isRunning) {  
    // read some messages from Kafka  
    inputMessages = consumer.poll();  
  
    // do some processing...  
  
    // send output messages back to Kafka  
    producer.send(outputMessages);  
}
```

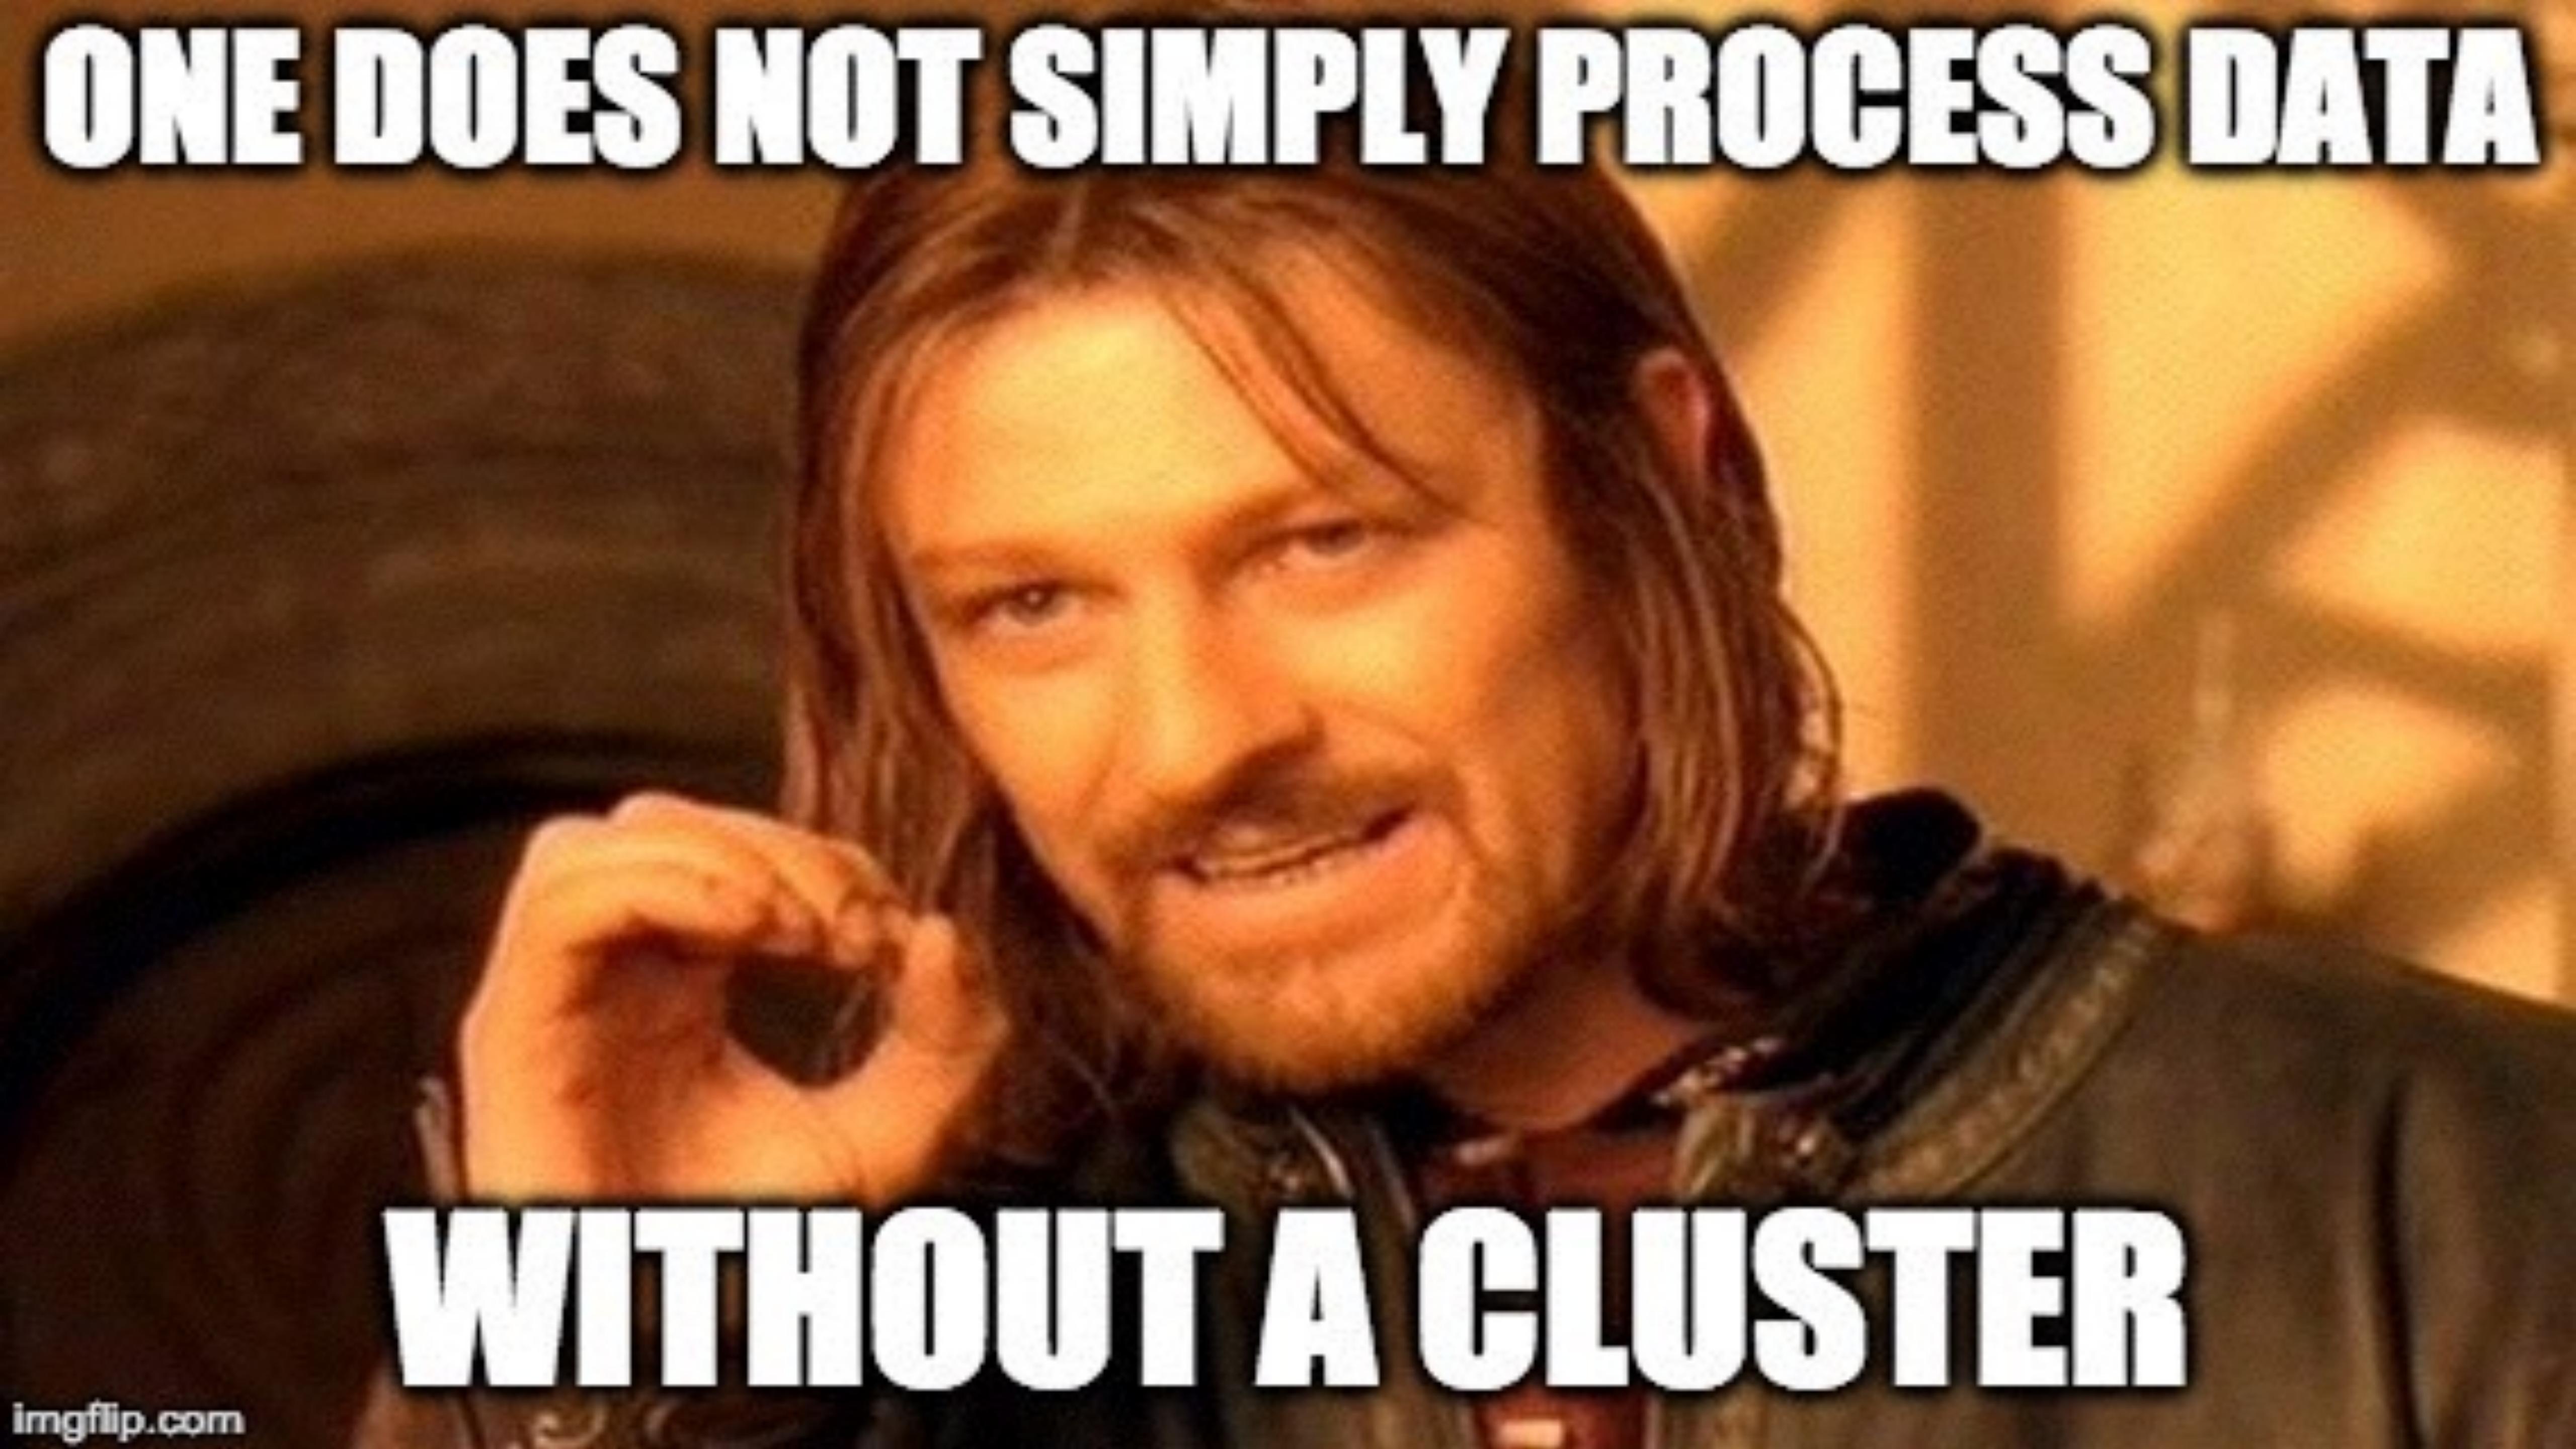
DIY Stream Processing is *Hard*

- *Ordering*
- *Partitioning & Scalability*
- *Fault tolerance*
- *State Management*
- *Time, Window & Out-of-order Data*
- *Re-processing*

Stream Processing with Kafka

- *Option I: Do It Yourself!*
- *Option II: full-fledged stream processing system*
 - *Storm, Spark, Flink, Samza, ..*

ONE DOES NOT SIMPLY PROCESS DATA

A close-up, slightly blurred portrait of Steve Jobs. He has long, light-colored hair and is looking directly at the camera with a thoughtful expression. His right hand is raised to his face, with his fingers resting near his eye. The background is dark and out of focus.

WITHOUT A CLUSTER

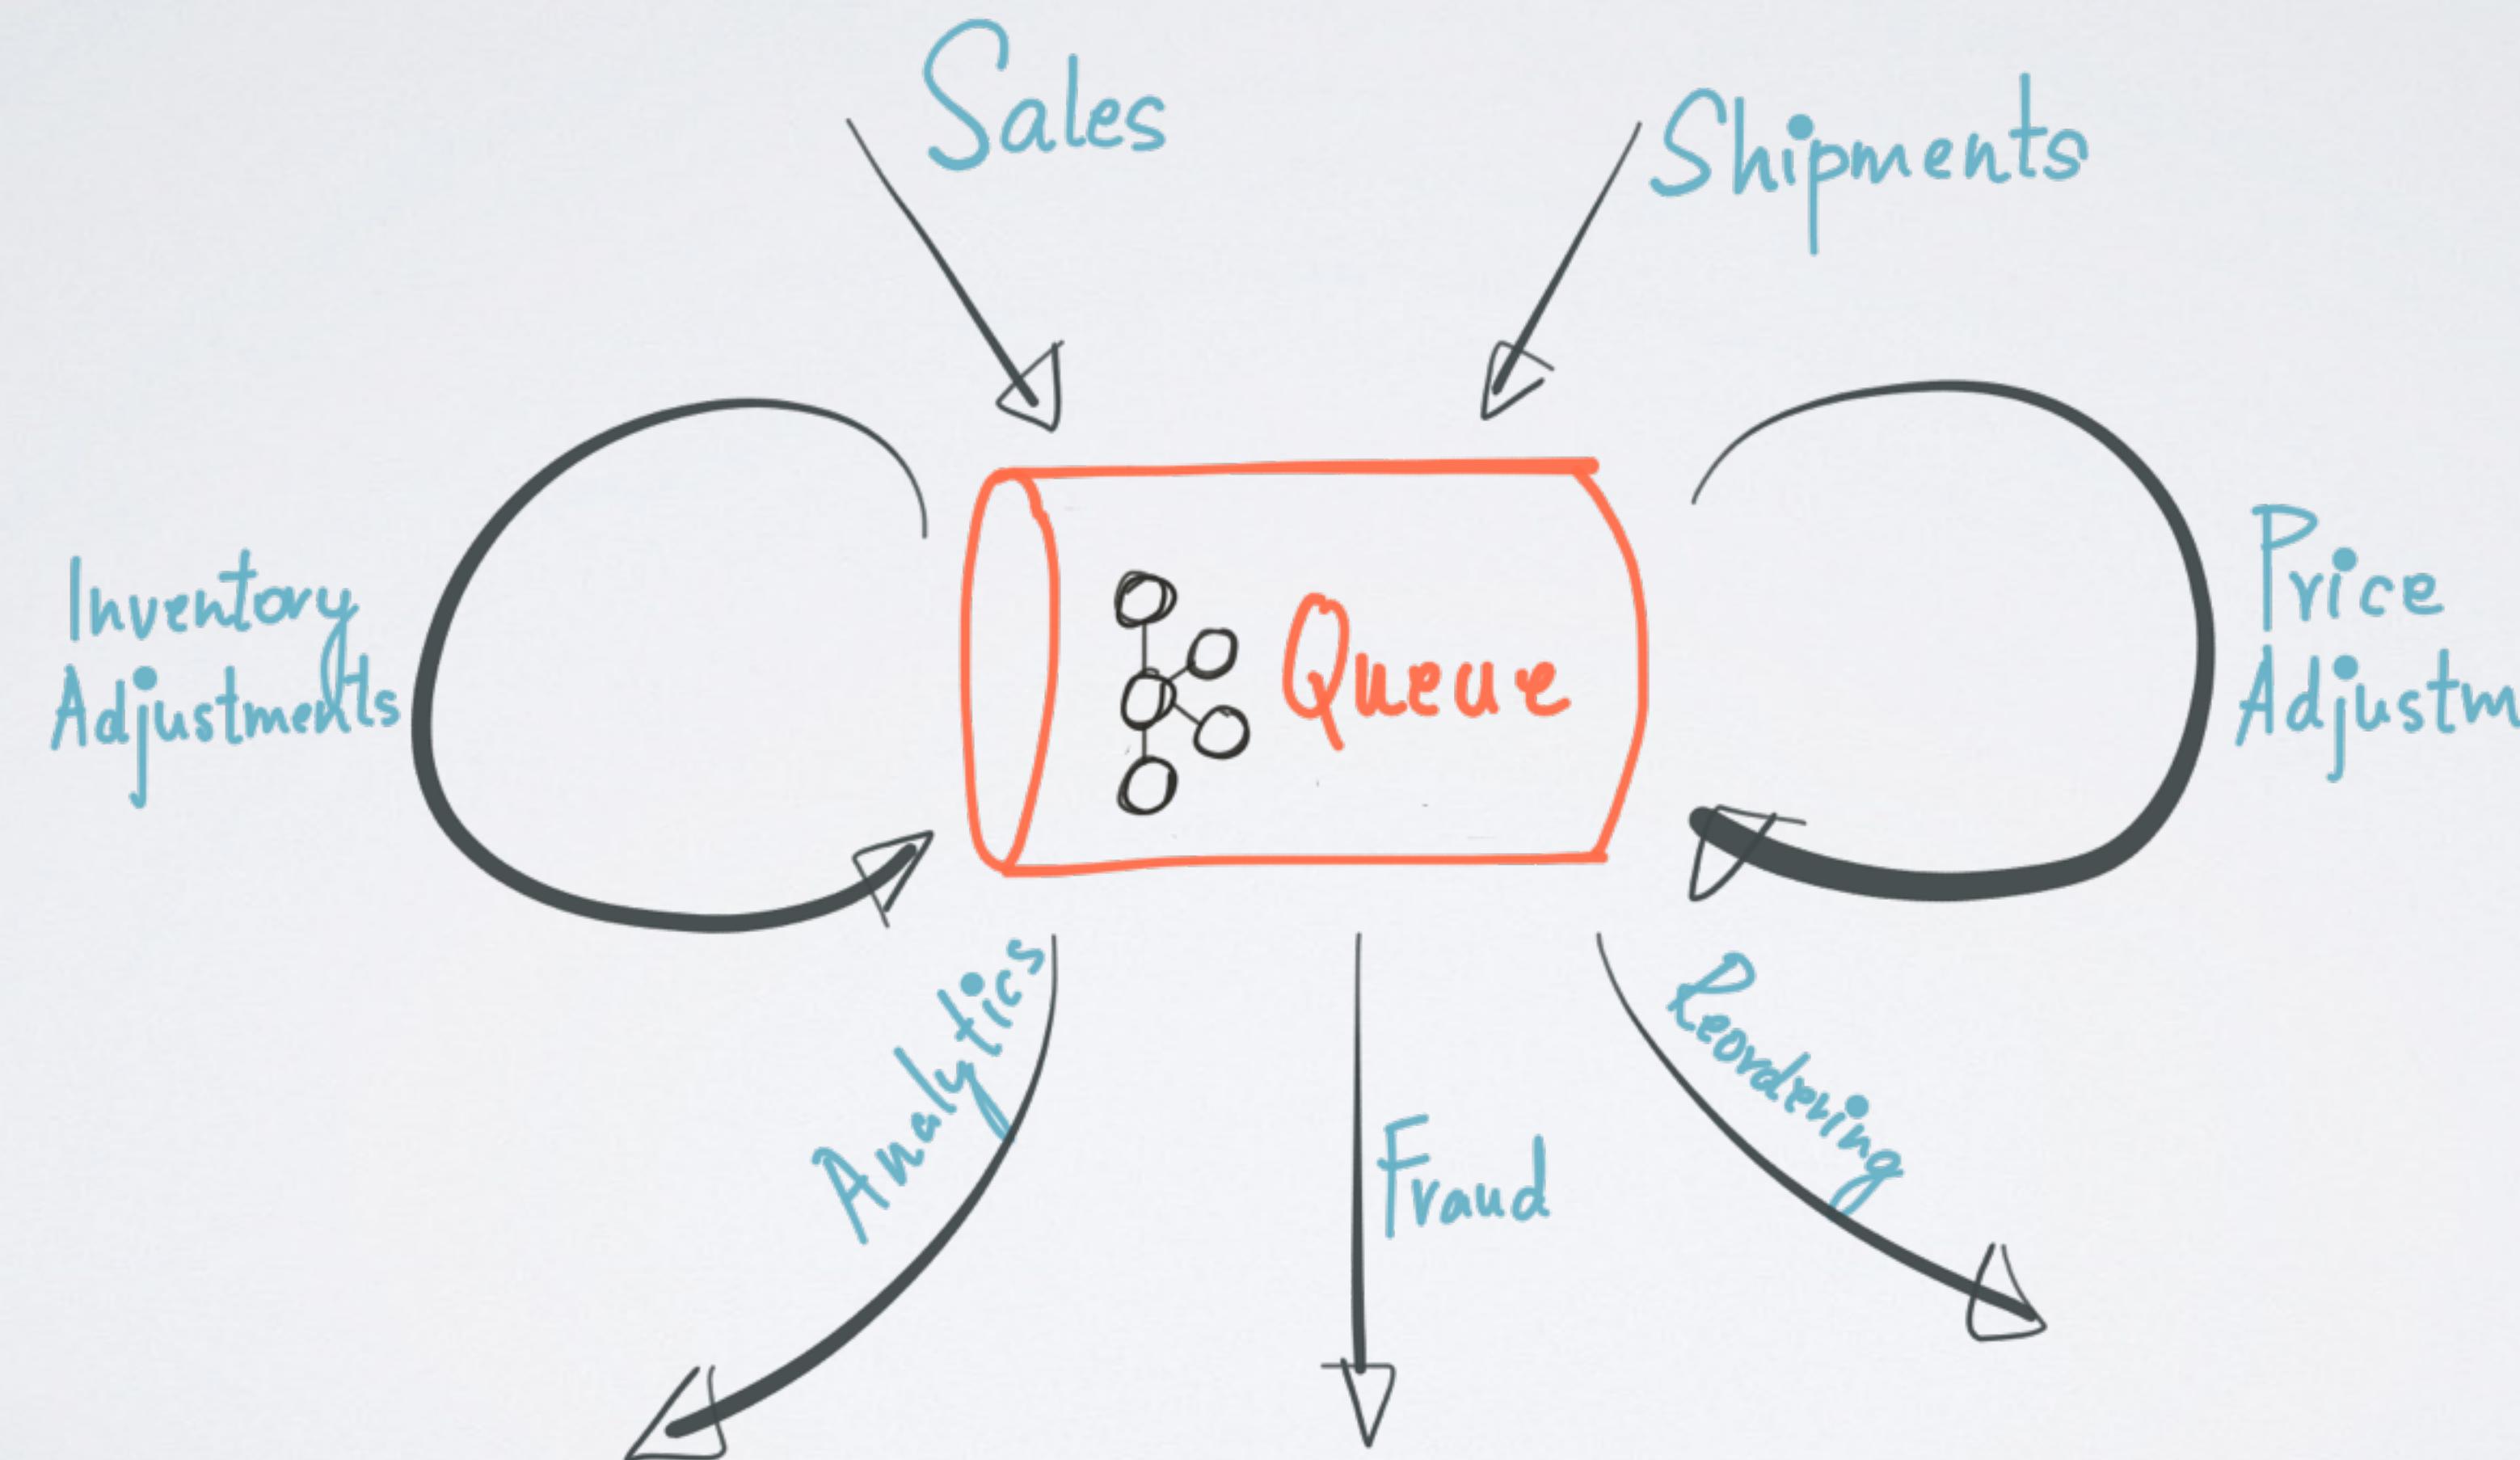
MapReduce Heritage?

- *Config Management*
- *Resource Management*
- *Deployment*
- *etc..*

MapReduce Heritage?

- *Config Management*
- ***Can I just use my own?!***
- *etc..*

Example: Async. Micro-Services



Stream Processing with Kafka

- *Option I: Do It Yourself!*
- *Option II: full-fledged stream processing system*
- *Option III: lightweight stream processing library*



Kafka Streams (0.10+)



- *New client library besides producer and consumer*
- *Powerful yet easy-to-use*
 - *Event-at-a-time, Stateful*
 - *Windowing with out-of-order handling*
 - *Highly scalable, distributed, fault tolerant*
 - *and more..*

Anywhere, anytime

Ok.



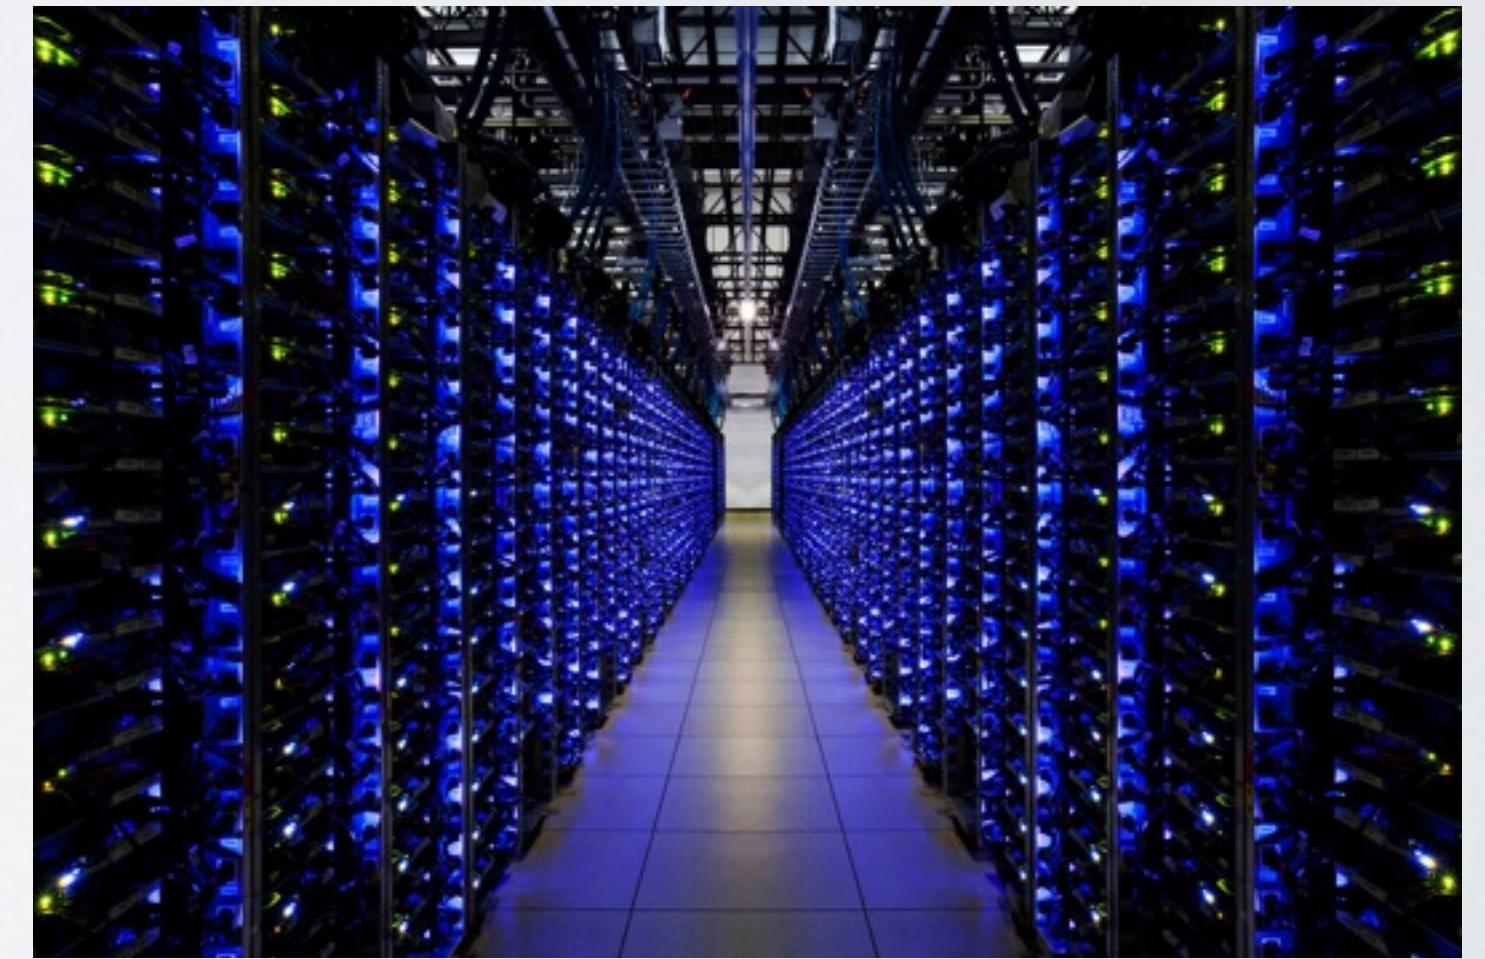
Ok.



Ok.



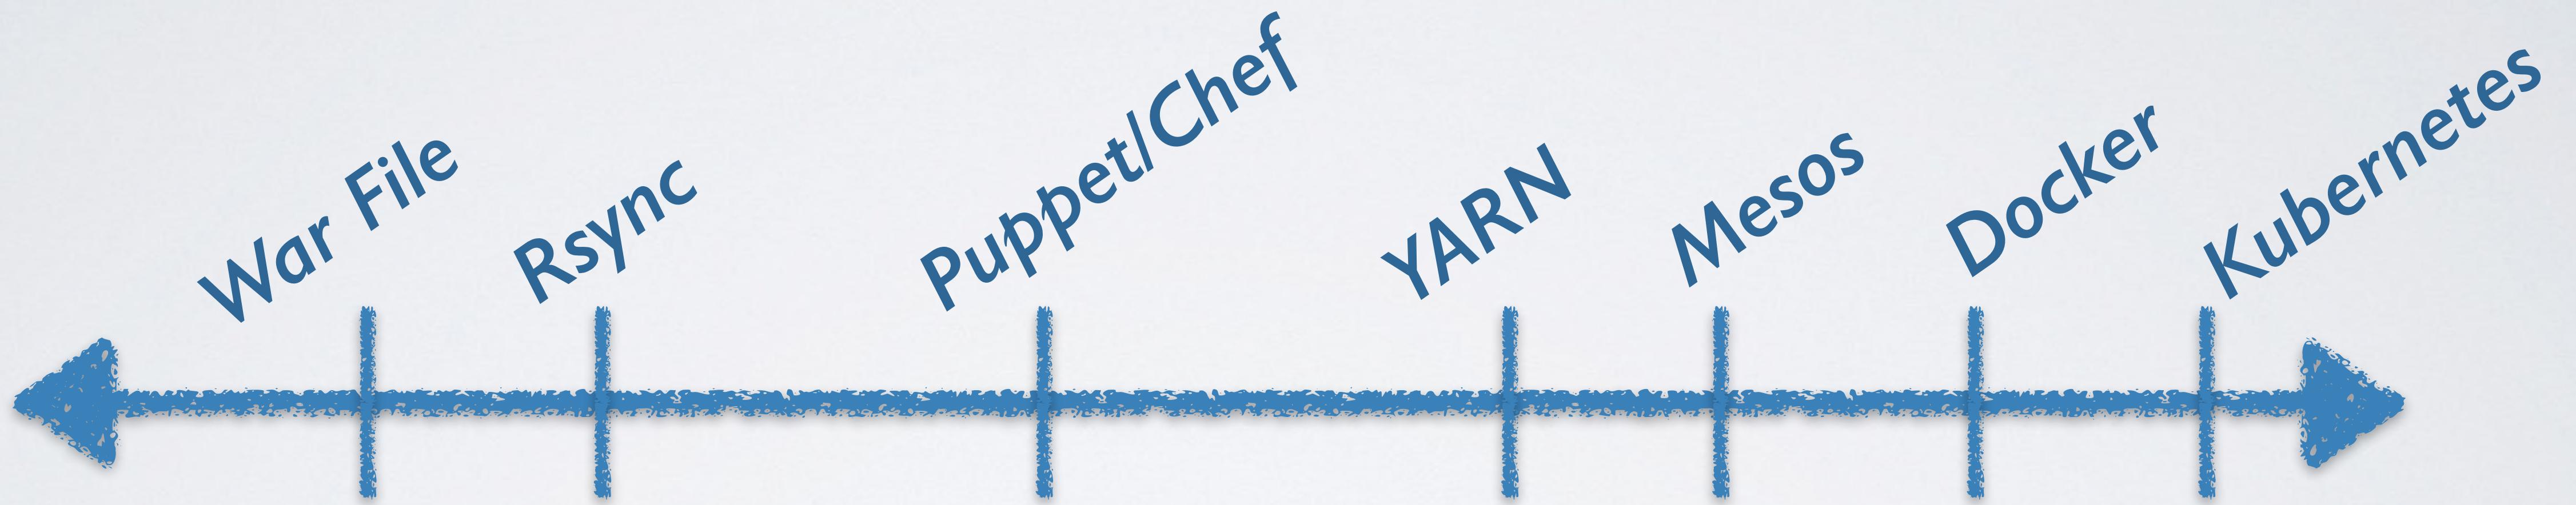
Ok.



Anywhere, anytime

```
<dependency>  
  <groupId>org.apache.kafka</groupId>  
  <artifactId>kafka-streams</artifactId>  
  <version>0.10.0.0</version>  
</dependency>
```

Anywhere, anytime



Simple is Beautiful

Kafka Streams DSL

```
public static void main(String[] args) {  
    // specify the processing topology by first reading in a stream from a topic  
    KStream<String, String> words = builder.stream("topic1");  
  
    // count the words in this stream as an aggregated table  
    KTable<String, Long> counts = words.countByKey("Counts");  
  
    // write the result table to a new topic  
    counts.to("topic2");  
  
    // create a stream processing instance and start running it  
    KafkaStreams streams = new KafkaStreams(builder, config);  
    streams.start();  
}
```

Kafka Streams DSL

```
public static void main(String[] args) {
    // specify the processing topology by first reading in a stream from a topic
    KStream<String, String> words = builder.stream("topic1");

    // count the words in this stream as an aggregated table
    KTable<String, Long> counts = words.countByKey("Counts");

    // write the result table to a new topic
    counts.to("topic2");

    // create a stream processing instance and start running it
    KafkaStreams streams = new KafkaStreams(builder, config);
    streams.start();
}
```

Kafka Streams DSL

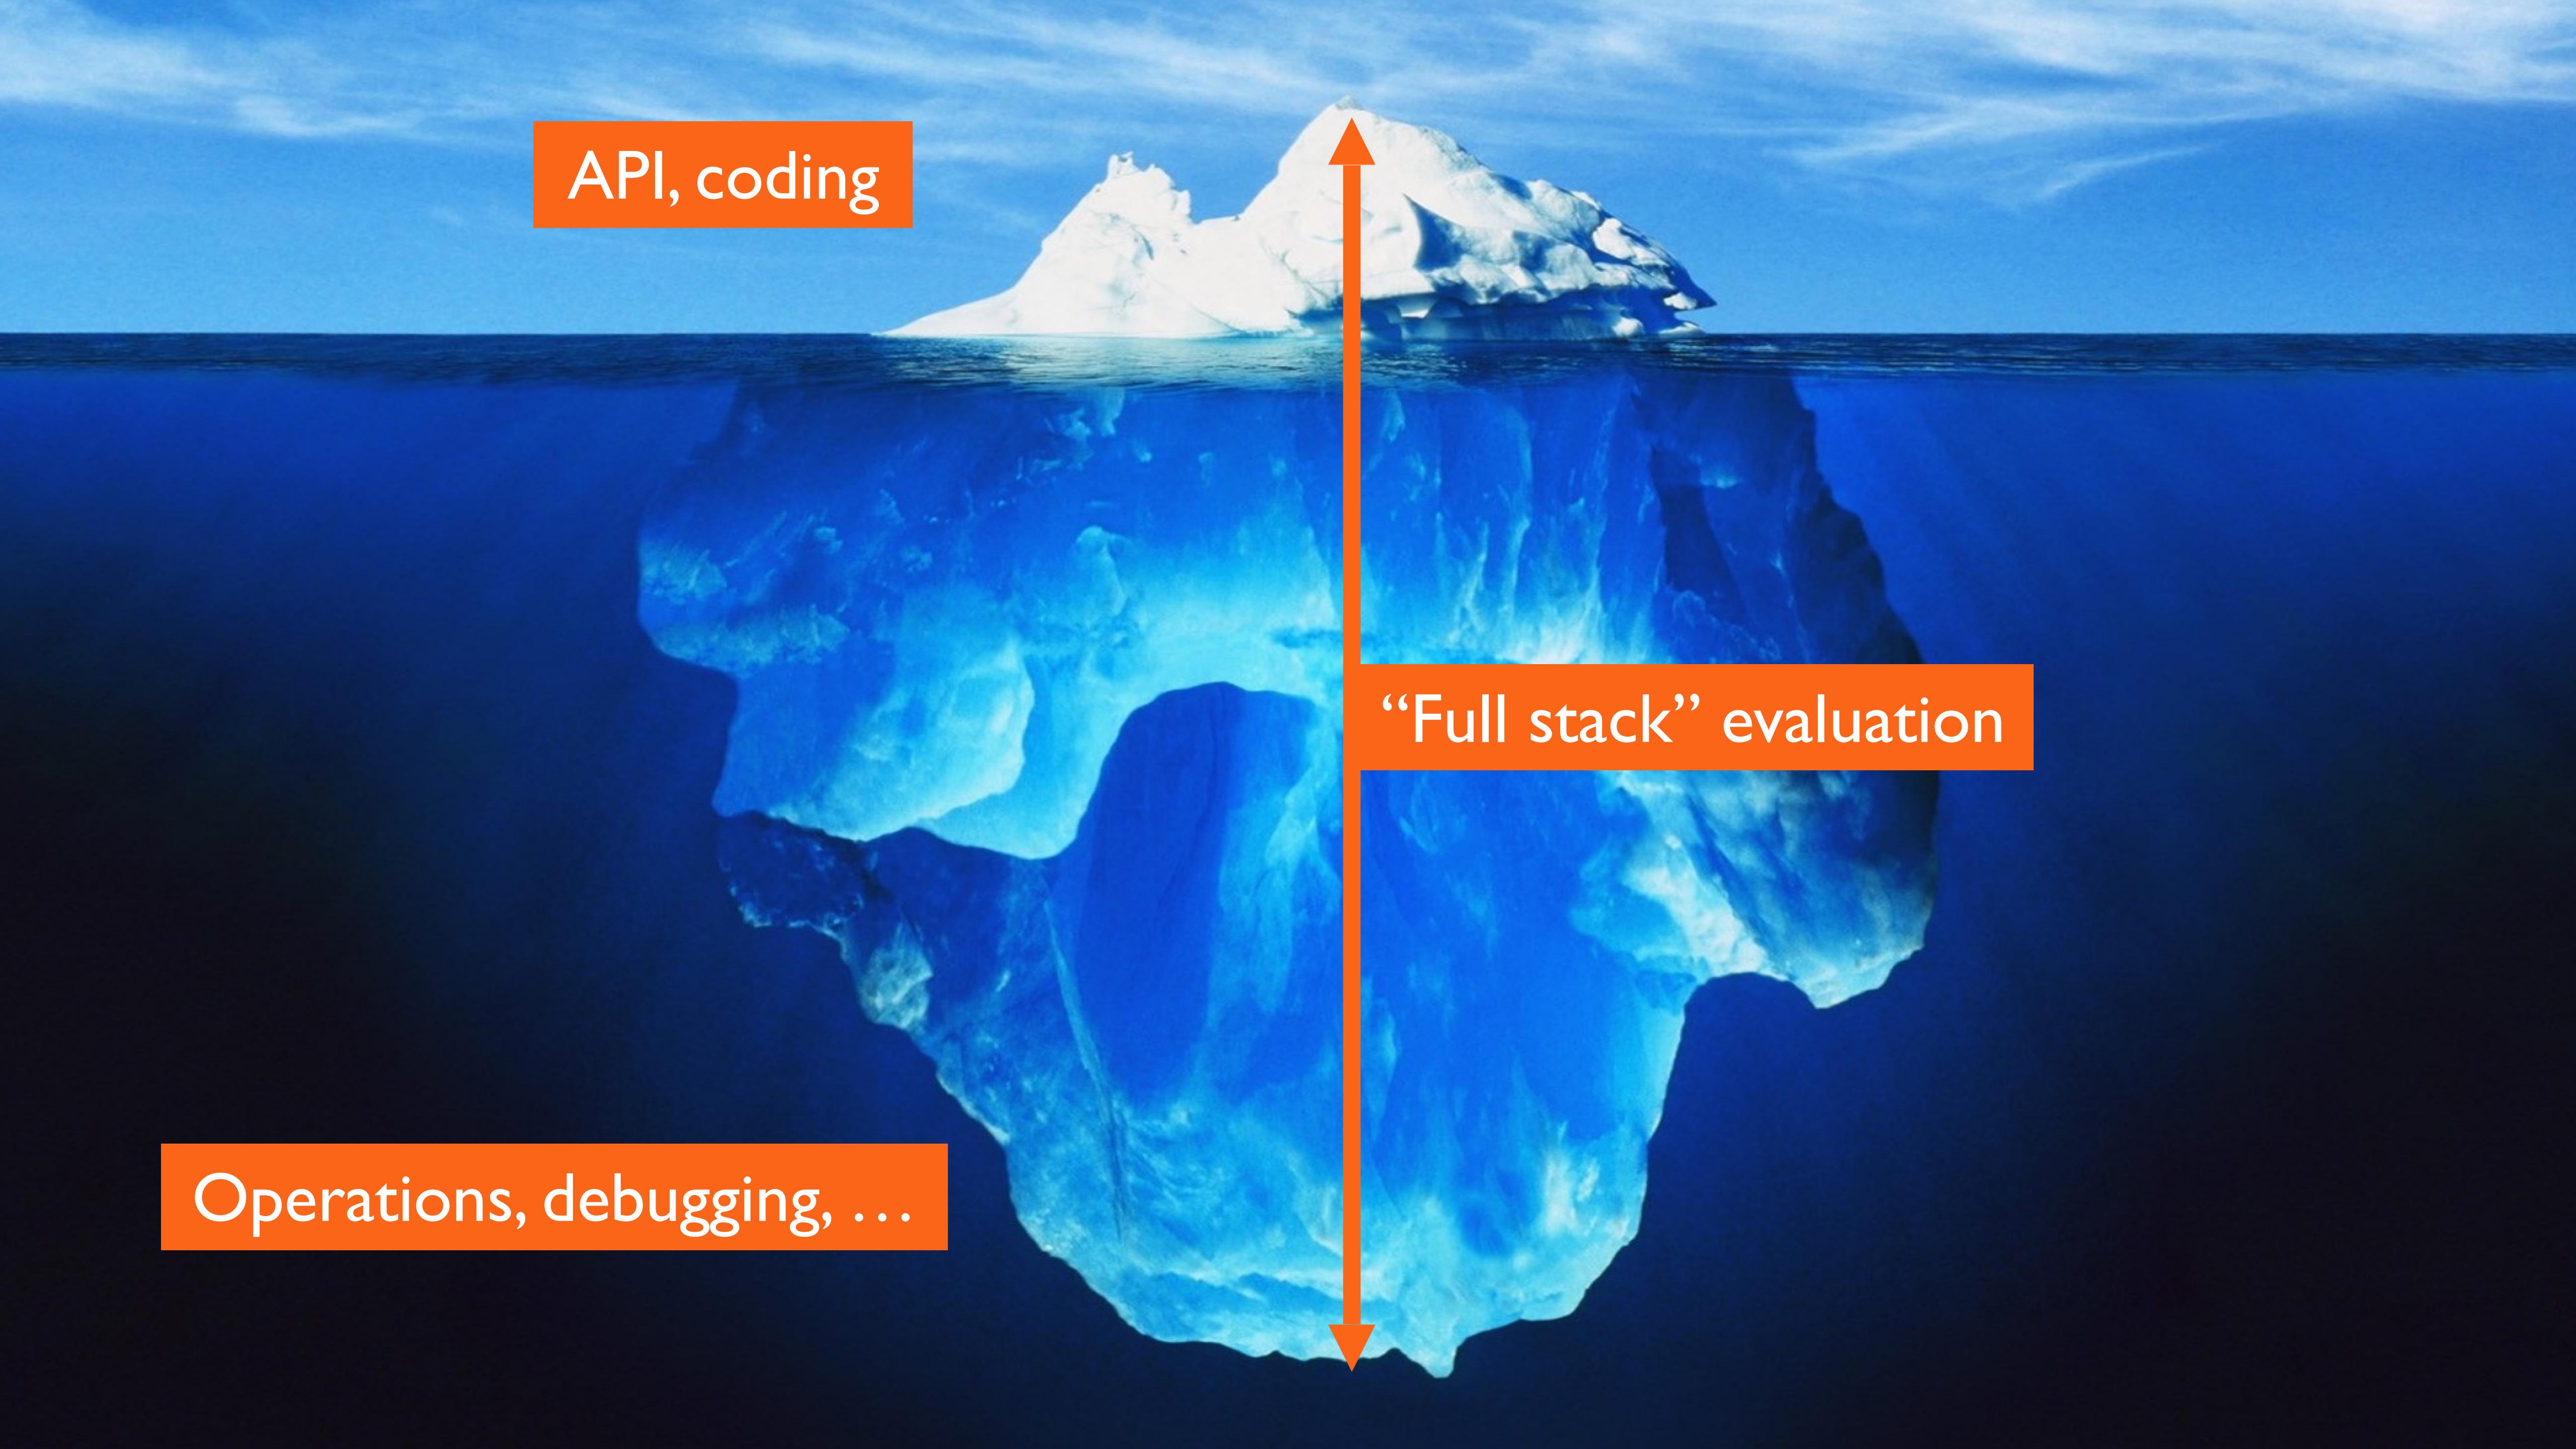
```
public static void main(String[] args) {  
    // specify the processing topology by first reading in a stream from a topic  
    KStream<String, String> words = builder.stream("topic1");  
  
    // count the words in this stream as an aggregated table  
    KTable<String, Long> counts = words.countByKey("Counts");  
  
    // write the result table to a new topic  
    counts.to("topic2");  
  
    // create a stream processing instance and start running it  
    KafkaStreams streams = new KafkaStreams(builder, config);  
    streams.start();  
}
```

Kafka Streams DSL

```
public static void main(String[] args) {  
    // specify the processing topology by first reading in a stream from a topic  
    KStream<String, String> words = builder.stream("topic1");  
  
    // count the words in this stream as an aggregated table  
    KTable<String, Long> counts = words.countByKey("Counts");  
  
    // write the result table to a new topic  
    counts.to("topic2");  
  
    // create a stream processing instance and start running it  
    KafkaStreams streams = new KafkaStreams(builder, config);  
    streams.start();  
}
```

Kafka Streams DSL

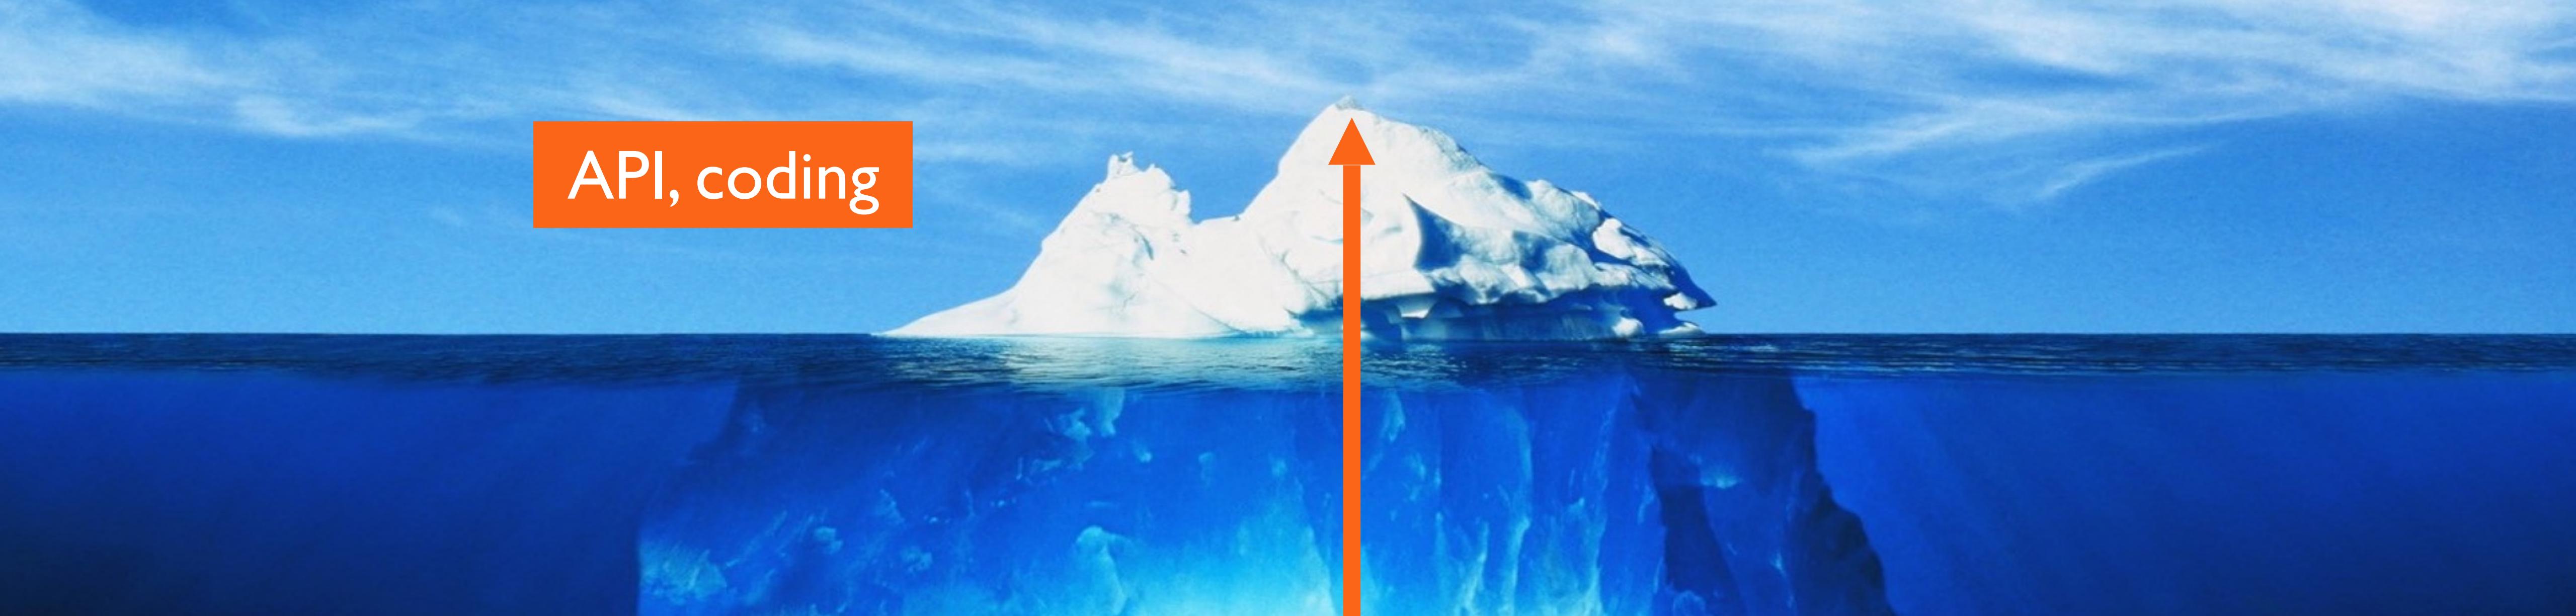
```
public static void main(String[] args) {  
    // specify the processing topology by first reading in a stream from a topic  
    KStream<String, String> words = builder.stream("topic1");  
  
    // count the words in this stream as an aggregated table  
    KTable<String, Long> counts = words.countByKey("Counts");  
  
    // write the result table to a new topic  
    counts.to("topic2");  
  
    // create a stream processing instance and start running it  
    KafkaStreams streams = new KafkaStreams(builder, config);  
    streams.start();  
}
```

A photograph of a large iceberg floating in a deep blue ocean under a clear blue sky. The visible portion of the iceberg above the water's surface is small compared to the massive, submerged portion below.

API, coding

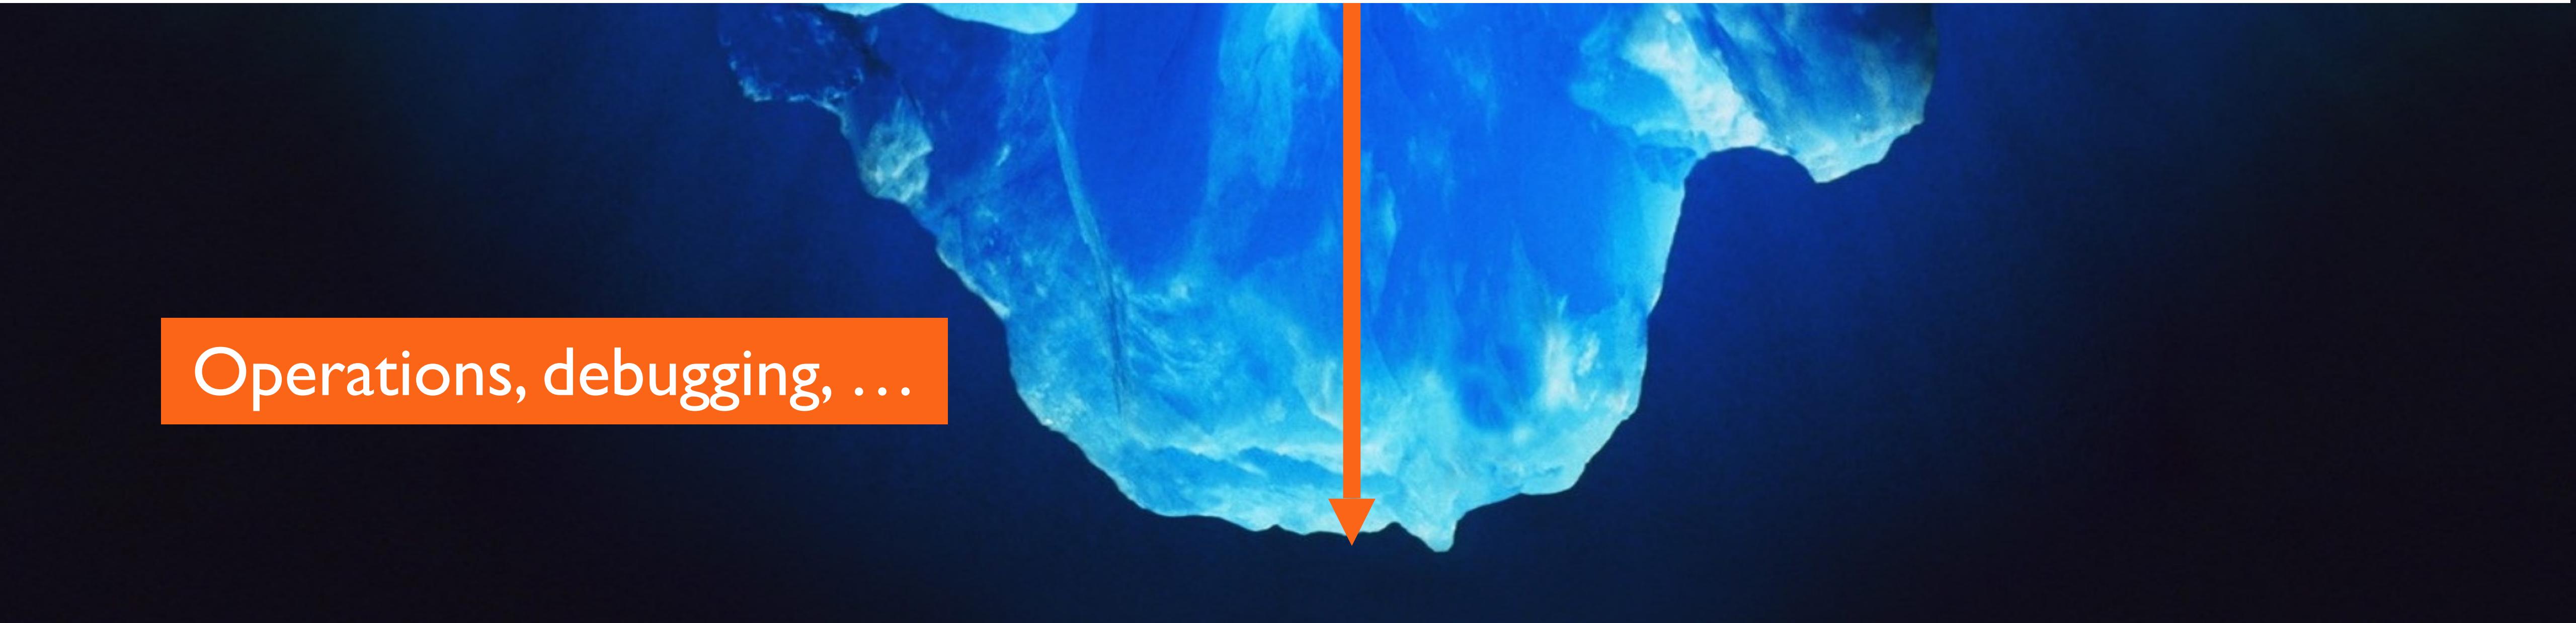
“Full stack” evaluation

Operations, debugging, ...

A photograph of a large iceberg floating in the ocean. The visible portion above the water's surface is white and jagged, while the submerged portion below is a translucent blue.

API, coding

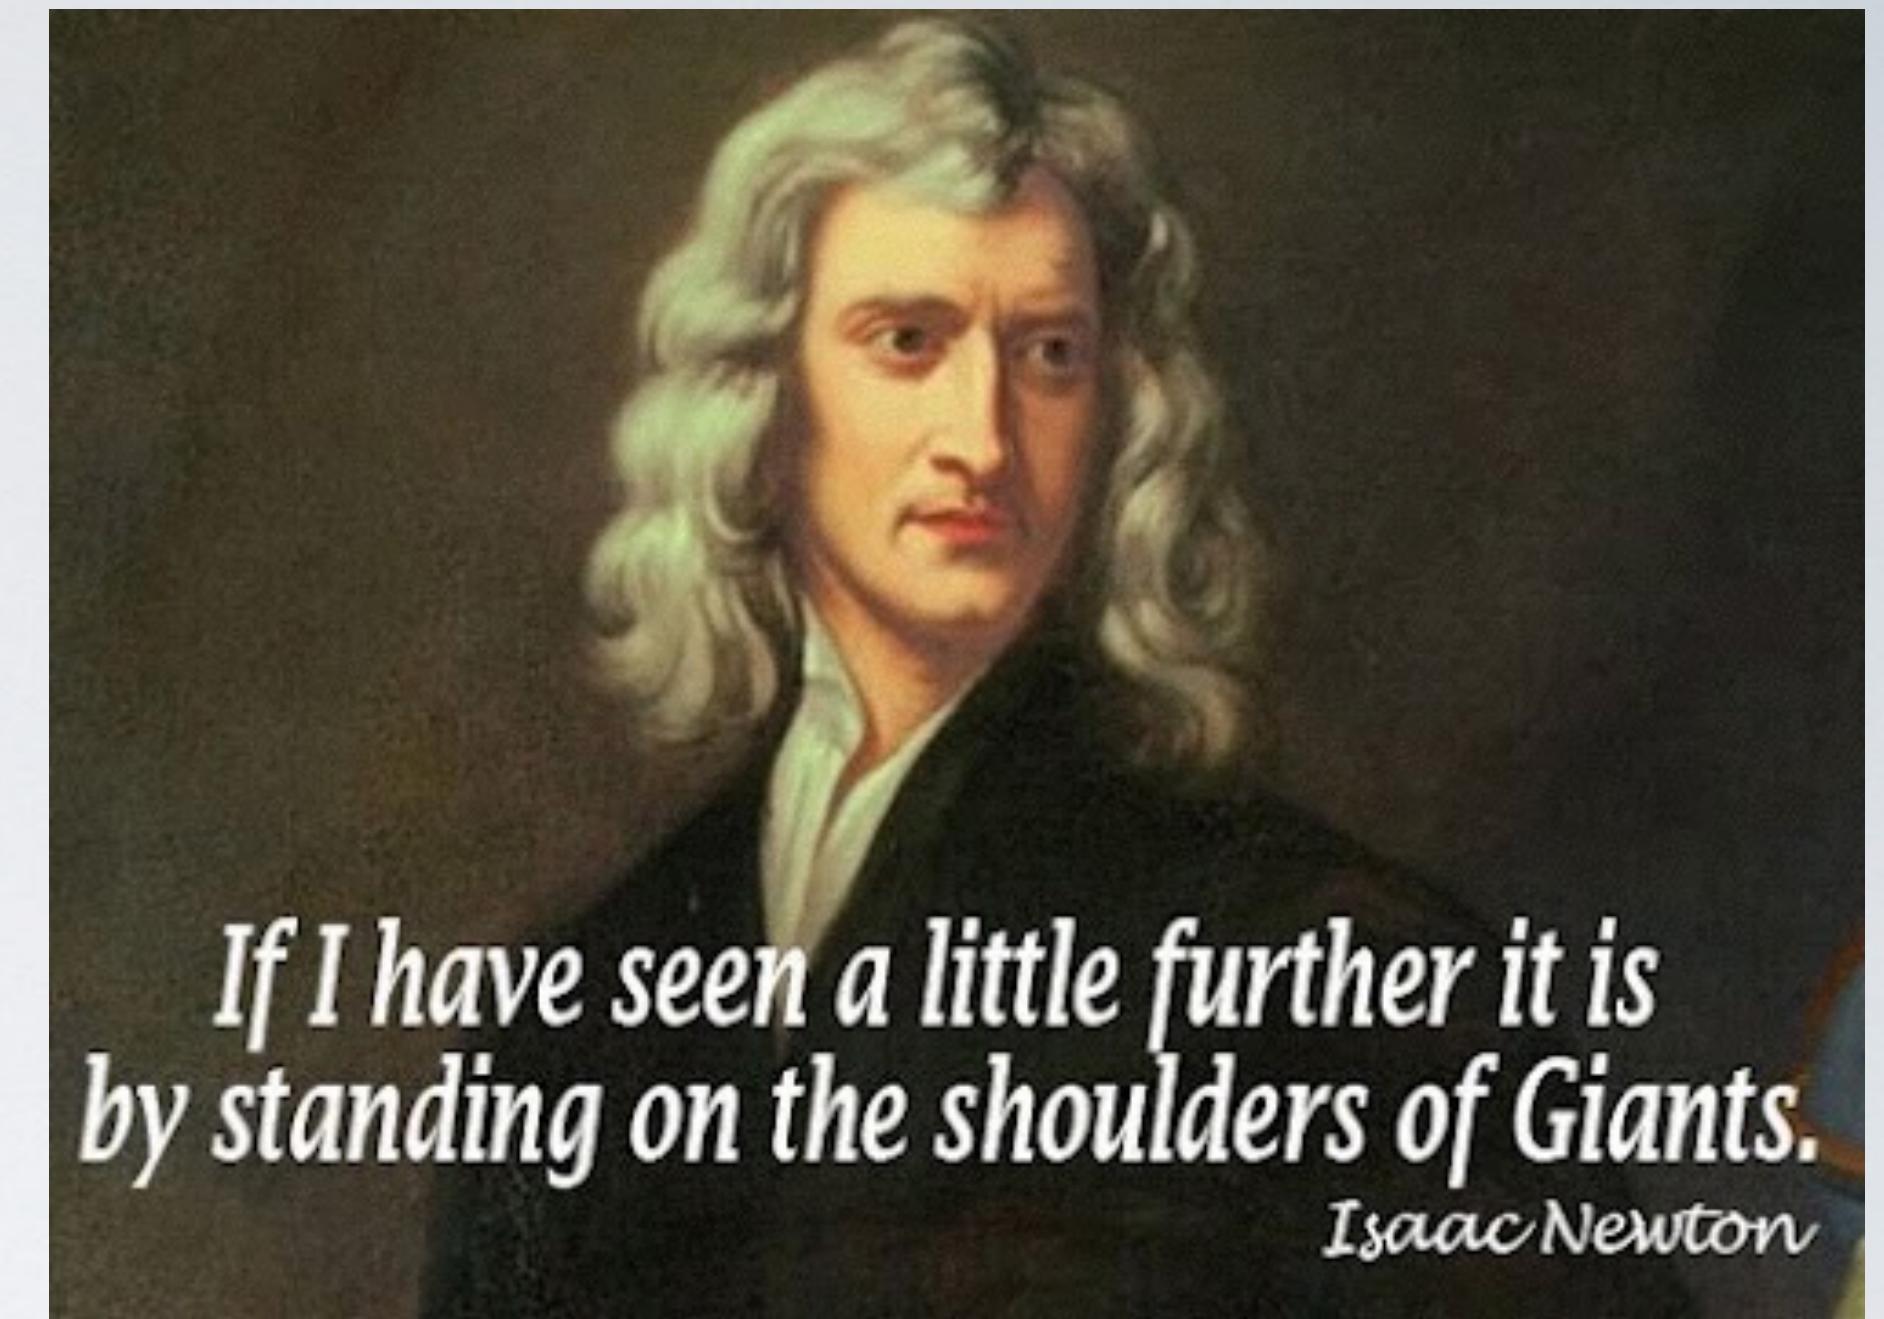
Simple is Beautiful

The same iceberg from the previous image, shown from a lower angle. The submerged part is now more prominent, appearing as a large, translucent blue mass against a dark blue background.

Operations, debugging, ...

Key Idea:

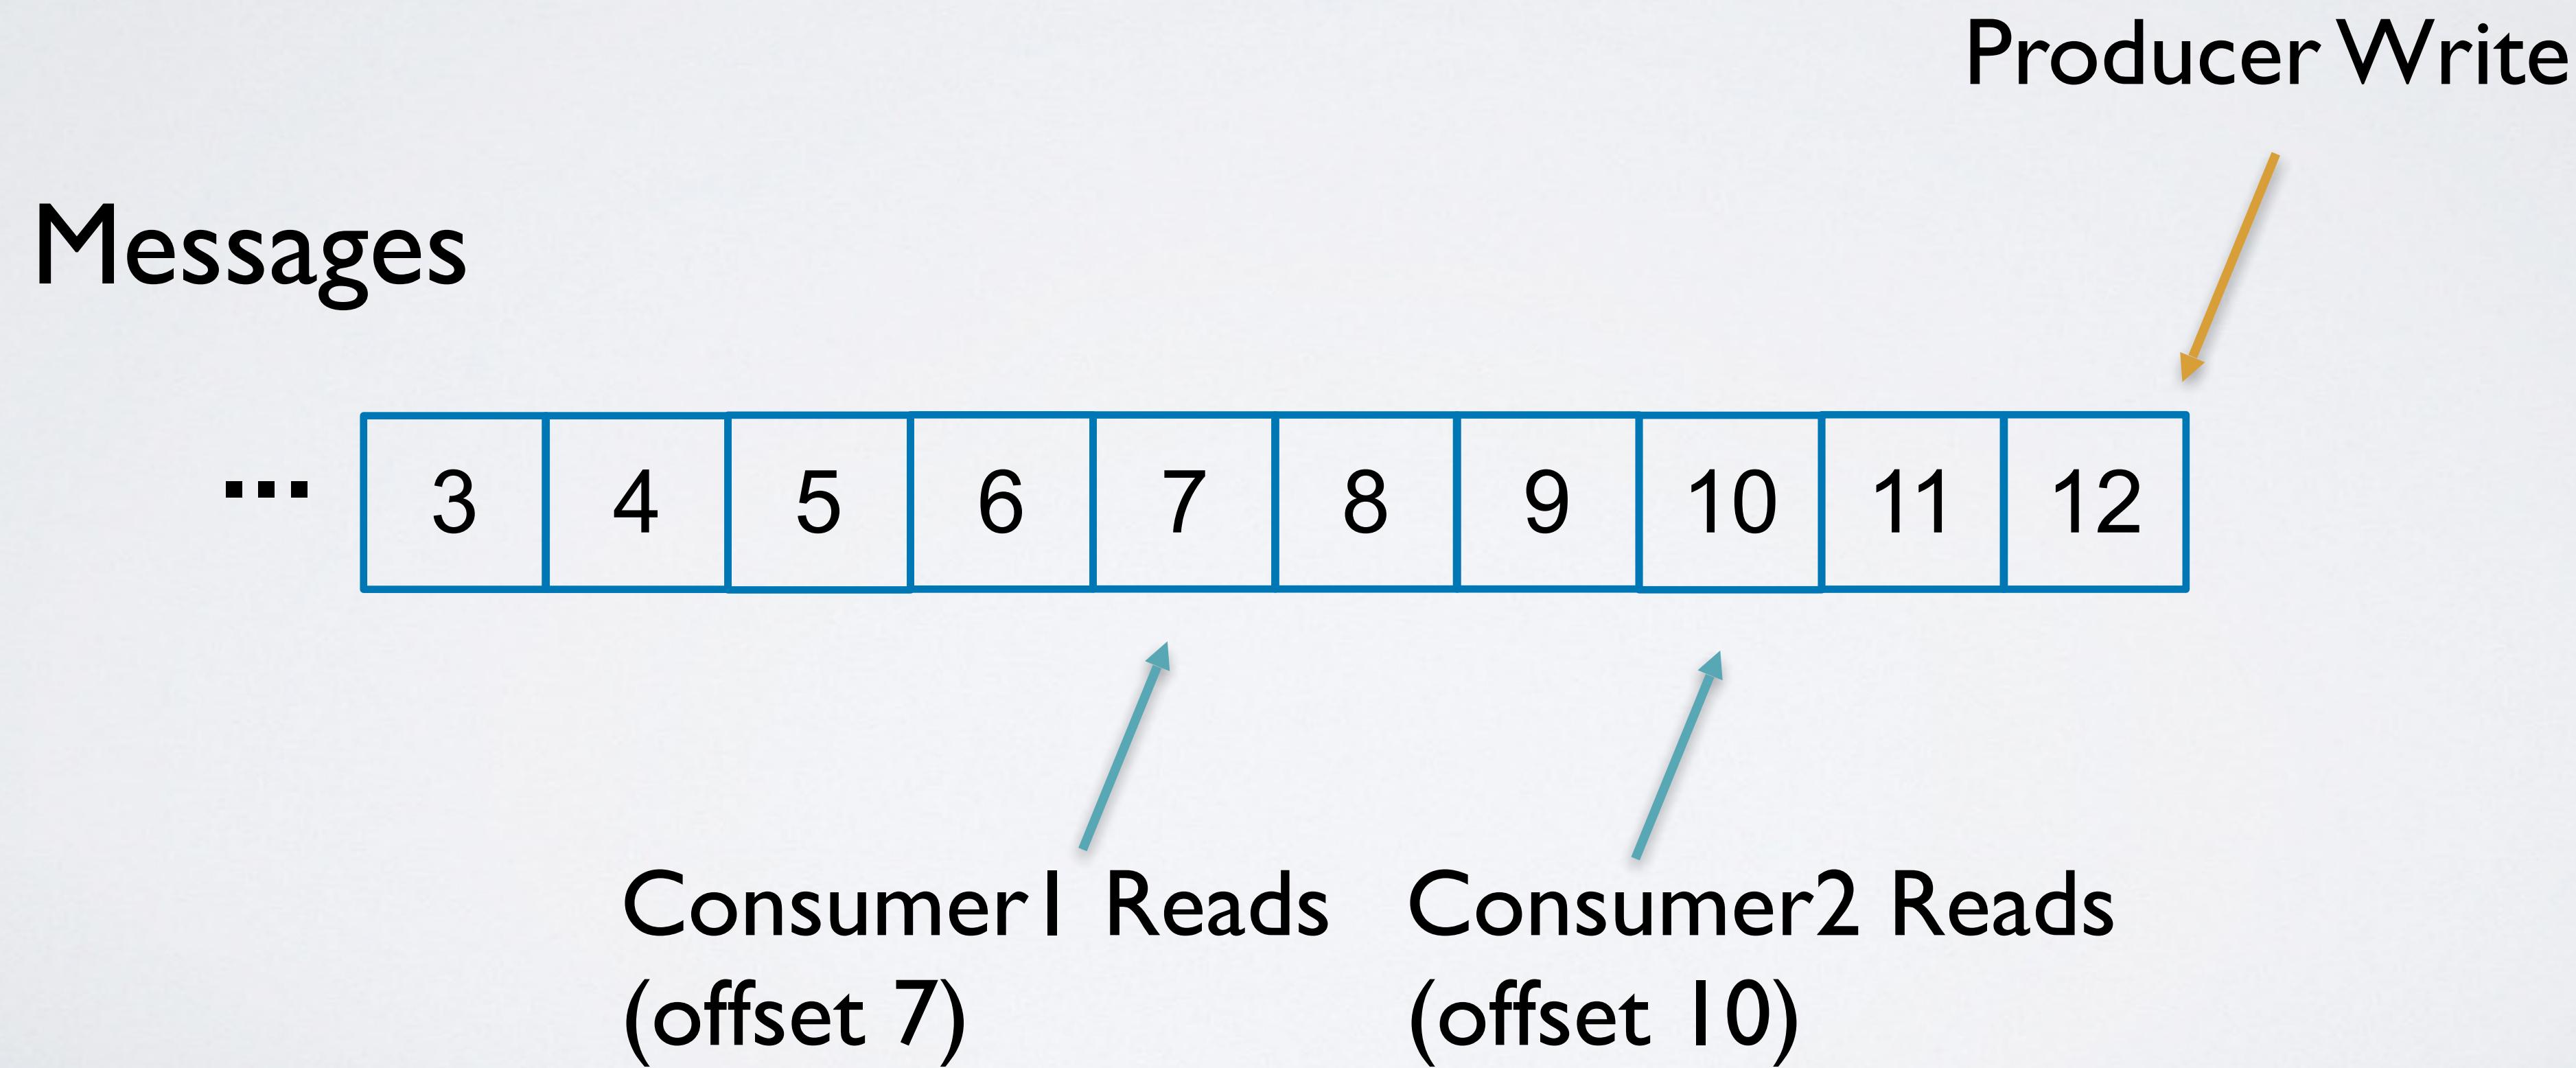
Outsource hard problems to Kafka!



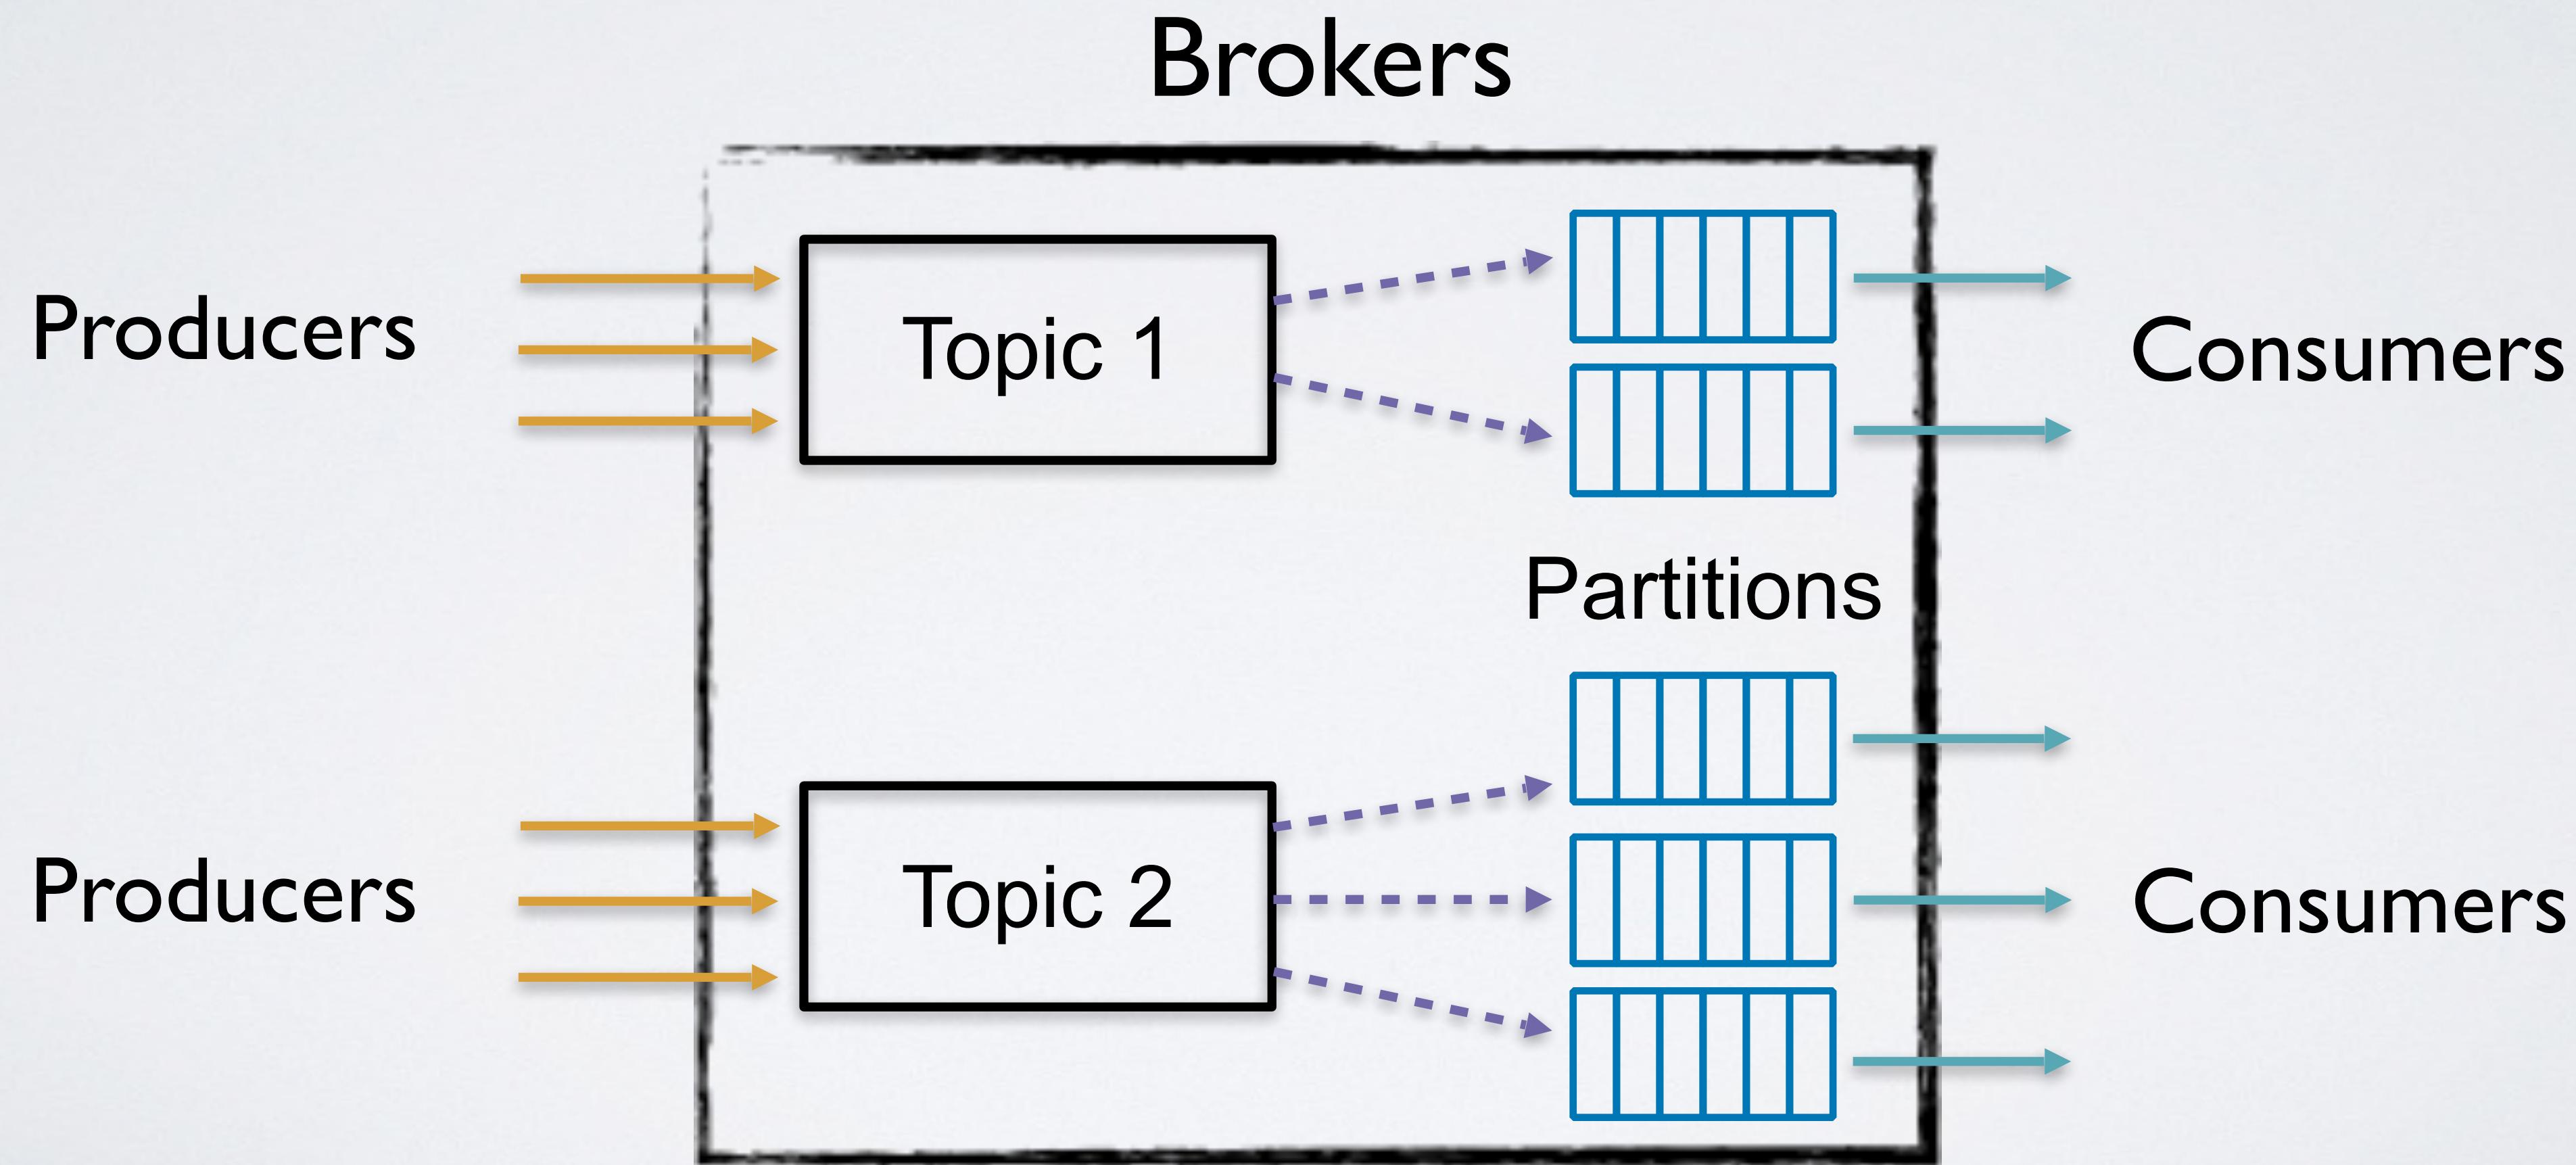
*If I have seen a little further it is
by standing on the shoulders of Giants.*

Isaac Newton

Kafka Concepts: the *Log*

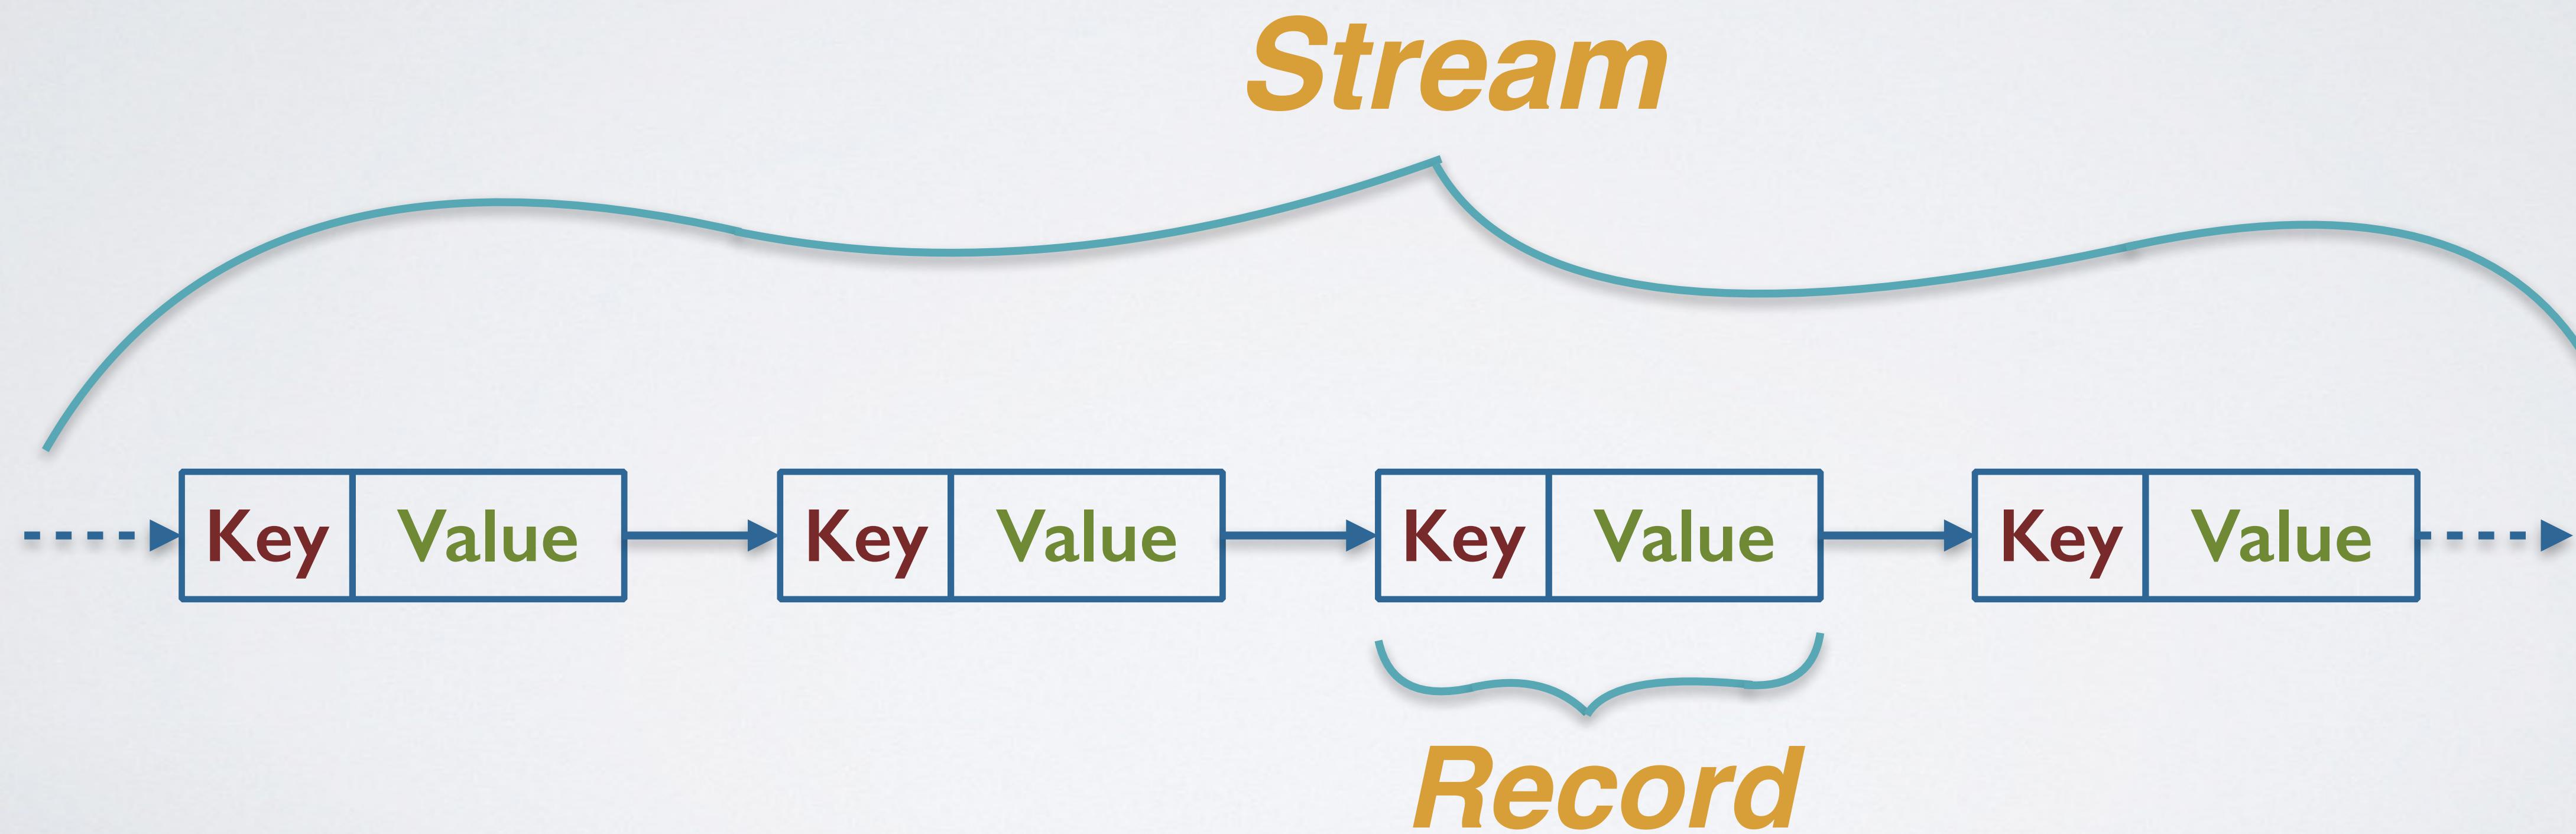


Kafka Concepts: the *Log*



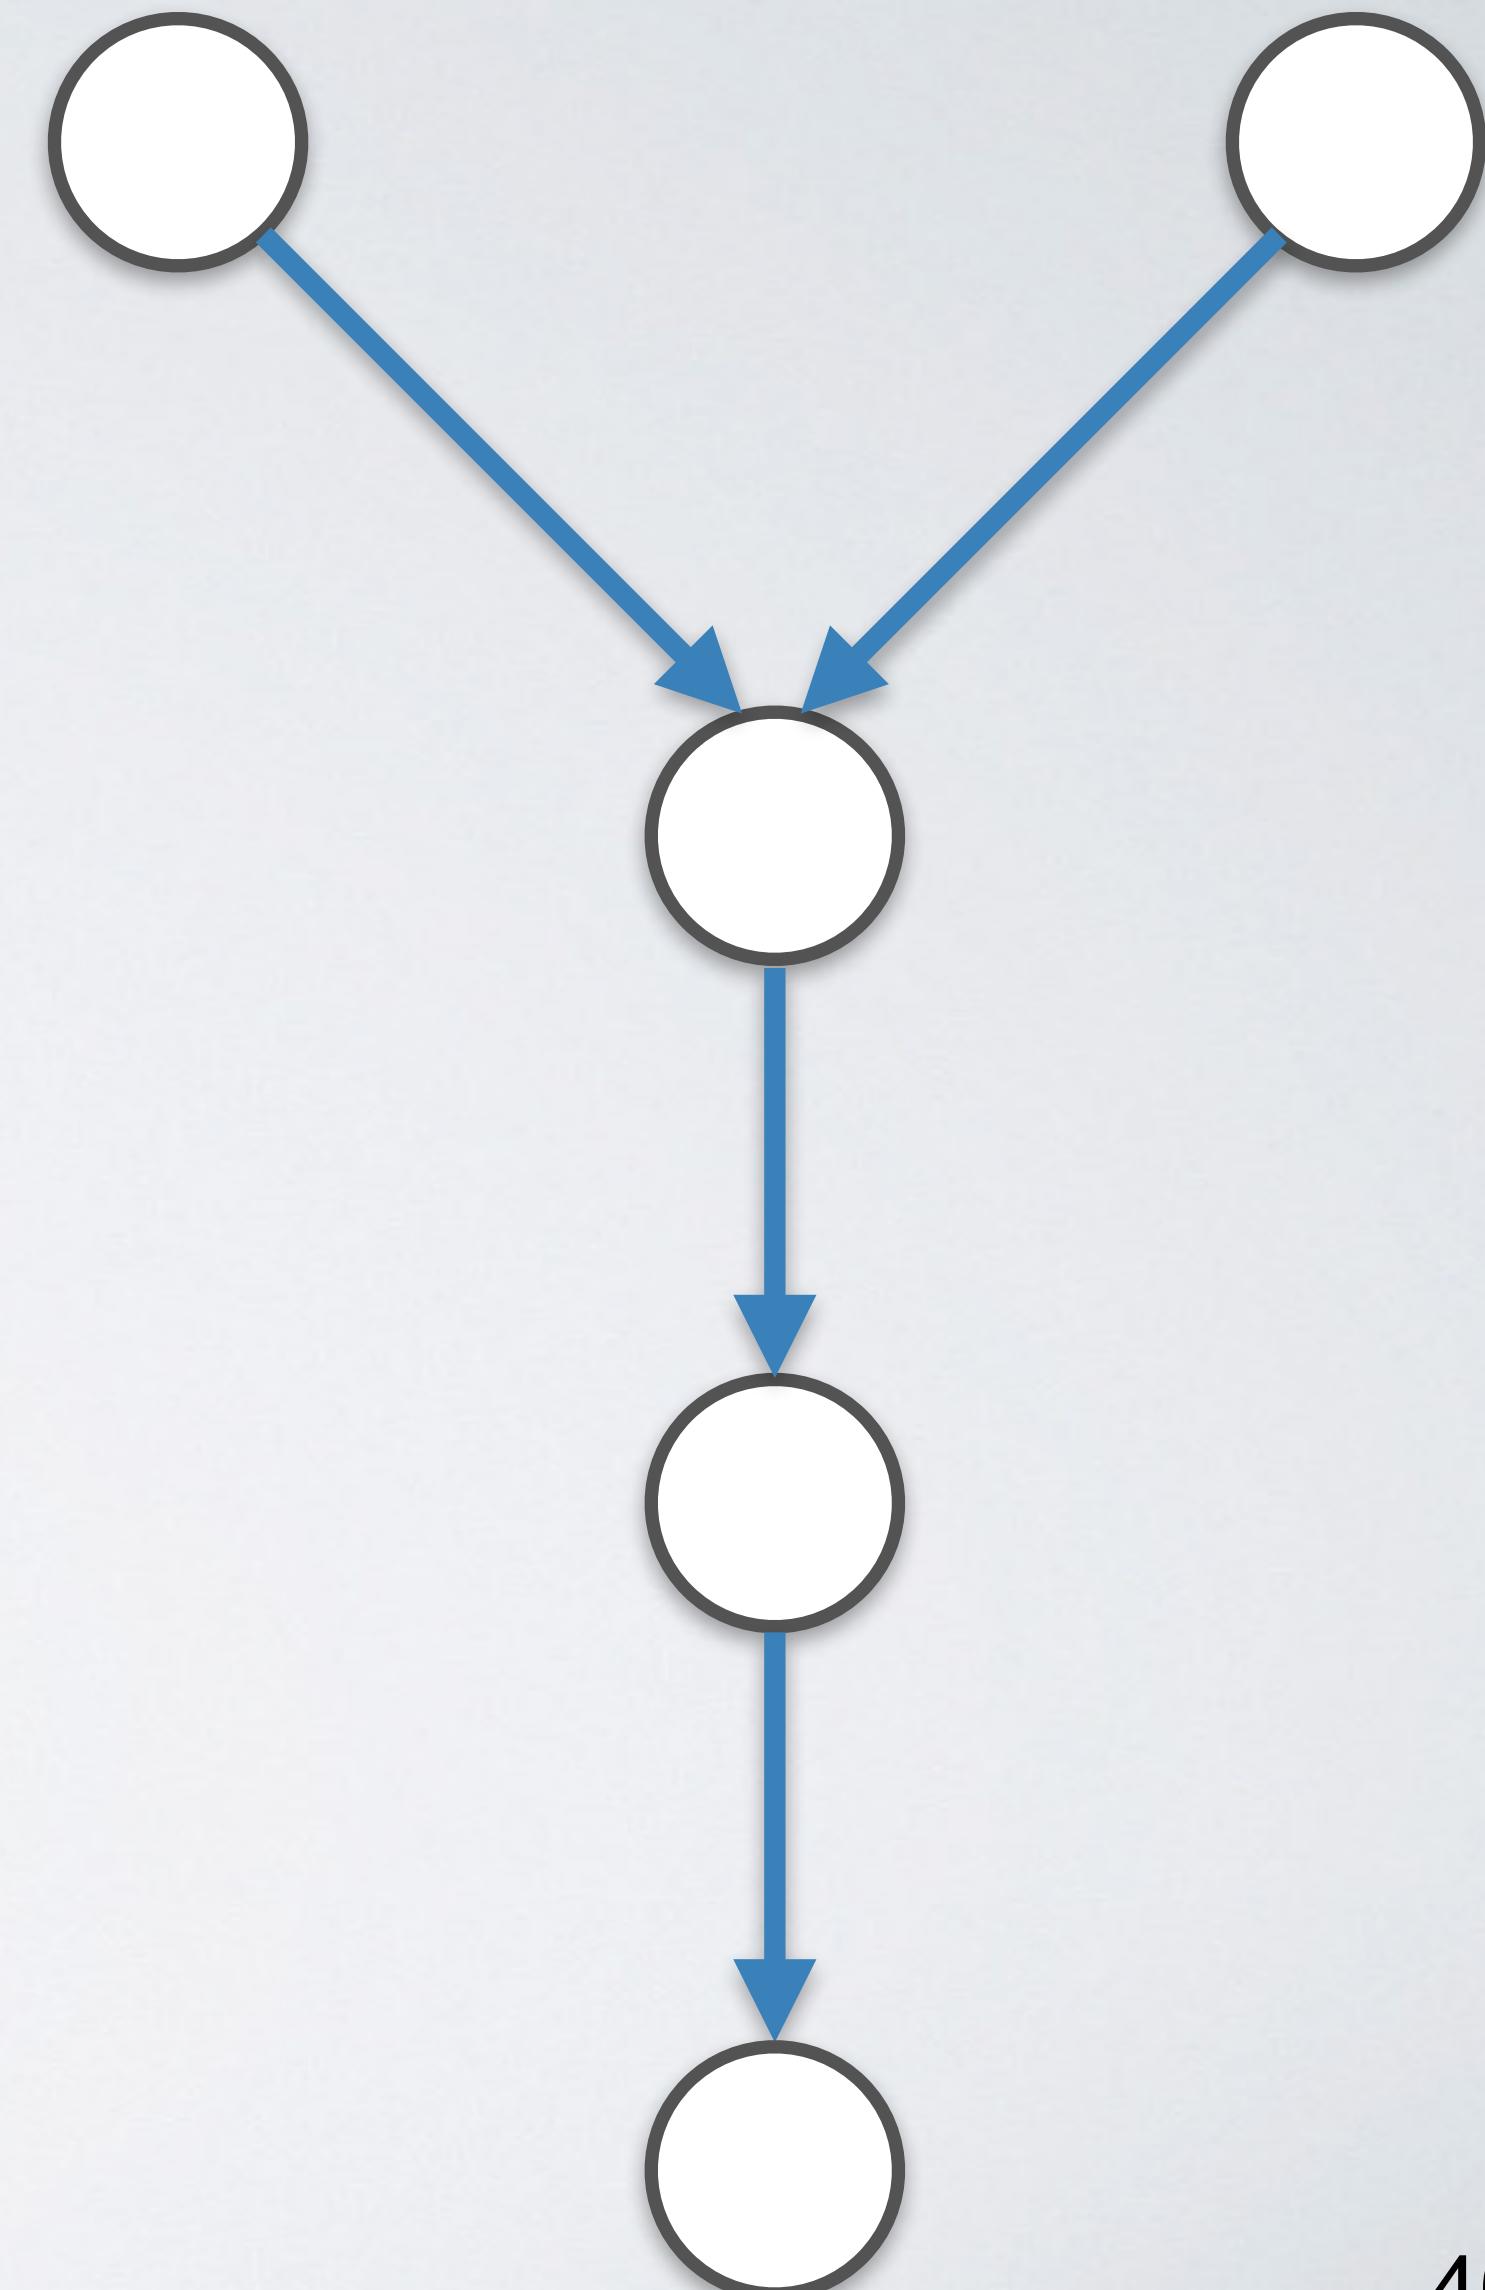
Kafka Streams: Key Concepts

Stream and Records

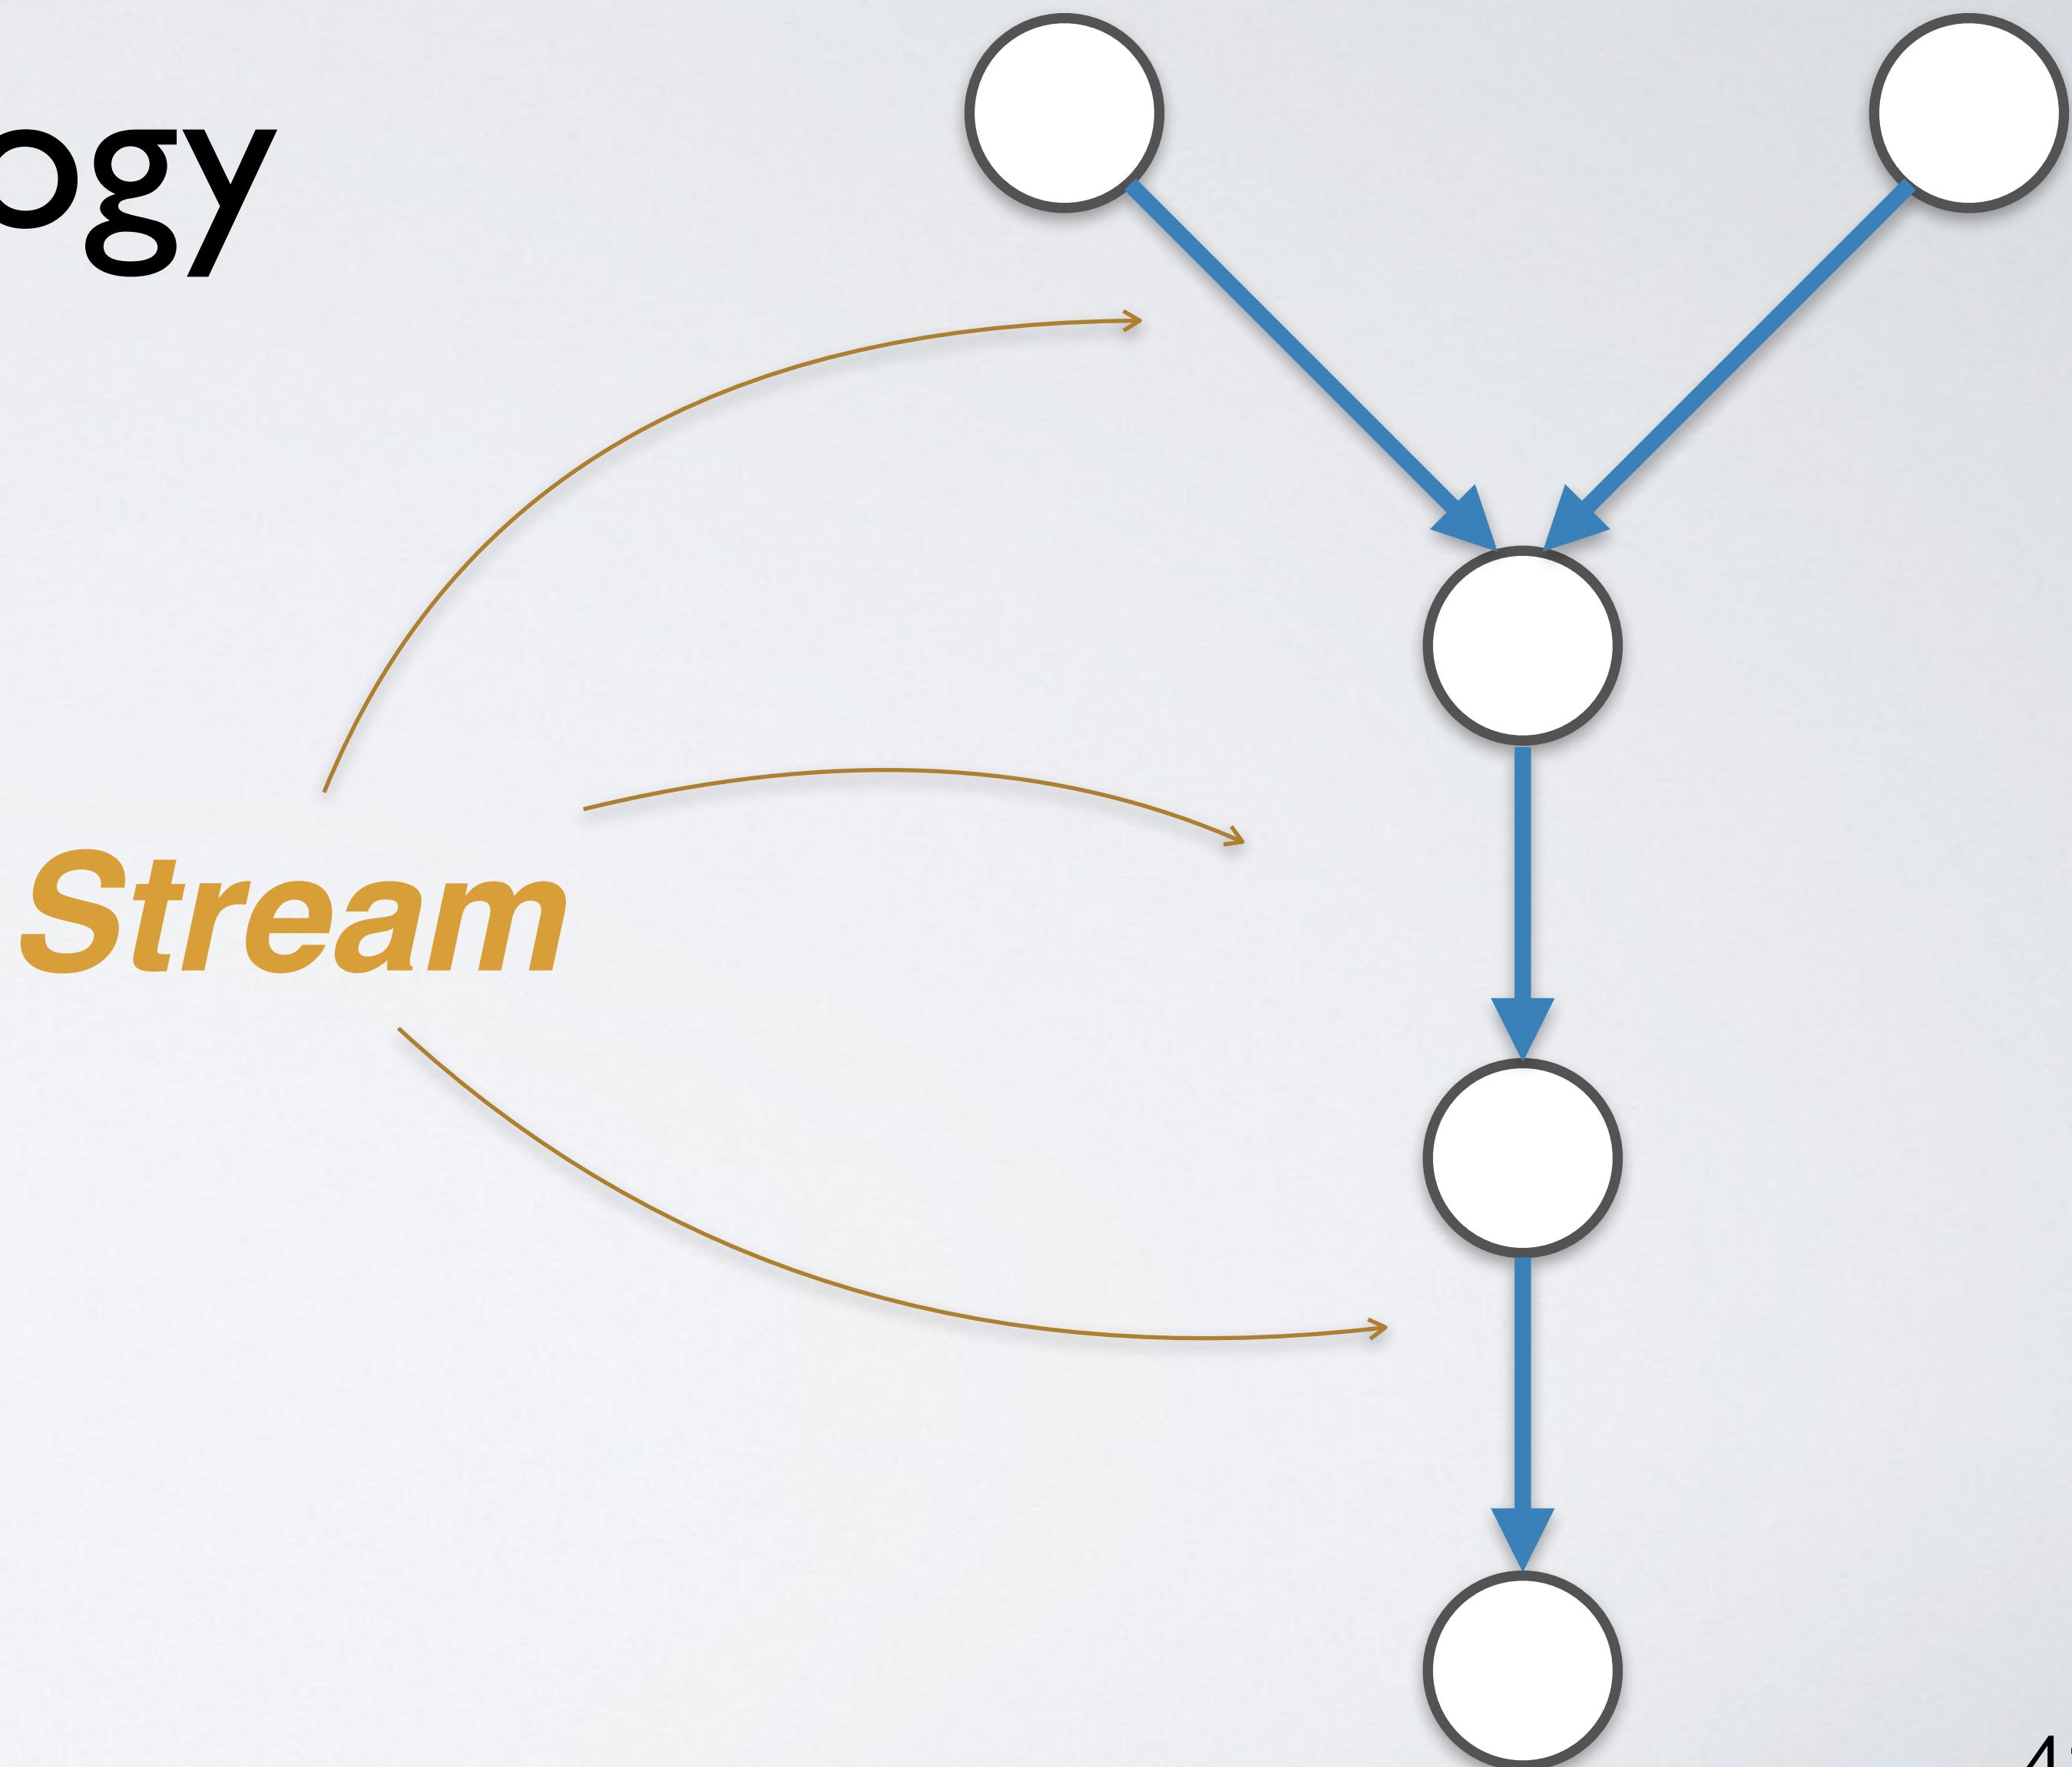


Processor Topology

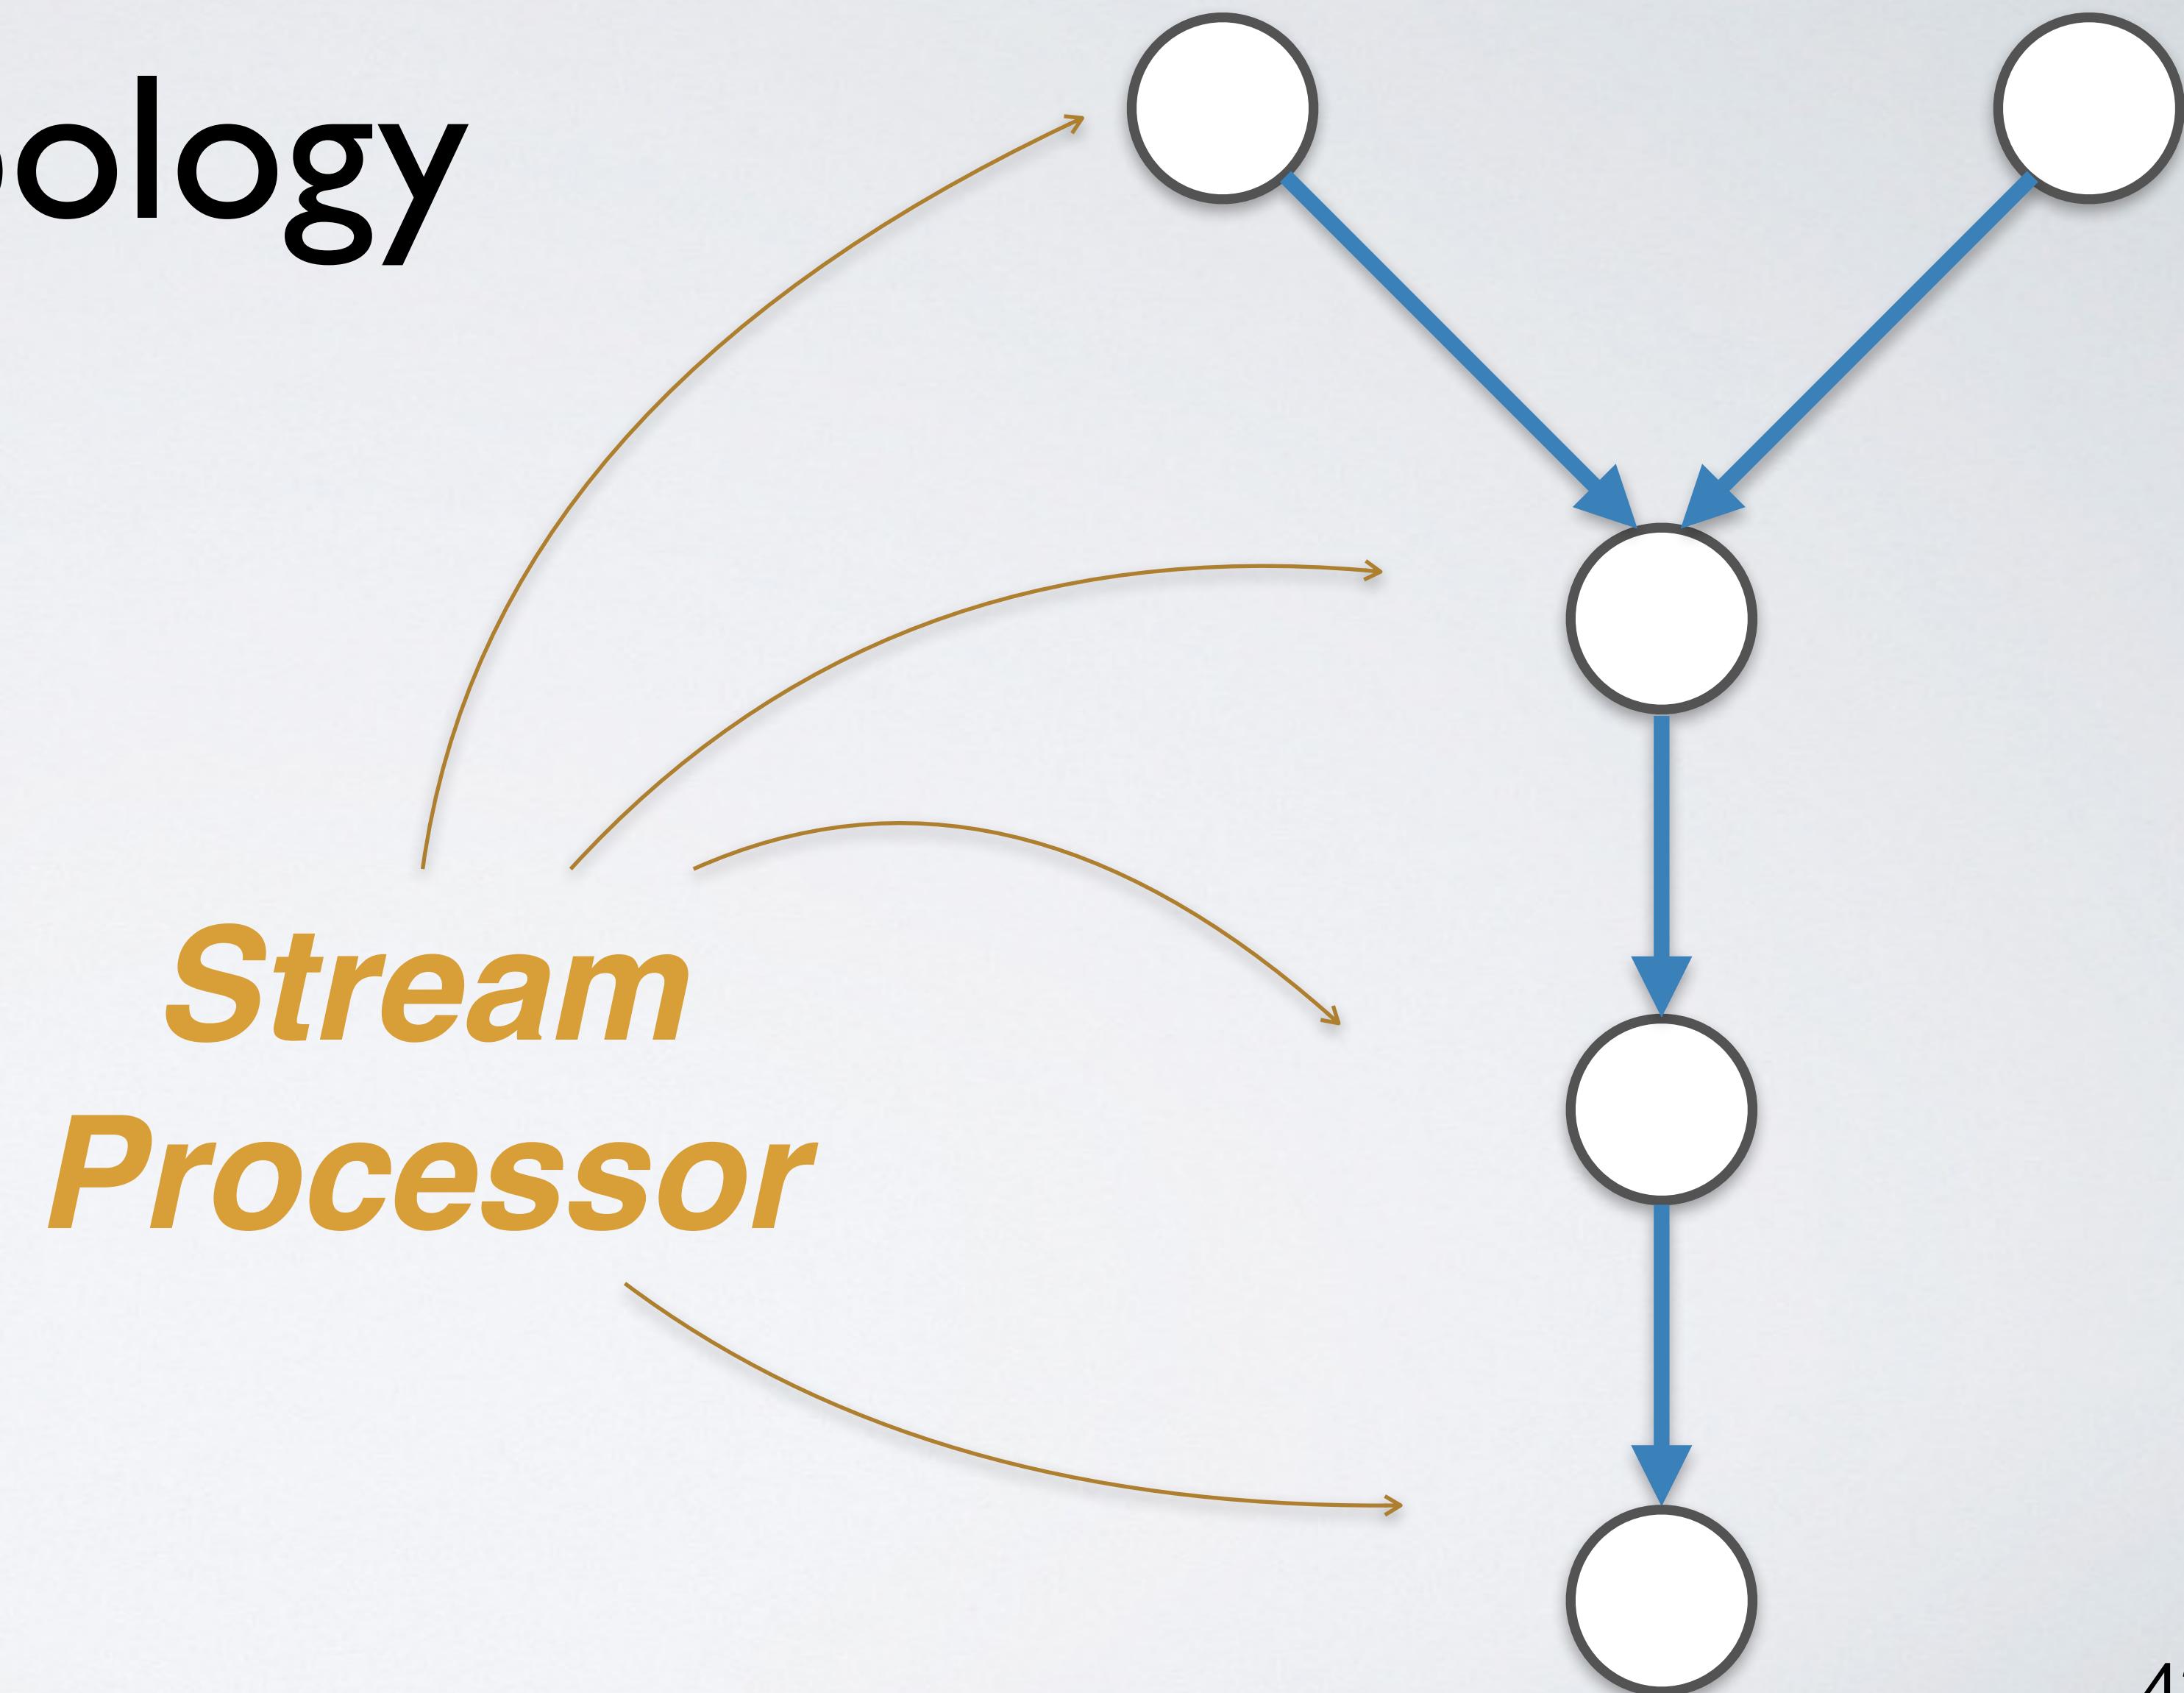
```
KStream<..> stream1 = builder.stream("topic1");
KStream<..> stream2 = builder.stream("topic2");
KStream<..> joined = stream1.leftJoin(stream2, ...);
KTable<..> aggregated = joined.aggregateByKey(...);
aggregated.to("topic3");
```



Processor Topology

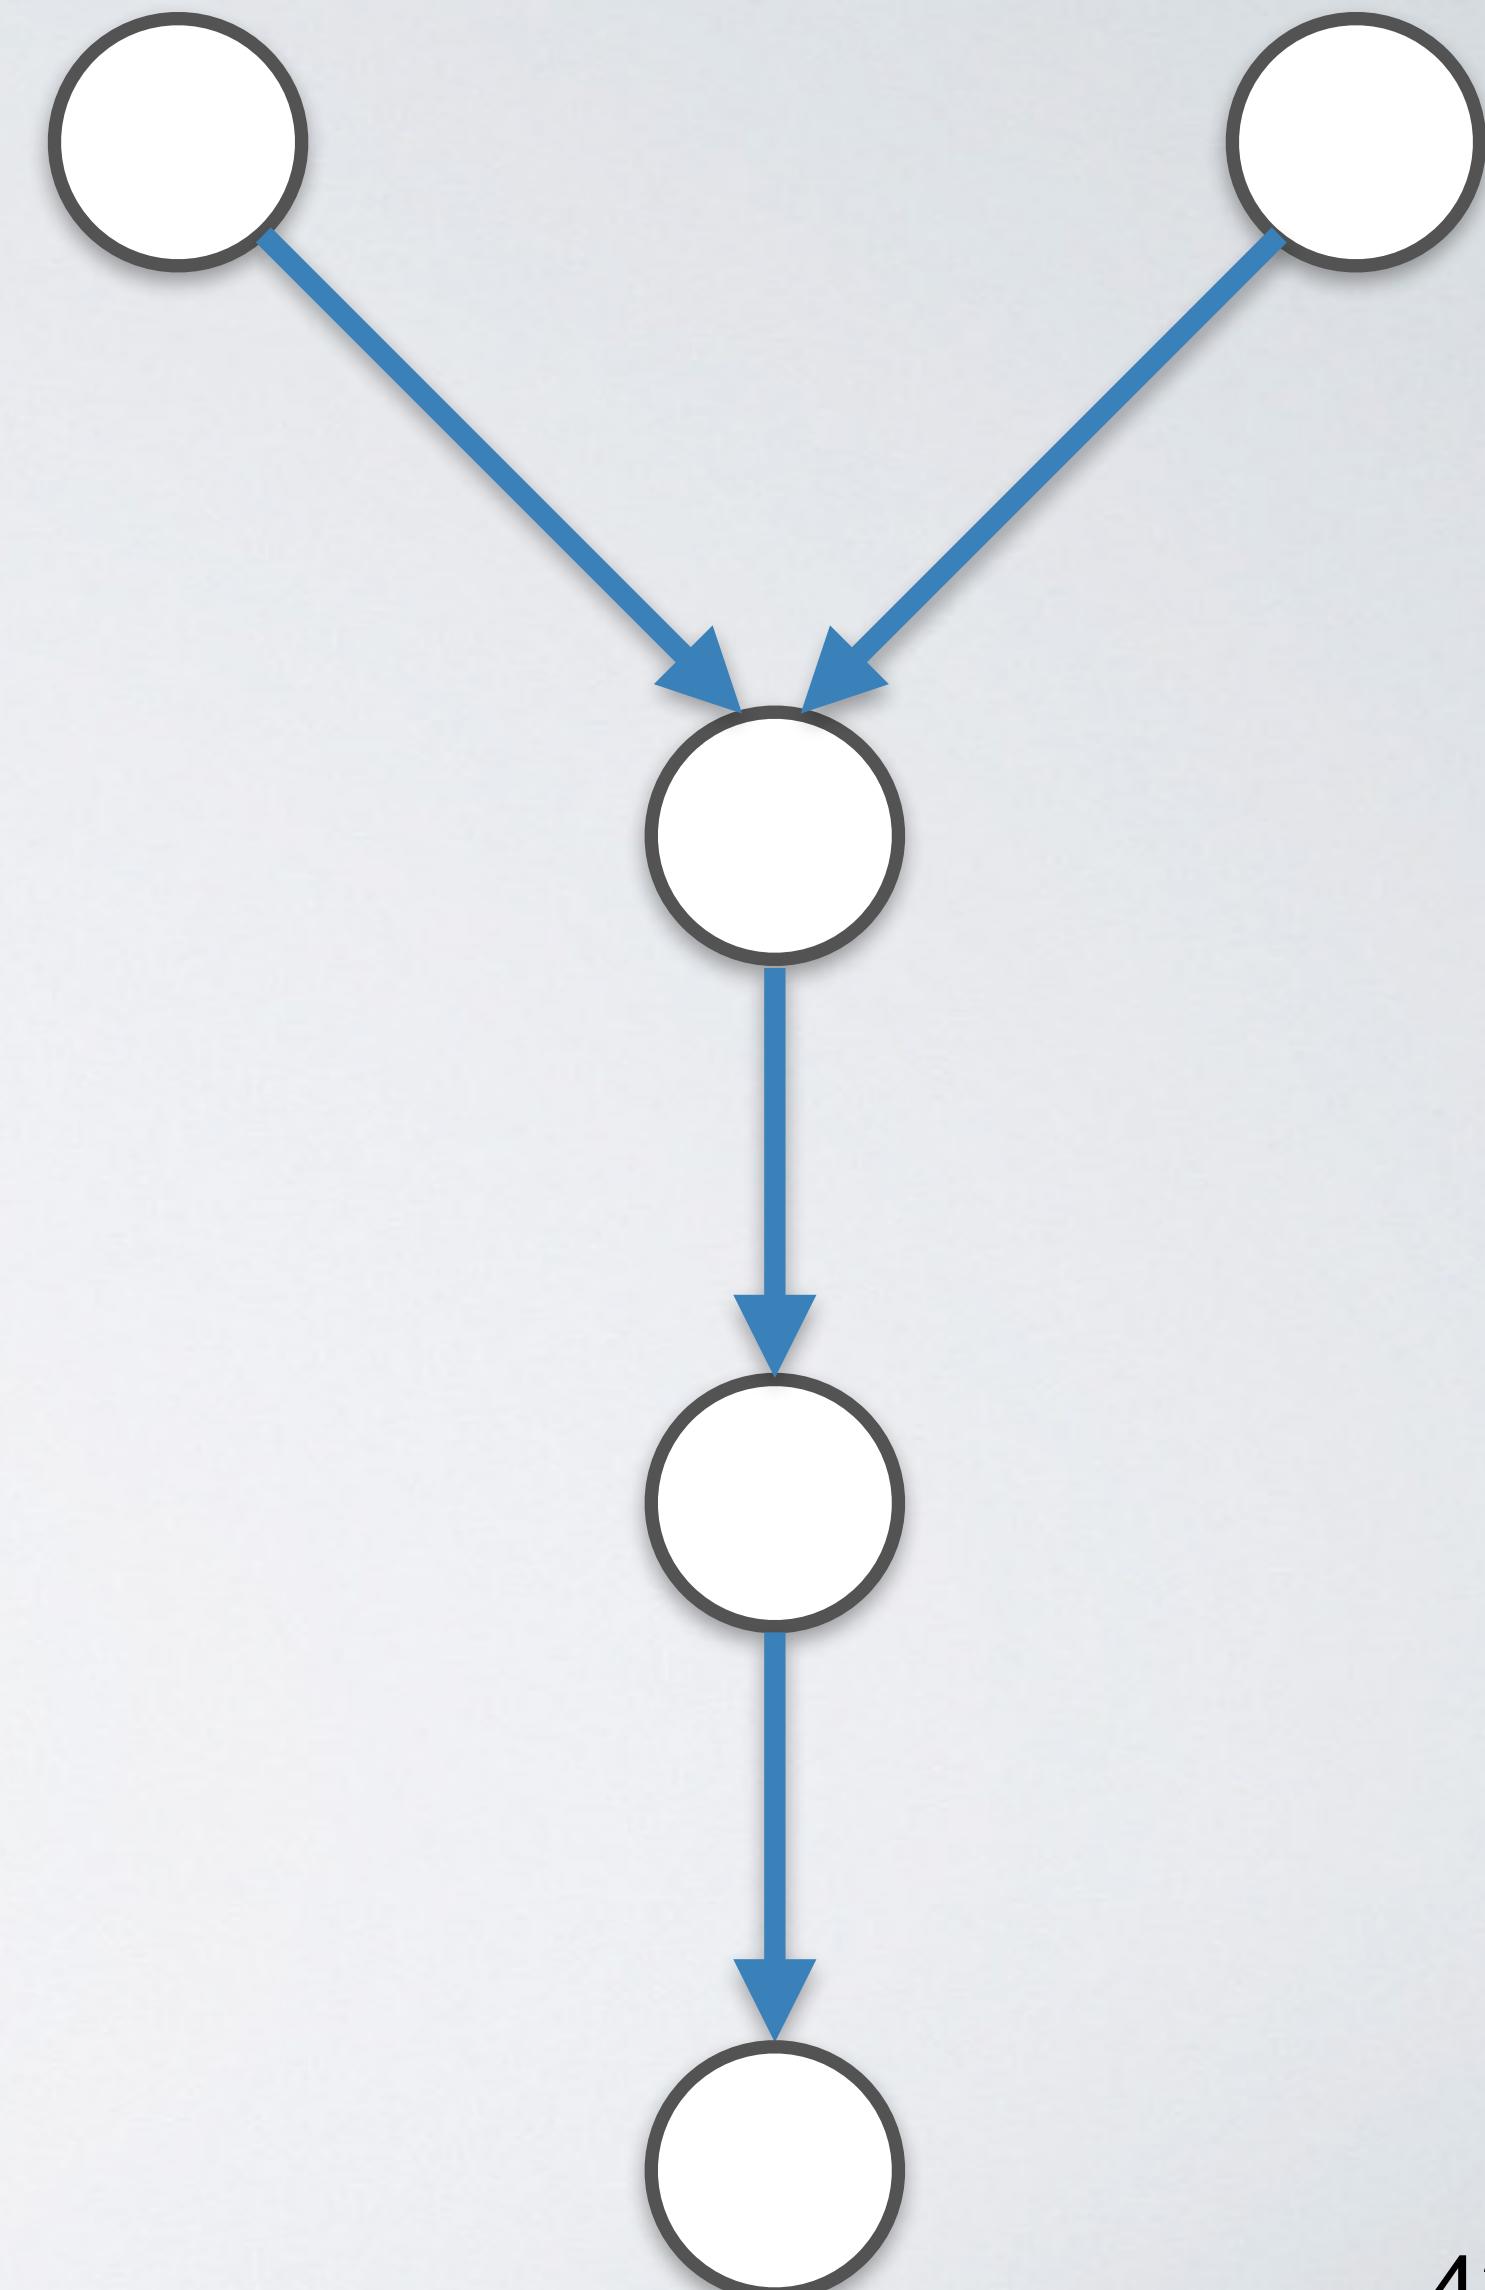


Processor Topology



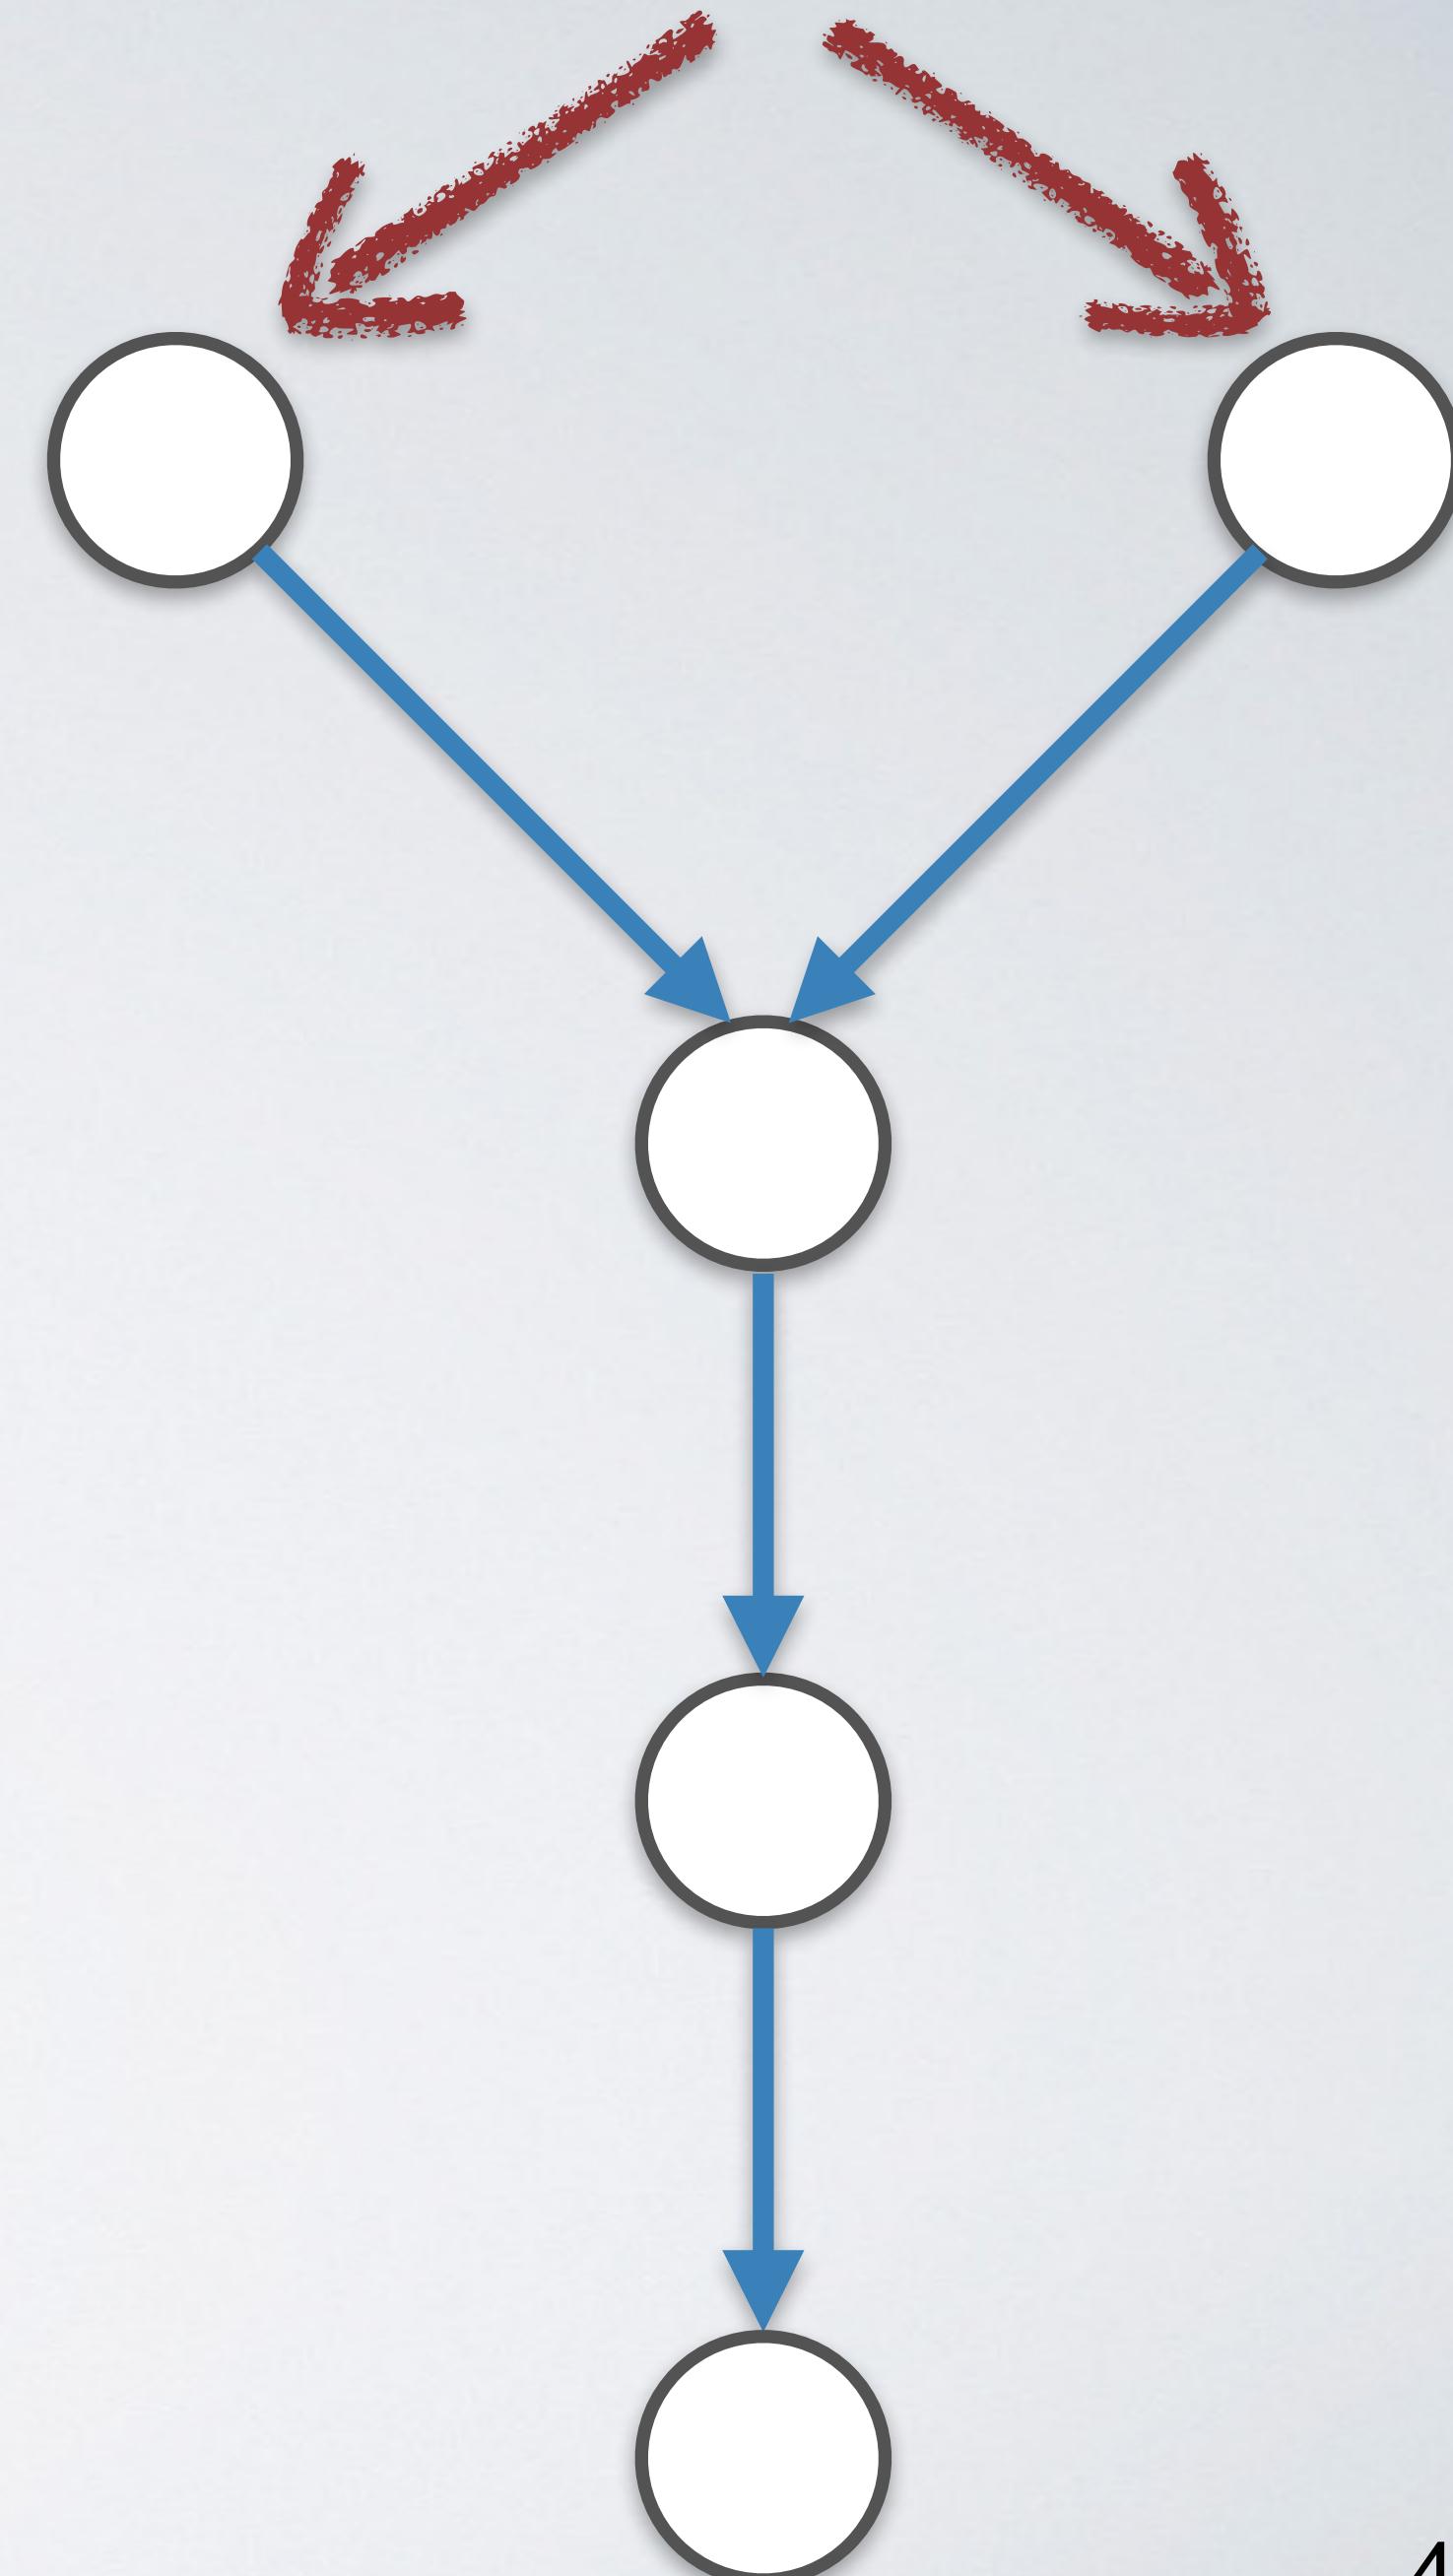
Processor Topology

```
KStream<..> stream1 = builder.stream("topic1");
KStream<..> stream2 = builder.stream("topic2");
KStream<..> joined = stream1.leftJoin(stream2, ...);
KTable<..> aggregated = joined.aggregateByKey(...);
aggregated.to("topic3");
```



Processor Topology

```
KStream<..> stream1 = builder.stream("topic1");
KStream<..> stream2 = builder.stream("topic2");
KStream<..> joined = stream1.leftJoin(stream2, ...);
KTable<..> aggregated = joined.aggregateByKey(...);
aggregated.to("topic3");
```



Processor Topology

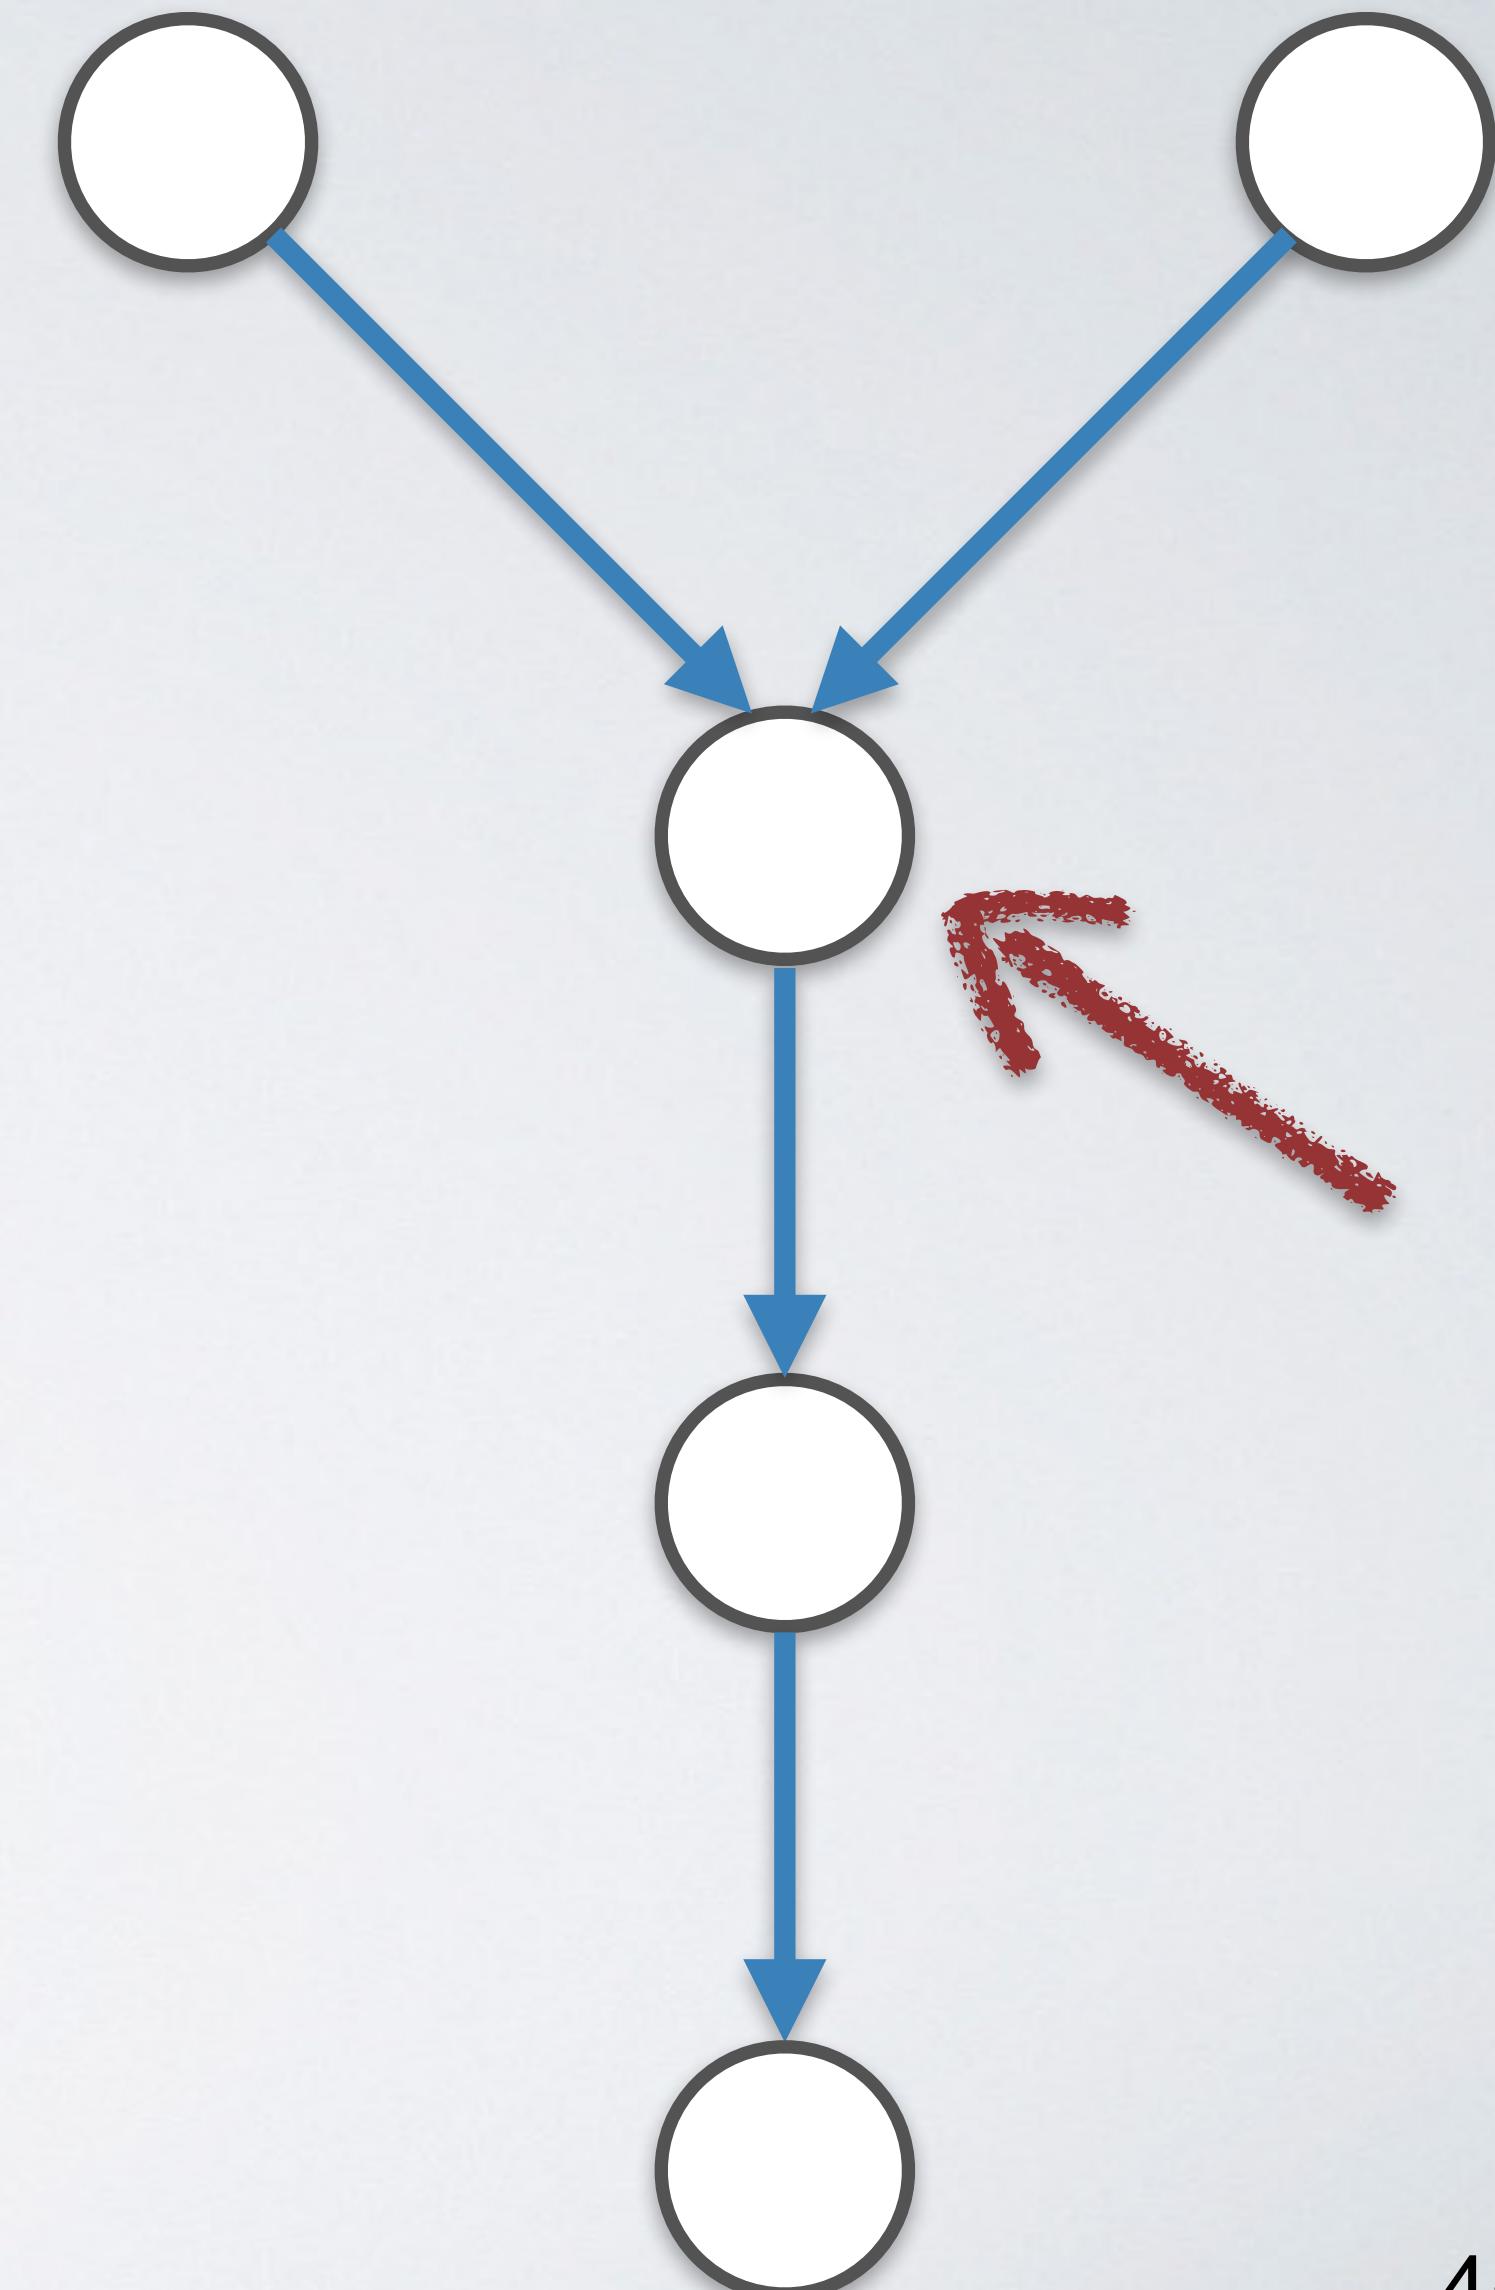
```
KStream<..> stream1 = builder.stream("topic1");
```

```
KStream<..> stream2 = builder.stream("topic2");
```

```
KStream<..> joined = stream1.leftJoin(stream2, ...);
```

```
KTable<..> aggregated = joined.aggregateByKey(...);
```

```
aggregated.to("topic3");
```



Processor Topology

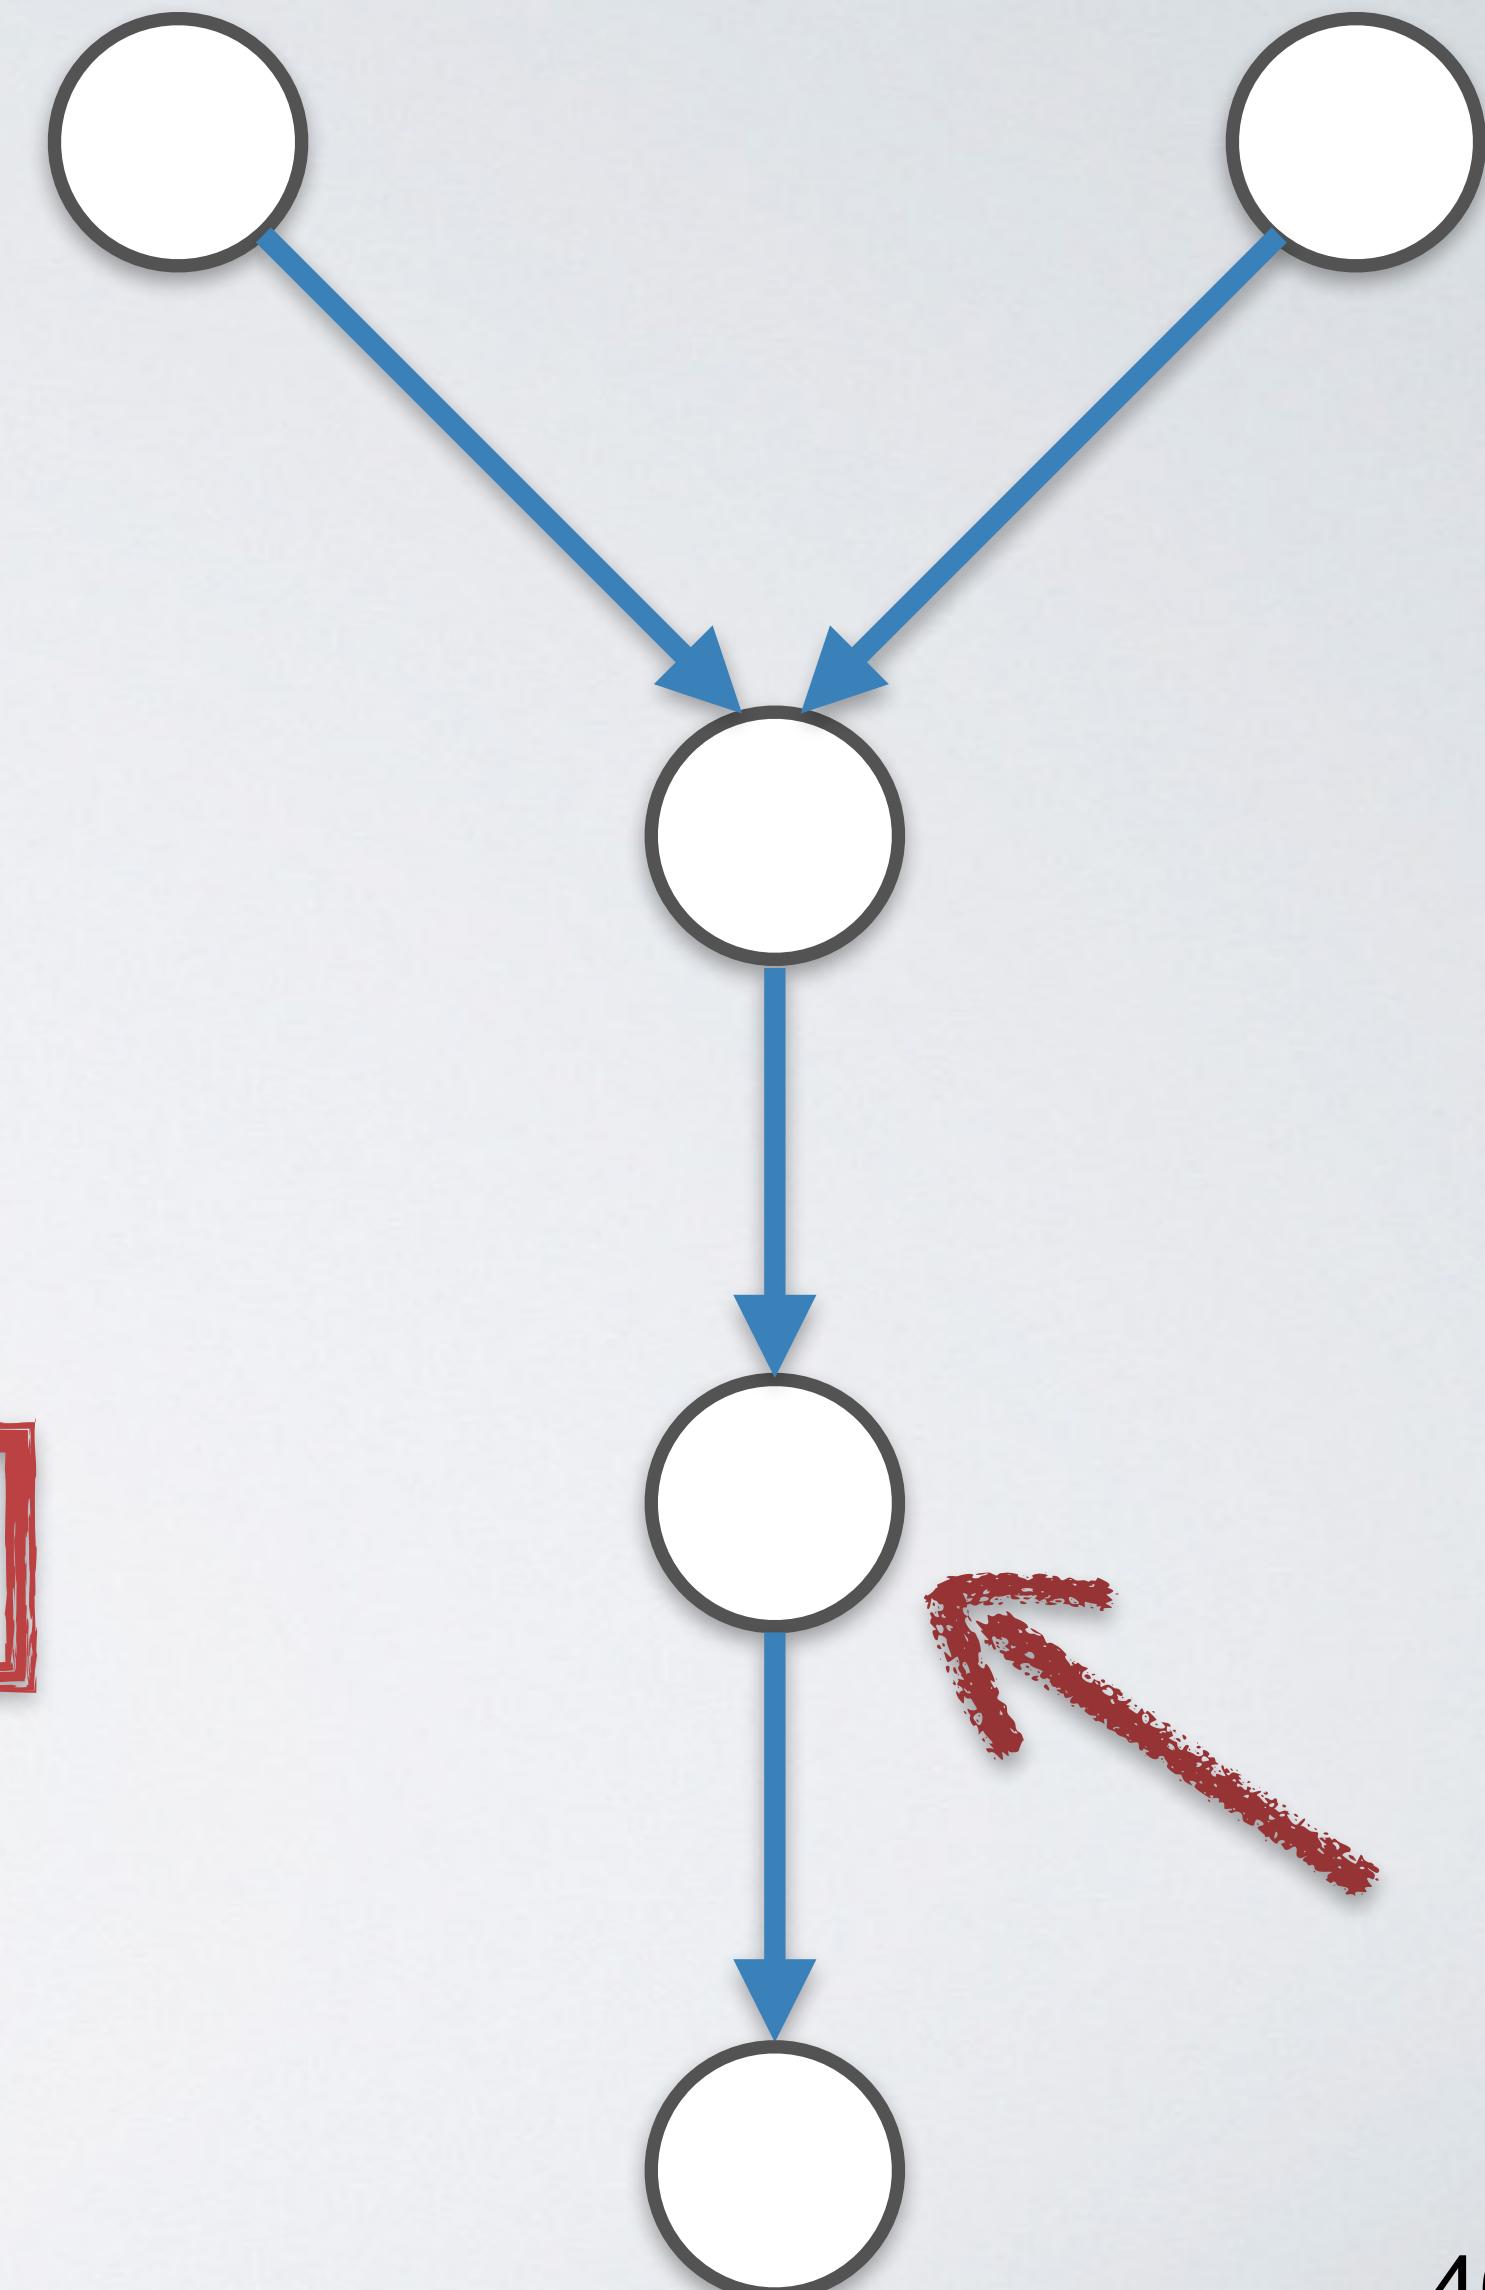
```
KStream<..> stream1 = builder.stream("topic1");
```

```
KStream<..> stream2 = builder.stream("topic2");
```

```
KStream<..> joined = stream1.leftJoin(stream2, ...);
```

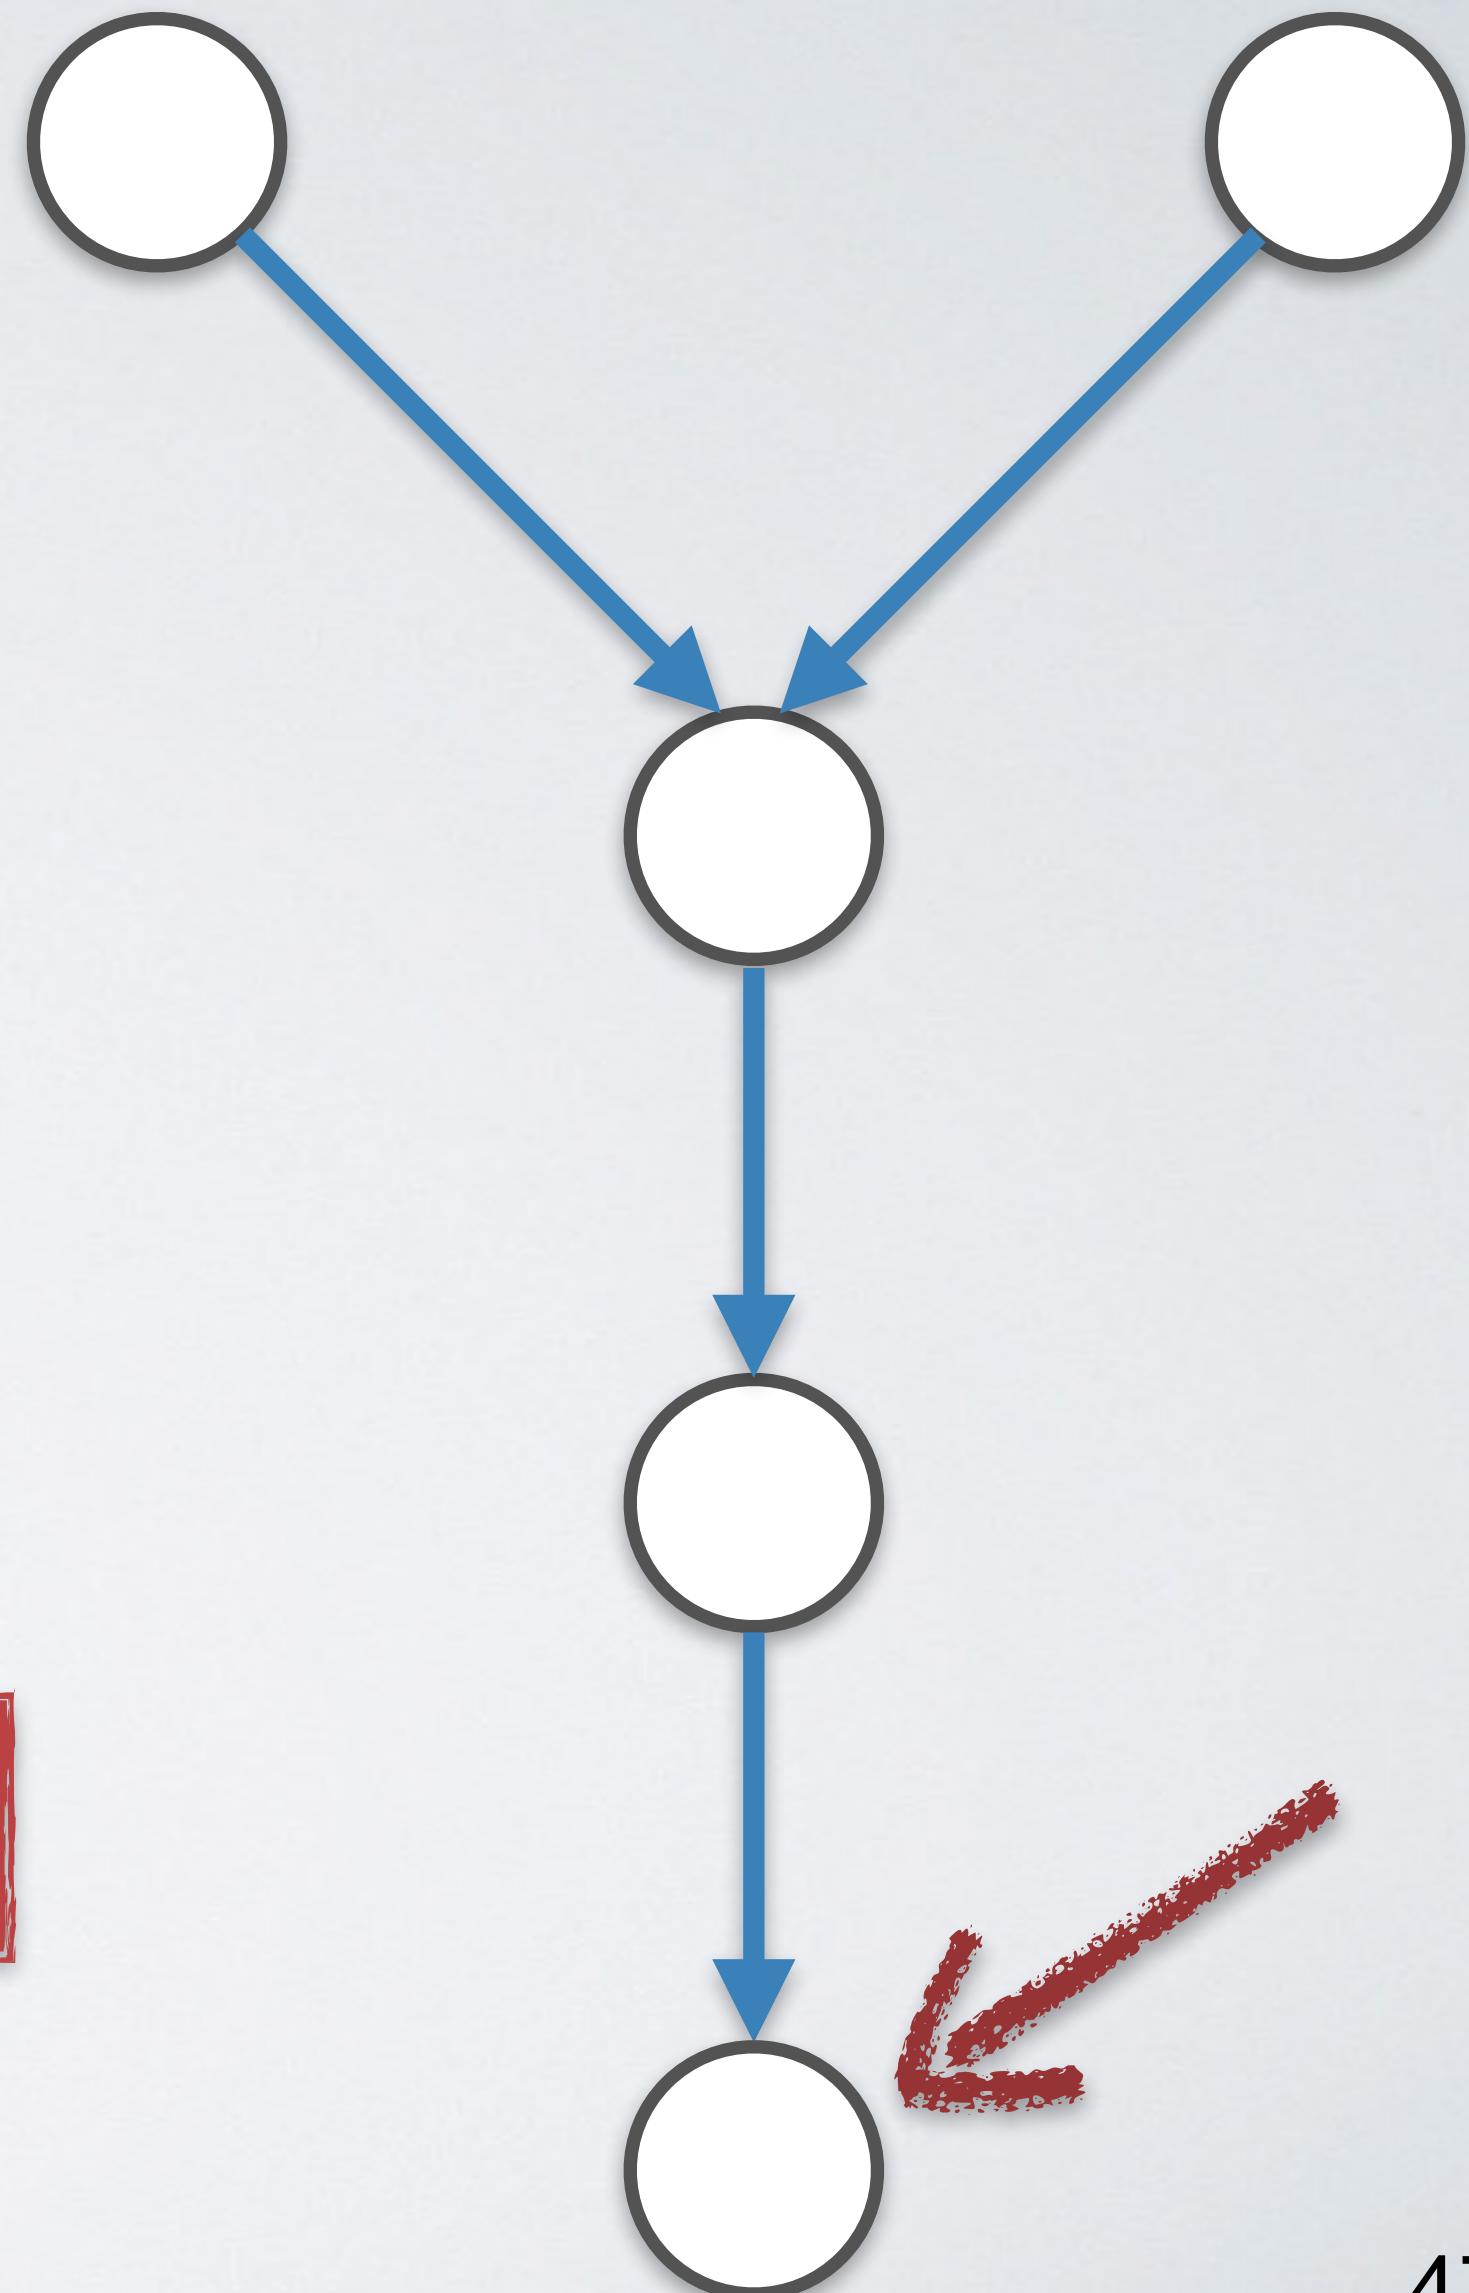
```
KTable<..> aggregated = joined.aggregateByKey(...);
```

```
aggregated.to("topic3");
```



Processor Topology

```
KStream<..> stream1 = builder.stream("topic1");
KStream<..> stream2 = builder.stream("topic2");
KStream<..> joined = stream1.leftJoin(stream2, ...);
KTable<..> aggregated = joined.aggregateByKey(...);
aggregated.to("topic3");
```



Processor Topology

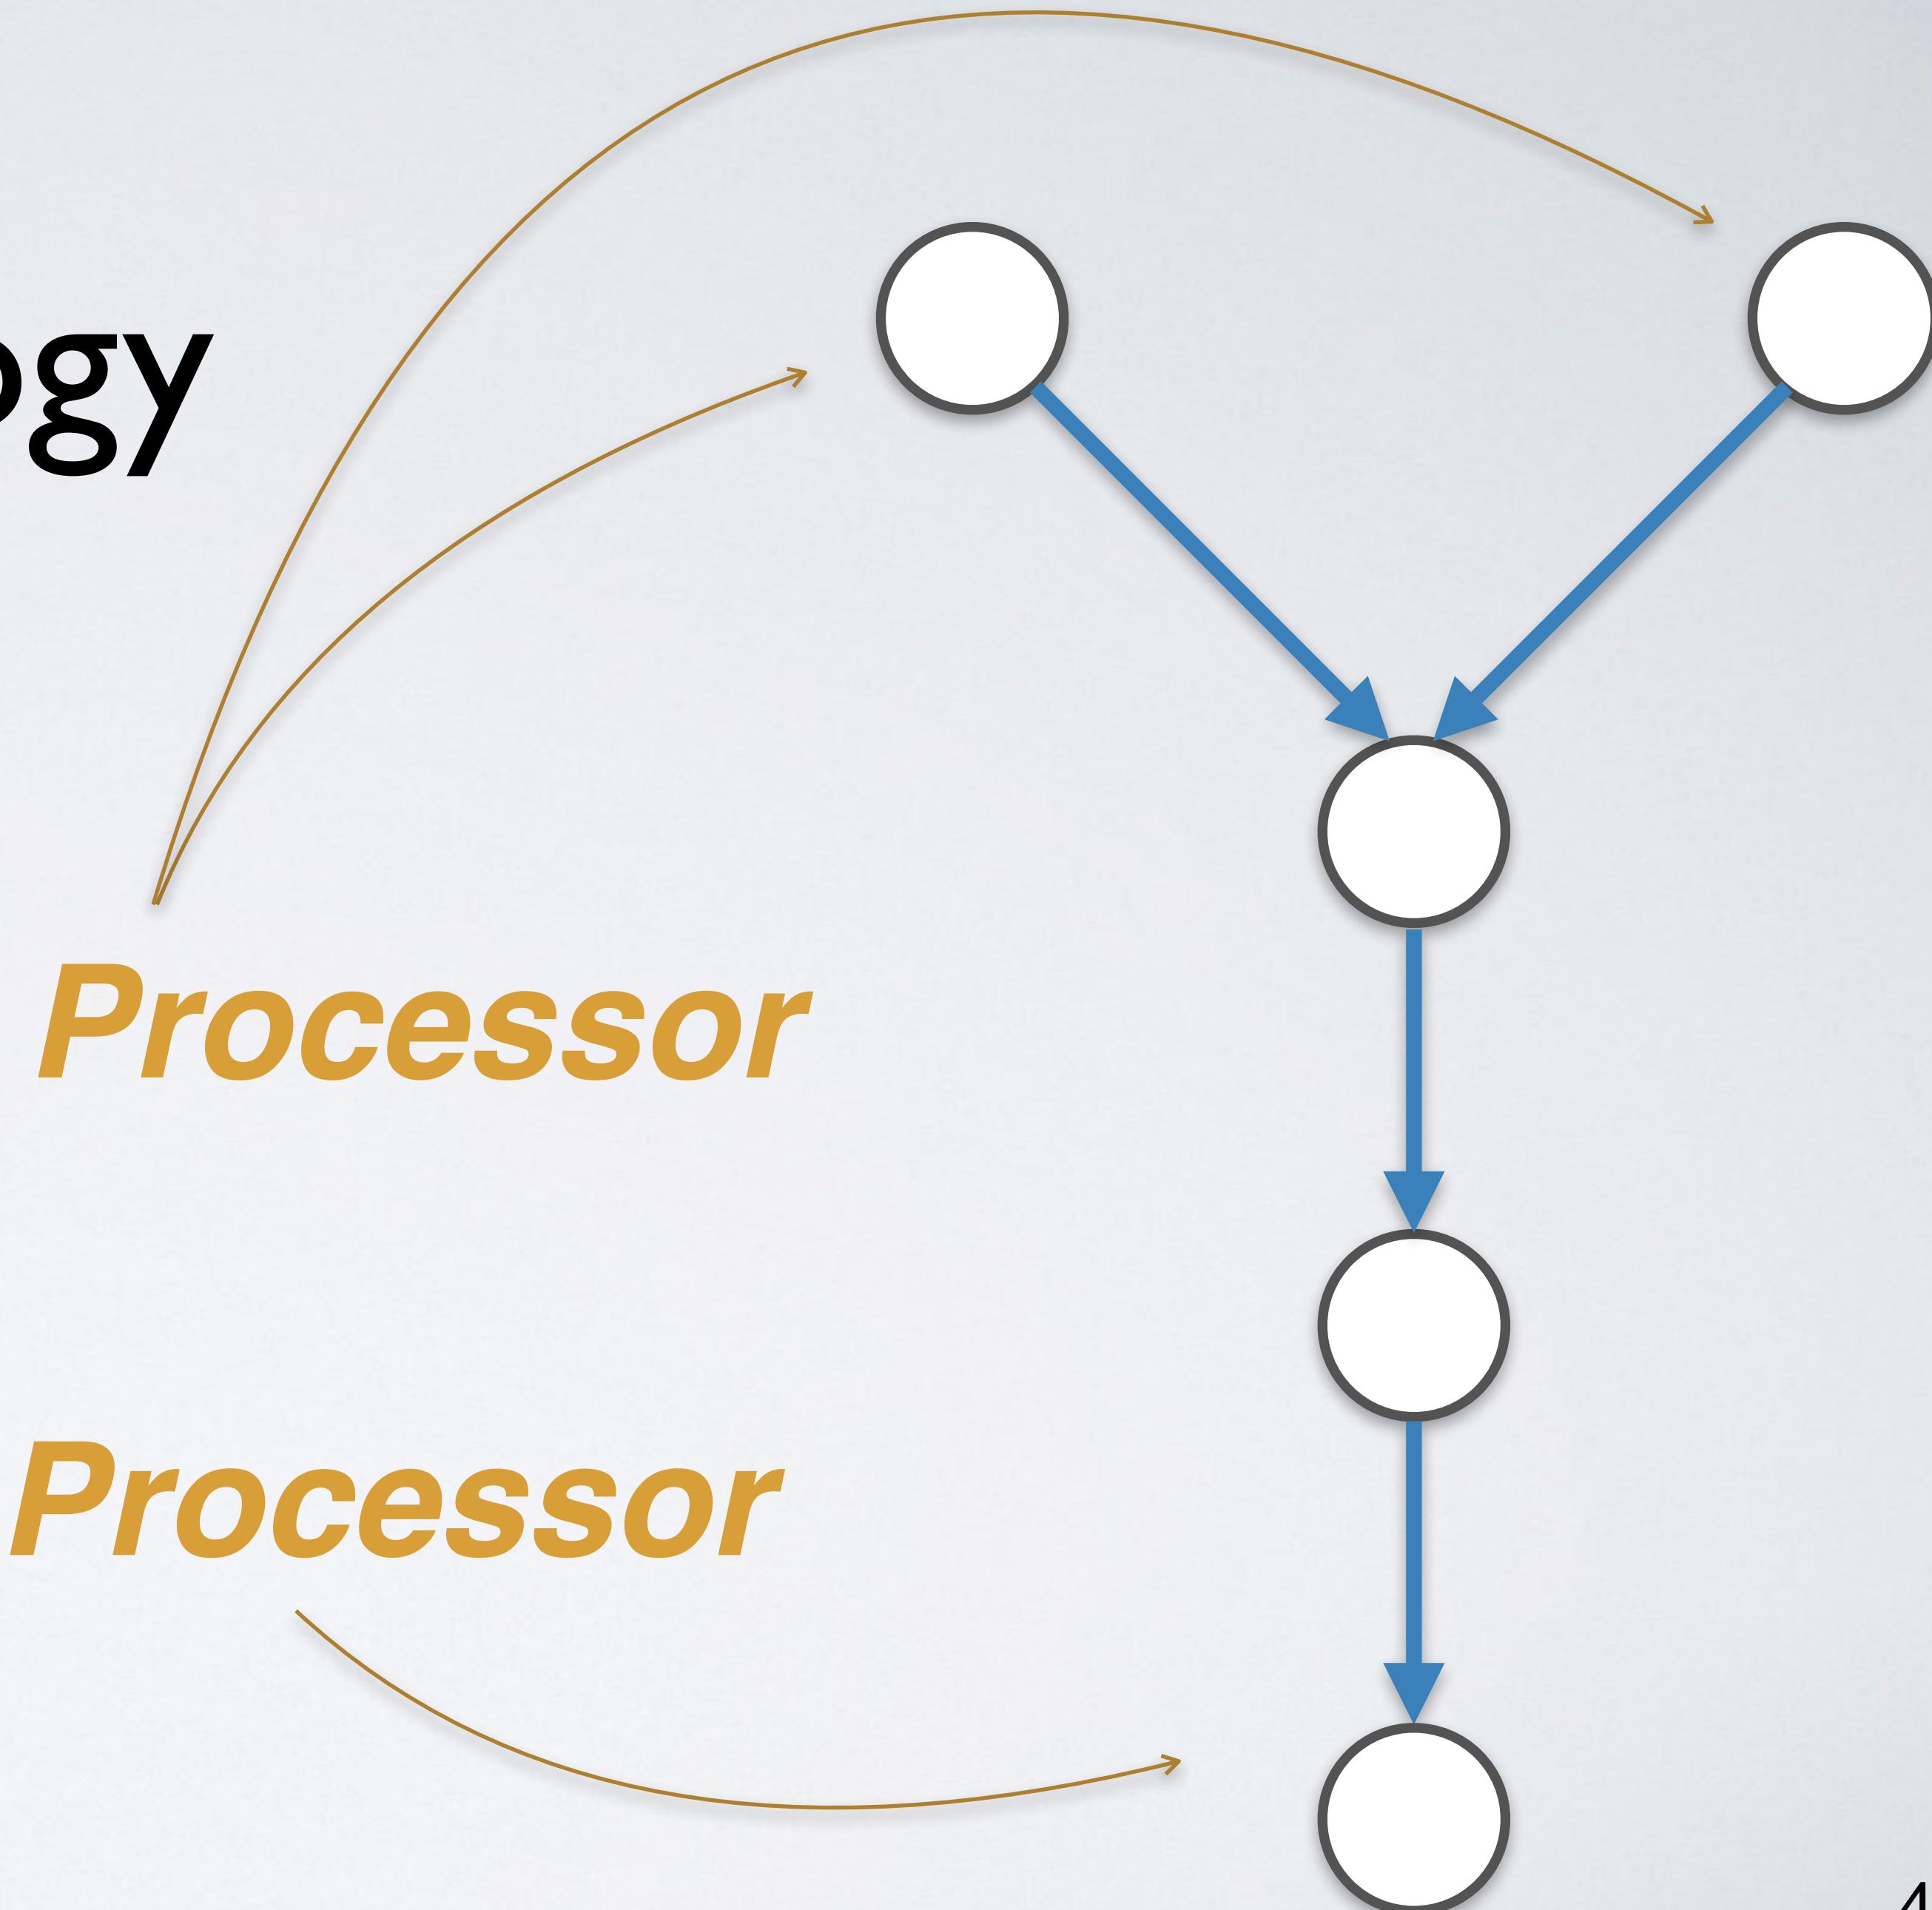
Source Processor

```
KStream<..> stream1 = builder.stream()
```

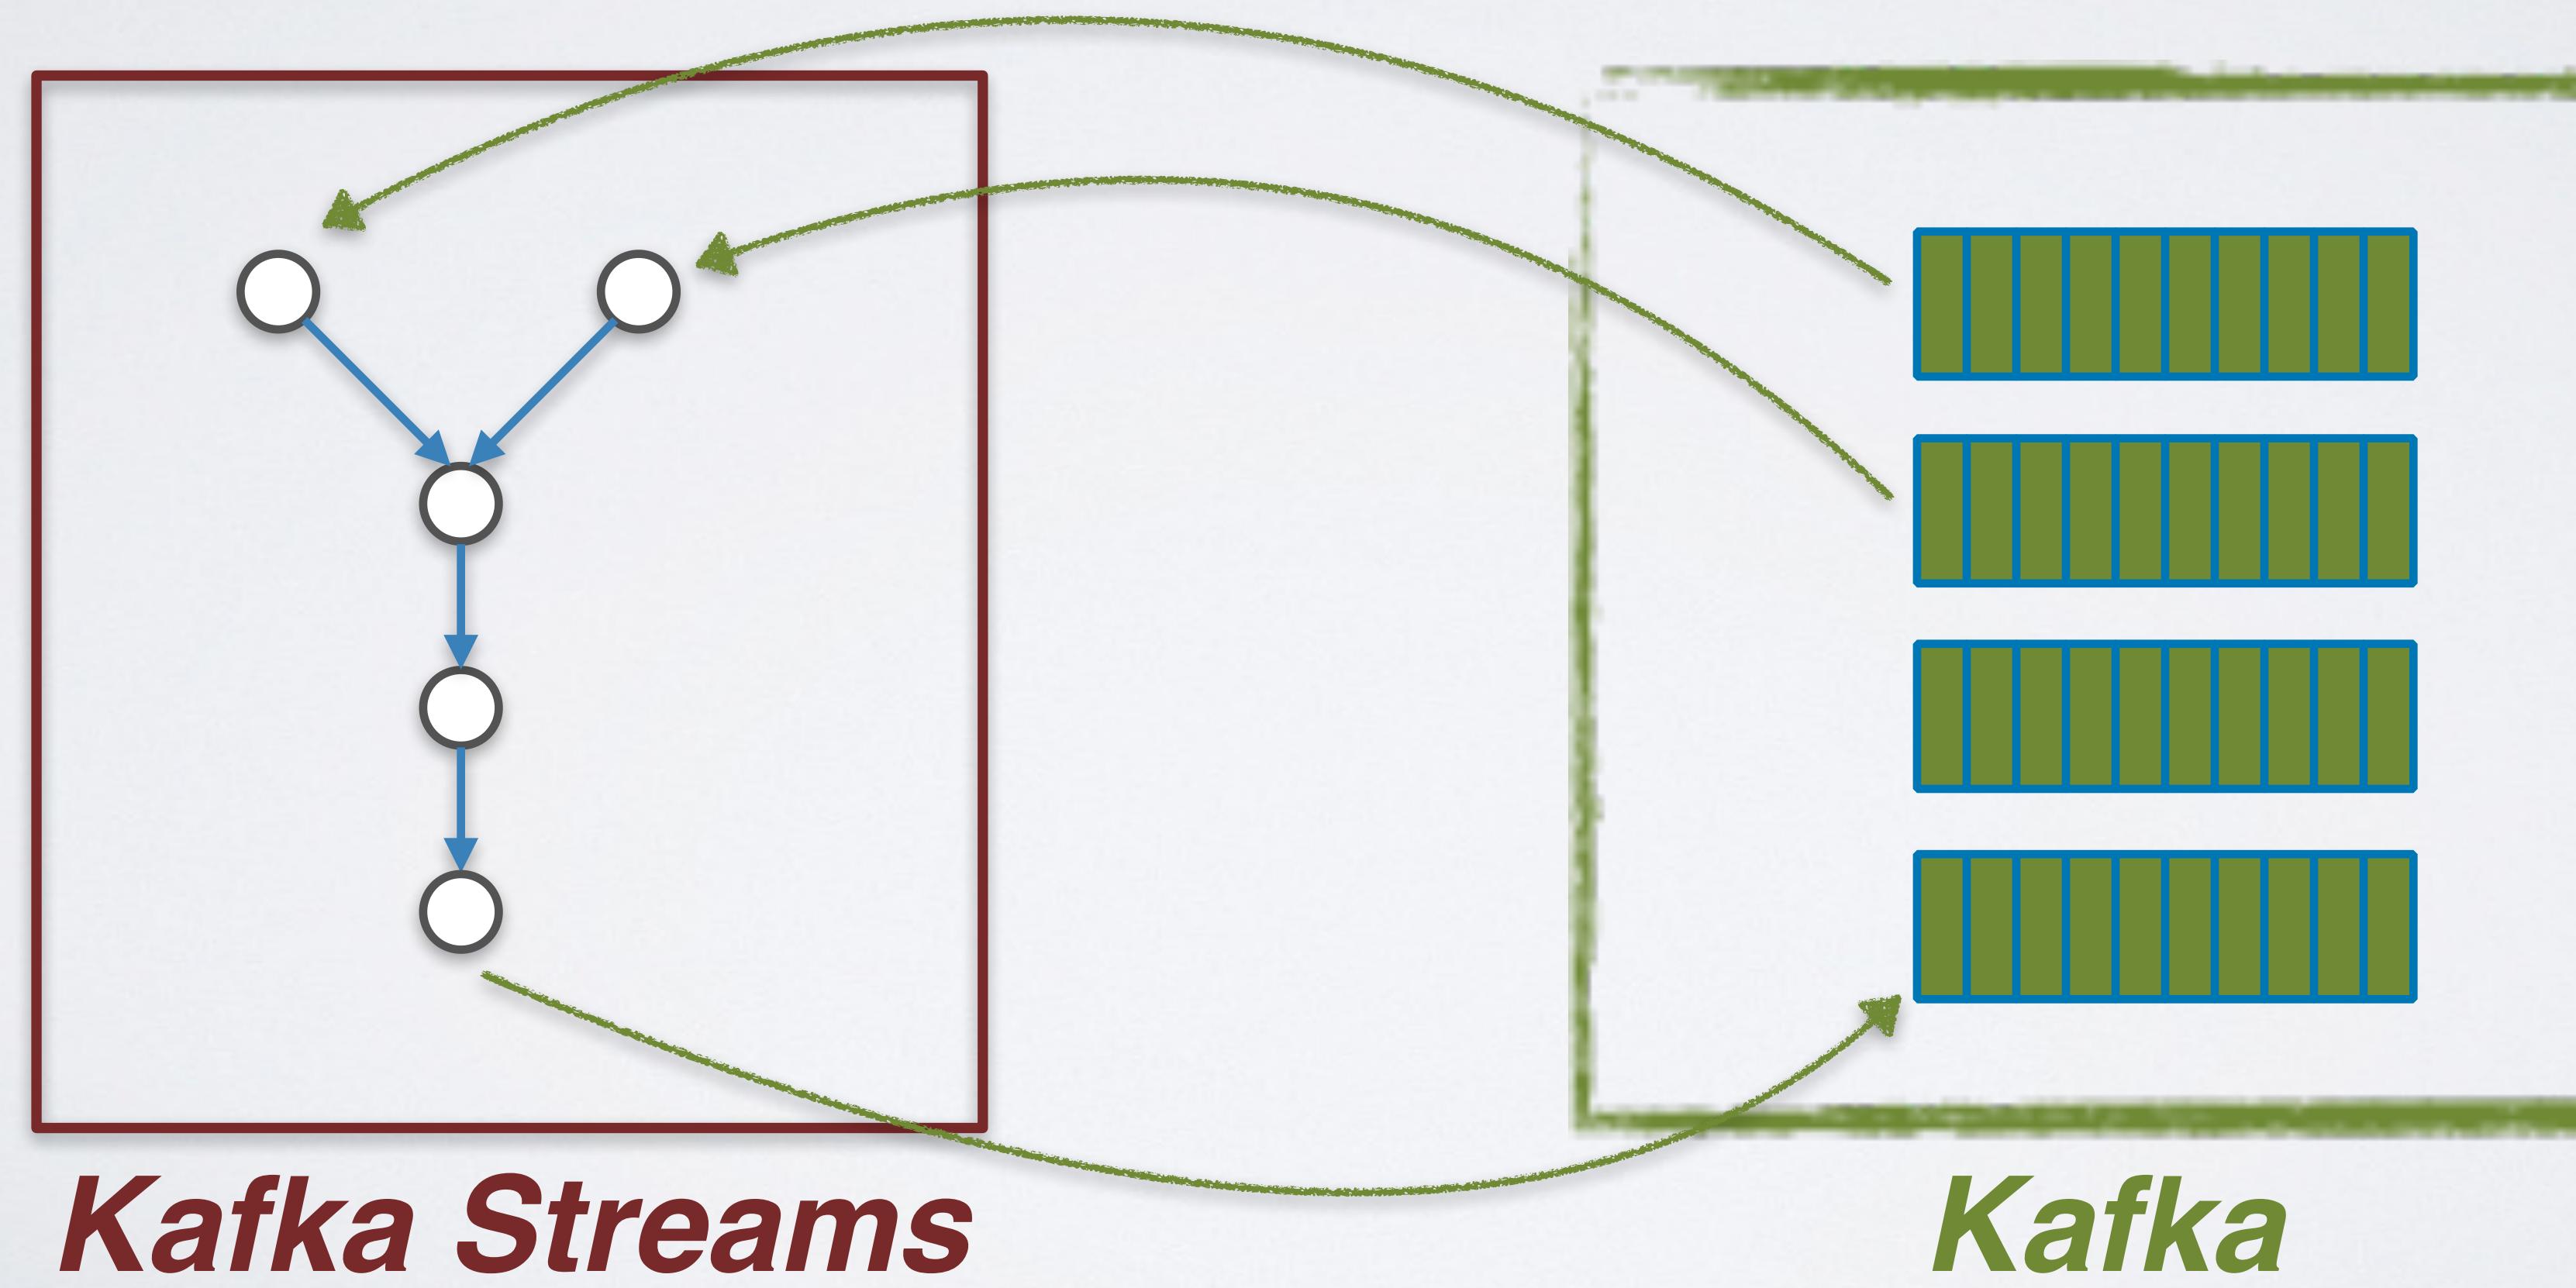
```
KStream<..> stream2 = builder.stream()
```

Sink Processor

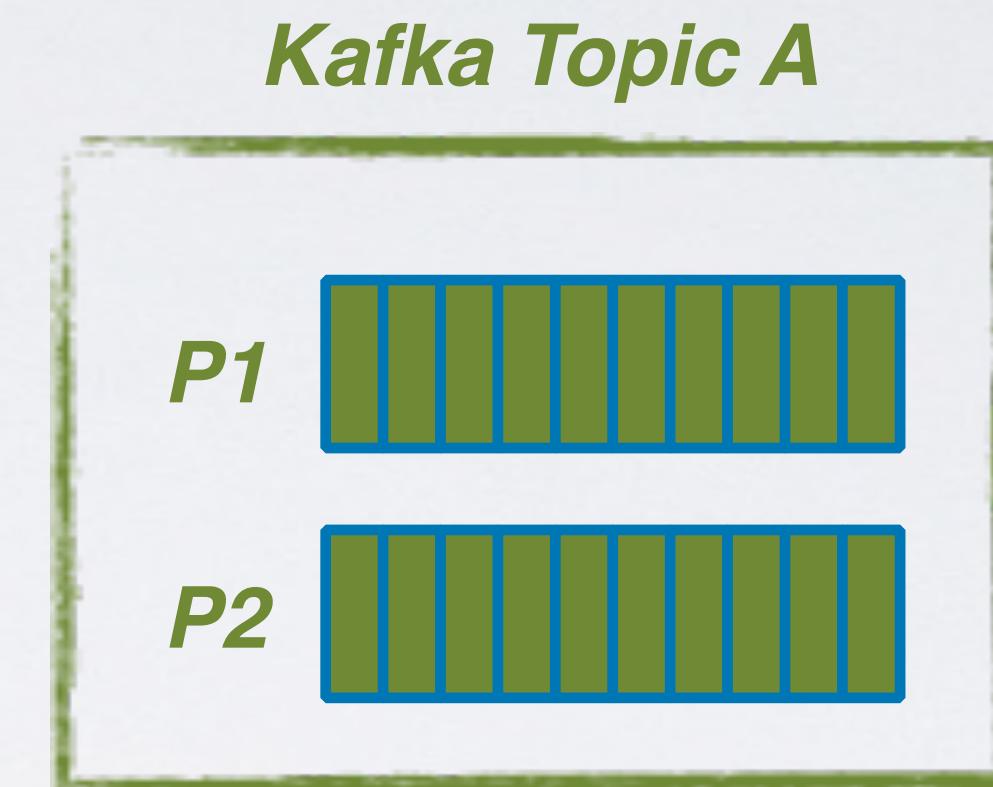
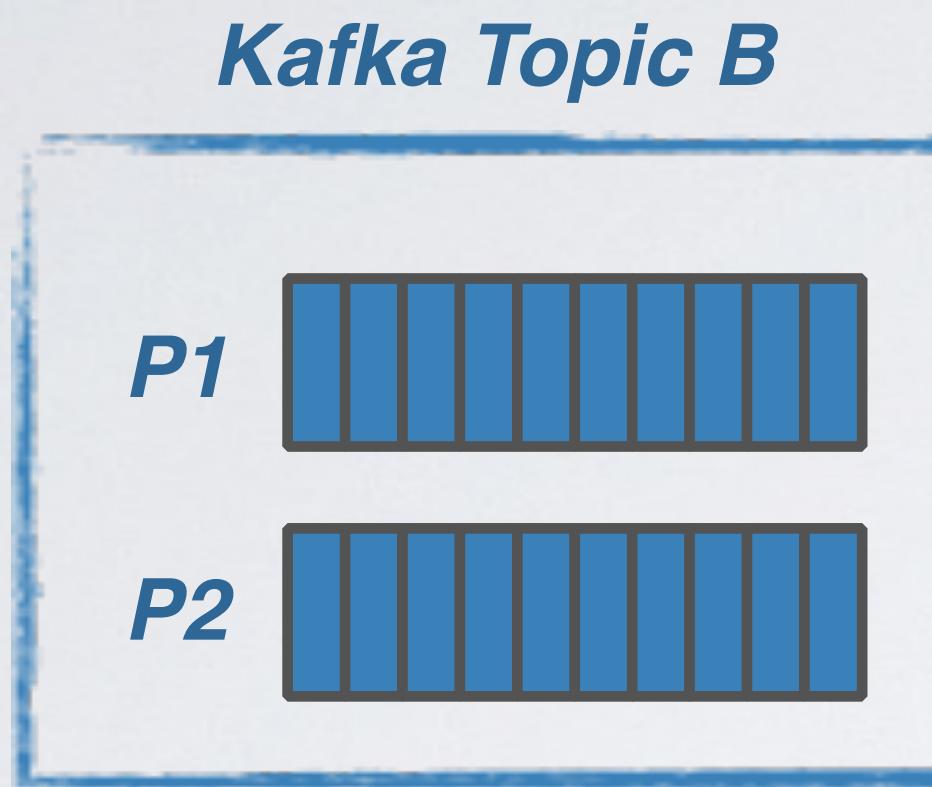
```
aggregated.to()
```



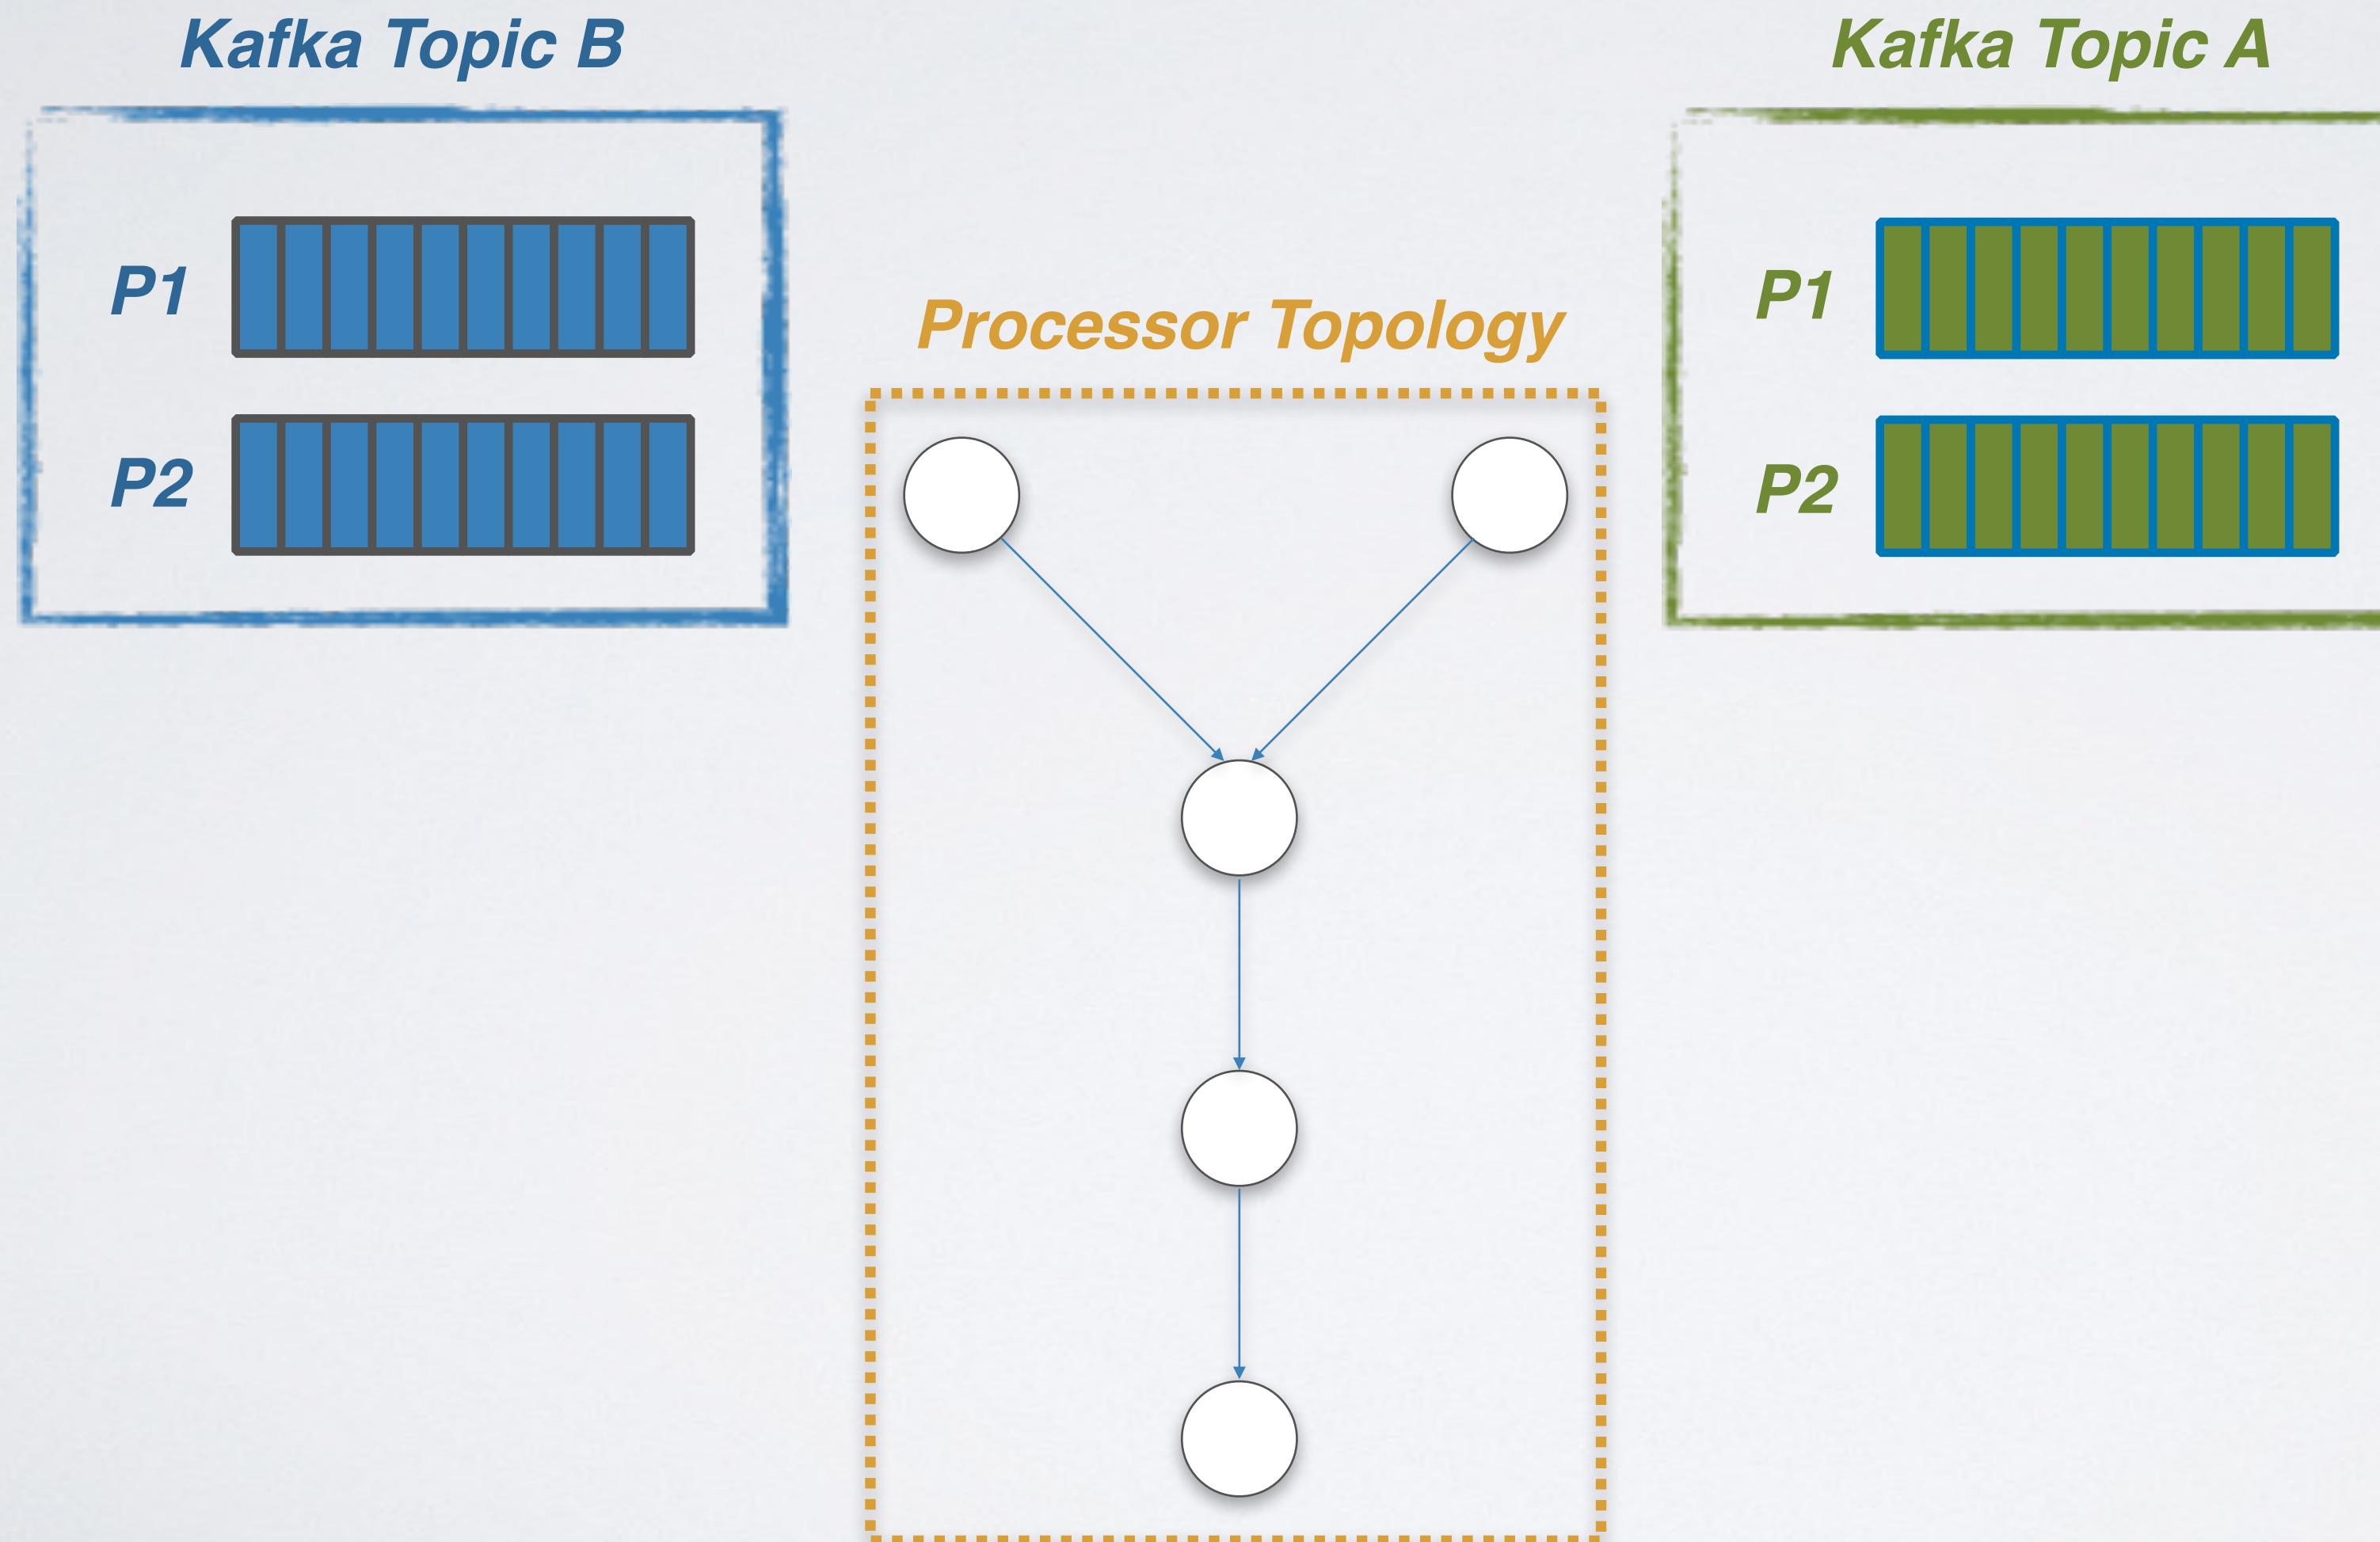
Processor Topology



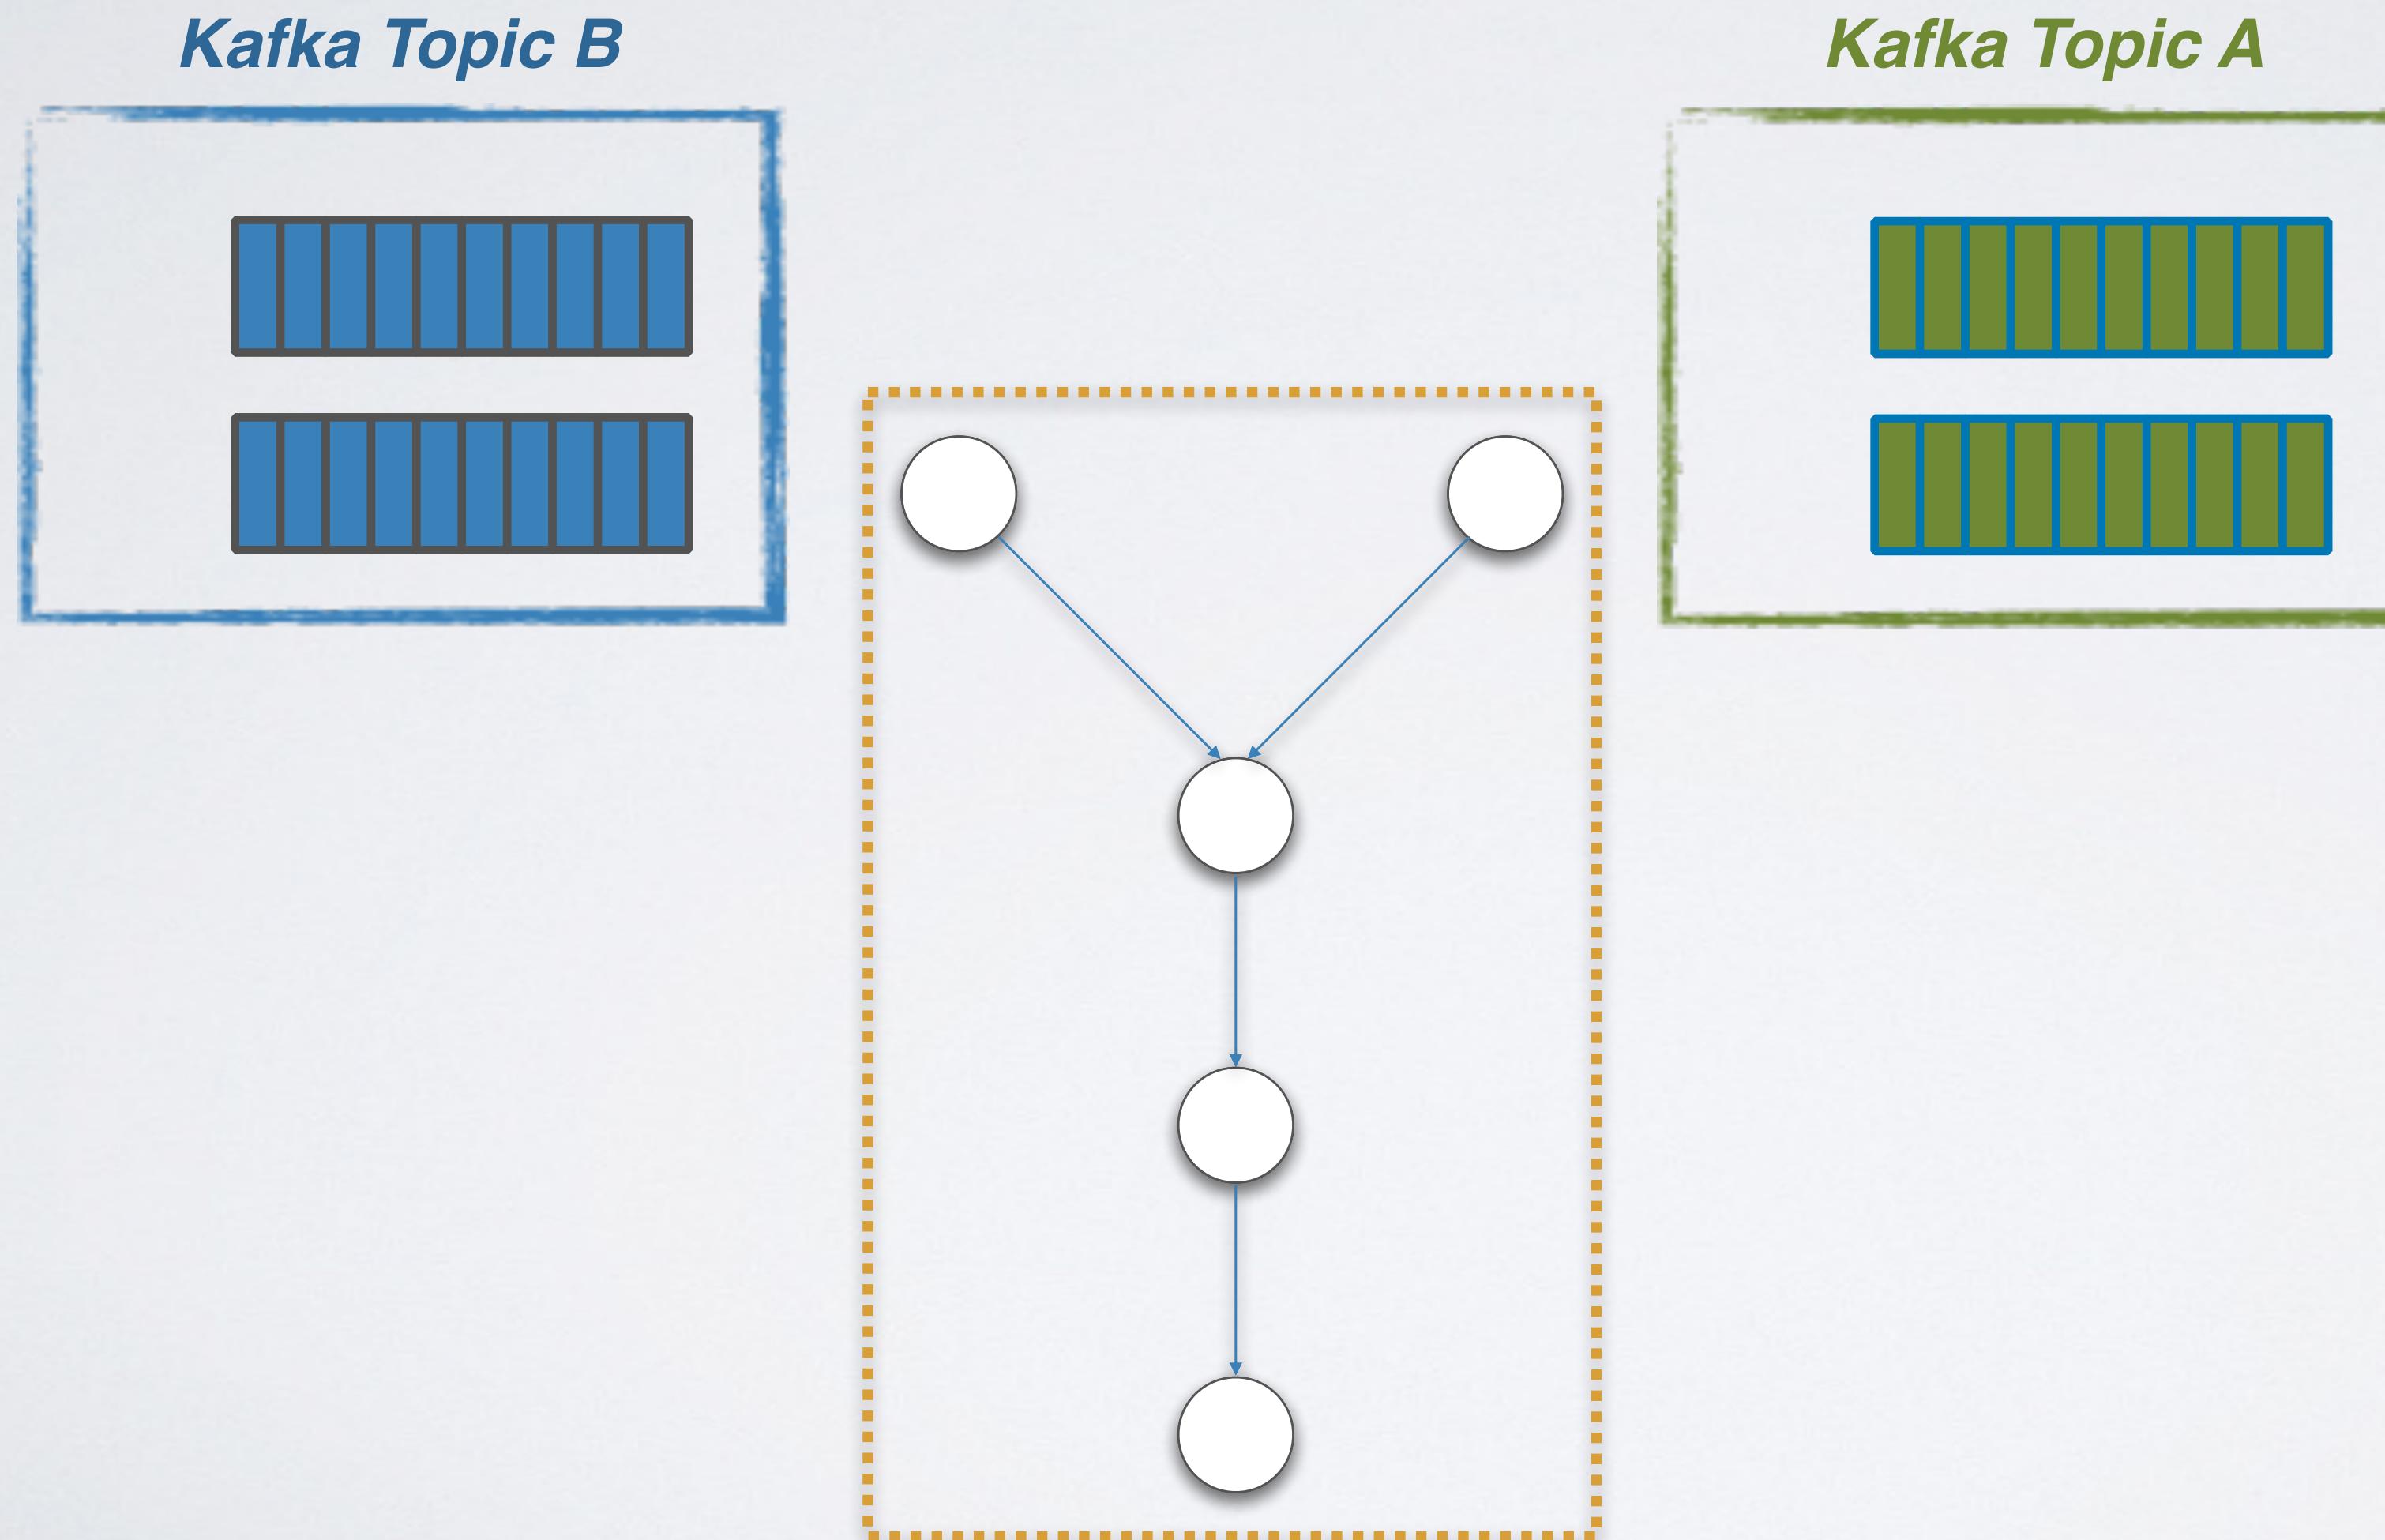
Stream Partitions and Tasks



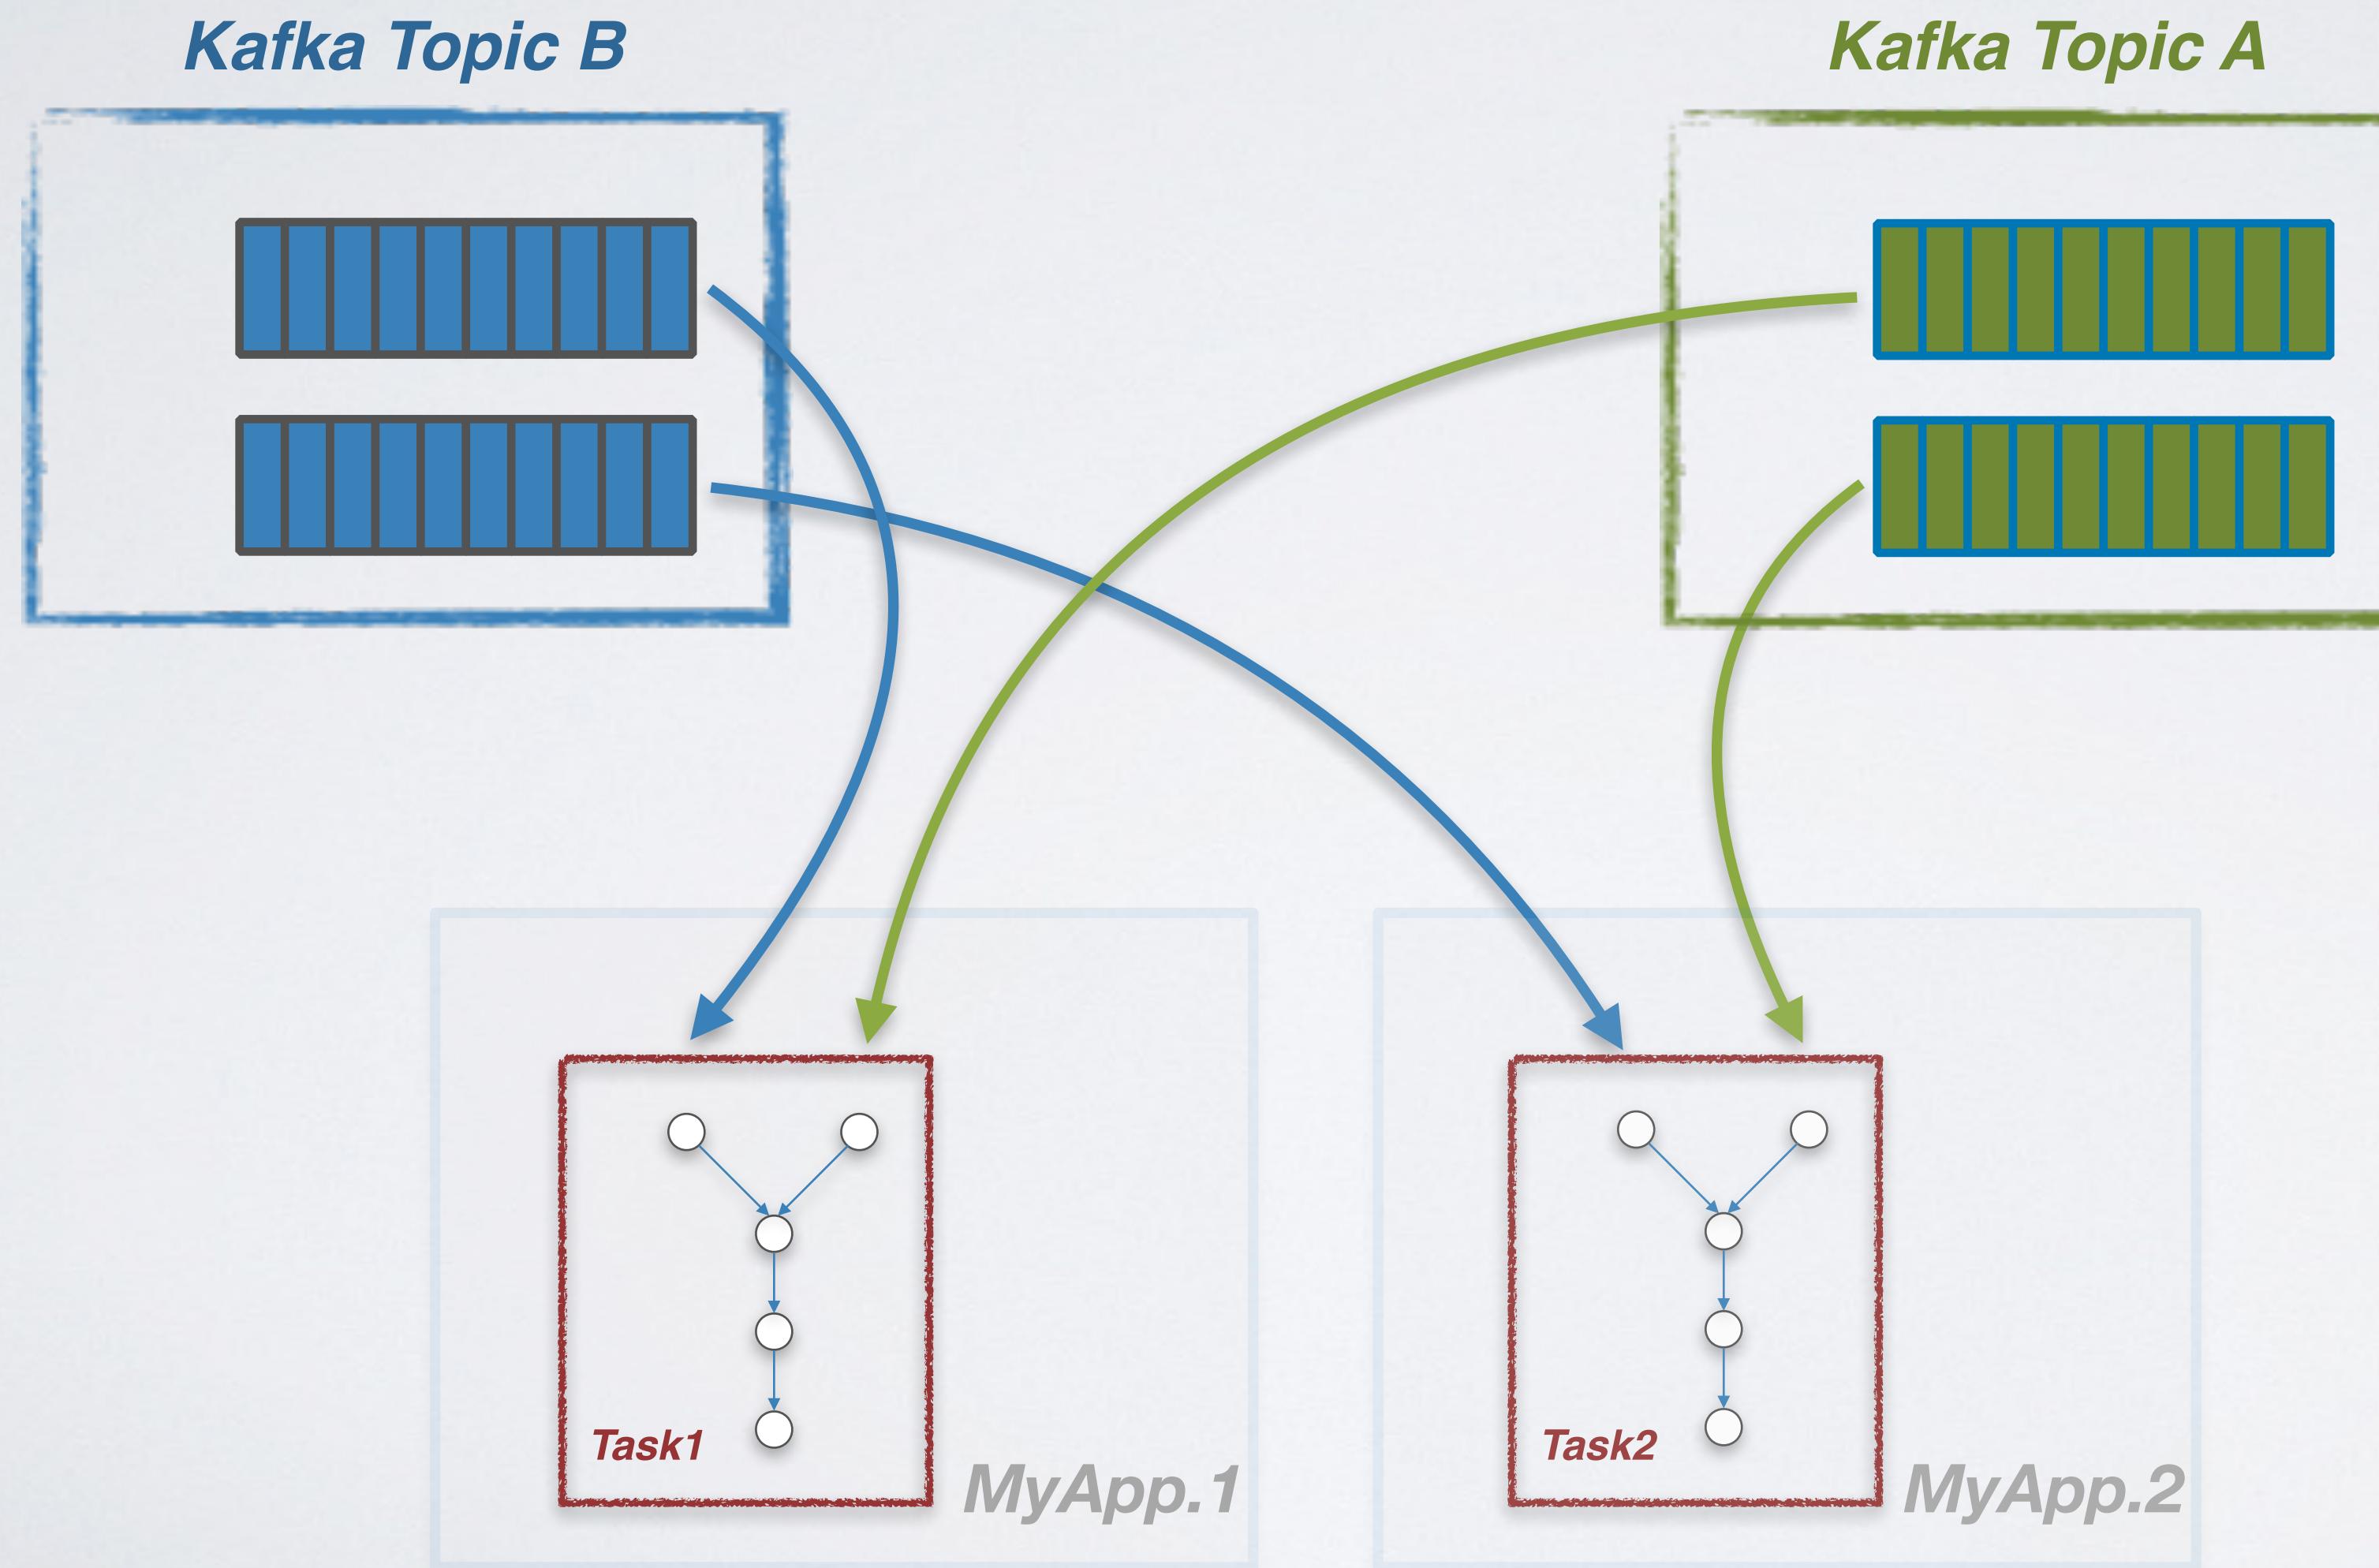
Stream Partitions and Tasks



Stream Partitions and Tasks



Stream Threads



States in Stream Processing

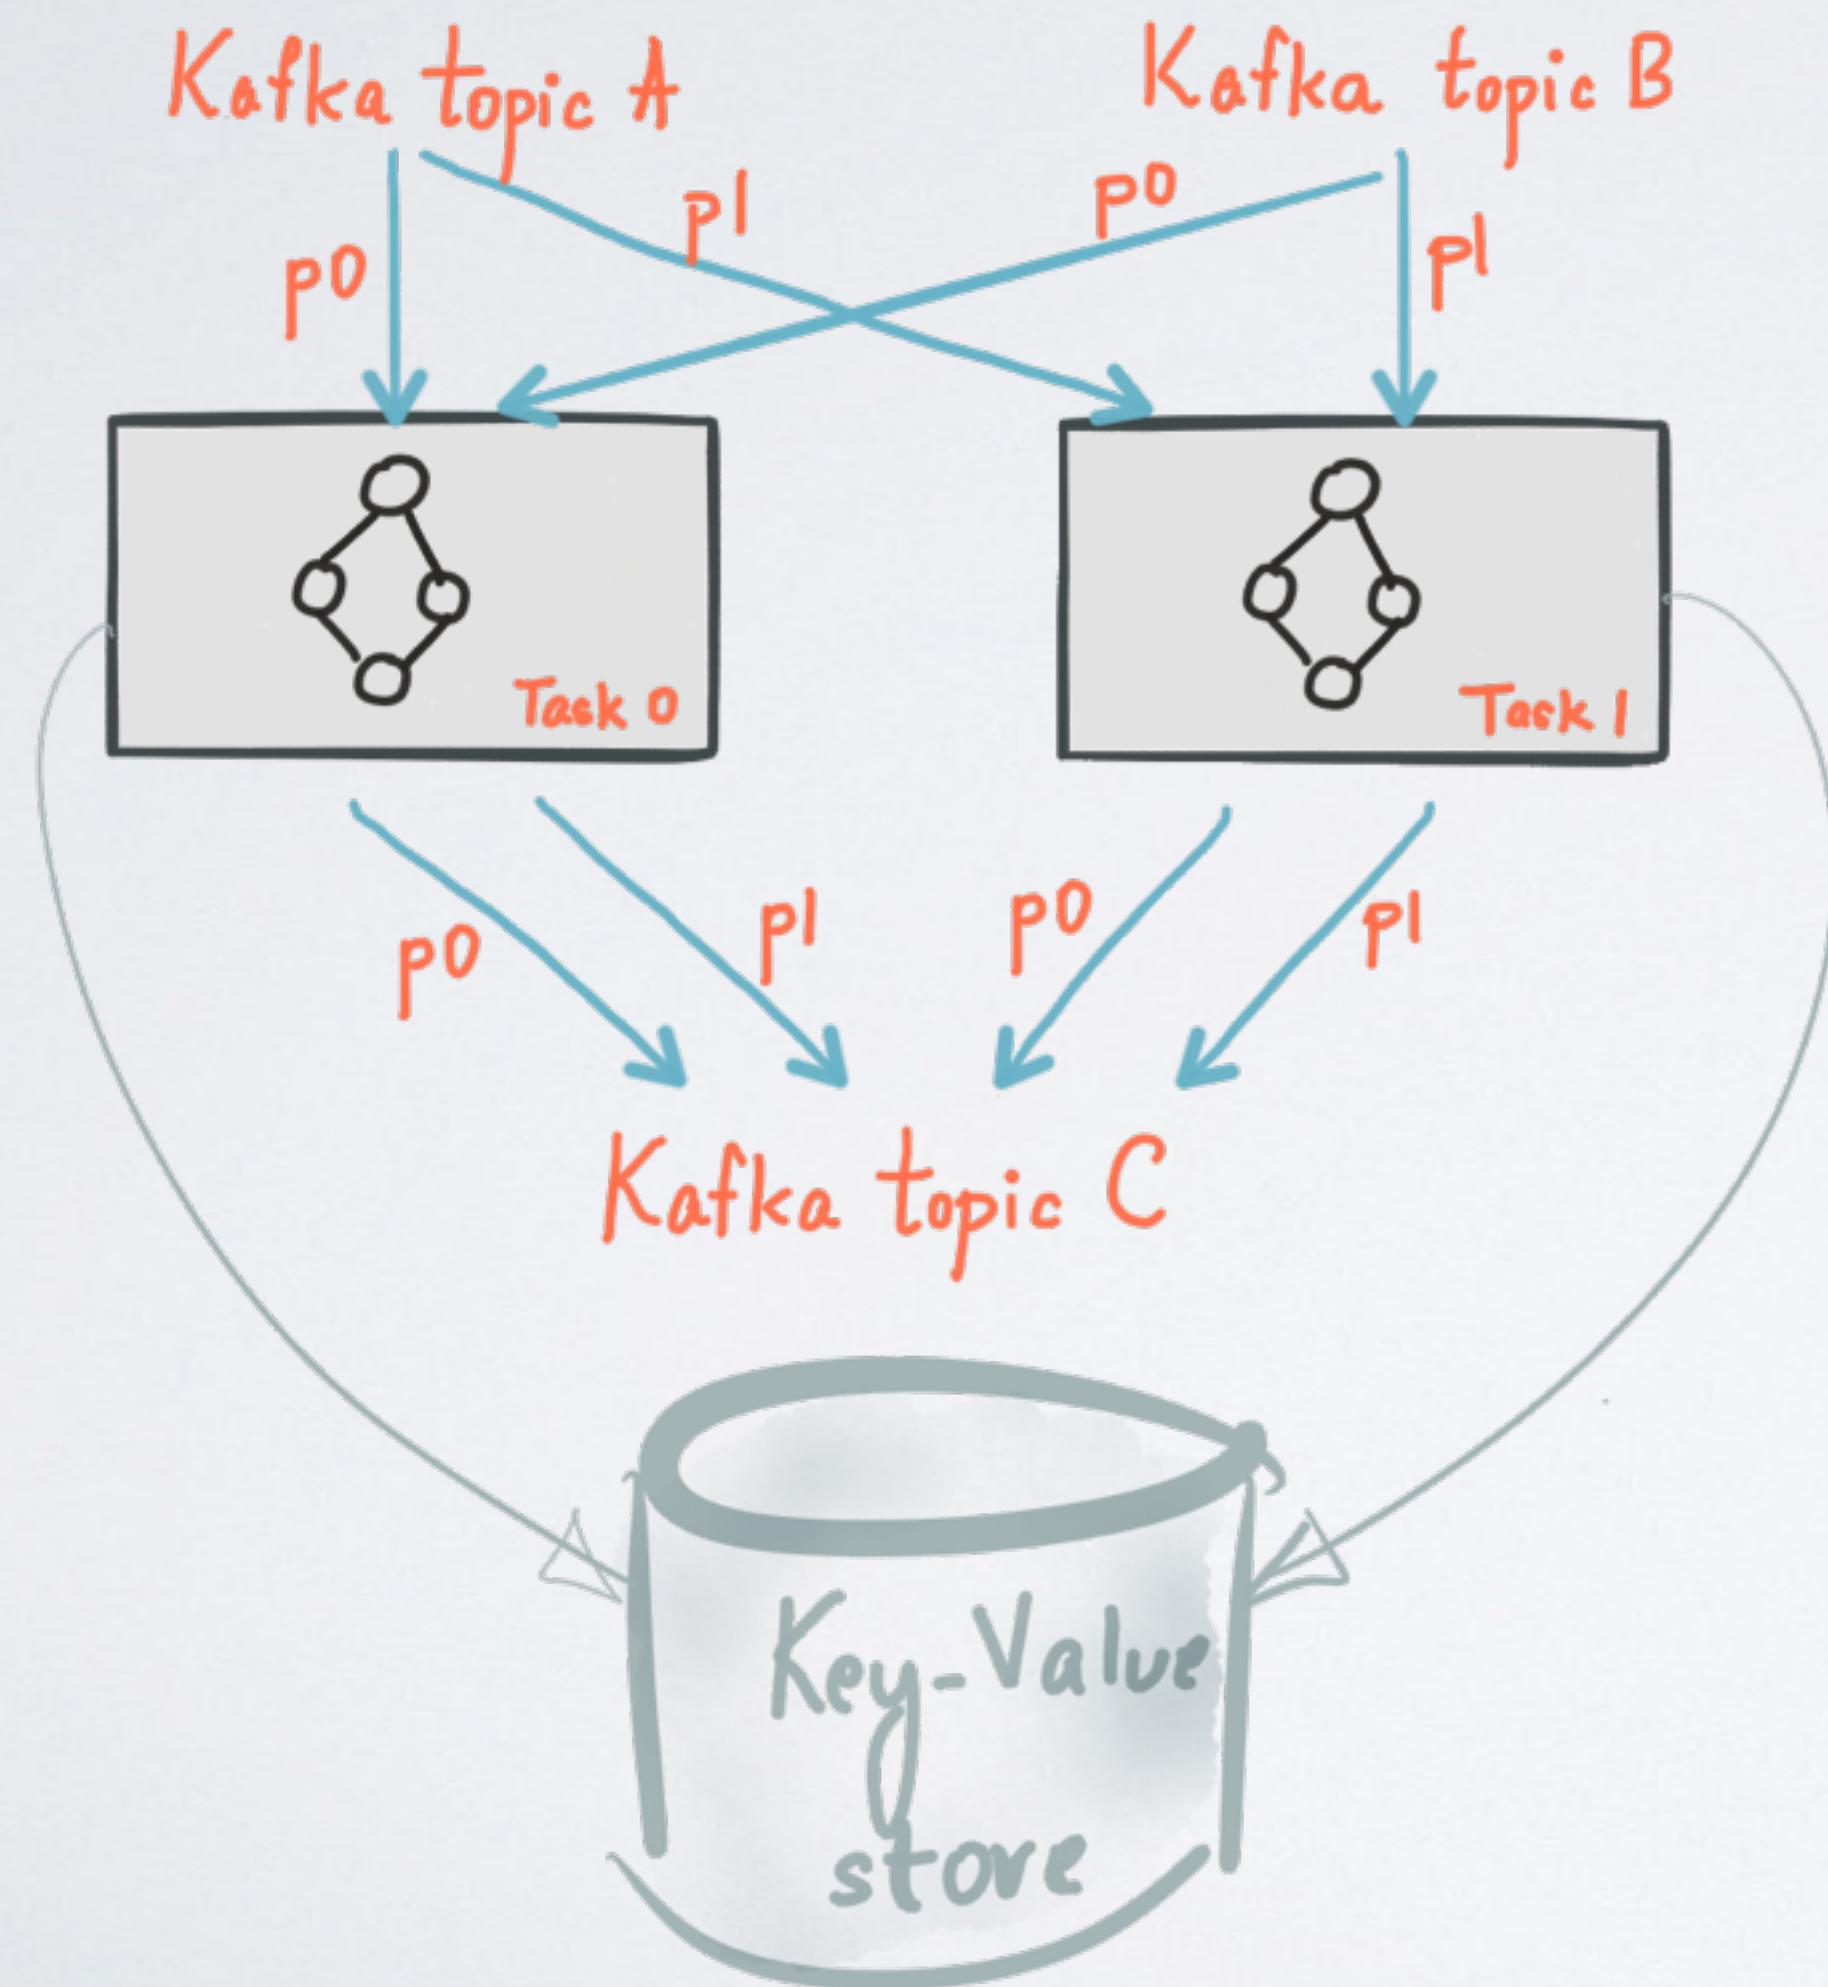
- *filter*
- *map*
- *join*
- *aggregate*



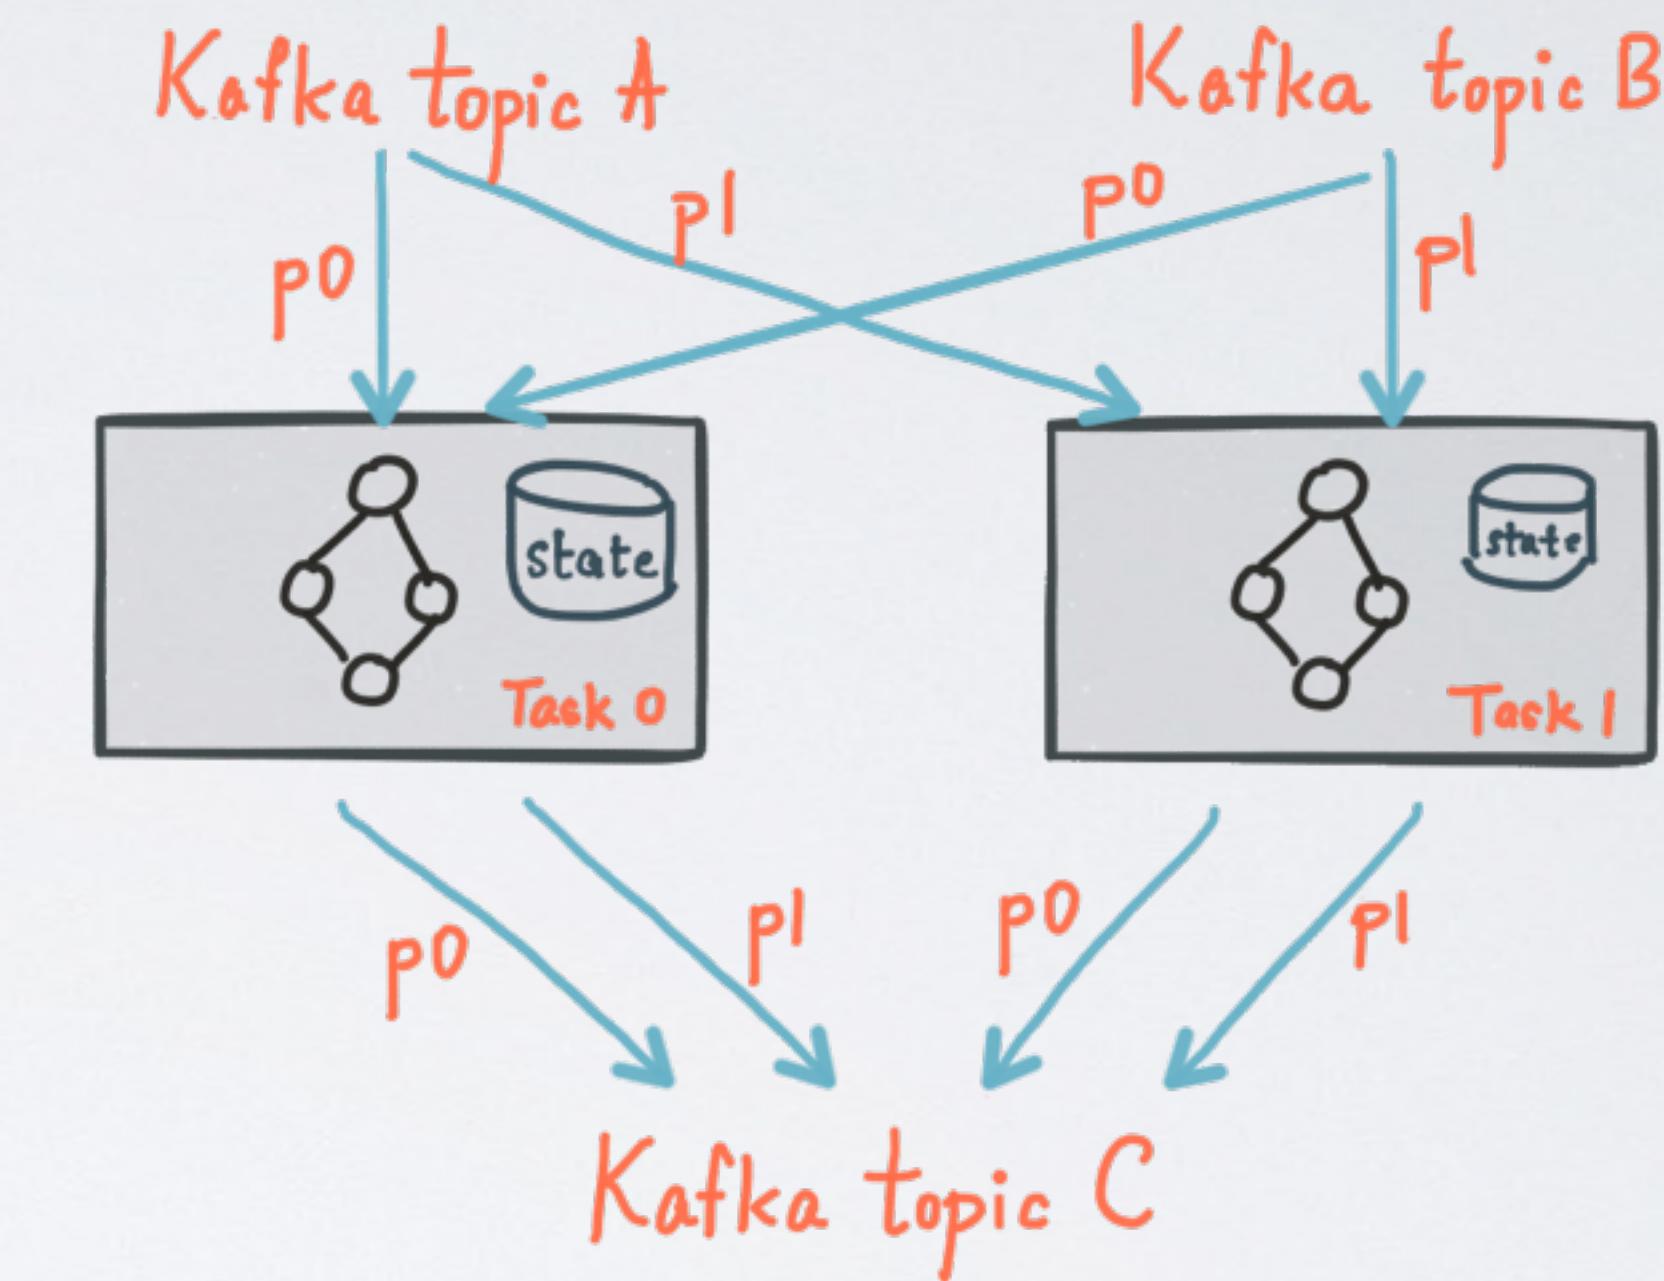
Stateless

Stateful

REMOTE STATE



LOCAL STATE



- faster
- better isolation
- flexible

States in Stream Processing

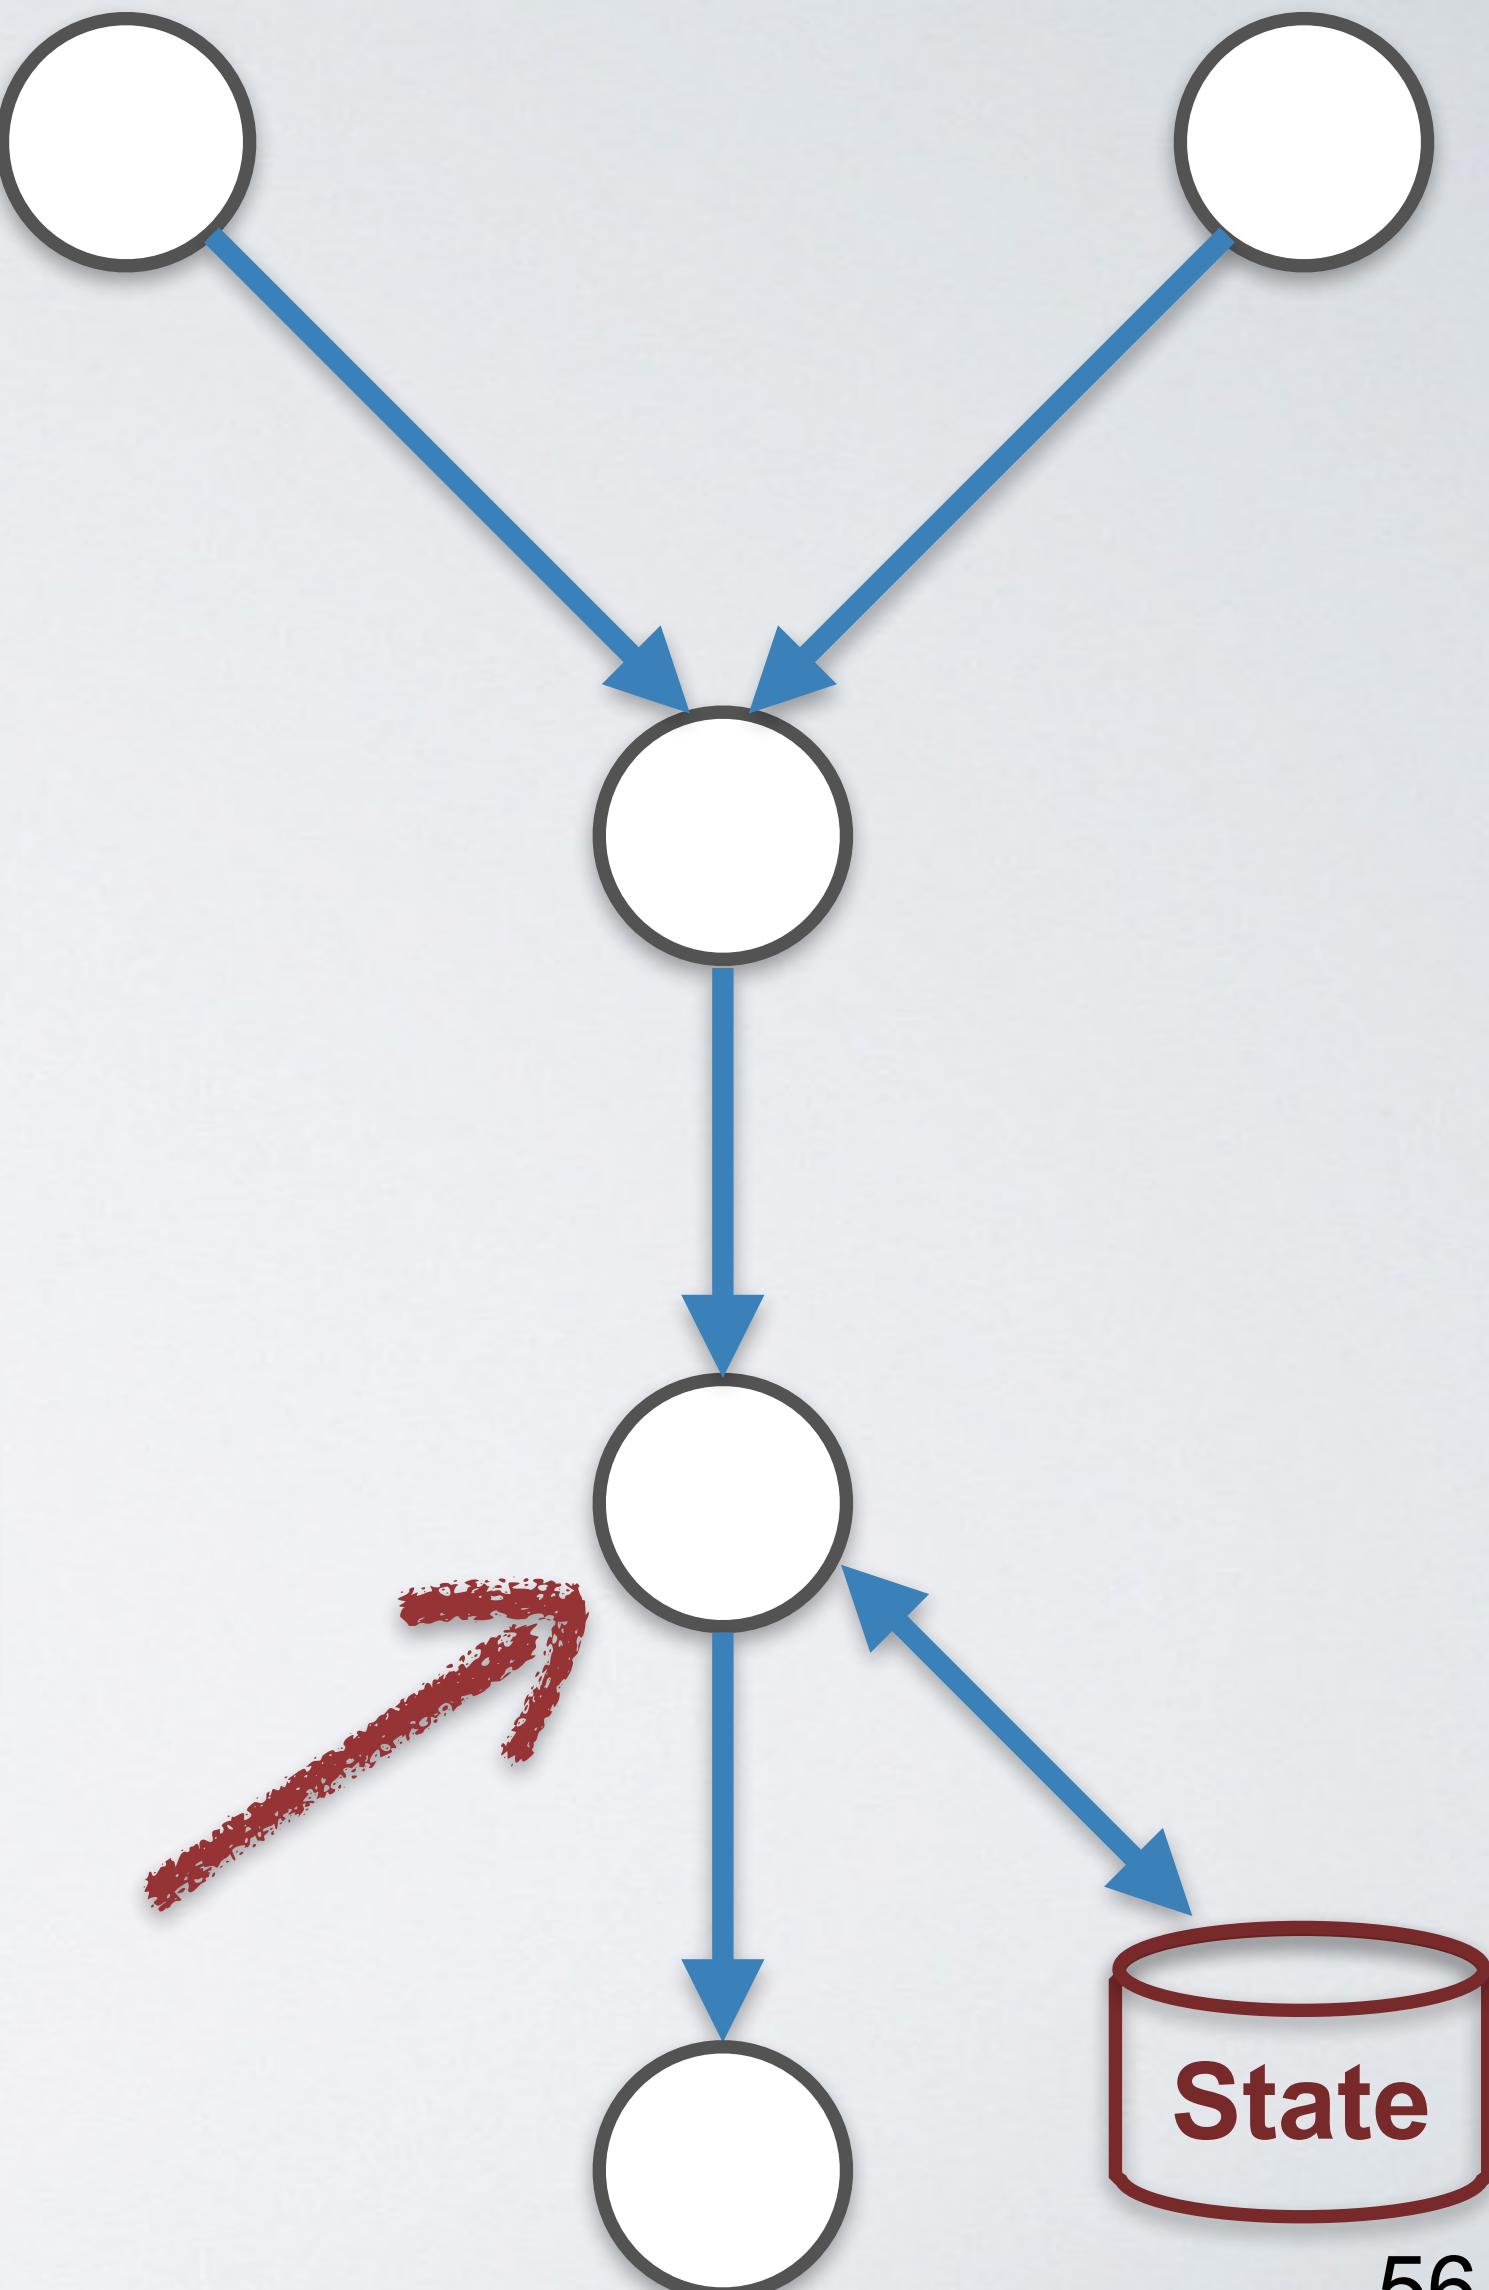
```
KStream<..> stream1 = builder.stream("topic1");
```

```
KStream<..> stream2 = builder.stream("topic2");
```

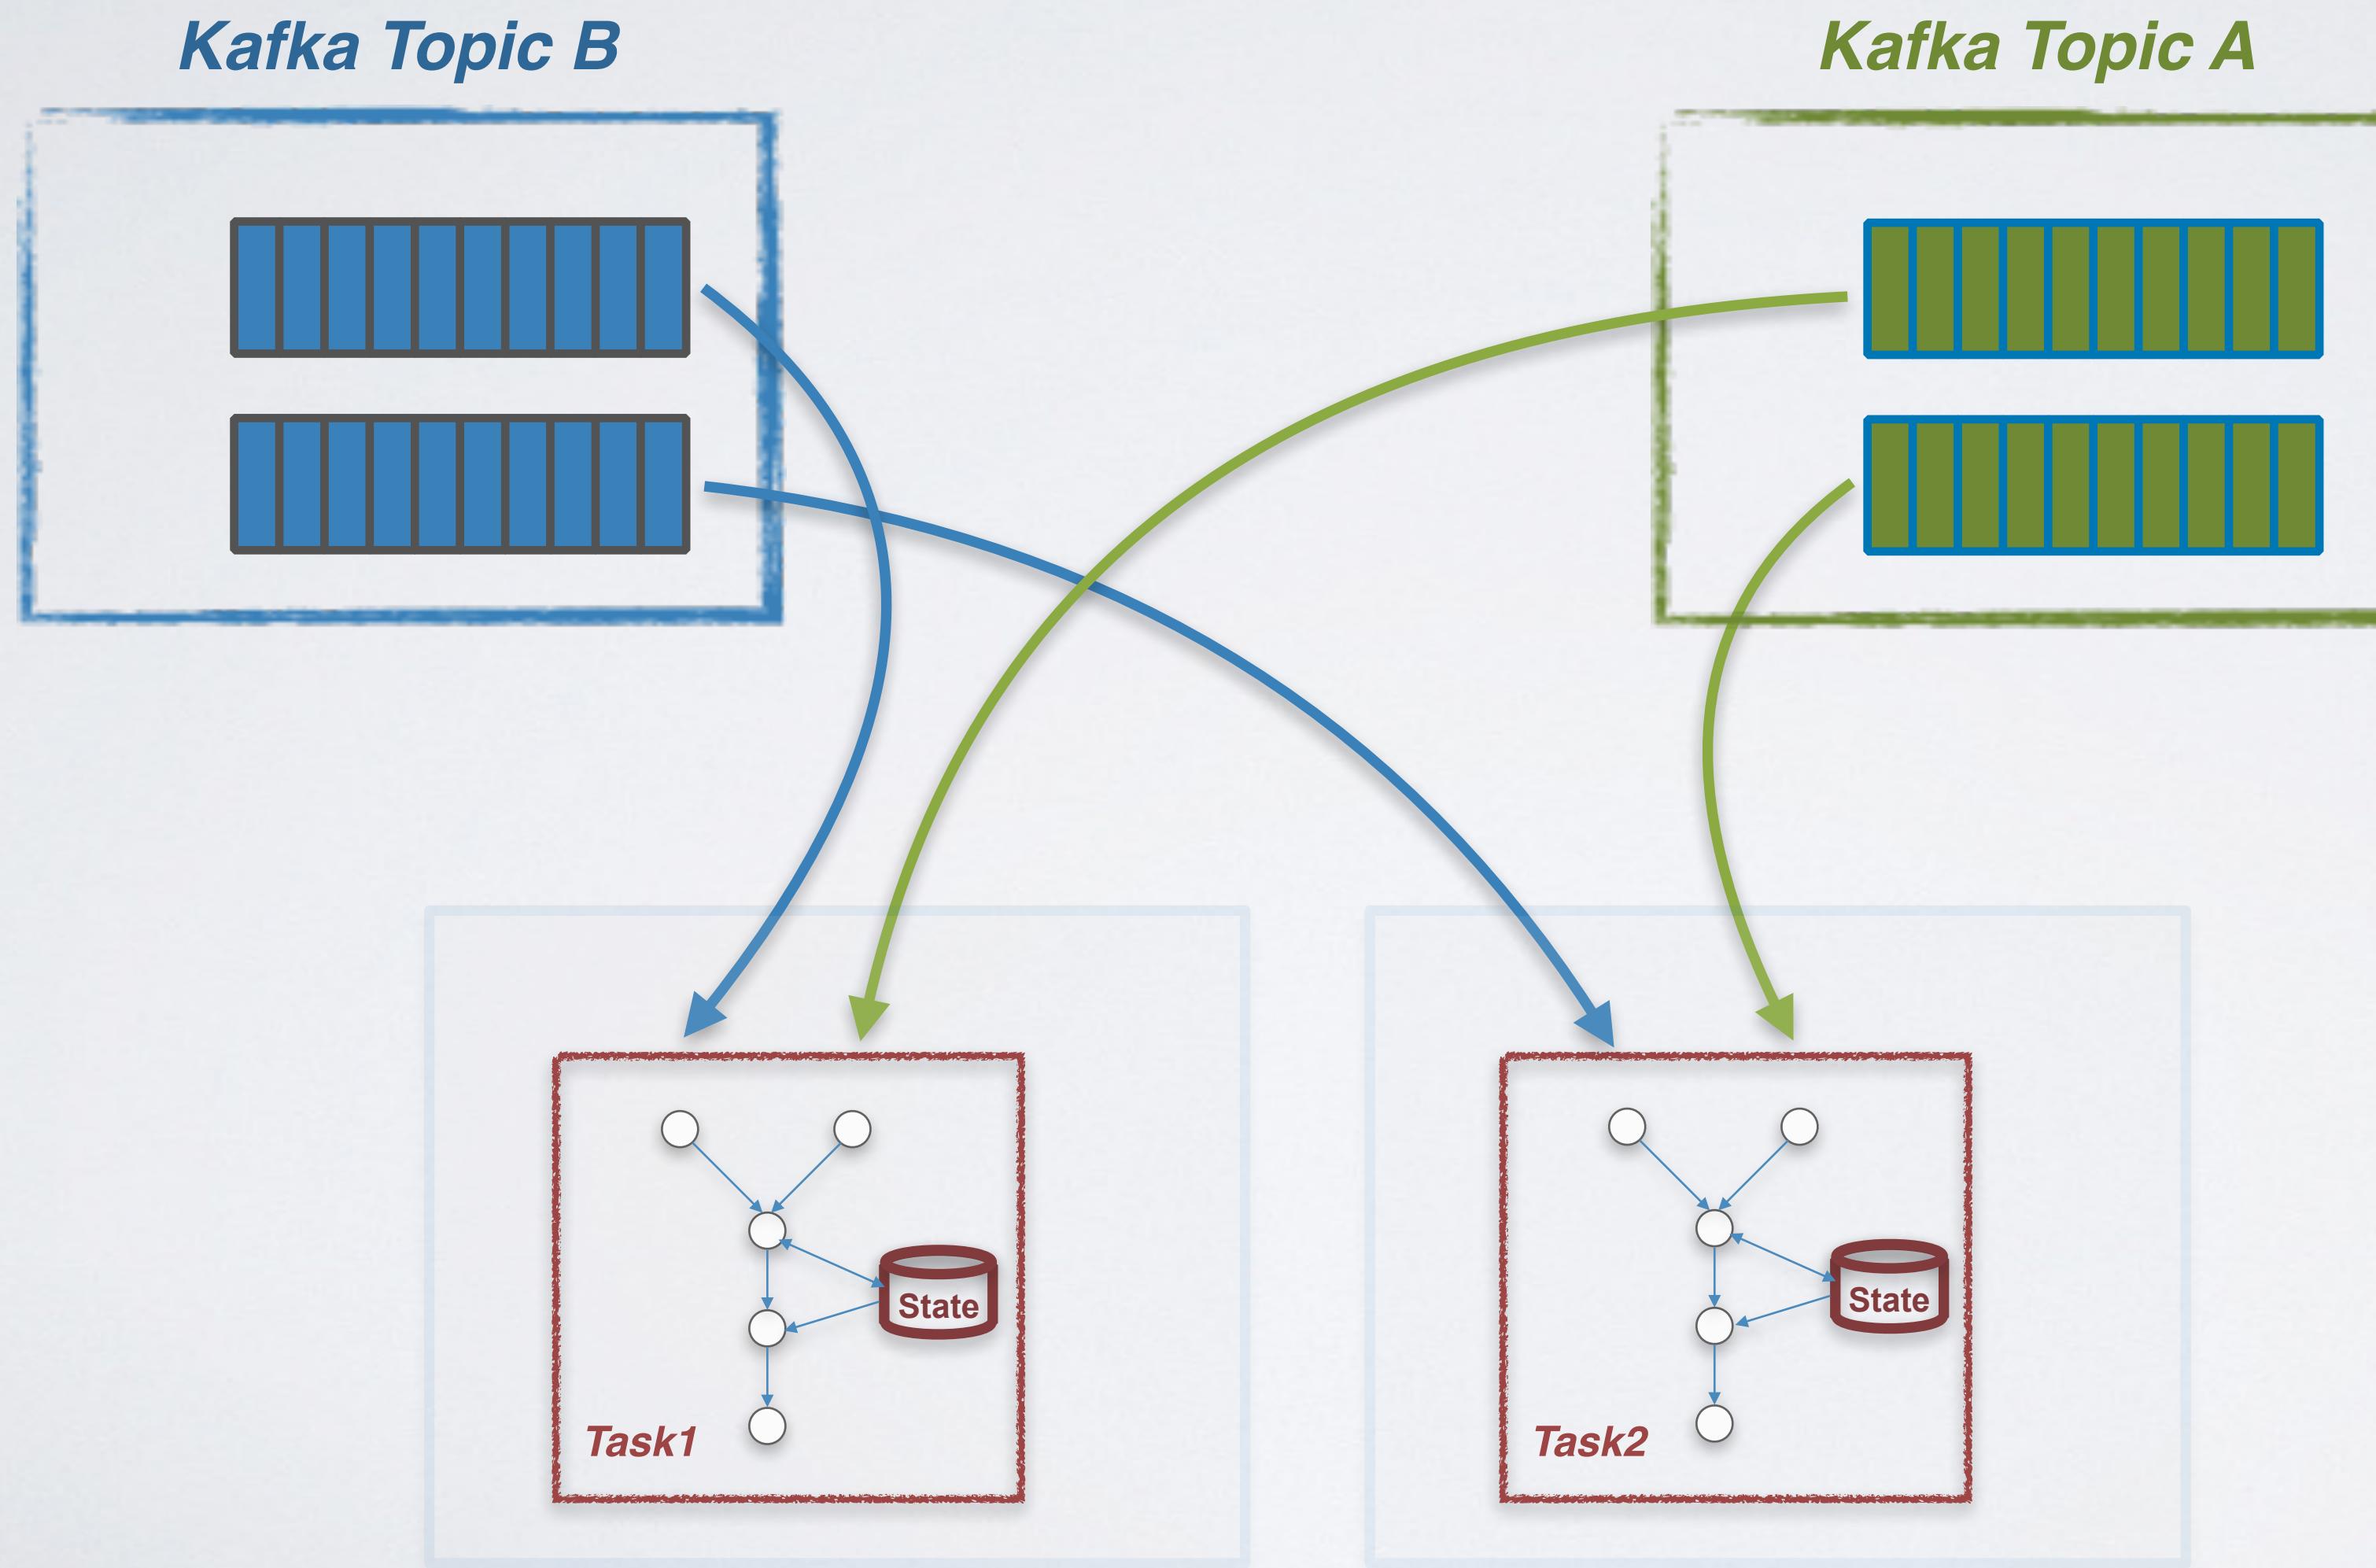
```
KStream<..> joined = stream1.leftJoin(stream2, ...);
```

```
KTable<..> aggregated = joined.aggregateByKey(...);
```

```
aggregated.to("topic2");
```



States in Stream Processing



Stream v.s. Table?

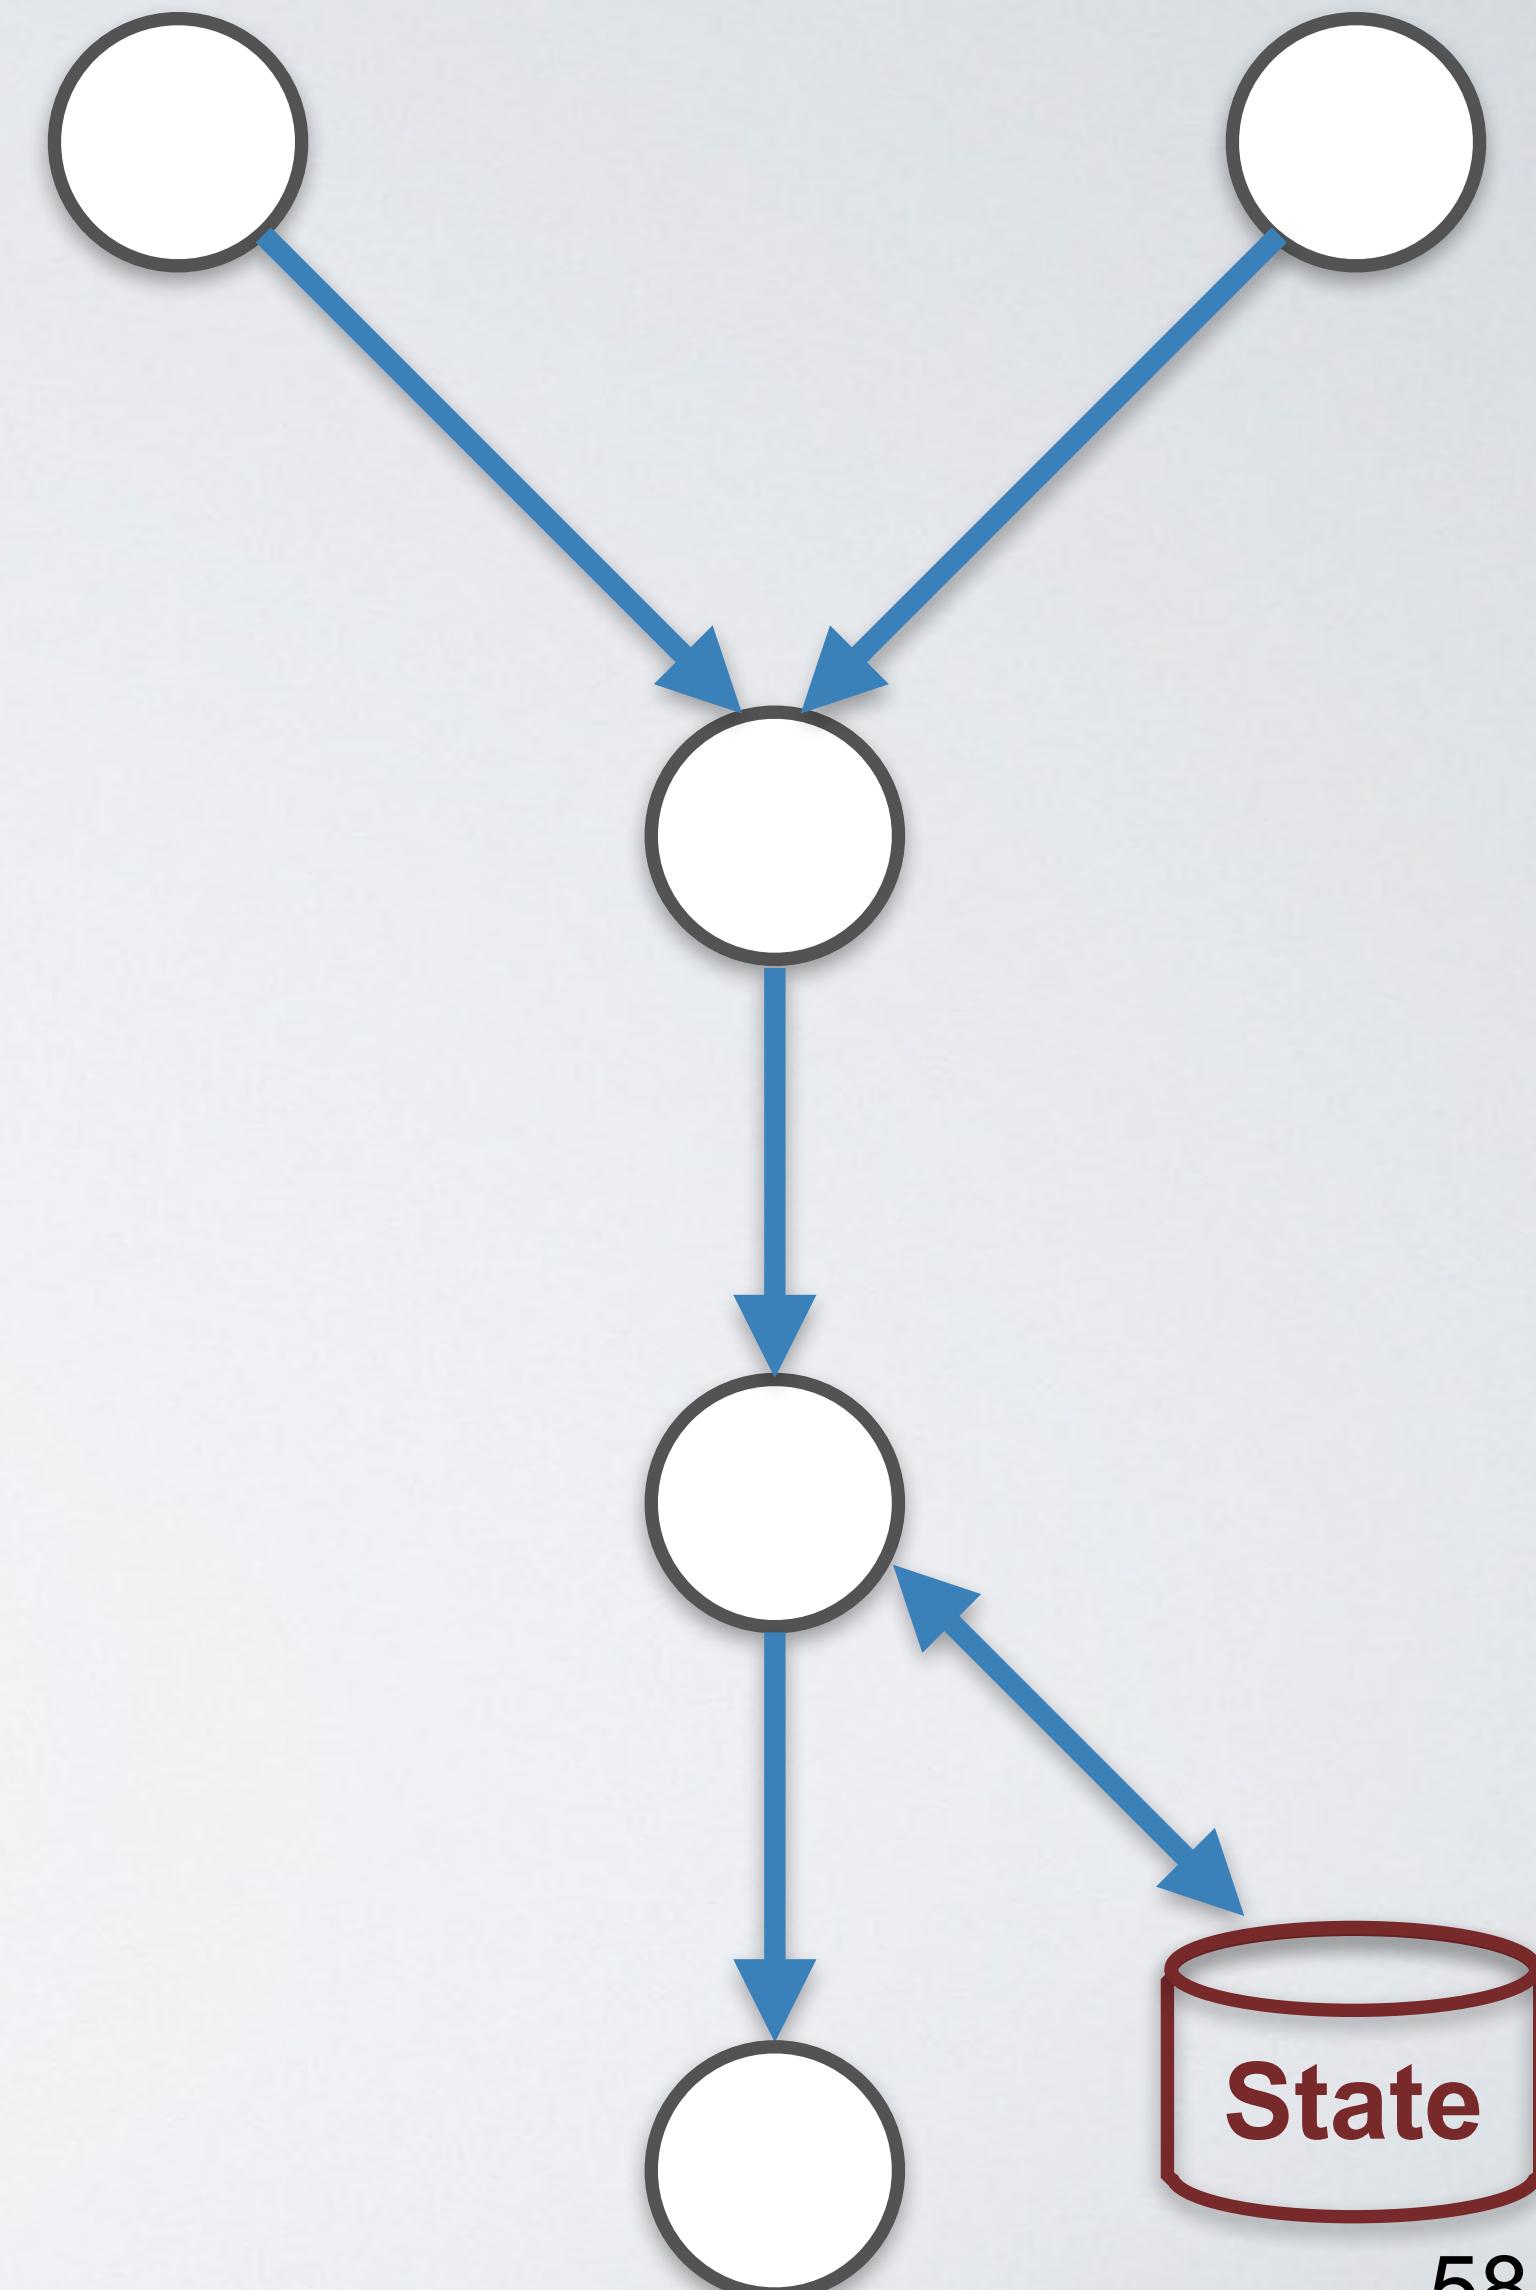
```
KStream<..> stream1 = builder.stream("topic1");
```

```
KStream<..> stream2 = builder.stream("topic2");
```

```
KStream<..> joined = stream1.leftJoin(stream2, ...);
```

```
KTable<..> aggregated = joined.aggregateByKey(...);
```

```
aggregated.to("topic2");
```



Tables \approx *Streams*

TABLES \approx STREAMS

(key₁, value₁)

(key₂, value₂)

(key₁, value₃)

.

.

.

TABLES \approx STREAMS

$(\text{key1}, \text{value1}) \rightarrow$

key1	value1
------	--------

$(\text{key2}, \text{value2}) \rightarrow$

key1	value1
key2	value2

$(\text{key1}, \text{value3}) \rightarrow$

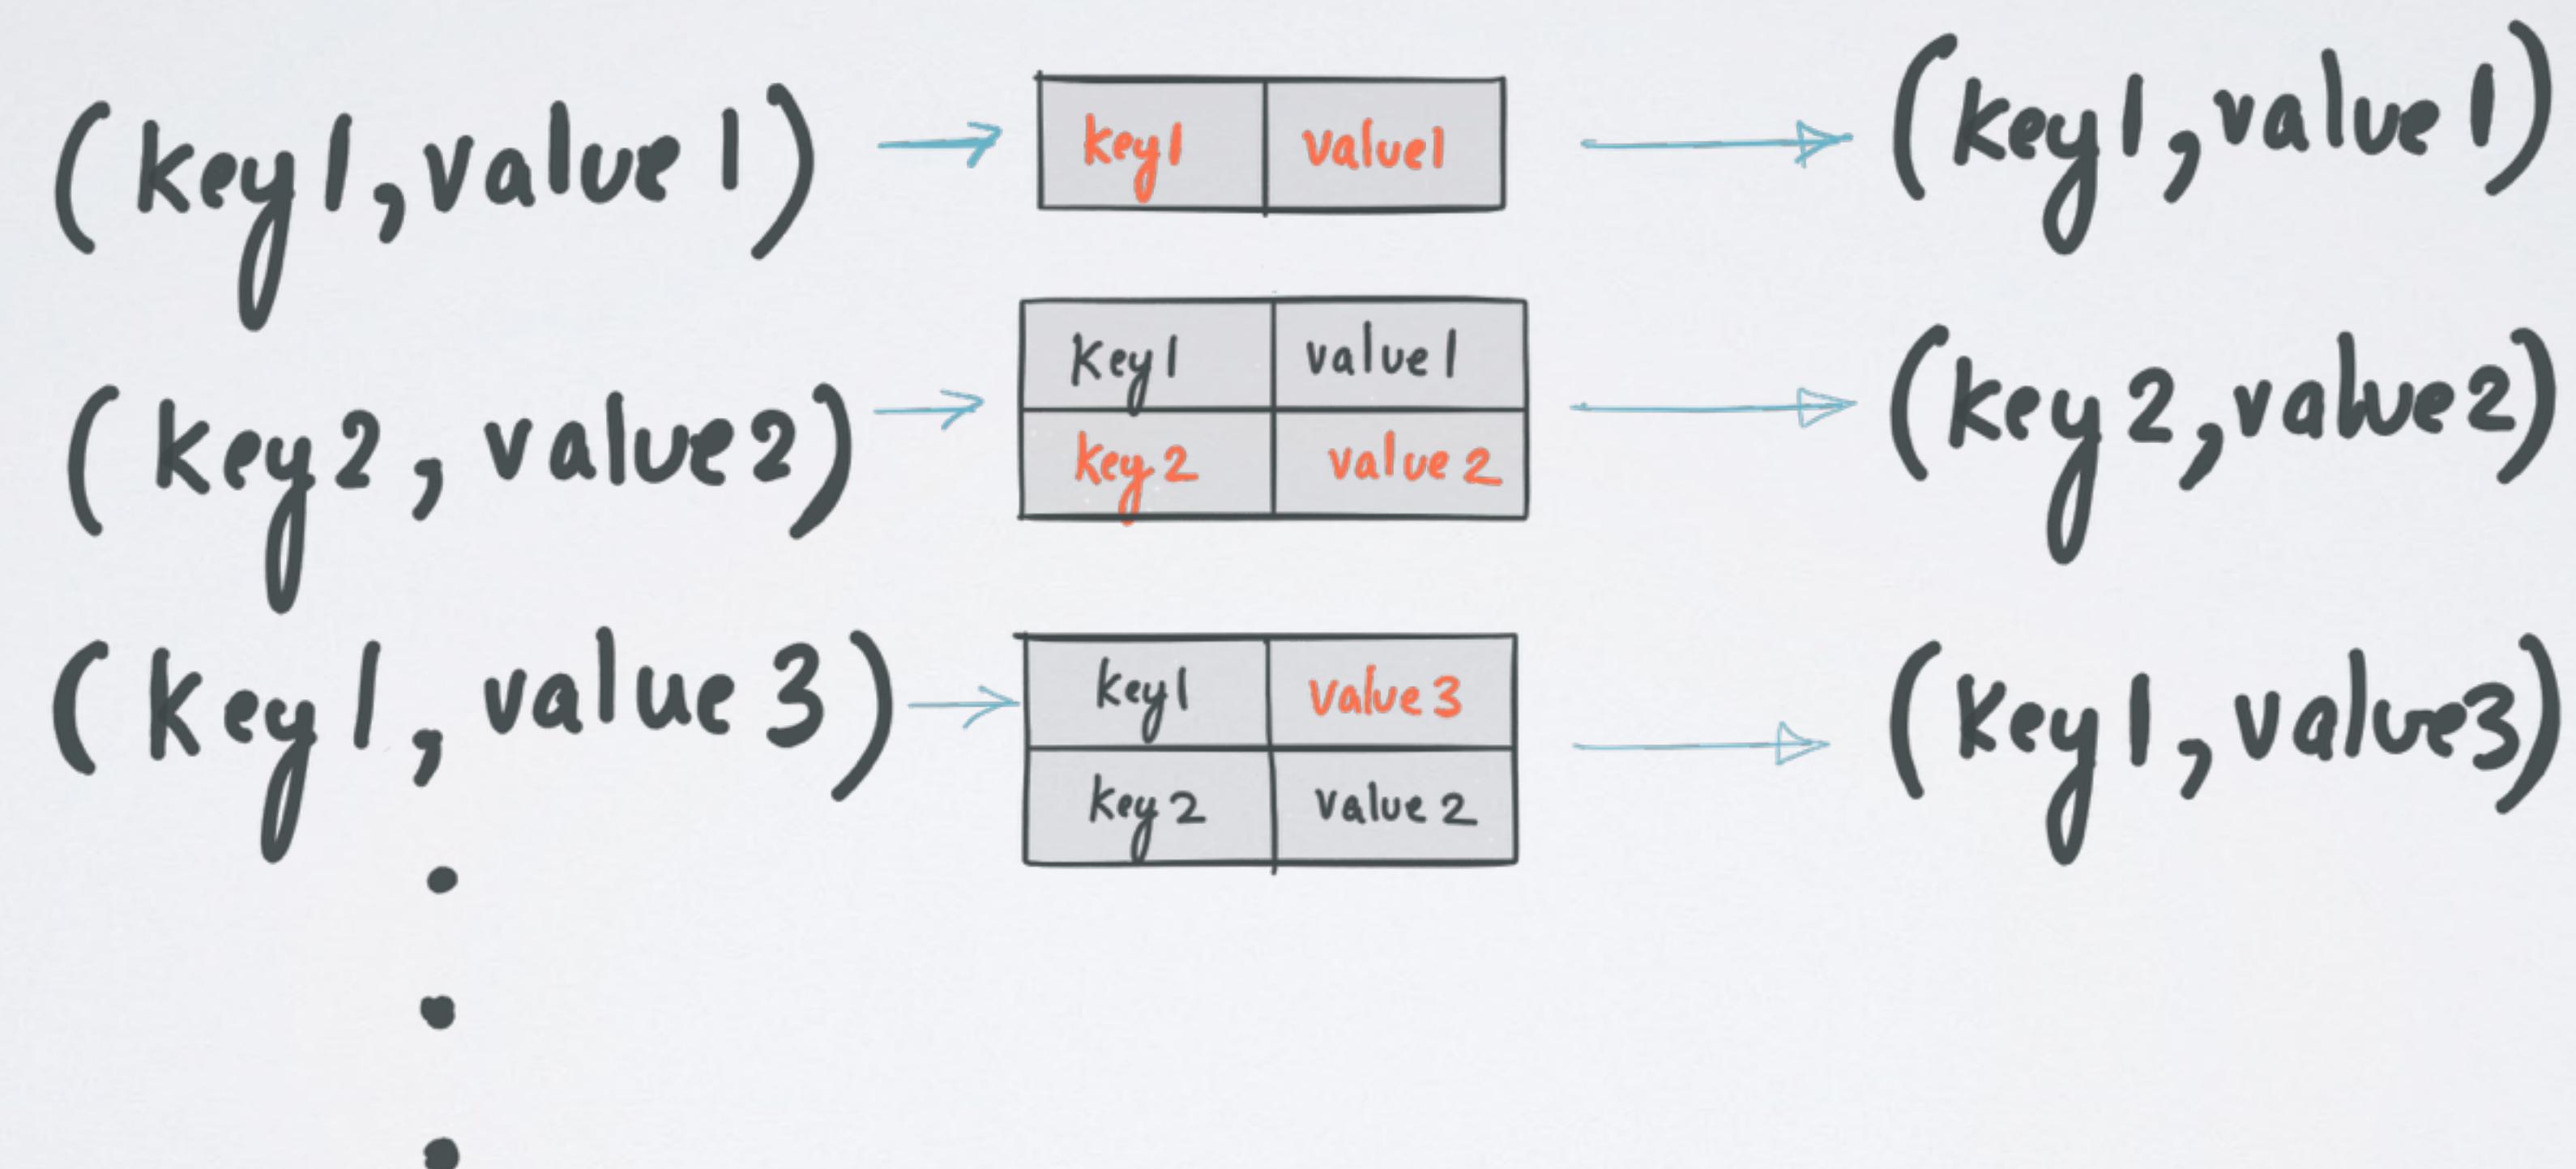
key1	value3
key2	value2

•

•

•

TABLES \approx STREAMS



The Stream-Table Duality

- A **stream** is a changelog of a **table**
- A **table** is a materialized view at time of a **stream**
- Example: change data capture (CDC) of databases

KStream = *interprets data as record stream*

~ think: “append-only”

KTable = *data as changelog stream*

~ *continuously updated materialized view*

KStream

User purchase history



KTable

User employment profile



KStream

User purchase history



KTable

User employment profile



KStream

User purchase history



time

"Alice bought **eggs** and **milk**."

KTable

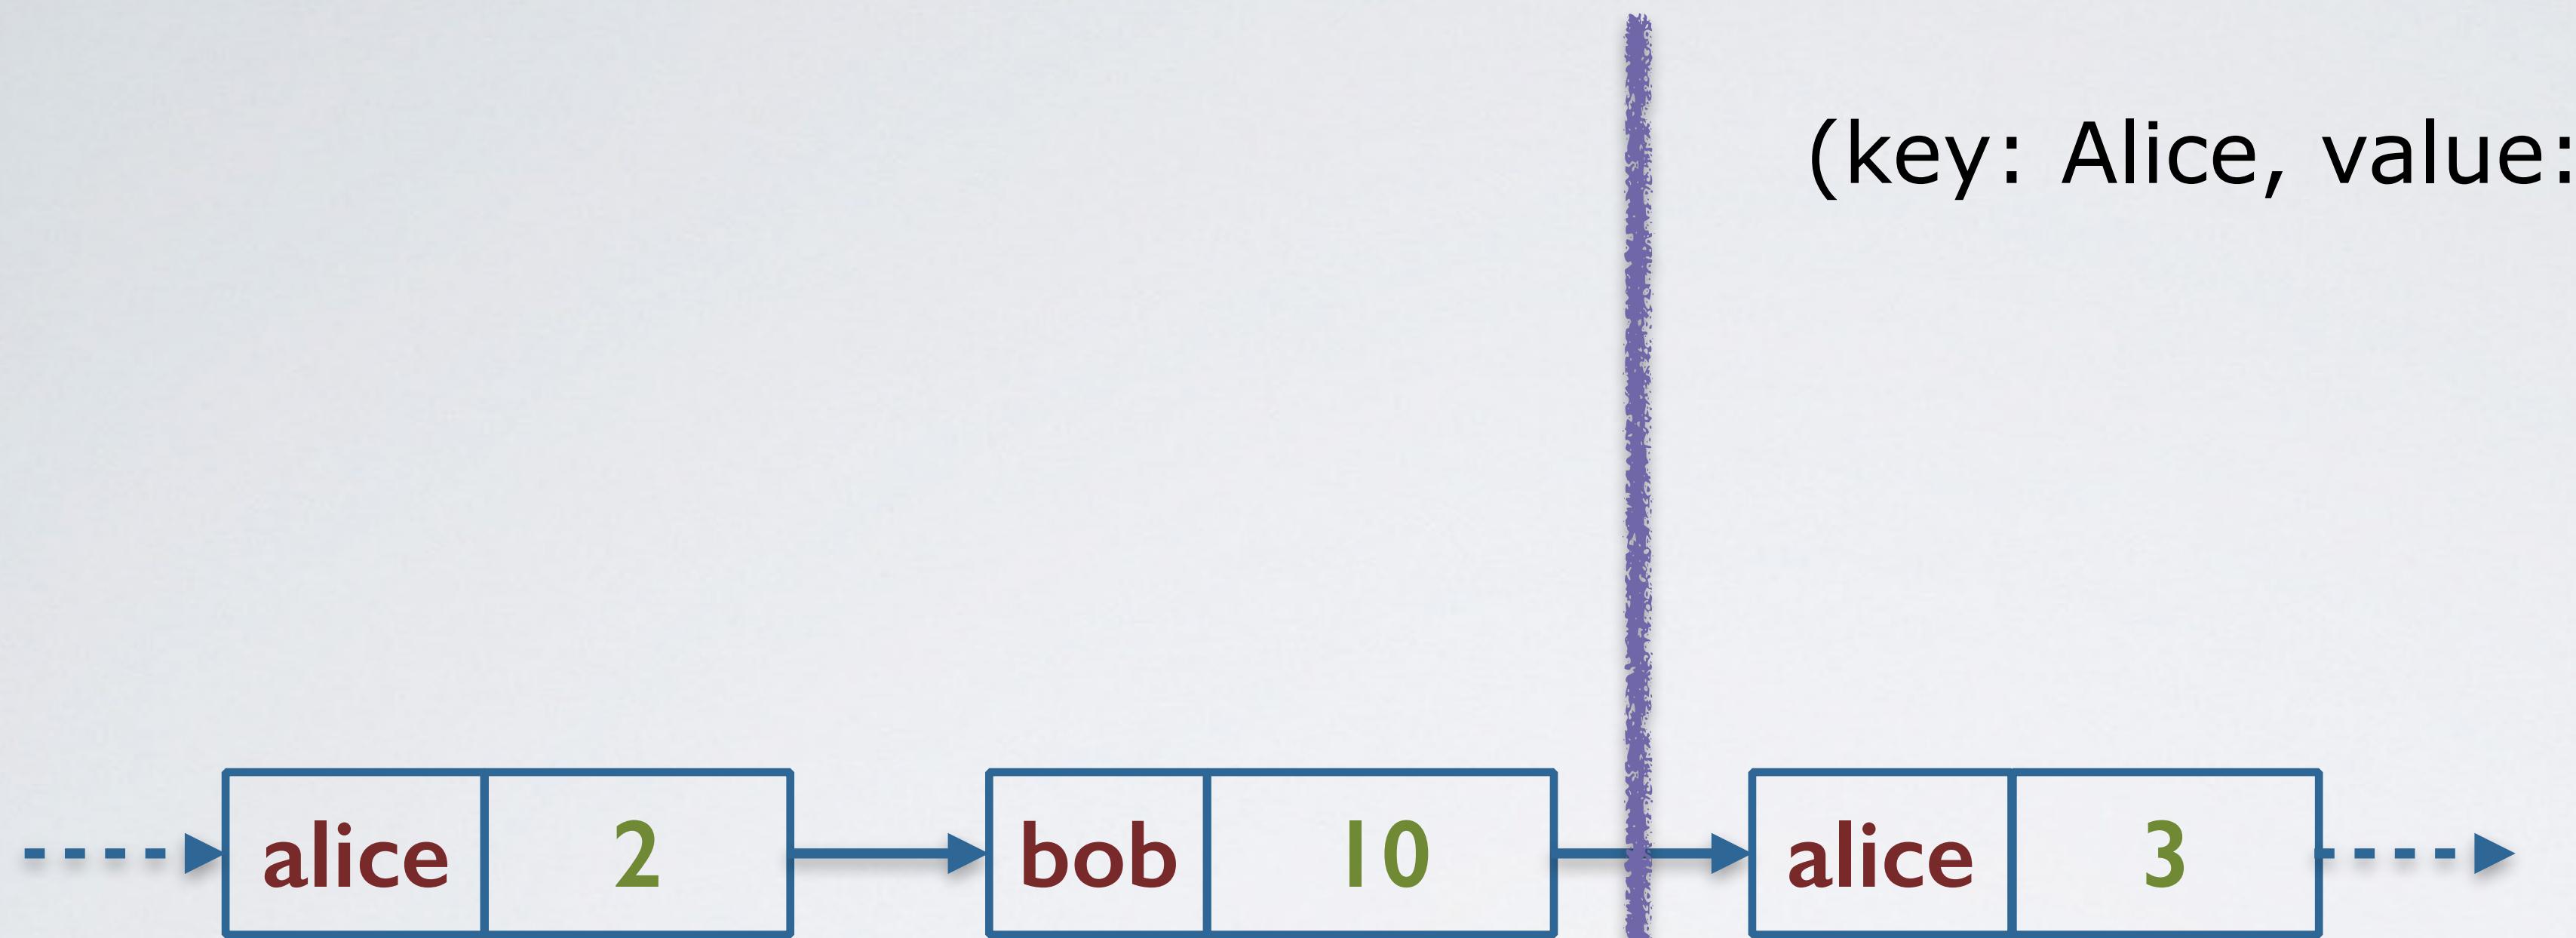
User employment profile



"Alice is now at ~~LinkedIn~~
Microsoft."

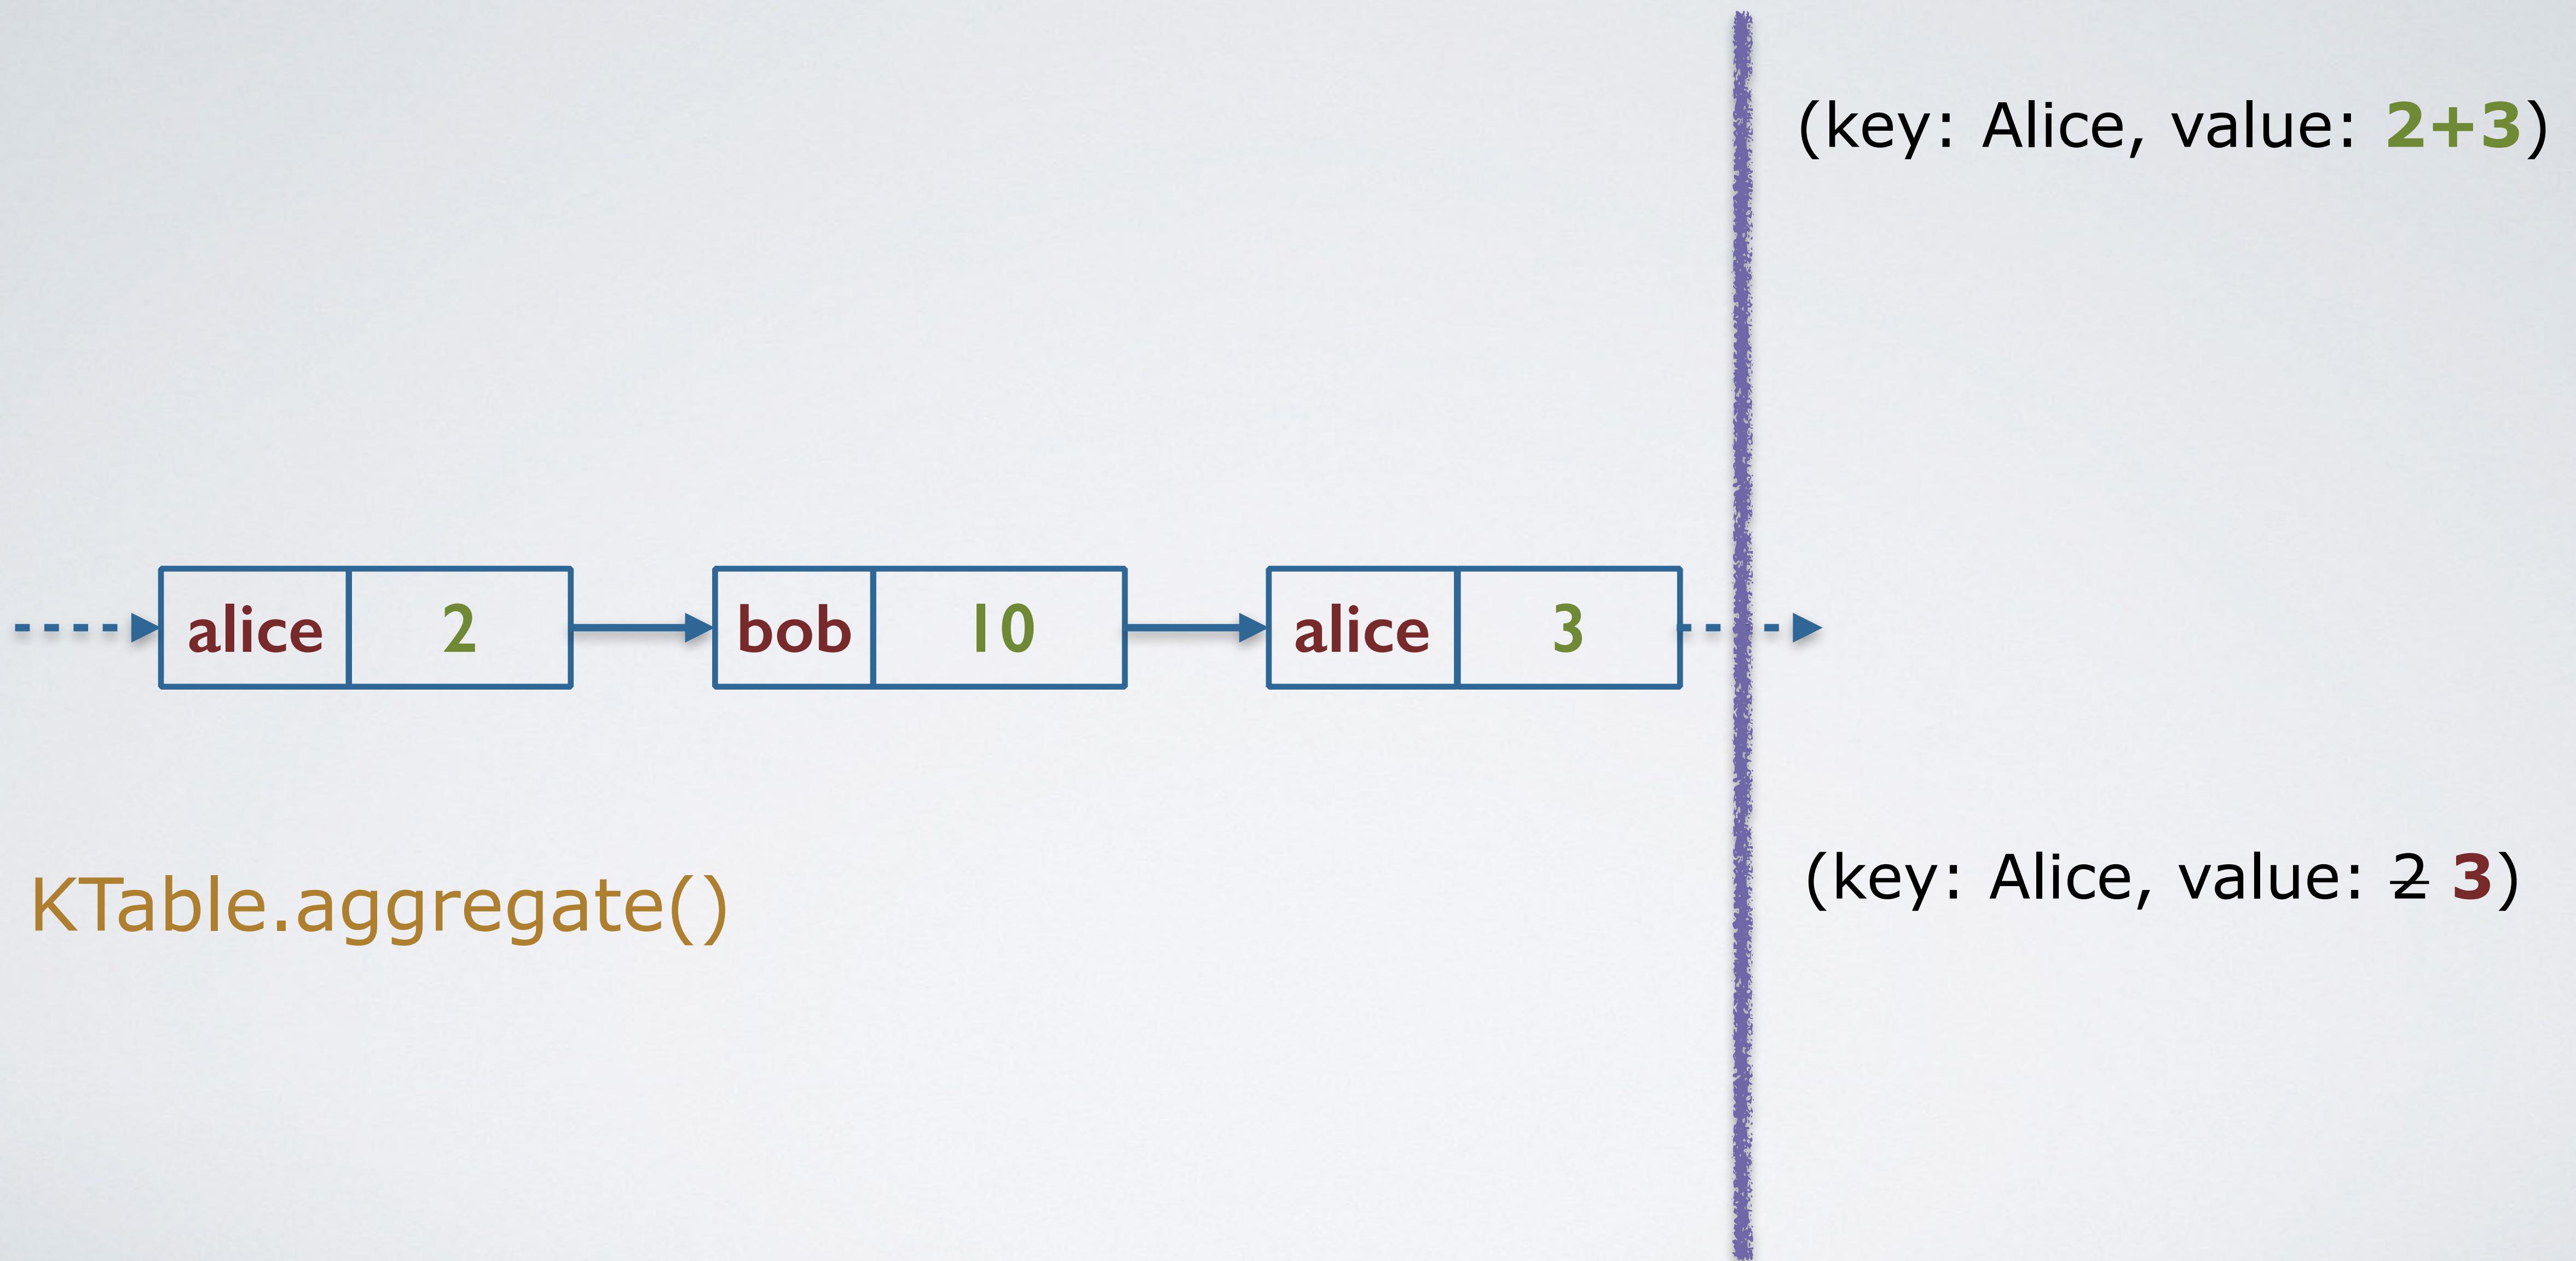
KStream.aggregate()

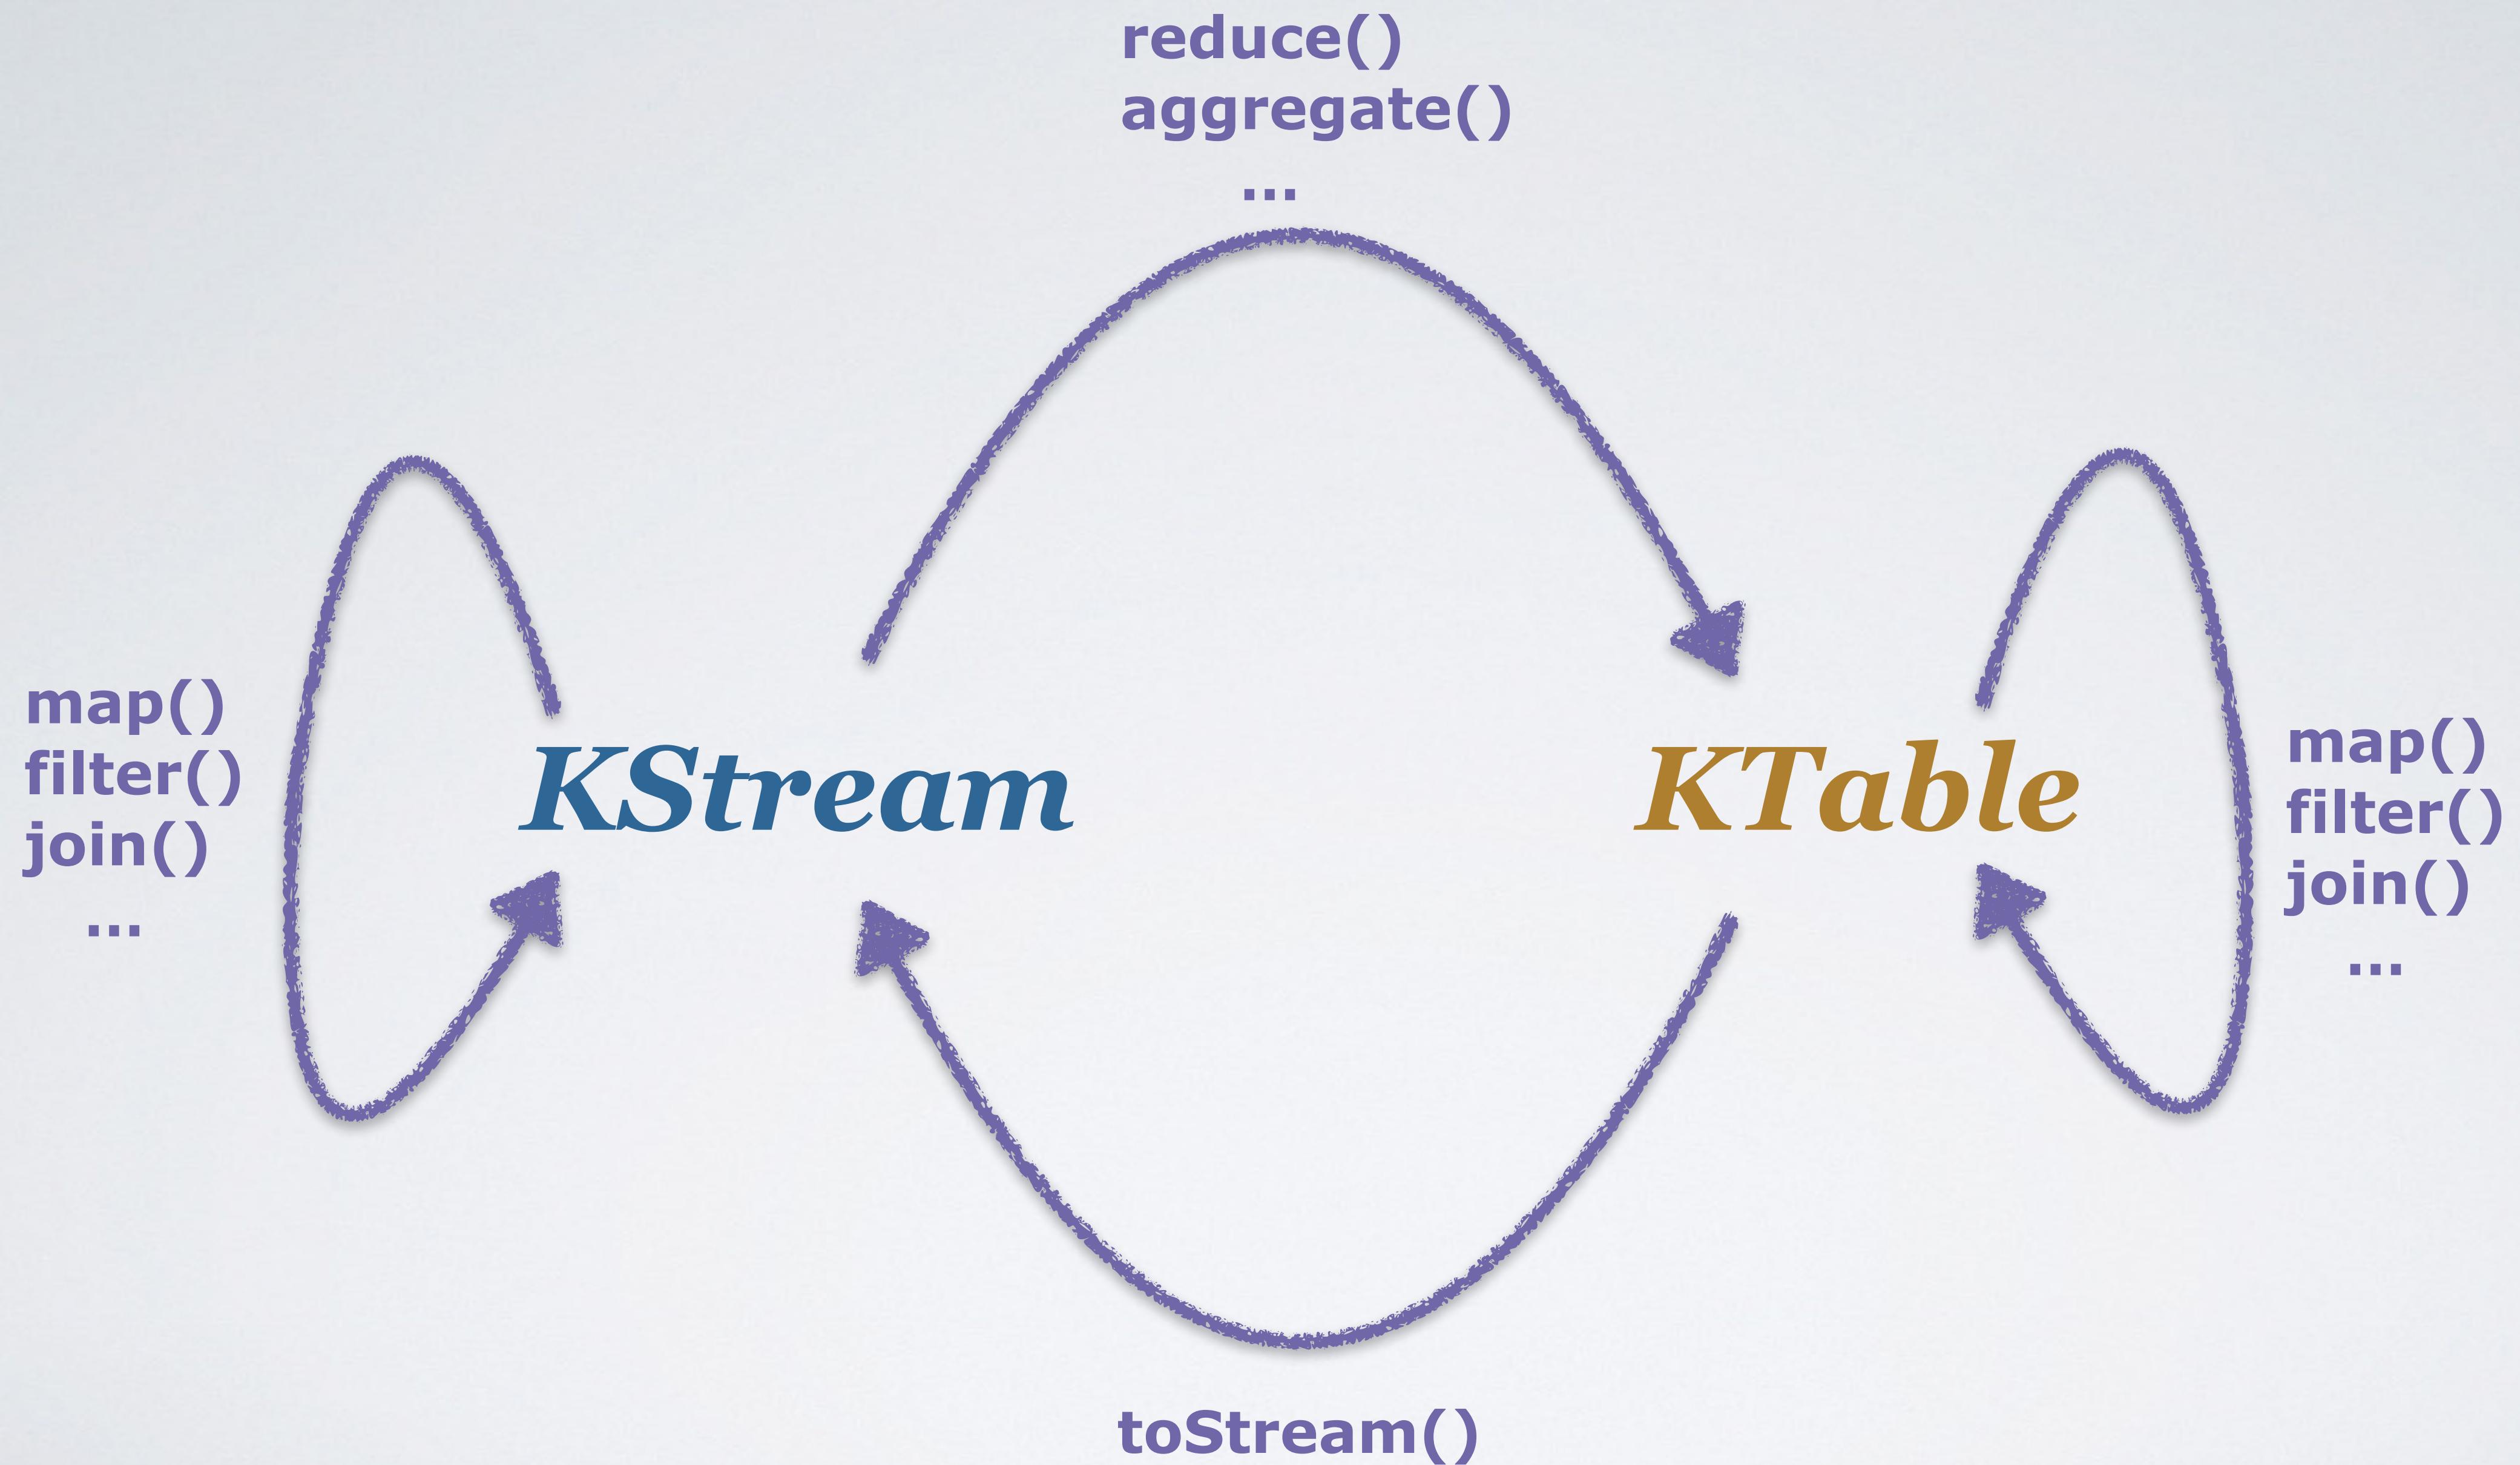
time



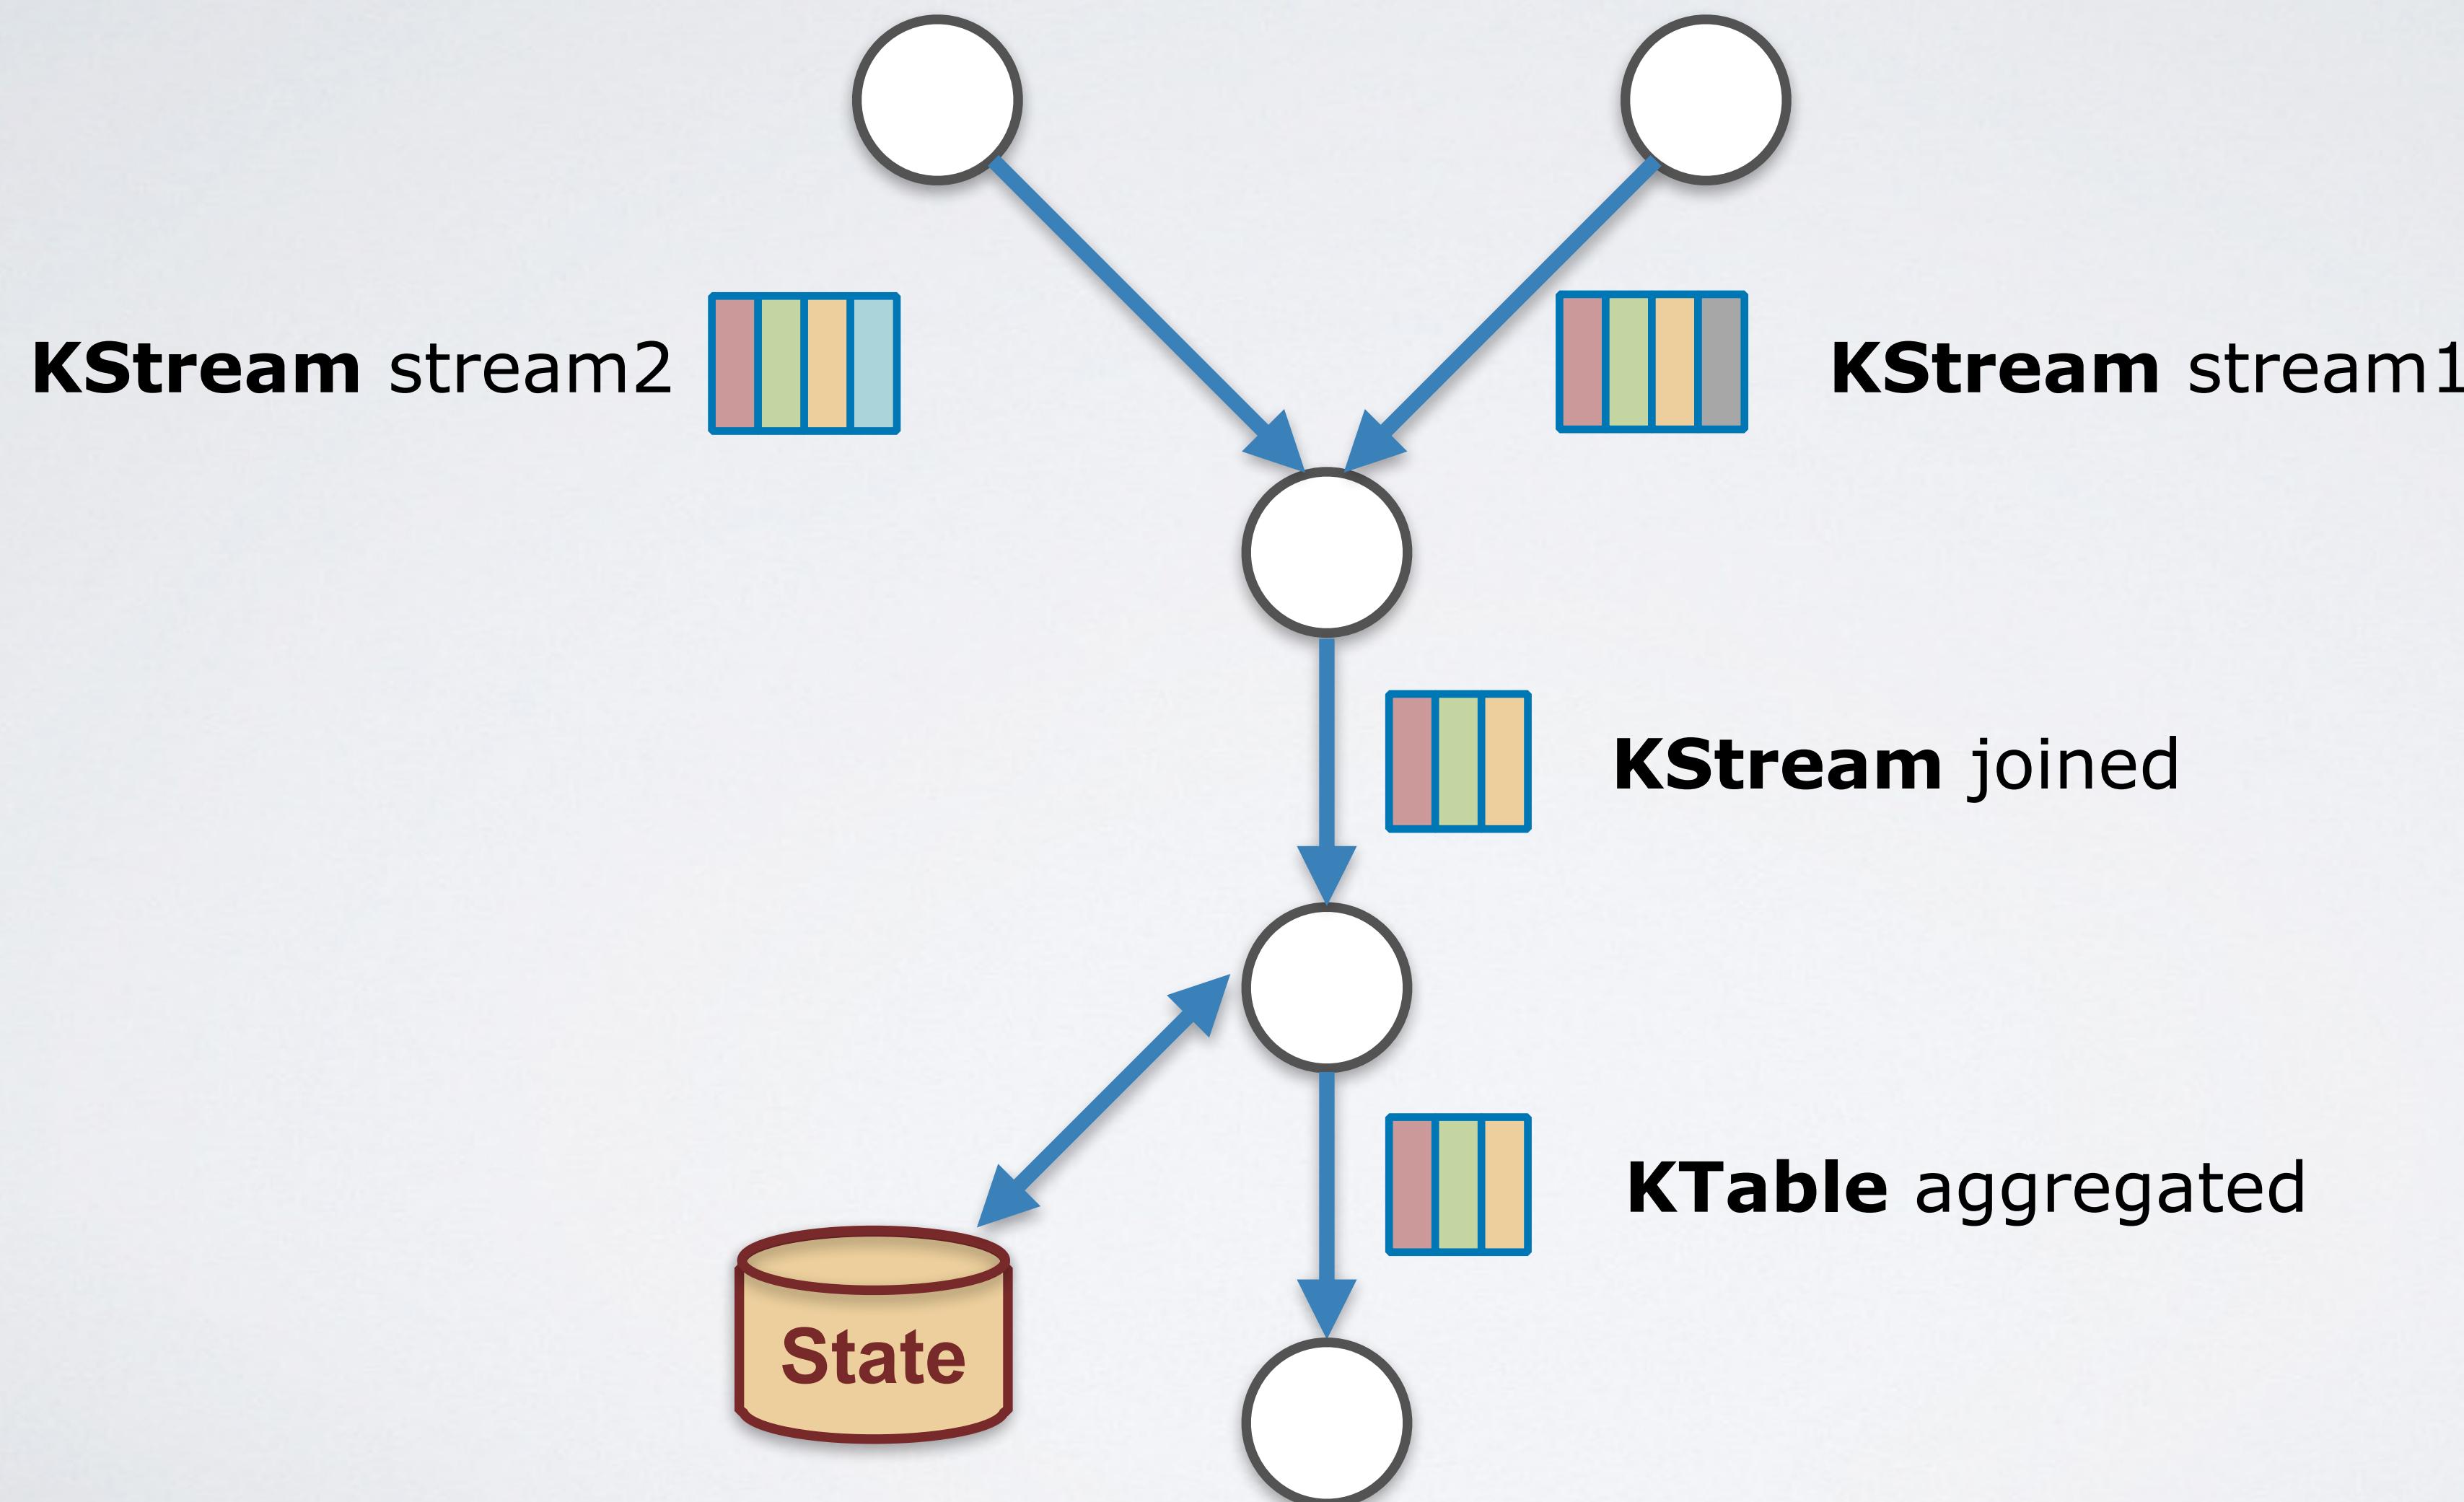
KTable.aggregate()

KStream.aggregate()

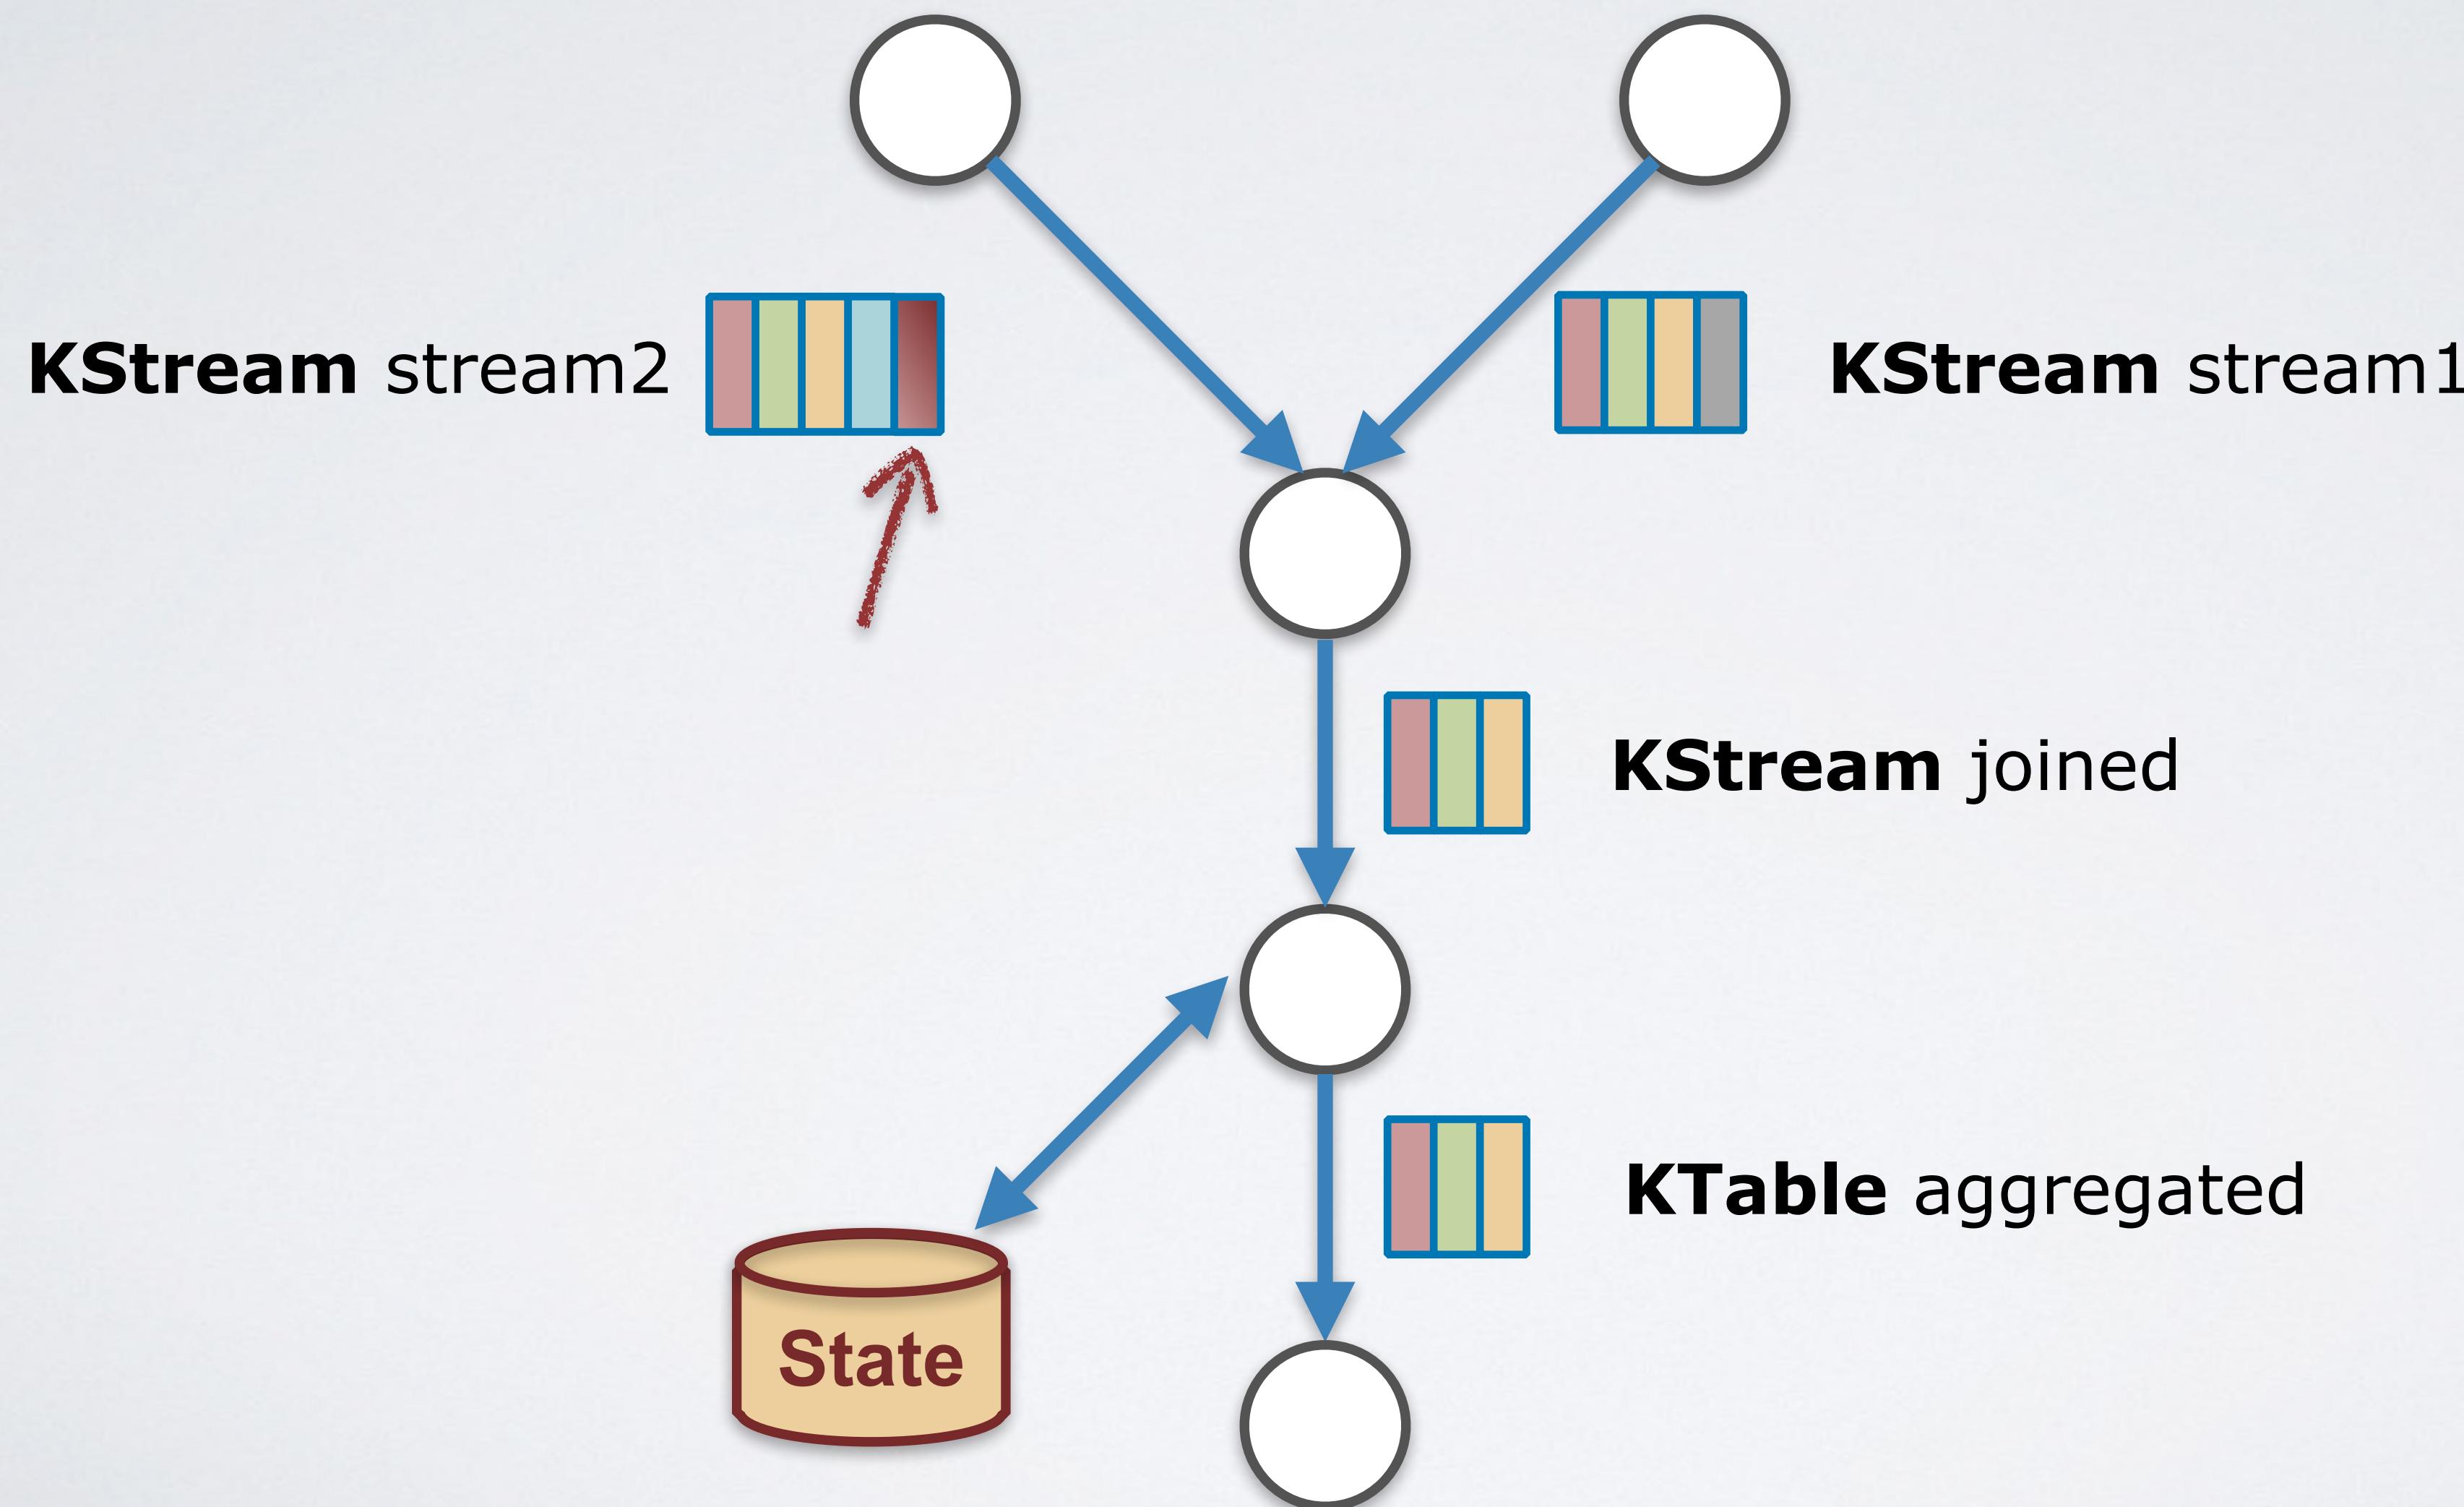




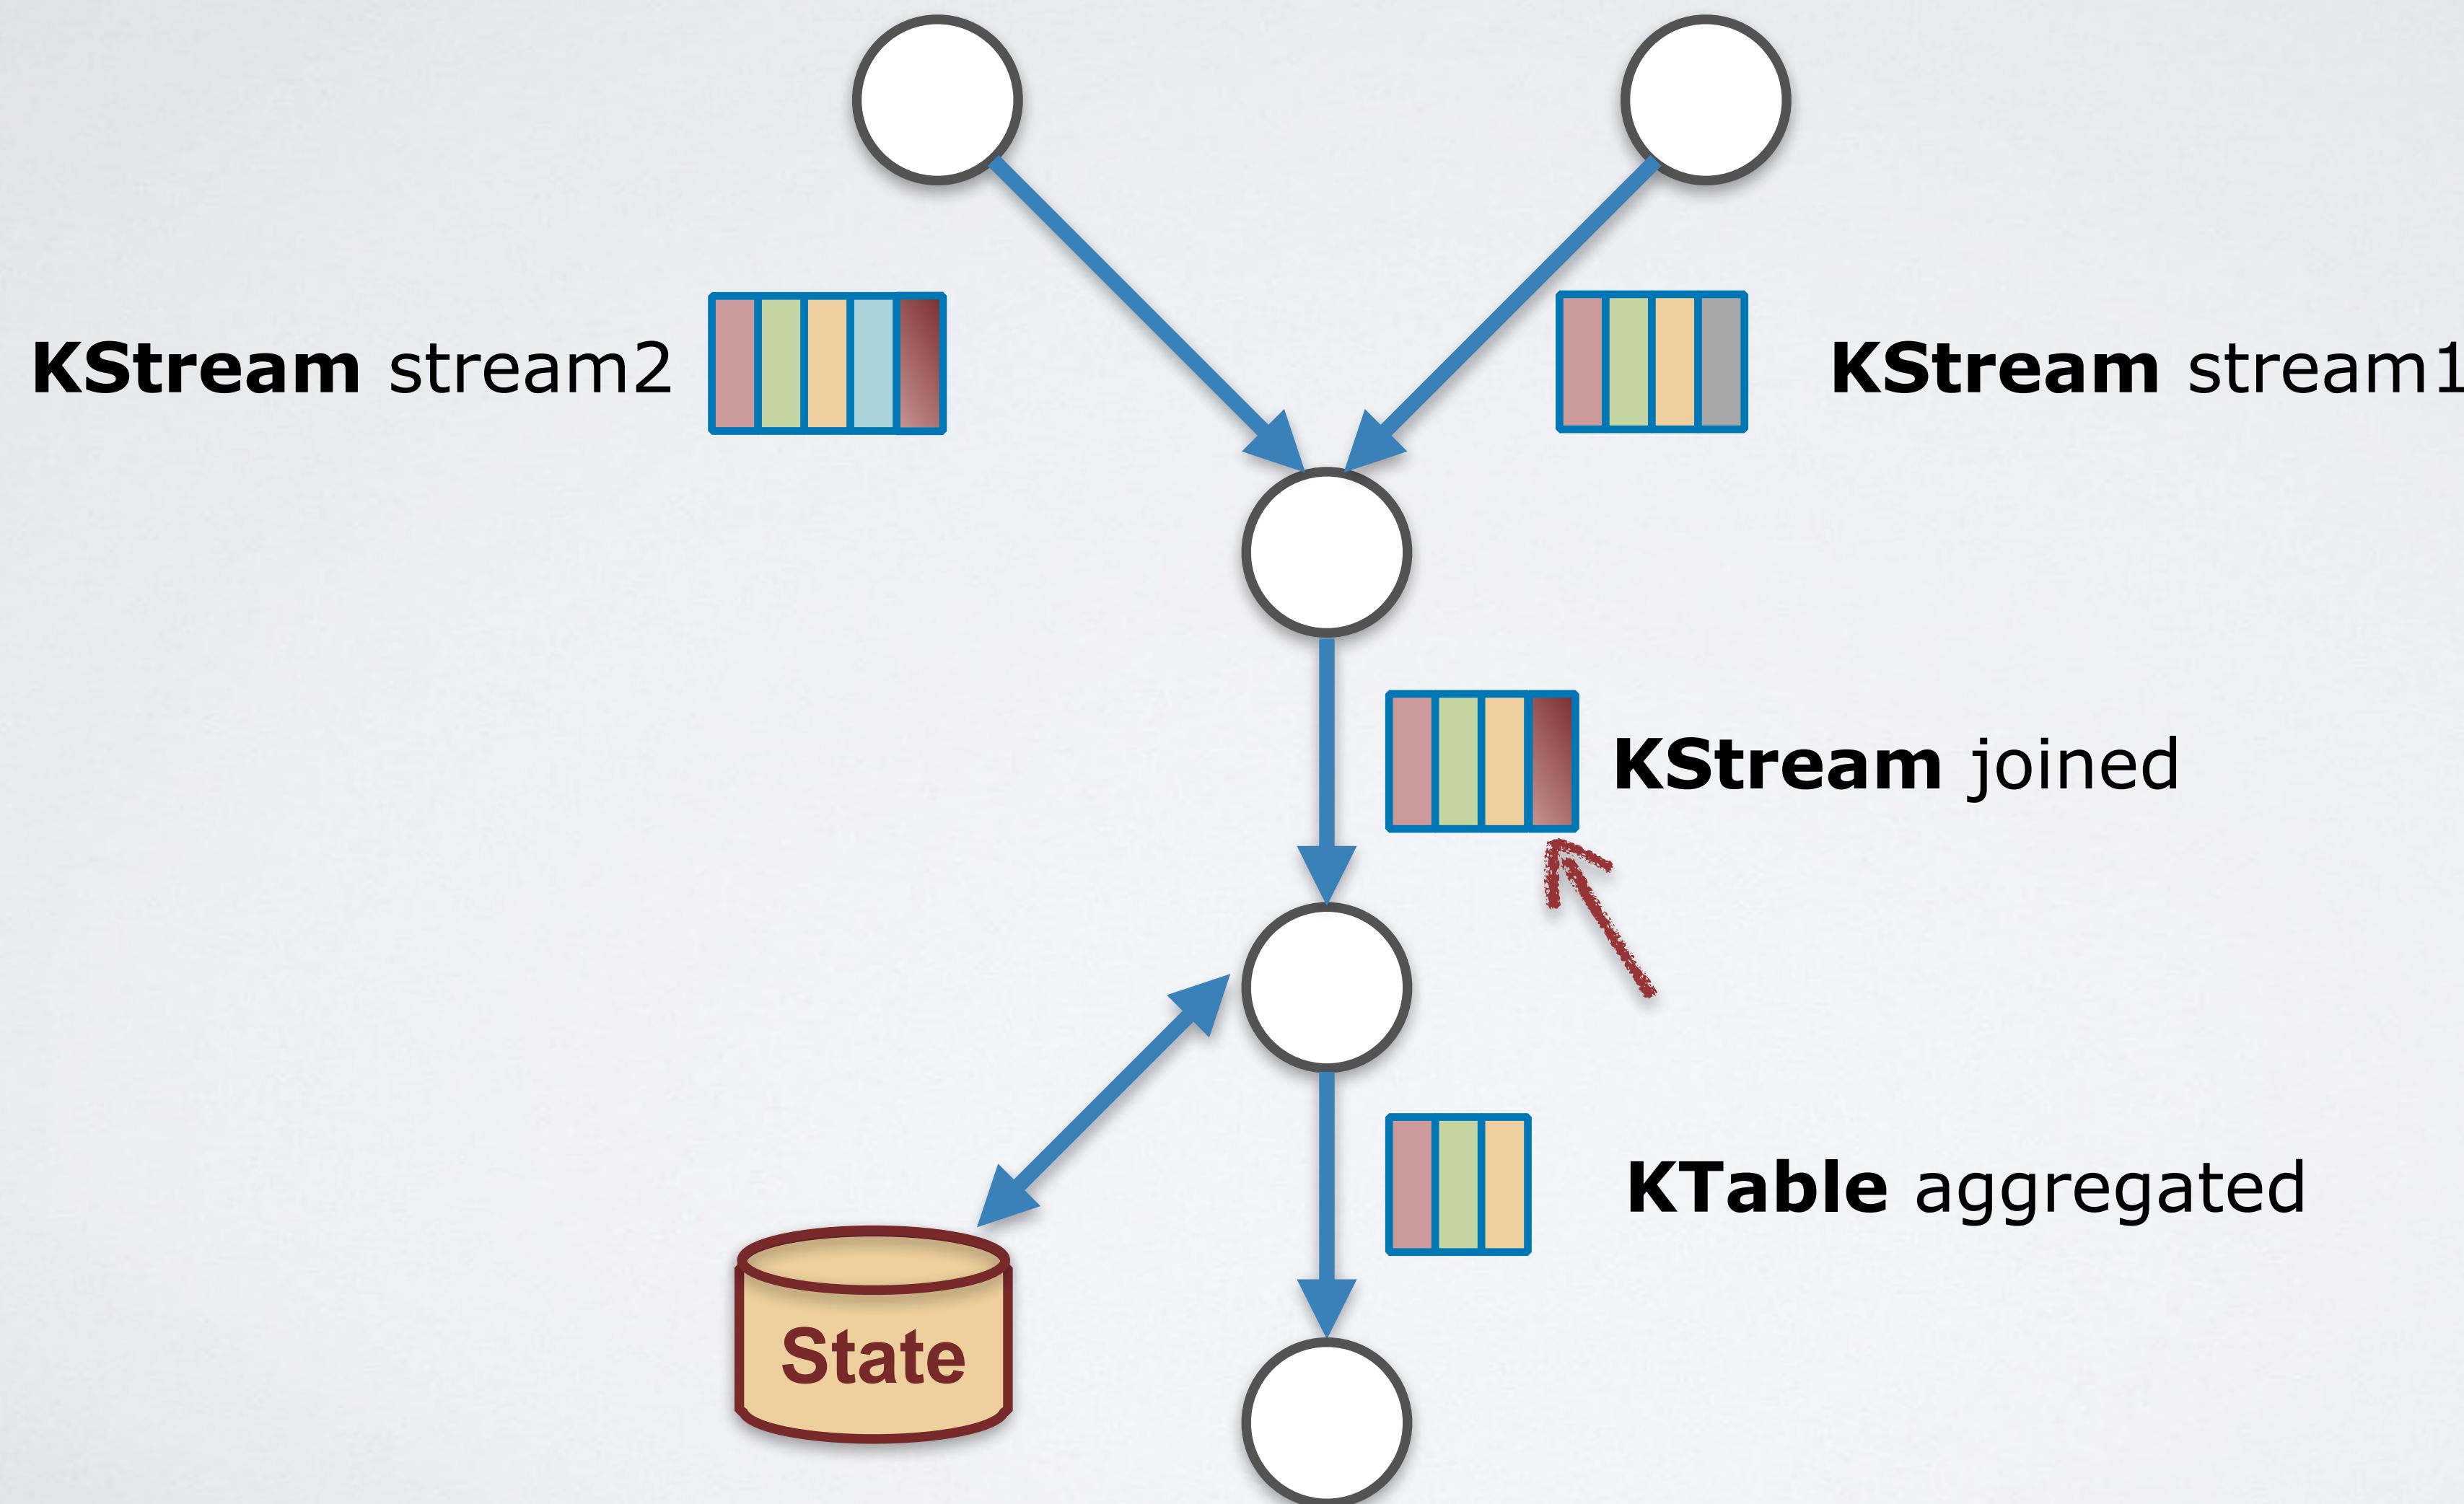
Updates Propagation in KTable



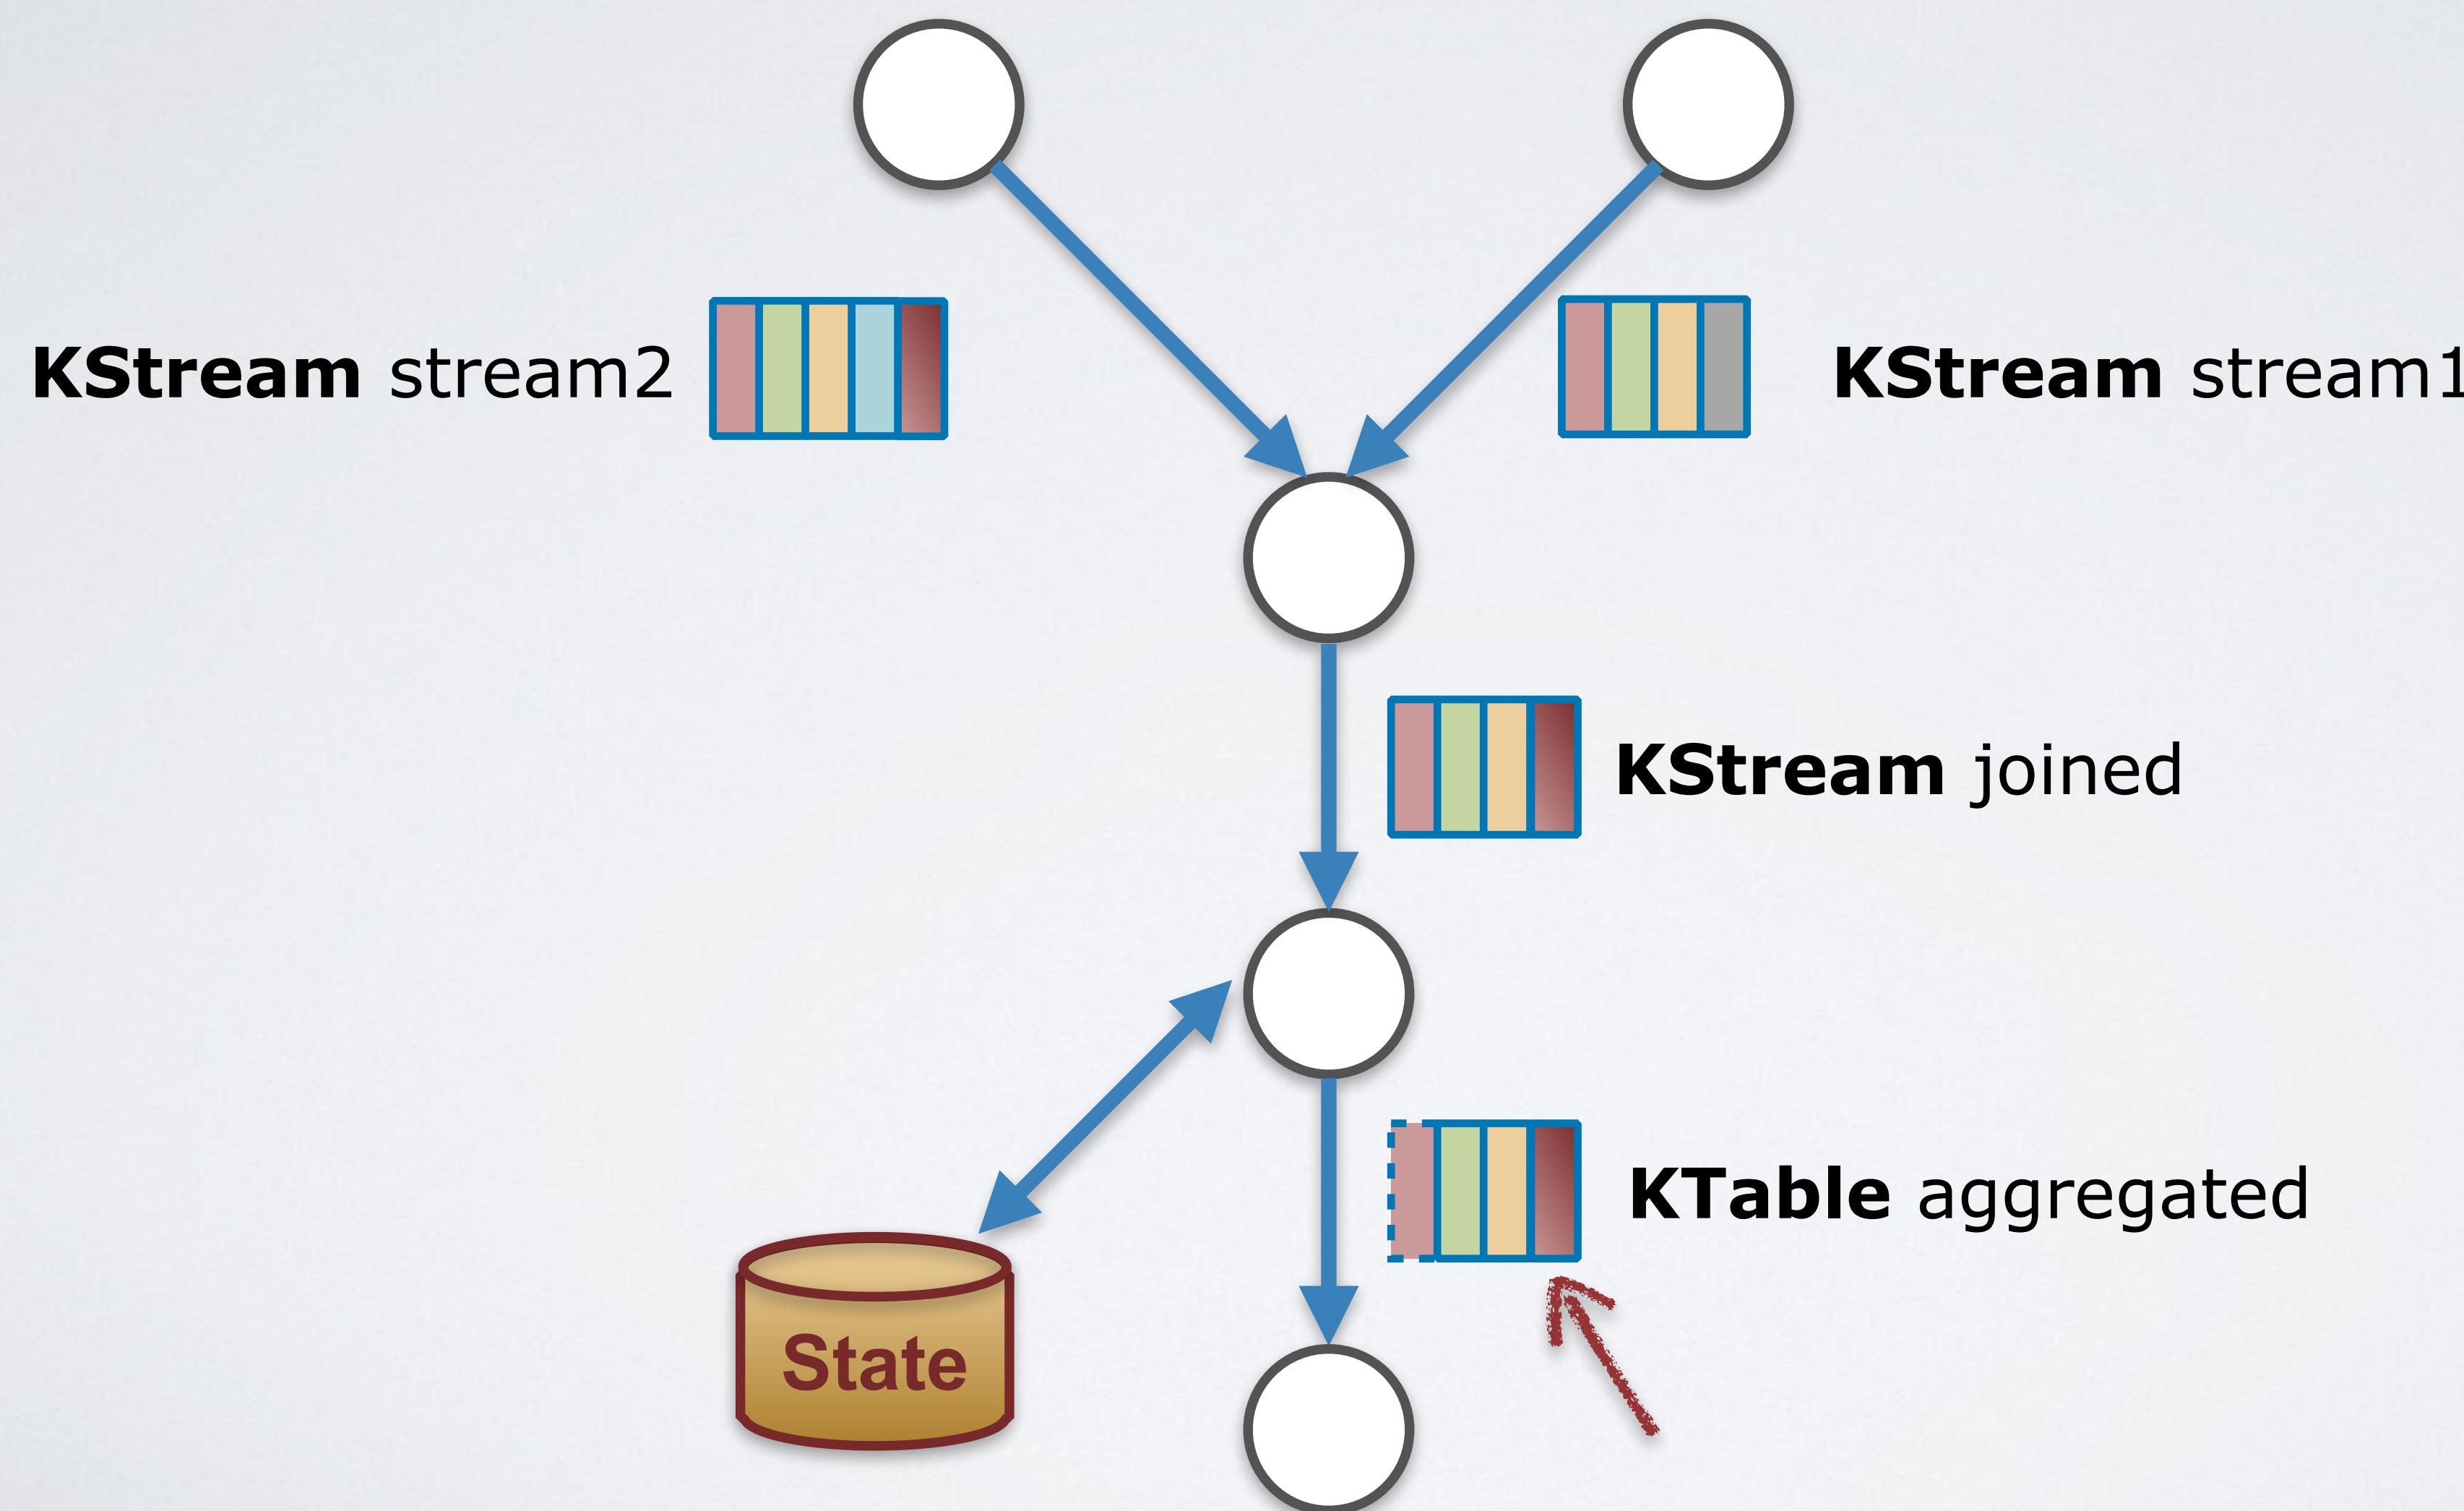
Updates Propagation in KTable



Updates Propagation in KTable



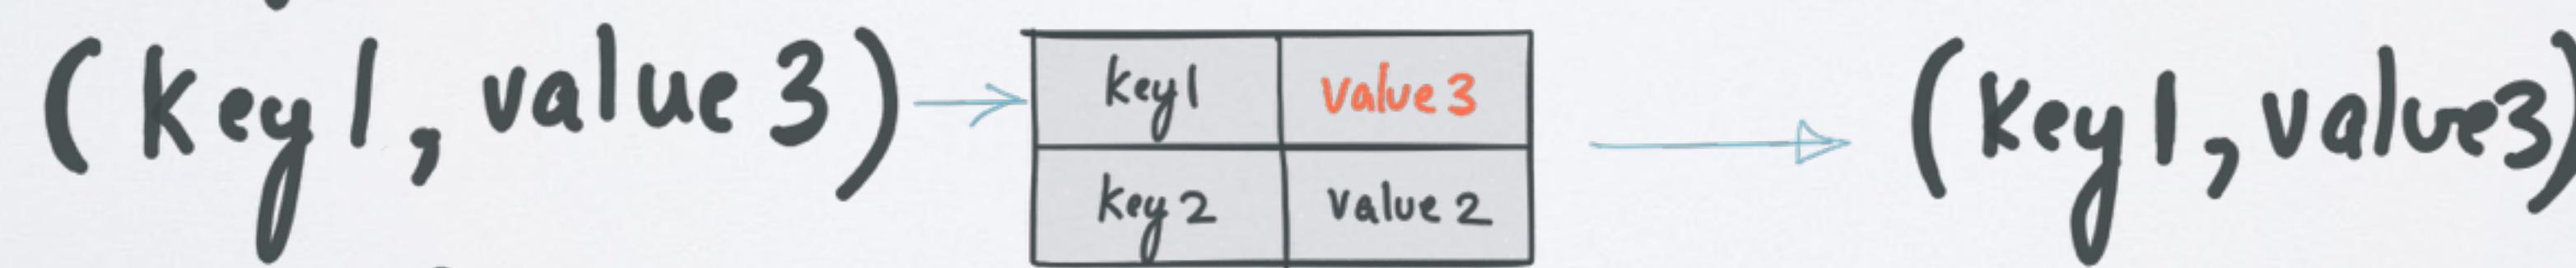
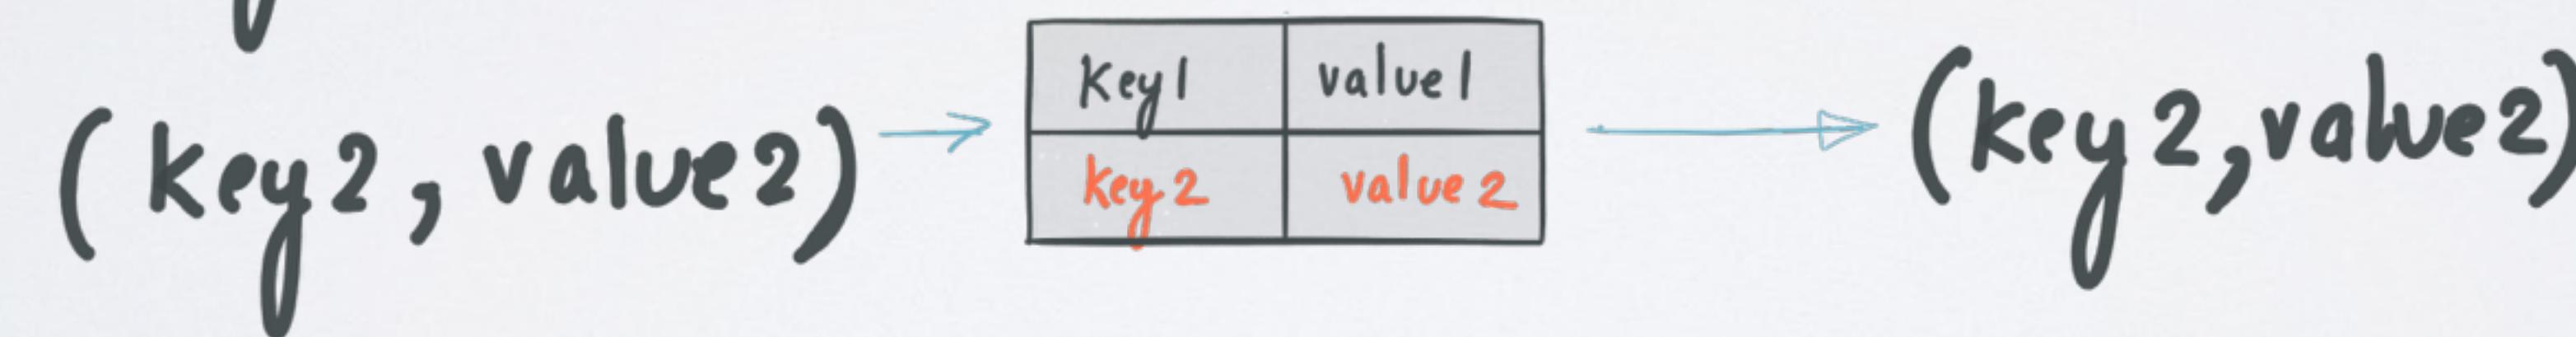
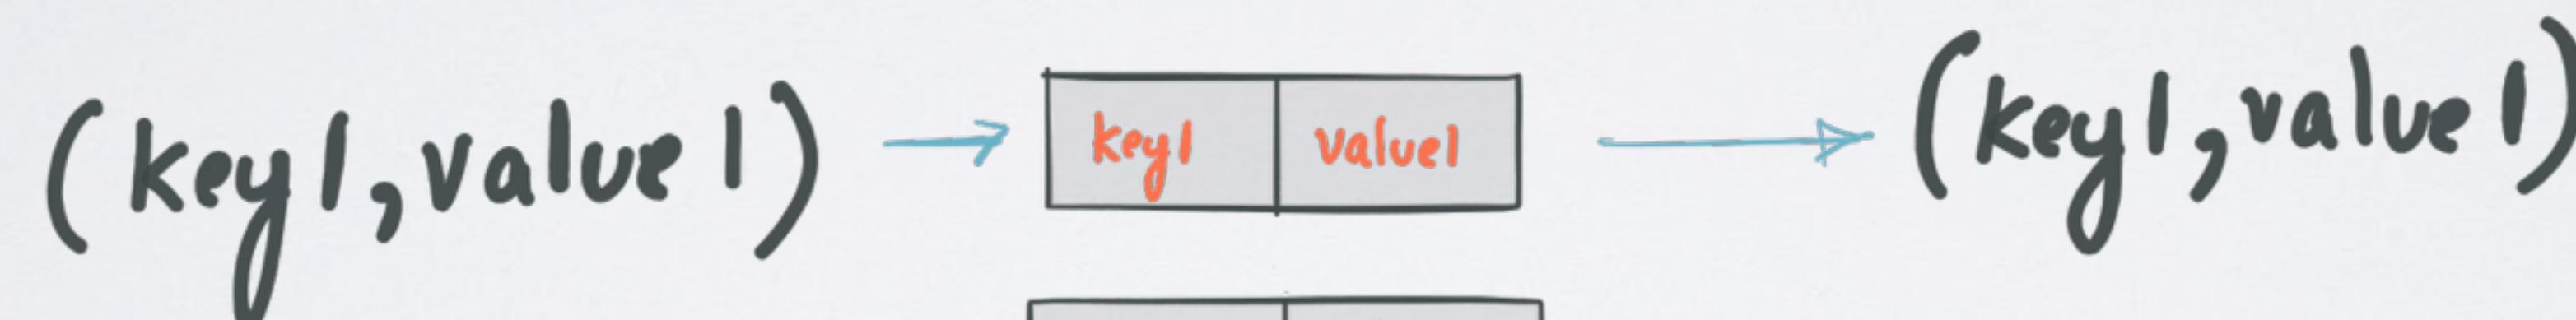
Updates Propagation in KTable



What about *Fault Tolerance*?

Remember?

TABLES \approx STREAMS



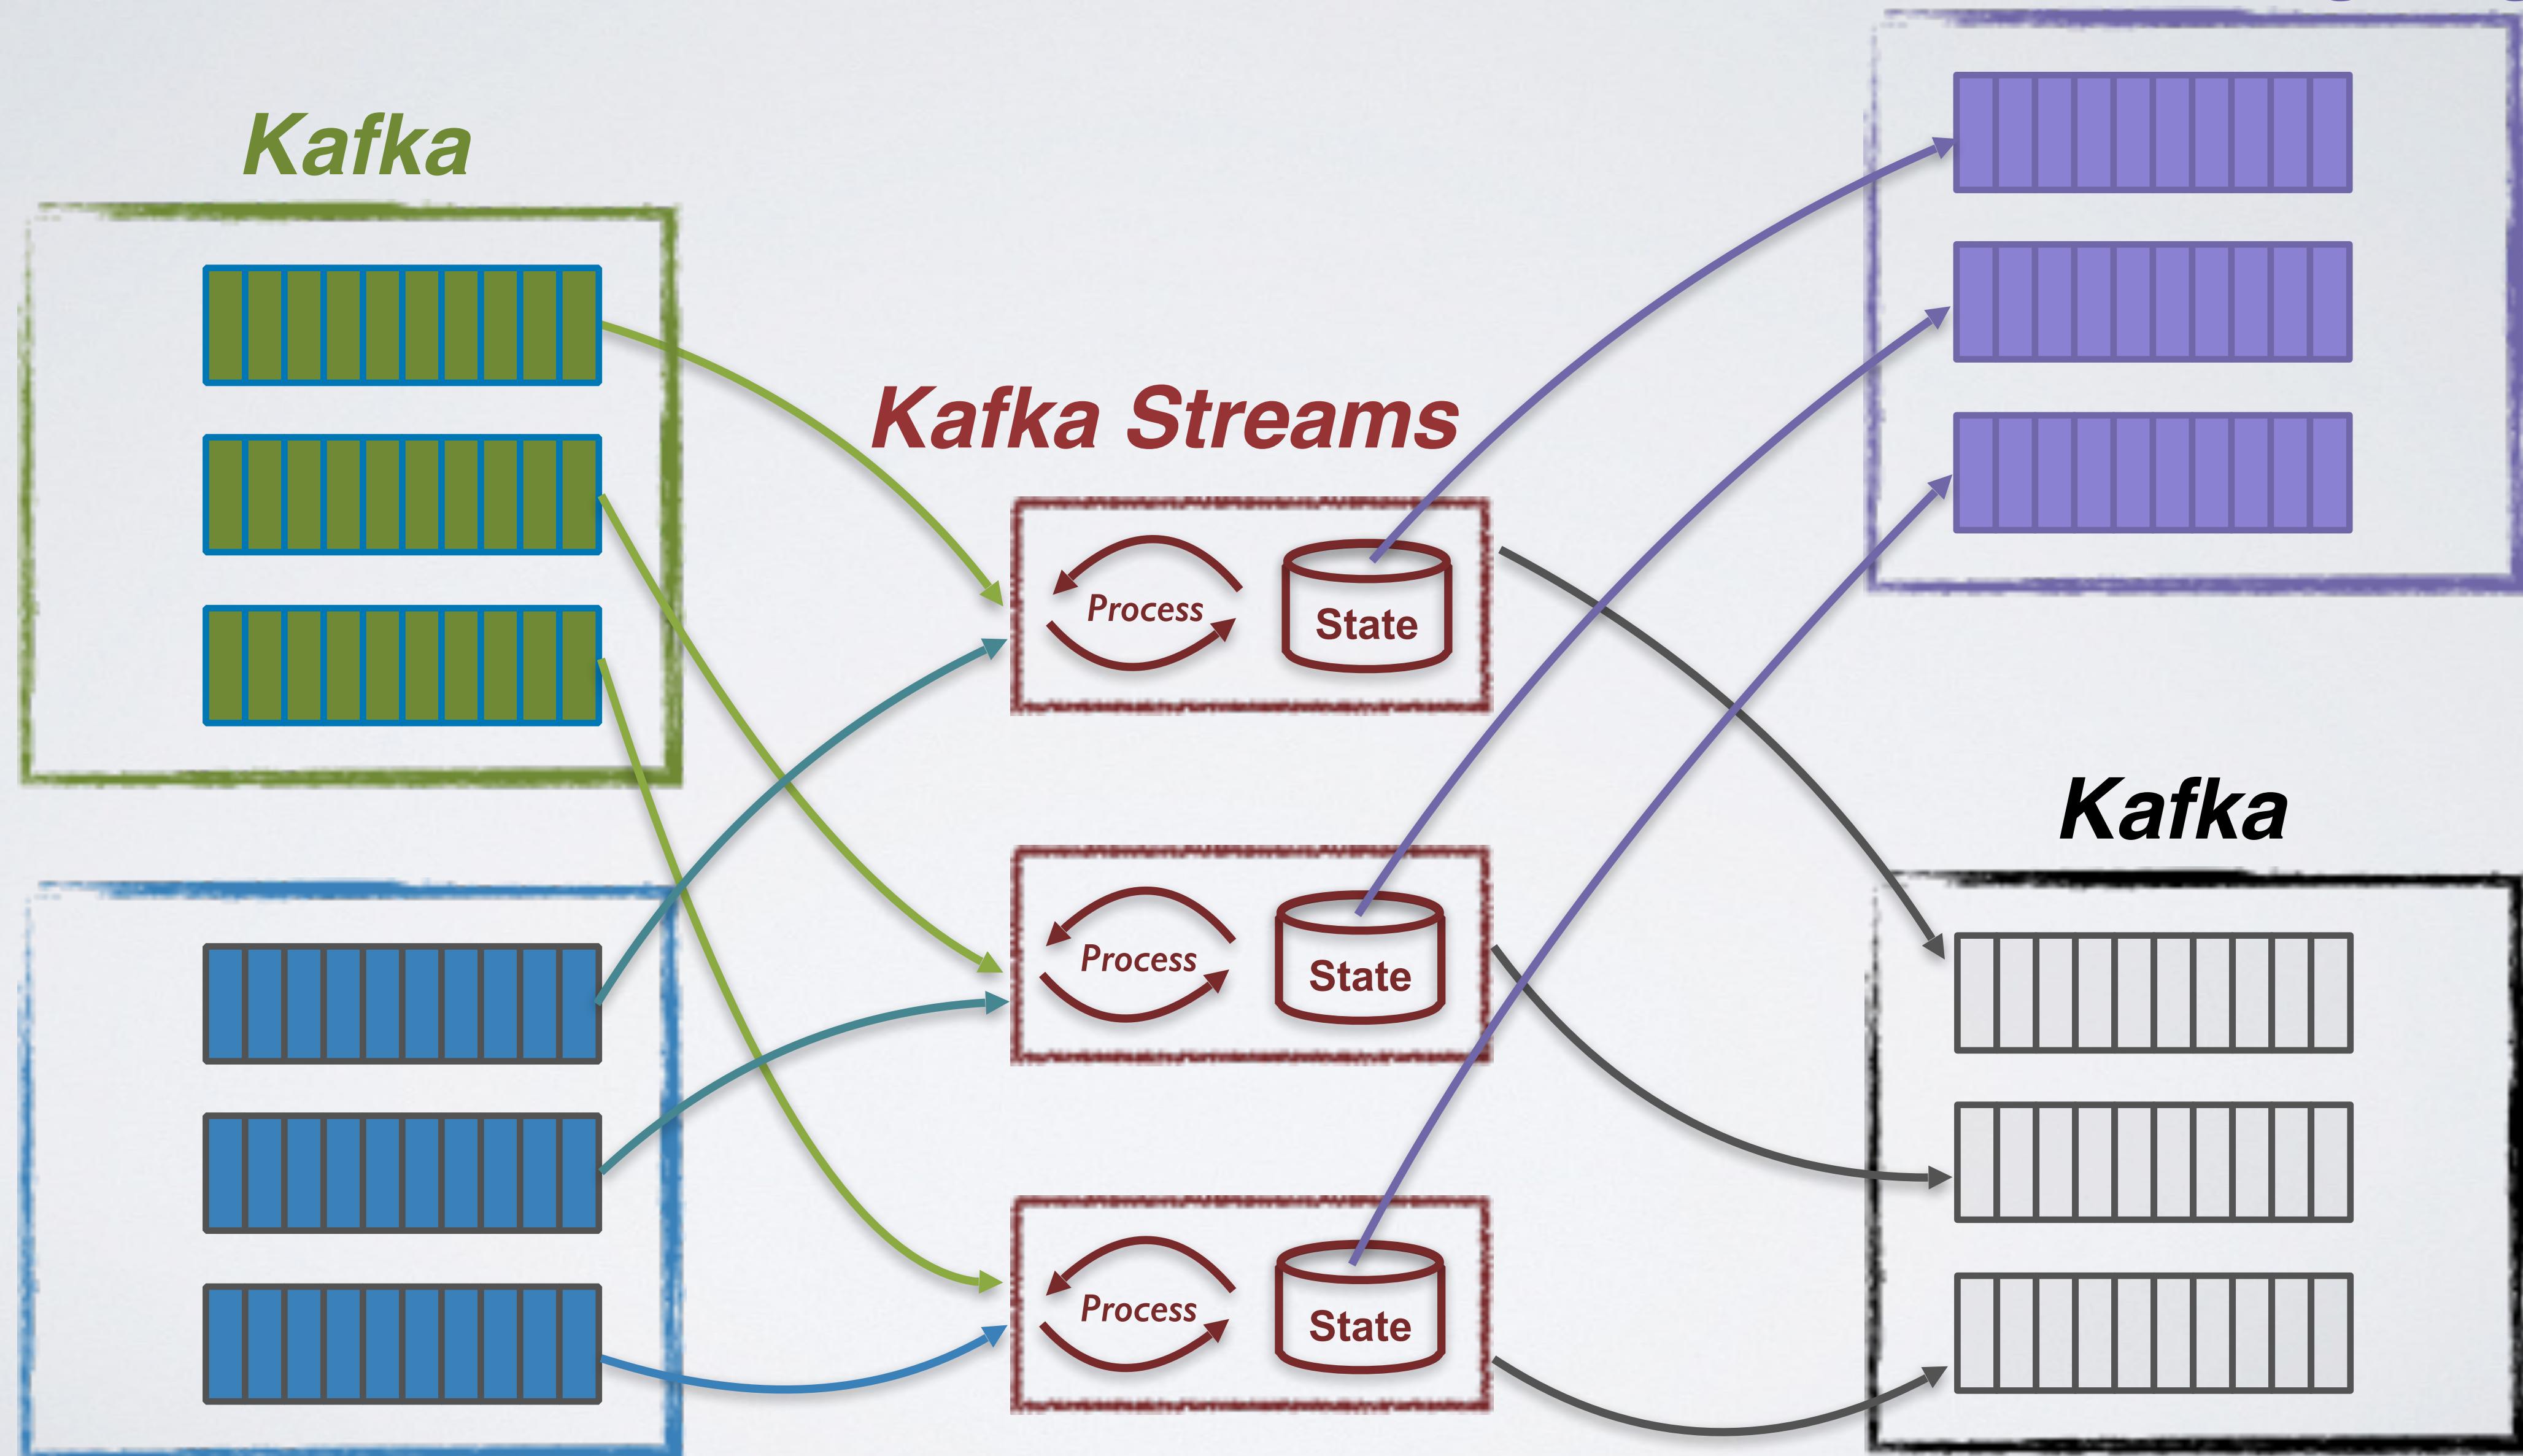
.

.

.

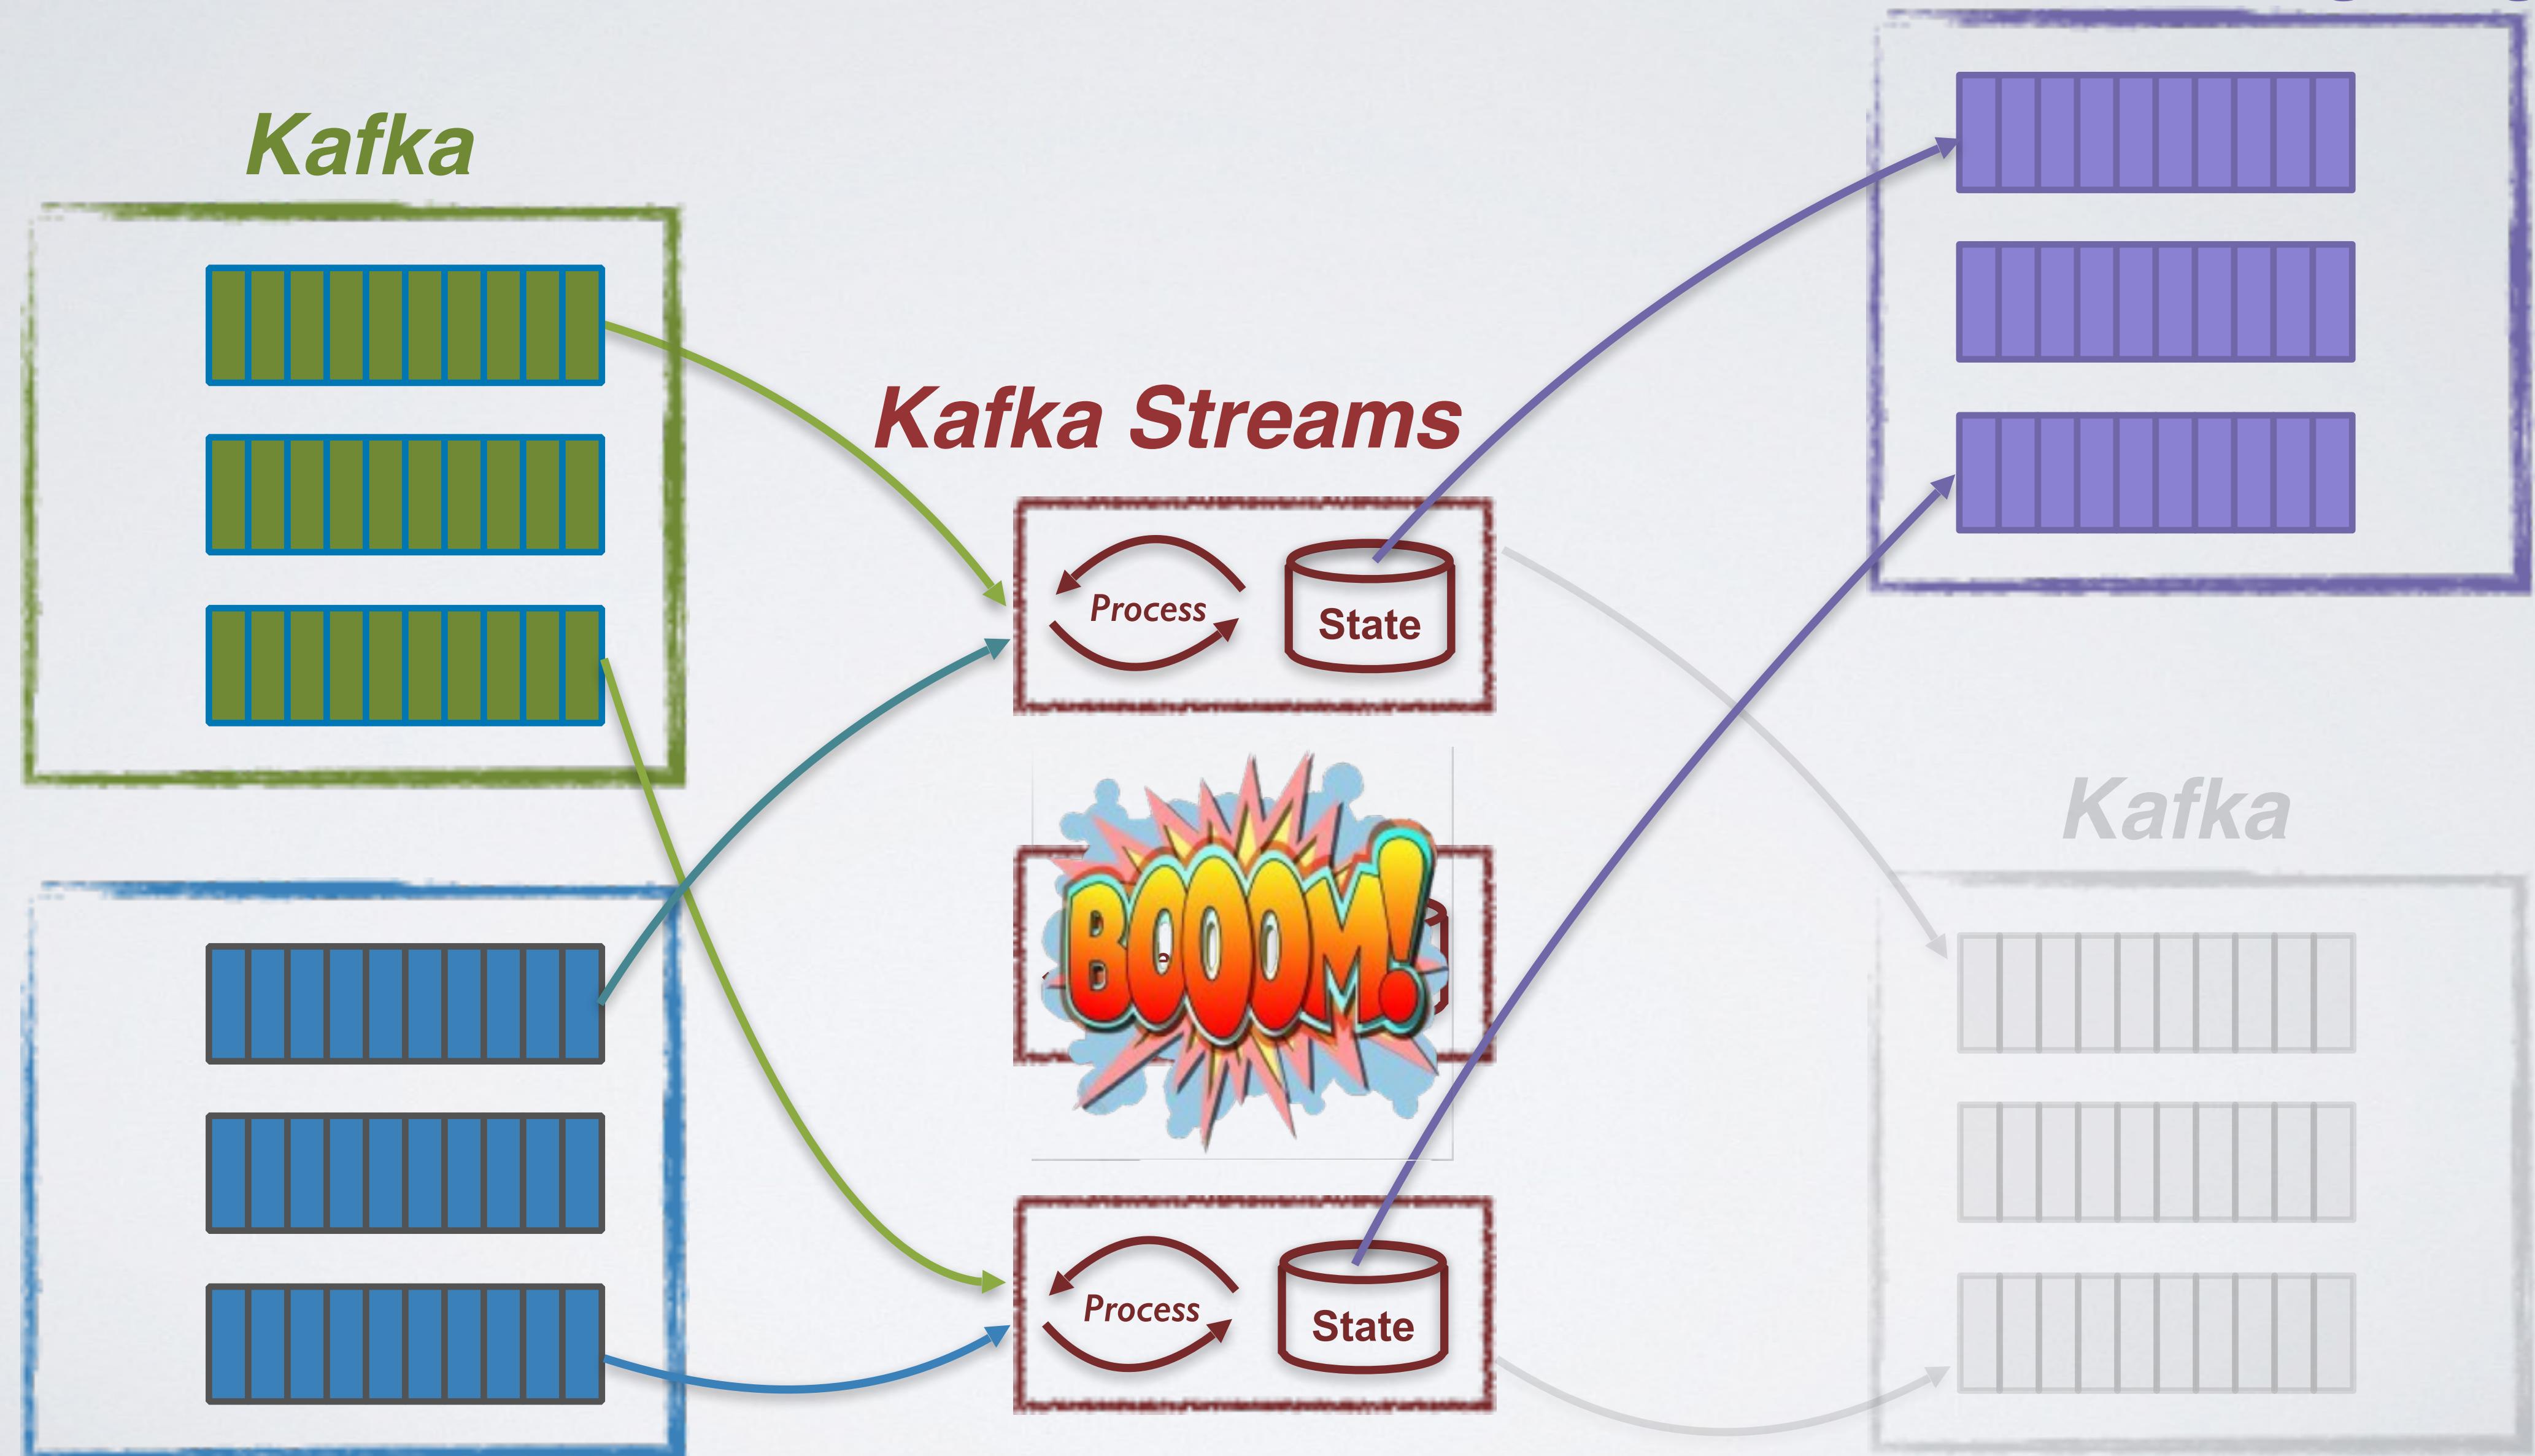
Fault Tolerance

Kafka Changelog



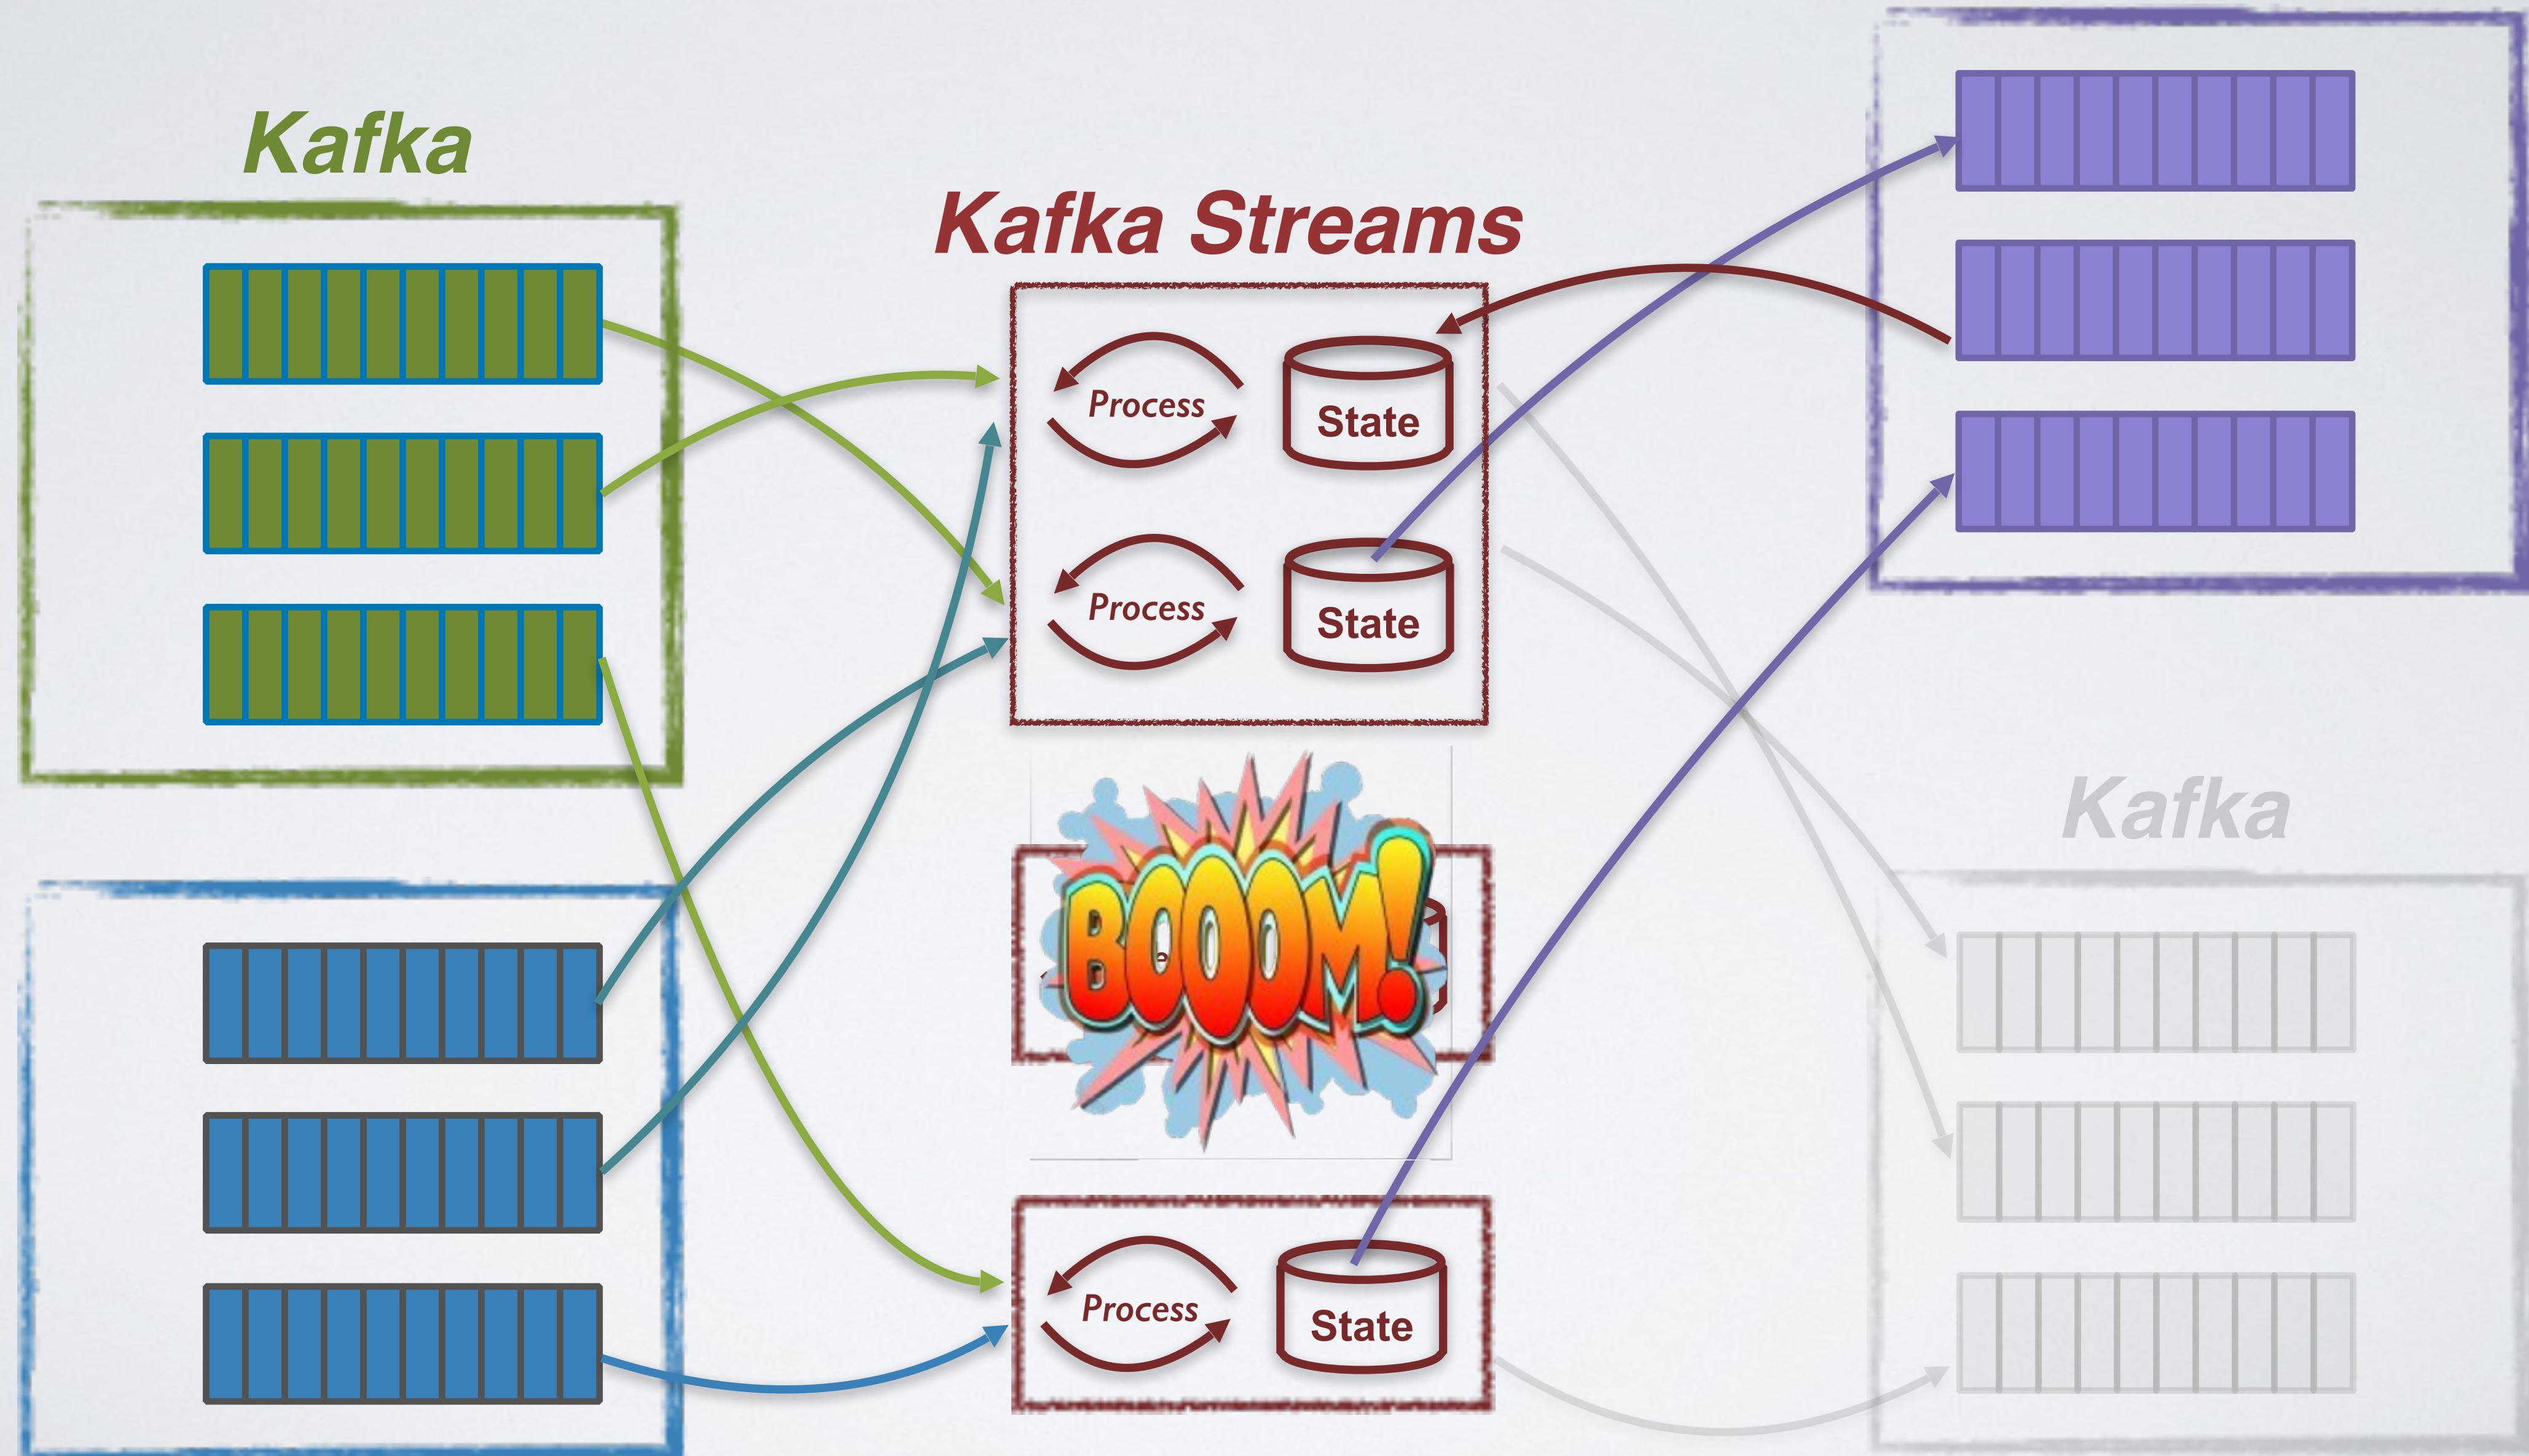
Fault Tolerance

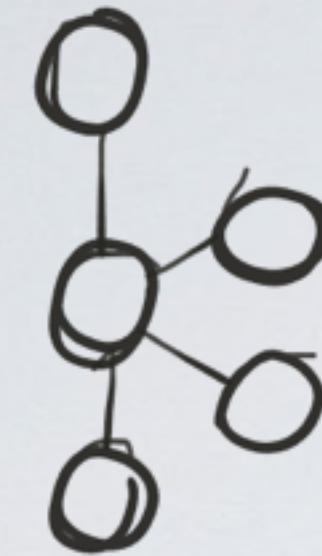
Kafka Changelog



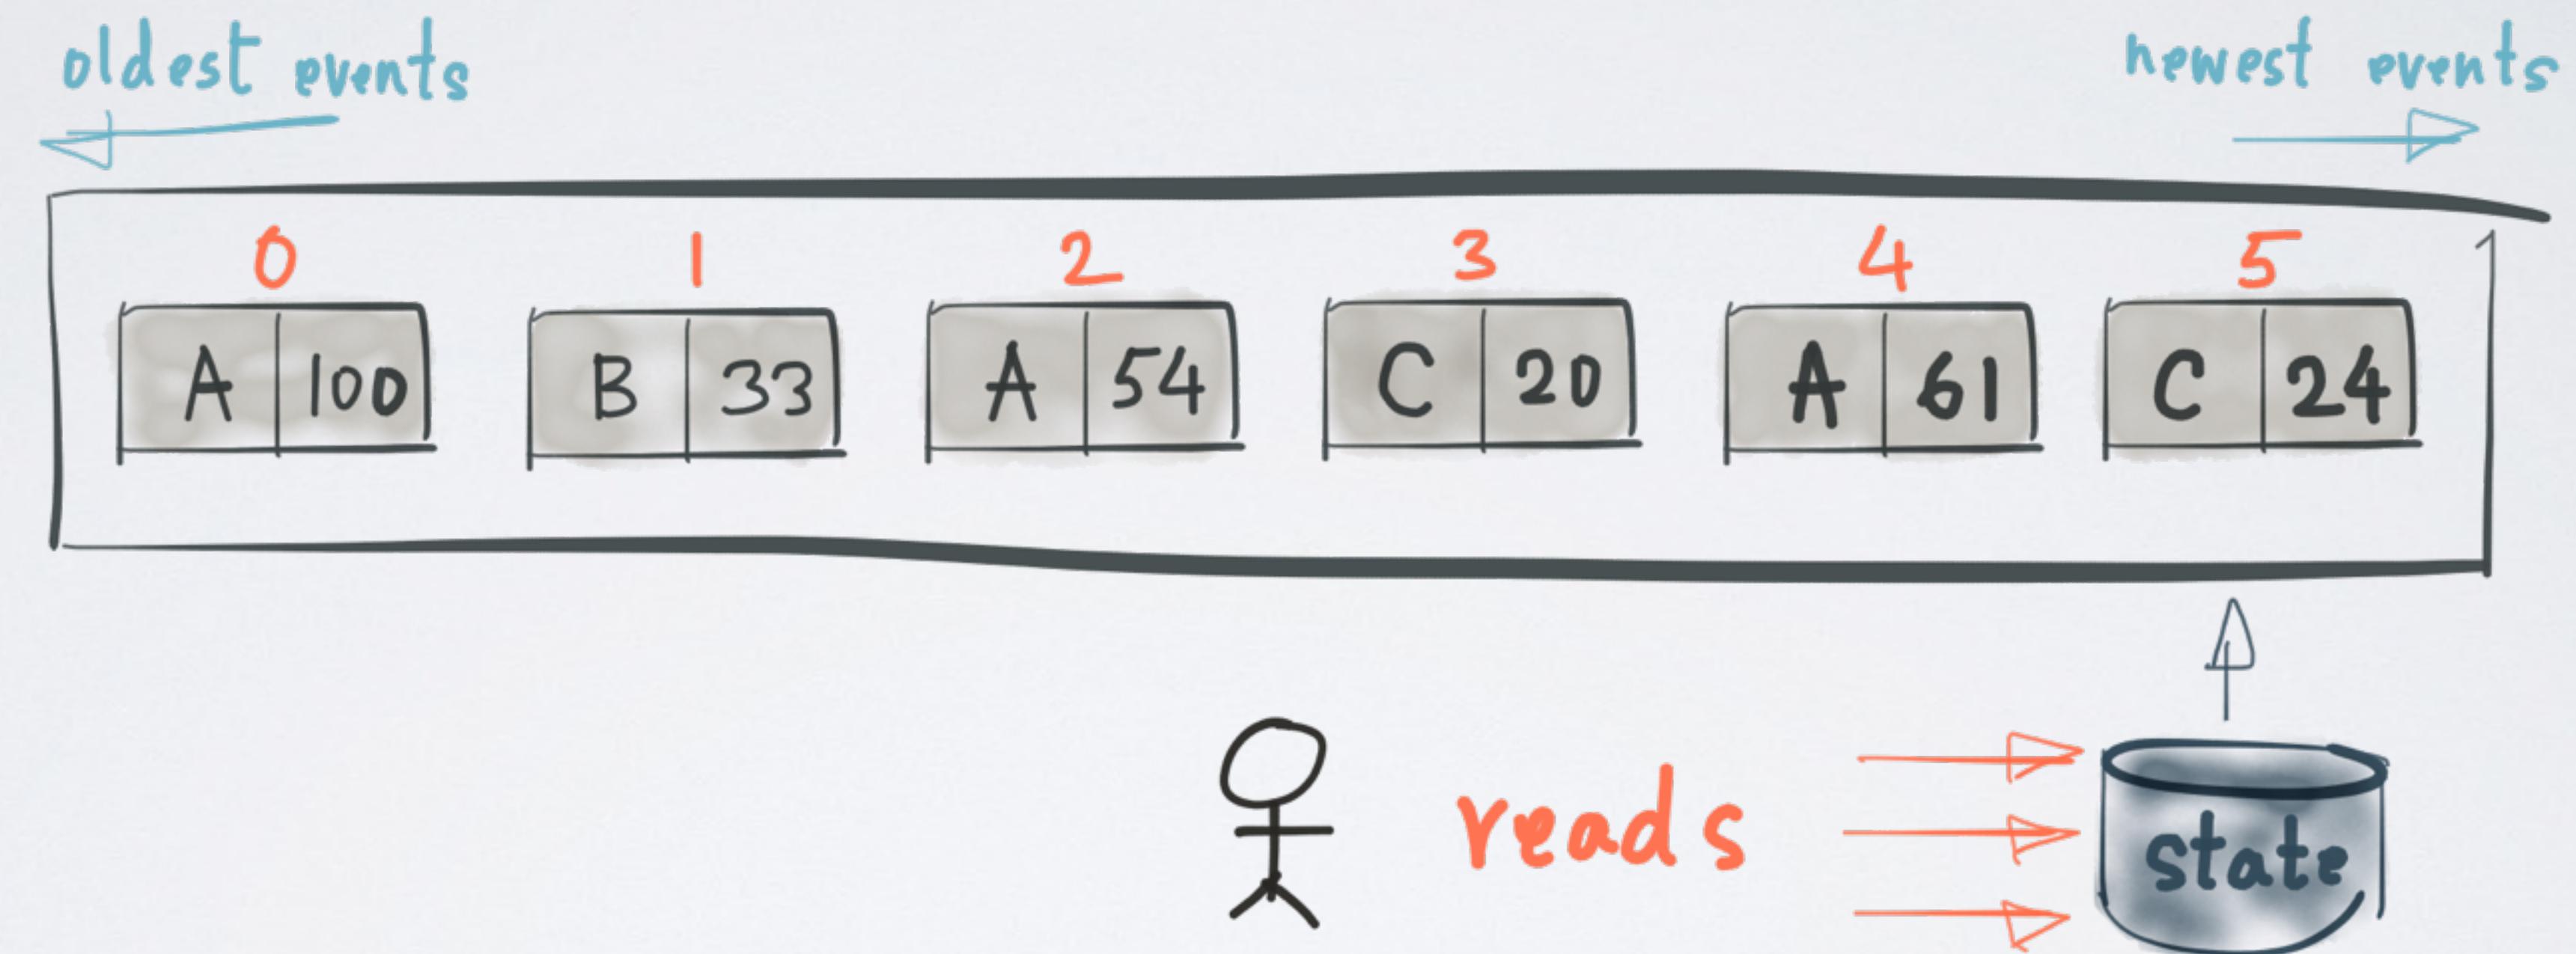
Fault Tolerance

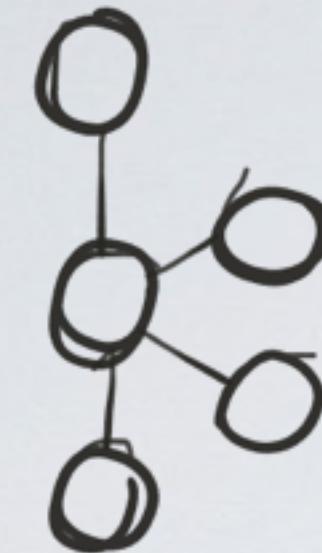
Kafka Changelog



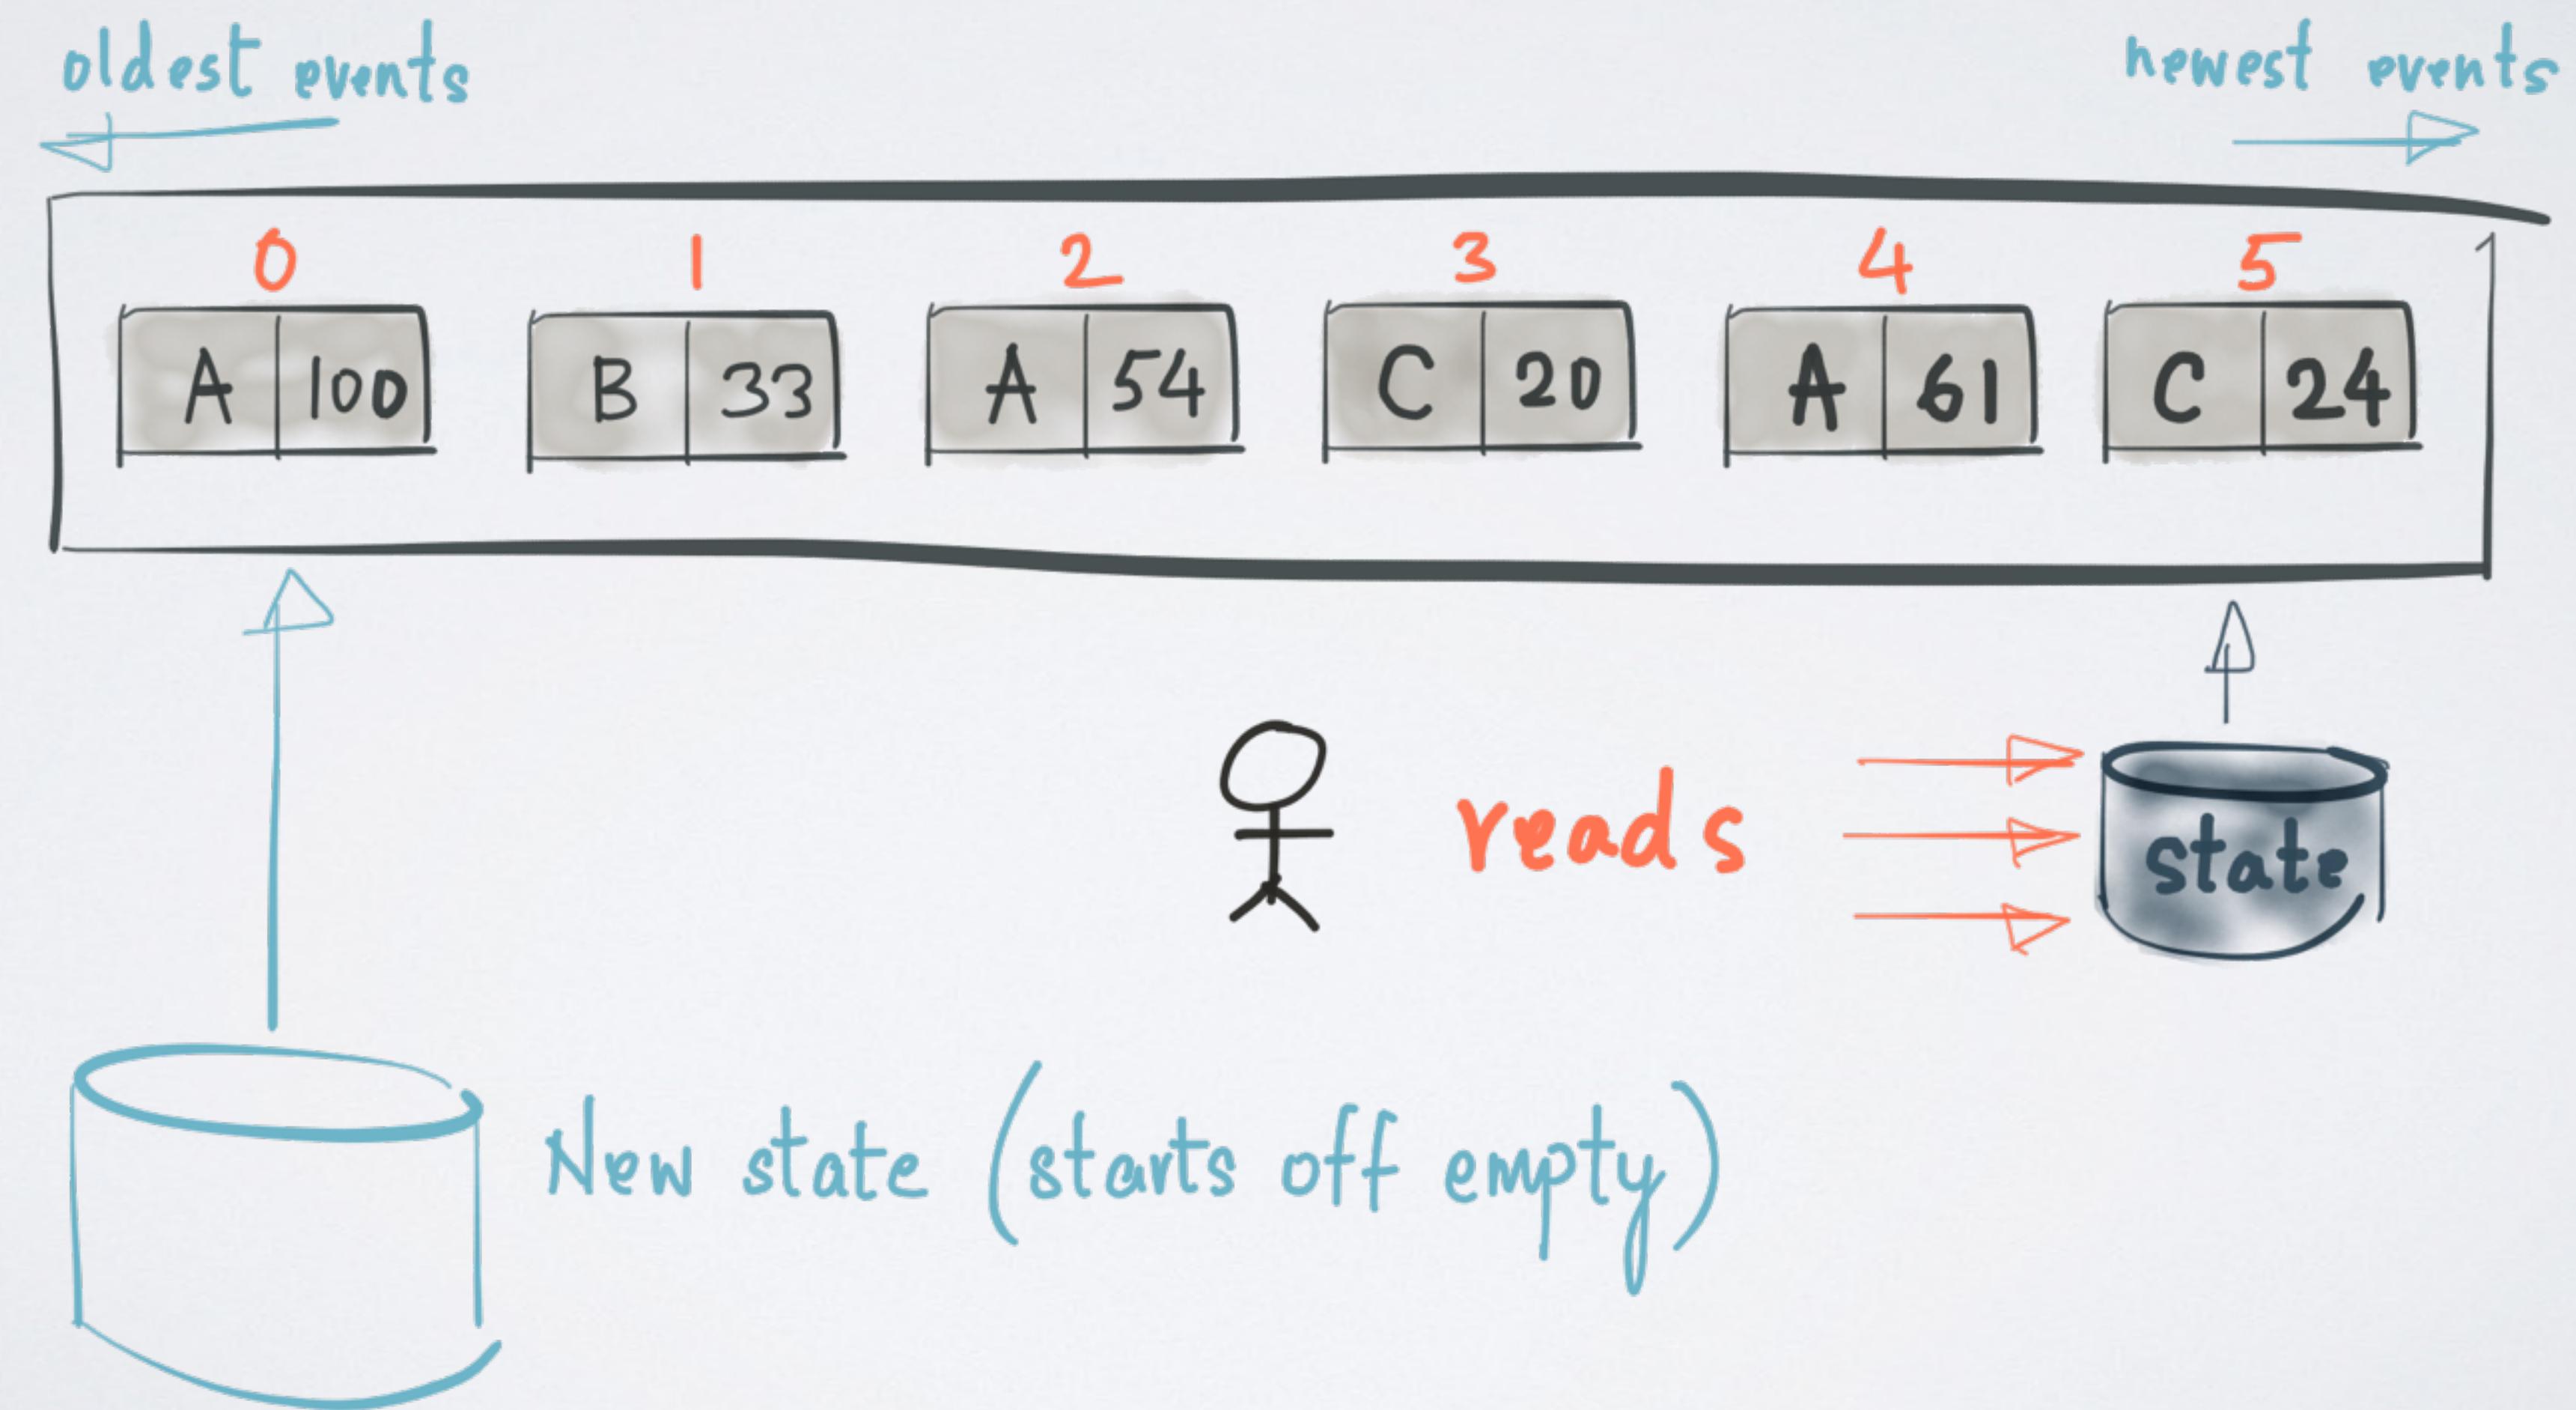


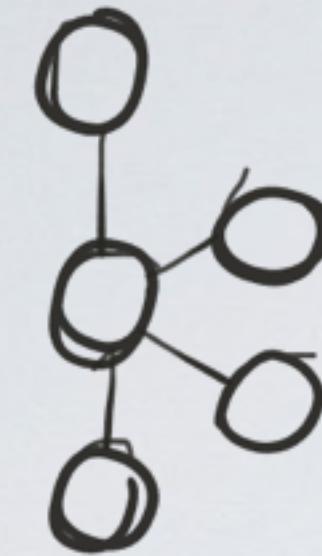
REPROCESSING



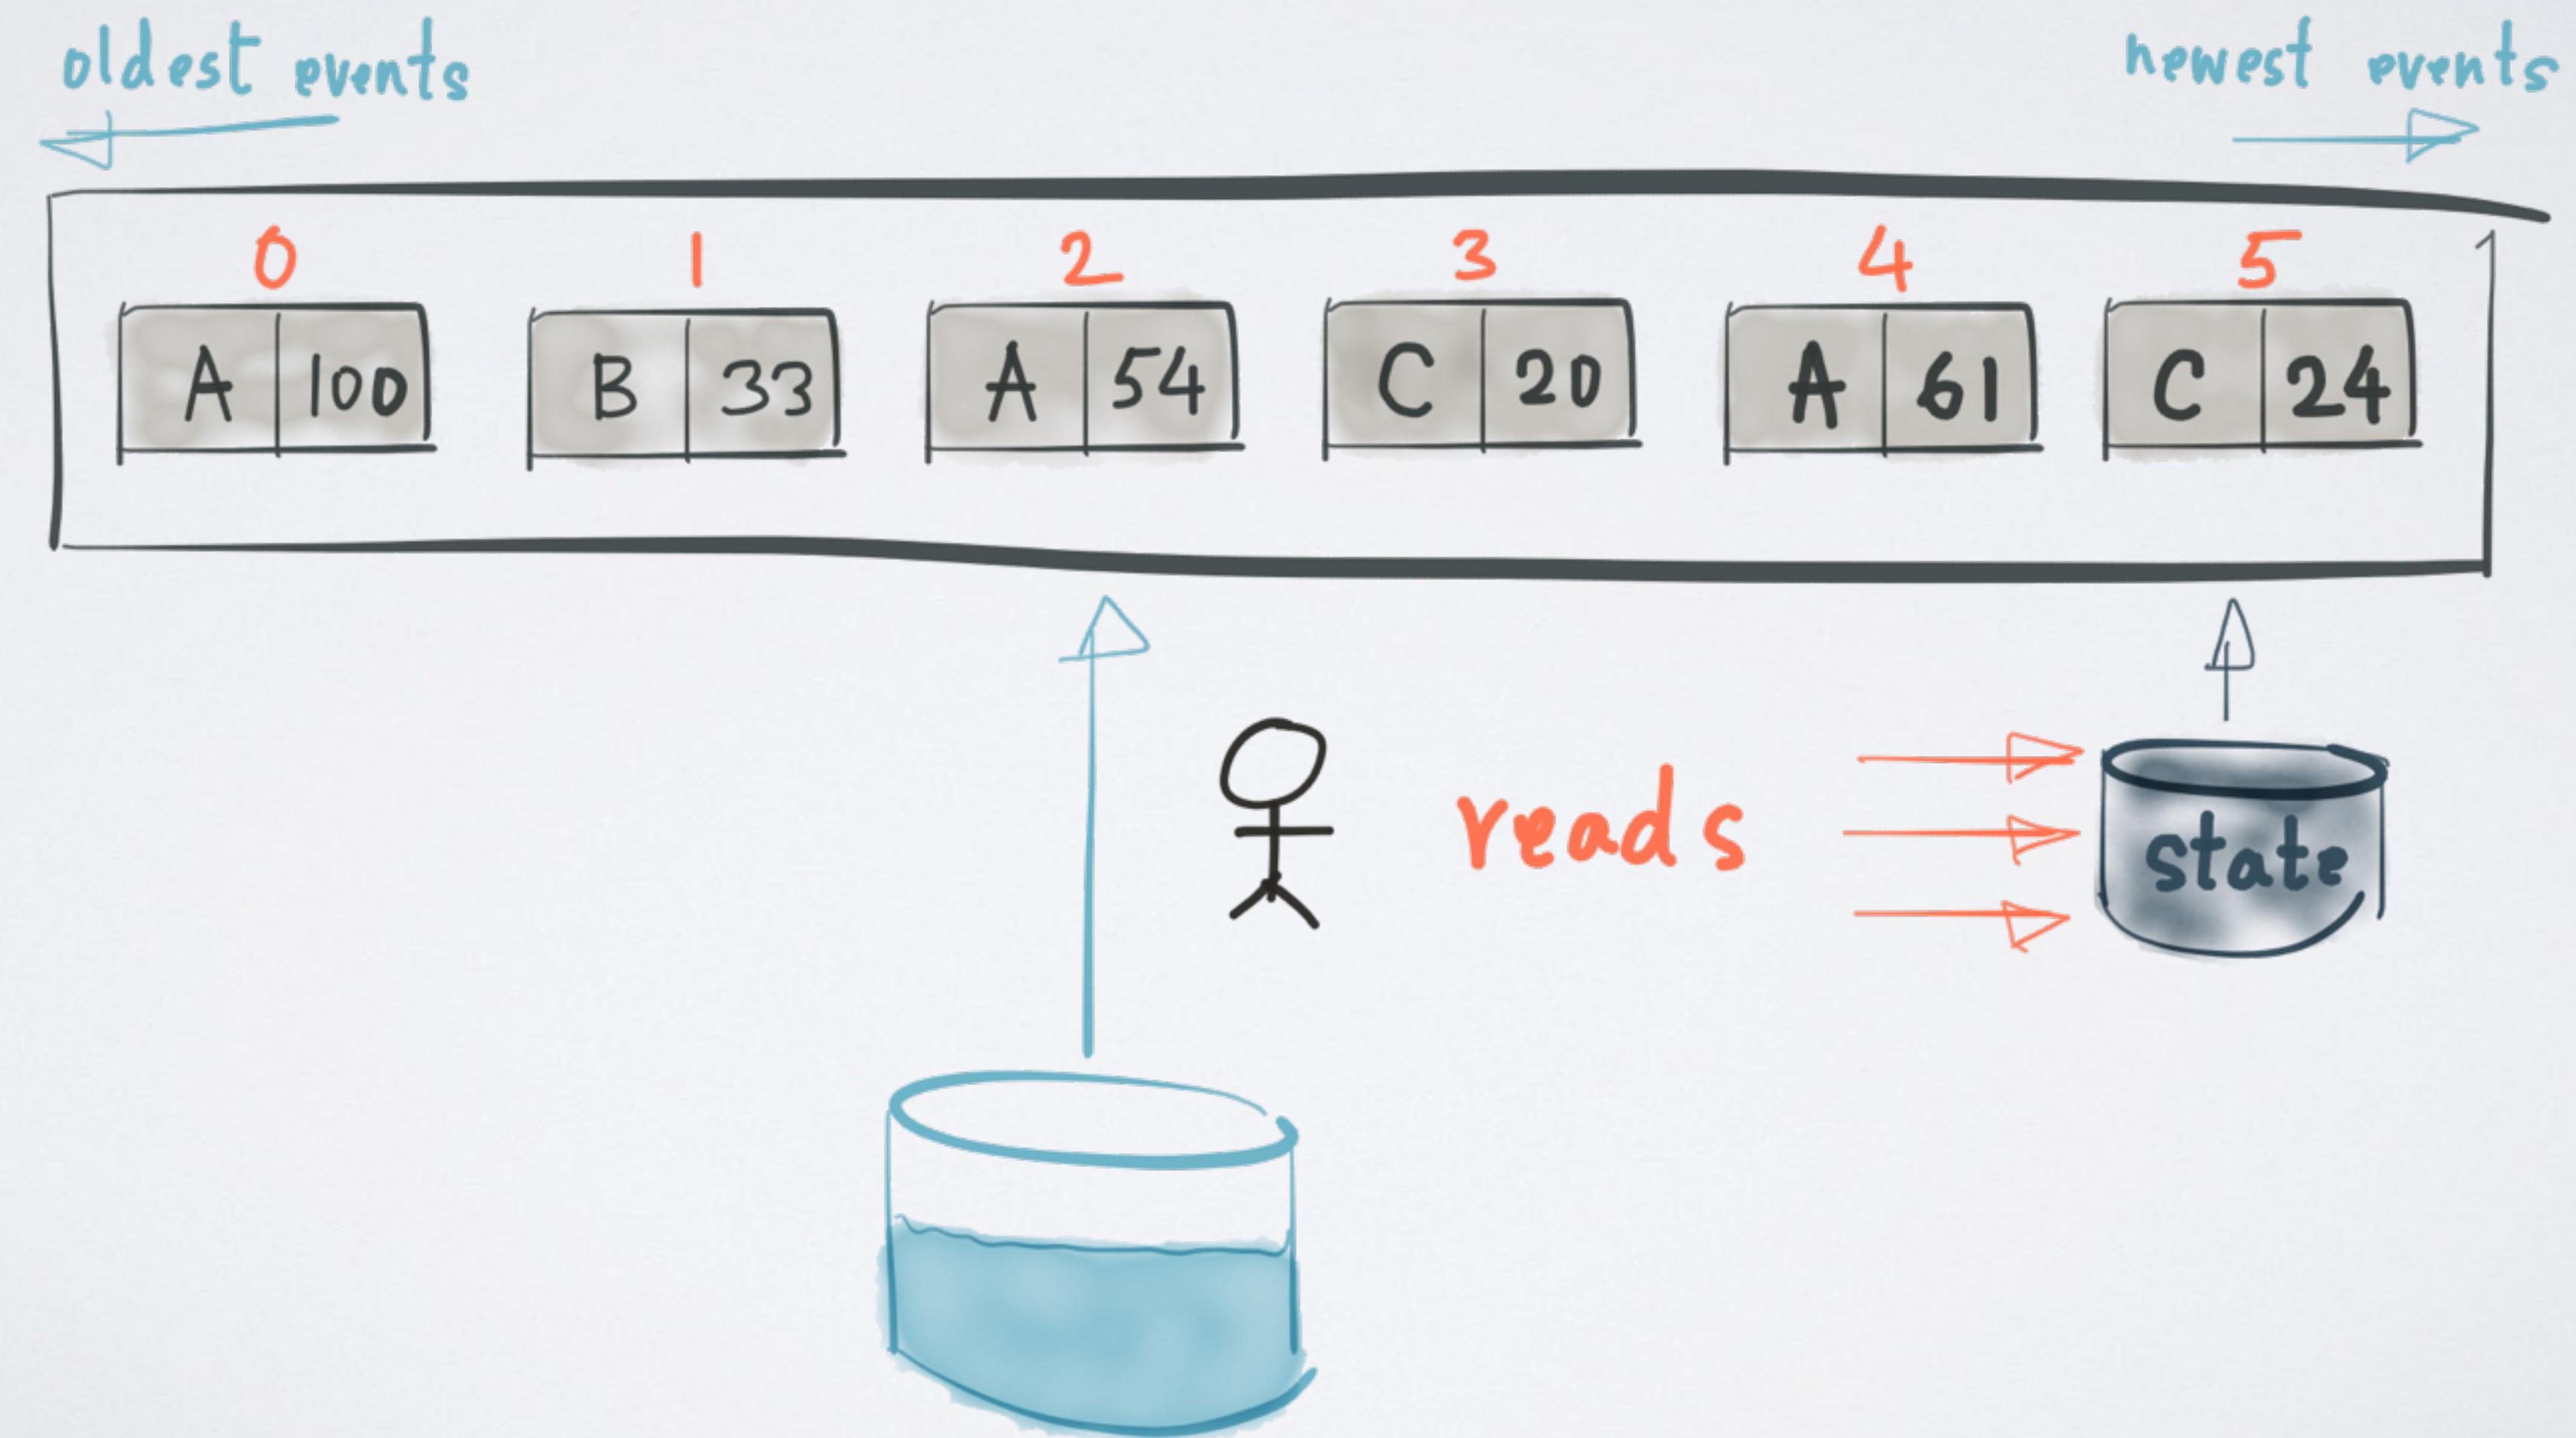


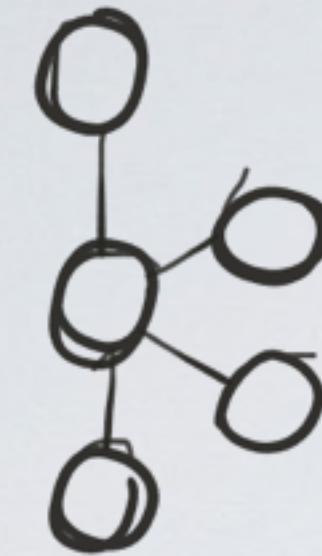
REPROCESSING



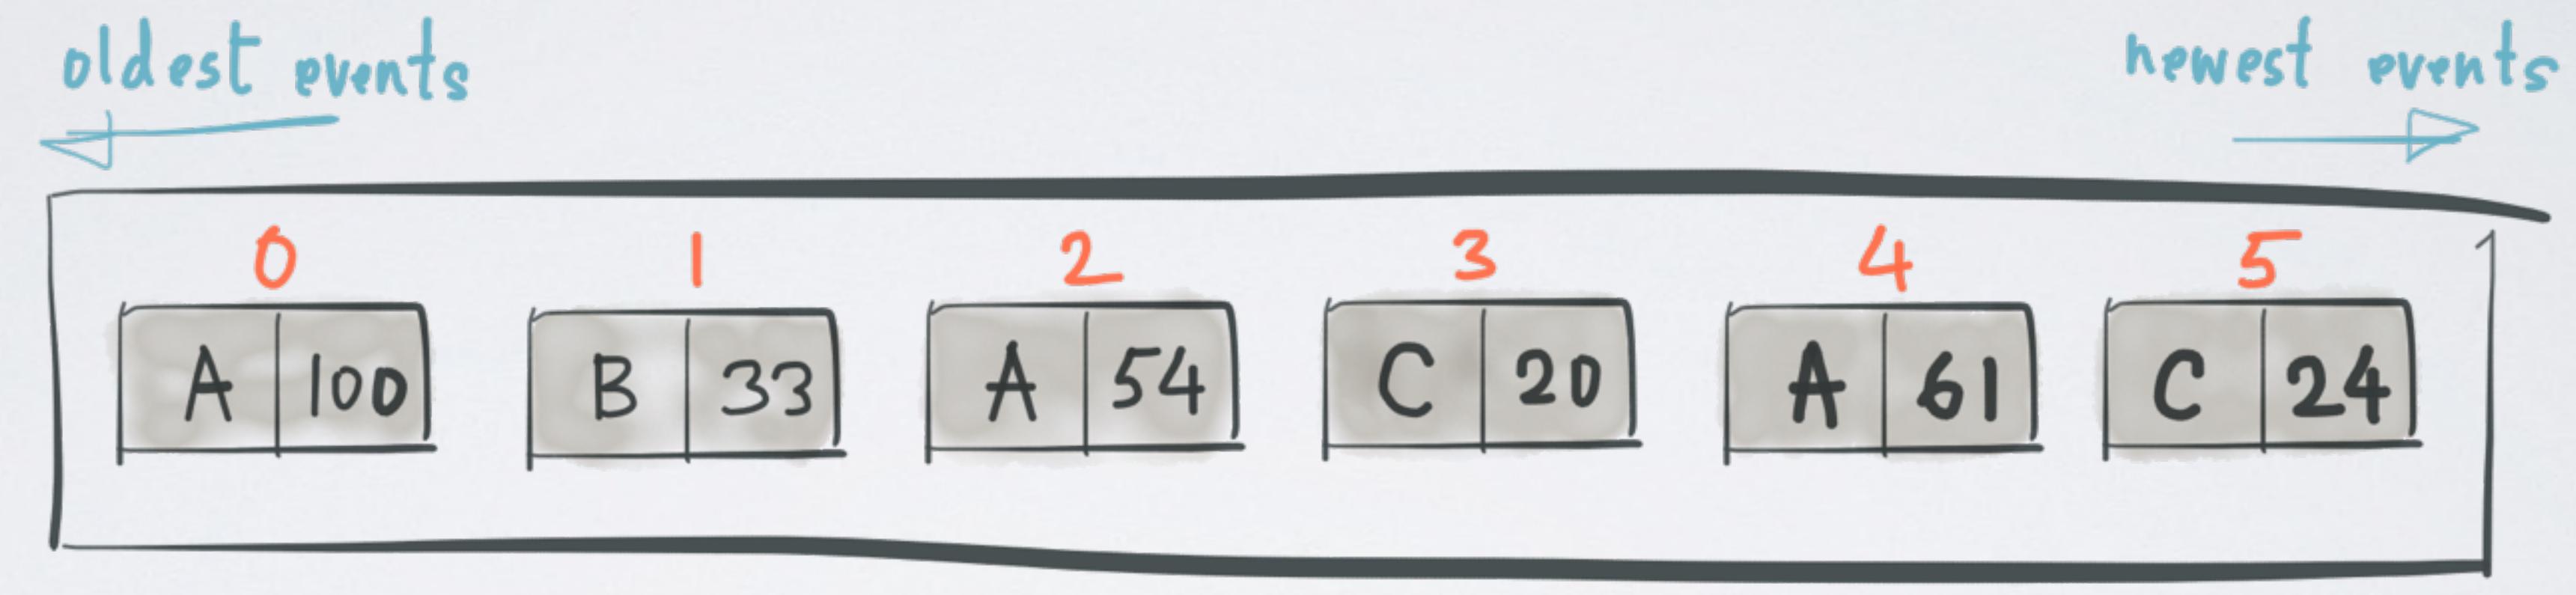


REPROCESSING

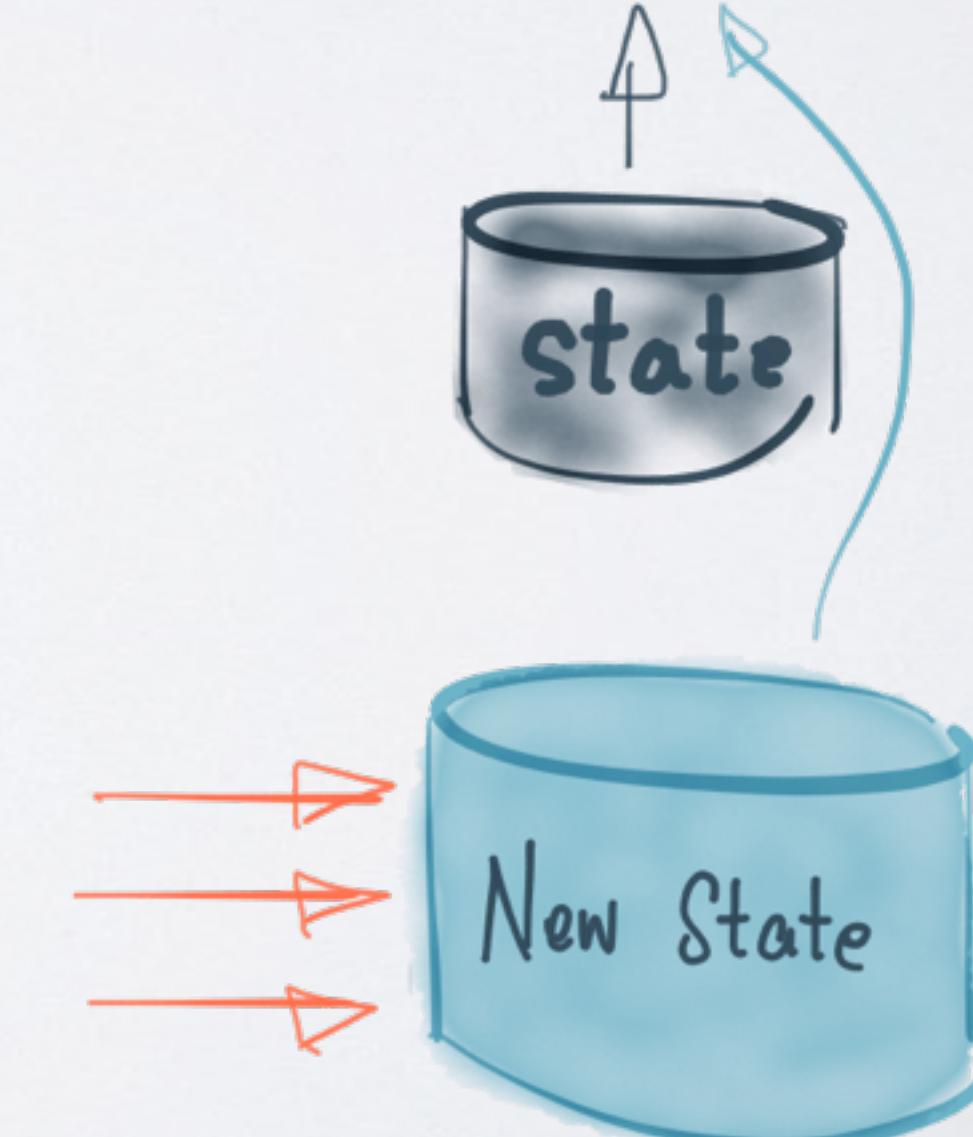




REPROCESSING



reads



It's all about *Time*

- ***Event-time*** (when an event is created)
- ***Processing-time*** (when an event is processed)



Out-of-Order



PHANTOM MENACE

ATTACK OF THE CLONES

REVENGE OF THE SITH

A NEW HOPE

THE EMPIRE STRIKES BACK

RETURN OF THE JEDI

THE FORCE AWAKENS



Event-time

1

2

3

4

5

6

7

Processing-time

1999

2002

2005

1977

1980

1983

2015

Timestamp Extractor

```
public long extract(ConsumerRecord<Object, Object> record) {  
    return System.currentTimeMillis();  
}
```

```
public long extract(ConsumerRecord<Object, Object> record) {  
    return record.timestamp();  
}
```

Timestamp Extractor

```
public long extract(ConsumerRecord<Object, Object> record) {  
    return System.currentTimeMillis();  
}
```

processing-time

```
public long extract(ConsumerRecord<Object, Object> record) {  
    return record.timestamp();  
}
```

Timestamp Extractor

```
public long extract(ConsumerRecord<Object, Object> record) {  
    return System.currentTimeMillis();  
}
```

processing-time

```
public long extract(ConsumerRecord<Object, Object> record) {  
    return record.timestamp();  
}
```

event-time

Timestamp Extractor

```
public long extract(ConsumerRecord<Object, Object> record) {  
    return System.currentTimeMillis();  
}
```

processing-time

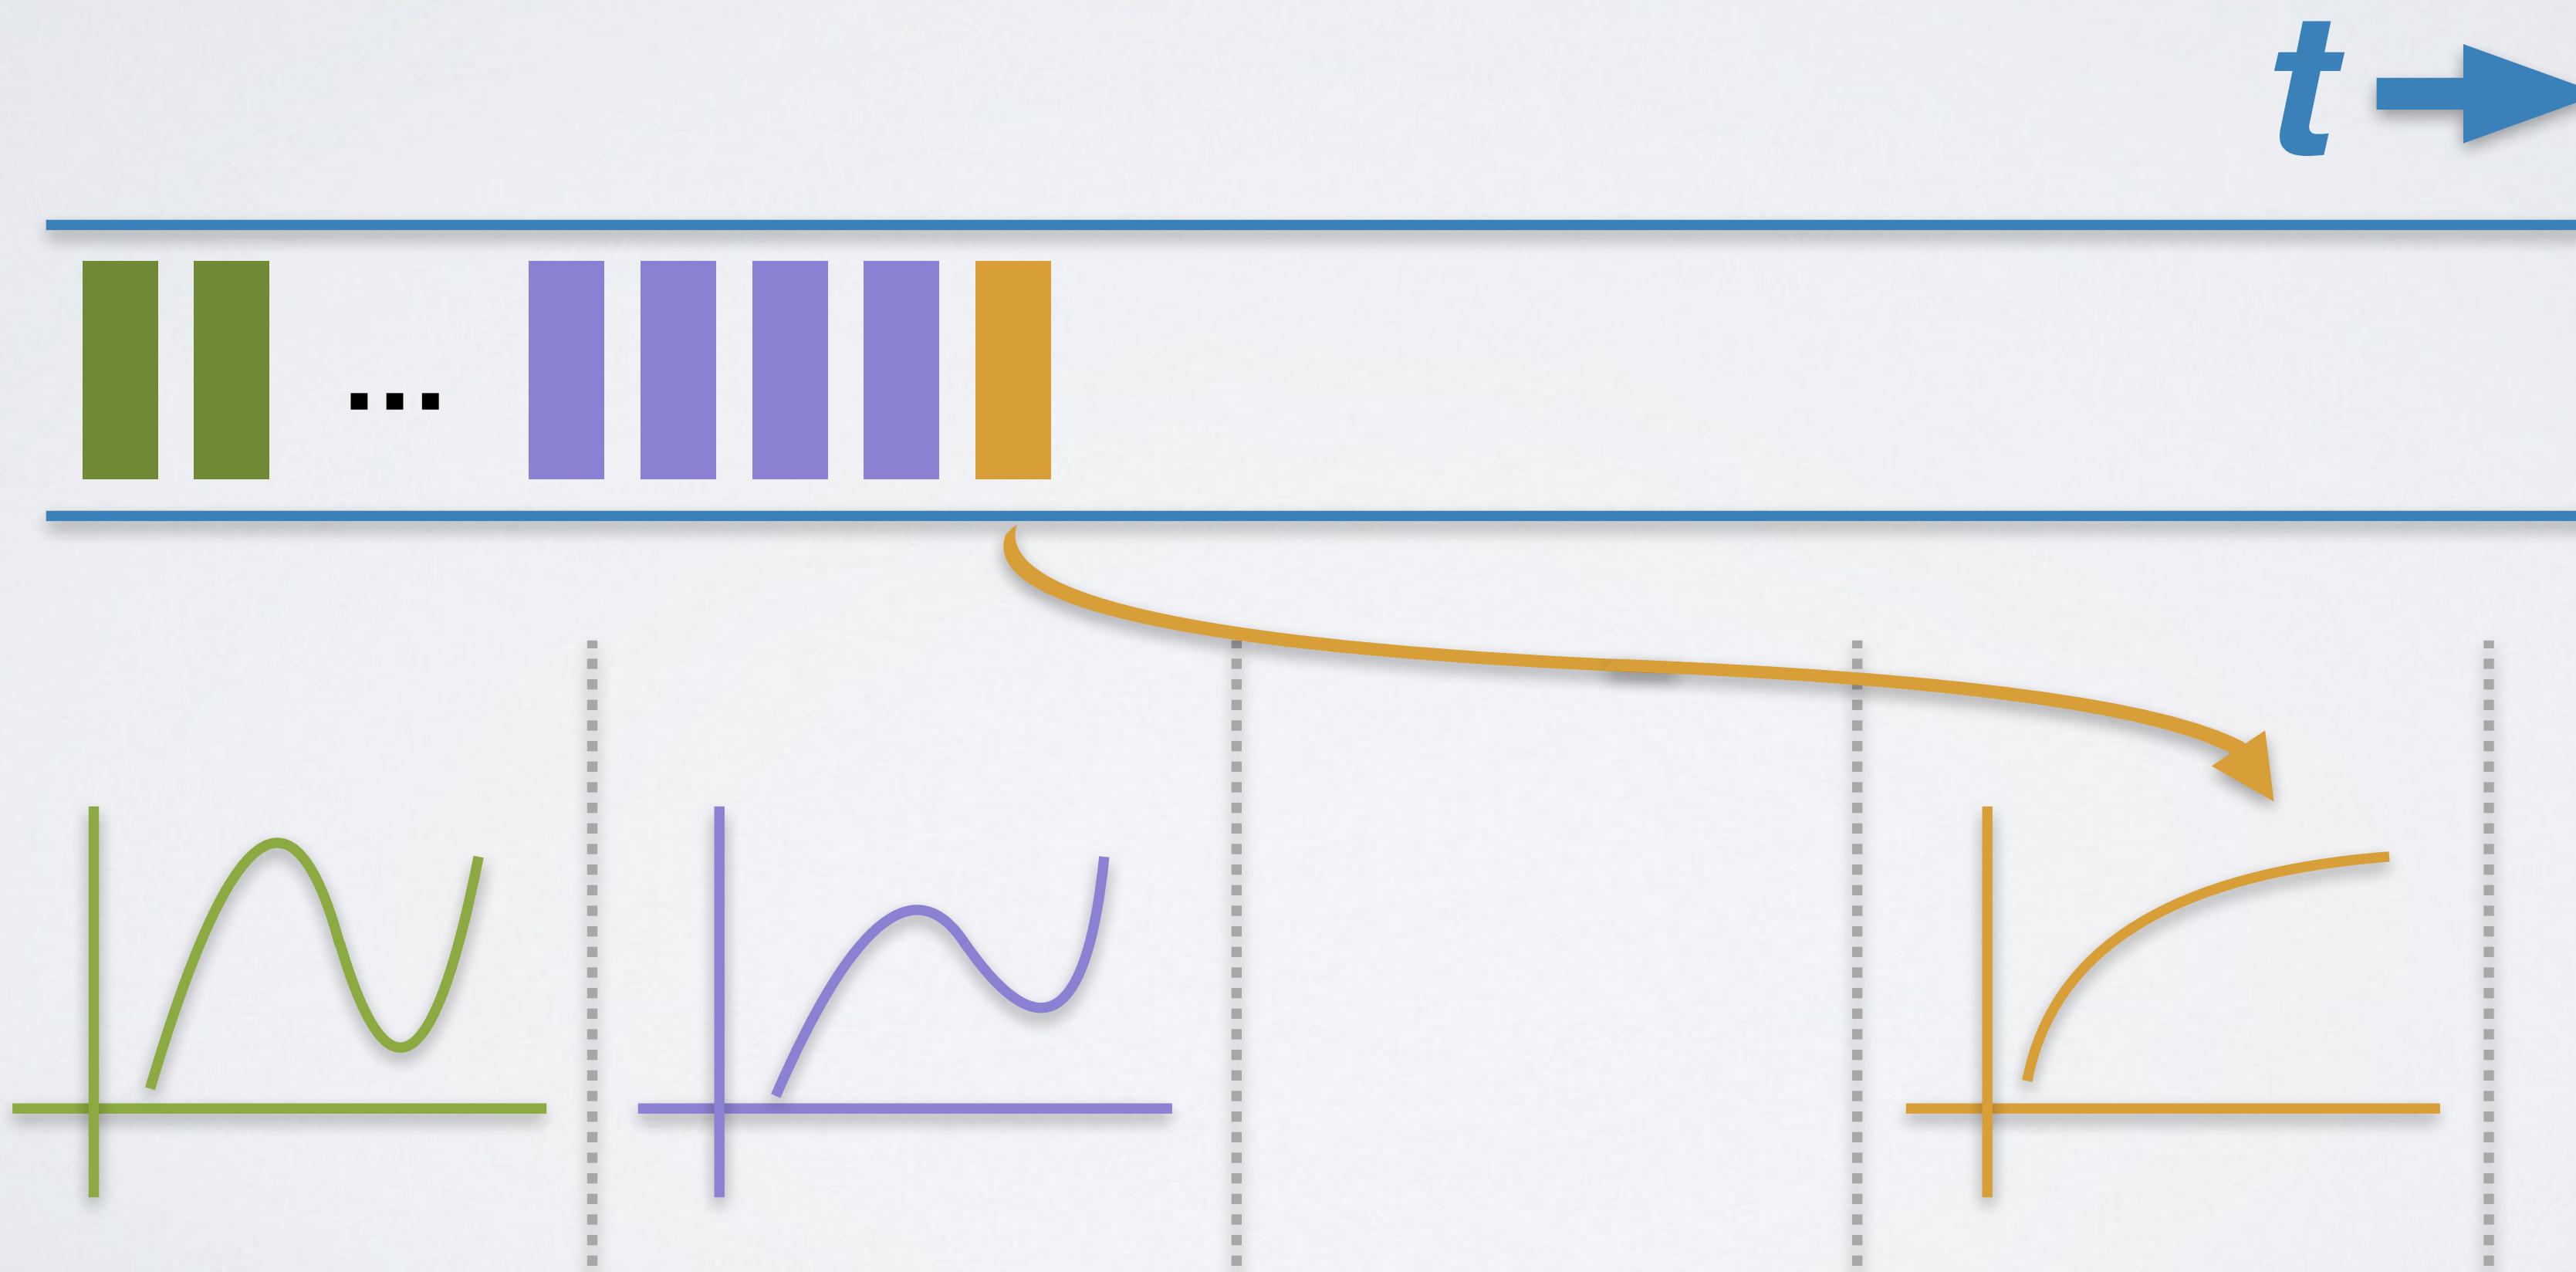
```
public long extract(ConsumerRecord<Object, Object> record) {  
    return ((JsonNode) record.value()).get("timestamp").longValue();  
}
```

event-time

Windowing



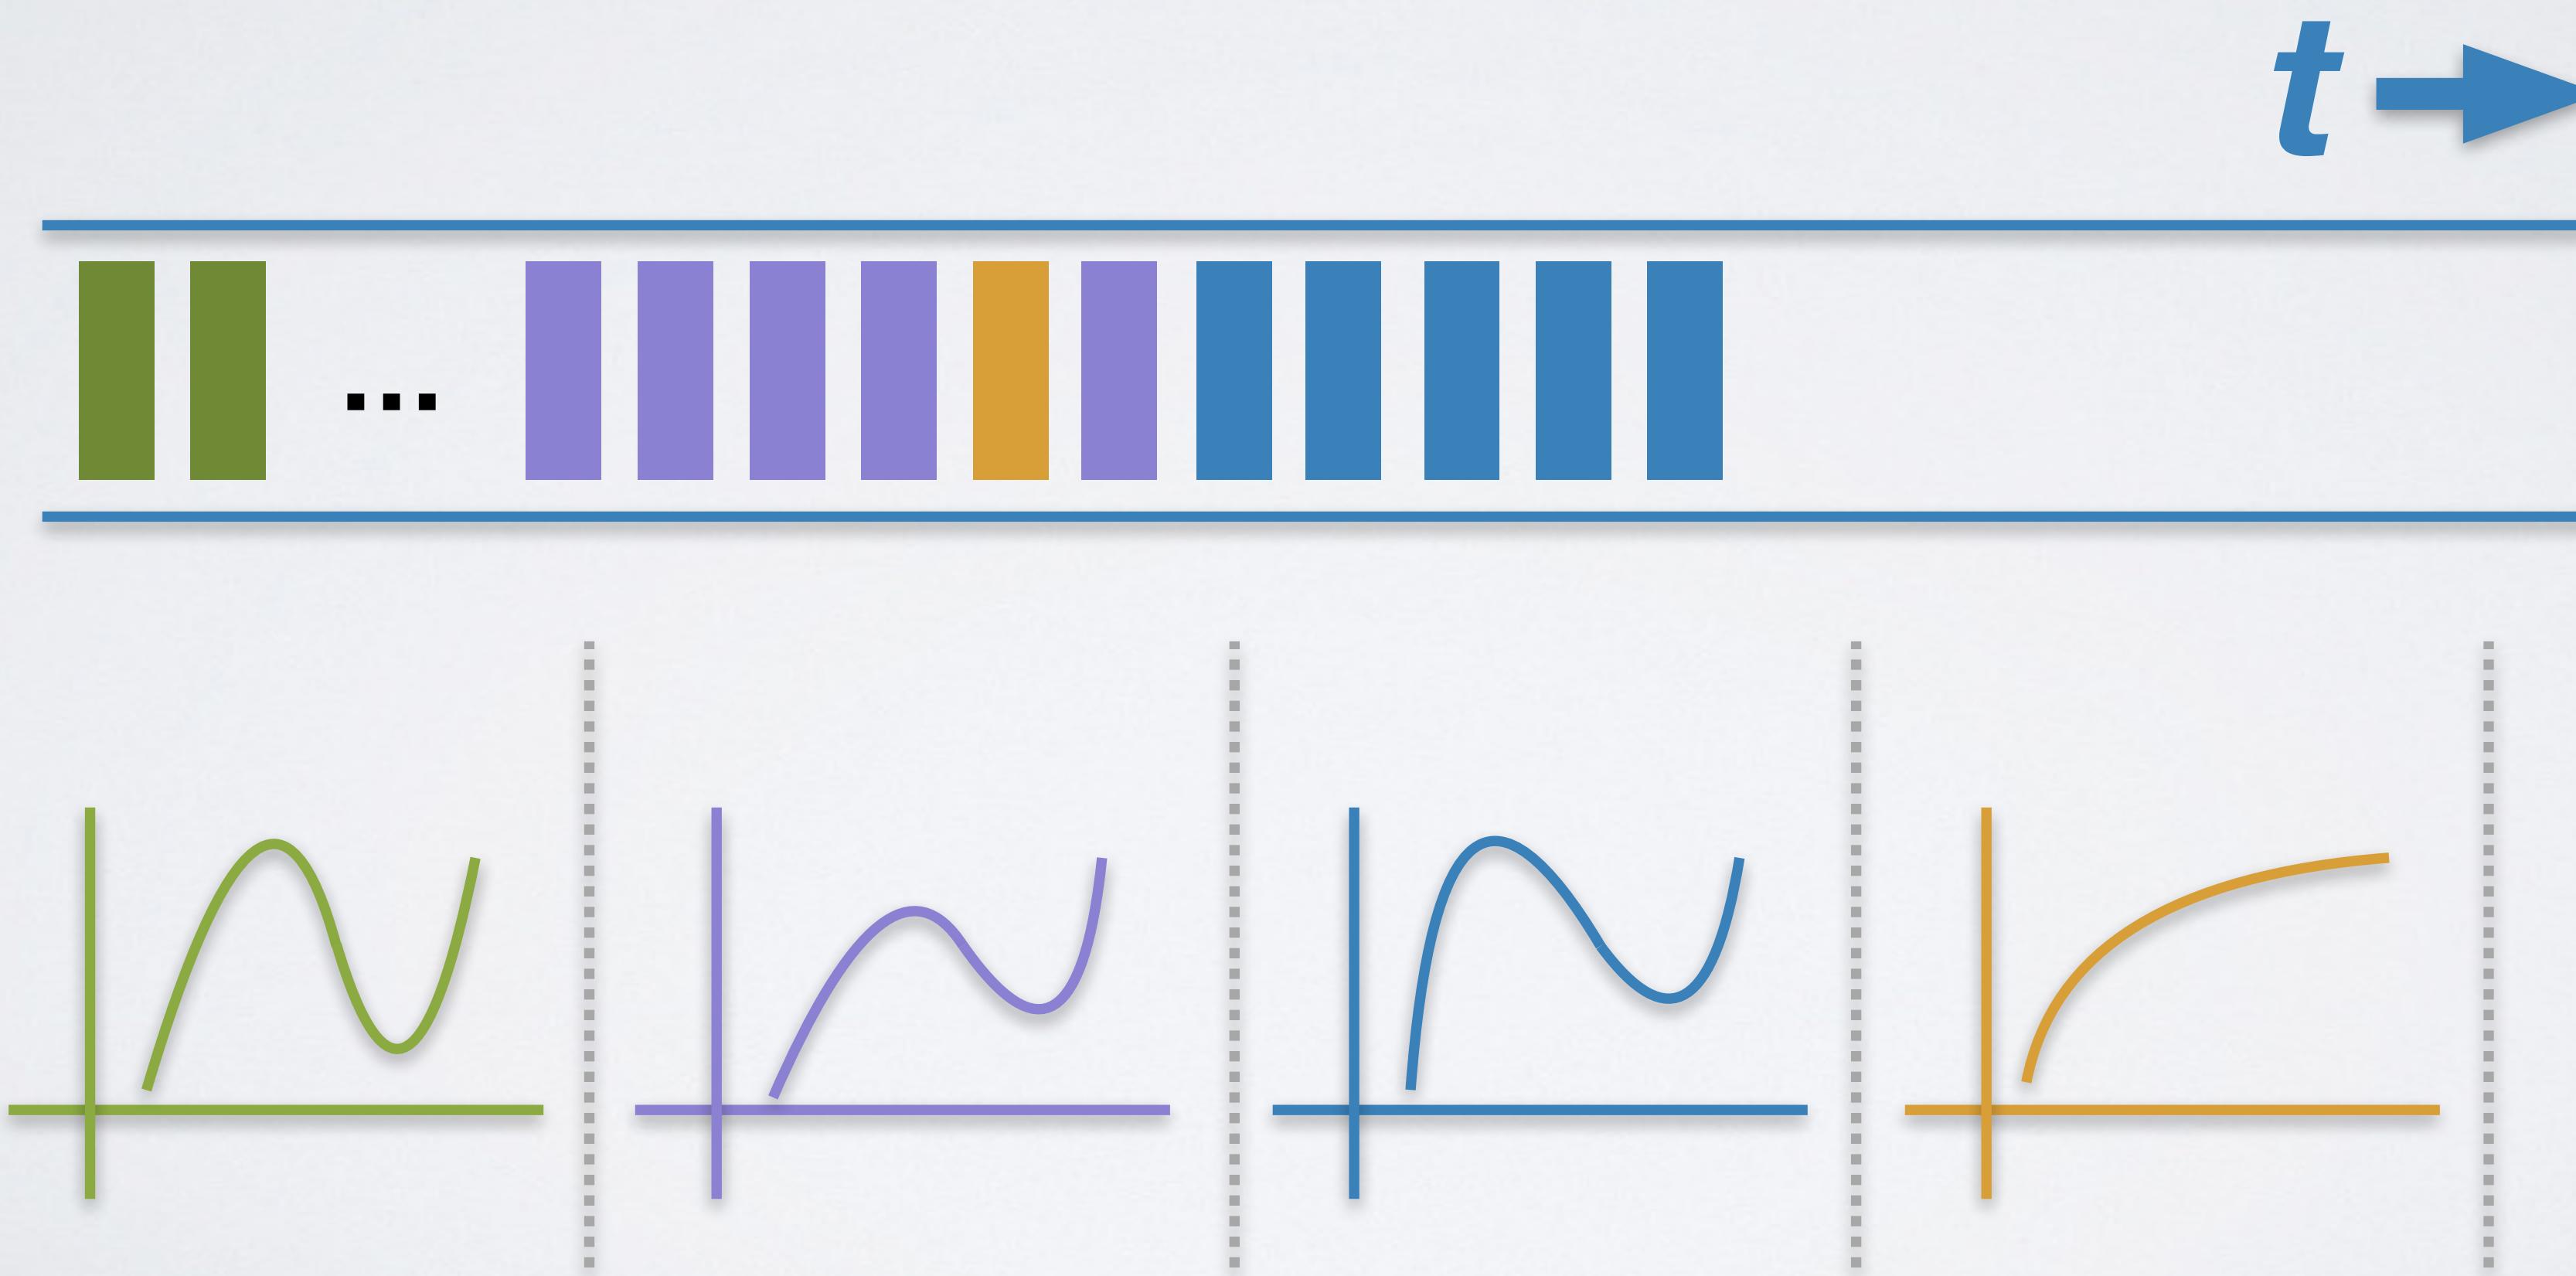
Windowing



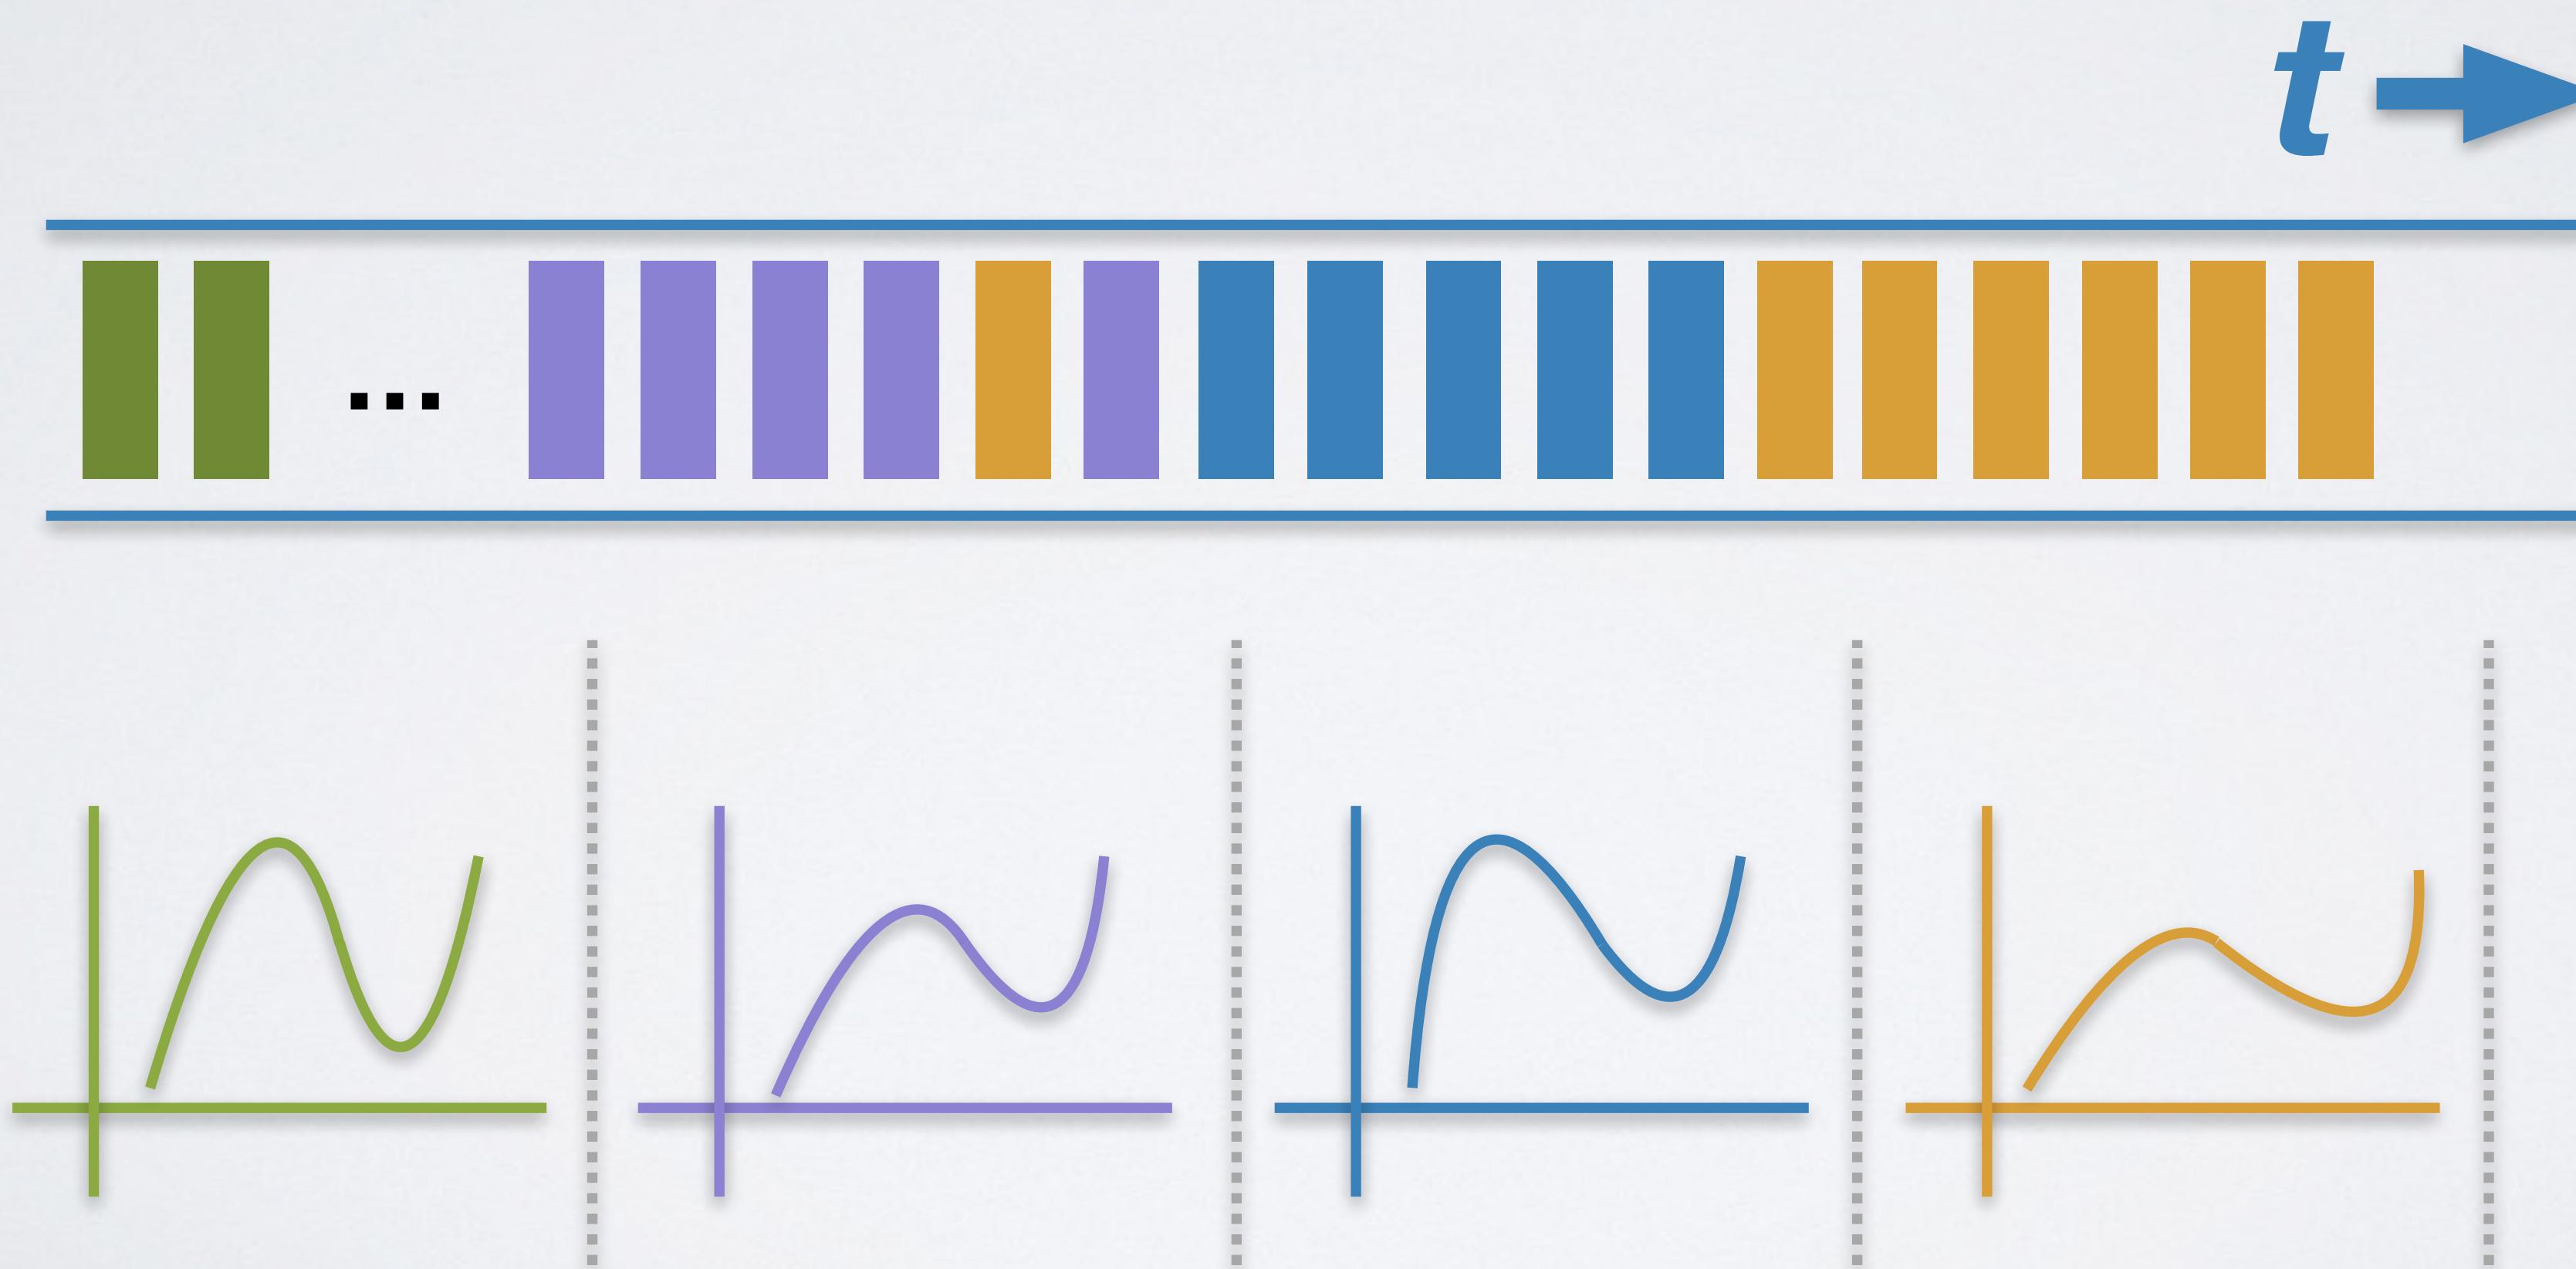
Windowing



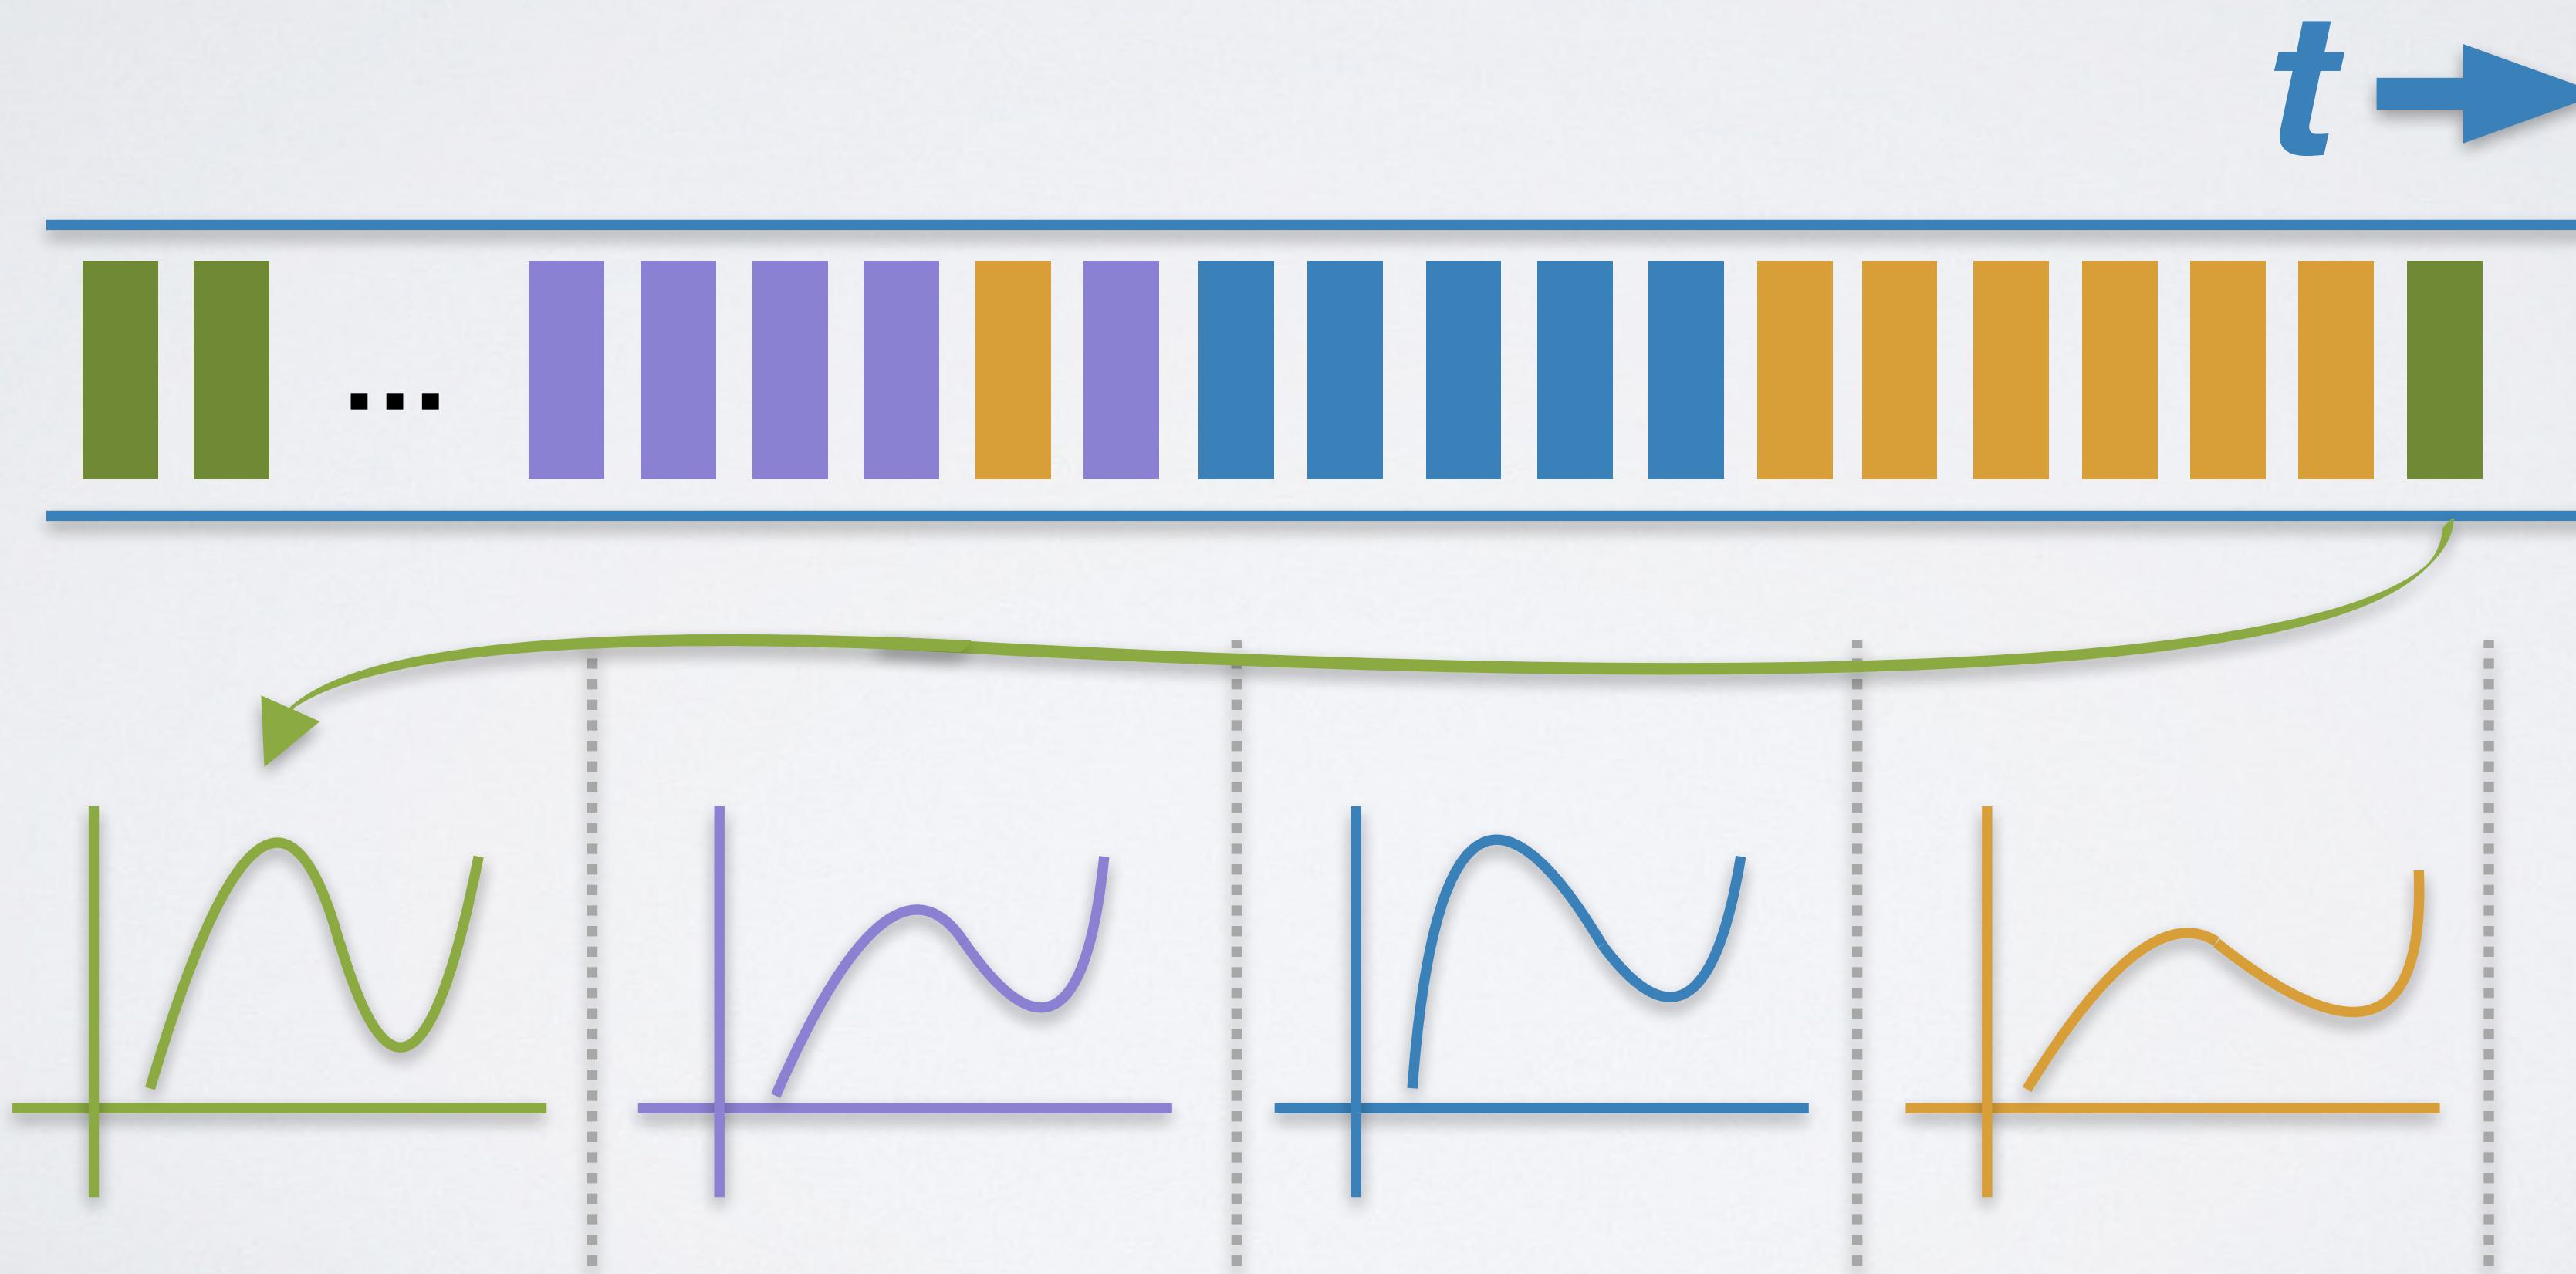
Windowing



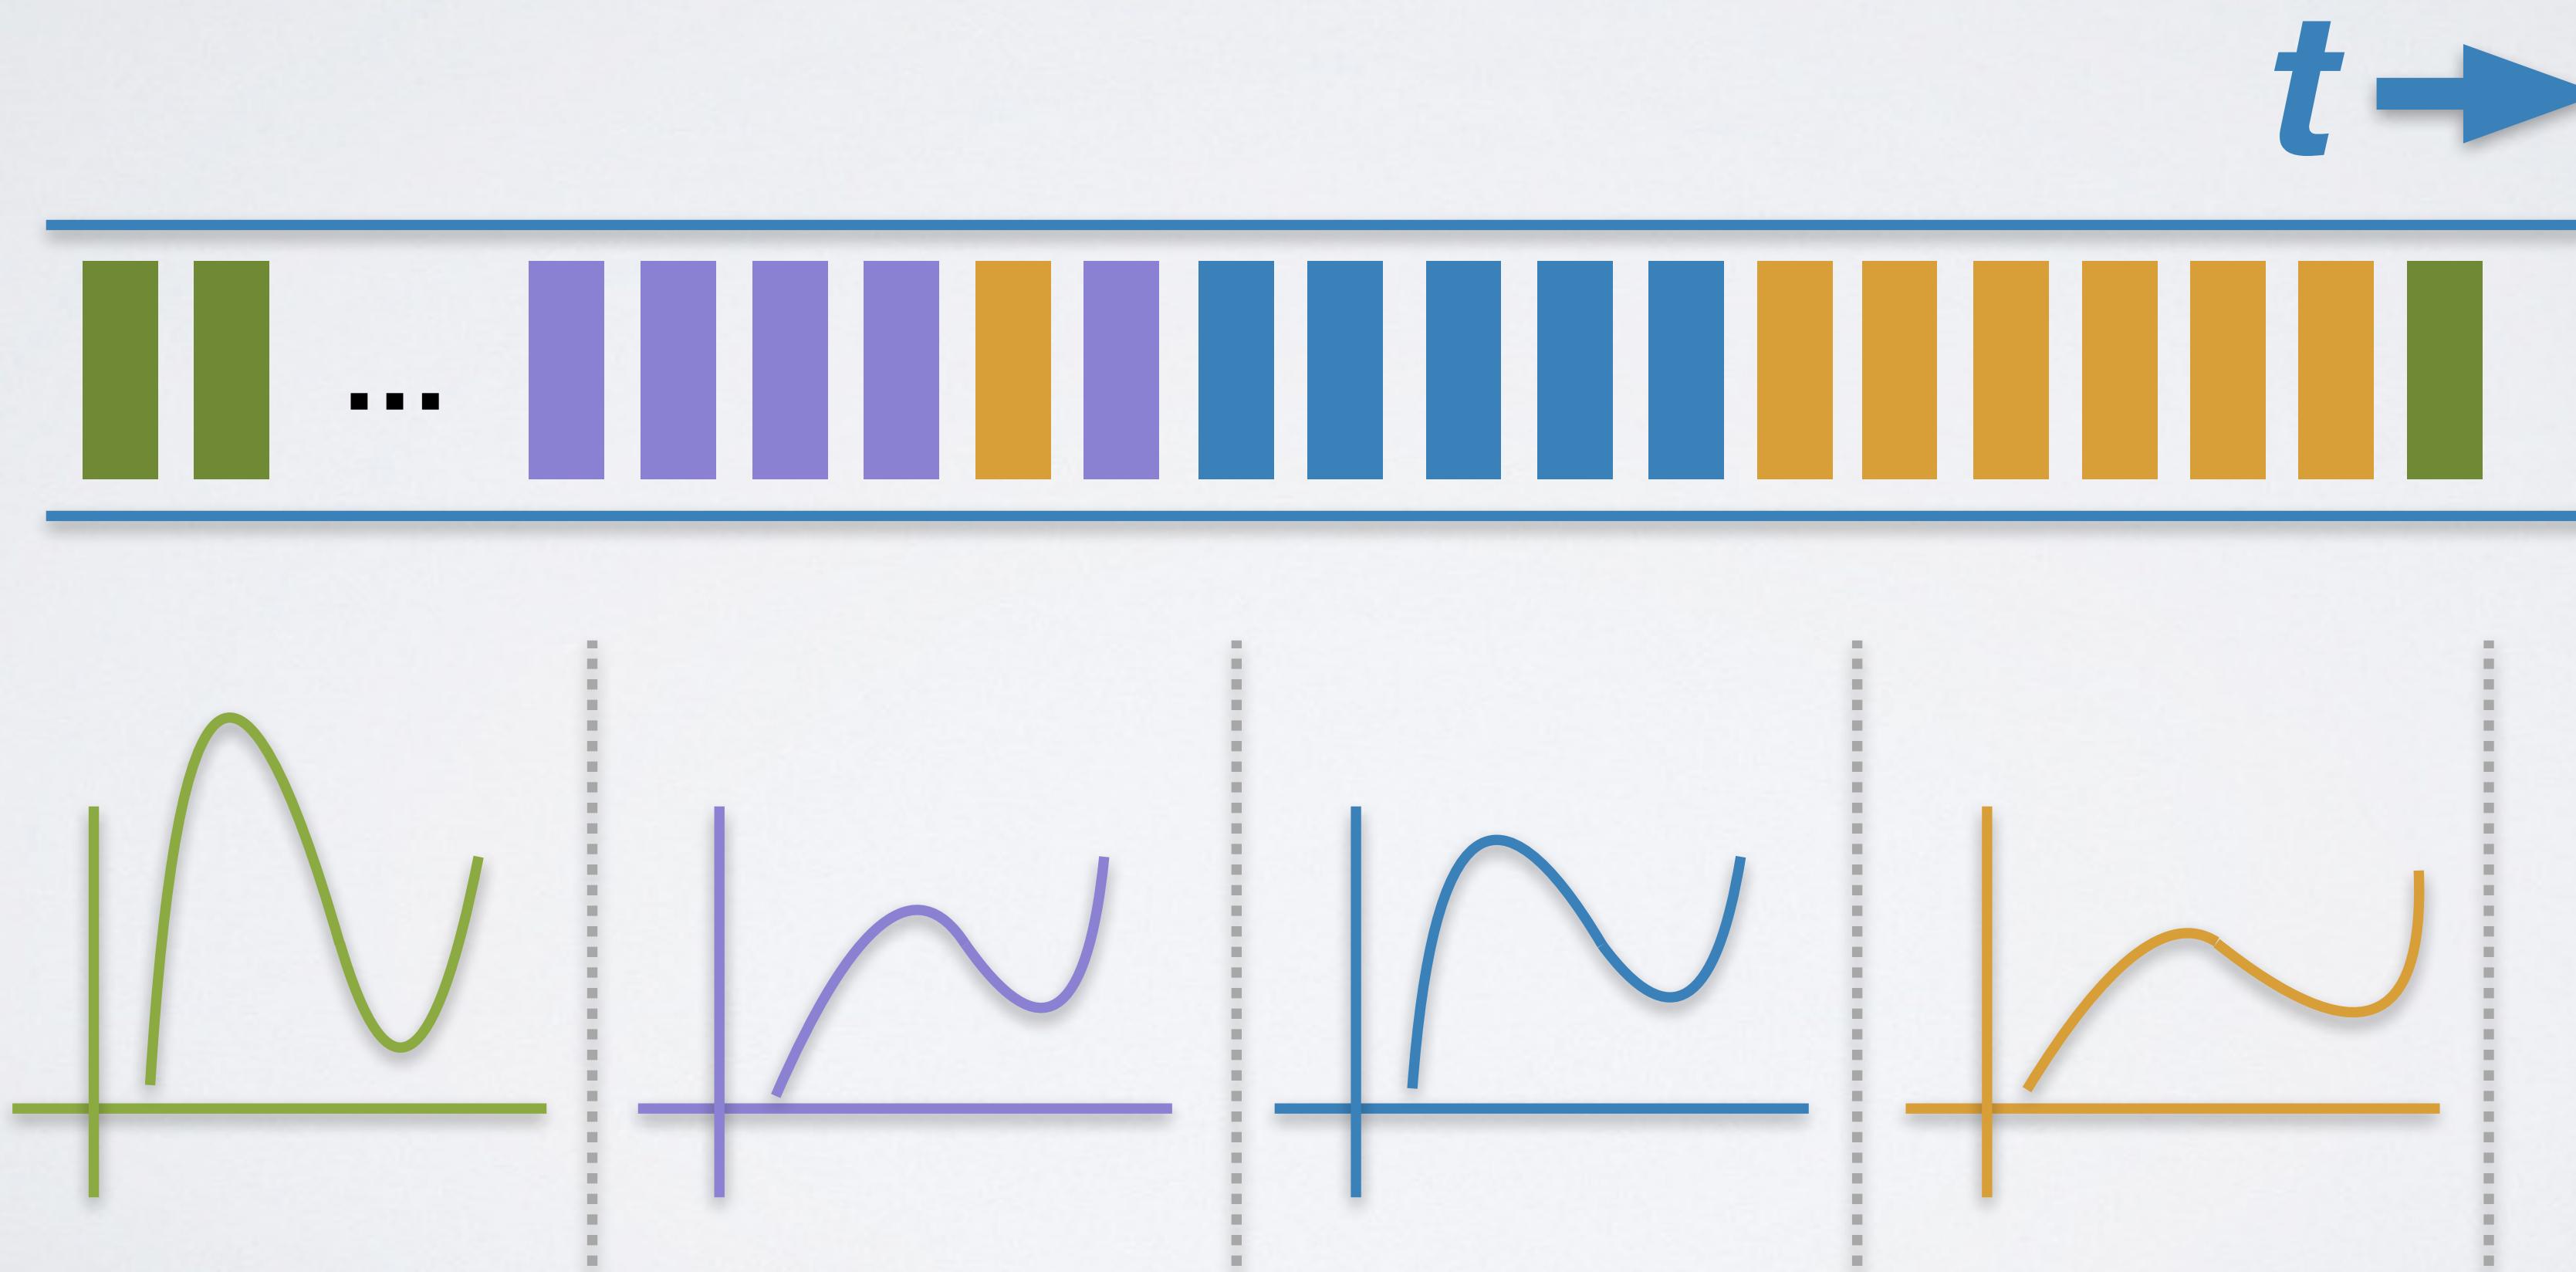
Windowing



Windowing



Windowing



Stream Processing *Hard Parts*

- *Ordering* ✓
- *Partitioning & Scalability* ✓
- *Fault tolerance* ✓
- *State Management* ✓
- *Time, Window & Out-of-order Data* ✓
- *Re-processing* ✓

Stream Processing *Hard Parts*

- *Ordering* ✓
- *State Management* ✓

Simple is Beautiful

- *Fault tolerance* ✓
- *Re-processing* ✓

Ongoing Work (0.10+)

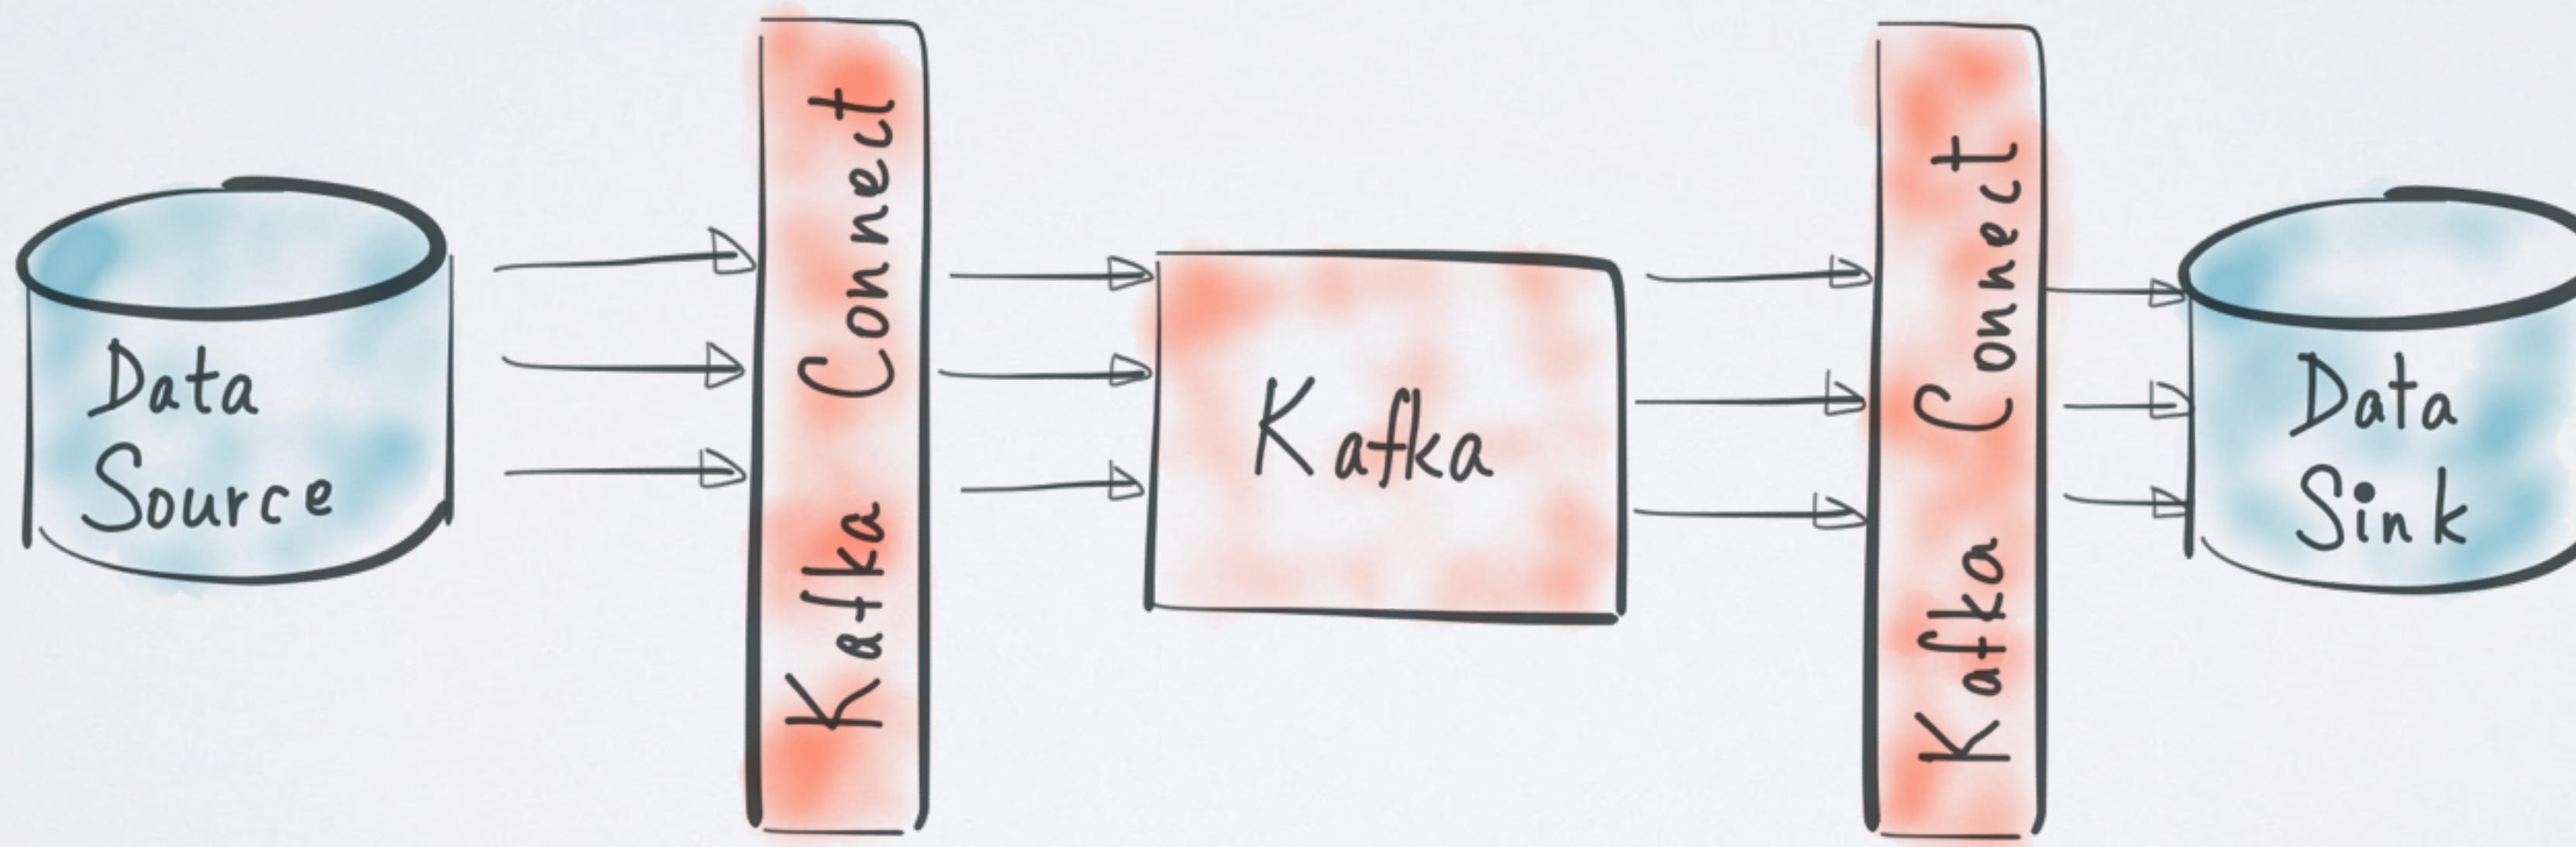
- *Beyond Java APIs*
- *Streaming SQL, Python client, etc*
- *Exactly-Once (end-to-end Semantics)*
- *... and more*



*But how to get data in / out **Kafka**?*



KAFKA CONNECT



Connectors

- 40+ since first release this Feb (0.9+)

- 13 from  & partners

Kafka Connect

Kafka Connect is a framework included in Apache Kafka that integrates Kafka with other systems. Its purpose is to make it easy to add new systems to your scalable and secure stream data pipelines.

To copy data between Kafka and another system, users instantiate Kafka Connectors for the systems they want to pull data from or push data to. **Source Connectors** import data from another system (e.g. a relational database into Kafka) and **Sink Connectors** export data (e.g. the contents of a Kafka topic to an HDFS file).

This page lists many of the notable connectors available.

Certified Connectors

Certified Connectors have been developed by vendors and/or Confluent utilizing the Kafka Connect framework. These Connectors have met criteria for code development best practices, schema registry integration, security, and documentation.

CONNECTOR	TAGS	DEVELOPER/SUPPORT	DOWNLOAD
HDFS (Sink)	HDFS, Hadoop, Hive	Confluent	Confluent
JDBC (Source)	JDBC, MySQL	Confluent	Confluent
Elasticsearch (Sink)	search, Elastic, log, analytics	Confluent	Confluent
DataStax (Sink)	Cassandra, DataStax	Data Mountaineer	Data Mountaineer
Attunity (Source)	CDC	Attunity	Attunity
Couchbase (Source)	Couchbase, NoSQL	Couchbase	Couchbase
GoldenGate (Source)	CDC, Oracle	Oracle	Community
JustOne (Sink)	Postgres	JustOne	JustOne
Striim (Source)	CDC, MS SQLServer, Oracle, MySQL	Striim	Striim
Syncsort DMX (Source)	DB2, IMS, VSAM, CICS	Syncsort	Syncsort
Syncsort DMX (Sink)	DB2, IMS, VSAM, CICS	Syncsort	Syncsort
Vertica (Source)	Vertica	HP Enterprise	HP Enterprise
Vertica (Sink)	Vertica	HP Enterprise	HP Enterprise

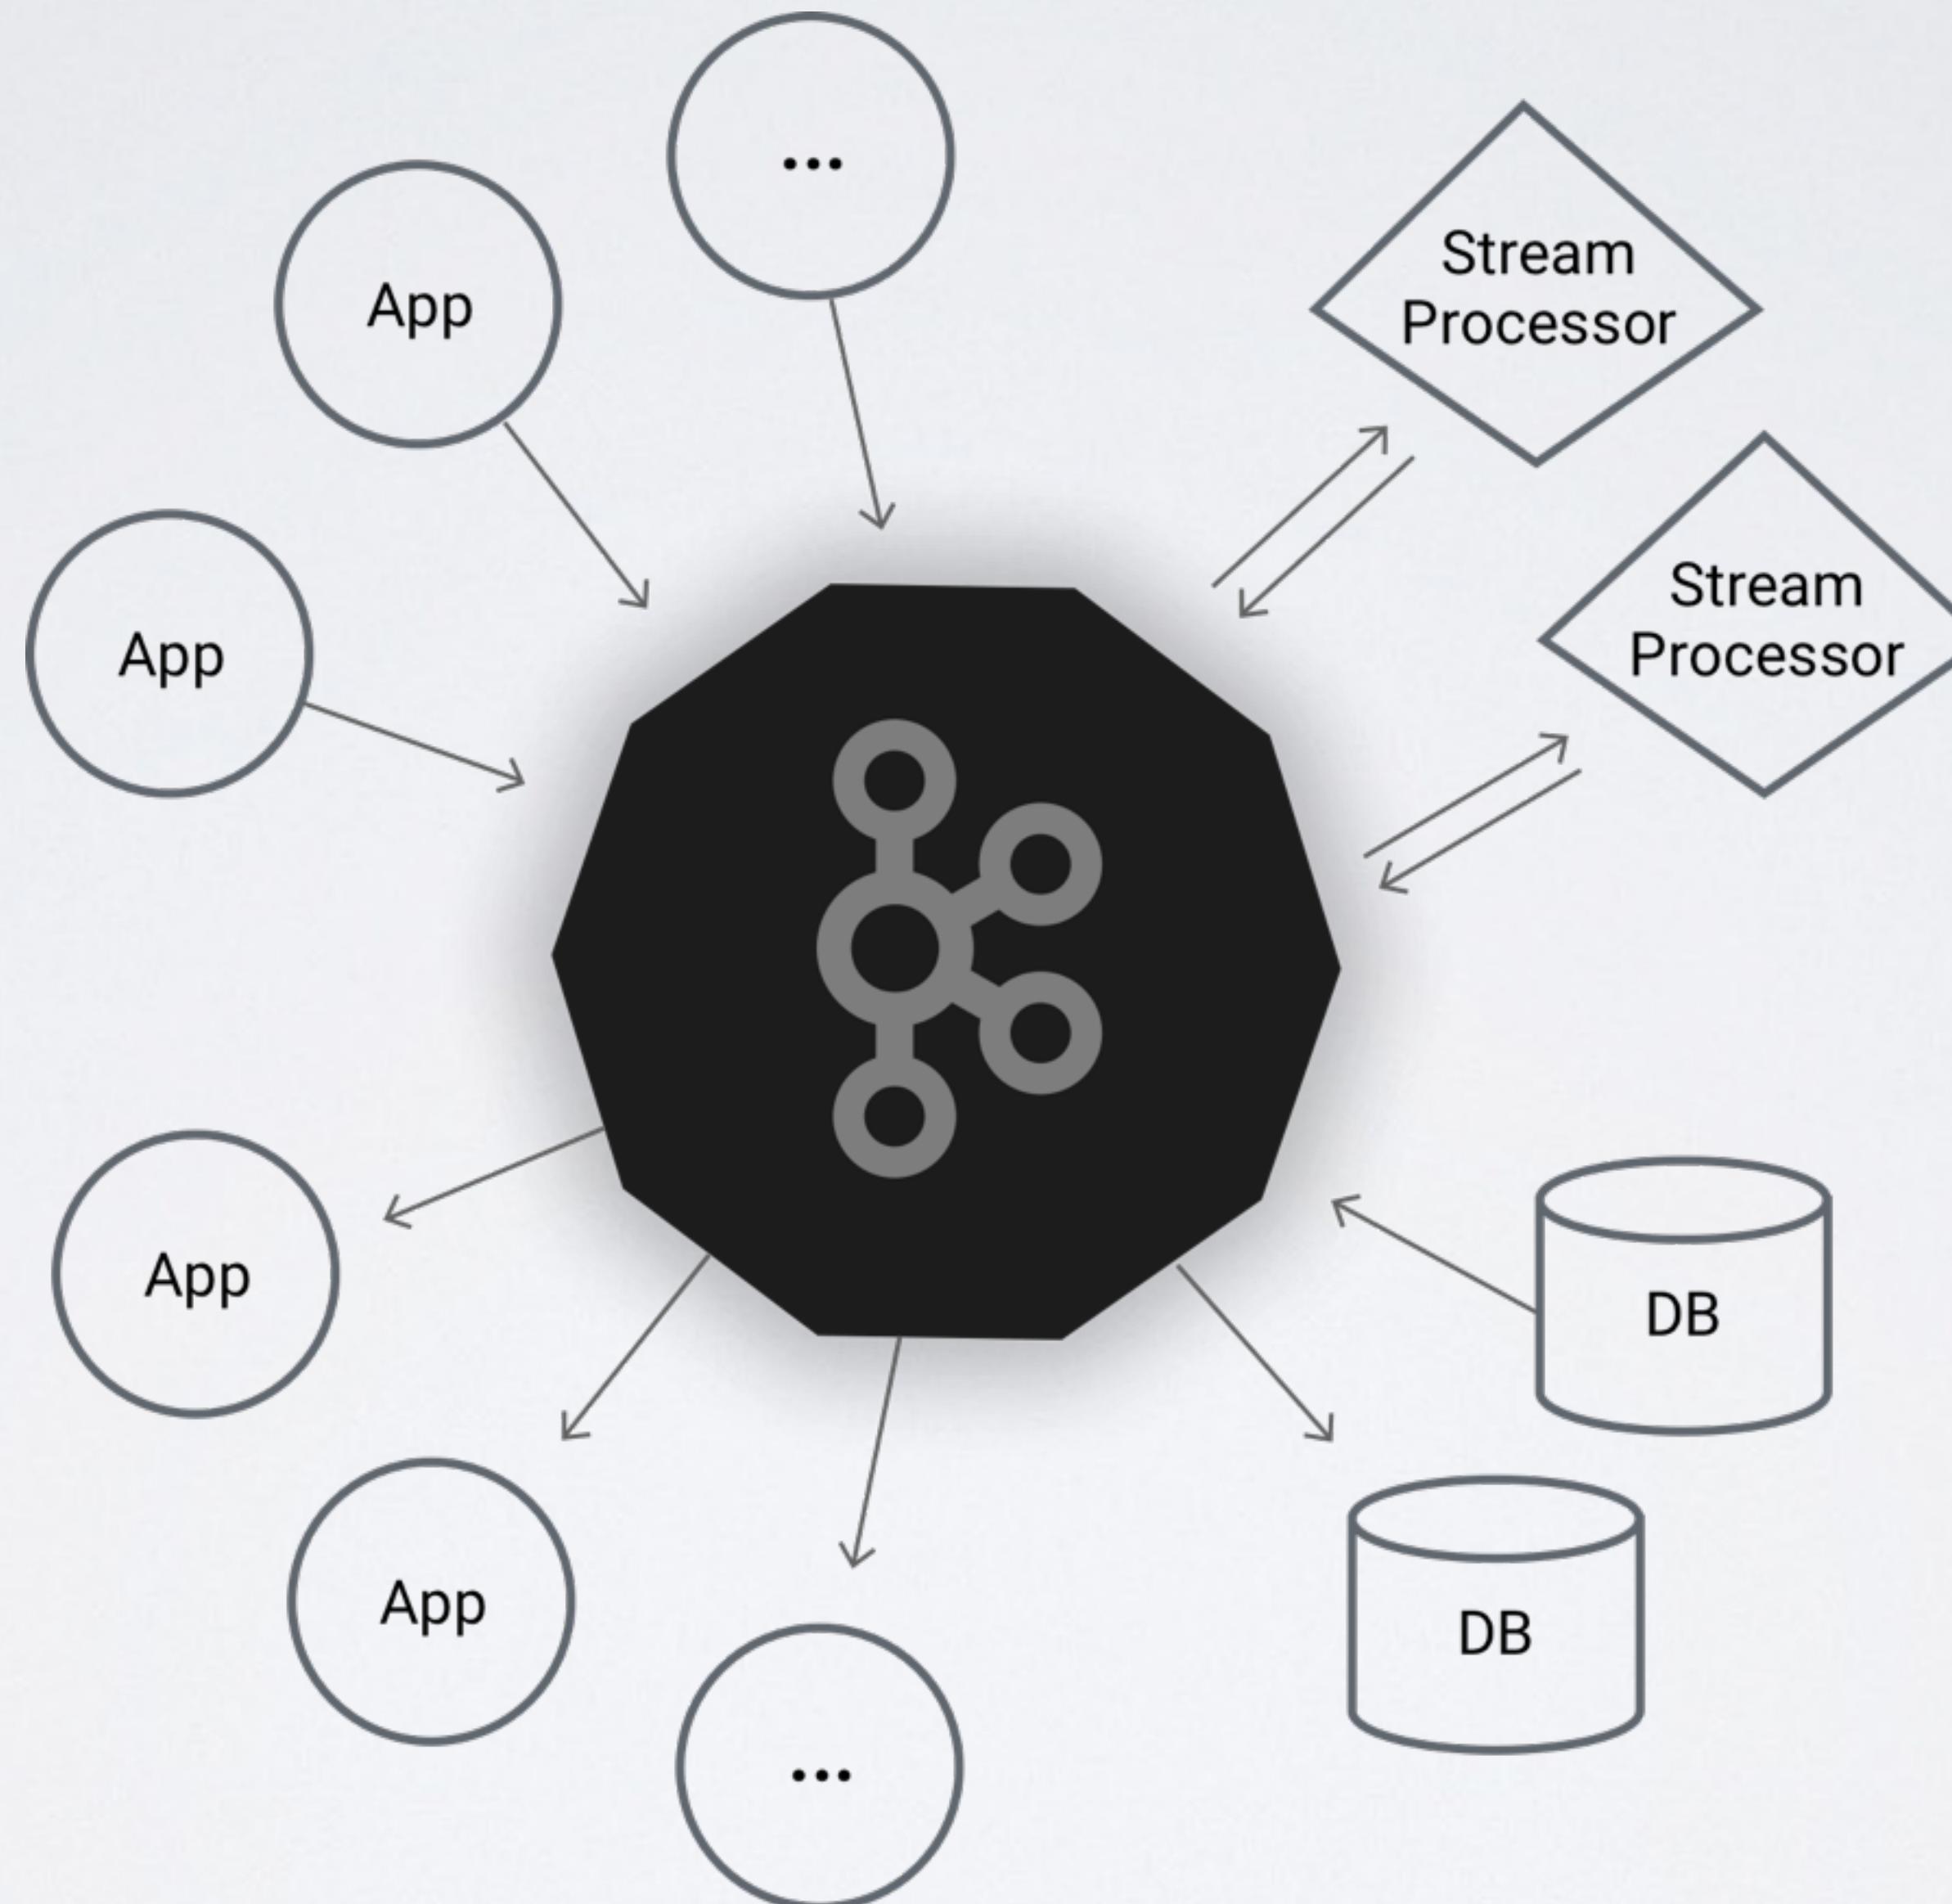
Additional Connectors Available

Other notable Connectors that have been developed utilizing the Kafka Connect framework.

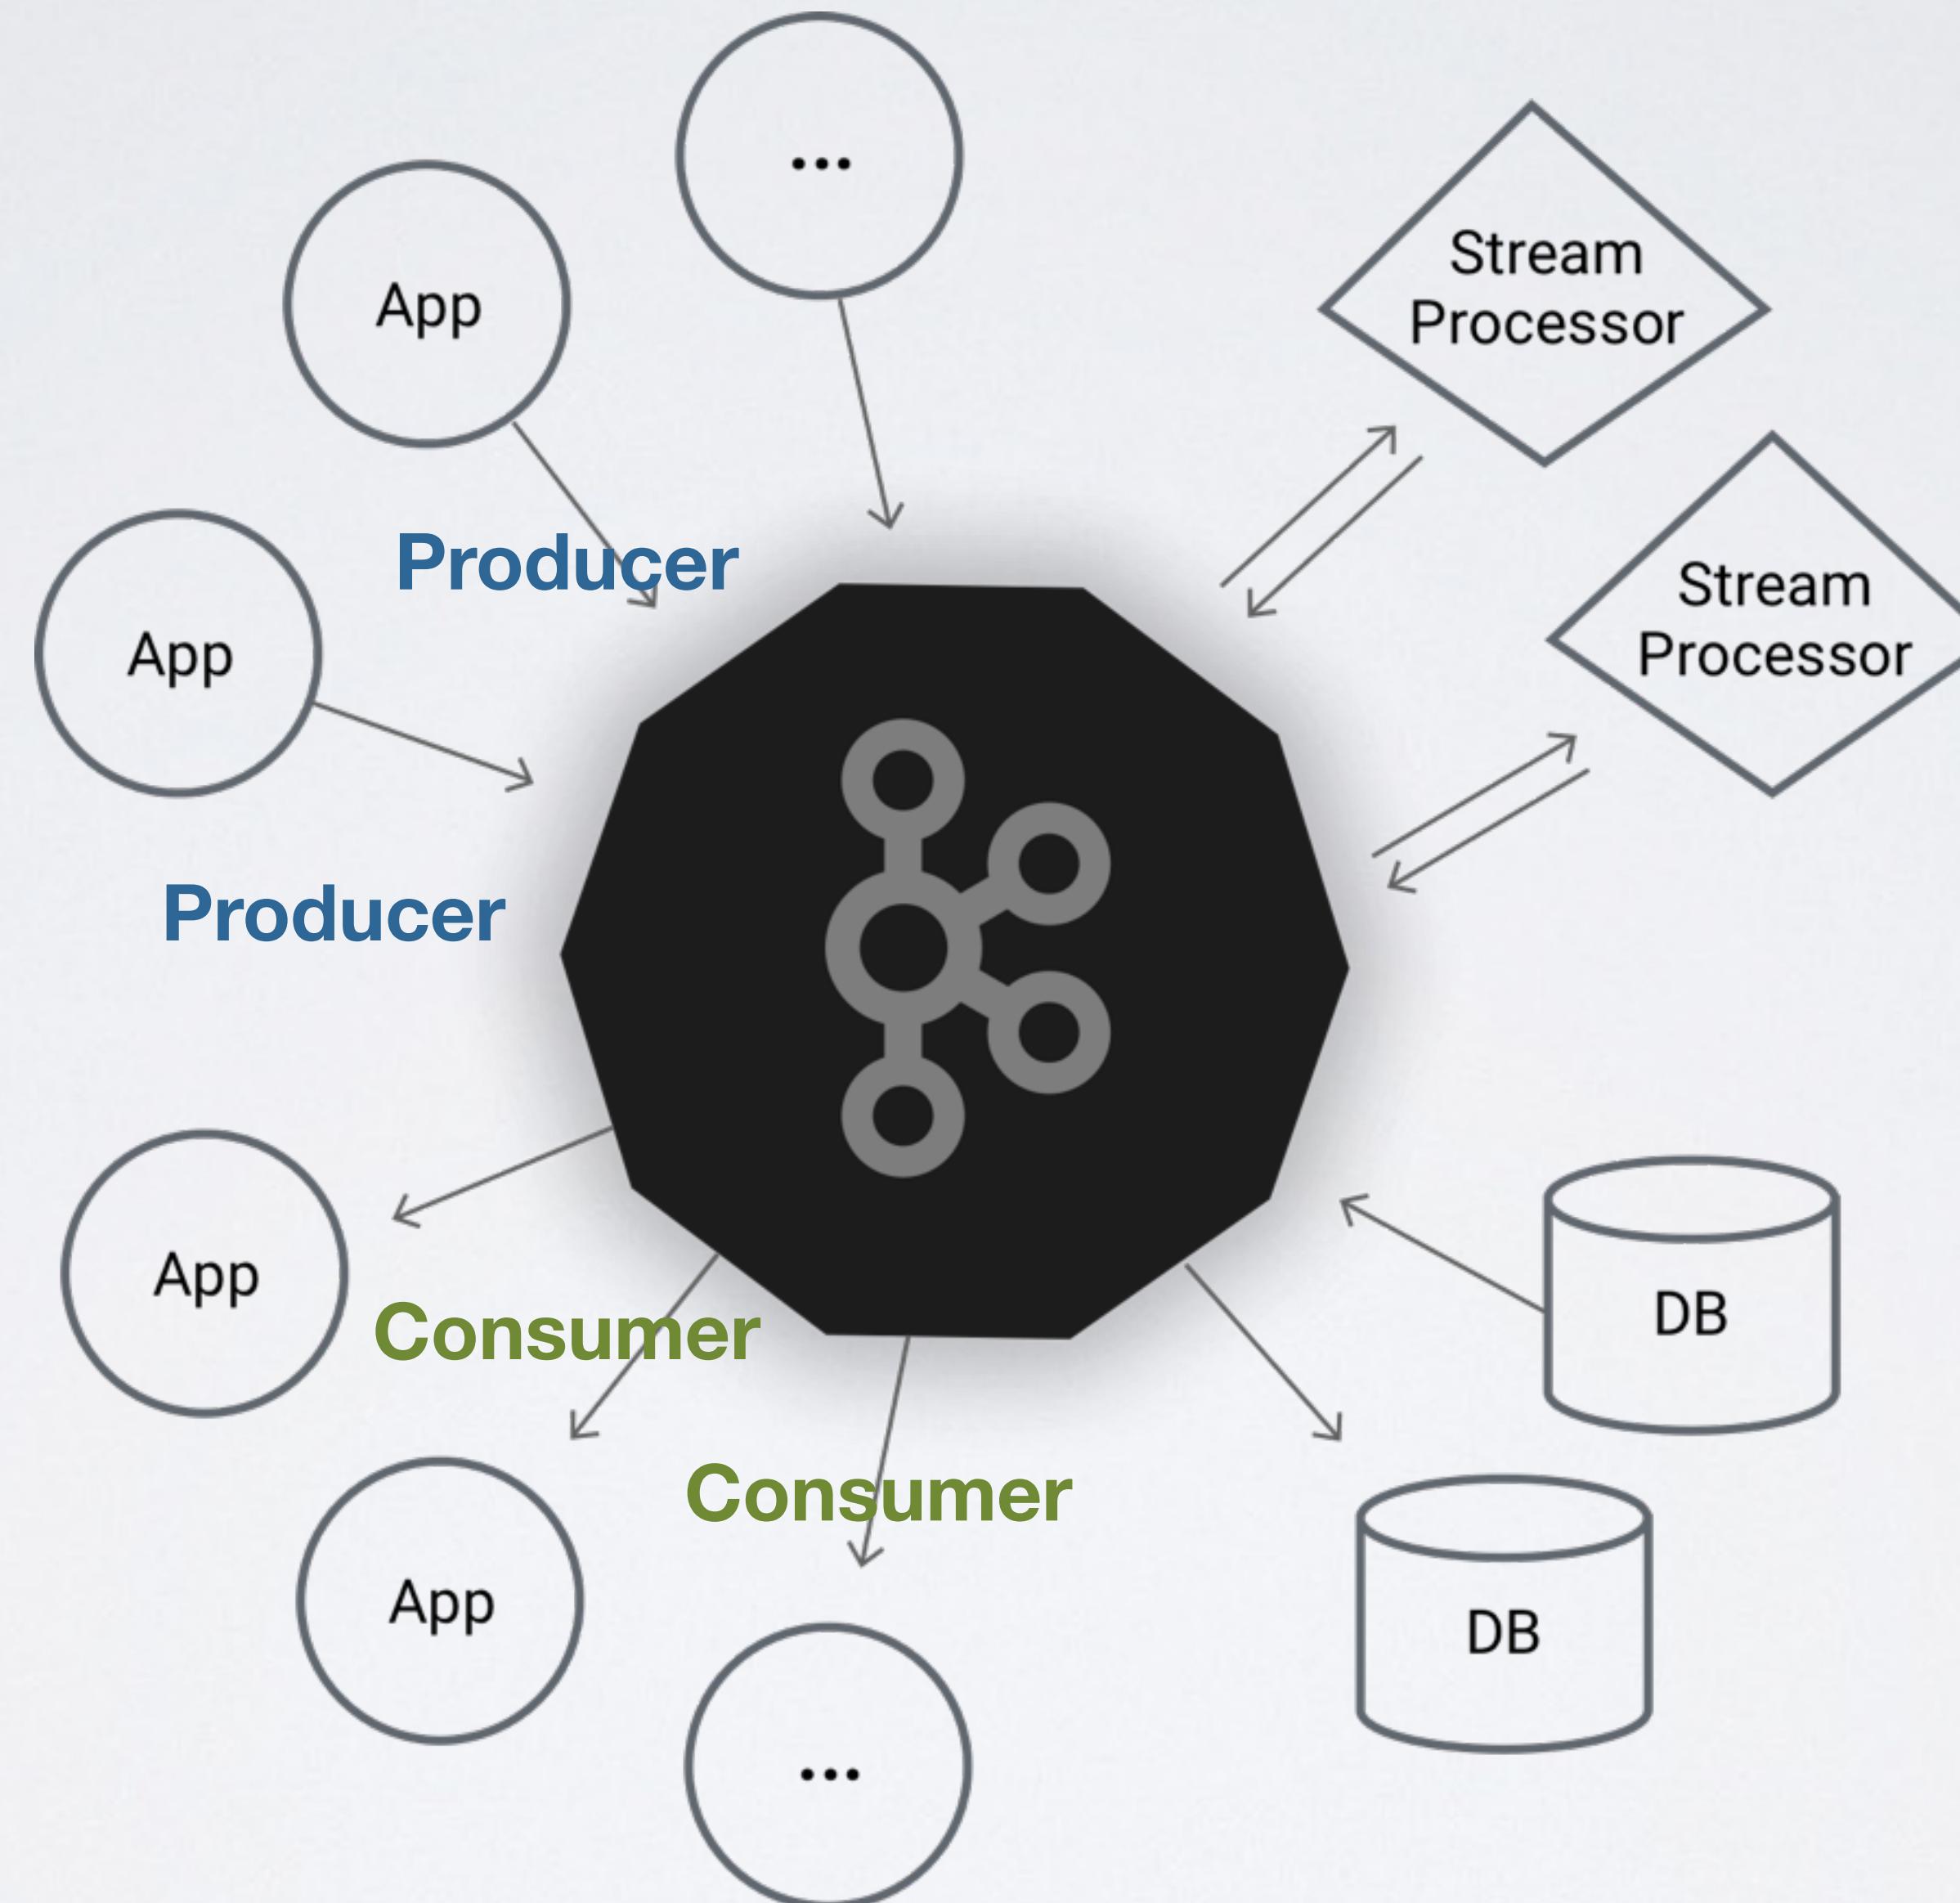
CONNECTOR	TAGS	DEVELOPER/SUPPORT	DOWNLOAD
Apache Ignite (Source)	File System	Community	Community
Apache Ignite (Sink)	File System	Community	Community
Bloomberg Ticker (Source)	Application feed	Community	Community
Cassandra (Source)	Cassandra	Community	Community 1
Cassandra (Sink)	Cassandra	Community	Community
DynamoDB	Dynamo, NoSQL	Community	Community
Elasticsearch (Sink)	Elastic, search, log, analytics	Community	Community 1 Community 2 Community 3
FTP (Source)	File System	Community	Community
Google PubSub (Source)	Messaging	Community	Community
Google PubSub (Sink)	Messaging	Community	Community
Hazelcast (Sink)	Datastore, In-memory	Community	Community
Hbase (Sink)	Hbase, NoSQL	Community	Community 1 Community 2
InfluxDB (Sink)	Datastore, Time-series	Community	Community
Jenkins (Source)	Application feed	Community	Community
JMS (Sink)	Messaging	Community	Community
Kudu (Sink)	Kudu	Community	Community
Mixpanel (Source)	analytics	Community	Community
MongoDB (Source)	Mongo, MongoDB, NoSQL	Community	Community
MongoDB CDC - Debezium (Source)	MongoDB, CDC	Community	Community
MQTT (Source)	MQTT, messaging	Community	Community
MySQL CDC - Debezium (Source)	MySQL, CDC, Oracle	Community	Community

102

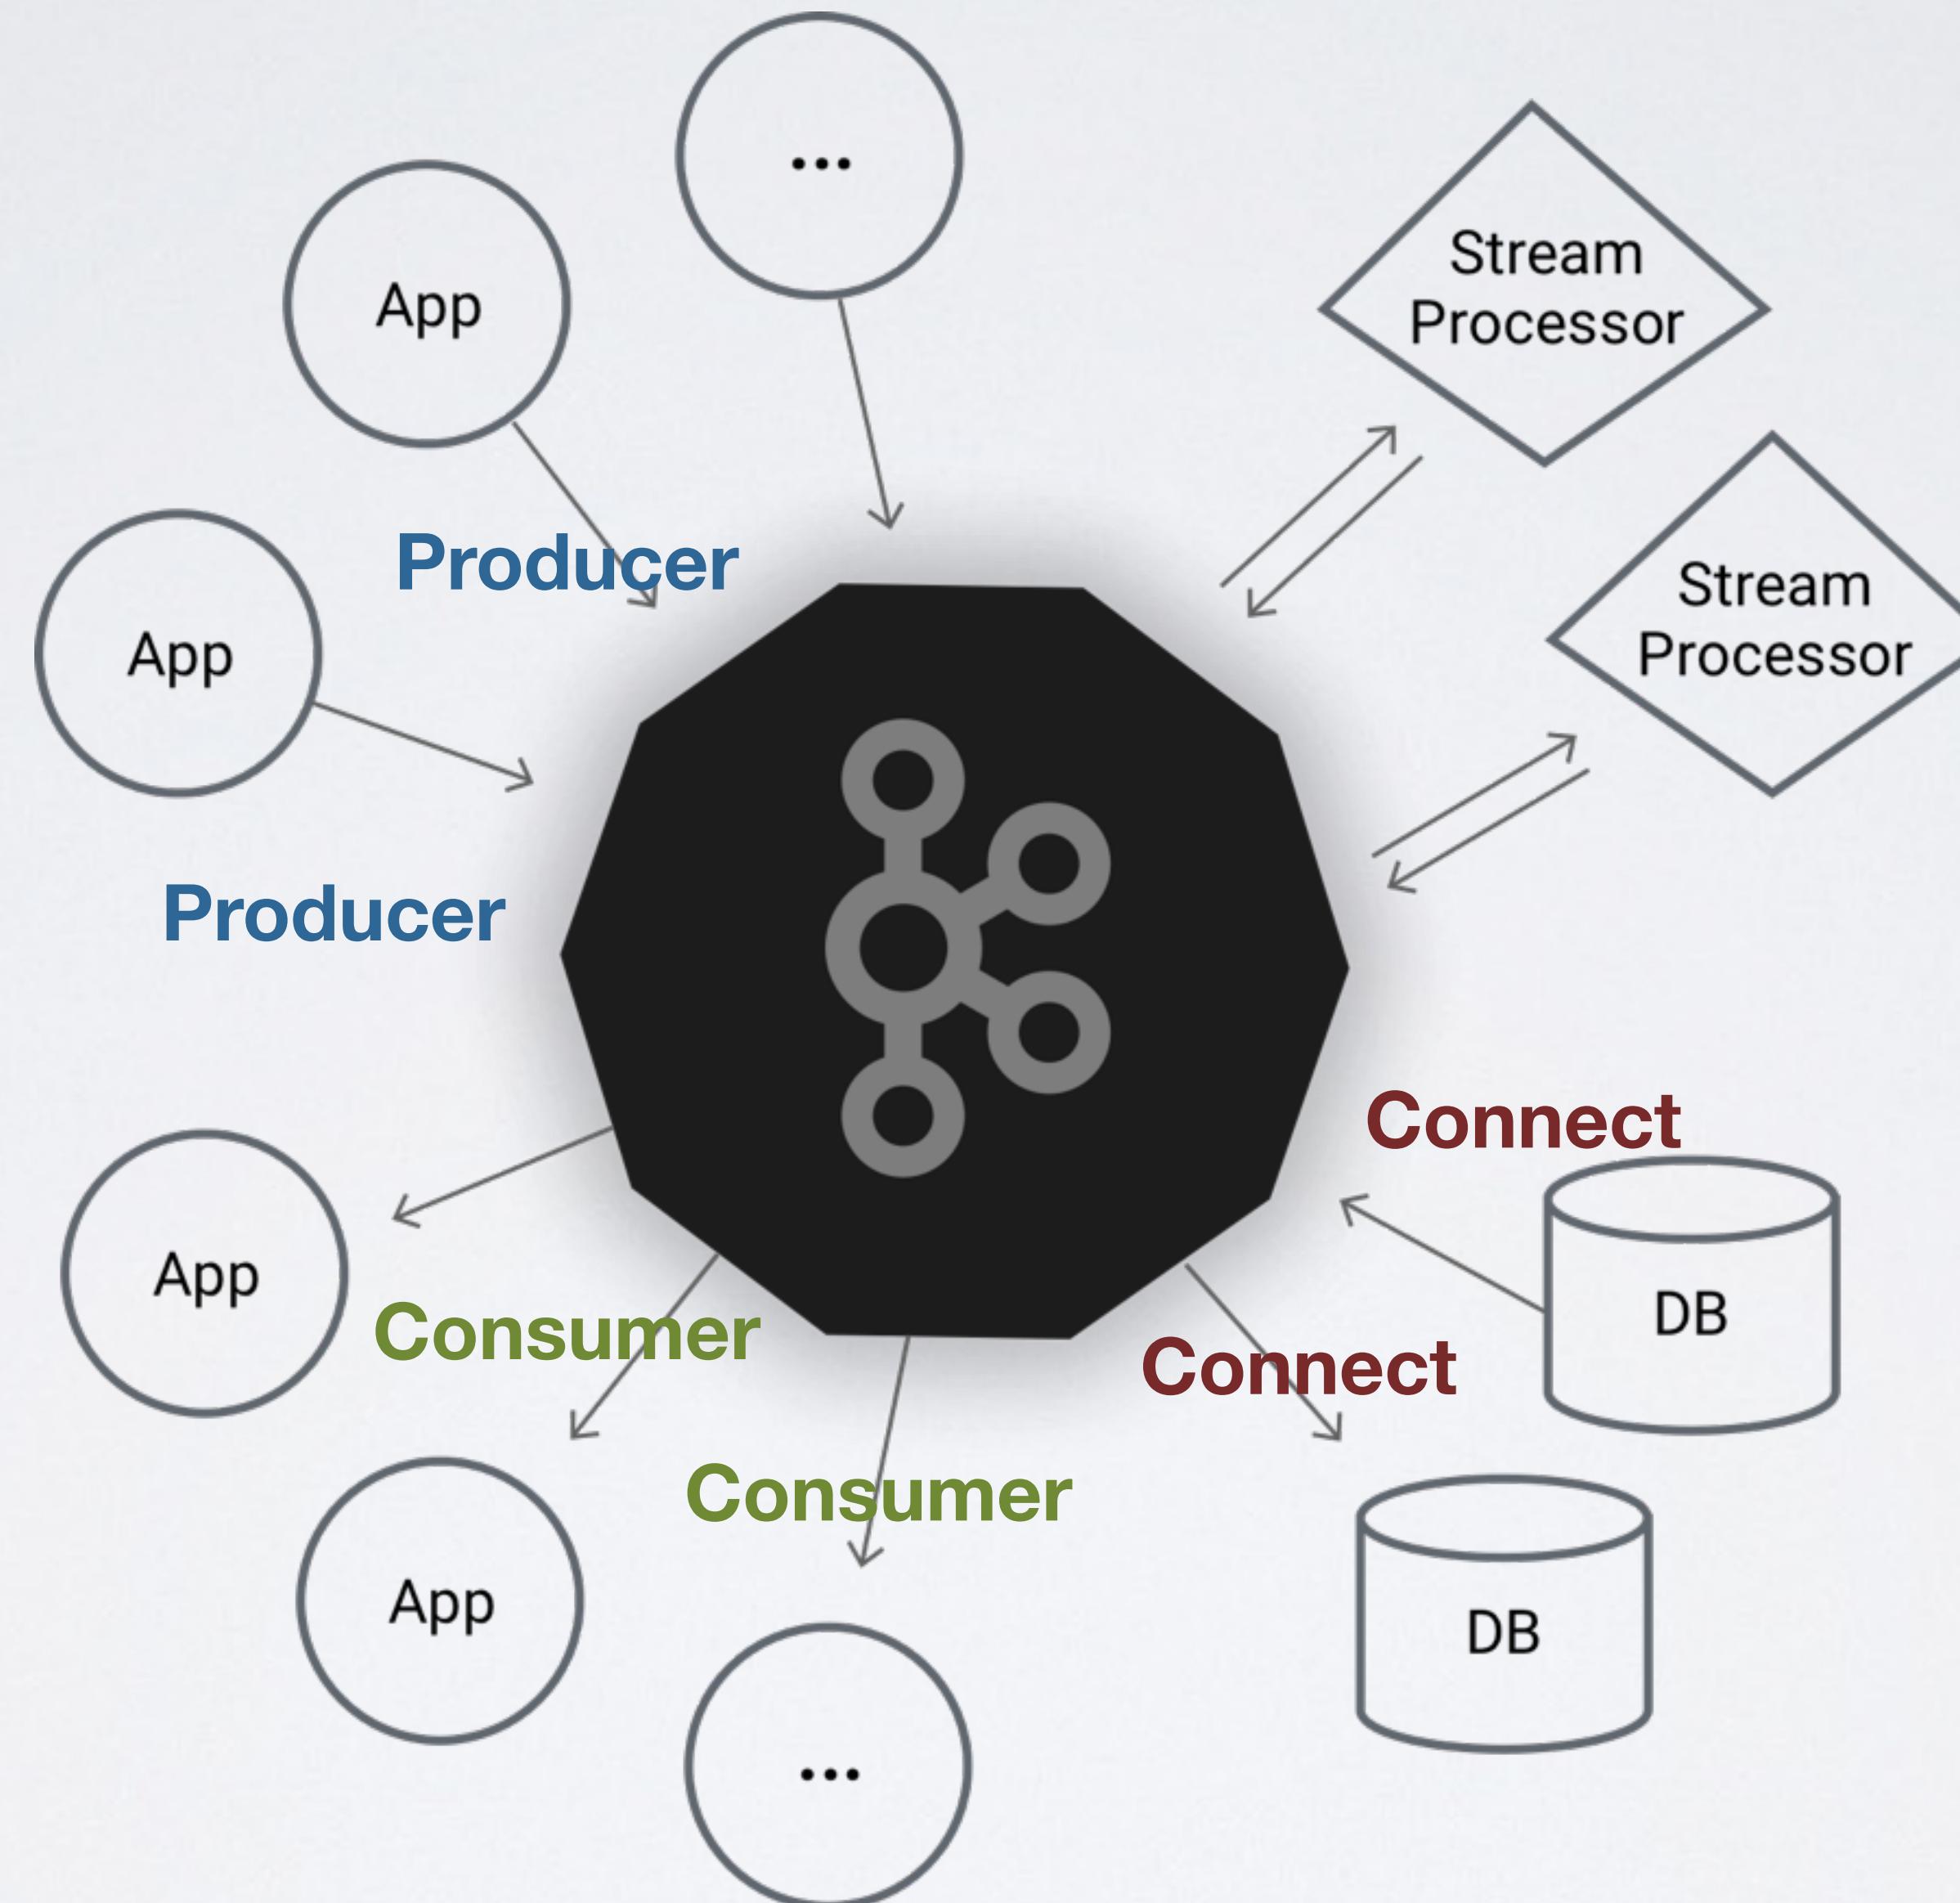
Data Streaming Platform



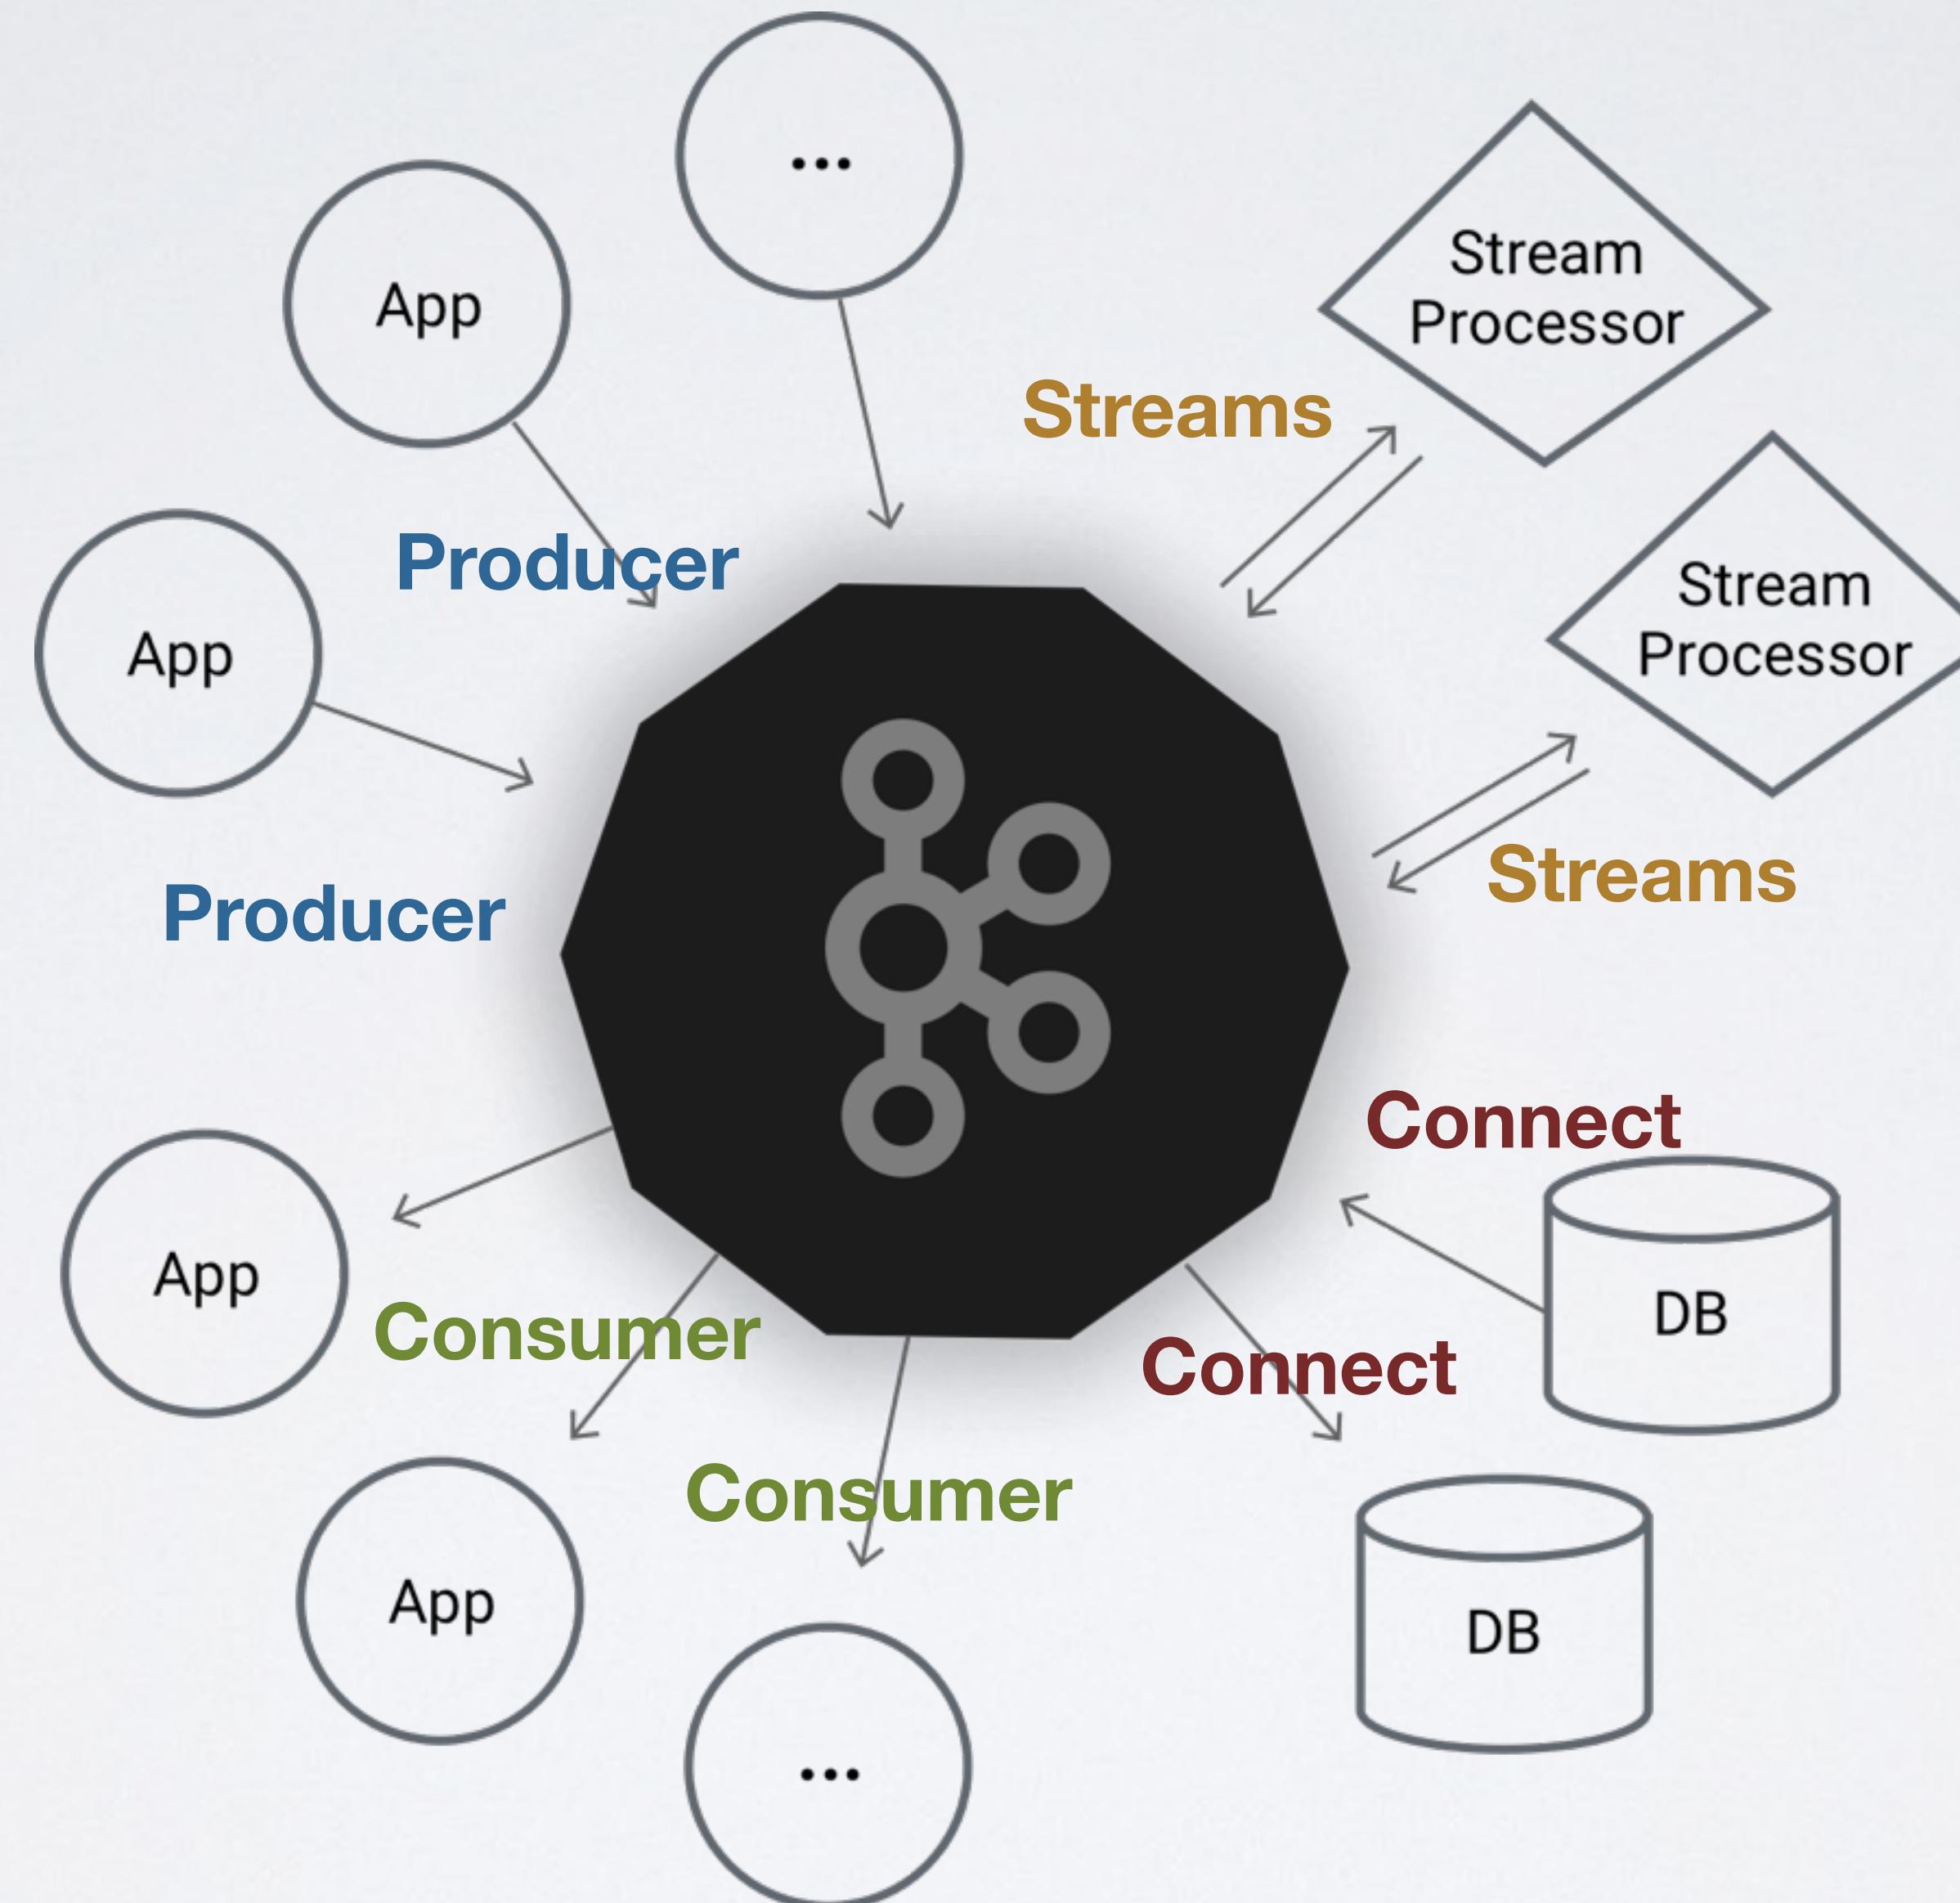
Data Streaming Platform



Data Streaming Platform



Data Streaming Platform



Take-aways

- **Stream Processing: a new programming paradigm**

Take-aways

- **Stream Processing:** *a new programming paradigm*
-  **Kafka Streams:** *stream processing made easy*

THANKS!

Take-aways

- **Stream Processing:** *a new programming paradigm*
-  **Kafka Streams:** *stream processing made easy*

Join Kafka Summit 2017 @ NYC & SF

Confluent Webinar: <http://www.confluent.io/resources>

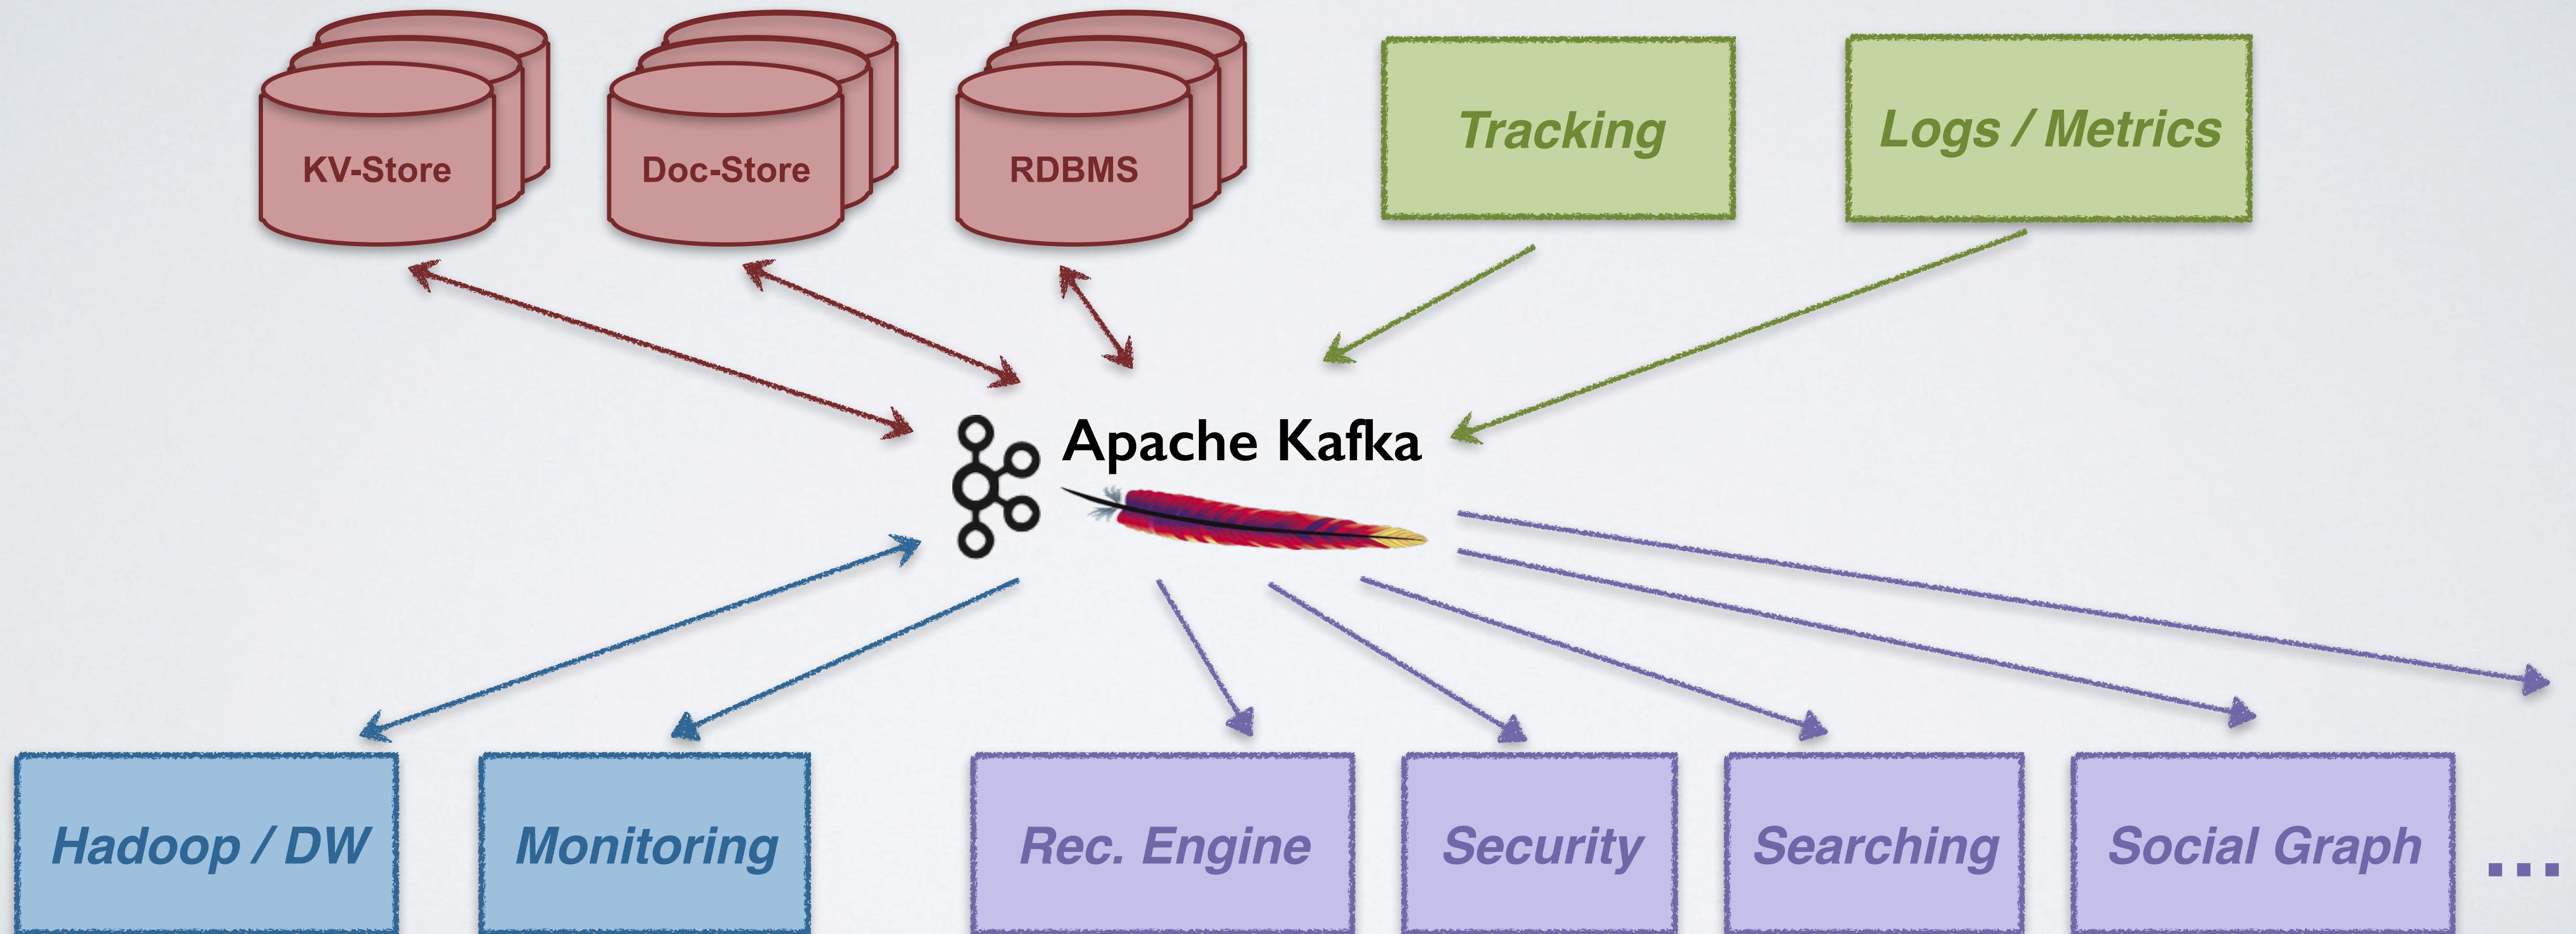
BACKUP SLIDES



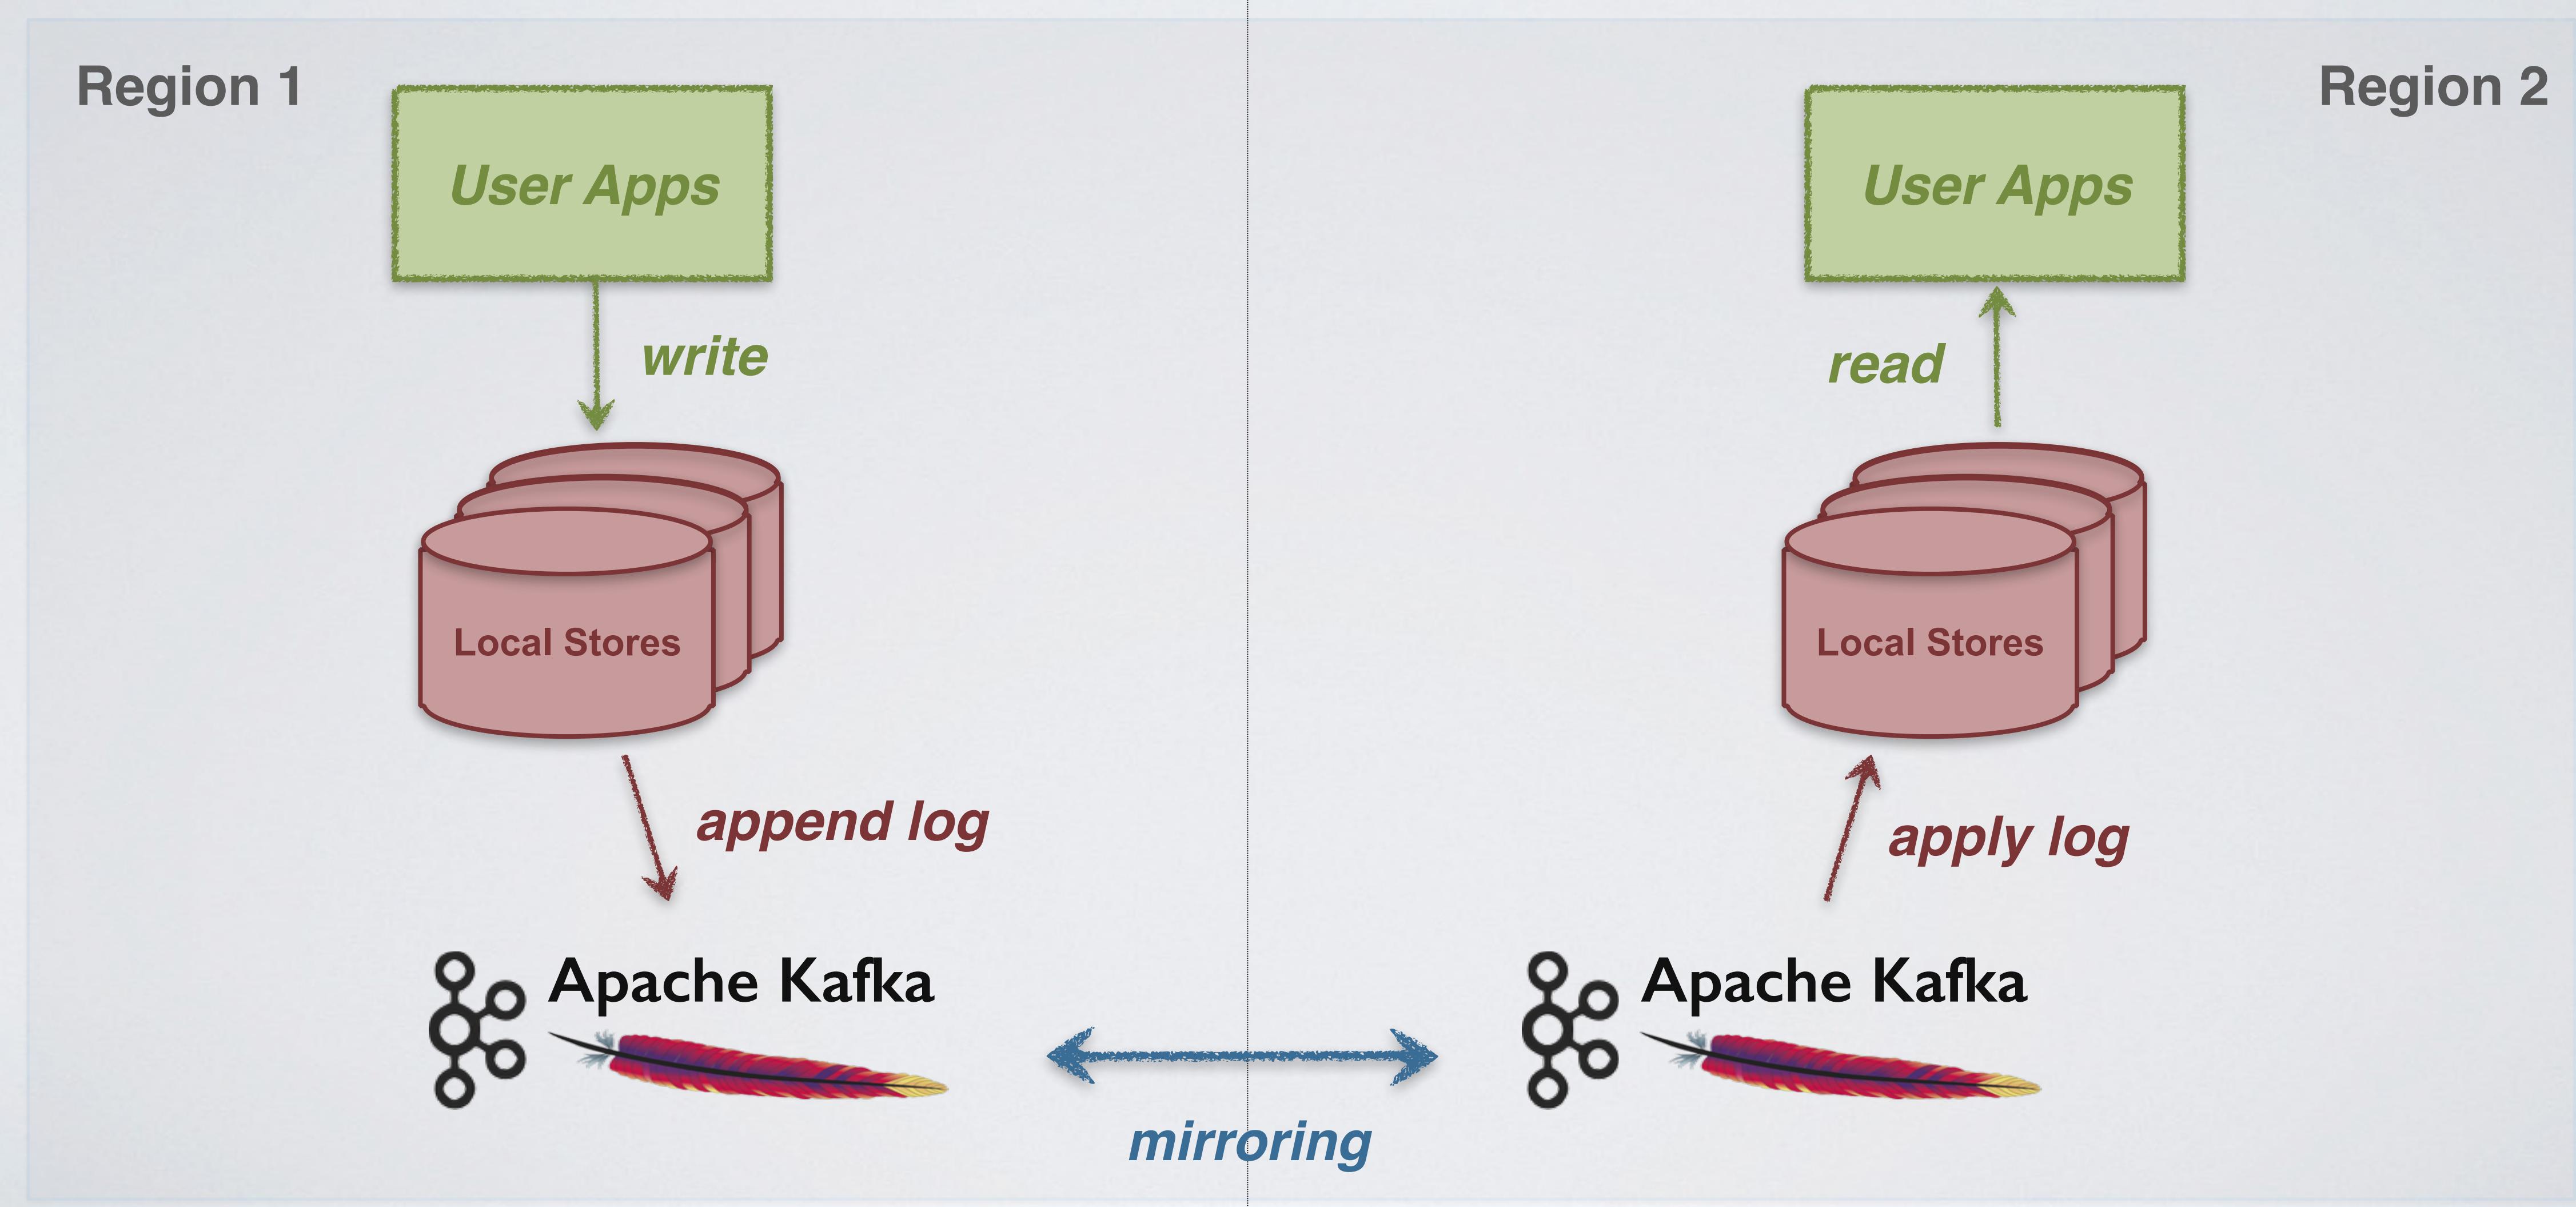
Real-time streams powered by Apache Kafka.

We are Hiring!

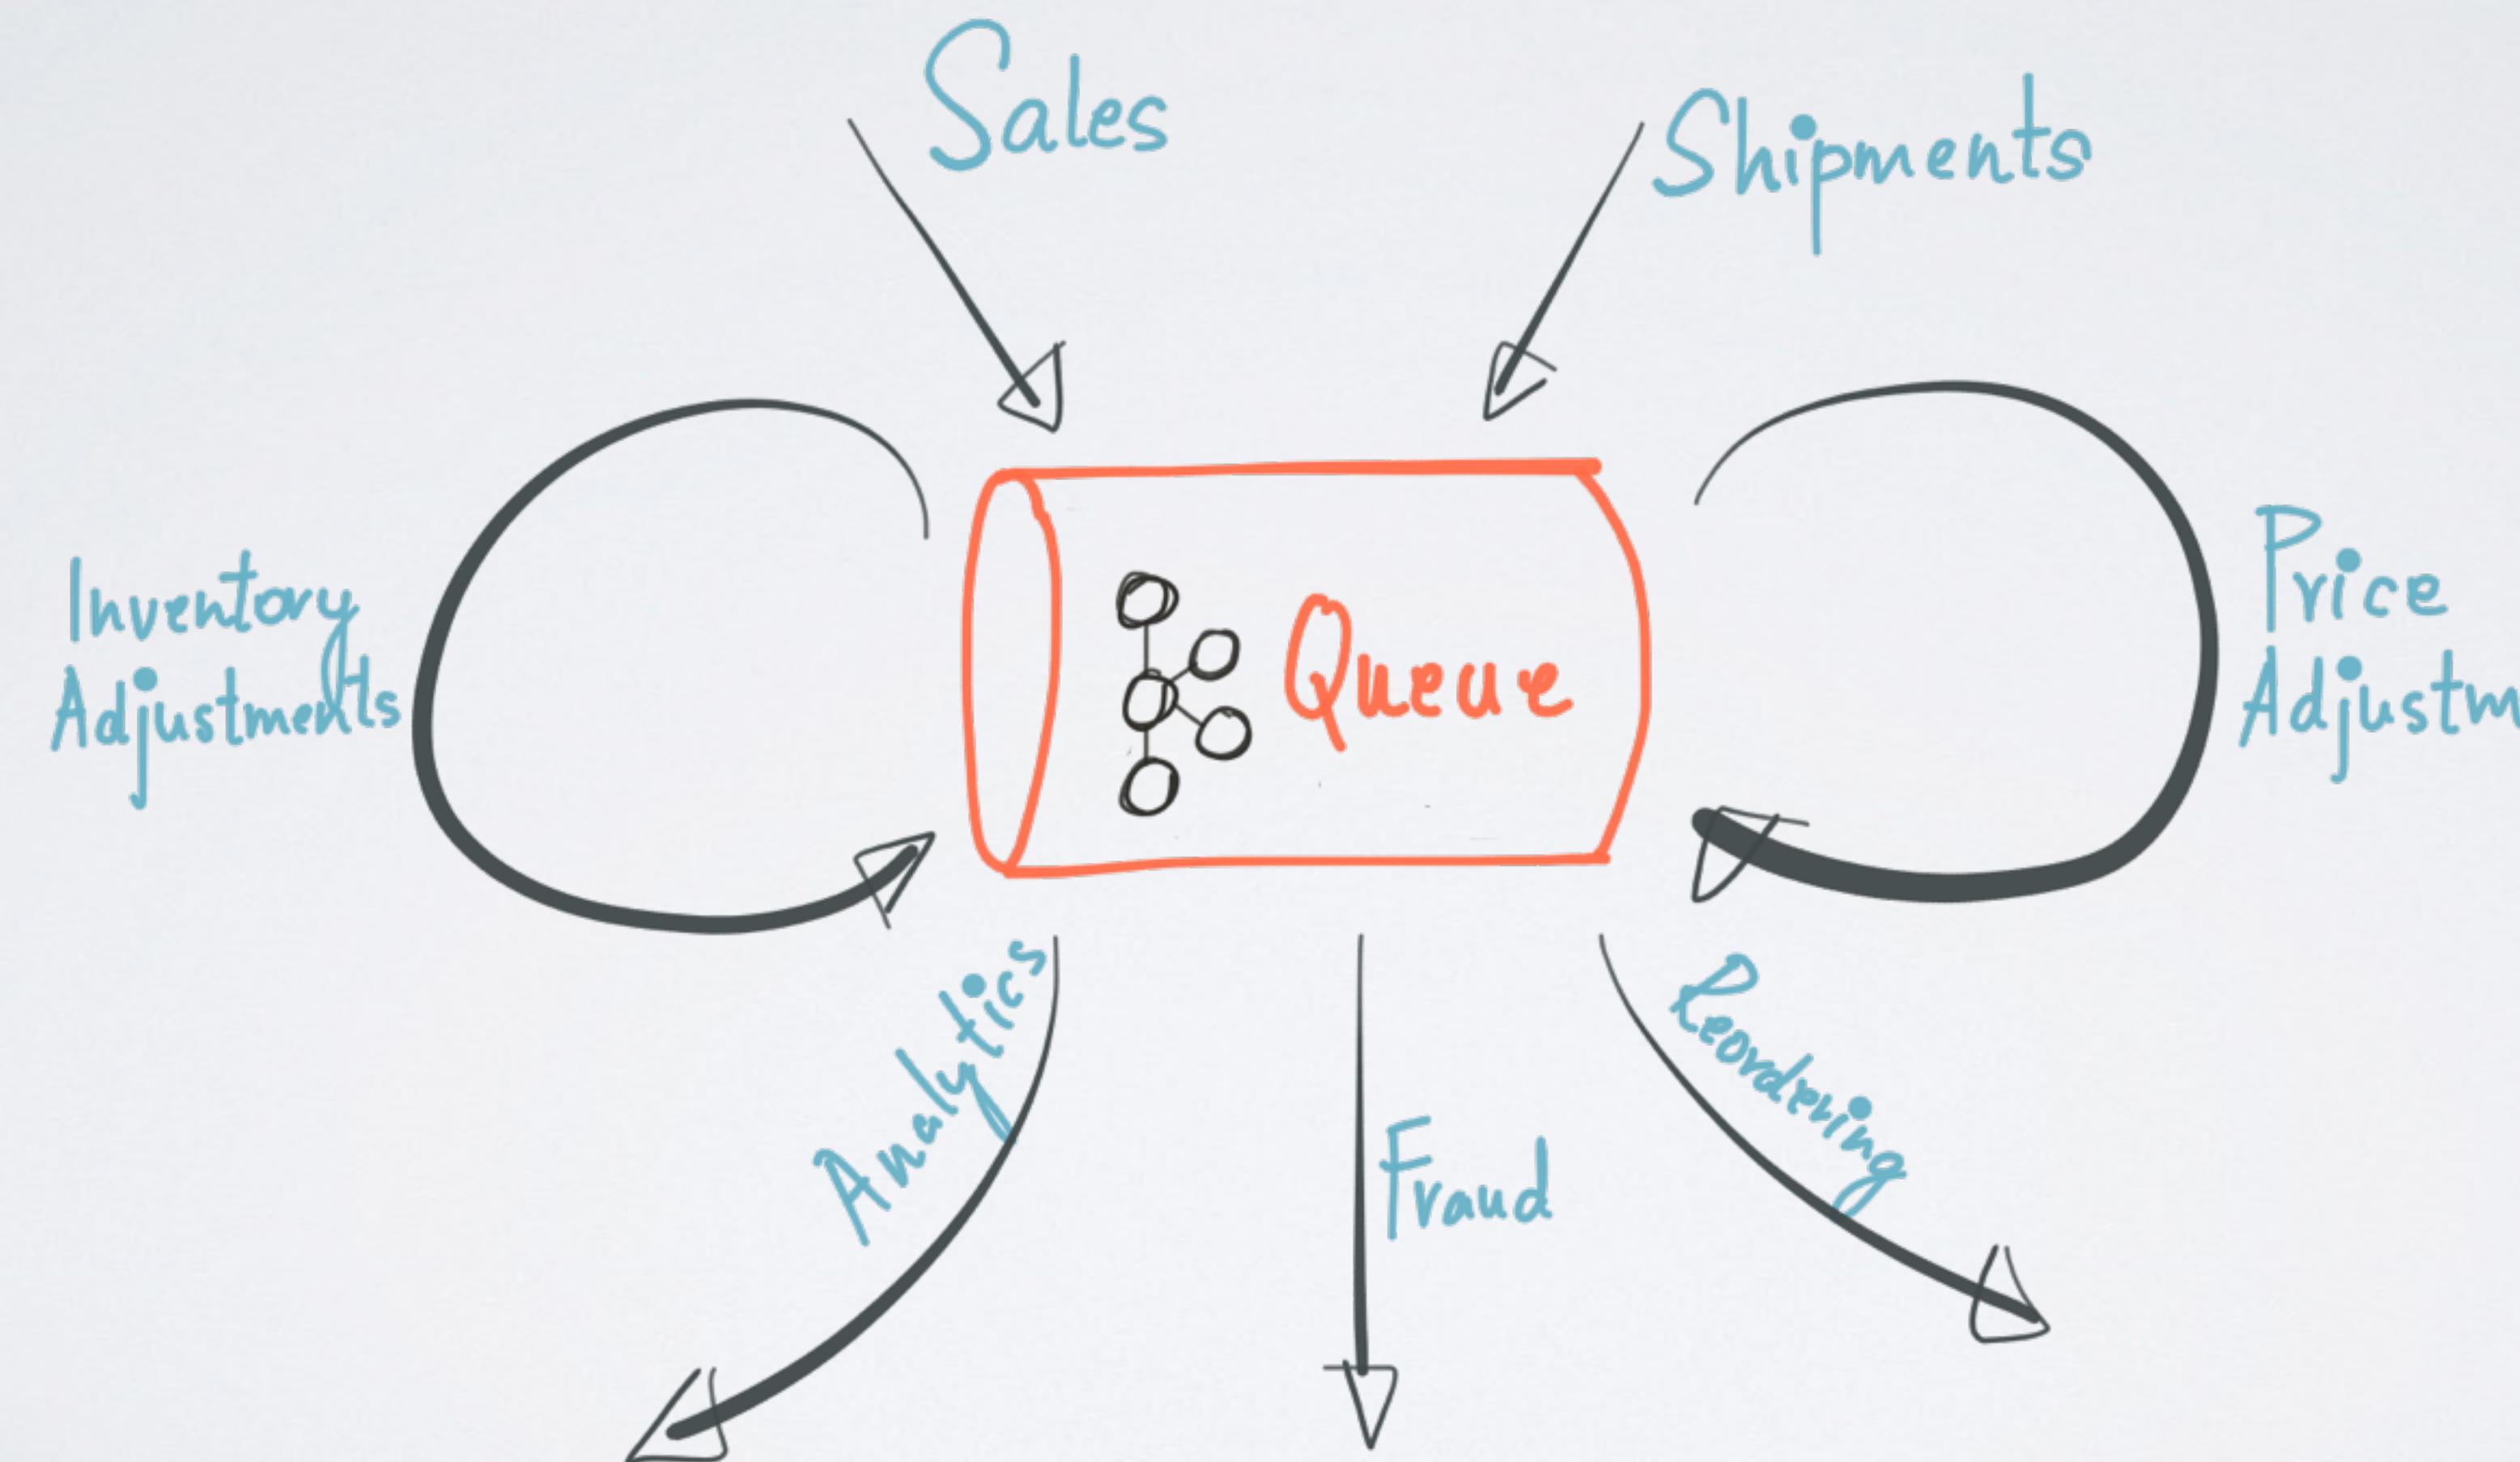
Example: Centralized Data Pipeline



Example: Data Store Geo-Replication



Example: Async. Micro-Services

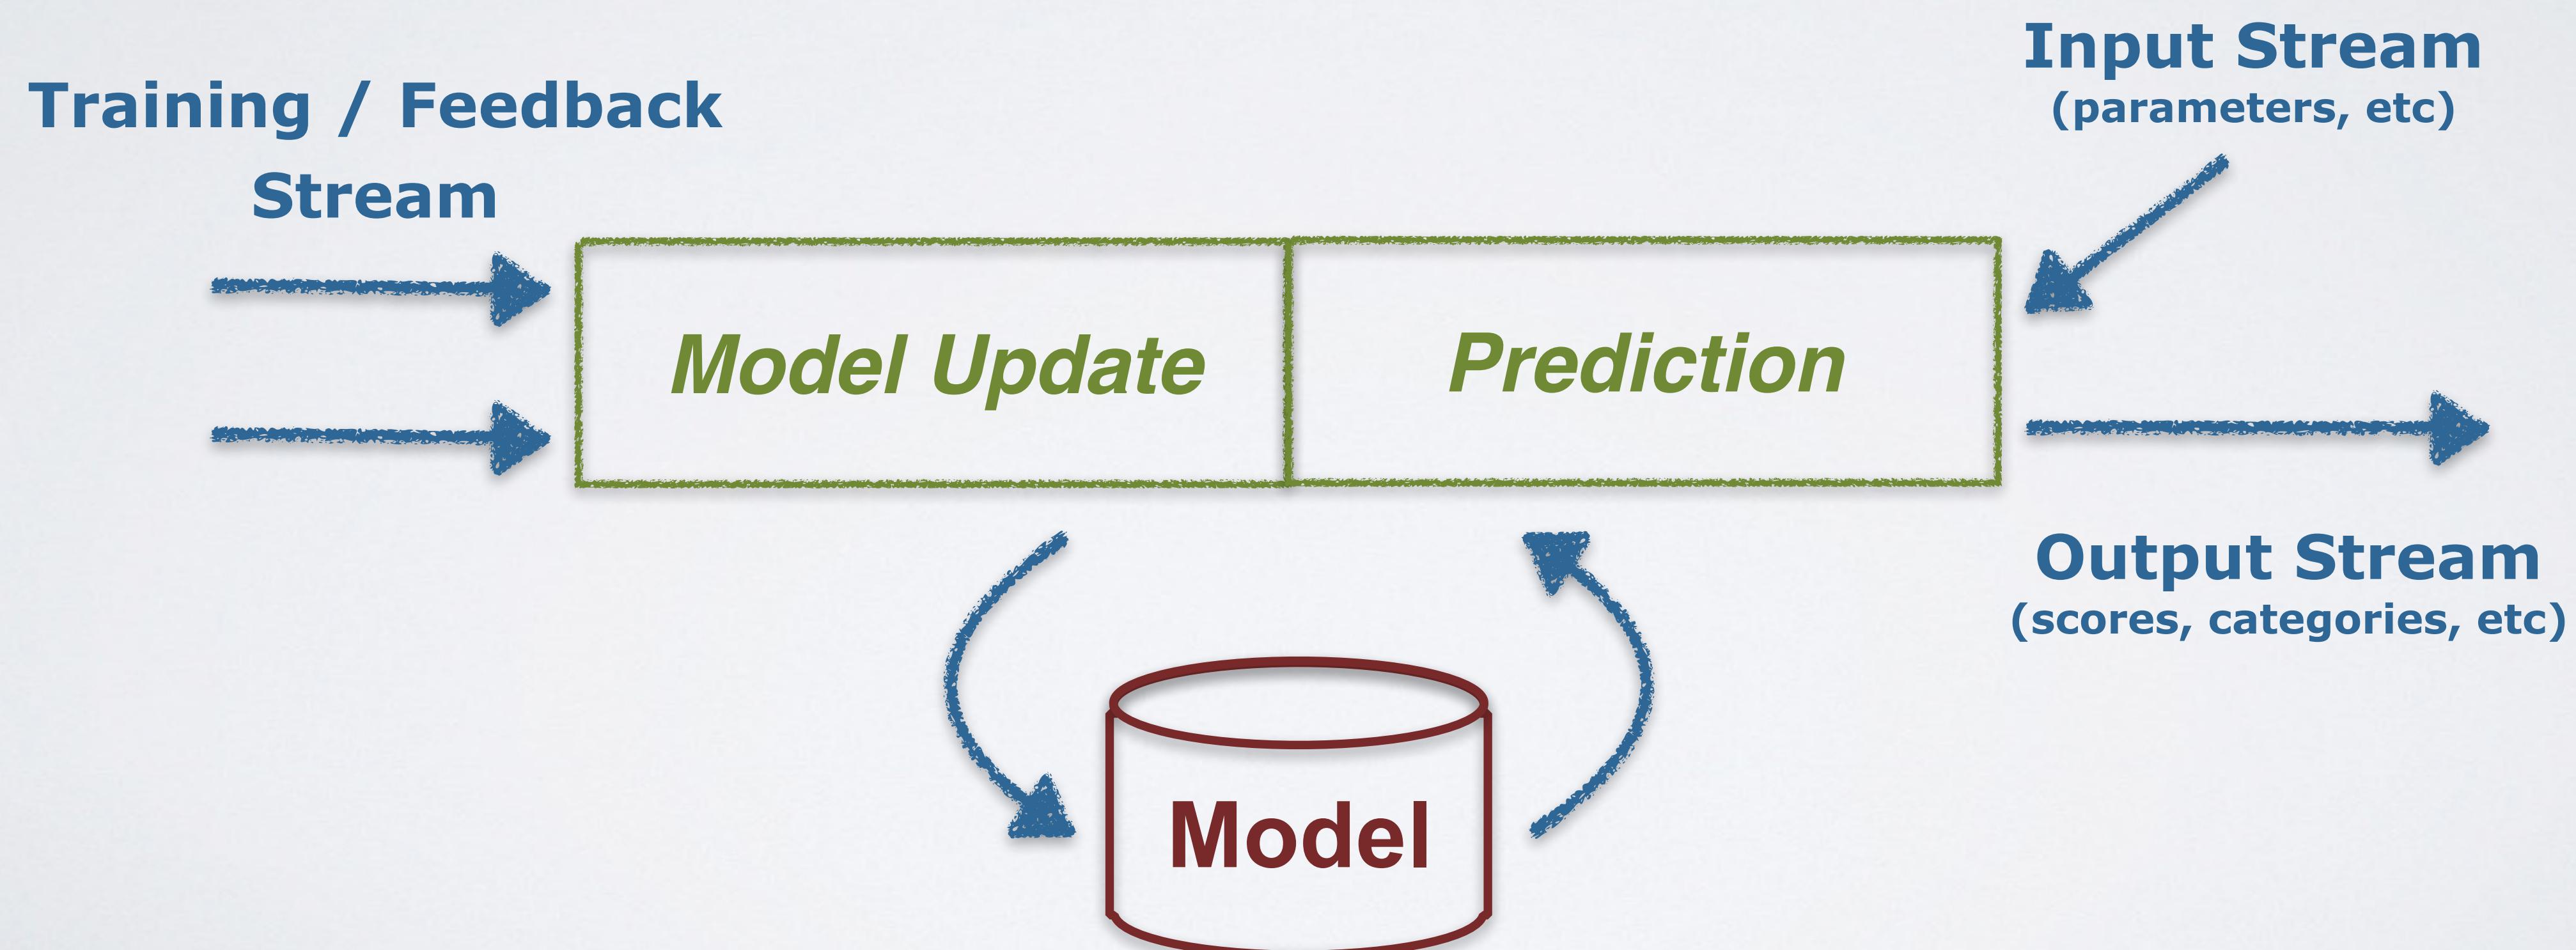


Ongoing Work (0.10+)

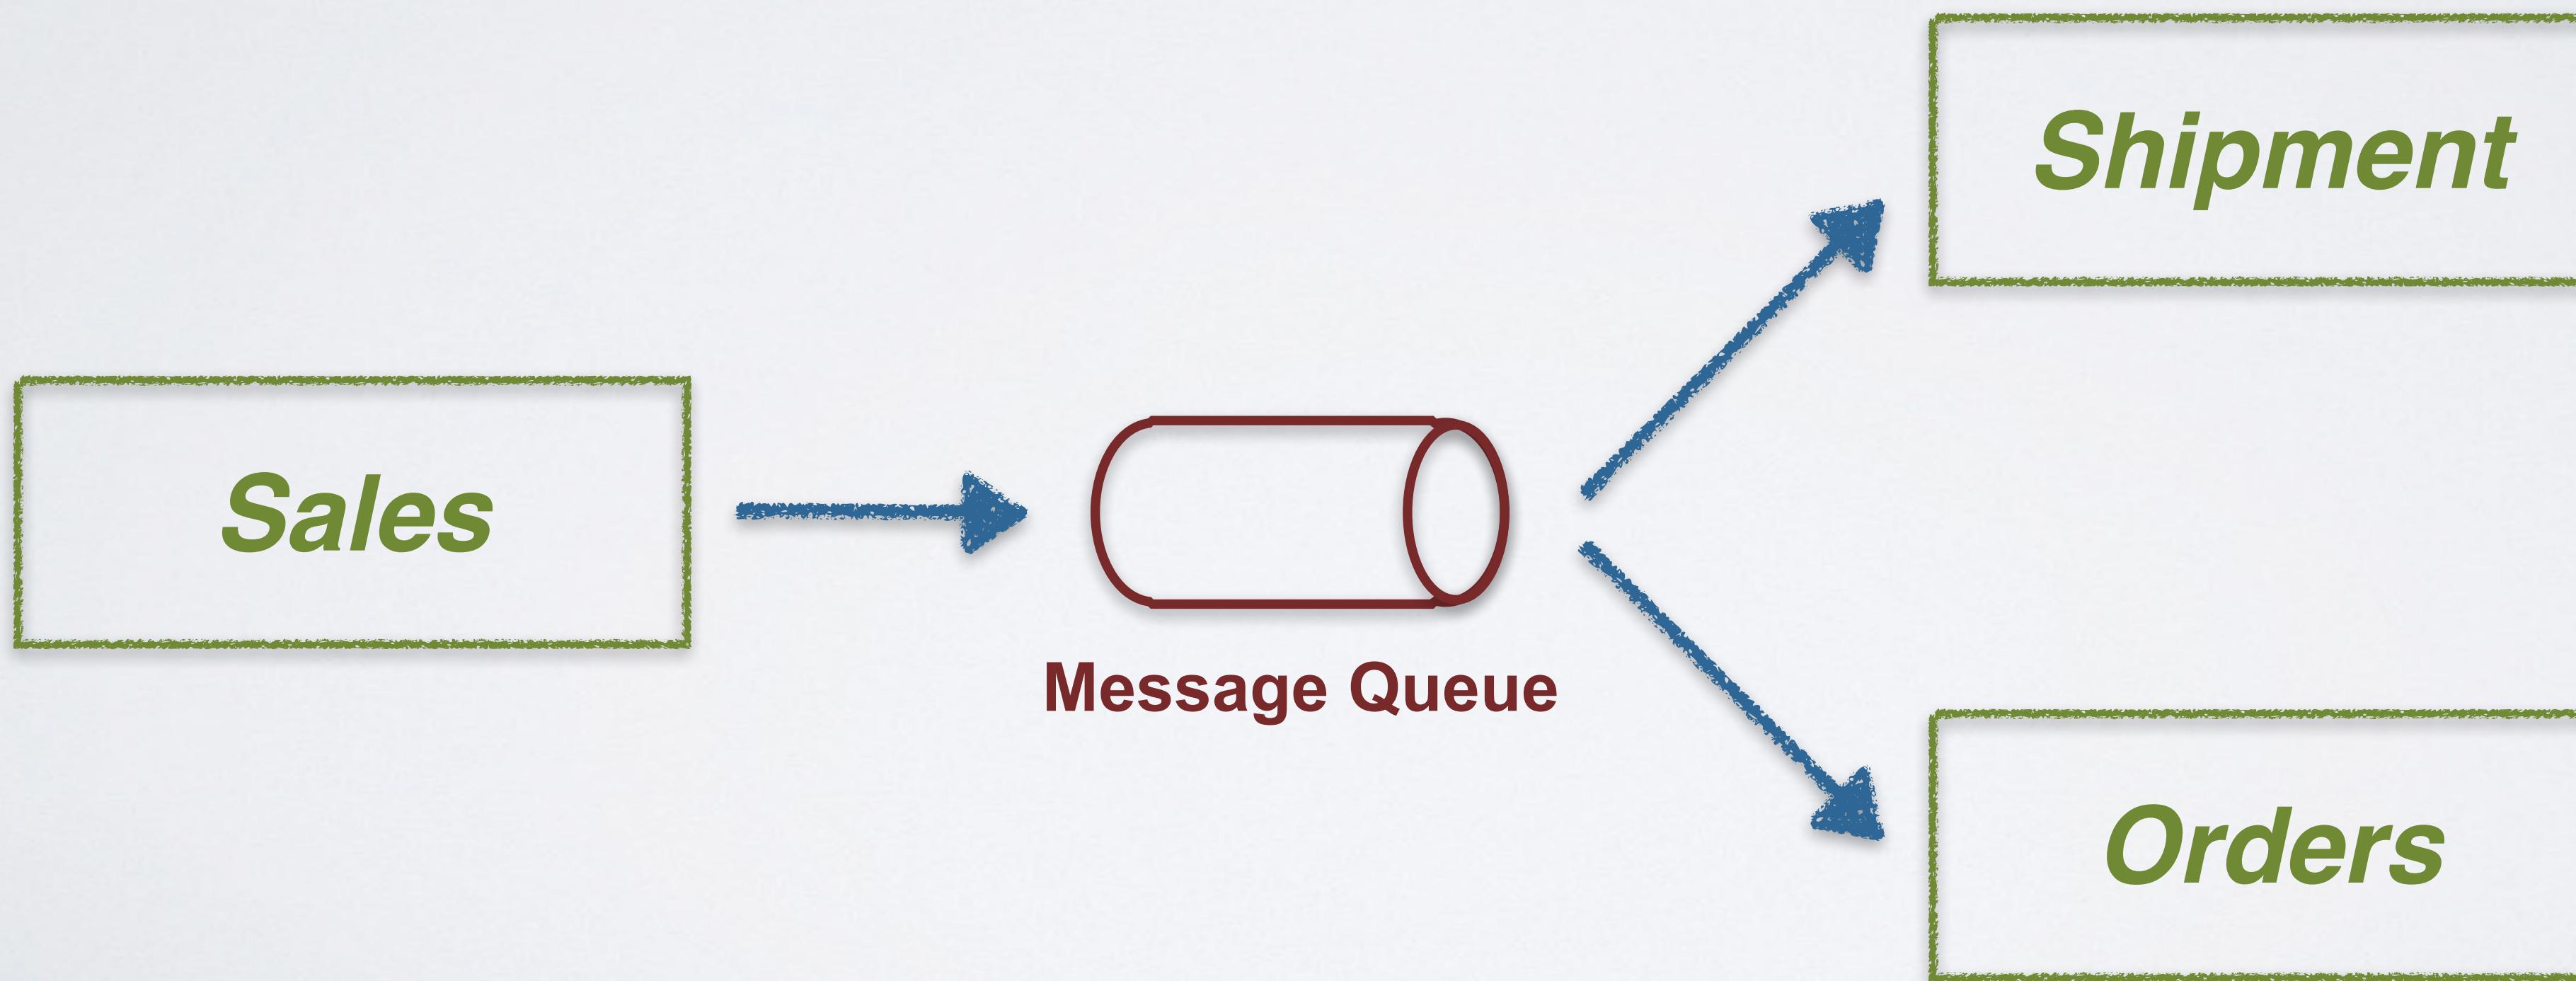
- *Beyond Java APIs*
- *SQL support, Python client, etc*
- *End-to-End Semantics* (exactly-once)
- *... and more*



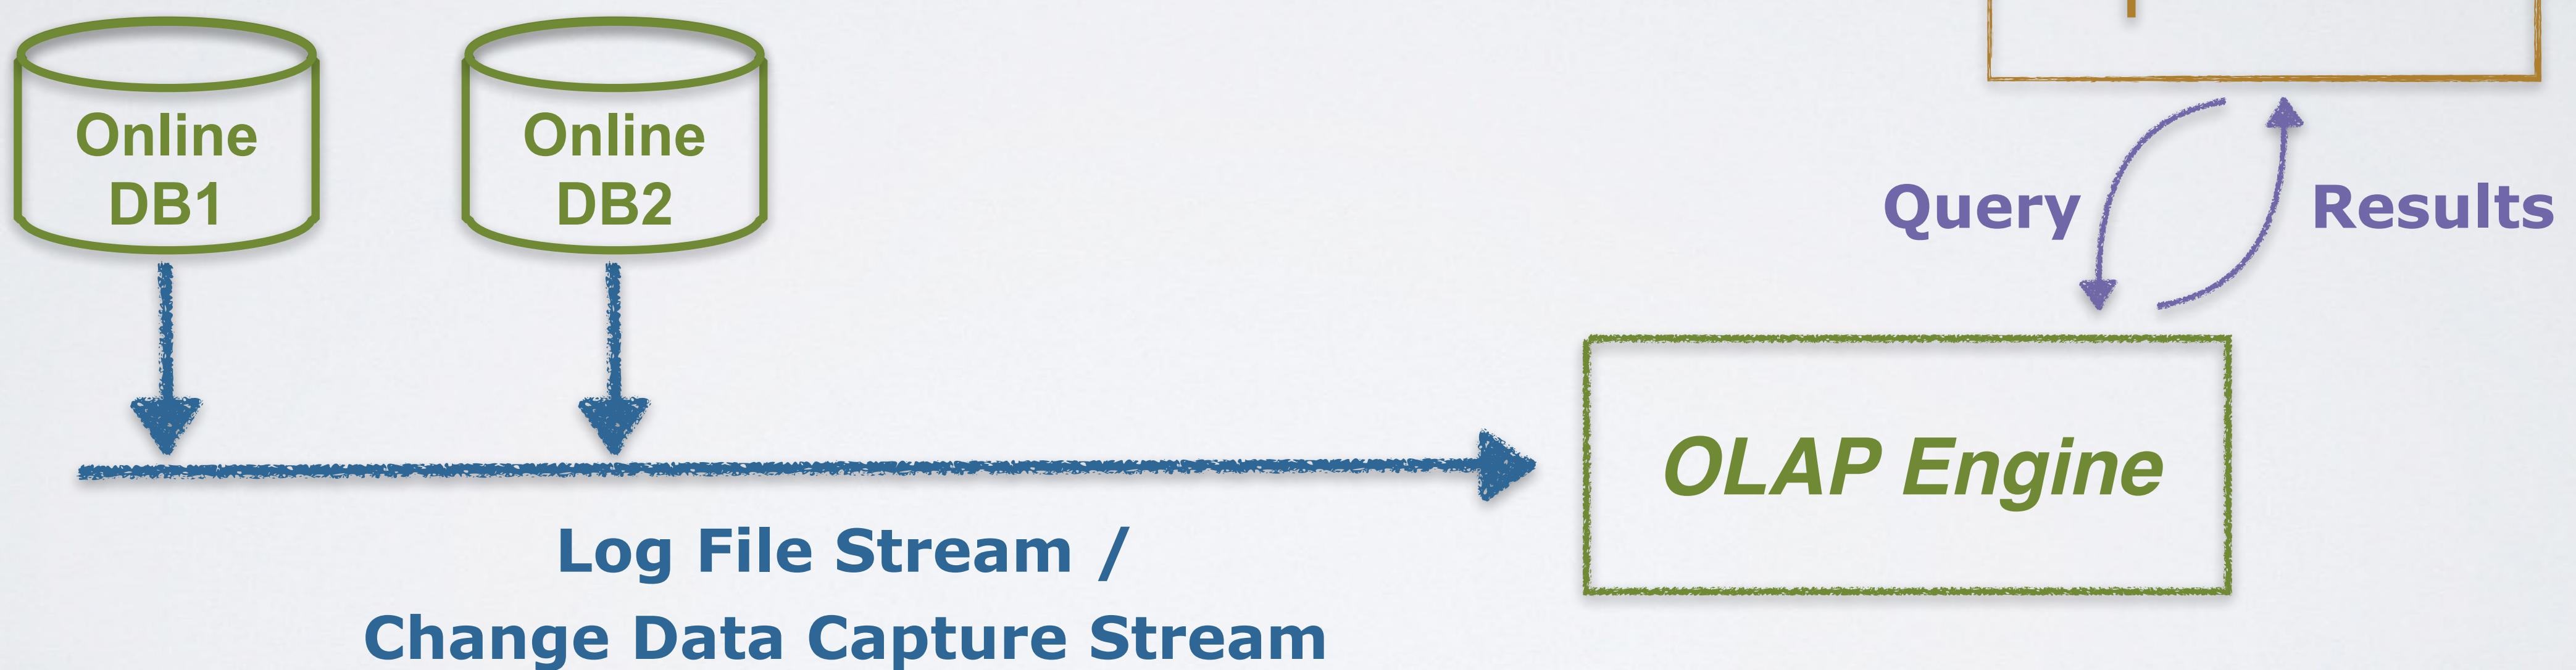
Online Machine Learning



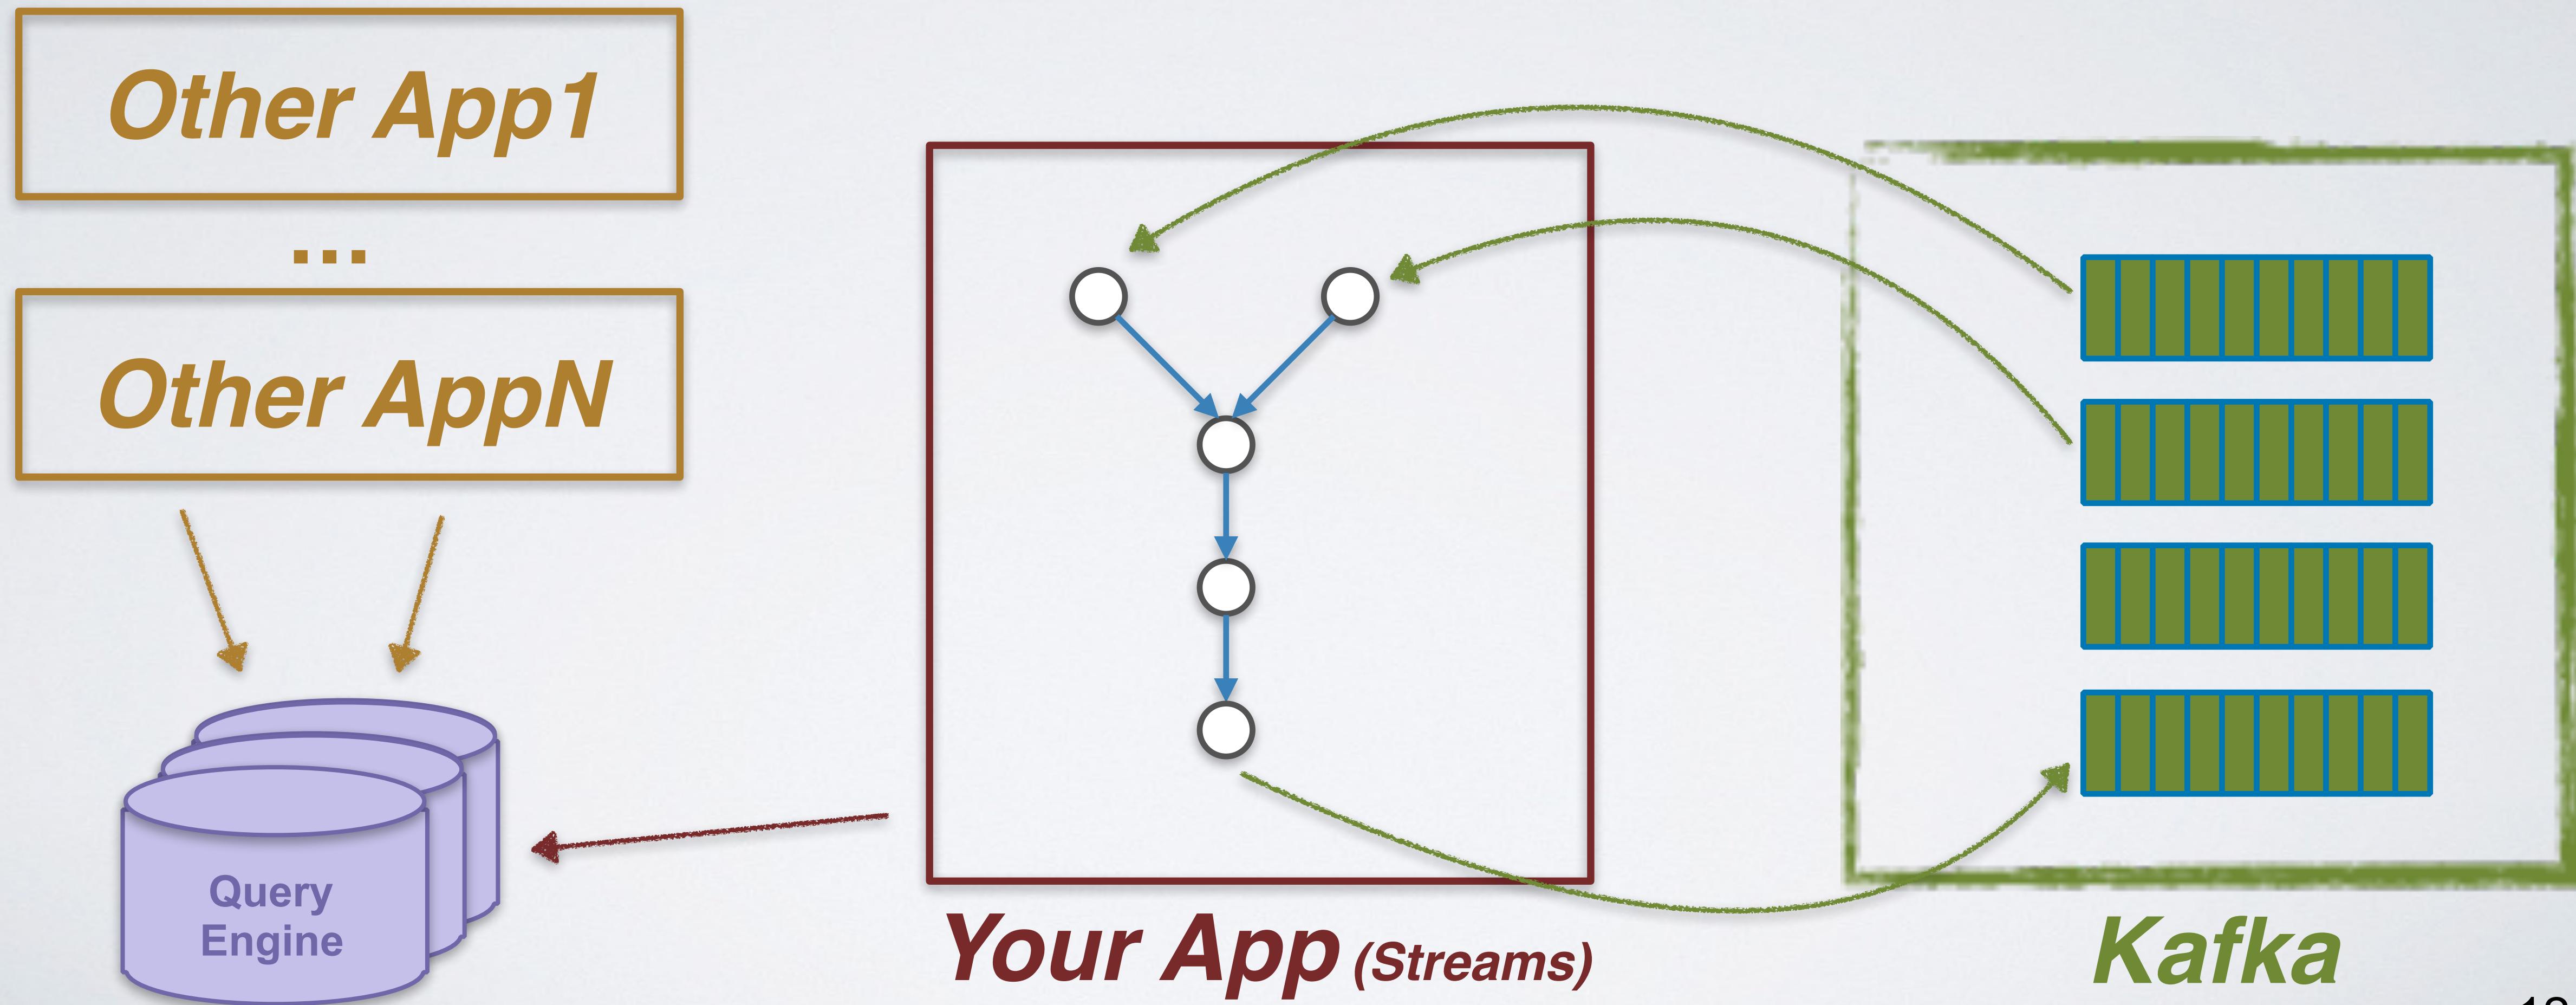
Async. Micro-services



Real-time Analytics



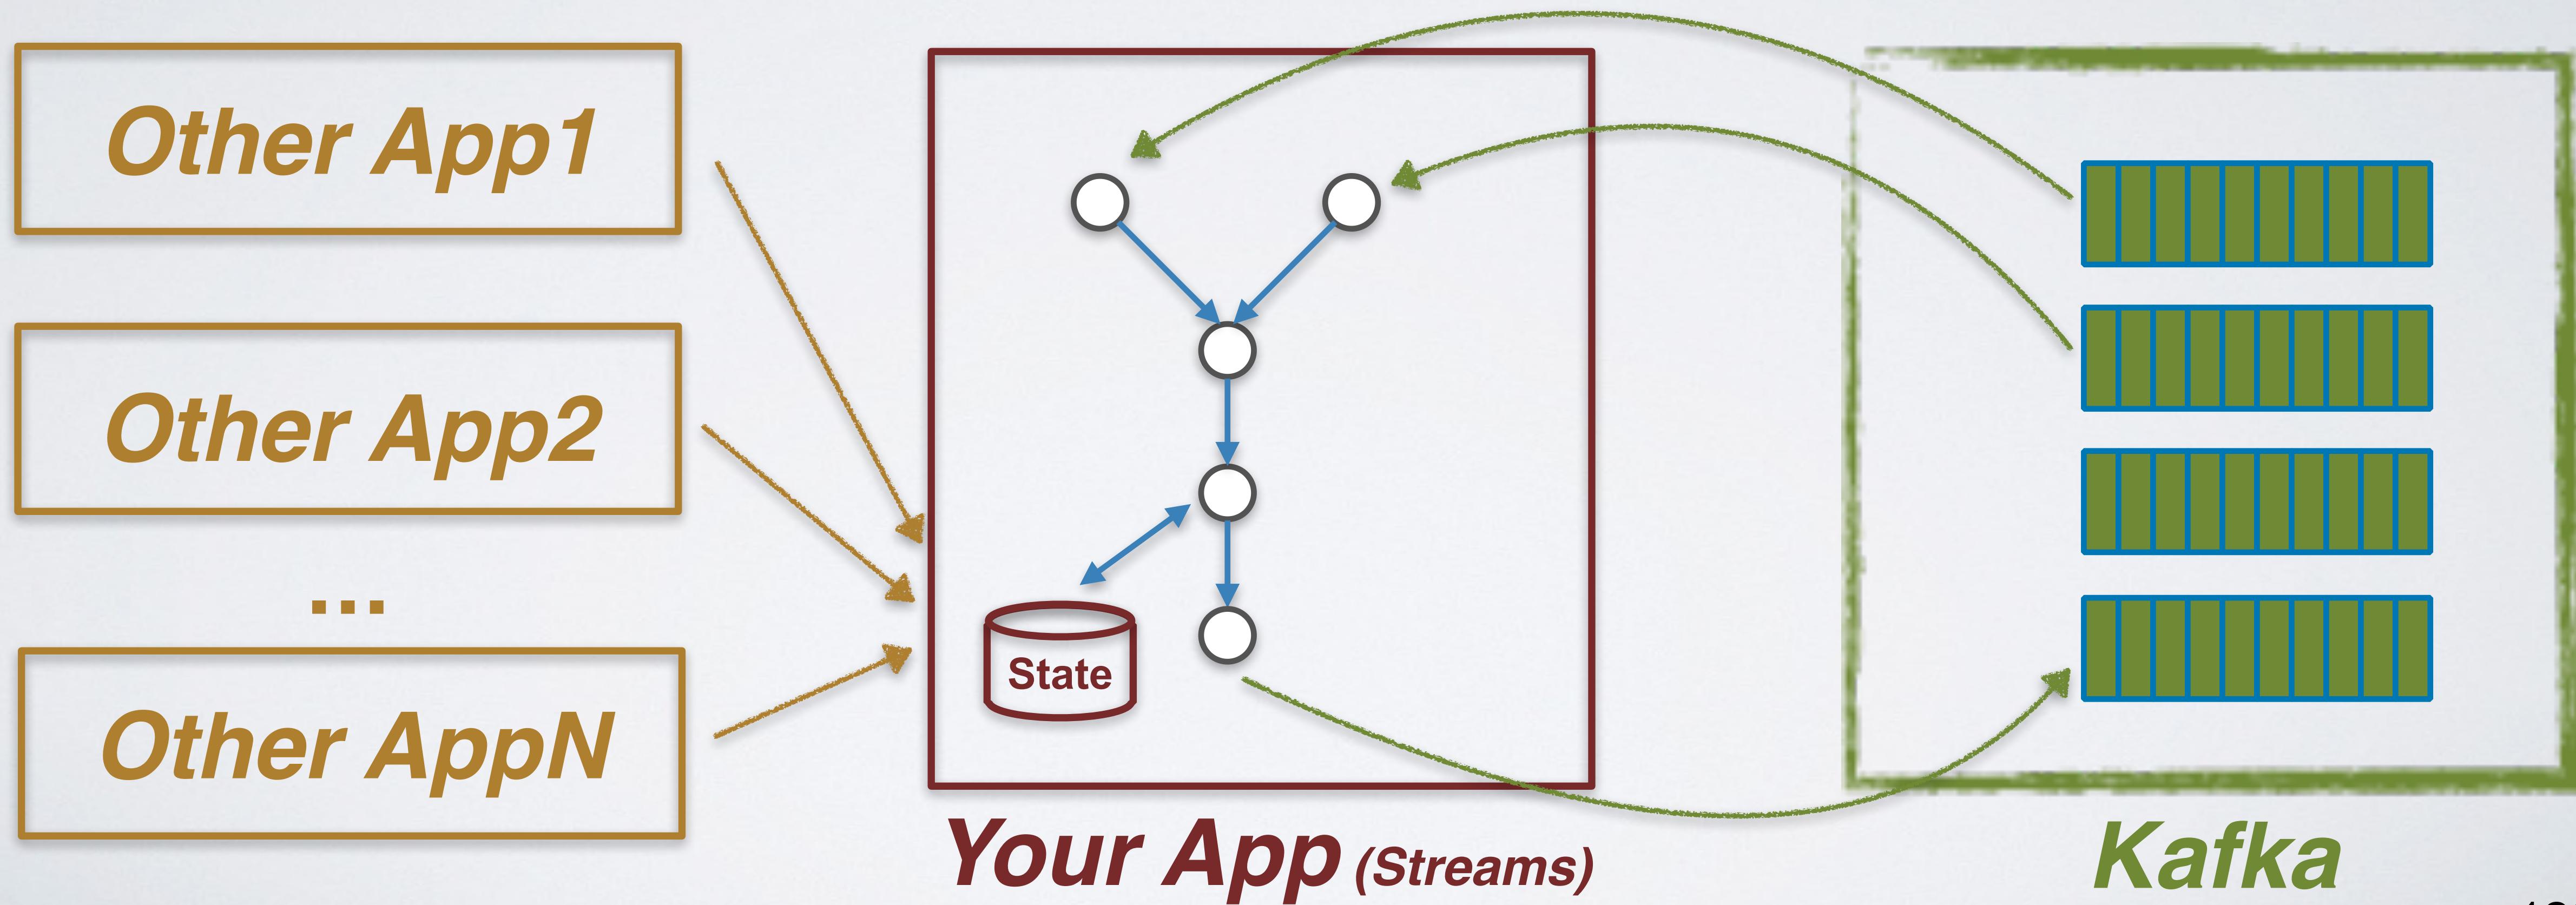
Interactive Queries on States



Interactive Queries on States



Interactive Queries on States (0.10.1+)



Processor Topology

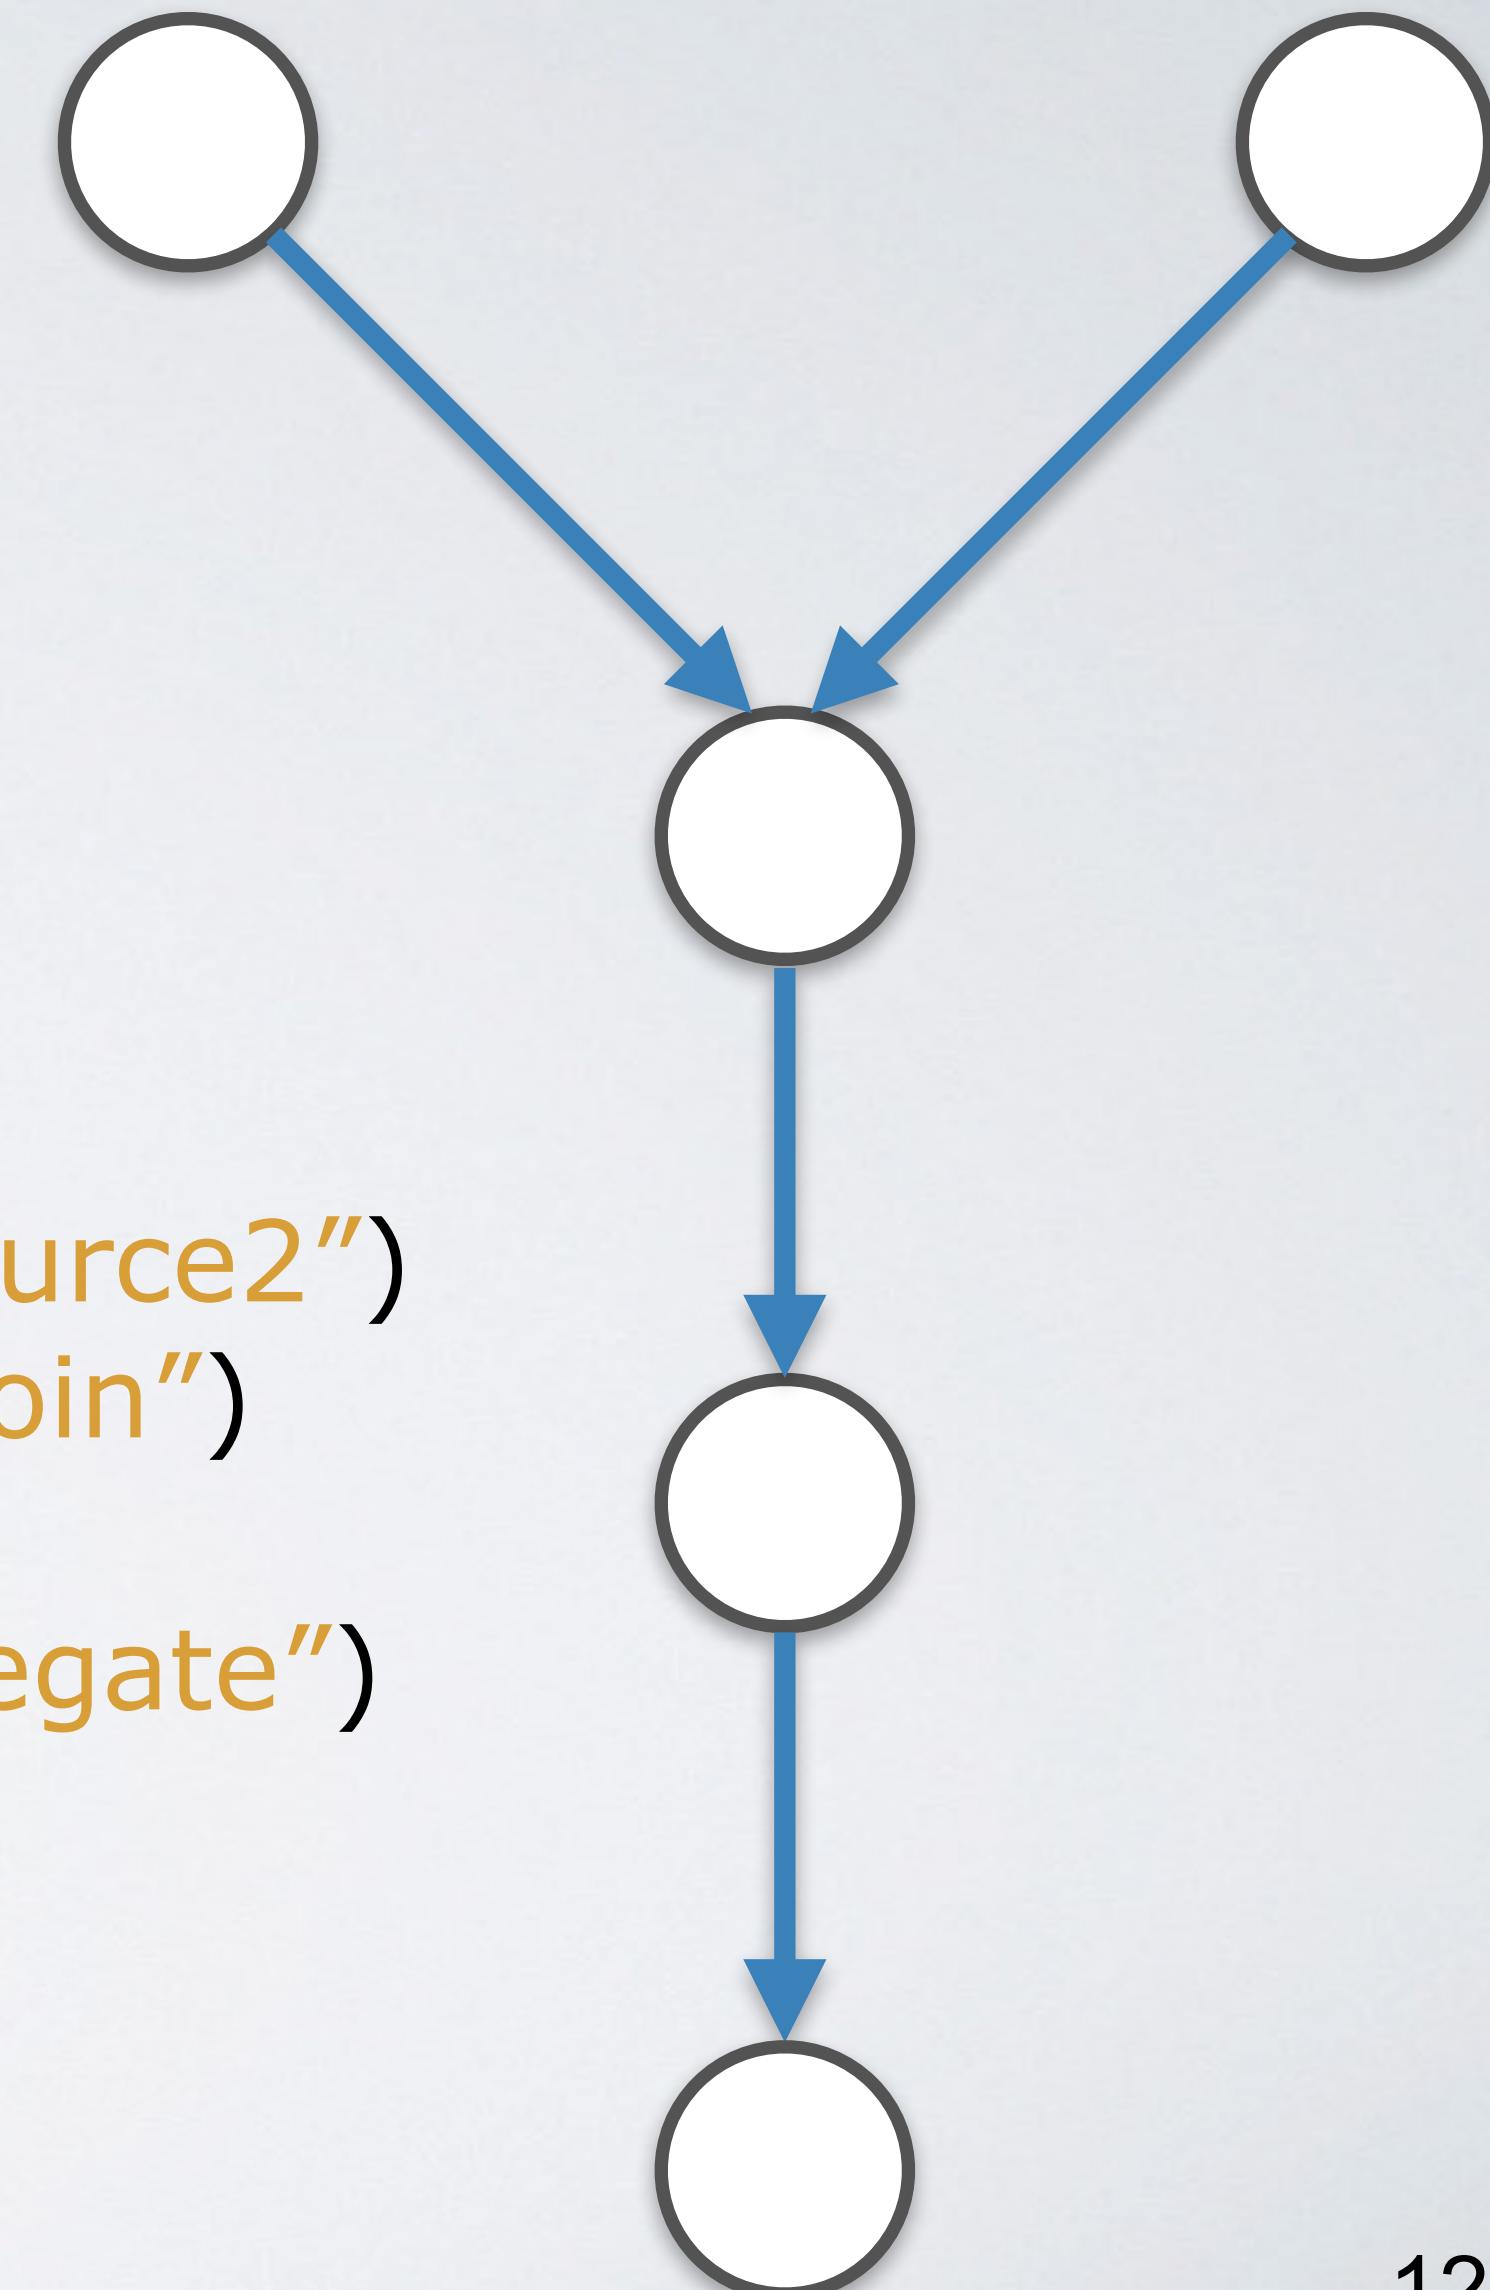
```
KStream<..> stream1 = builder.stream("topic1");  
builder.addSource("Source1", "topic1")
```

```
KStream<..> stream2 = builder.table("topic2");  
addSource("Source2", "topic2")
```

```
KStream<..> joined = stream1.leftJoin(stream2, ...);  
.addProcessor("Join", MyJoin::new, "Source1", "Source2")  
.addProcessor("Aggregate", MyAggregate::new, "Join")
```

```
KTable<..> aggregated = joined.aggregateByKey(...);  
.addStateStore(Stores.persistent().build(), "Aggregate")
```

```
aggregated.to("topic3");  
.addSink("Sink", "topic3", "Aggregate")
```



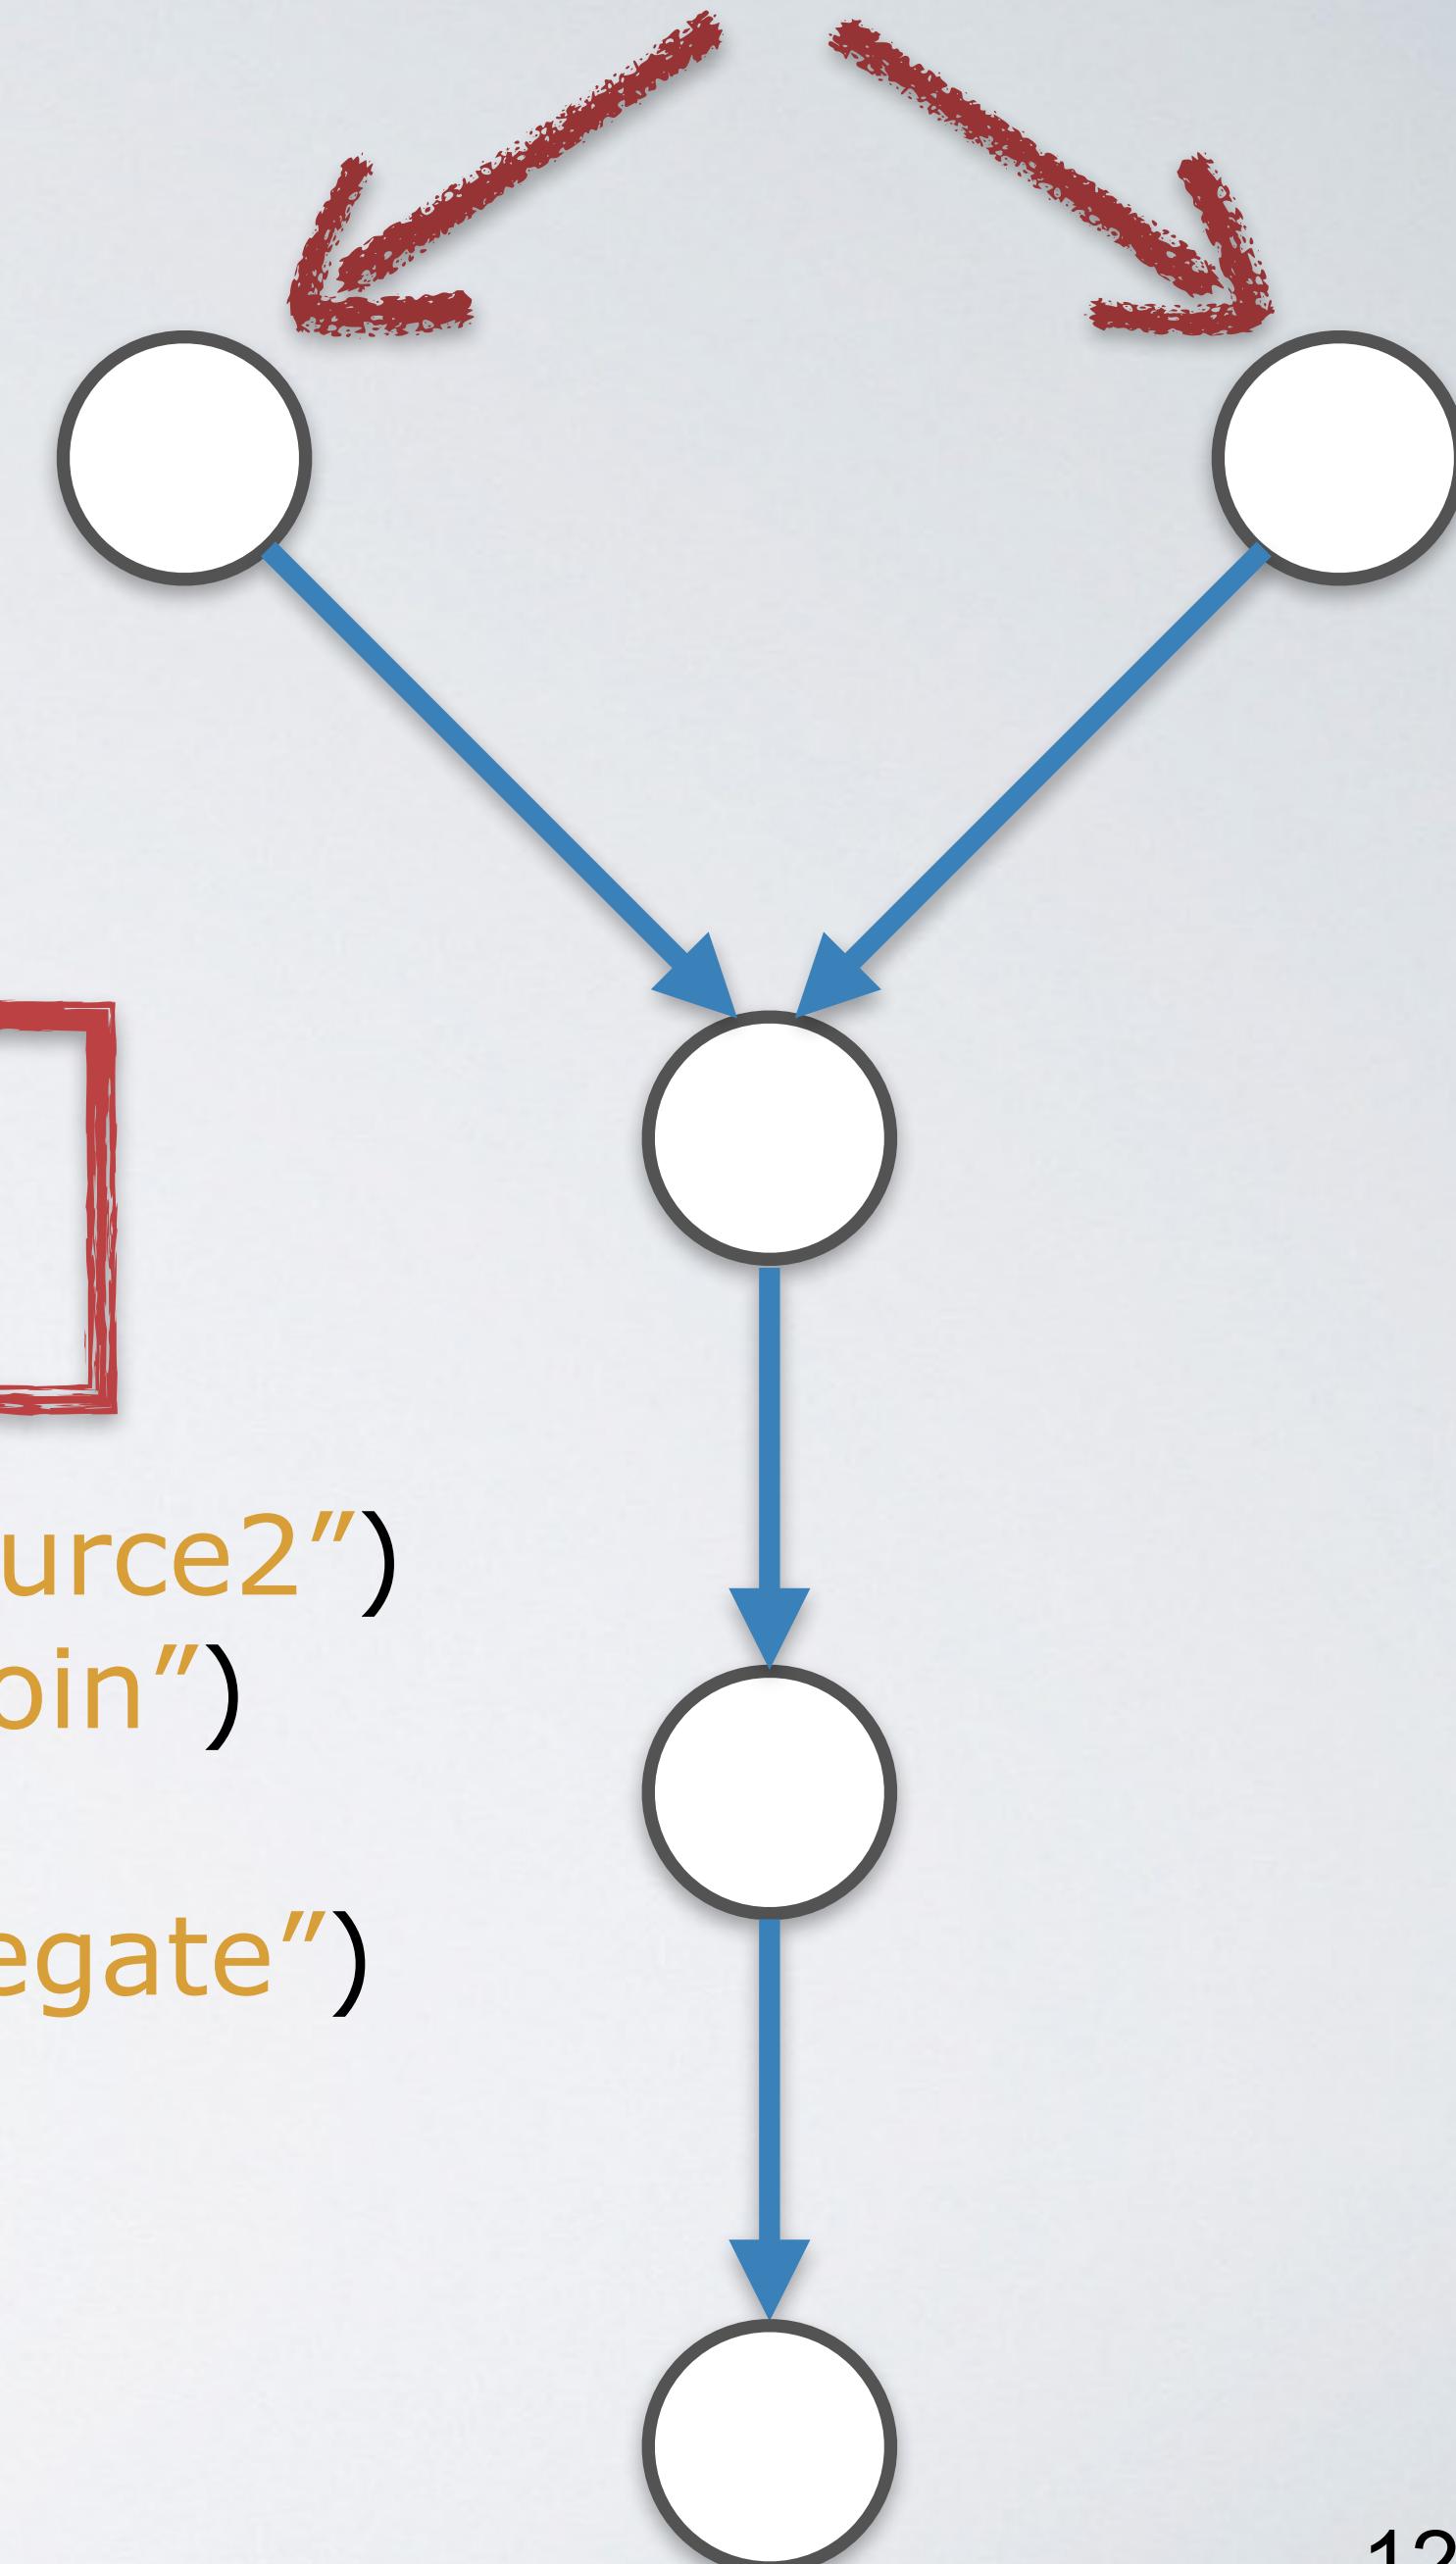
Processor Topology

```
builder.addSource("Source1", "topic1")
    .addSource("Source2", "topic2")

    .addProcessor("Join", MyJoin:new, "Source1", "Source2")
    .addProcessor("Aggregate", MyAggregate:new, "Join")

    .addStateStore(Stores.persistent().build(), "Aggregate")

    .addSink("Sink", "topic3", "Aggregate")
```



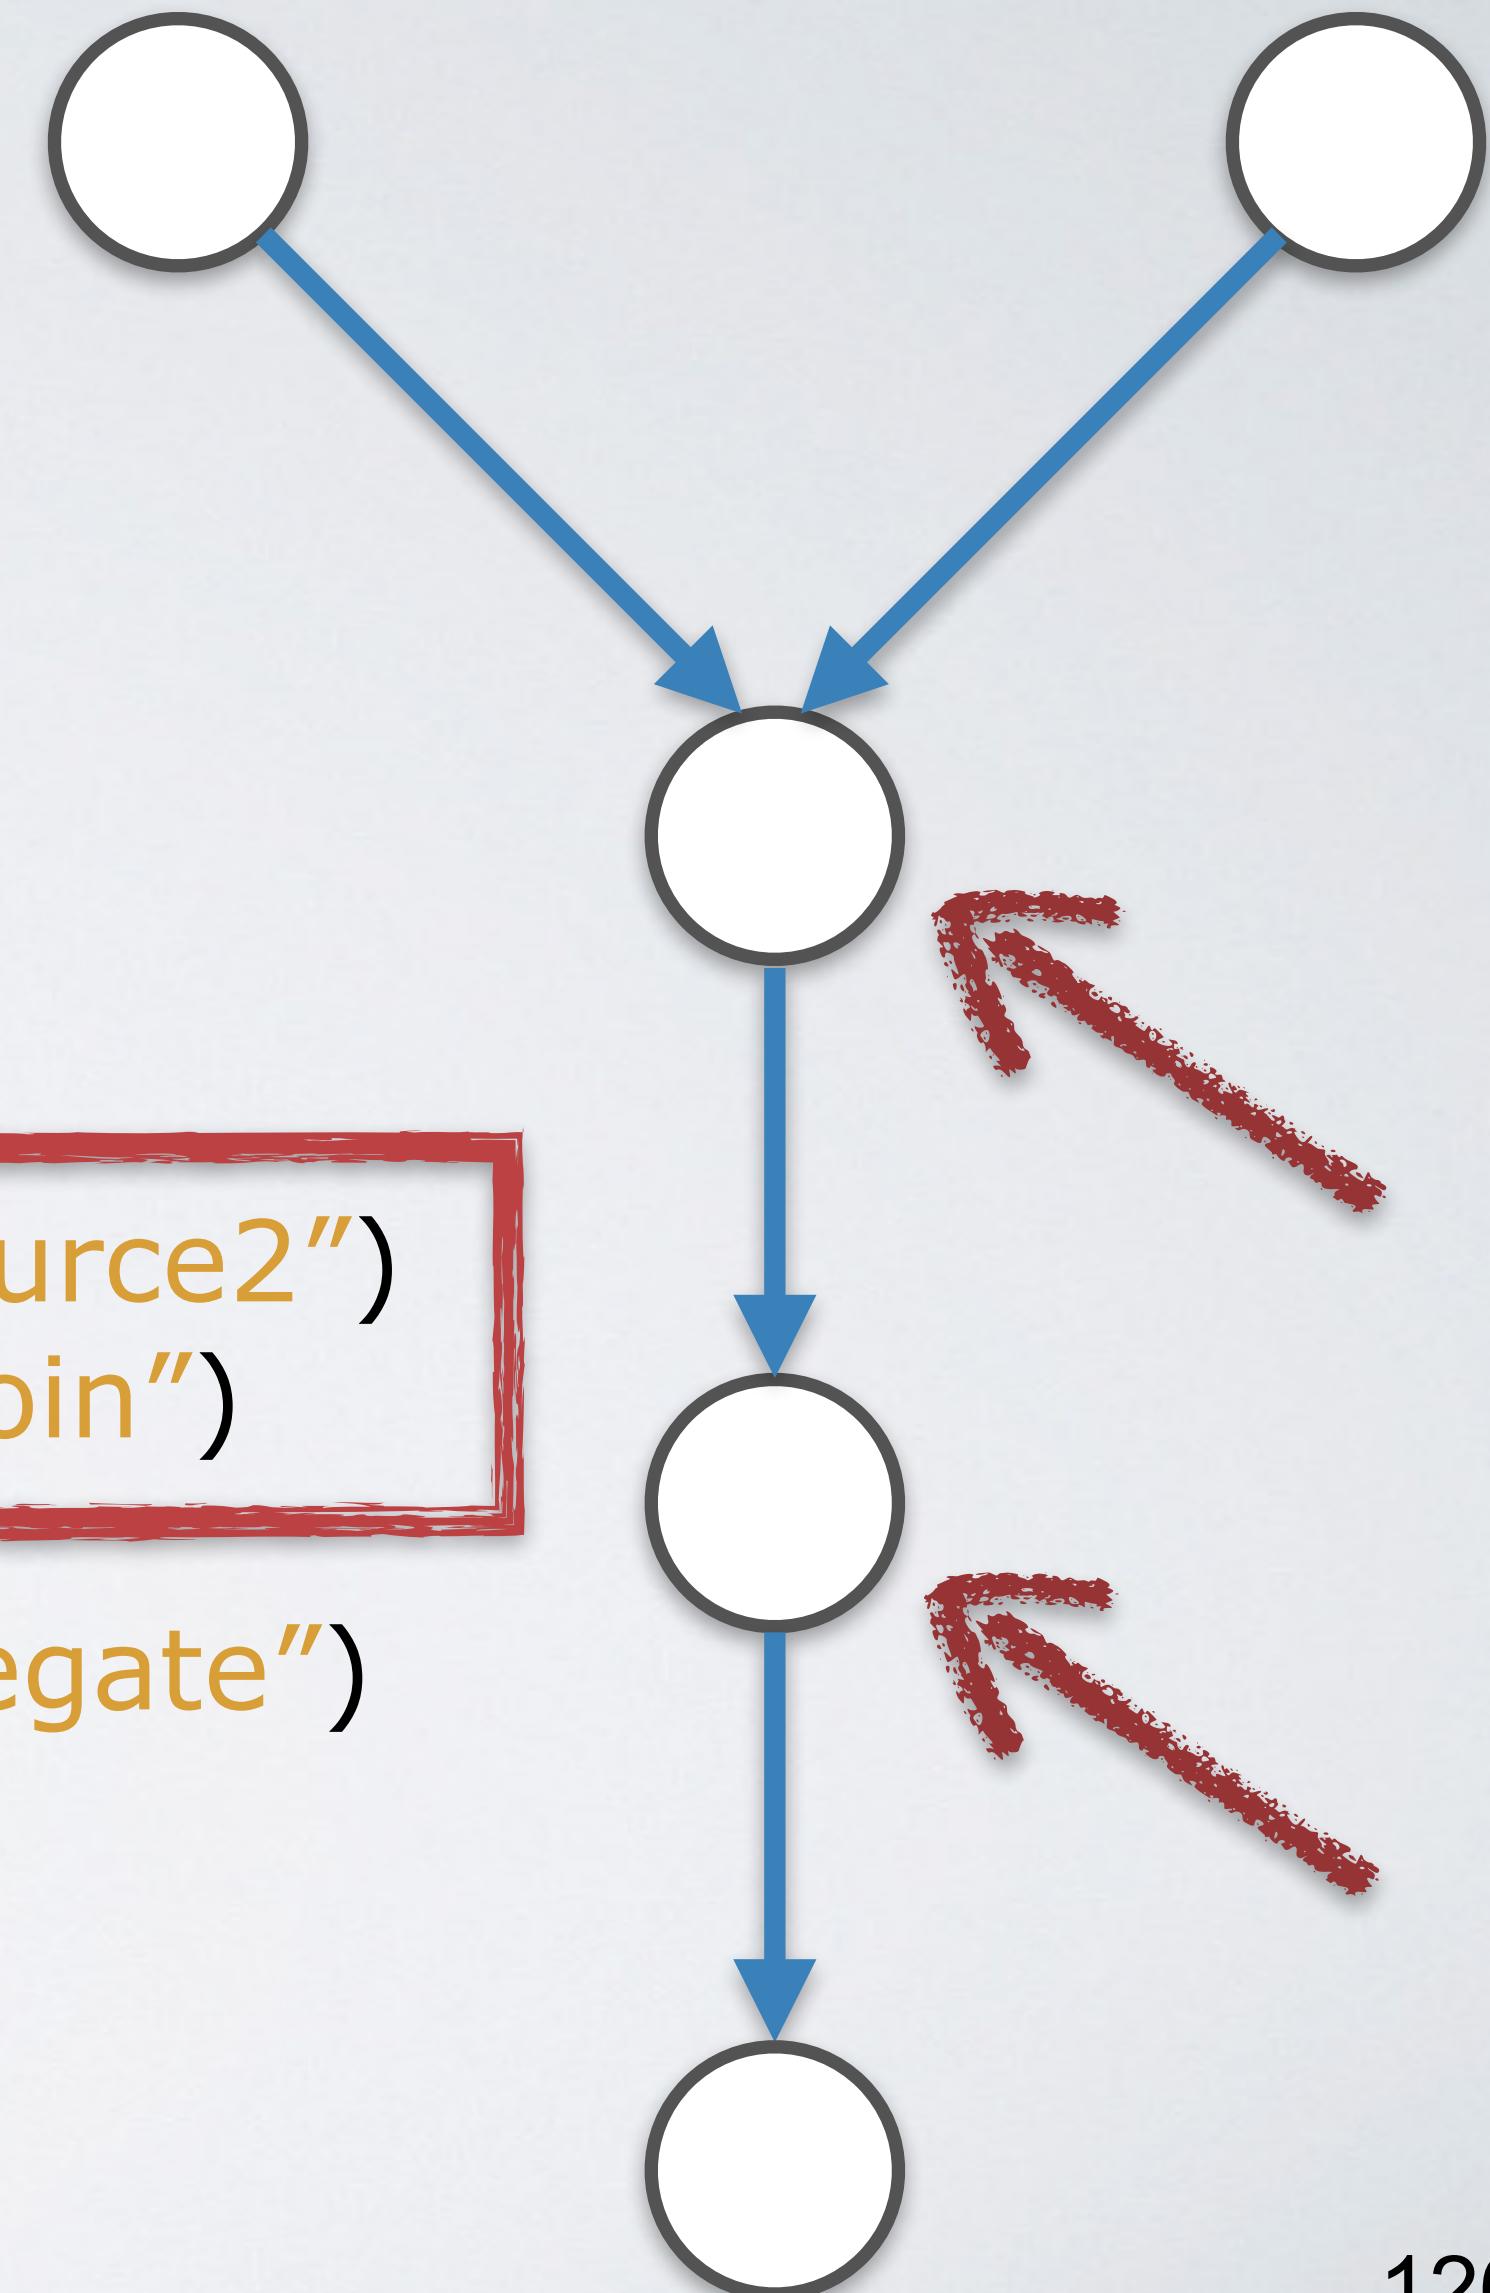
Processor Topology

```
builder.addSource("Source1", "topic1")
    .addSource("Source2", "topic2")

    .addProcessor("Join", MyJoin:new, "Source1", "Source2")
    .addProcessor("Aggregate", MyAggregate:new, "Join")

    .addStateStore(Stores.persistent().build(), "Aggregate")

    .addSink("Sink", "topic3", "Aggregate")
```



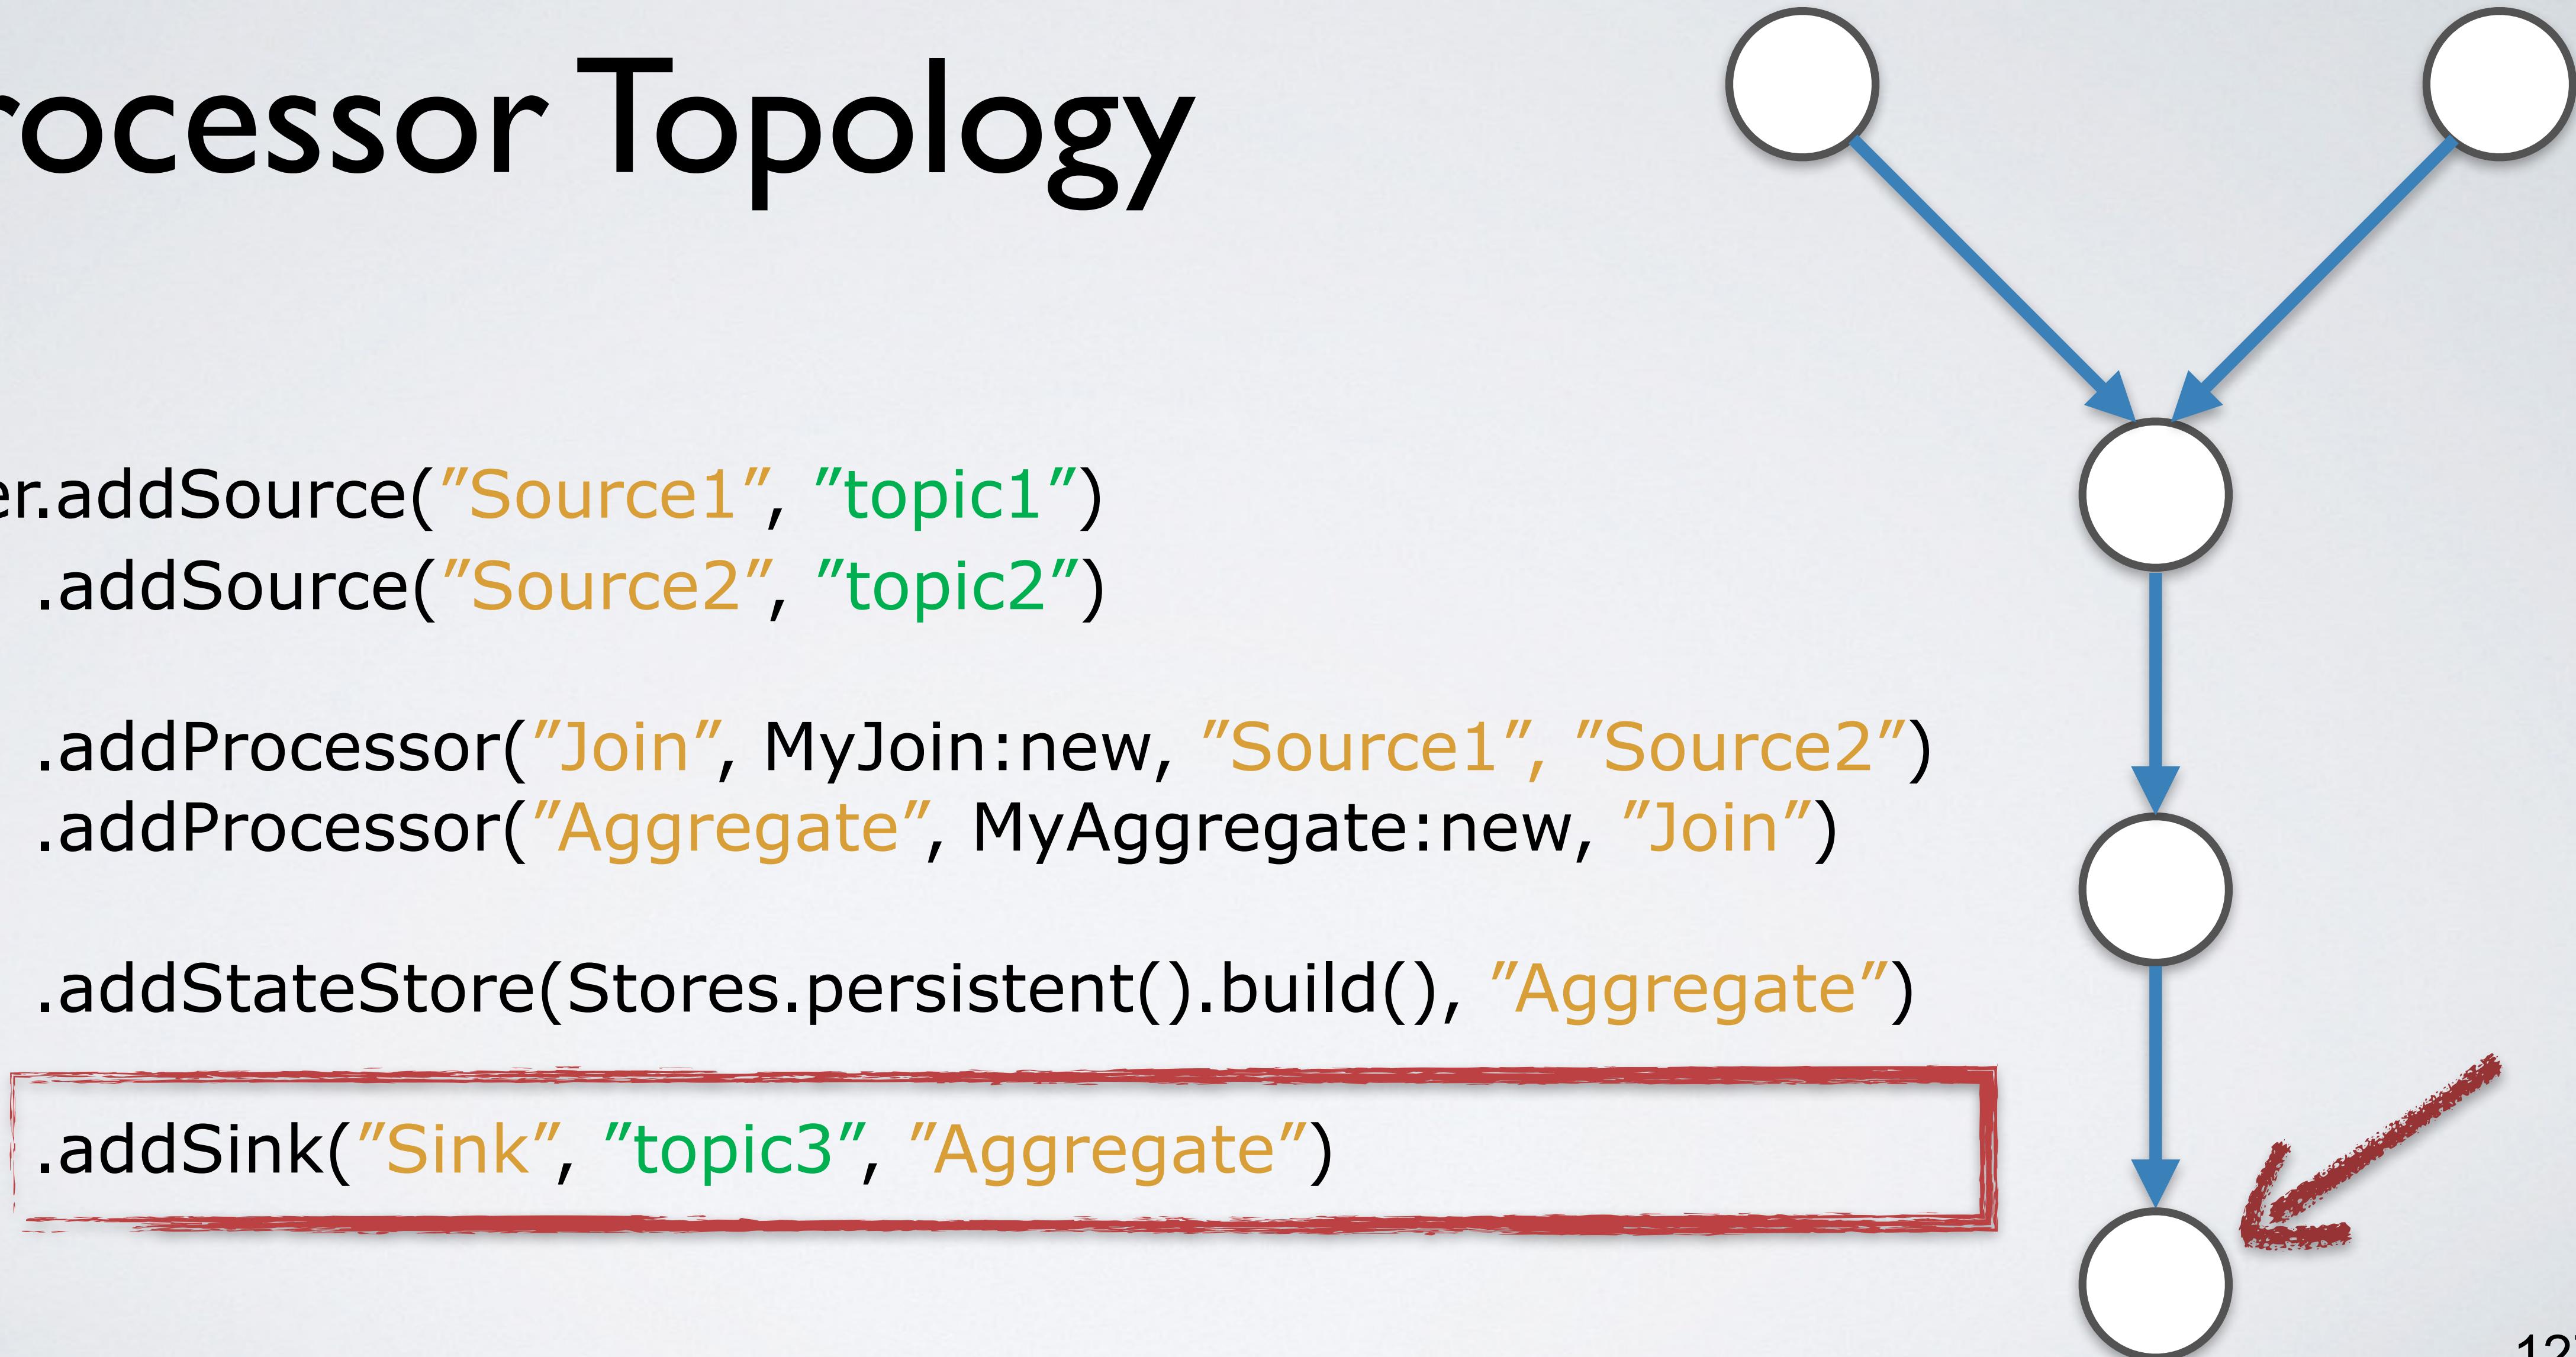
Processor Topology

```
builder.addSource("Source1", "topic1")
    .addSource("Source2", "topic2")

    .addProcessor("Join", MyJoin:new, "Source1", "Source2")
    .addProcessor("Aggregate", MyAggregate:new, "Join")

    .addStateStore(Stores.persistent().build(), "Aggregate")

    .addSink("Sink", "topic3", "Aggregate")
```



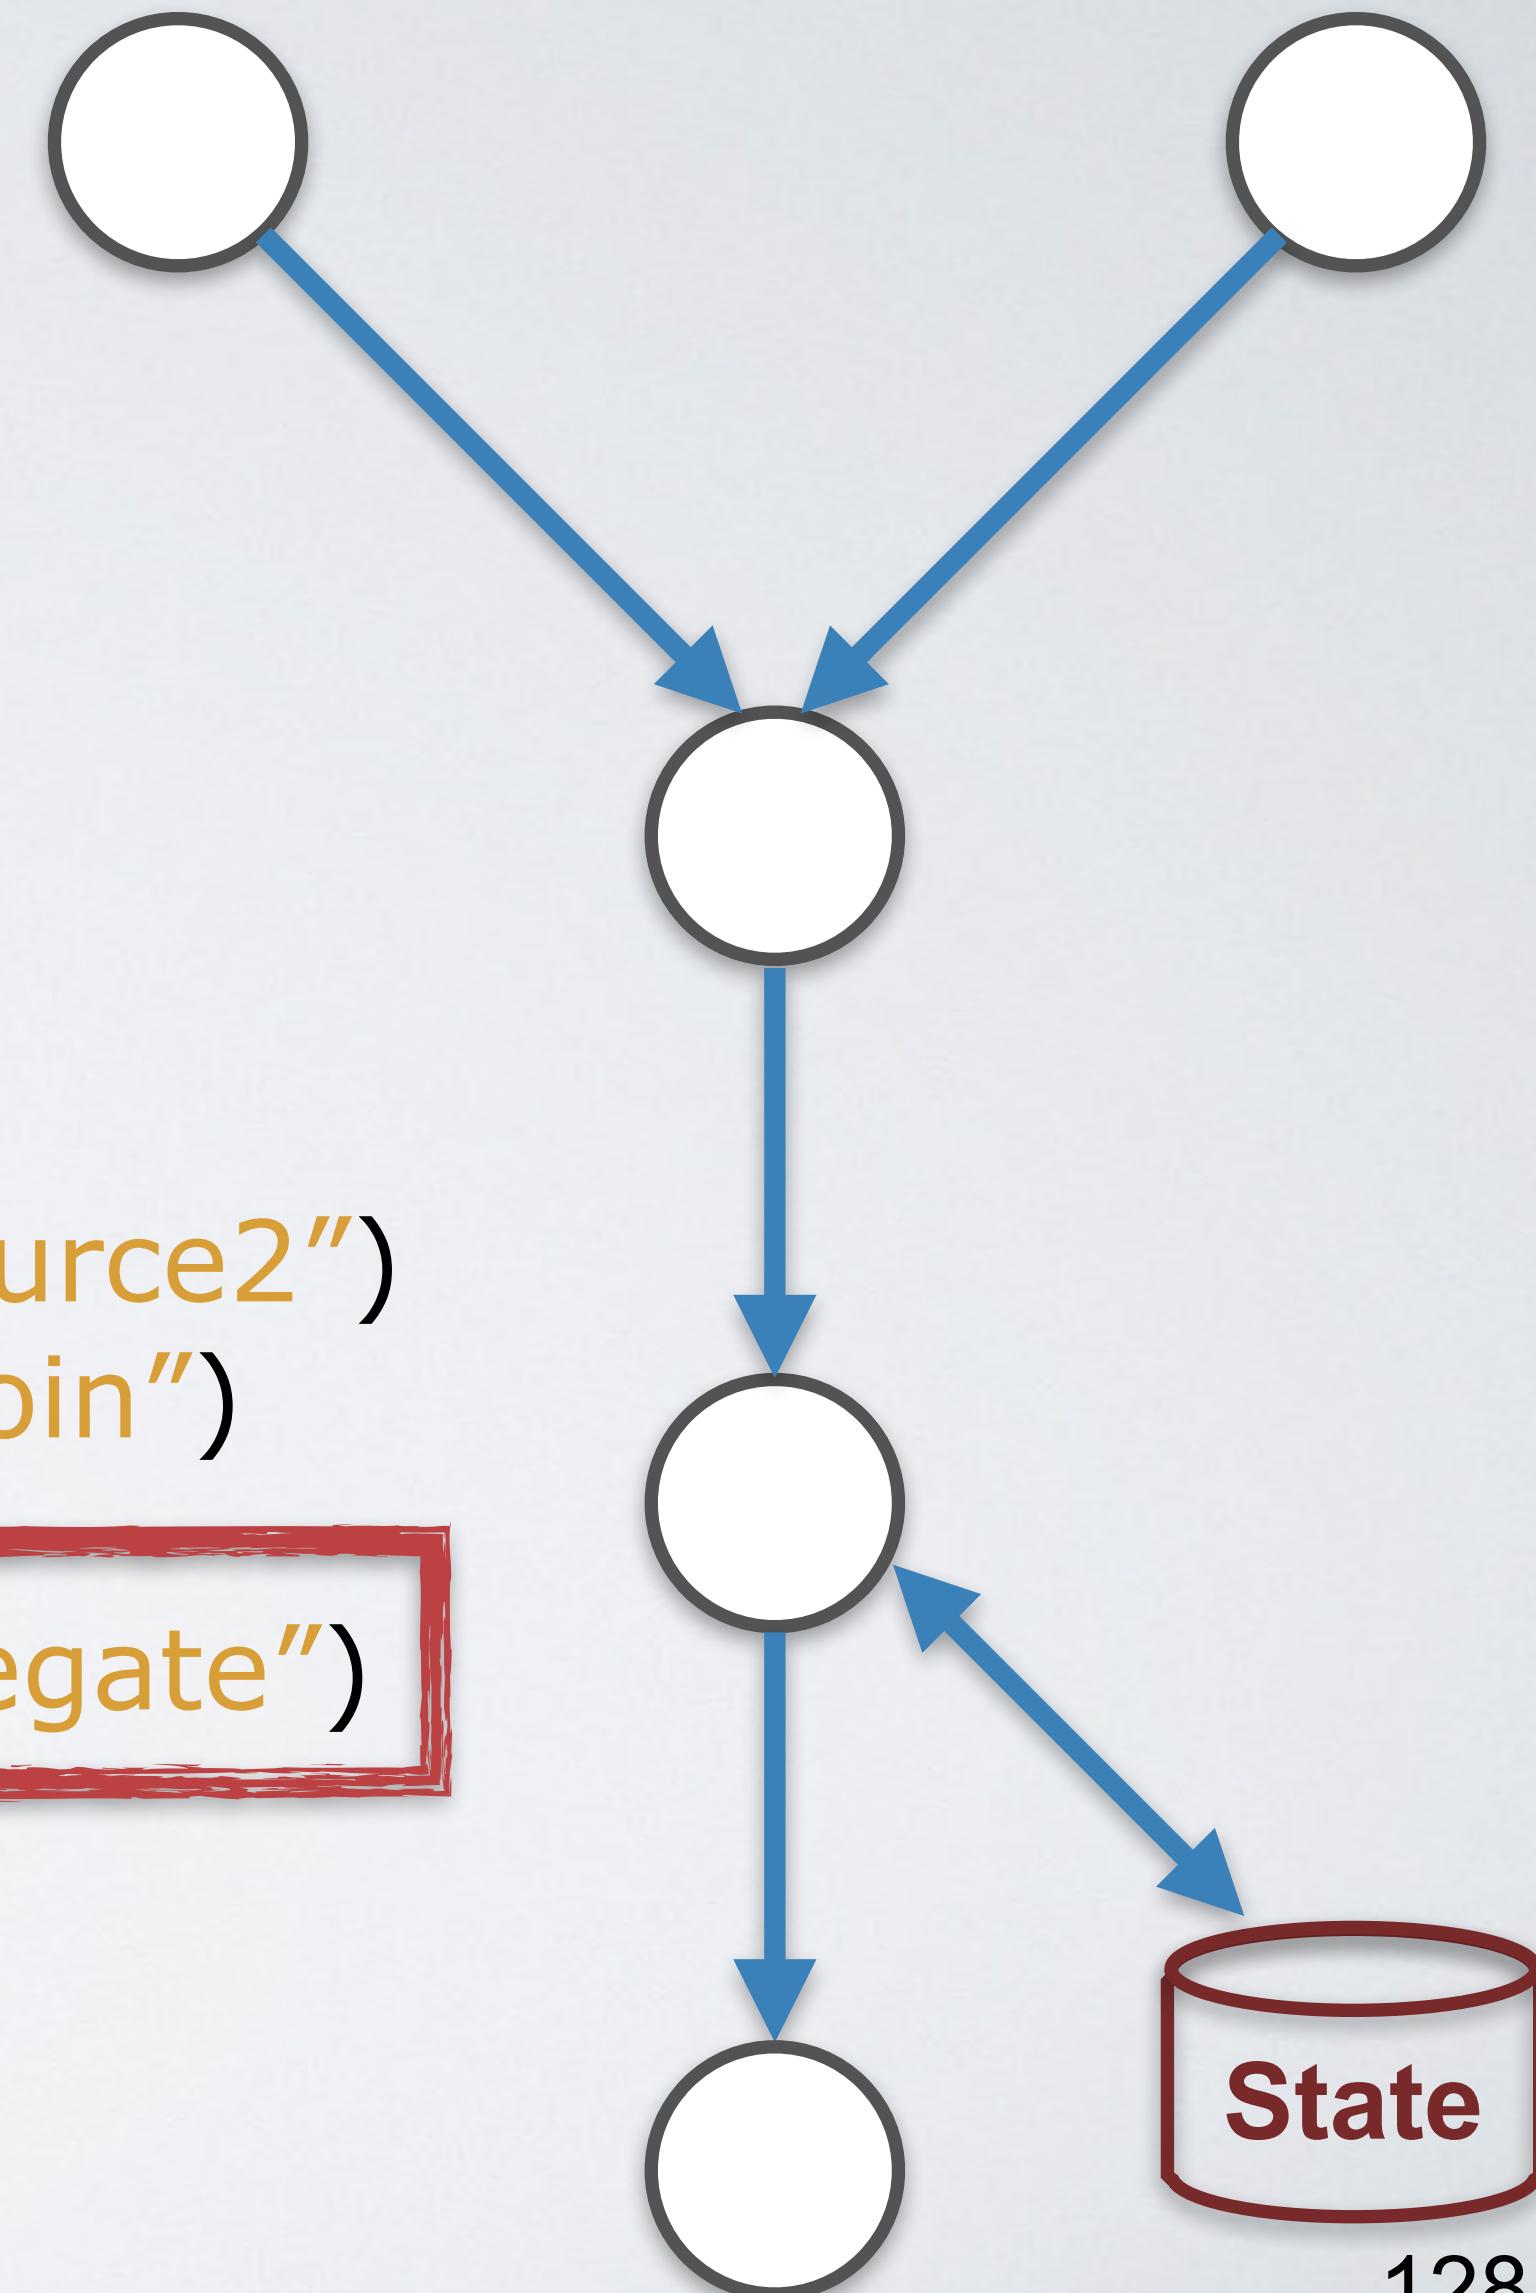
States in Stream Processing

```
builder.addSource("Source1", "topic1")
    .addSource("Source2", "topic2")

    .addProcessor("Join", MyJoin:new, "Source1", "Source2")
    .addProcessor("Aggregate", MyAggregate:new, "Join")

    .addStateStore(Stores.persistent().build(), "Aggregate")

    .addSink("Sink", "topic3", "Aggregate")
```



Processor Topology

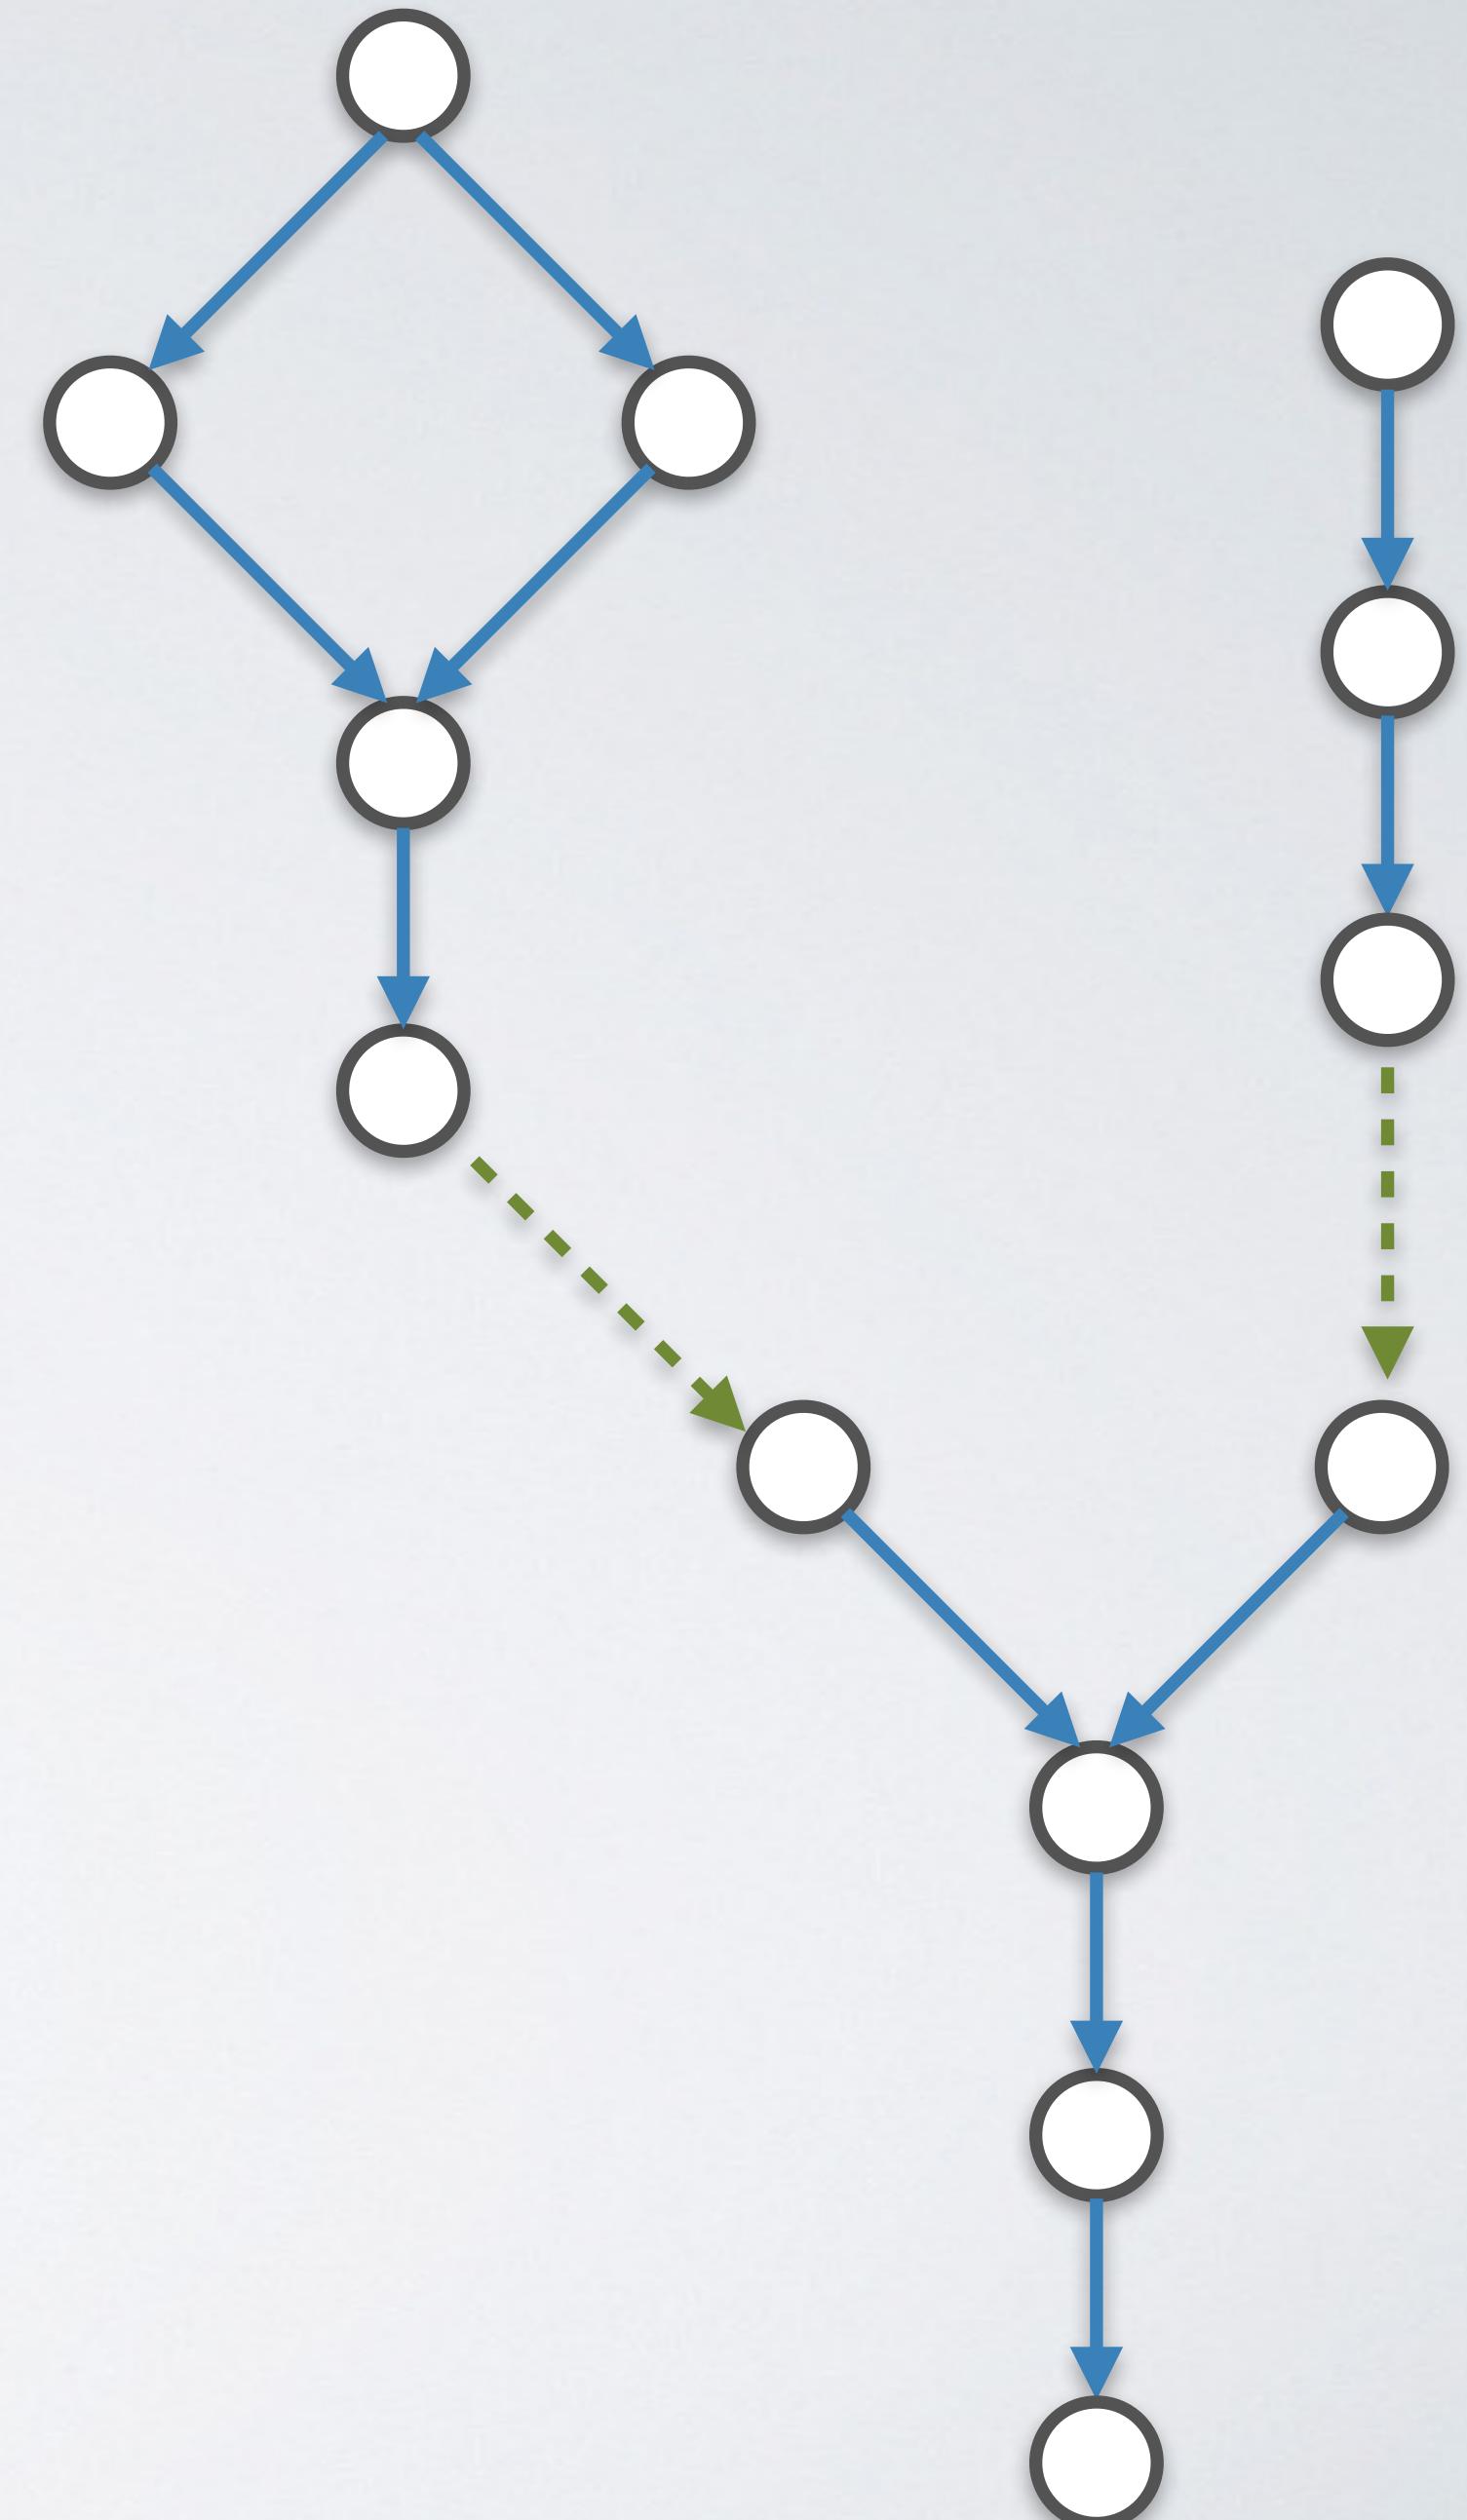
...

```
sink1.to("topic1");
```

```
source1 = builder.table("topic1");
```

```
source2 = sink1.through("topic2");
```

...



Processor Topology

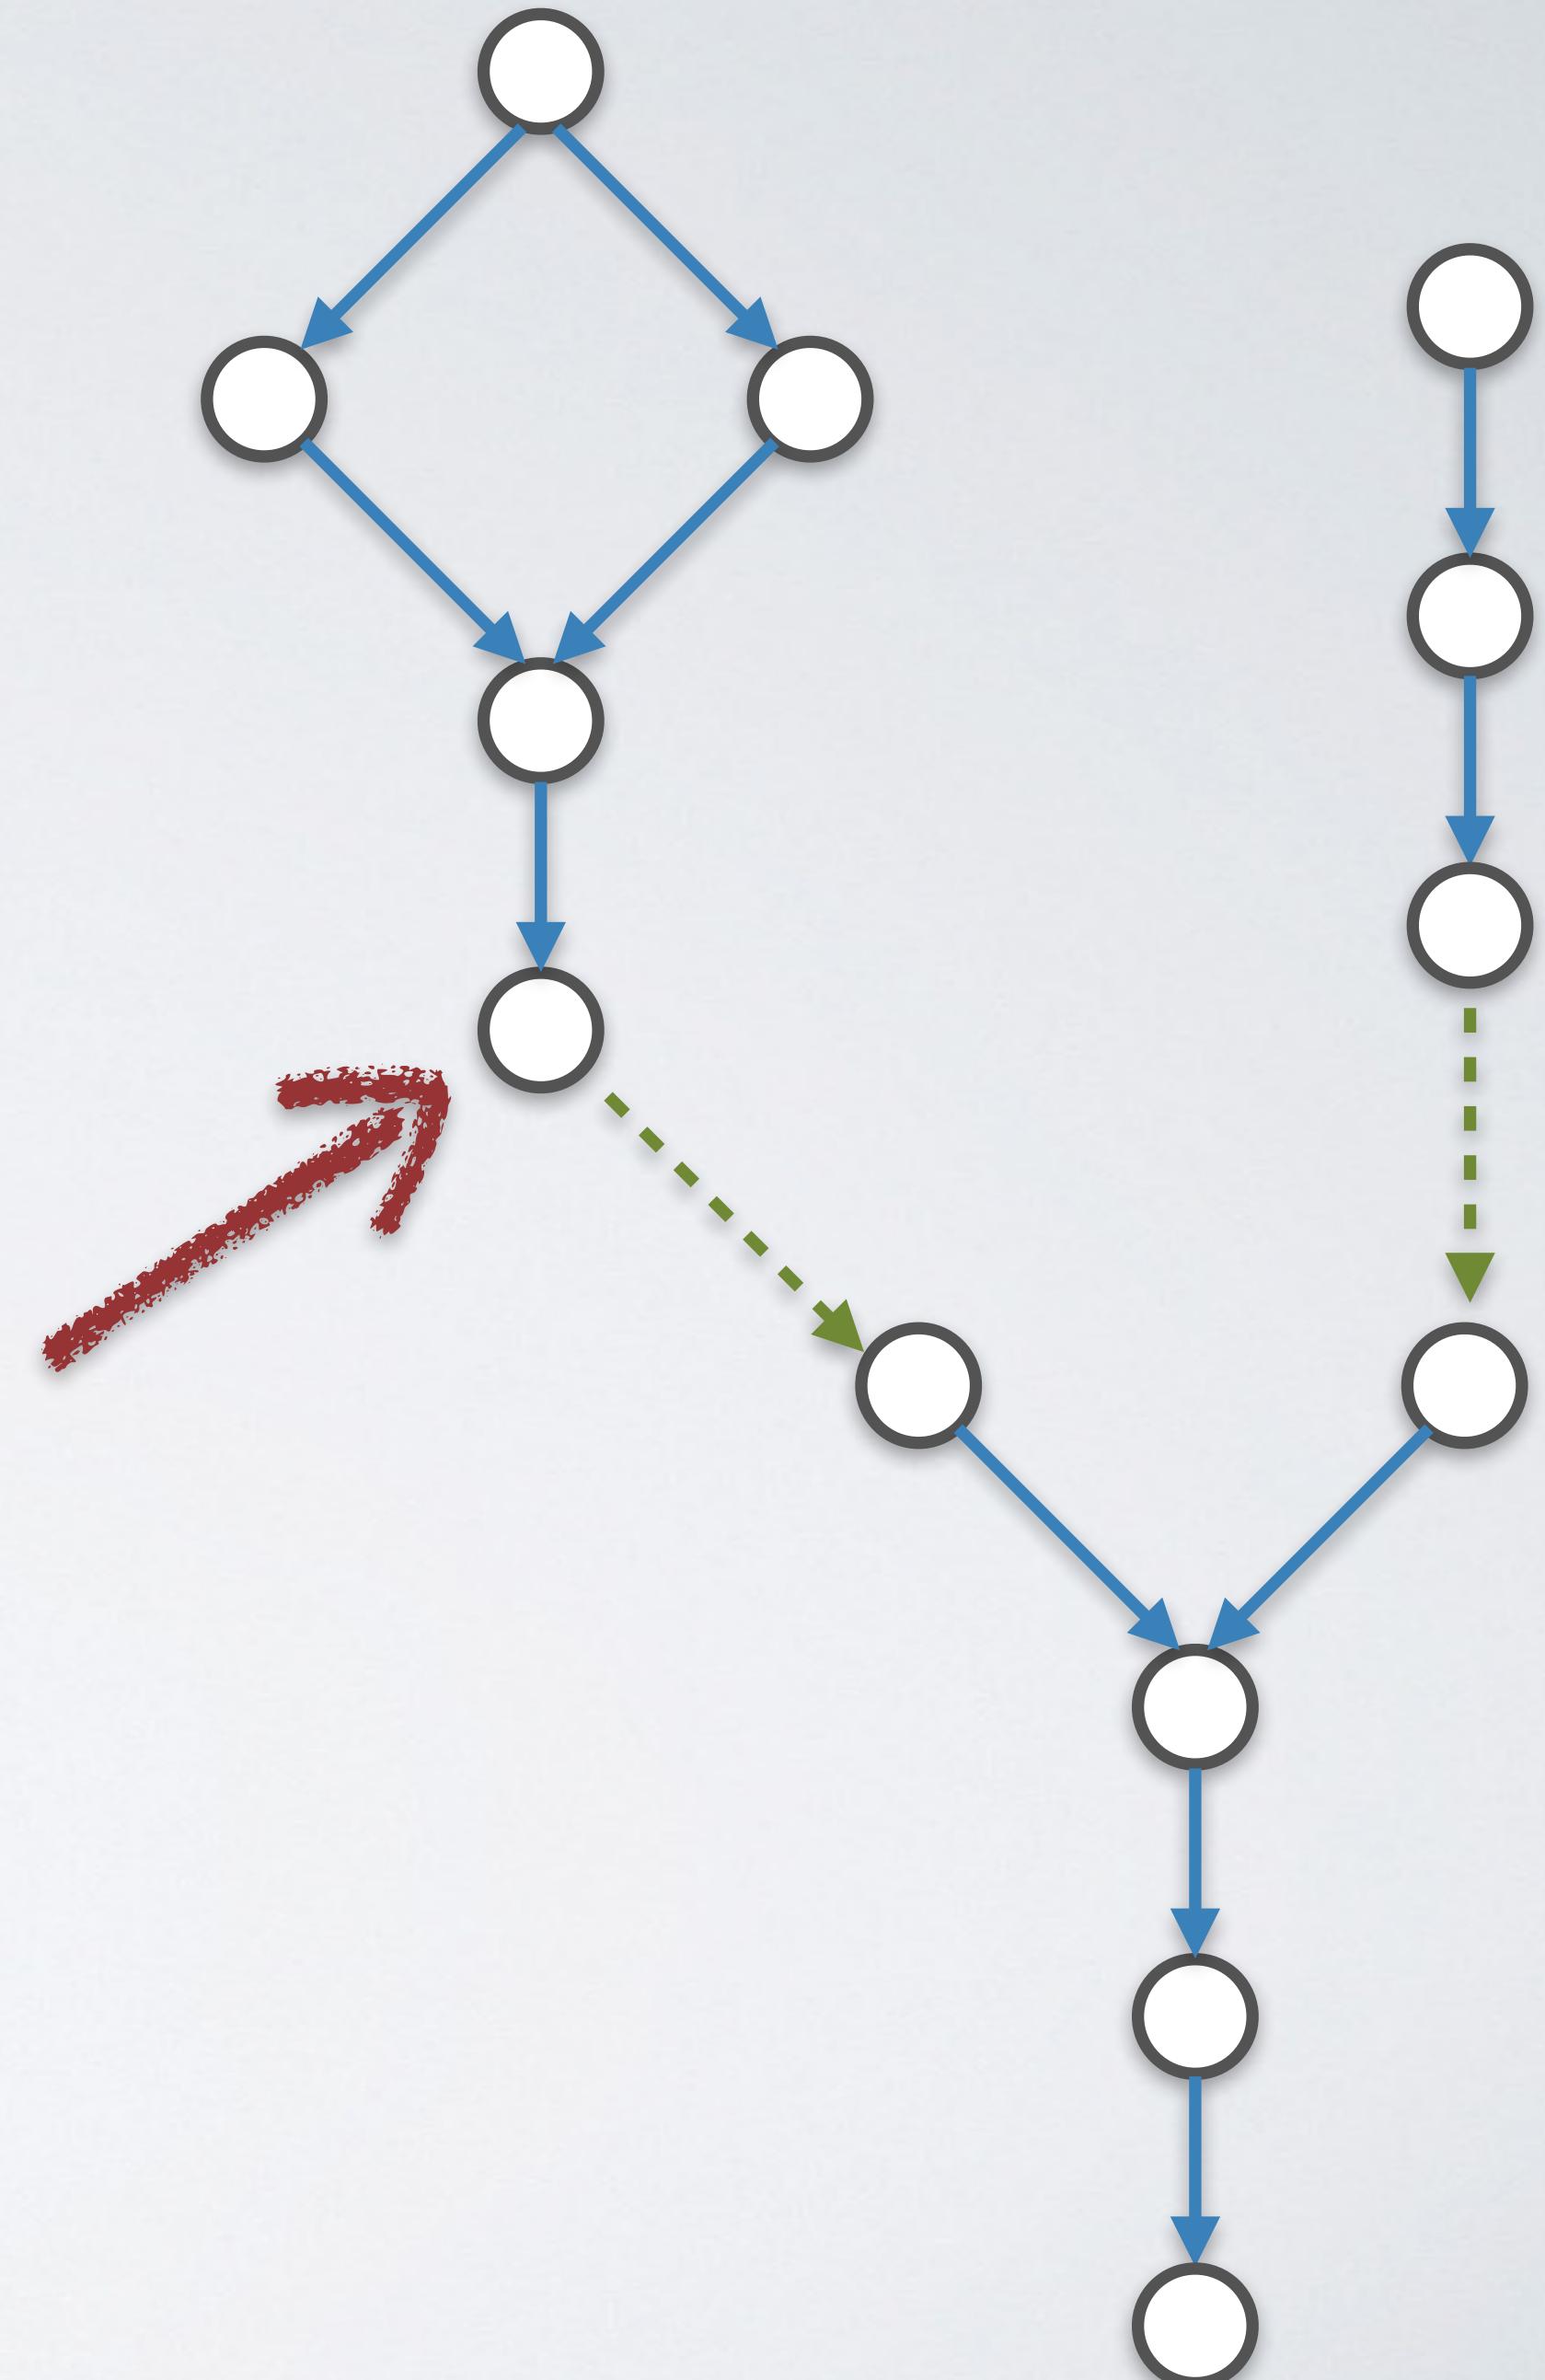
...

```
sink1.to("topic1");
```

```
source1 = builder.table("topic1");
```

```
source2 = sink1.through("topic2");
```

...



Processor Topology

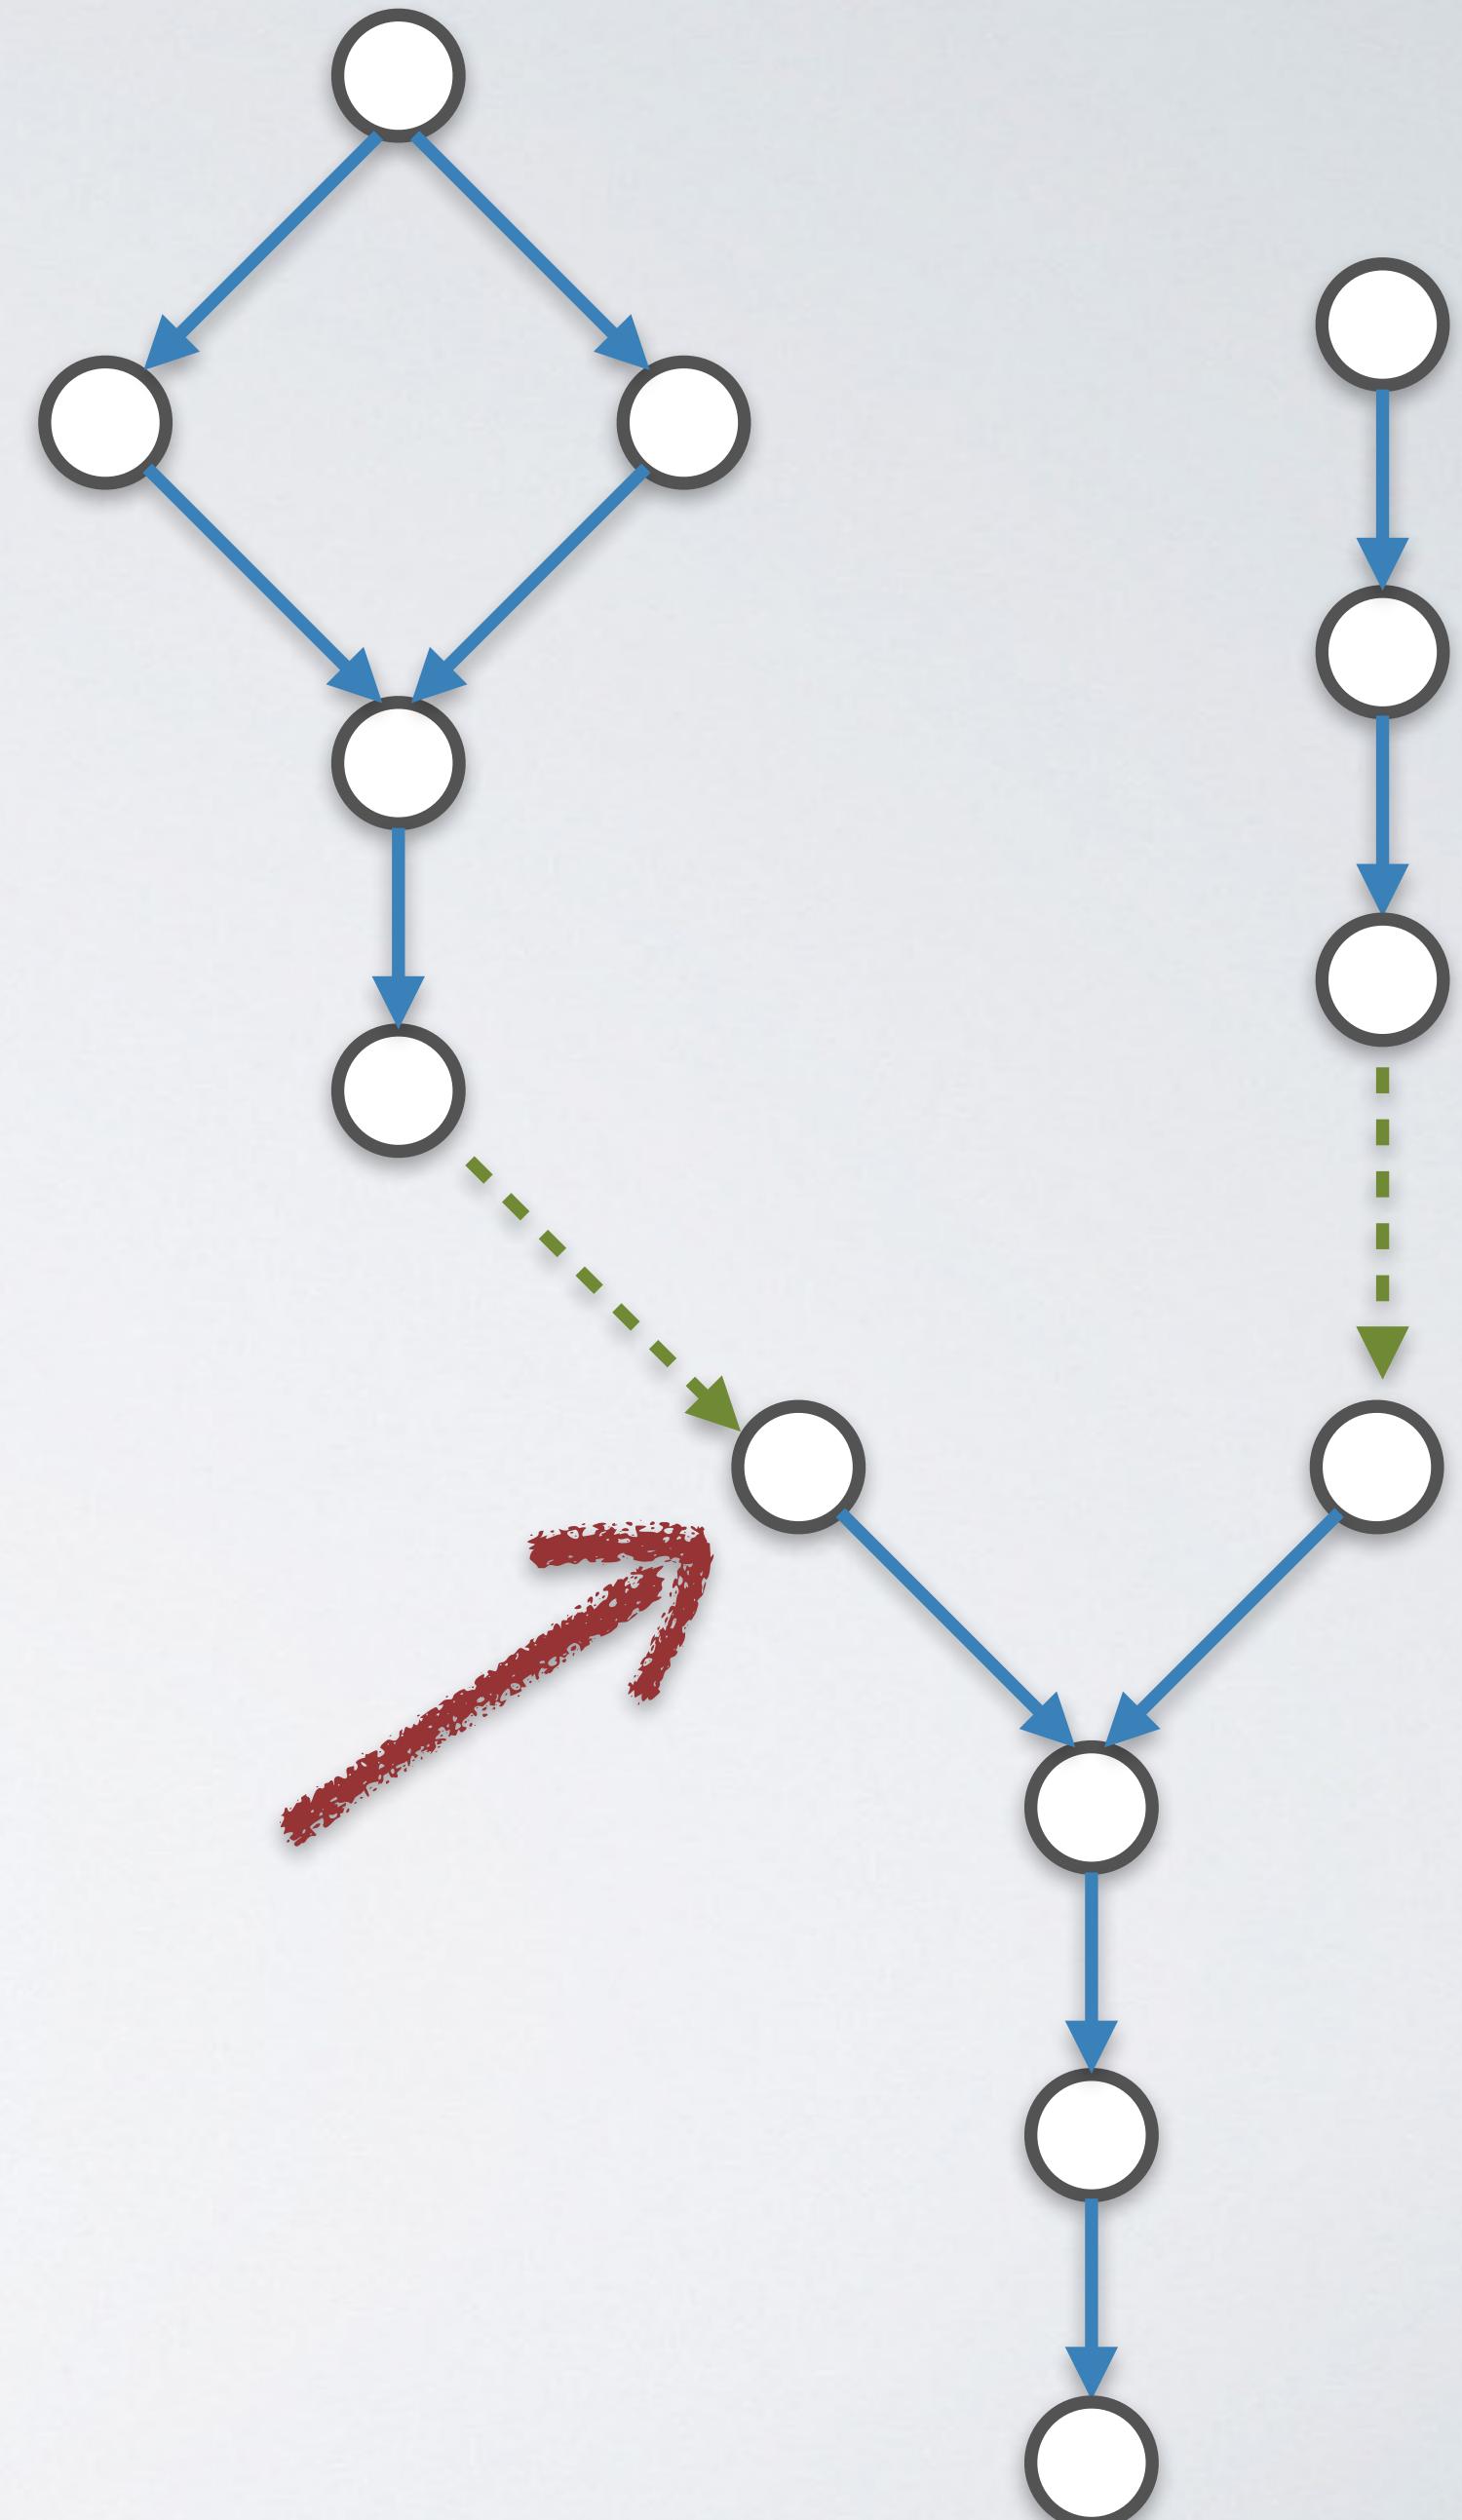
...

```
sink1.to("topic1");
```

```
source1 = builder.table("topic1");
```

```
source2 = sink1.through("topic2");
```

...



Processor Topology

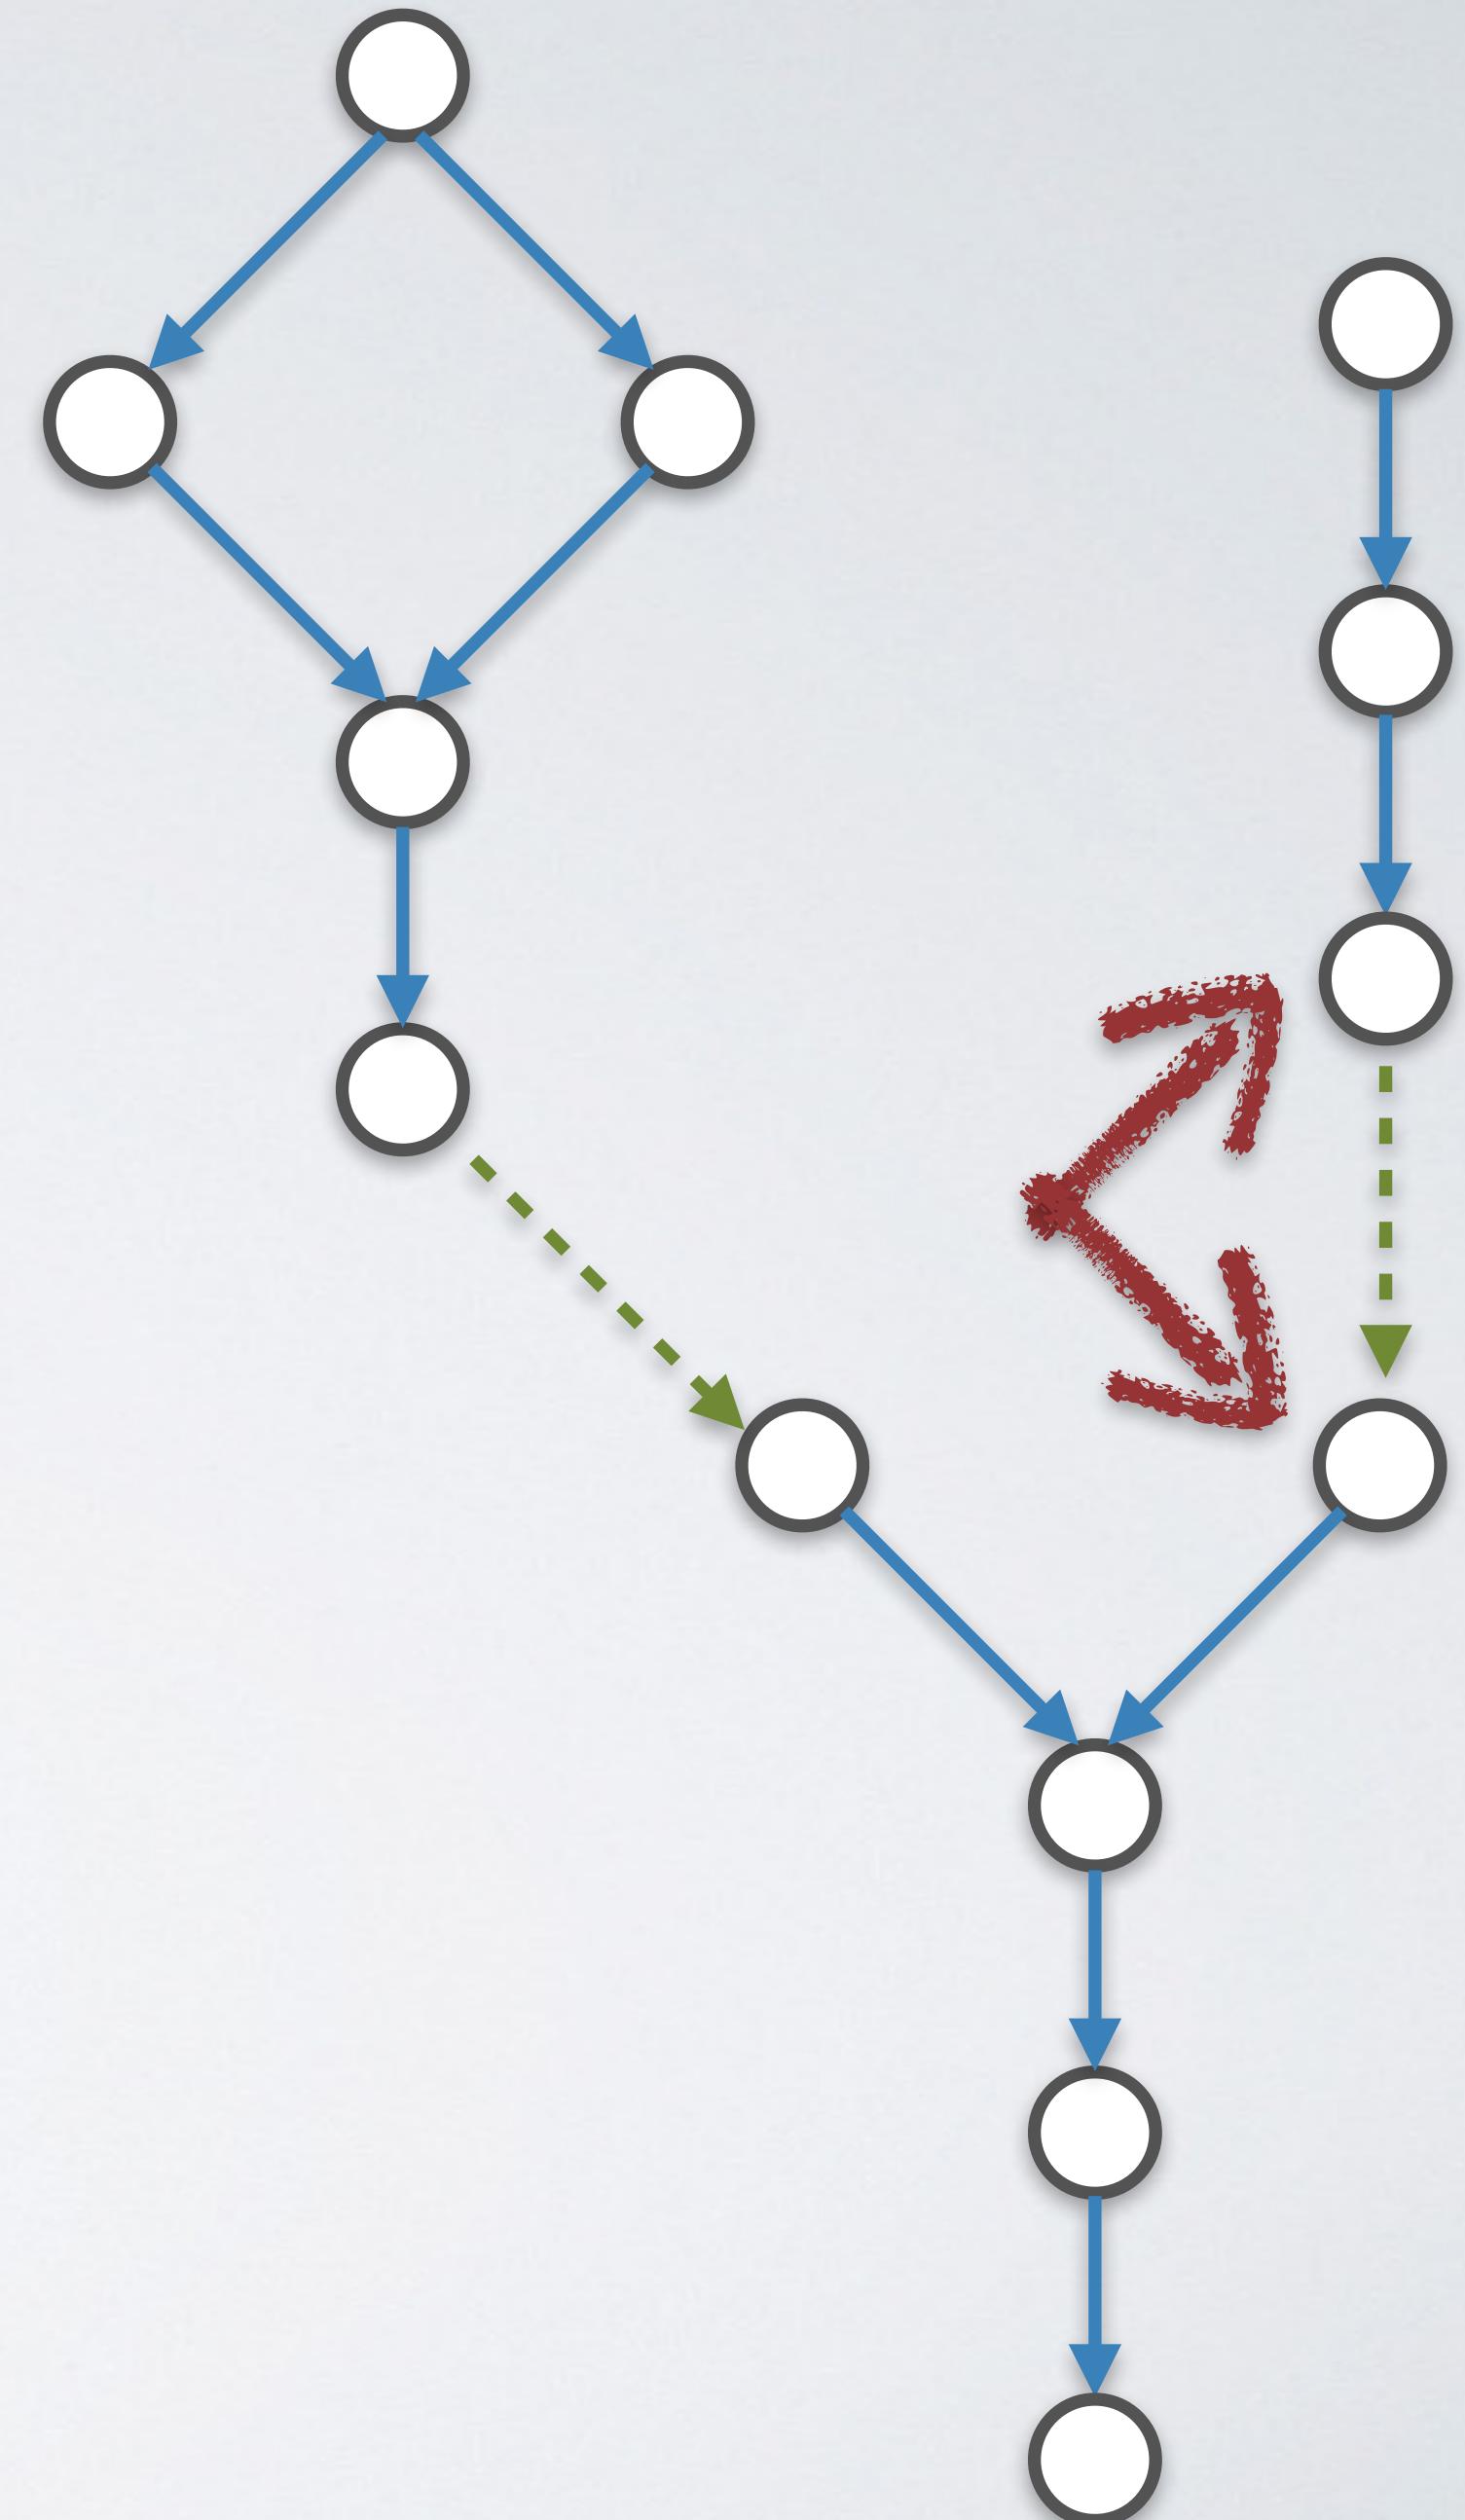
...

```
sink1.to("topic1");
```

```
source1 = builder.table("topic1");
```

```
source2 = sink1.through("topic2");
```

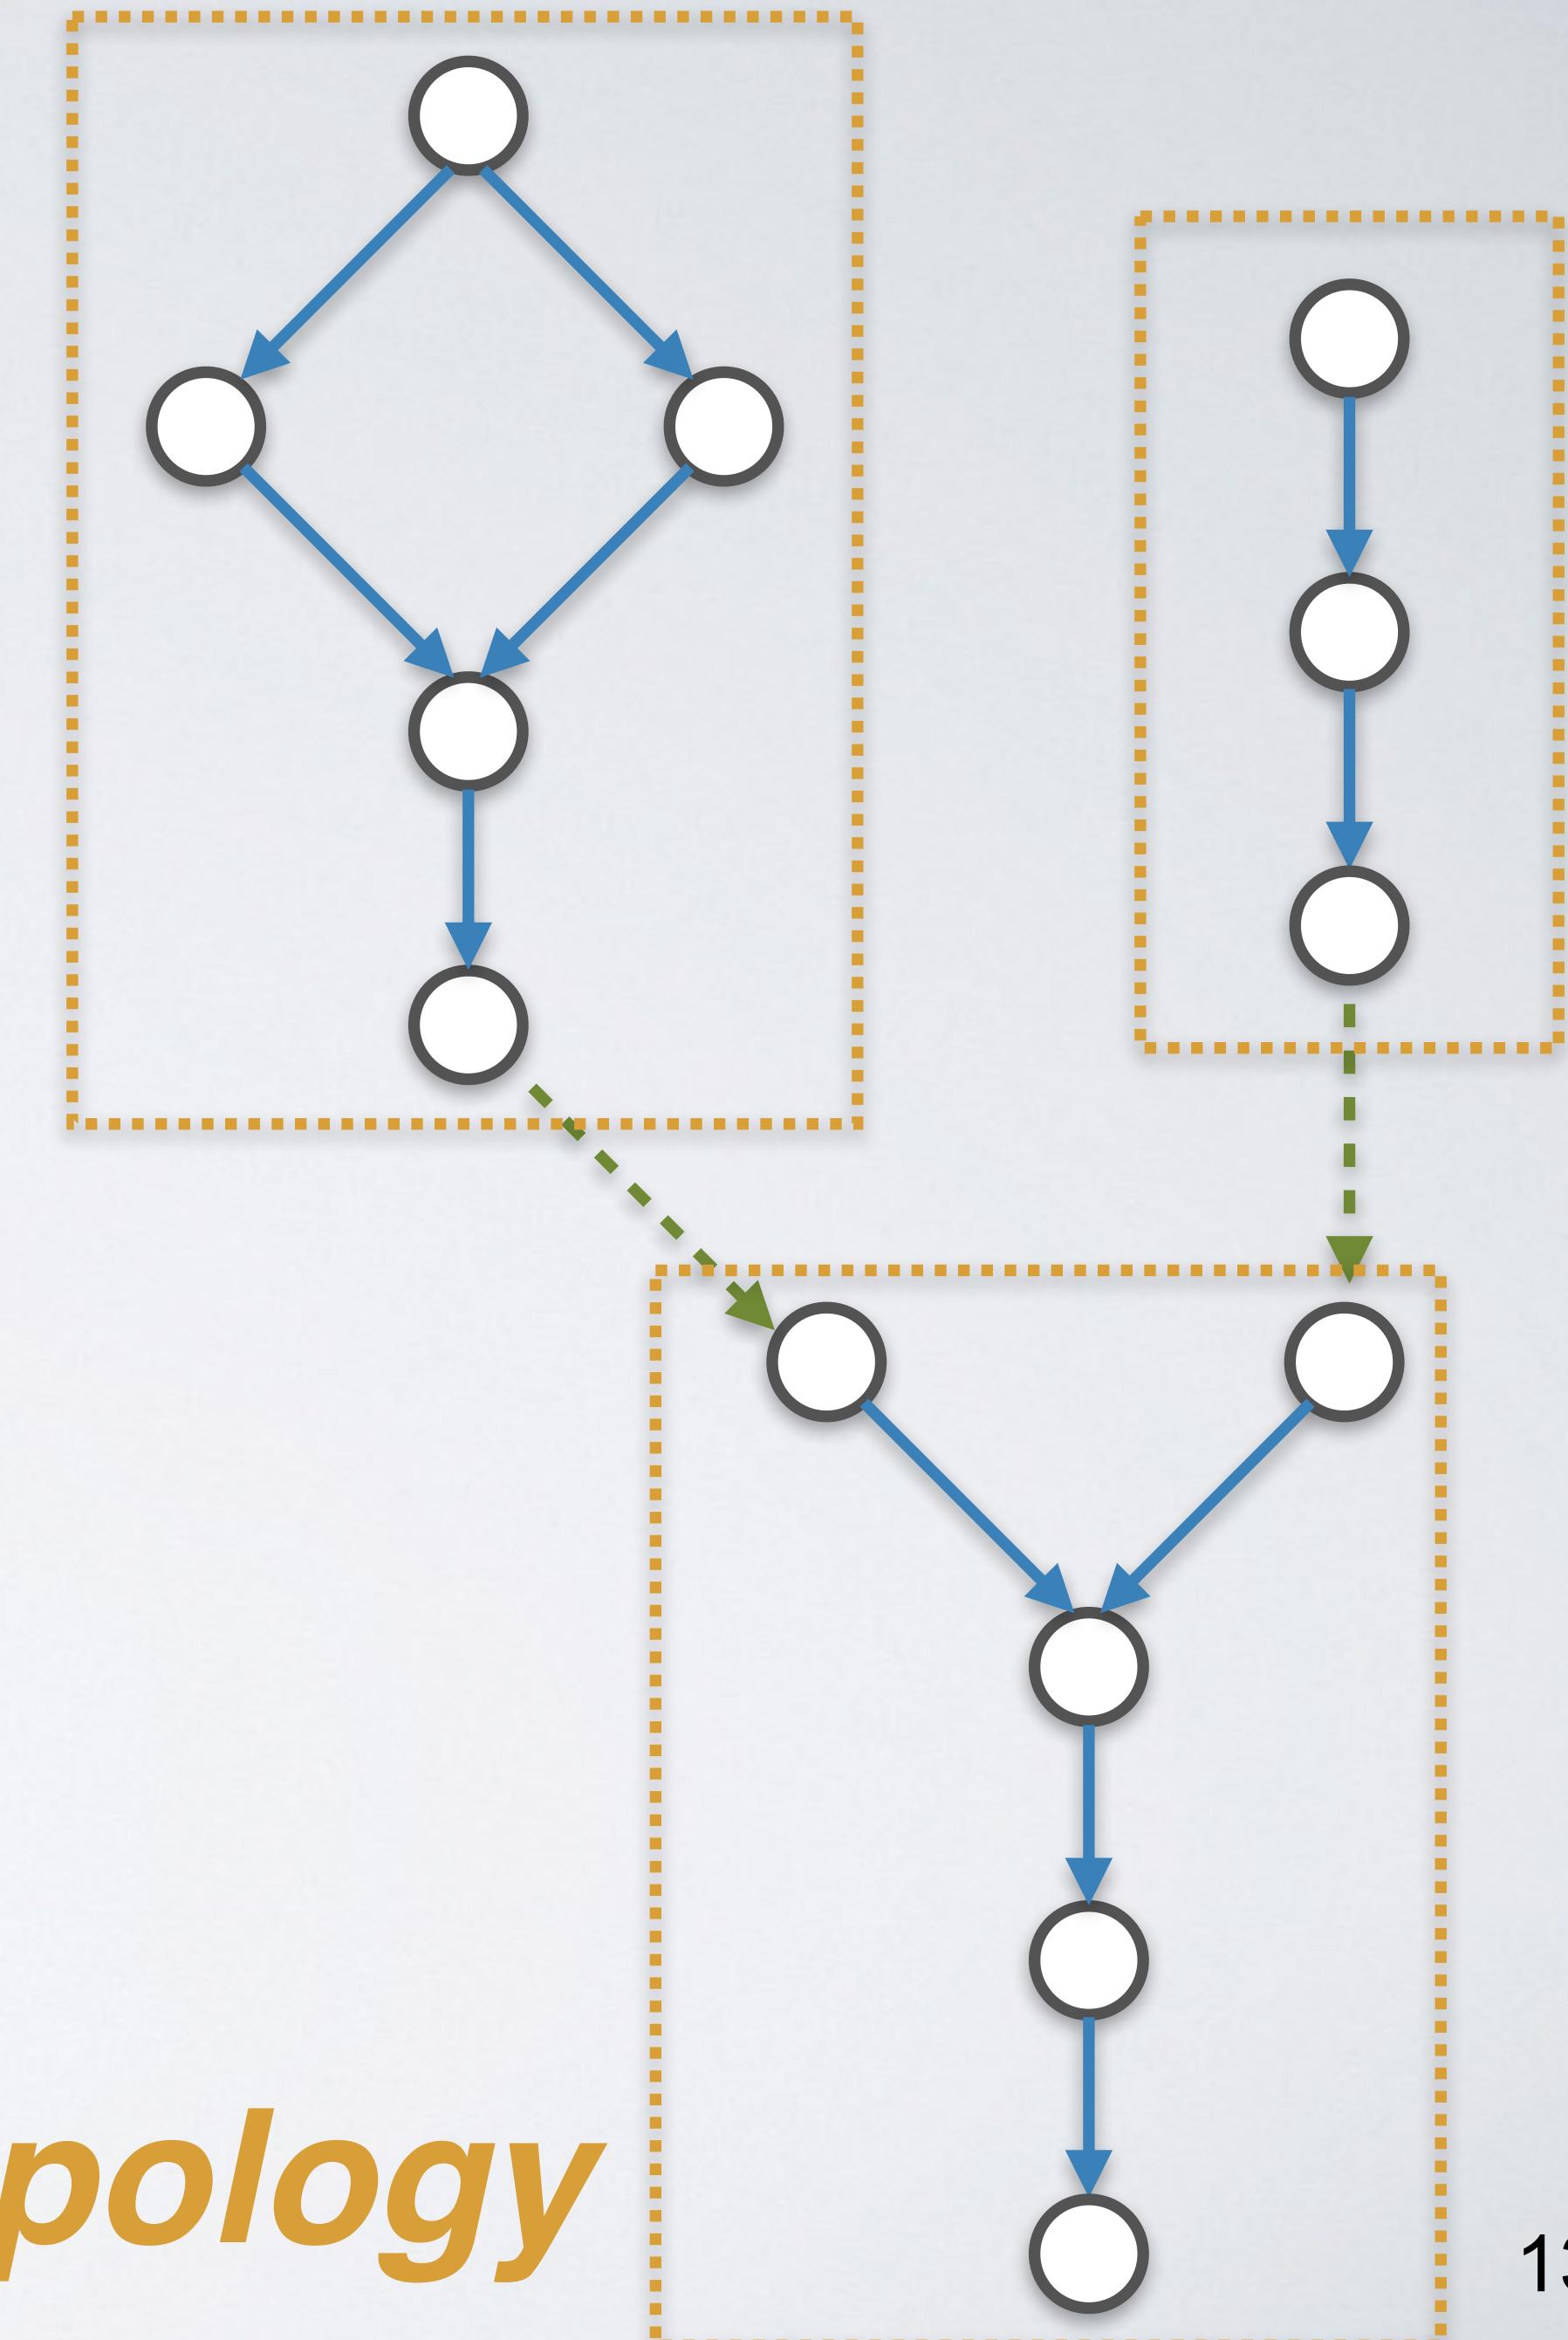
...



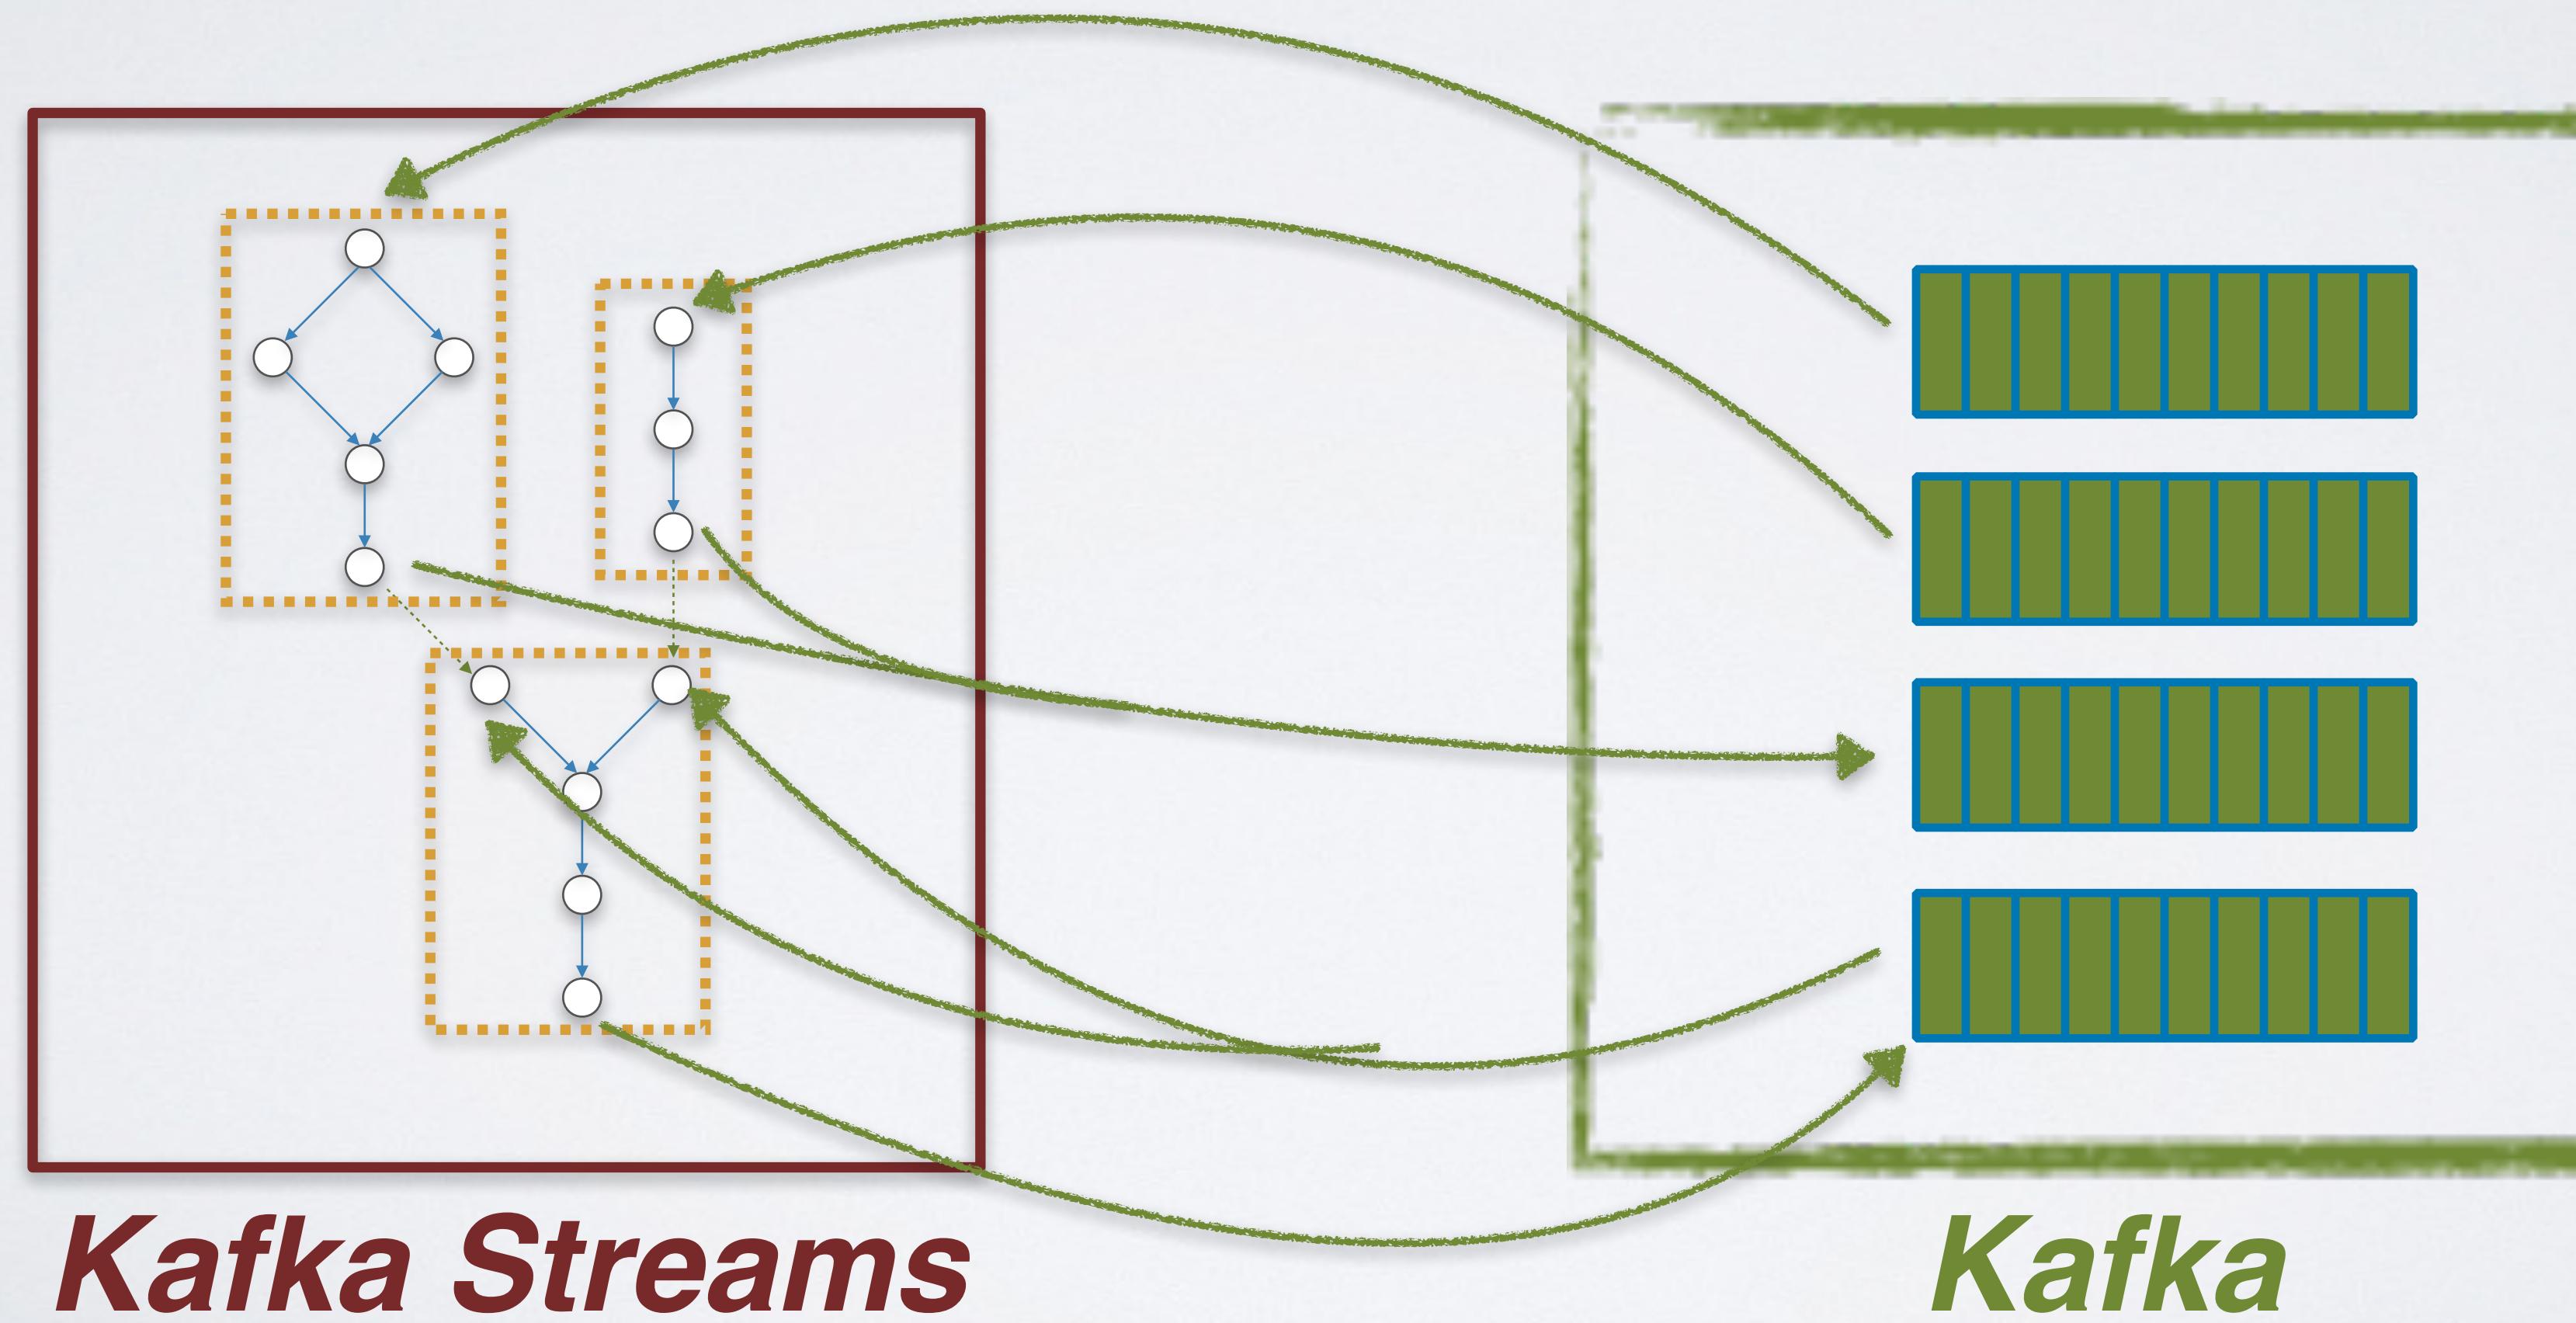
Processor Topology

```
...  
sink1.to("topic1");  
  
source1 = builder.table("topic1");  
  
source2 = sink1.through("topic2");  
...
```

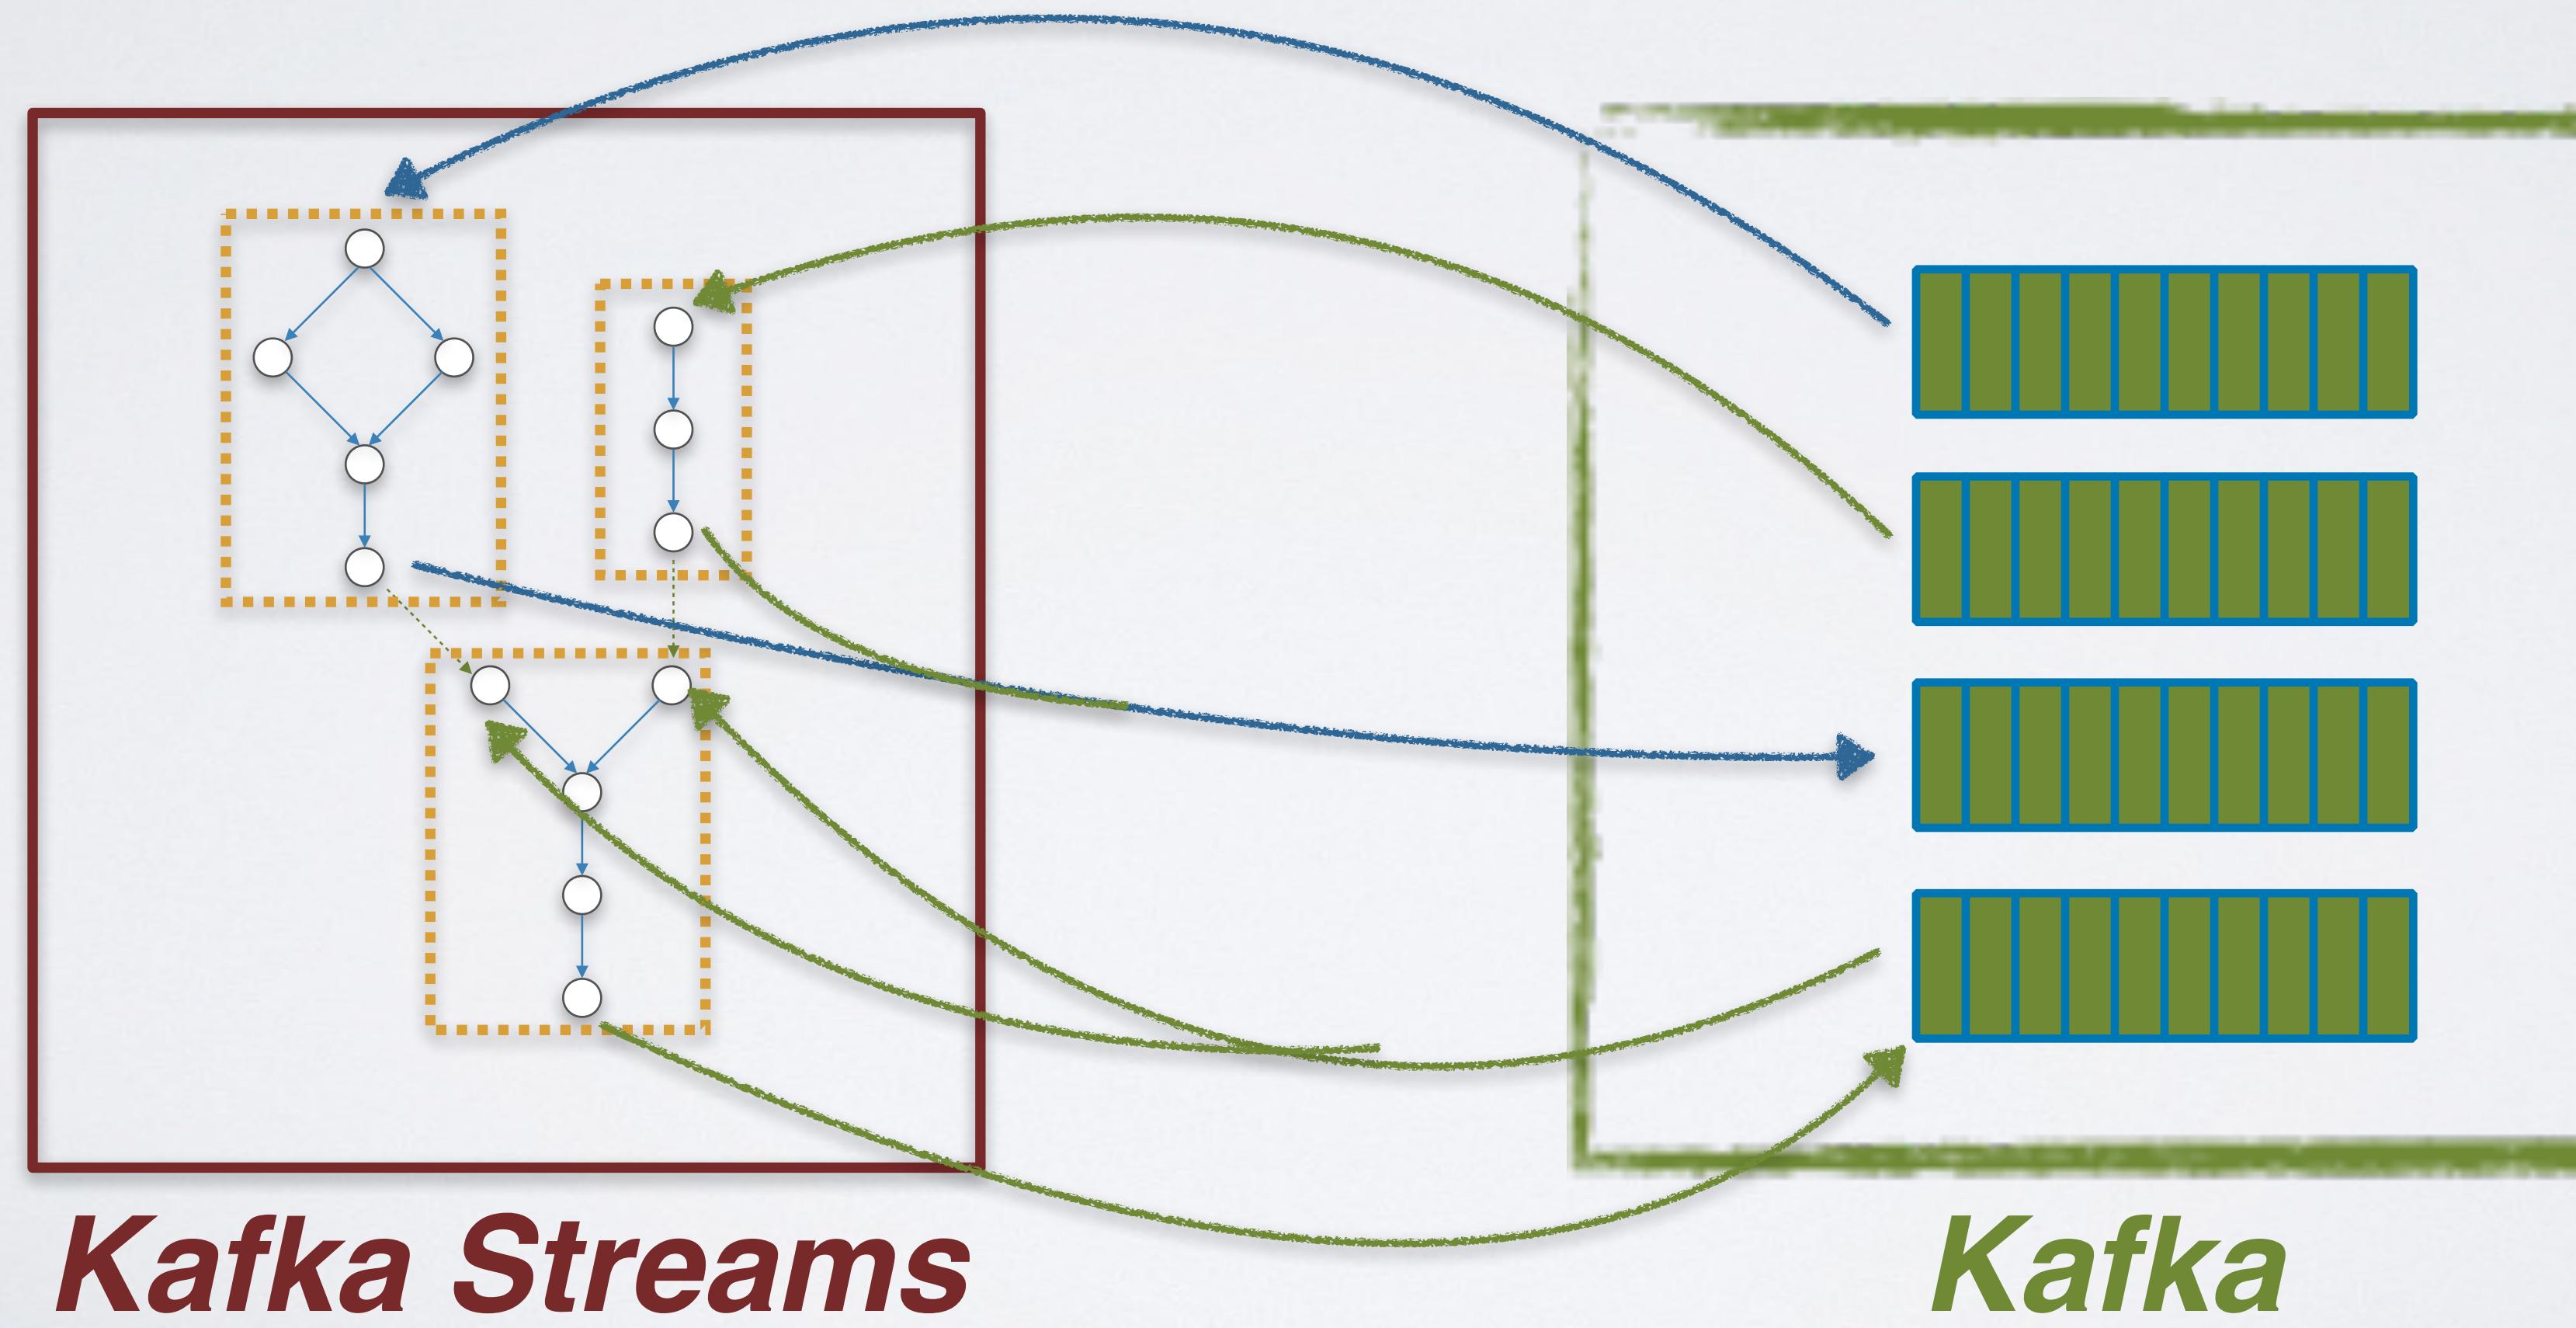
Sub-Topology



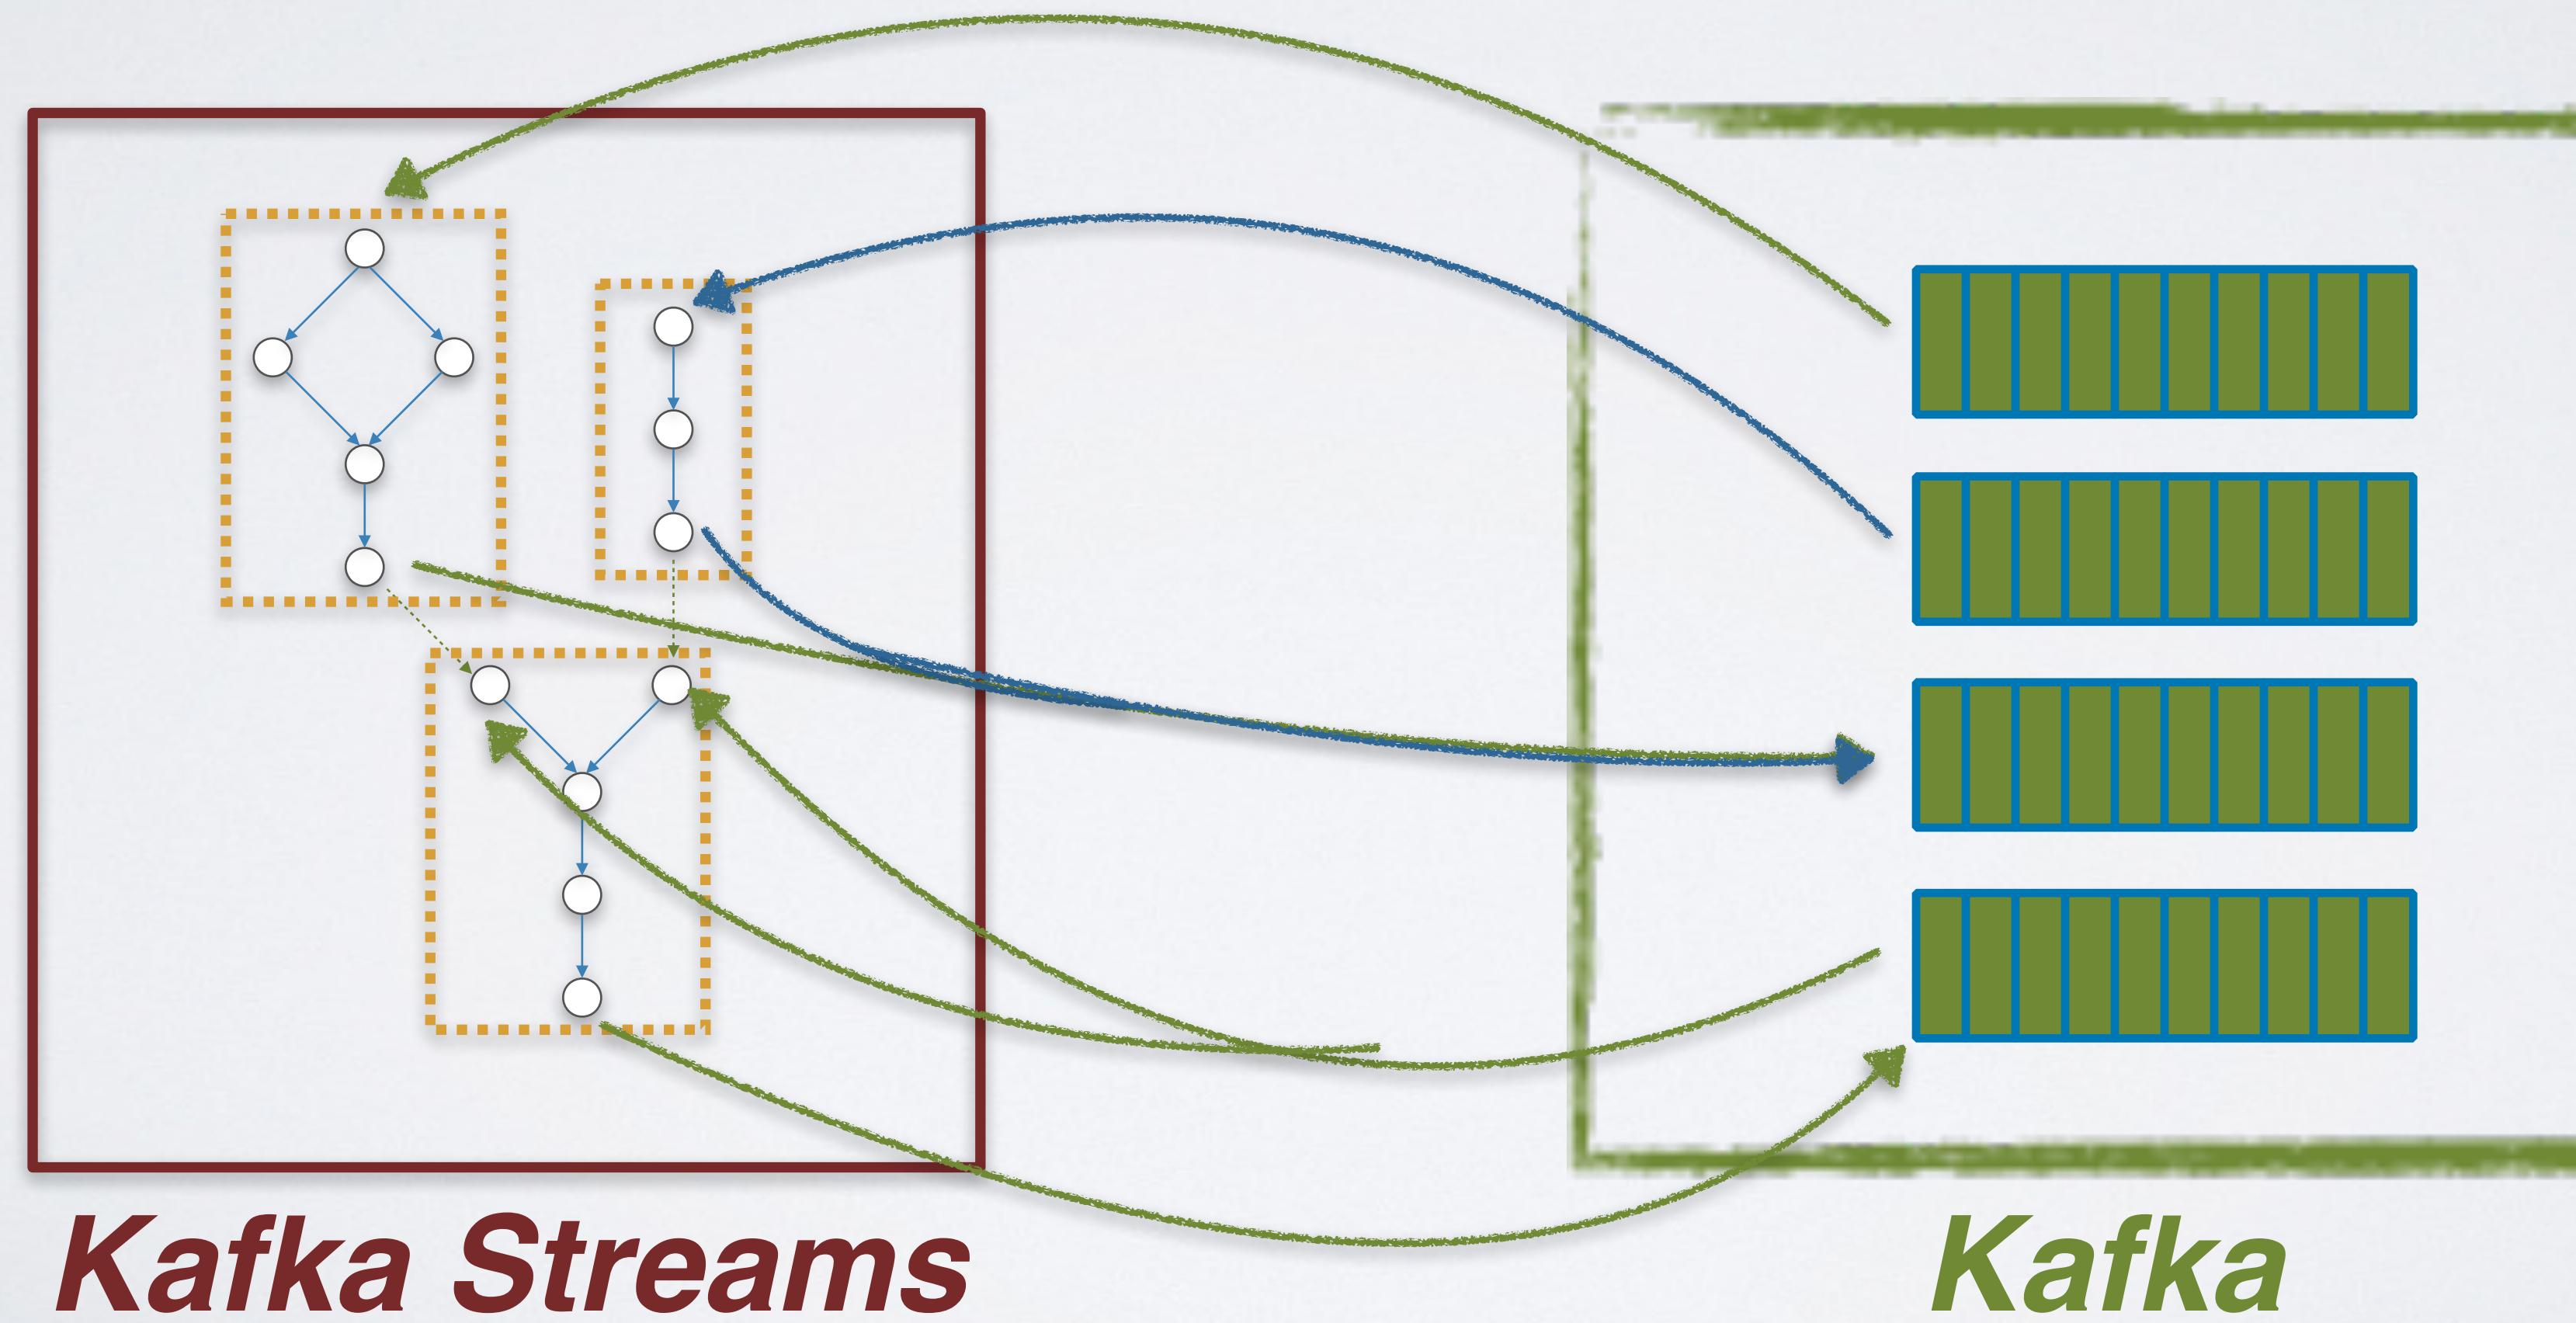
Processor Topology



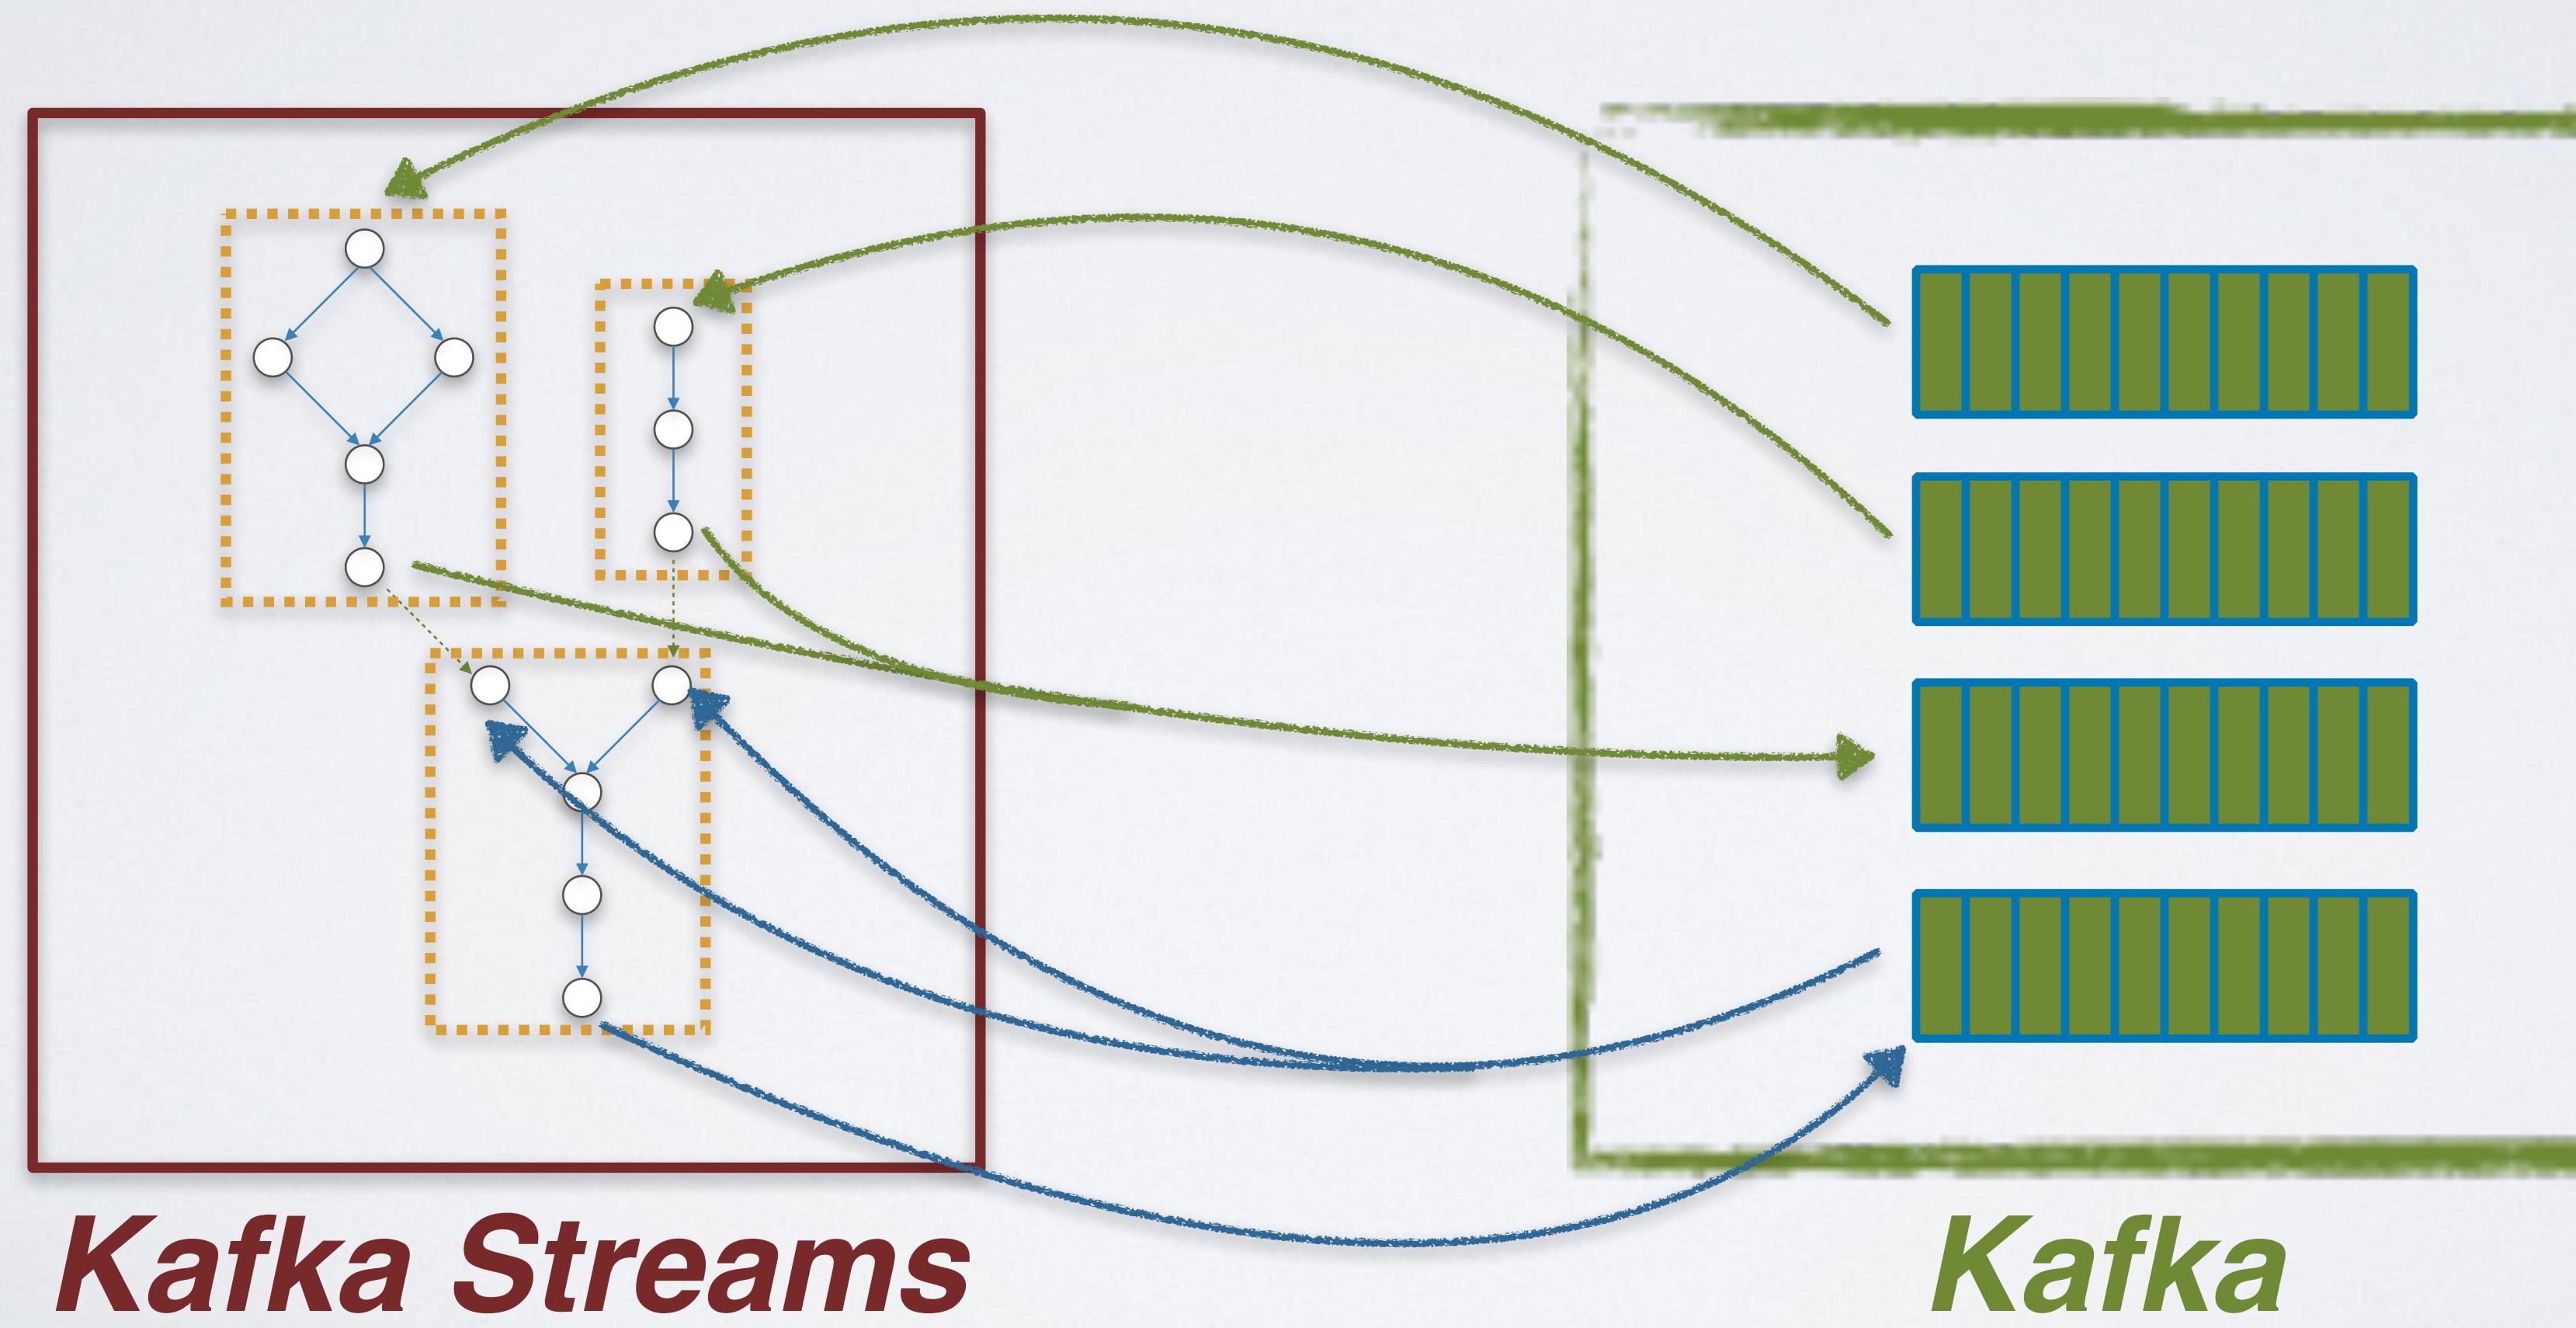
Processor Topology



Processor Topology

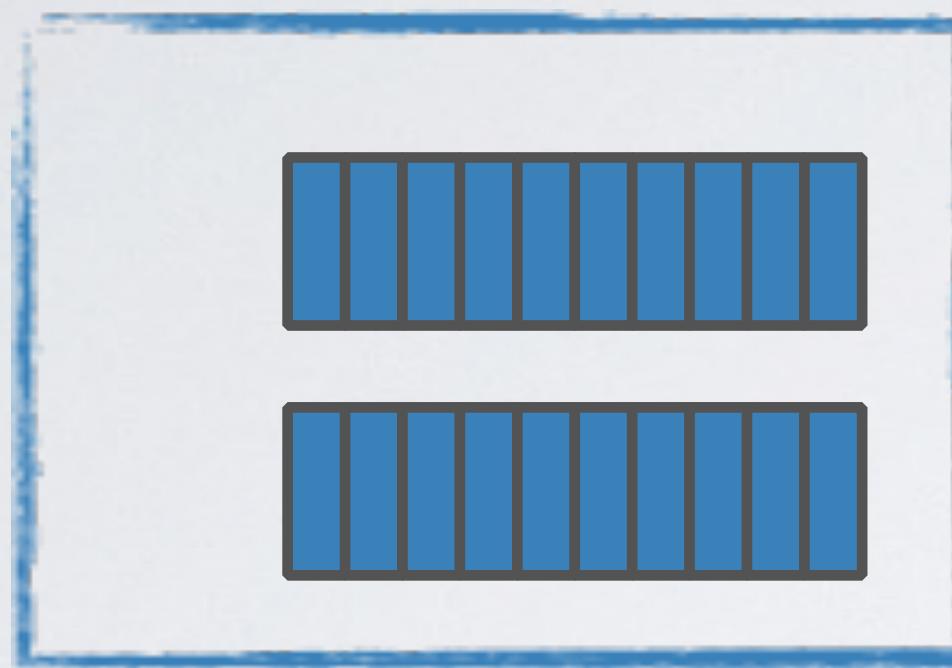


Processor Topology

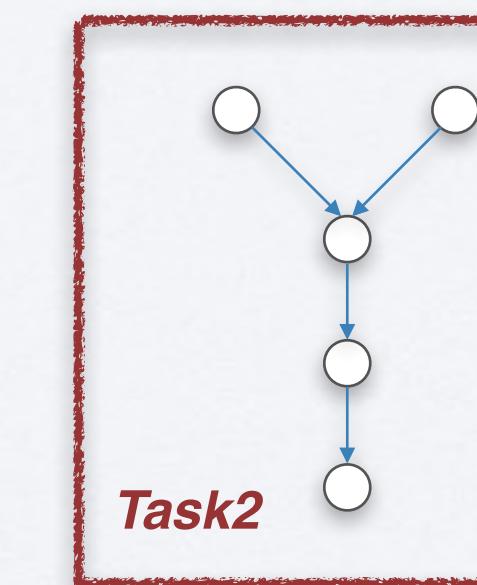
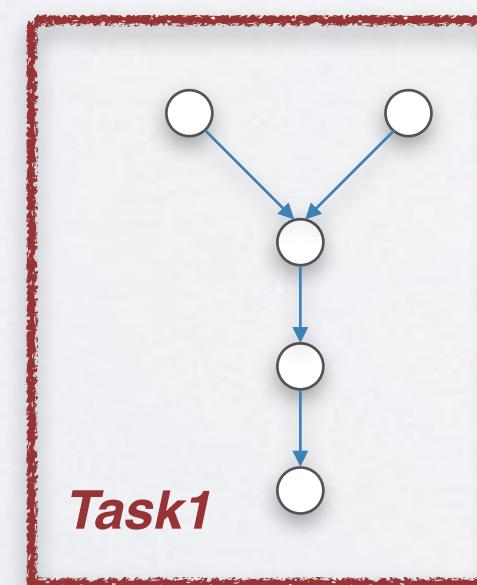
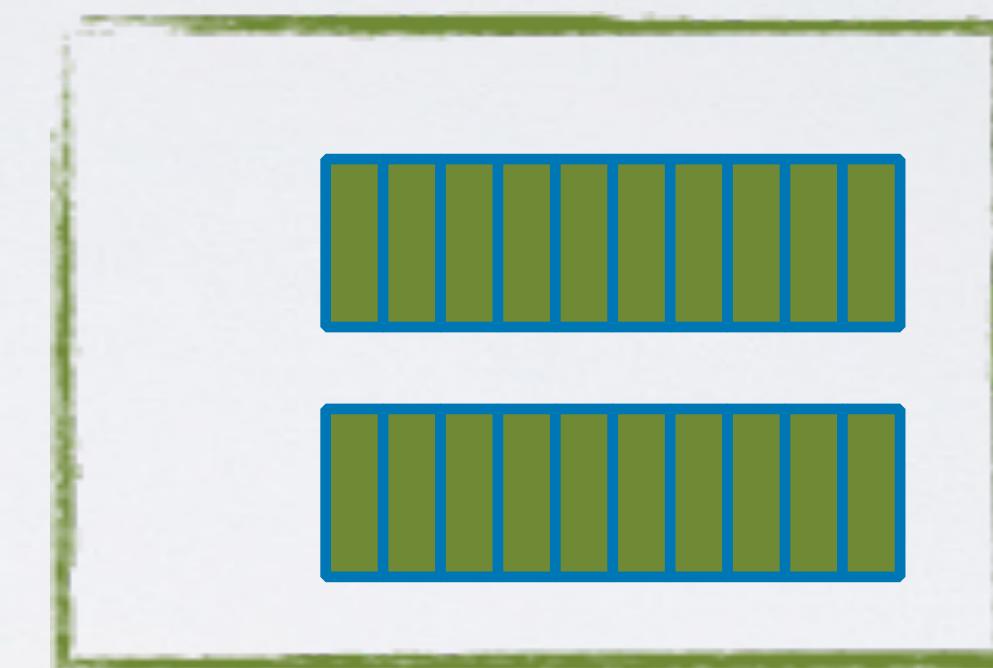


Stream Partitions and Tasks

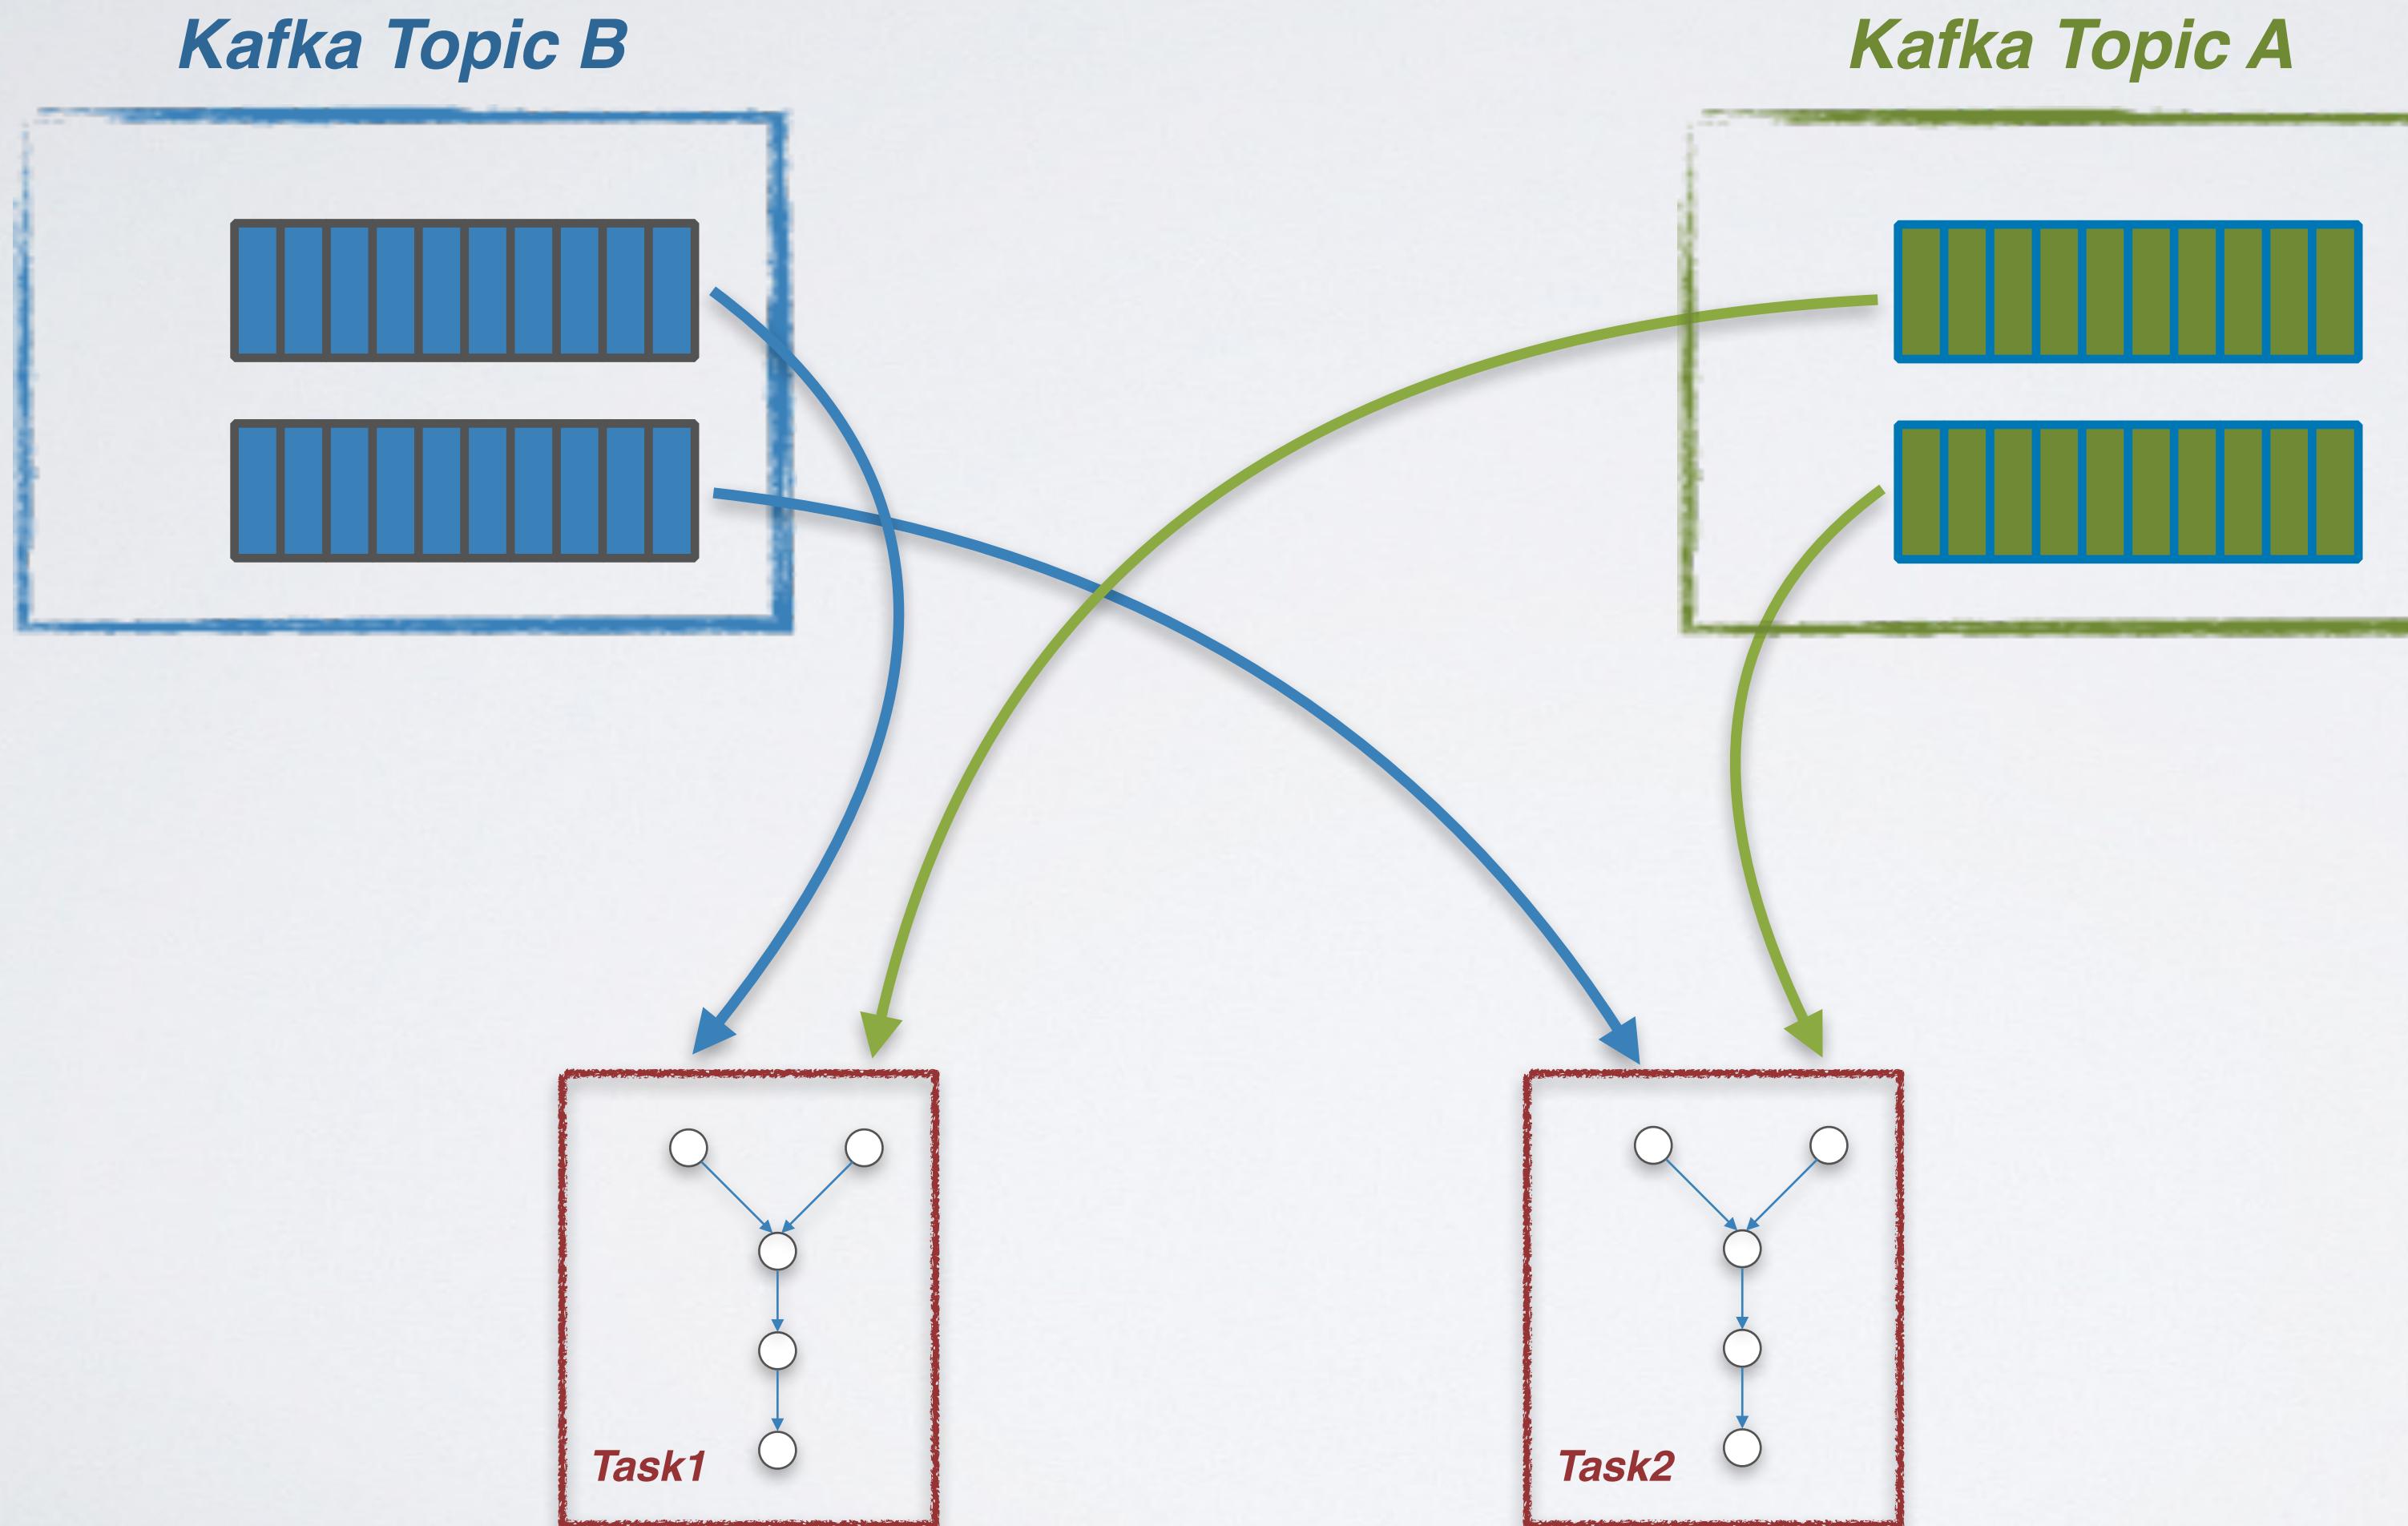
Kafka Topic B



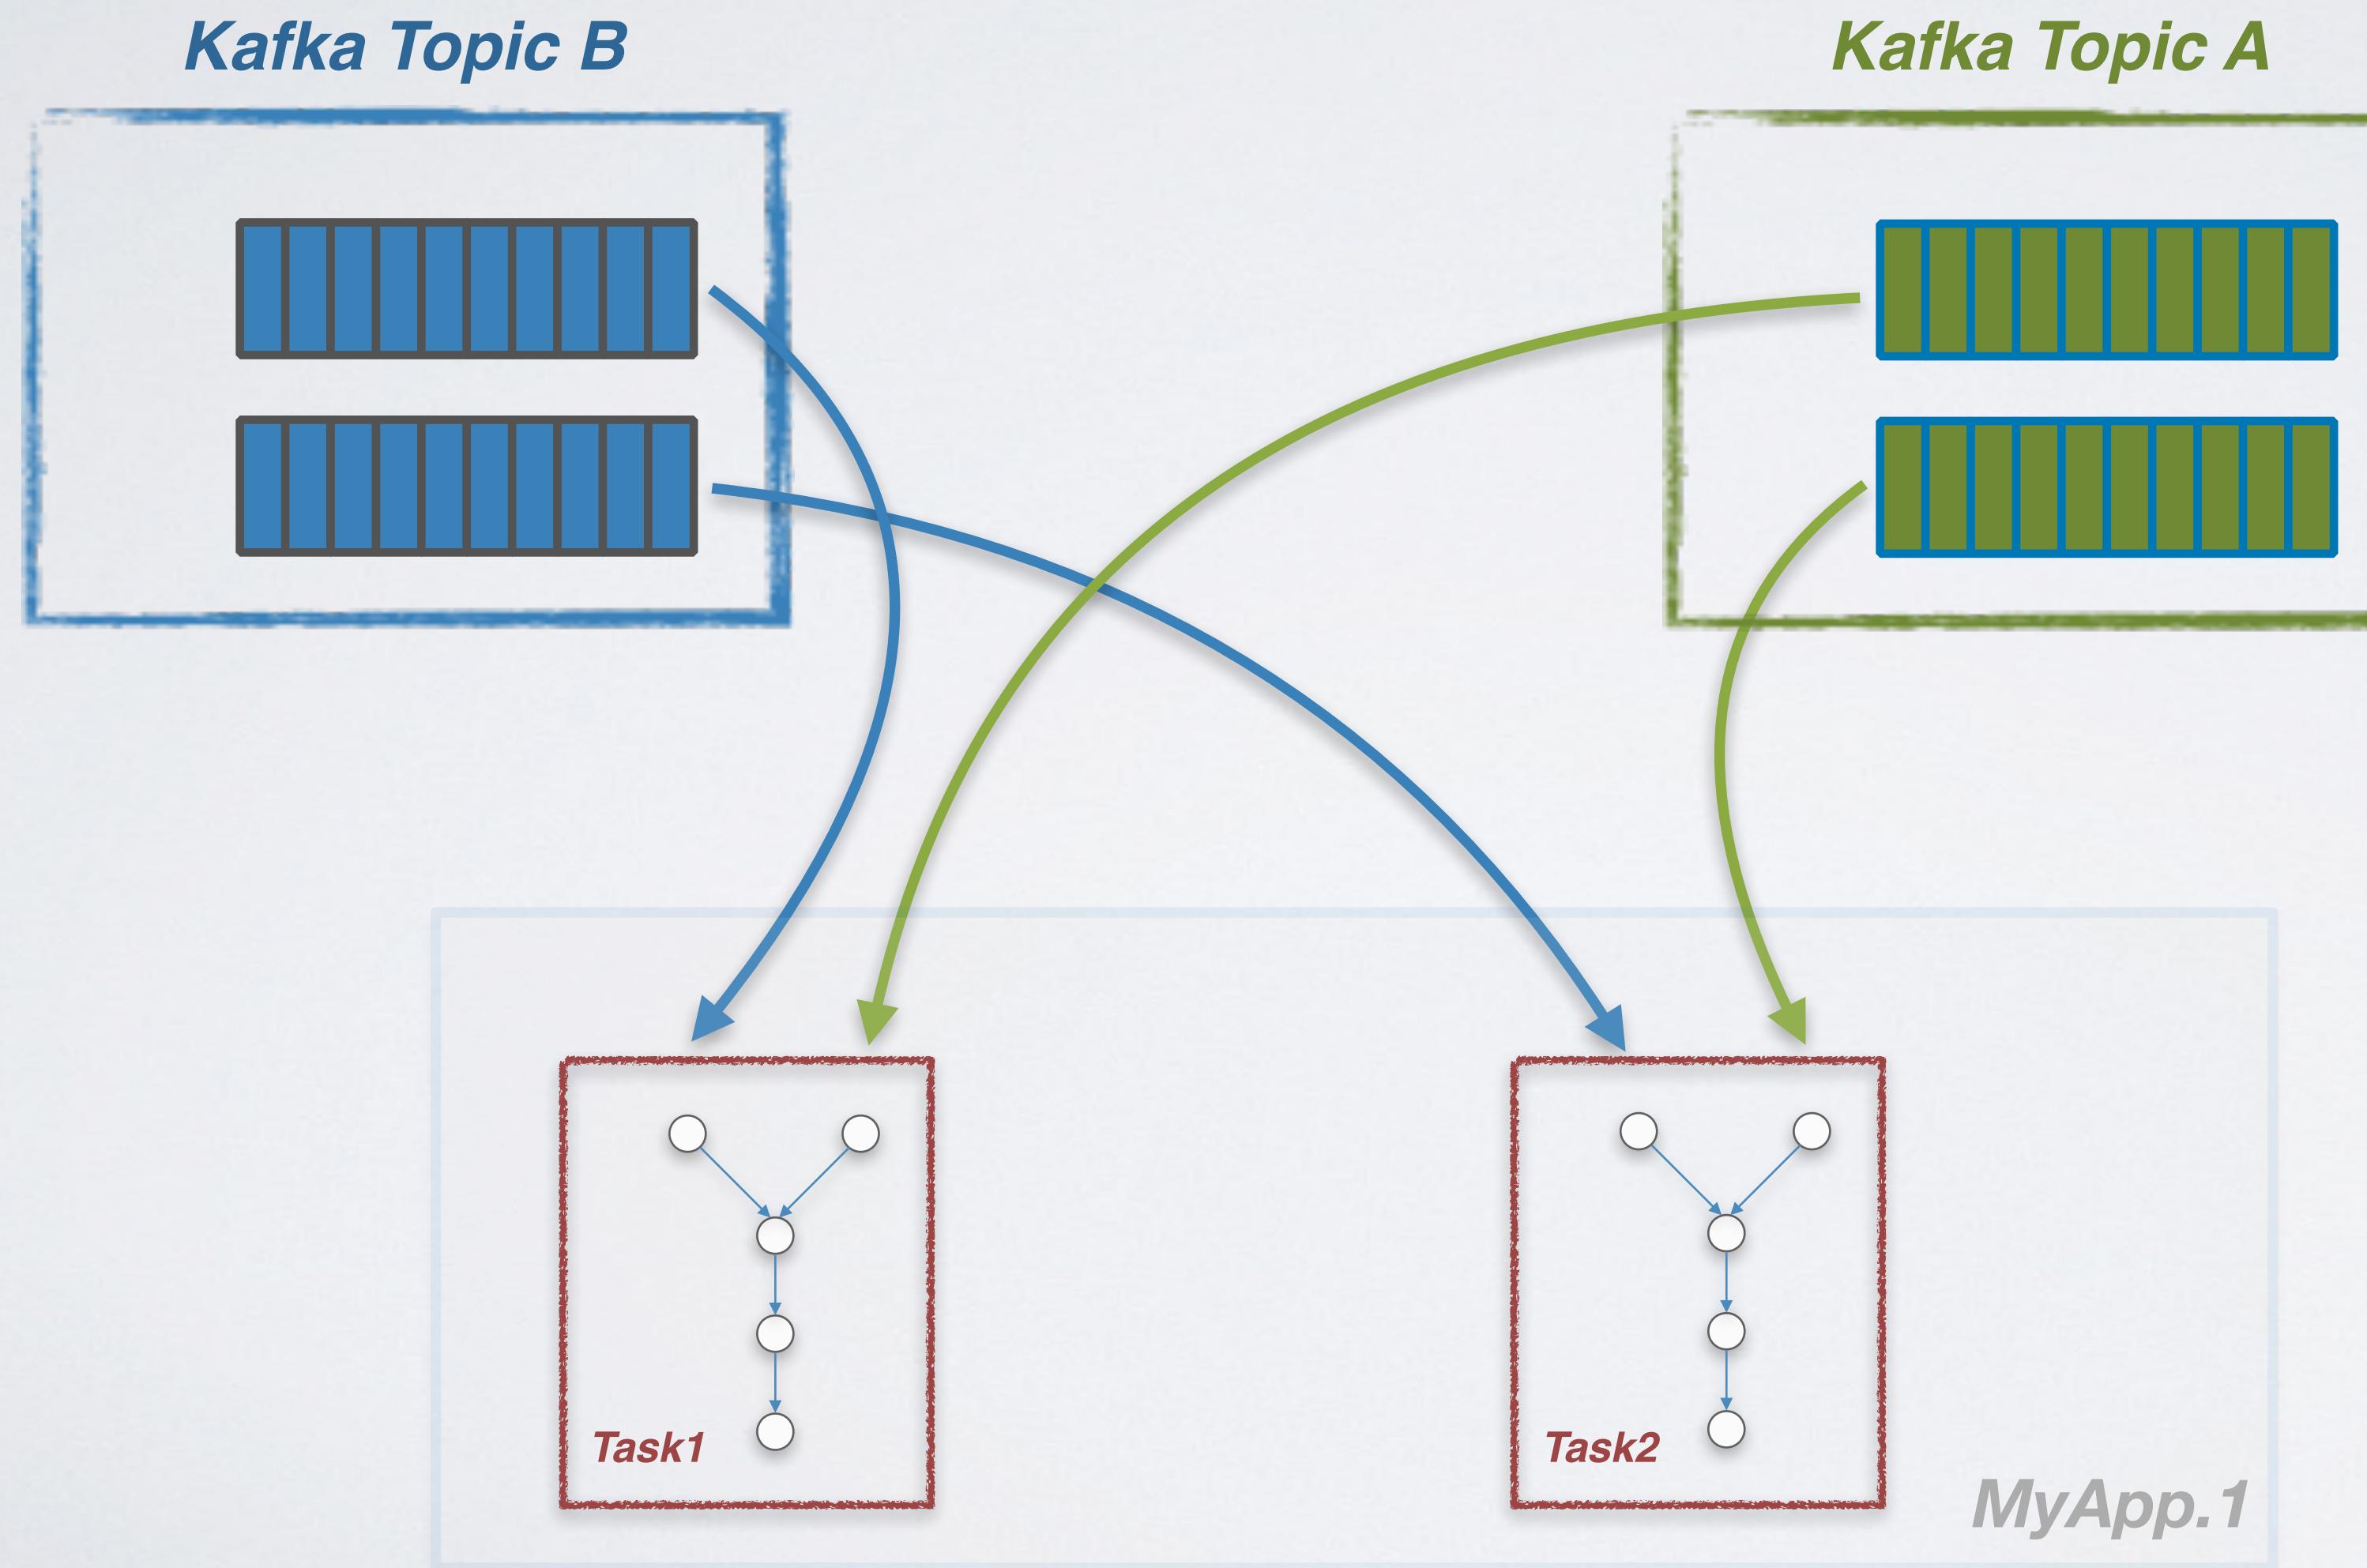
Kafka Topic A



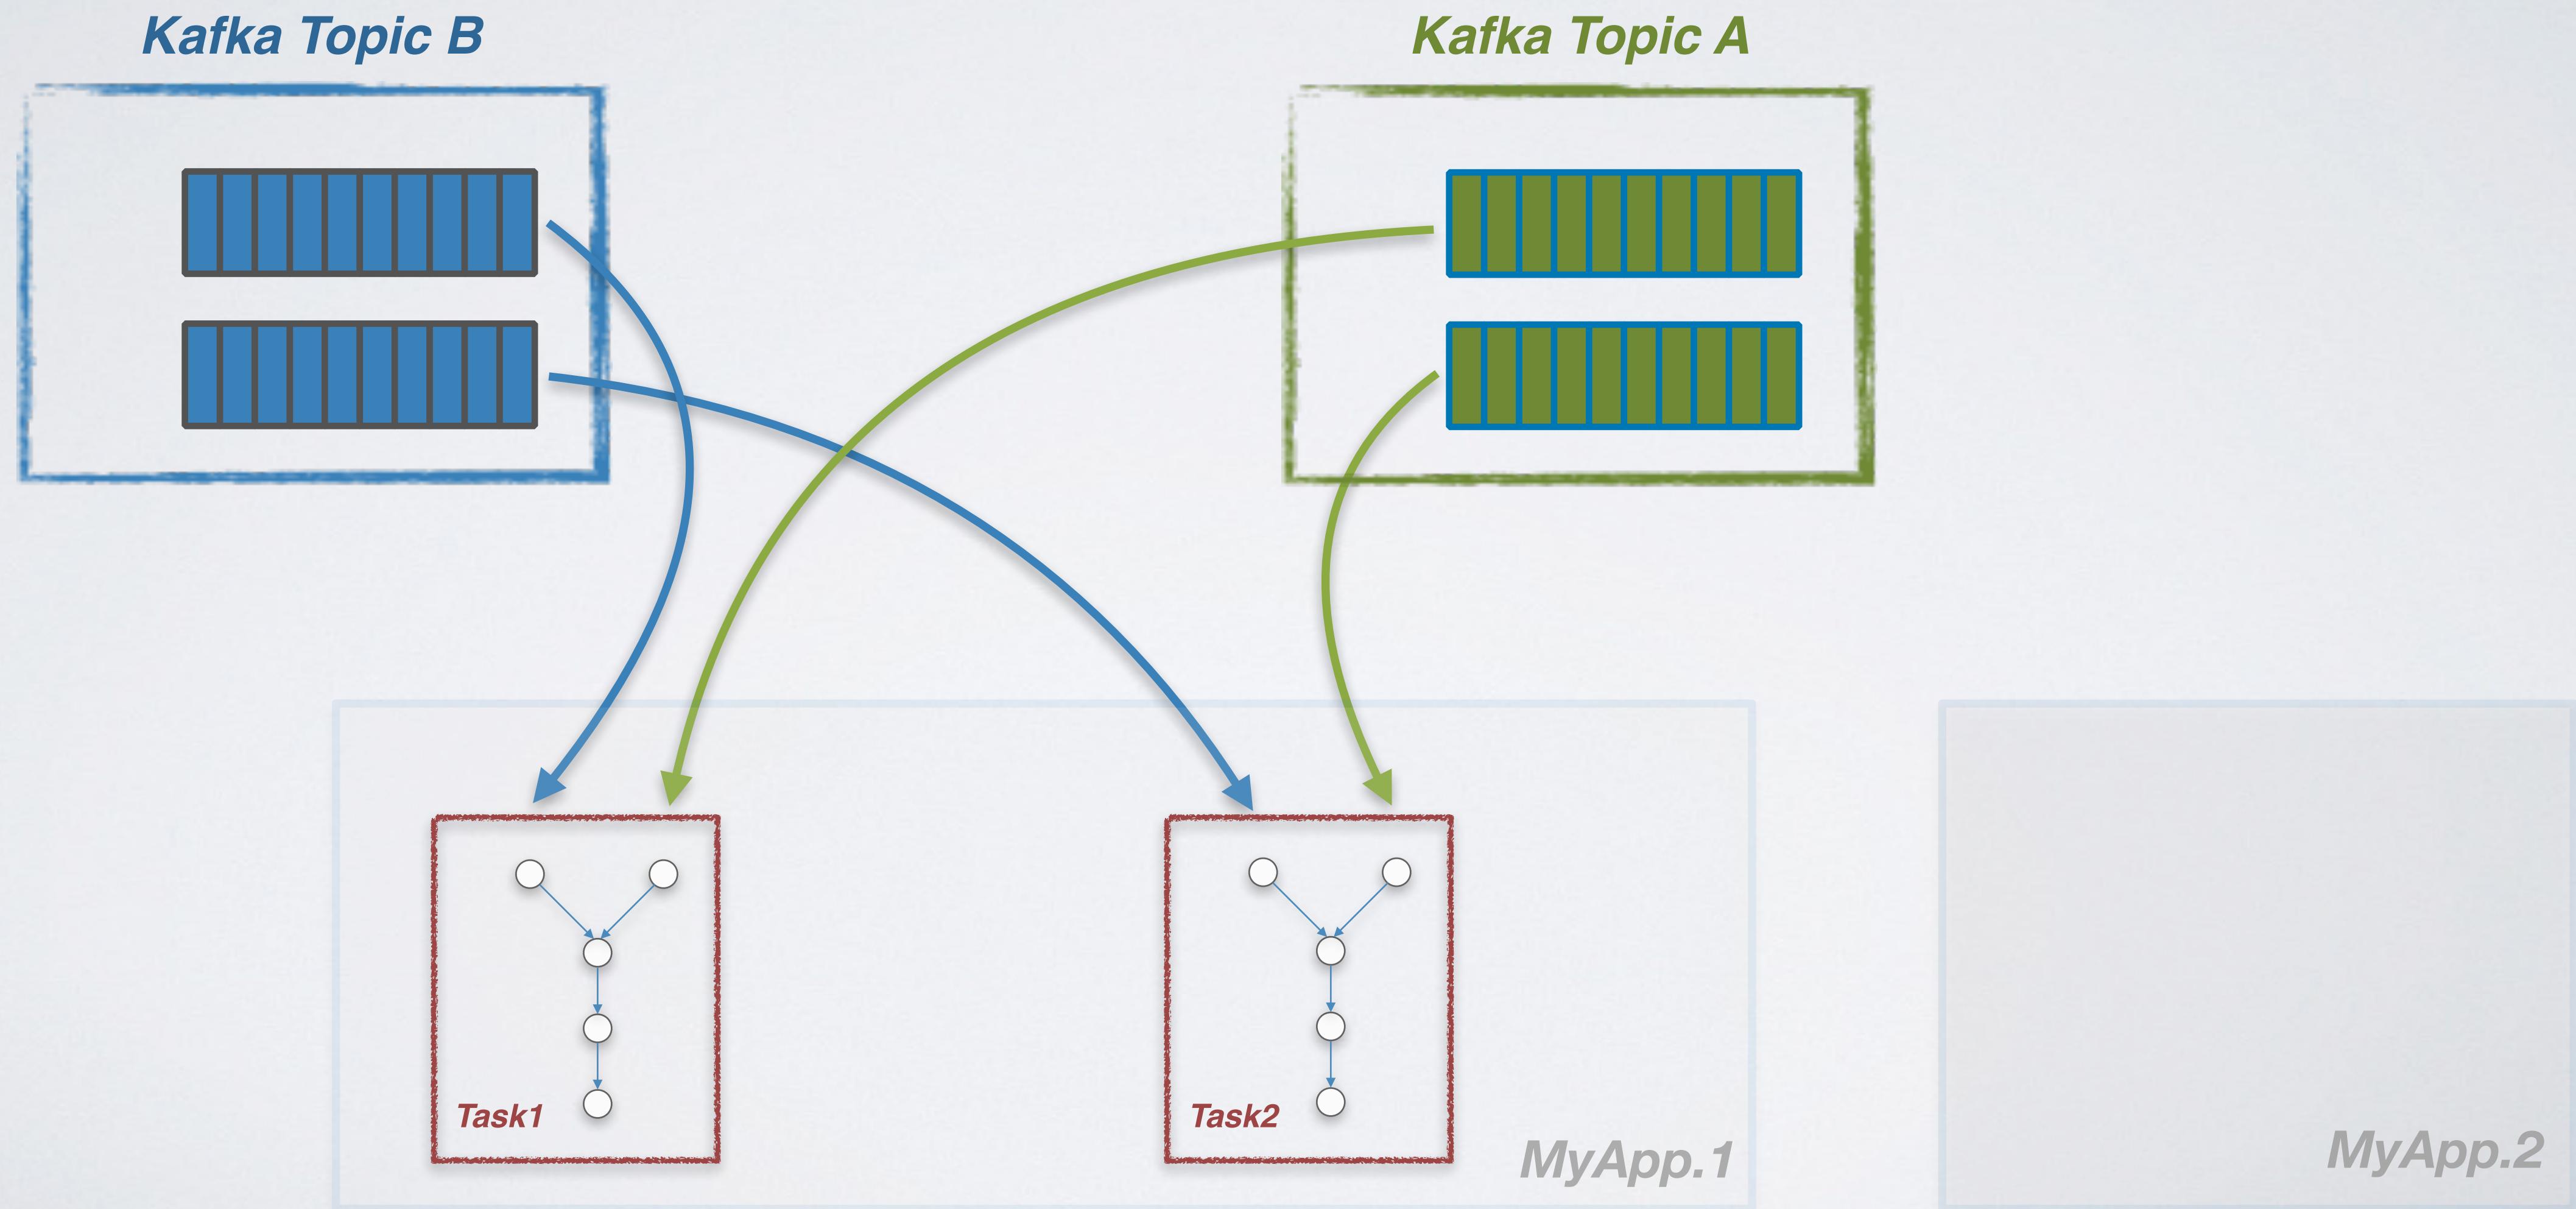
Stream Partitions and Tasks



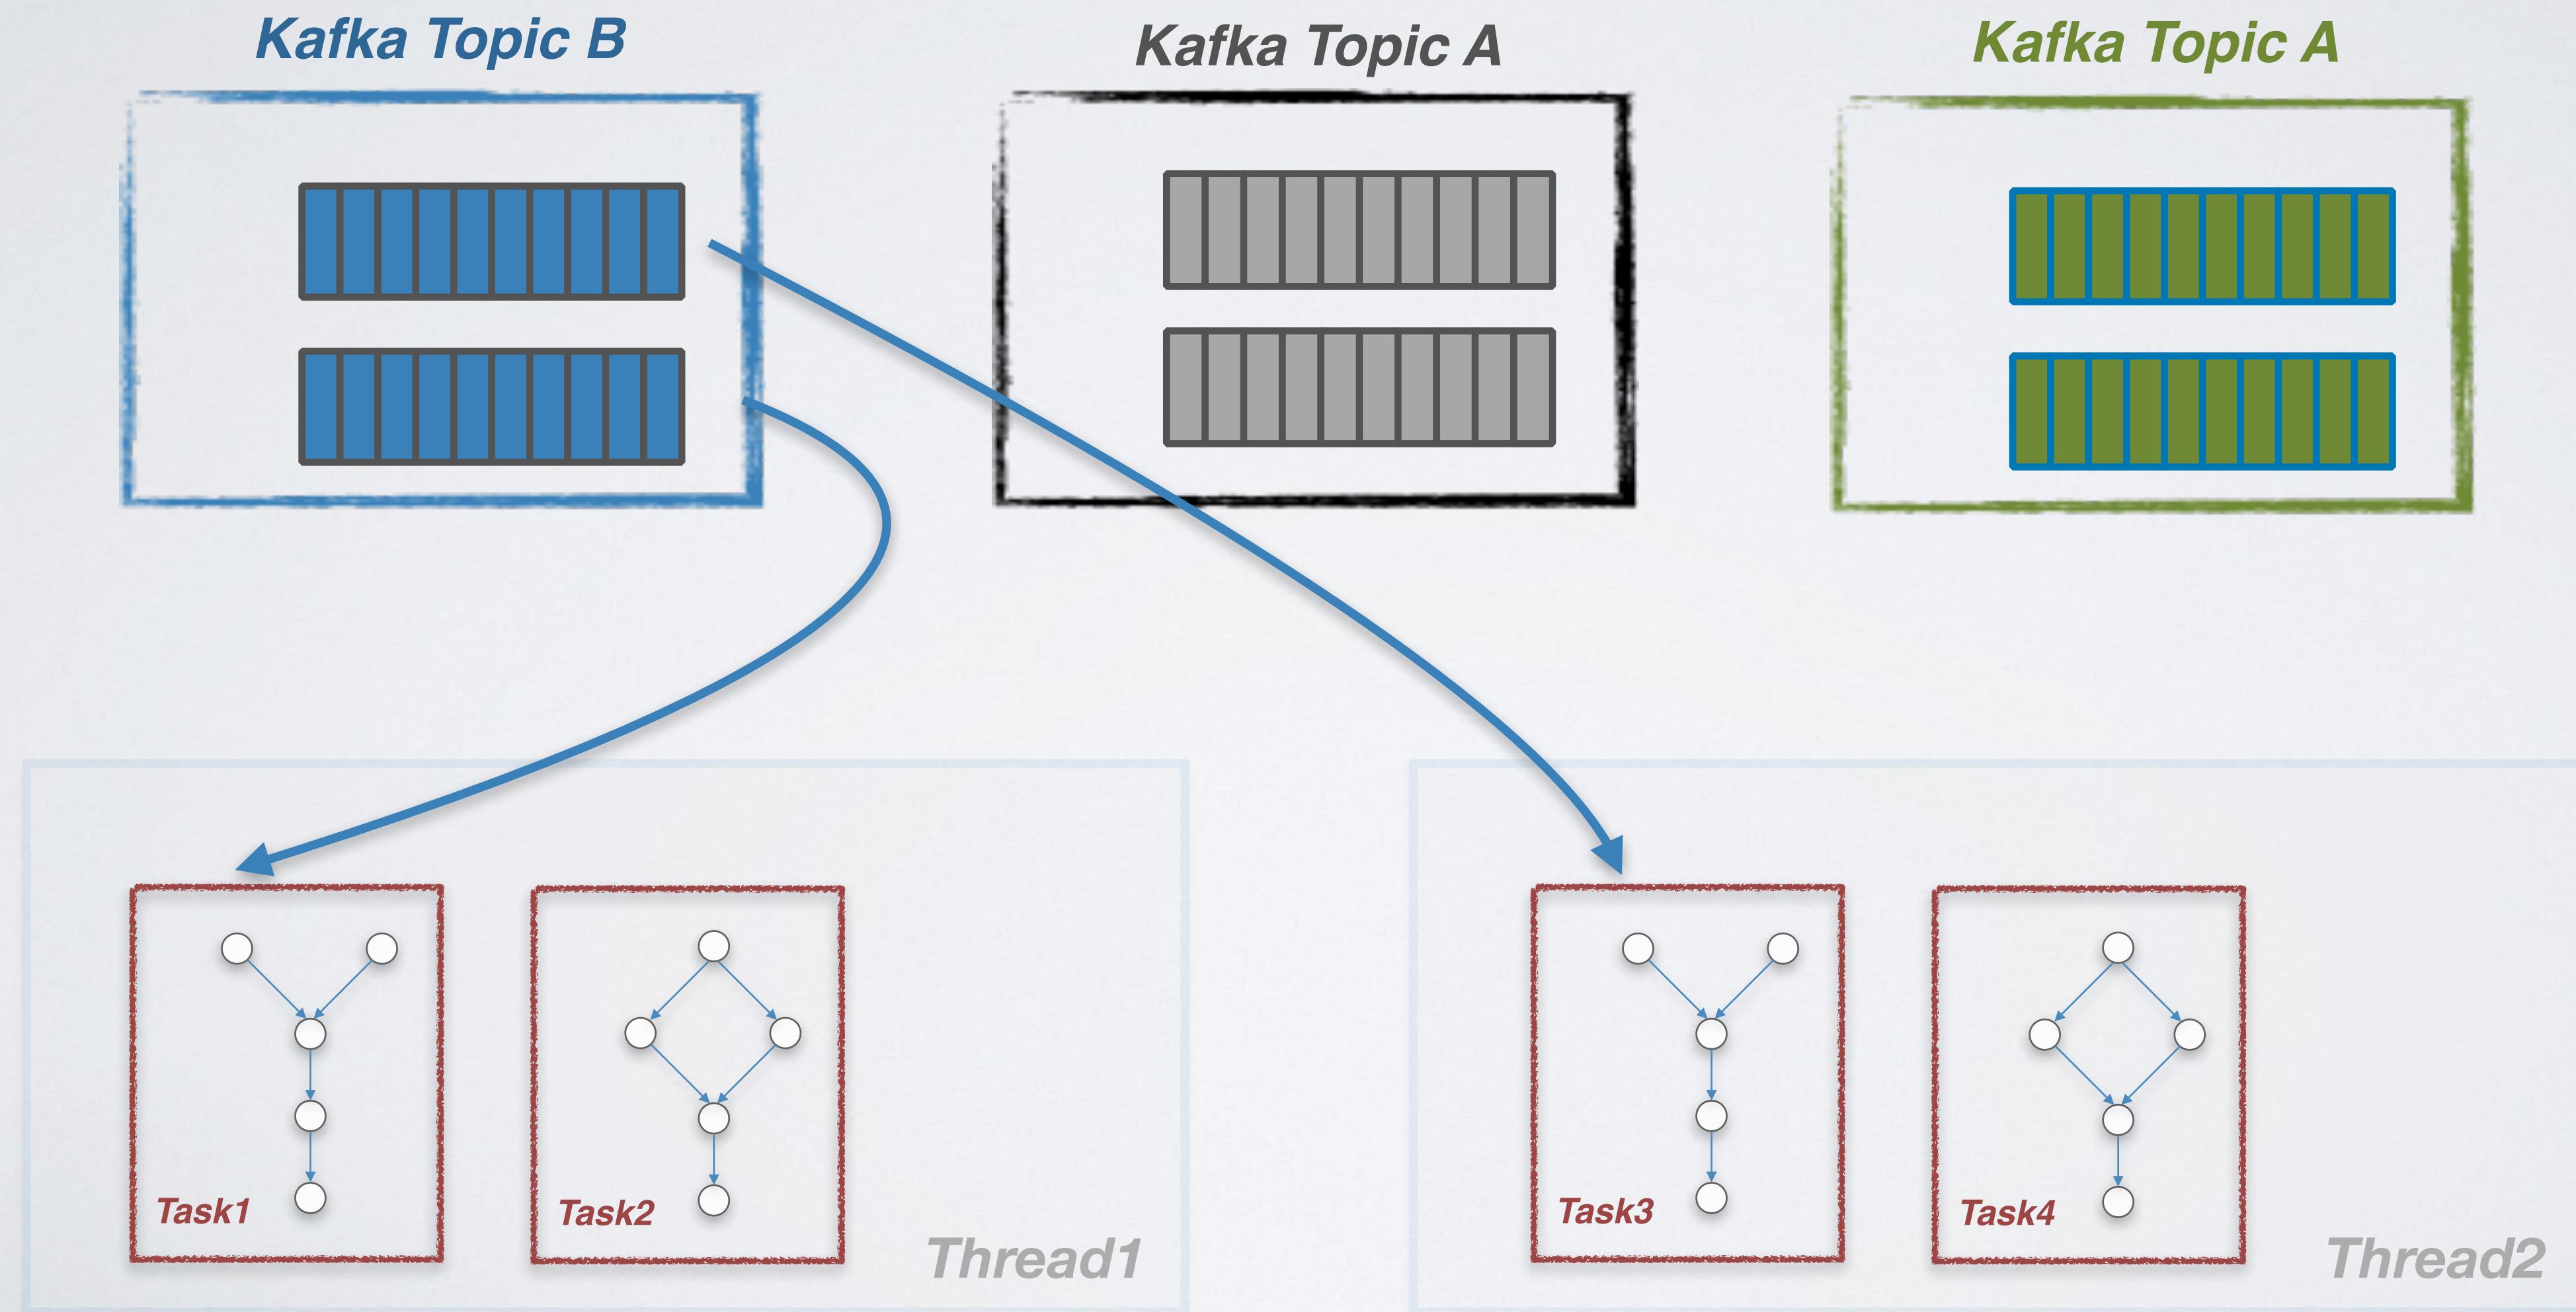
Stream Threads



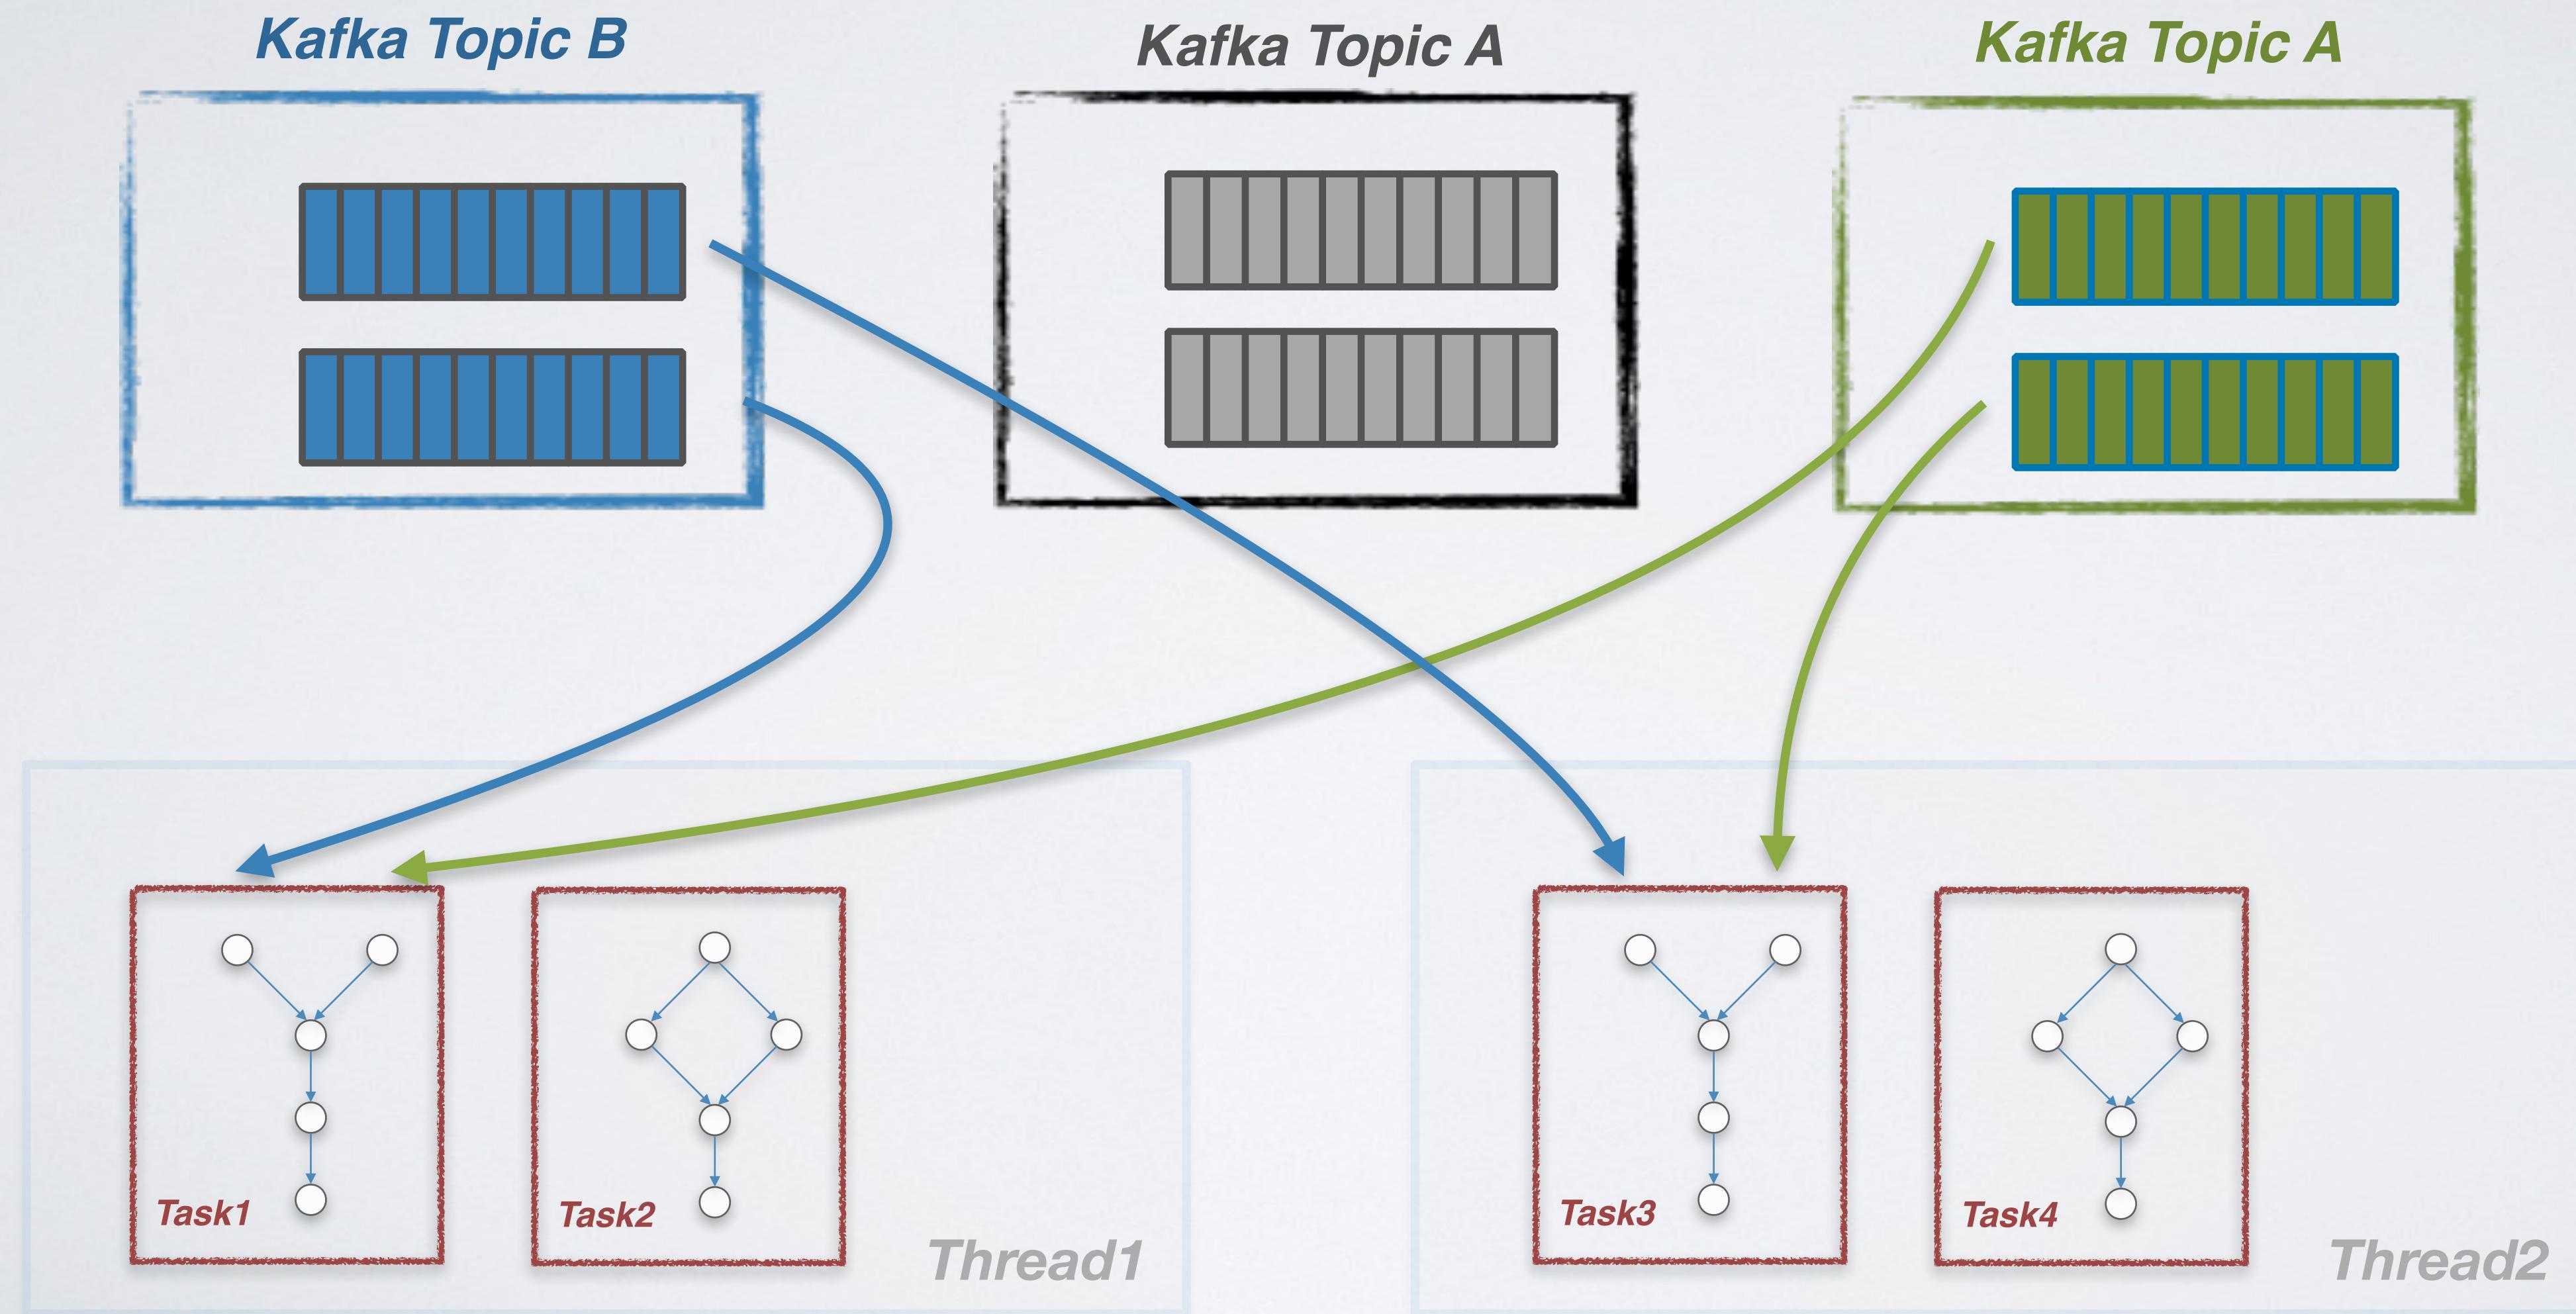
Stream Threads



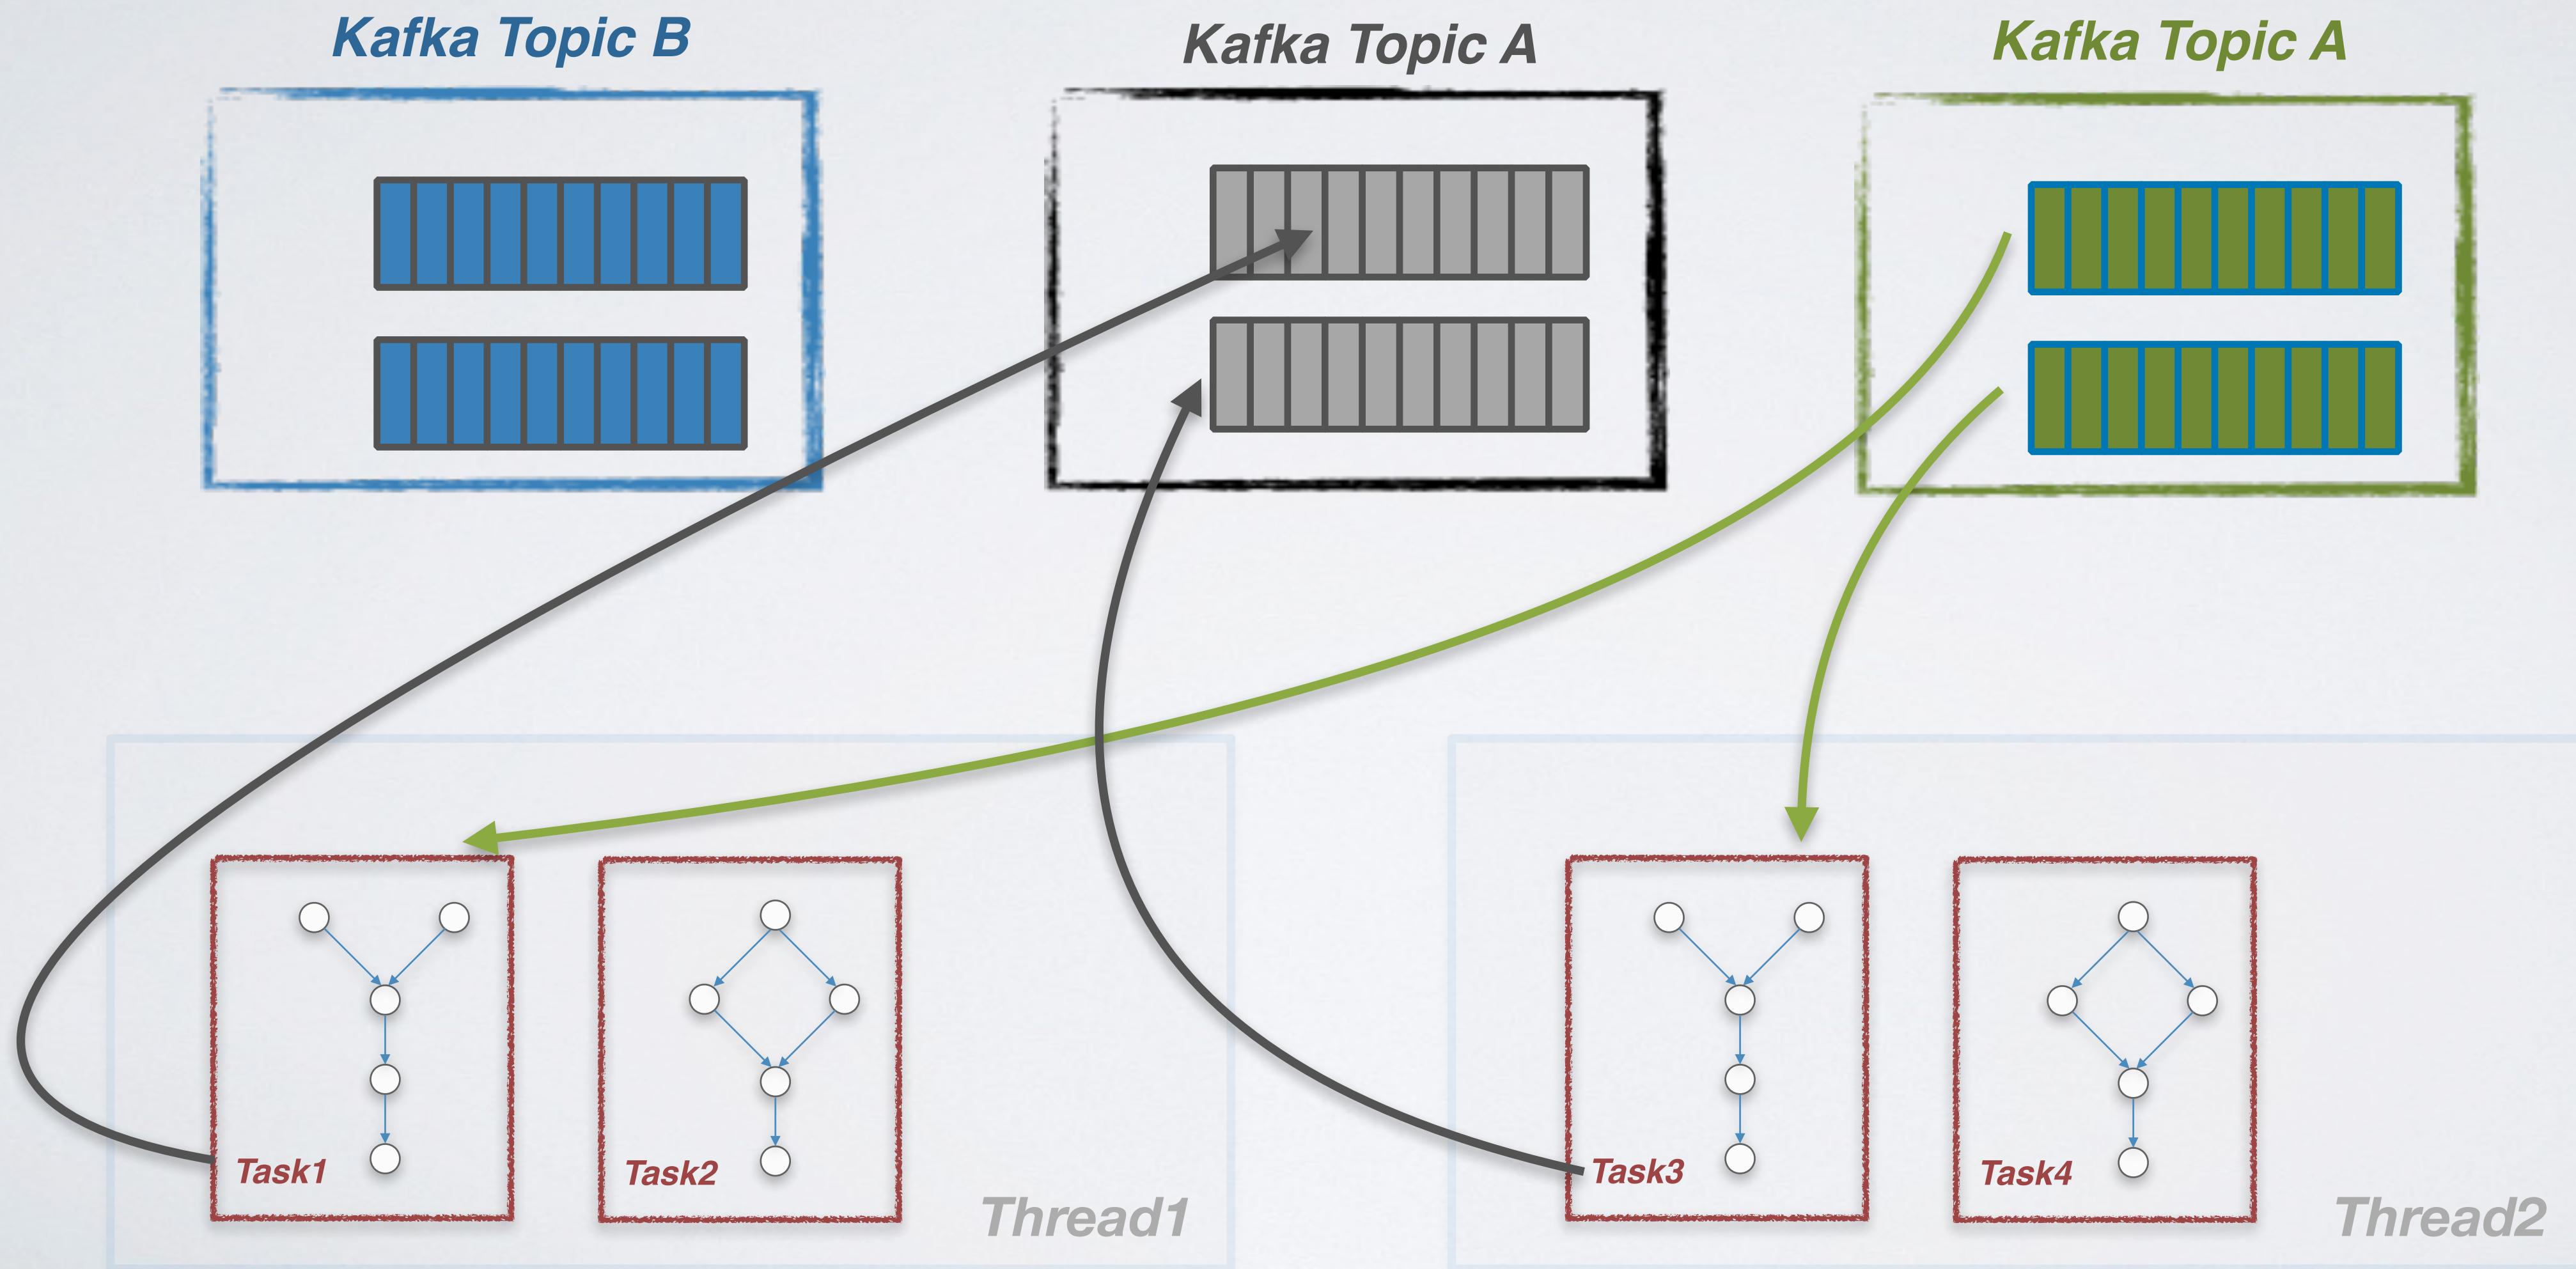
Stream Threads



Stream Threads



Stream Threads



Stream Threads

